

Topics:

- SVM - dual (recap) ✓
- Kernel-Trick → polynomial Kernel
- ① → RBF - kernel vs KNN
- SV-Regression
- SVMs in practice
- Naïve-Bayes (this & next class)

SVM-Dual

$$\begin{array}{l} \text{Max} \\ \boxed{\alpha_i} \end{array}$$

$$\sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y_i y_j \underbrace{x_i^T x_j}_{\equiv K(x_i, x_j)} \rightarrow \text{Similarity}$$

s.t. $0 \leq \alpha_i \leq C$

$$\sum_{i=1}^n \alpha_i y_i = 0$$

Primal-Dual

(proof)



$\alpha_i = 0$ for non SVs

- ① $x_i \rightarrow \alpha_i$
- ② $x_i^T x_j$ form $\underbrace{K(x_i, x_j)}$
- ③ $f(\underline{\alpha}) = \underbrace{y_0}_{\uparrow} \leftarrow \sum_{i=1}^n \alpha_i y_i x_i^T x_0 + b$
- ④ $\alpha_i > 0$ for SVs only

~~Recap~~

$$K(x_1, x_2) = \underbrace{(x_1^\top x_2)^2}_{\in \mathbb{R}^2} : \text{Quadratic Kernel}$$

\parallel

$$x_1' = [1, x_{11}^2, x_{12}^2, \sqrt{2}x_{11}, \sqrt{2}x_{12}, \sqrt{2}x_{11}x_{12}]^T$$

$$x_2' = [1, x_{21}^2, x_{22}^2, \sqrt{2}x_{21}, \sqrt{2}x_{22}, \sqrt{2}x_{21}x_{22}]^T$$

6-dim

\parallel

$$x_1' = [1, x_{11}^2, x_{12}^2, \sqrt{2}x_{11}, \sqrt{2}x_{12}, \sqrt{2}x_{11}x_{12}]^T$$

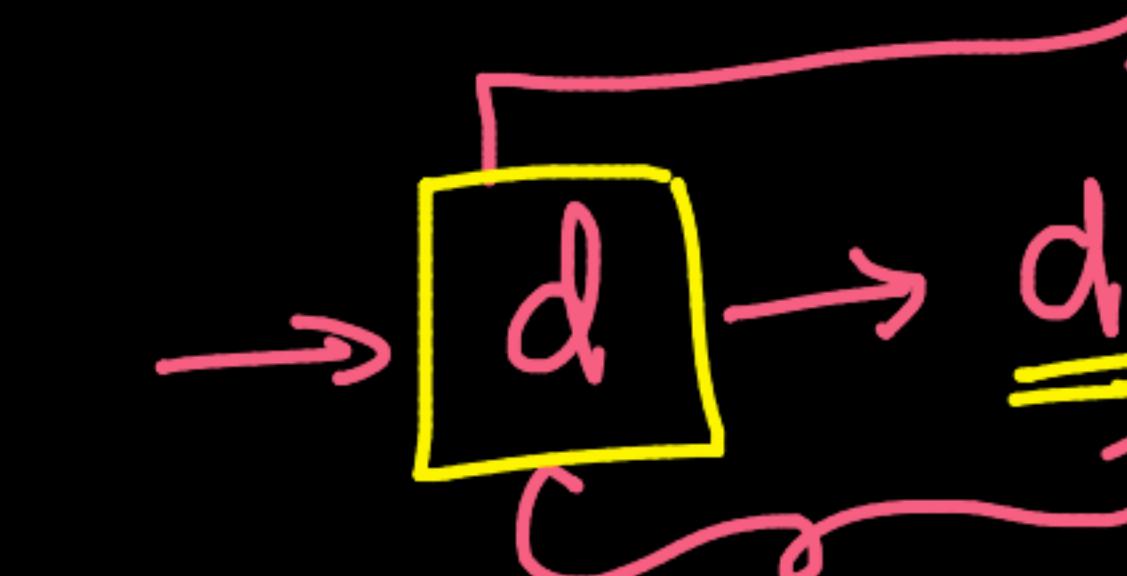
$$x_2' = [1, x_{21}^2, x_{22}^2, \sqrt{2}x_{21}, \sqrt{2}x_{22}, \sqrt{2}x_{21}x_{22}]^T$$

6-dim

Q

Task: Visualization

PCA



preserving variance

$2\pi^3$

$\sqrt{d'}$

$d' < d$

↳ visualize

Task: Classification

SVM

$d \xrightarrow{\text{Kernel}} d'$

$d' > d$

it is easier to find ω

$\overline{\omega}^{d'}$
sep

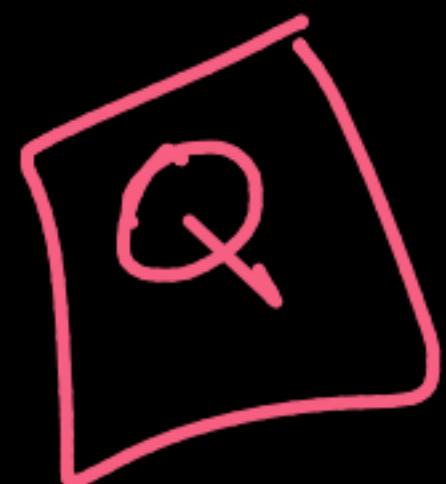
Lagrange-Mul \rightarrow concept in Optim²ⁿ

PCA
 \Leftrightarrow

$\min_x f(x) + \lambda_1 \boxed{\quad} + \lambda_2 \boxed{\quad}$

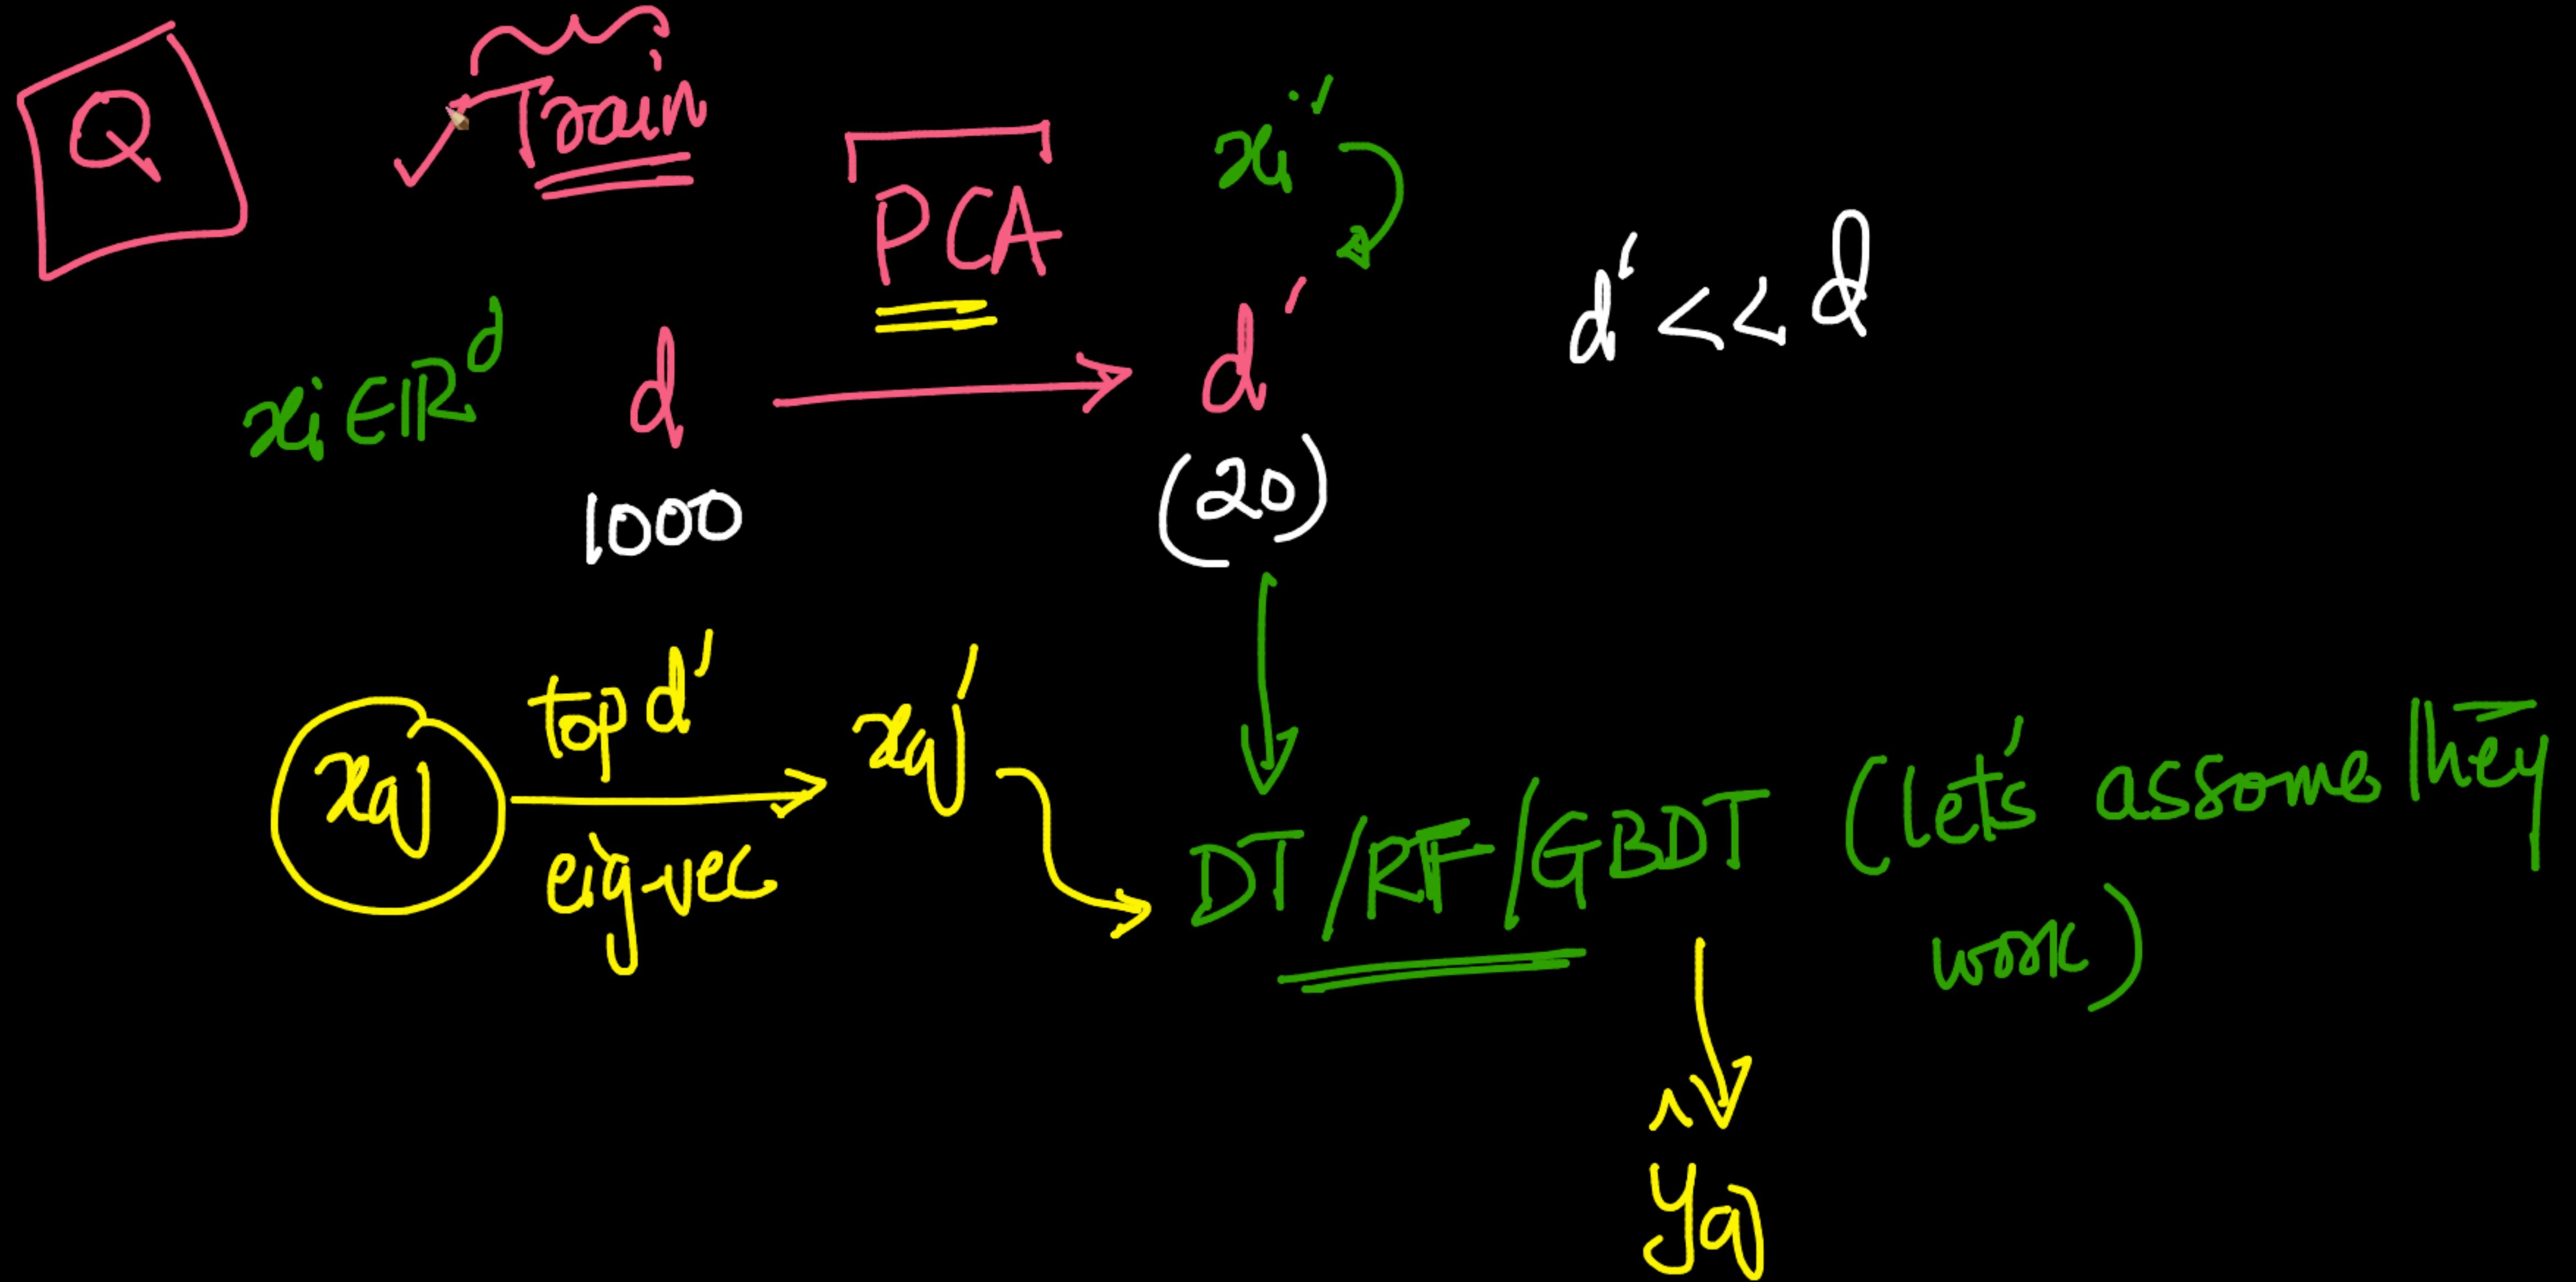
s.t. $g(x) = 0$

$h(x) \geq 0$



$x_i \in \mathbb{R}^d$ $d = 1000$ $\xrightarrow{\text{PCA}} \quad x_i' \quad d' \ll d$
 (20)





$\{D_{Train} \rightarrow \text{PCA } (d \rightarrow d')$
top d' eig-vec

$\{ D_{Test} \rightarrow \text{top } d' - \text{eigvec} \rightarrow \underline{\underline{x_i}}$

✓ Radial Basis fn (RBF / Gaussian)

- most popular
- general default if you don't have a domain specific kernel

$$K_{RBF}(x_1, x_2) = \exp\left\{-\frac{\text{euc.dist}(x_1, x_2)}{2\sigma^2}\right\}$$

euc.dist(x_1, x_2)
 $\overbrace{\|x_1 - x_2\|}^2$
 σ^2
hyper-param

$$K_{RBF}(x_1, x_2) = \exp\left(\frac{-d_{12}^2}{2\sigma^2}\right)$$



$$d_{12}^2 = \|x_1 - x_2\|^2$$

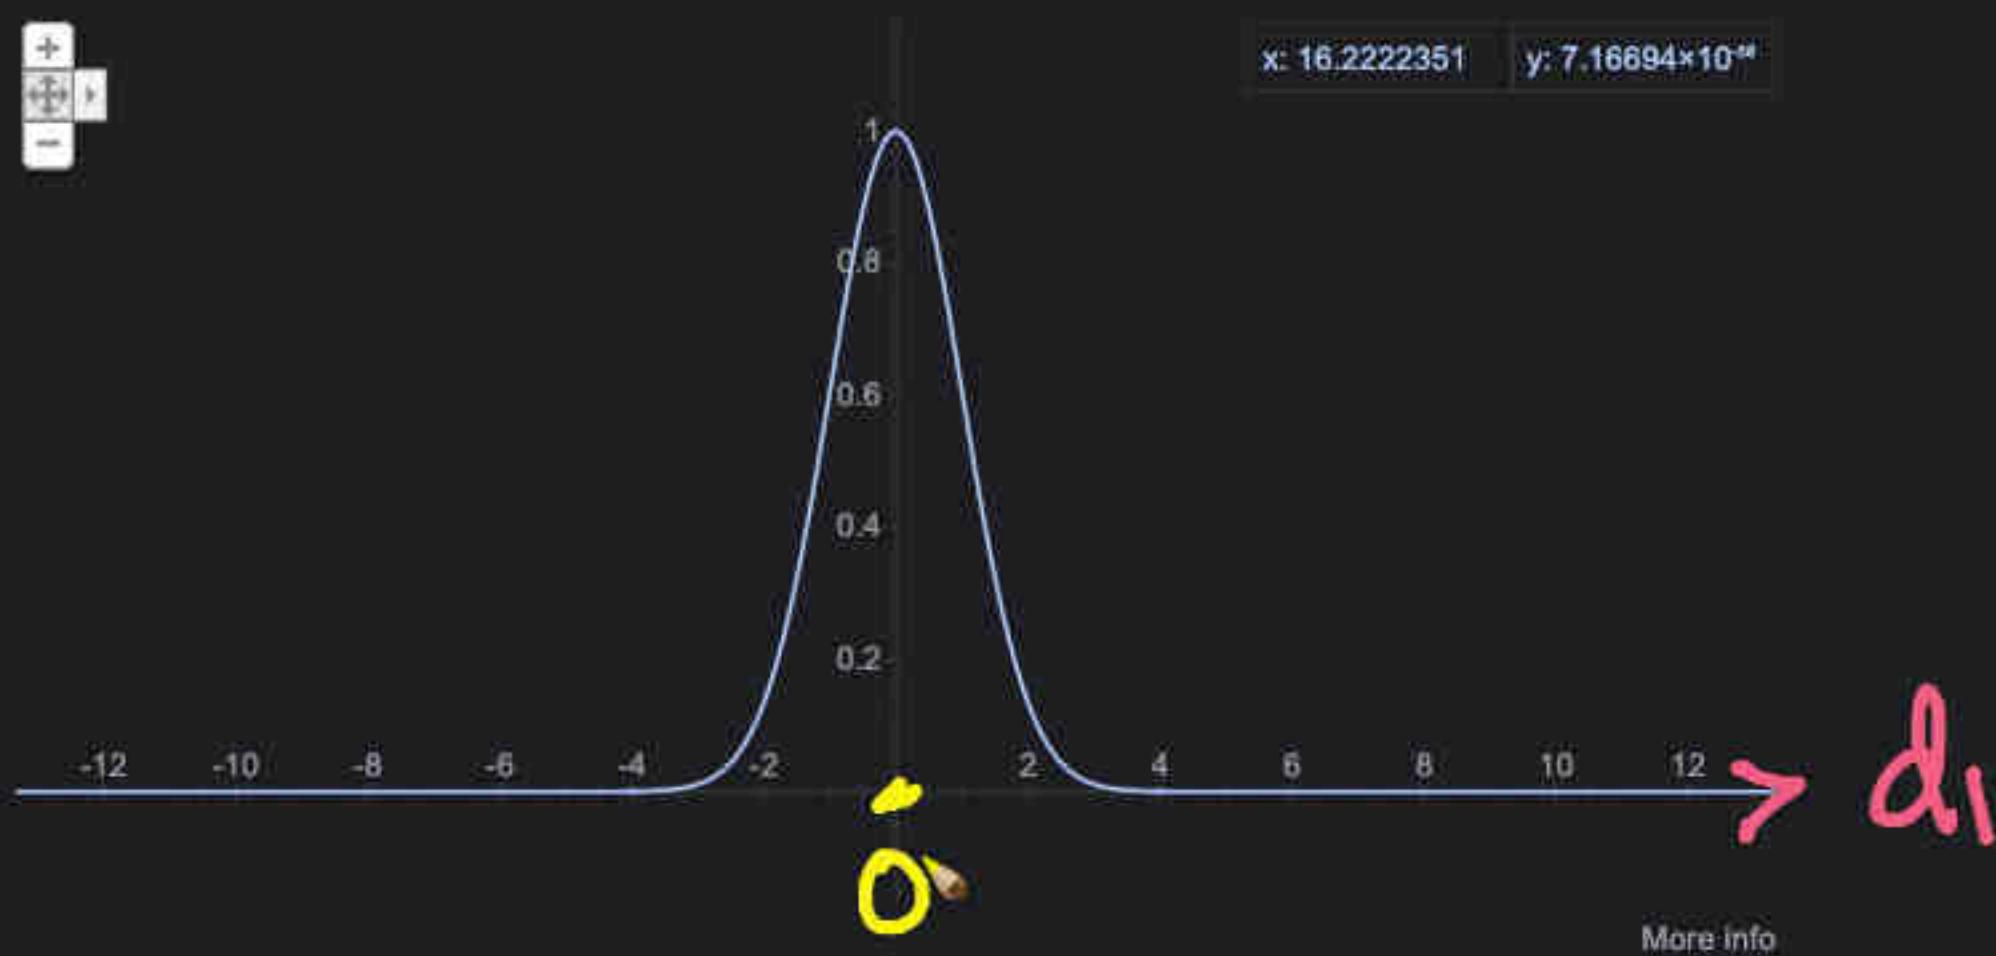
Google plot(exp(-x^2/(2*1*1)))

plot(exp(-x^2/(2*1*1)))

All Shopping Maps Videos News More Tools

About 2,65,00,00,000 results (1.09 seconds)

Graph for $\exp(-(x^2)/(2*1*1))$



<https://www.mathway.com/popular-problems/Calculus/>

Graph e^{-x^2} | Mathway

Graph e^{-x^2} . e-x² e - x 2. Find where the expression e-x² e - x 2 is undefined. The domain of the expression is all real numbers except where the ...

<https://www.wolframalpha.com/examples/mathematics/>

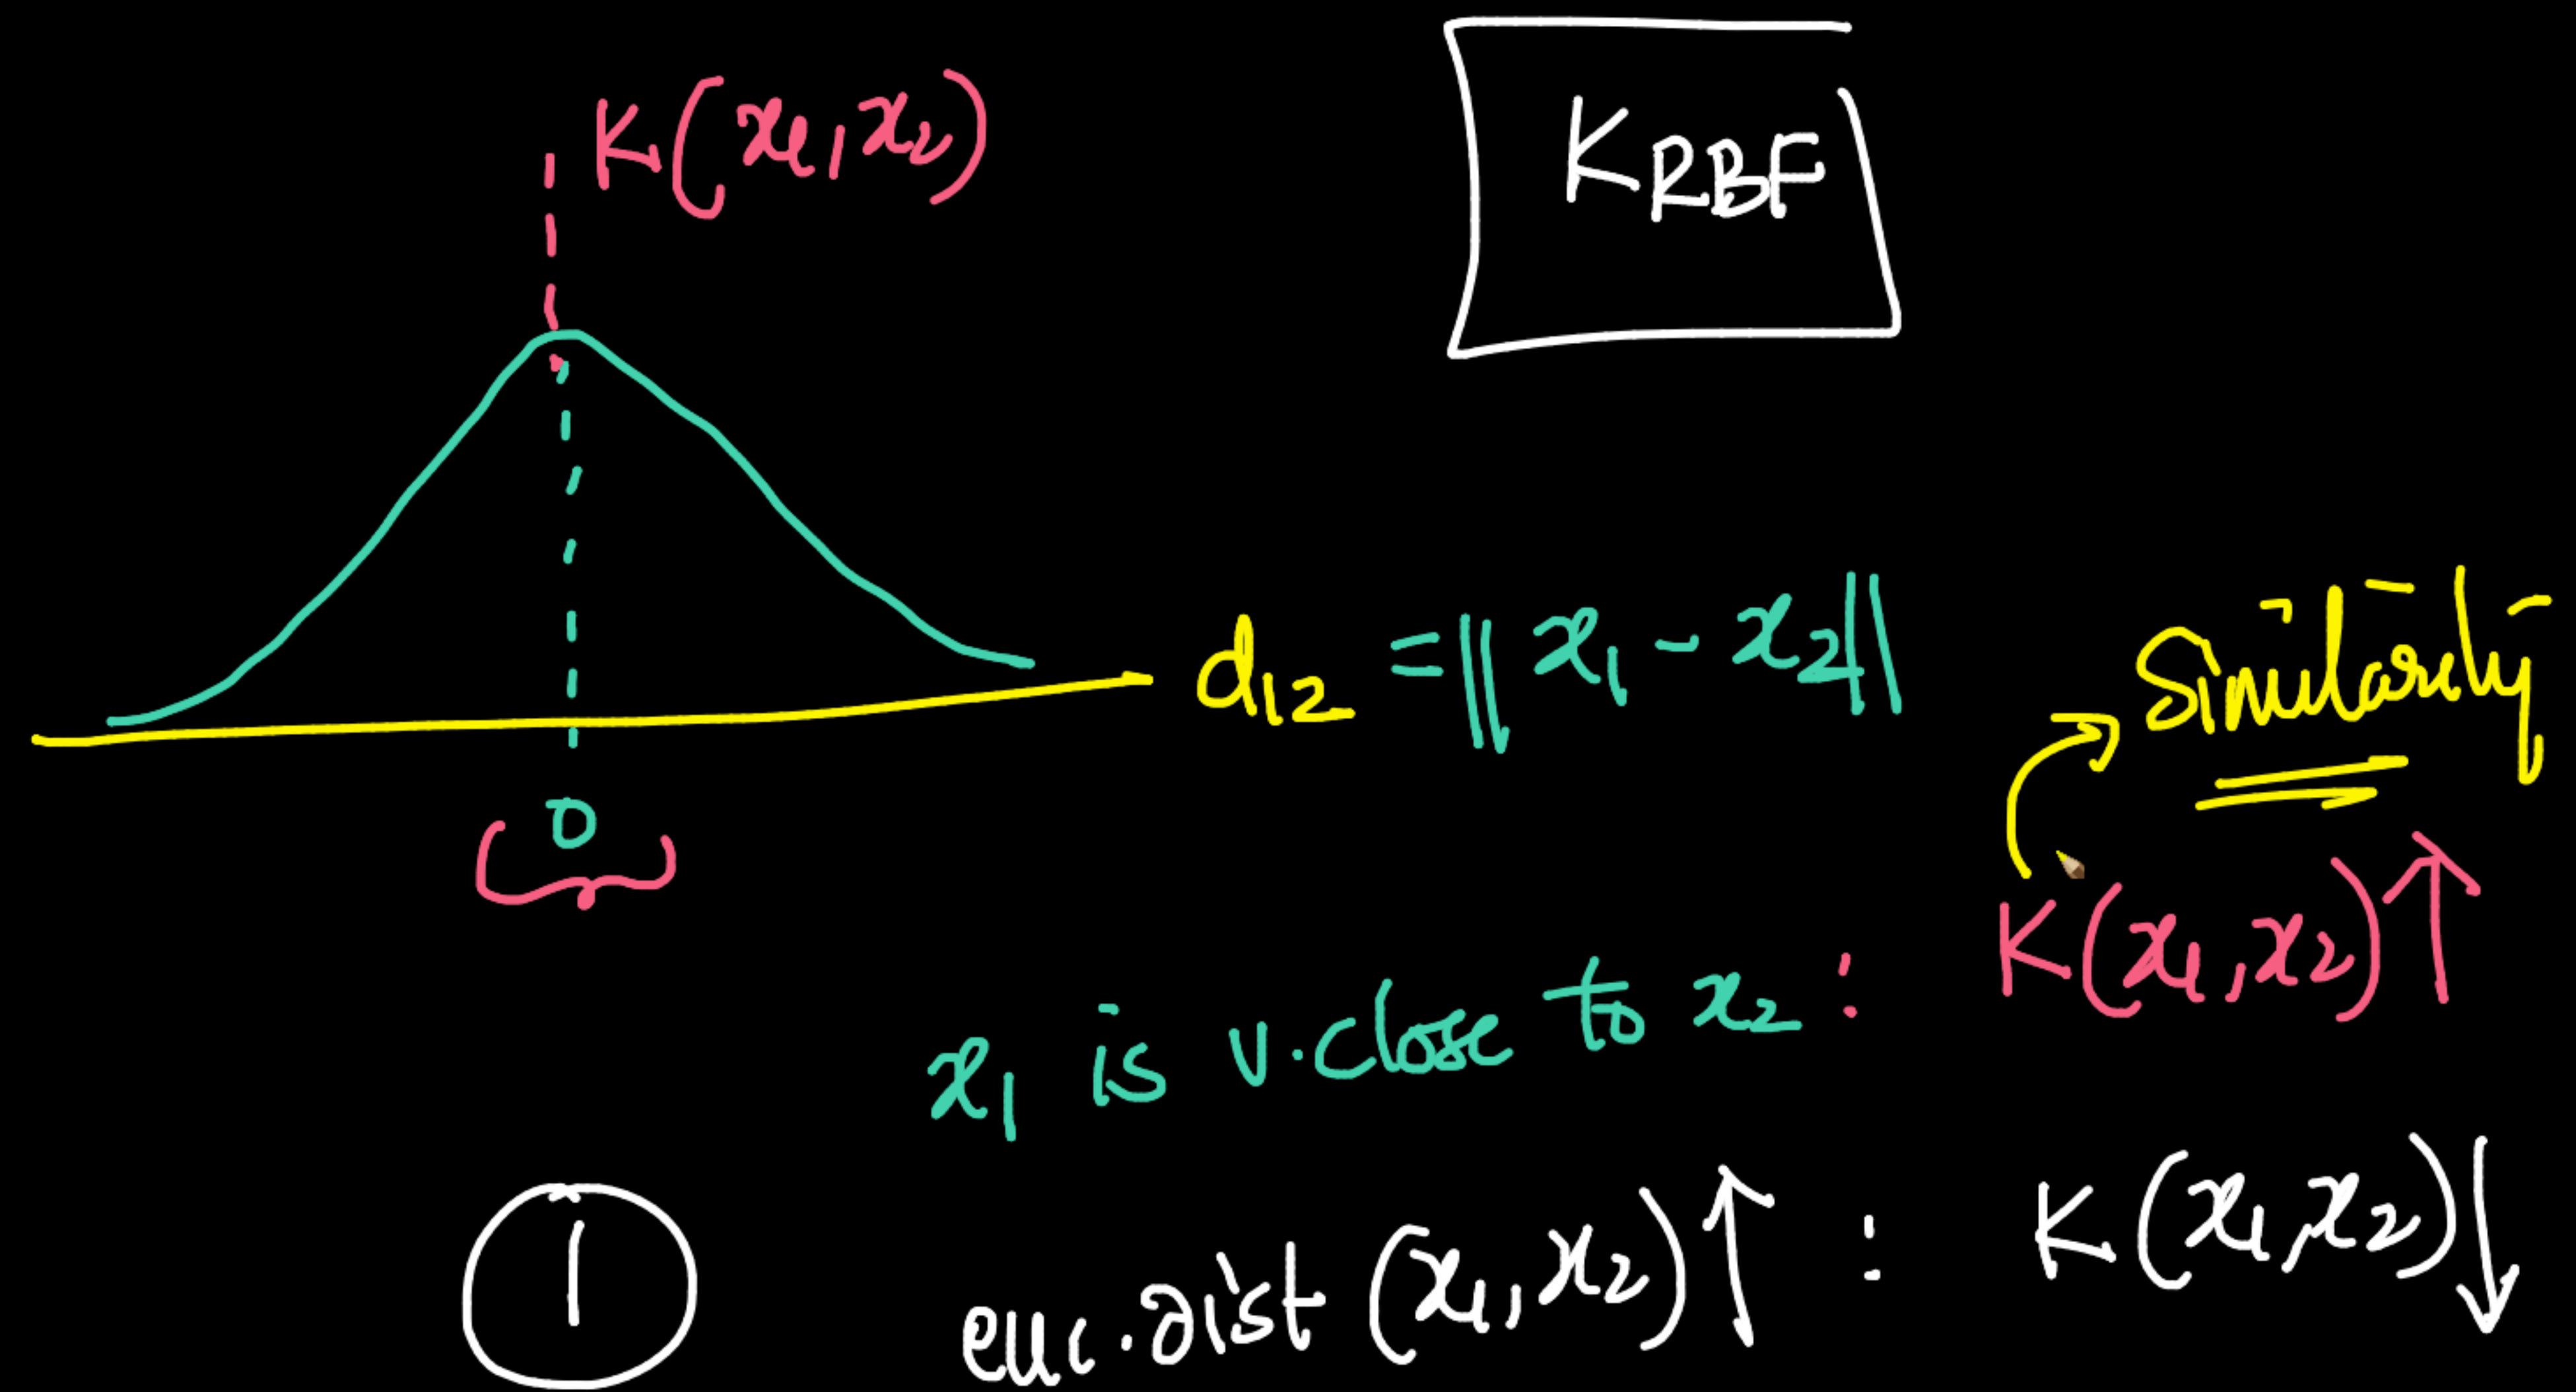
Plotting & Graphics - Wolfram|Alpha Examples

Use interactive calculators to plot and graph functions. Try 3D plots, equations, inequalities, polar and parametric plots. Specify ranges for variables.

Plot sin x, cos x, tan x · Plot x^2 y^3, x=-1..1, y=0..3 · Log-linear plot x^2 log x, x=1...

$$\exp\left(-\frac{x^2}{2r^2}\right)$$

- ① Gaussian
- ② $r \downarrow$; Gaussian flinney



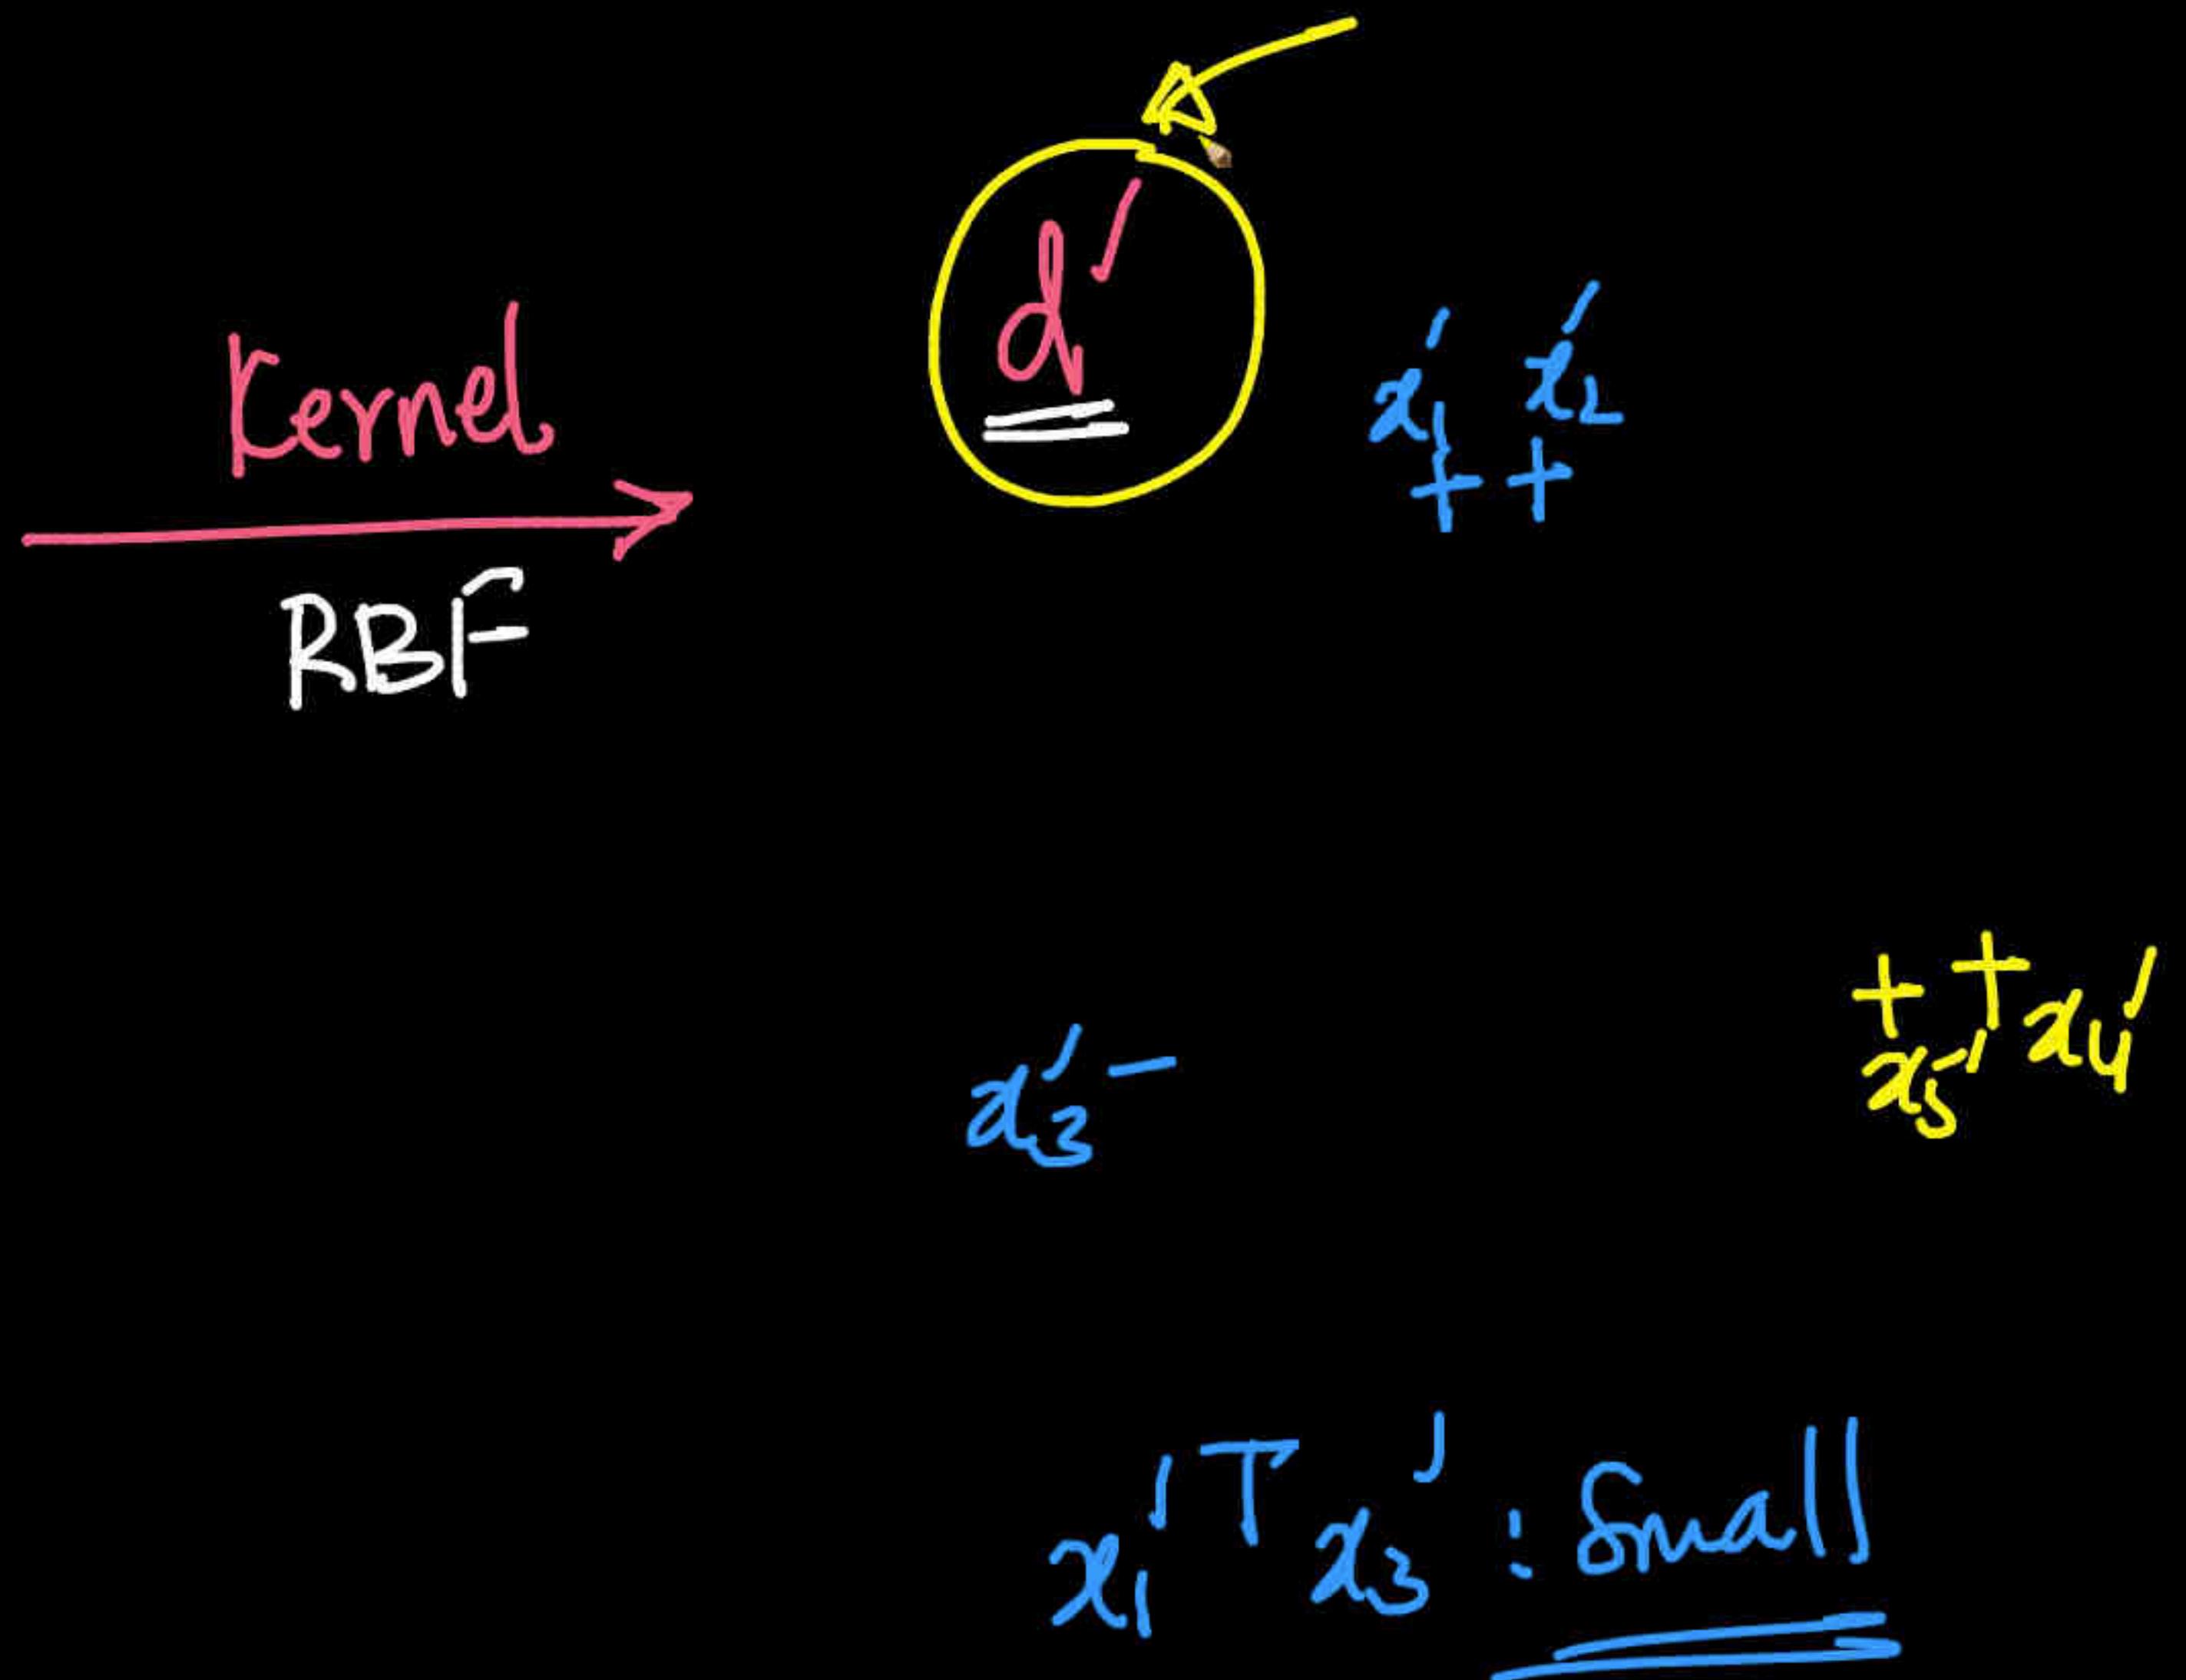
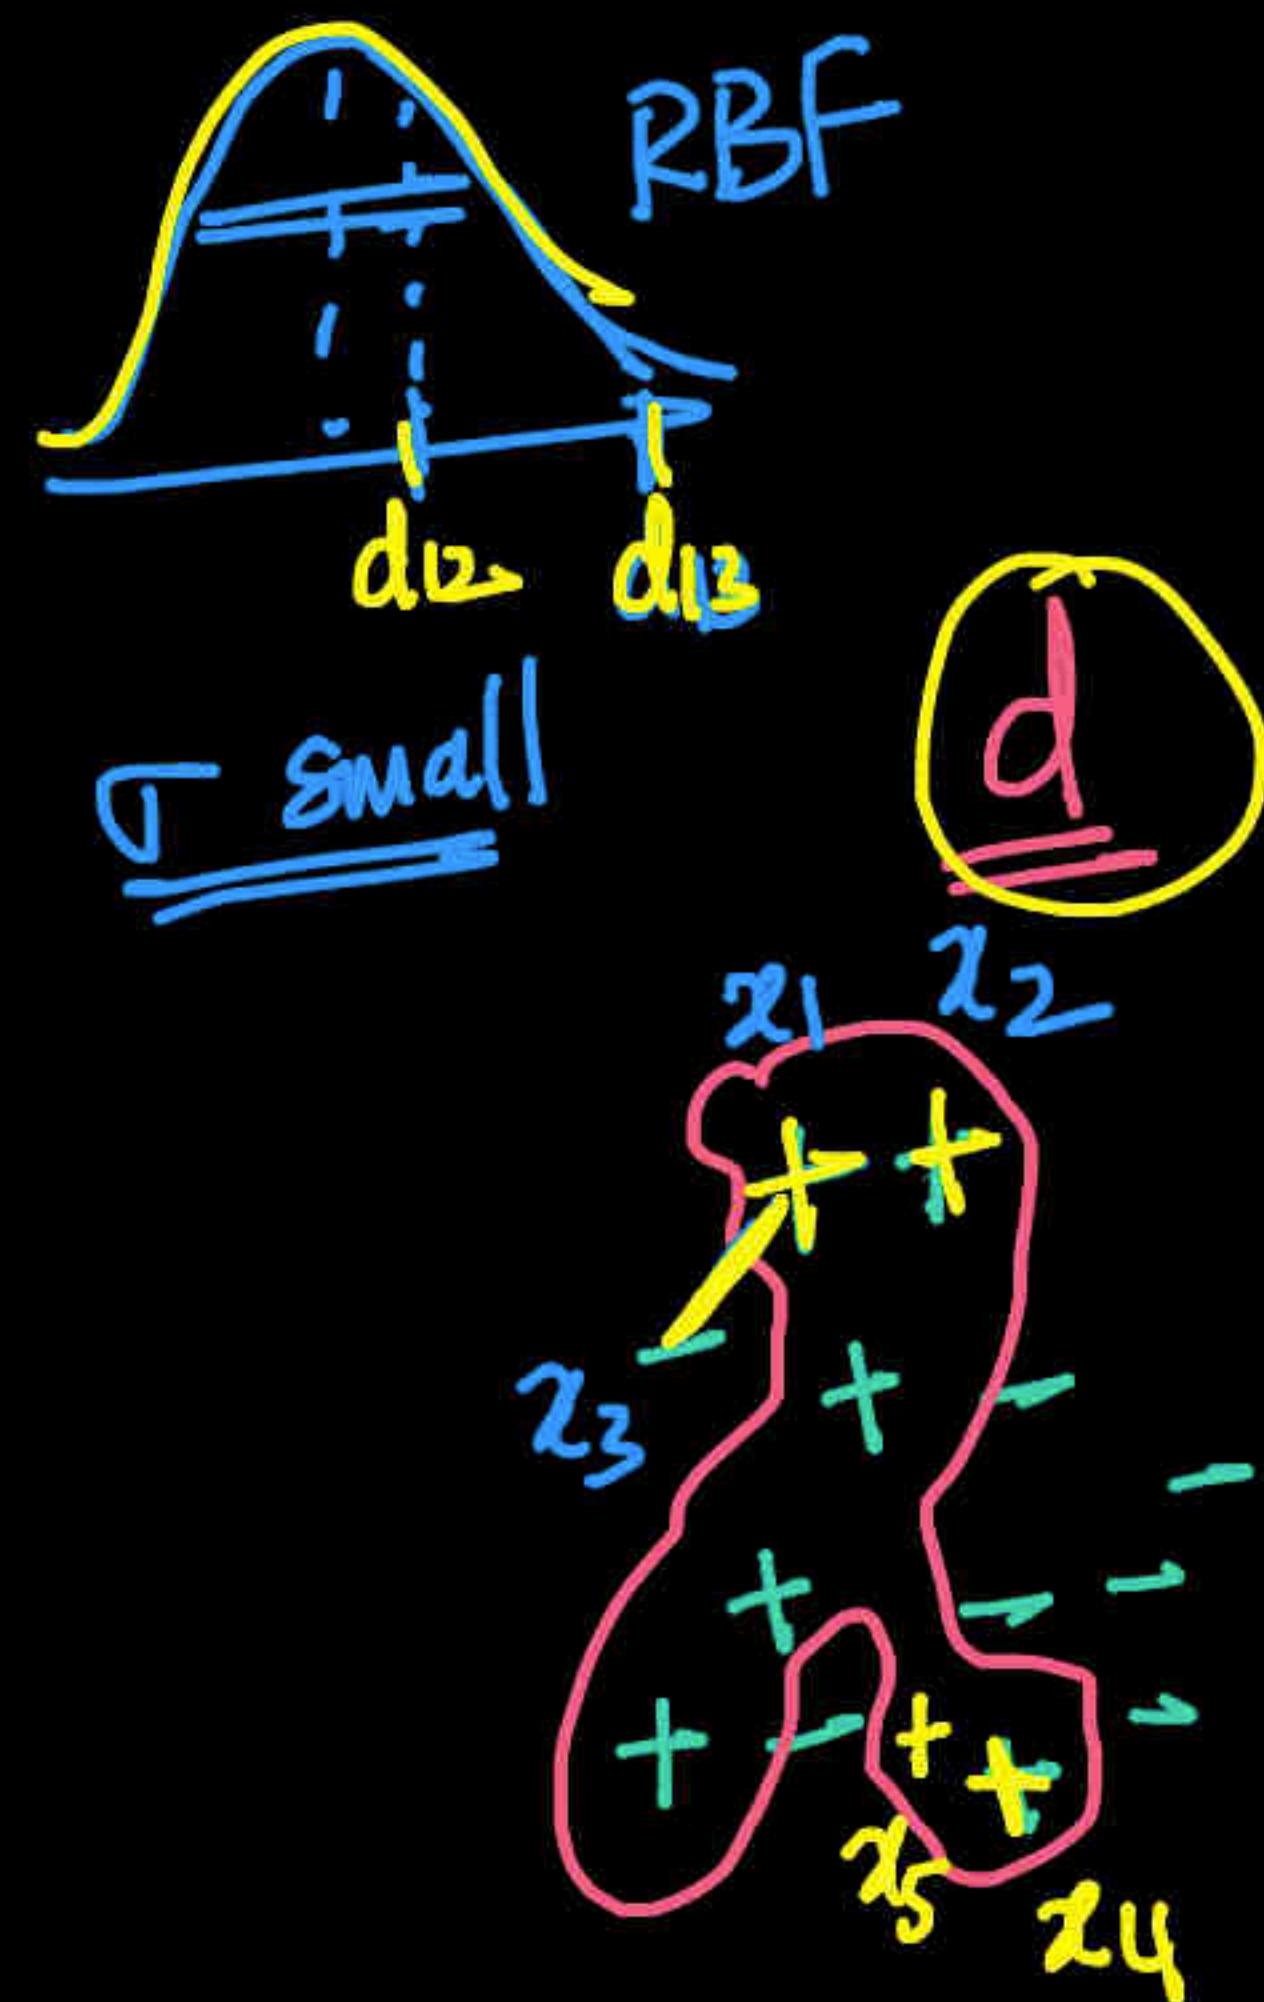
- ① $\text{euc.dist}(x_1, x_2) \downarrow \Rightarrow K_{RBF}(x_1, x_2) \uparrow$
(Sim)
- ② $\text{euc.dist}(x_1, x_2) \uparrow \Rightarrow \text{Sim or } K_{RBF} \text{ falls exponentially}$

SVM + RBF

{ dual
formulation

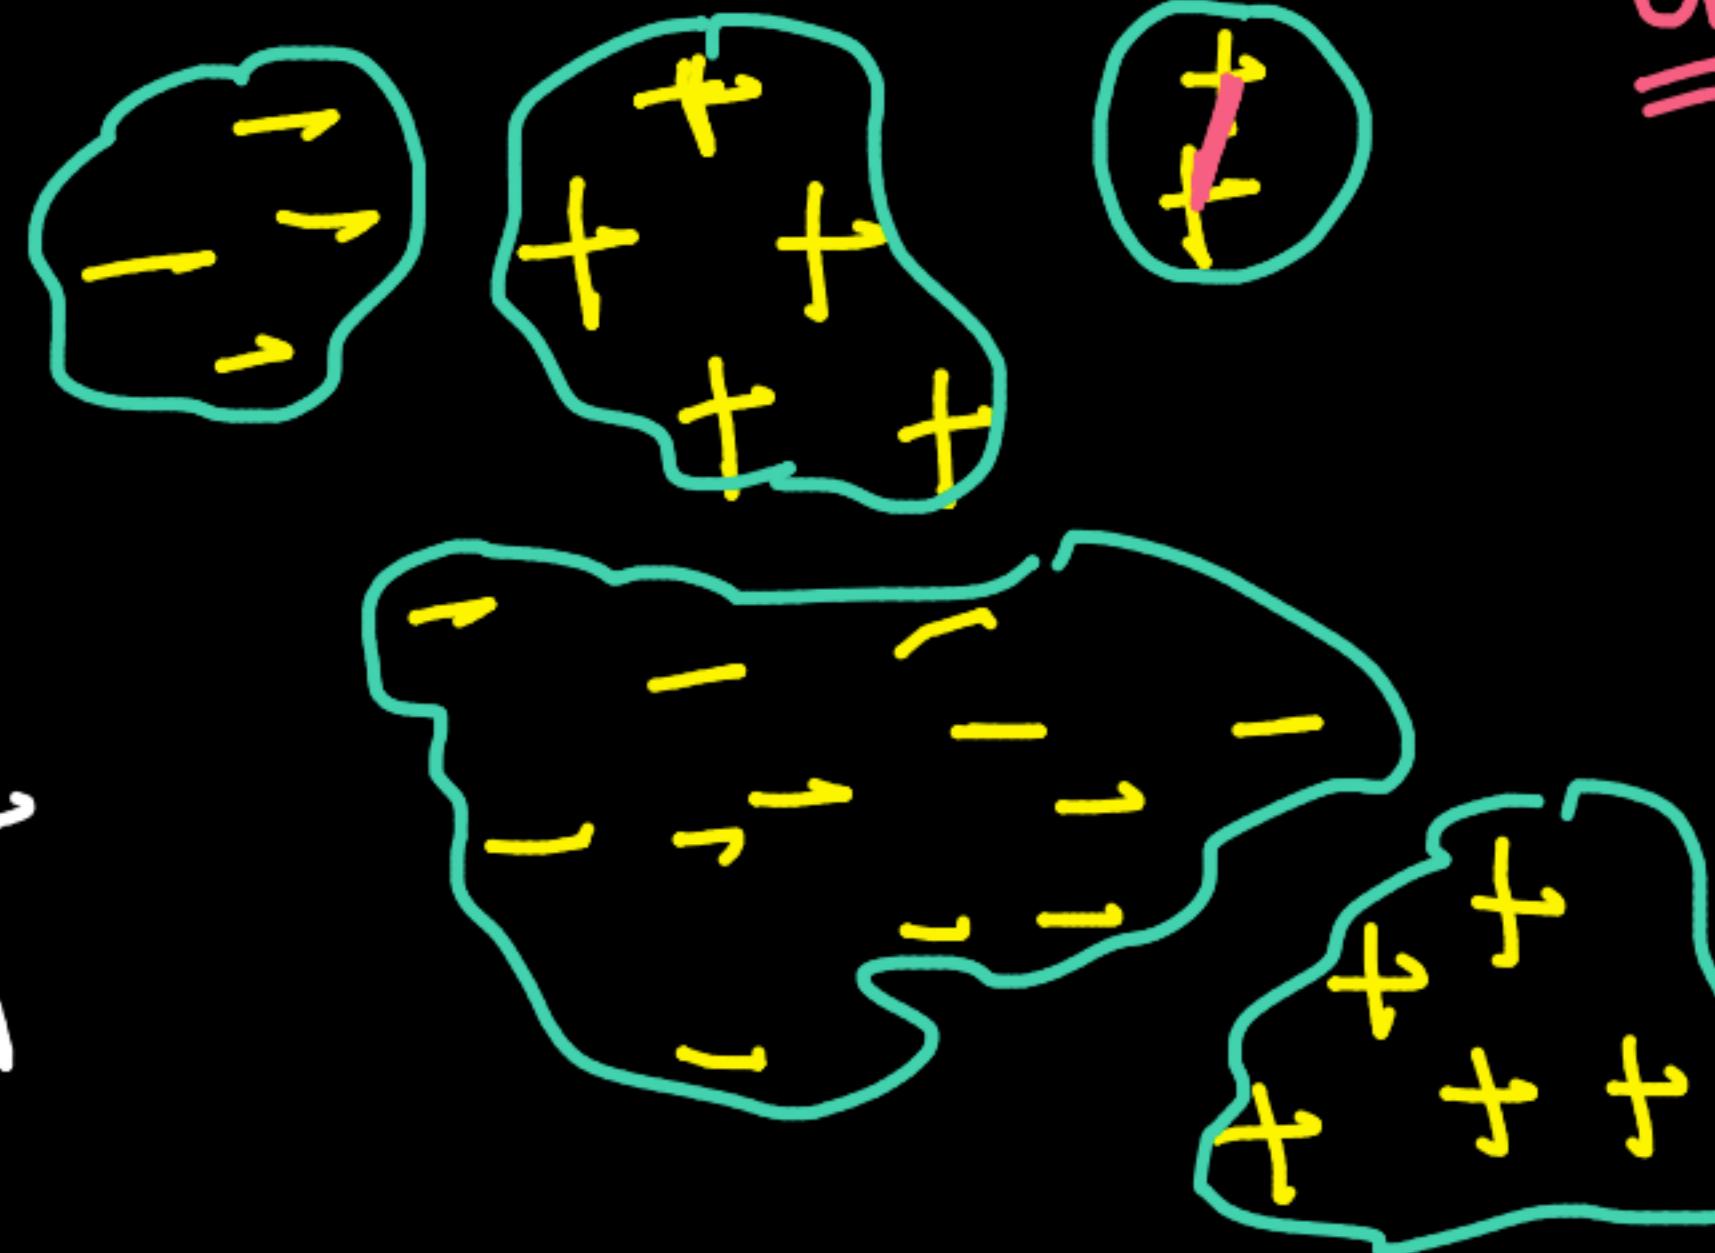
$$x_i^T x_j : K_{RBF}(x_i, x_j)$$

x_i is close to $x_j \rightarrow$



d -dim

proximity
based
Model



{ RBF-SVM
with small σ

overfit

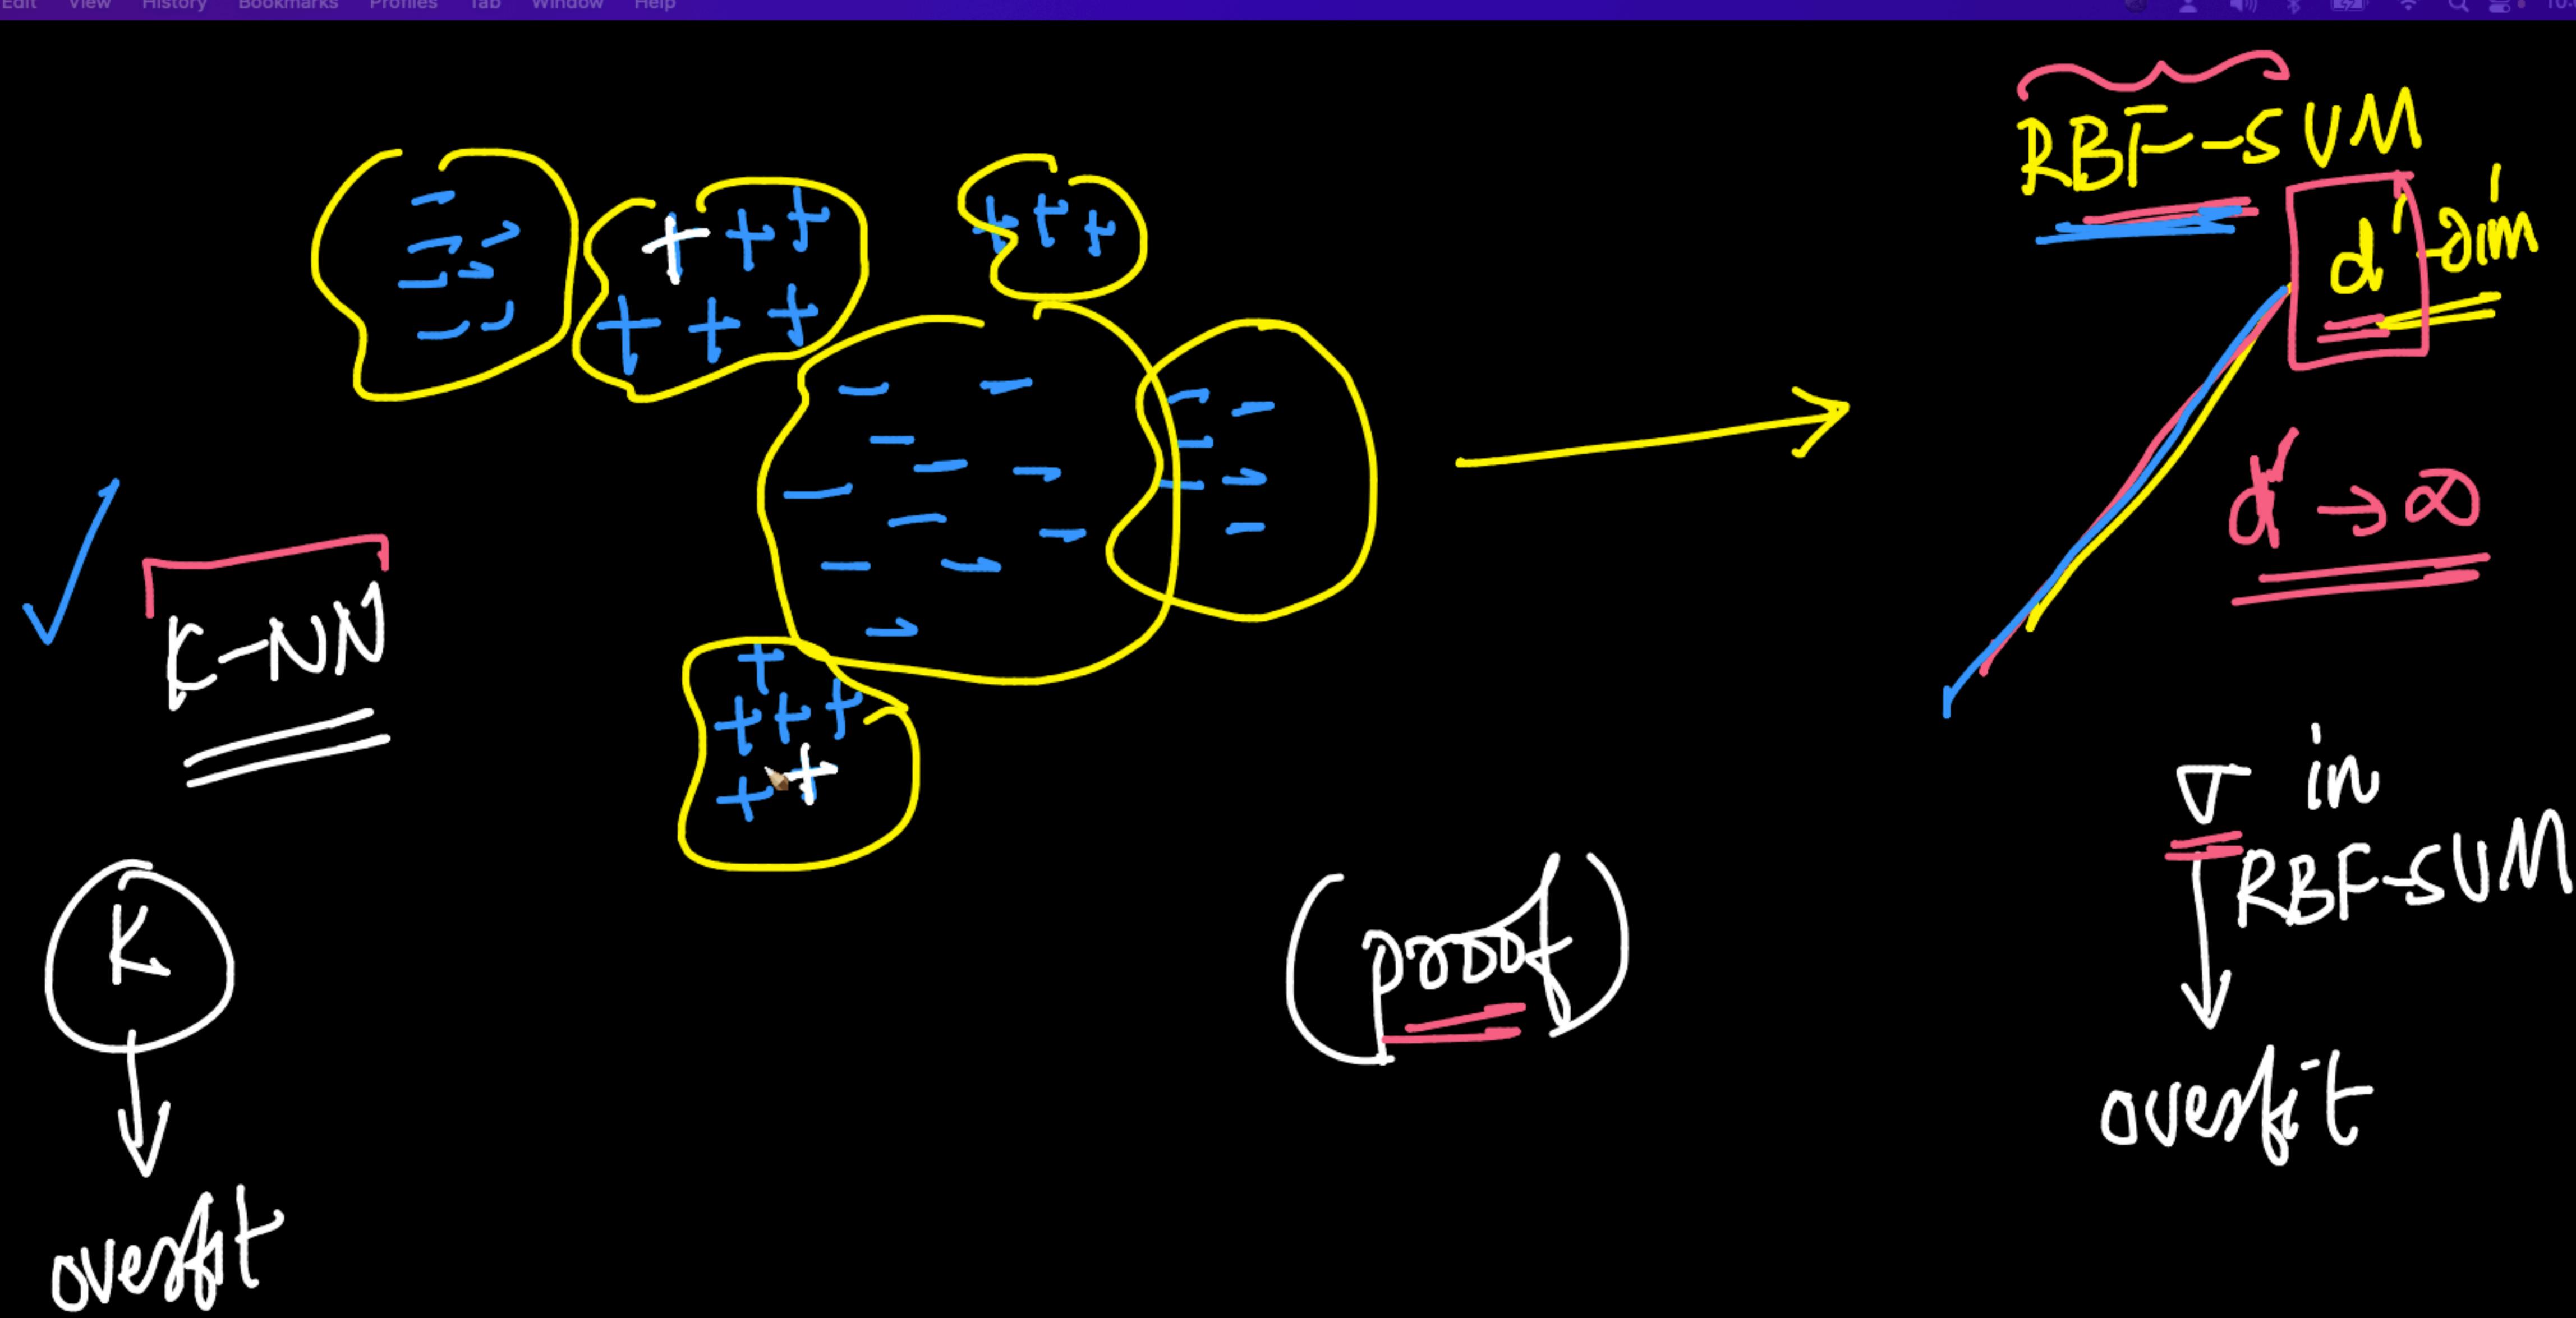
d'

σ is small

KNN with small

K

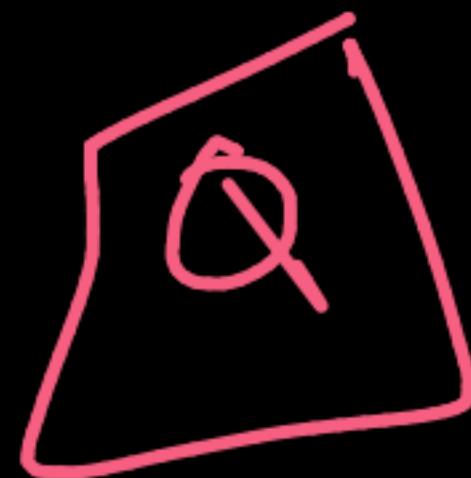
outlier
overfit



K=1 kNN

++ - - ++ - - ++ - - 

- - ++ - -

KNN \rightarrow $d \uparrow$

curse of dim

Euc. dist doesn't work

RBF-SUM!

$$\| \underline{x}_1 - \underline{x}_2 \| = d_{12} \xrightarrow{\text{RBF}} \text{func}(d_{12})$$

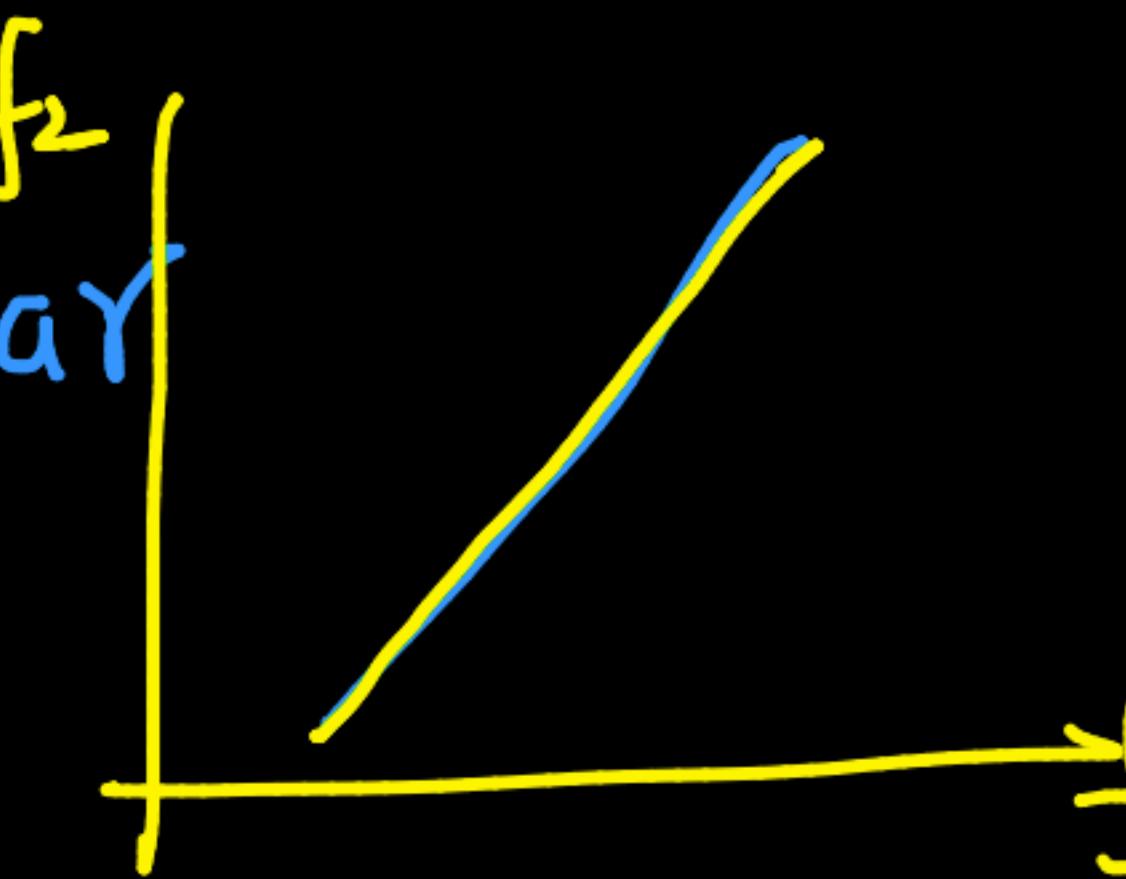
$\underline{x}_i \in \mathbb{R}^{\hat{d}}$ $\underline{d=10}$

INTUITION



$f_2 \rightarrow$ linear

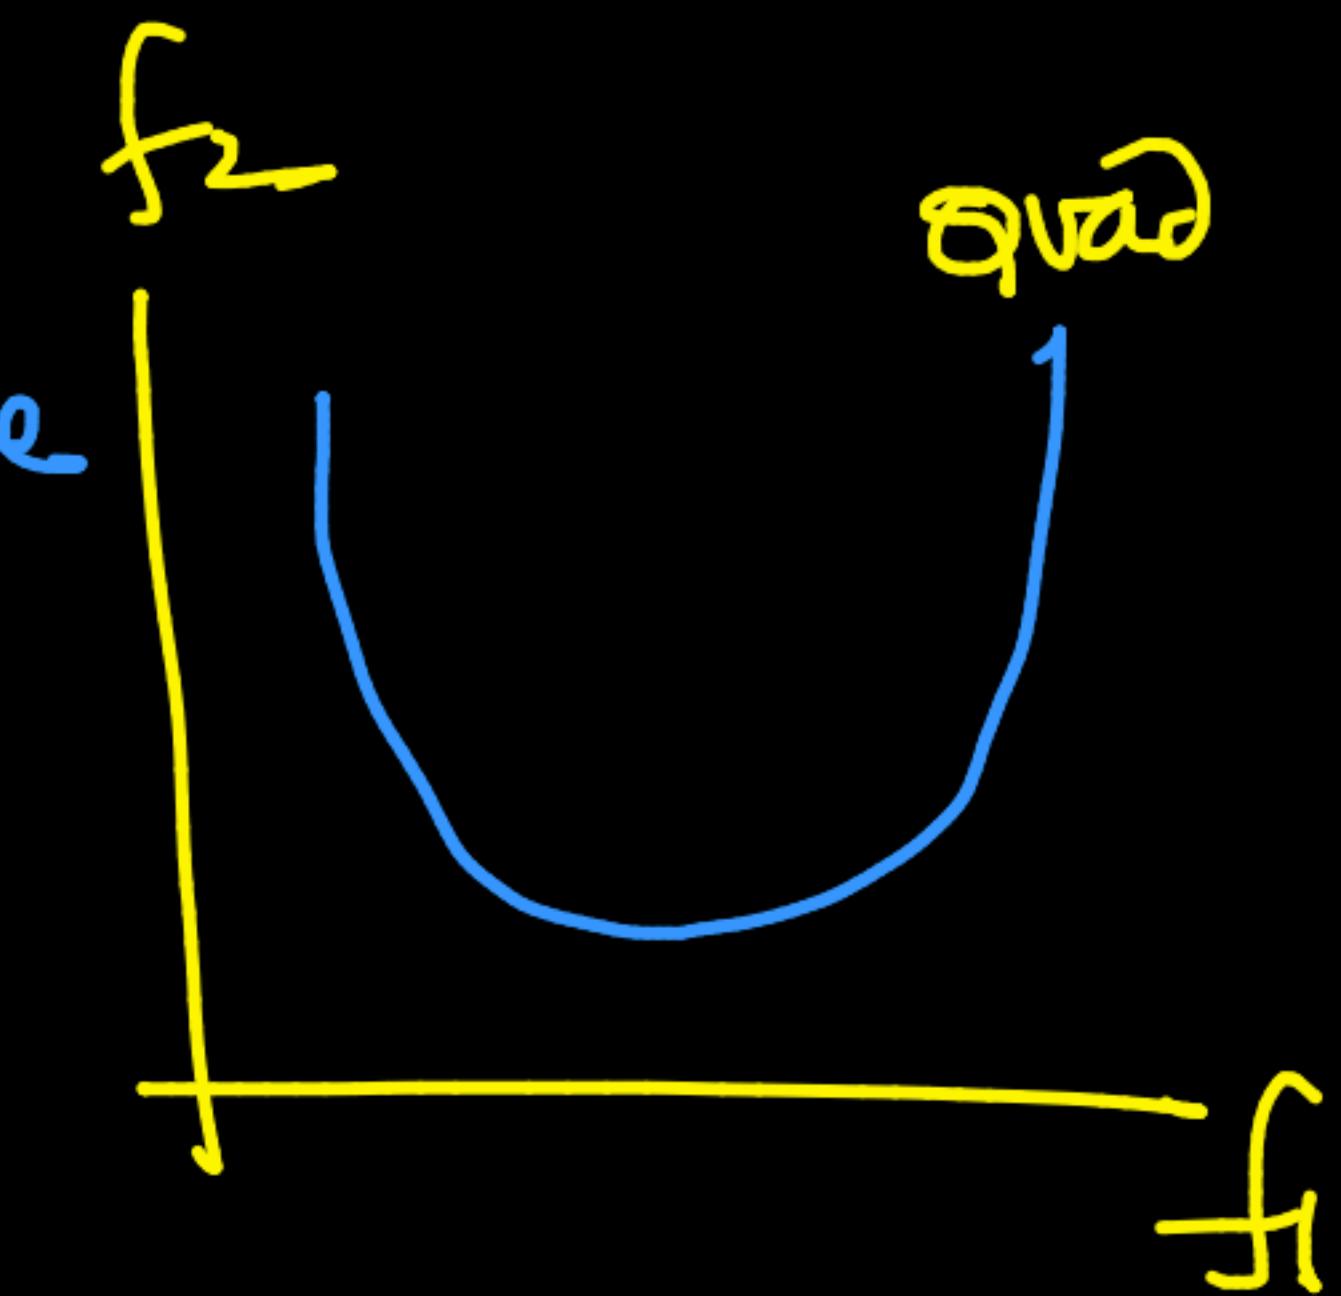
② f_1, f_2



Linear-model

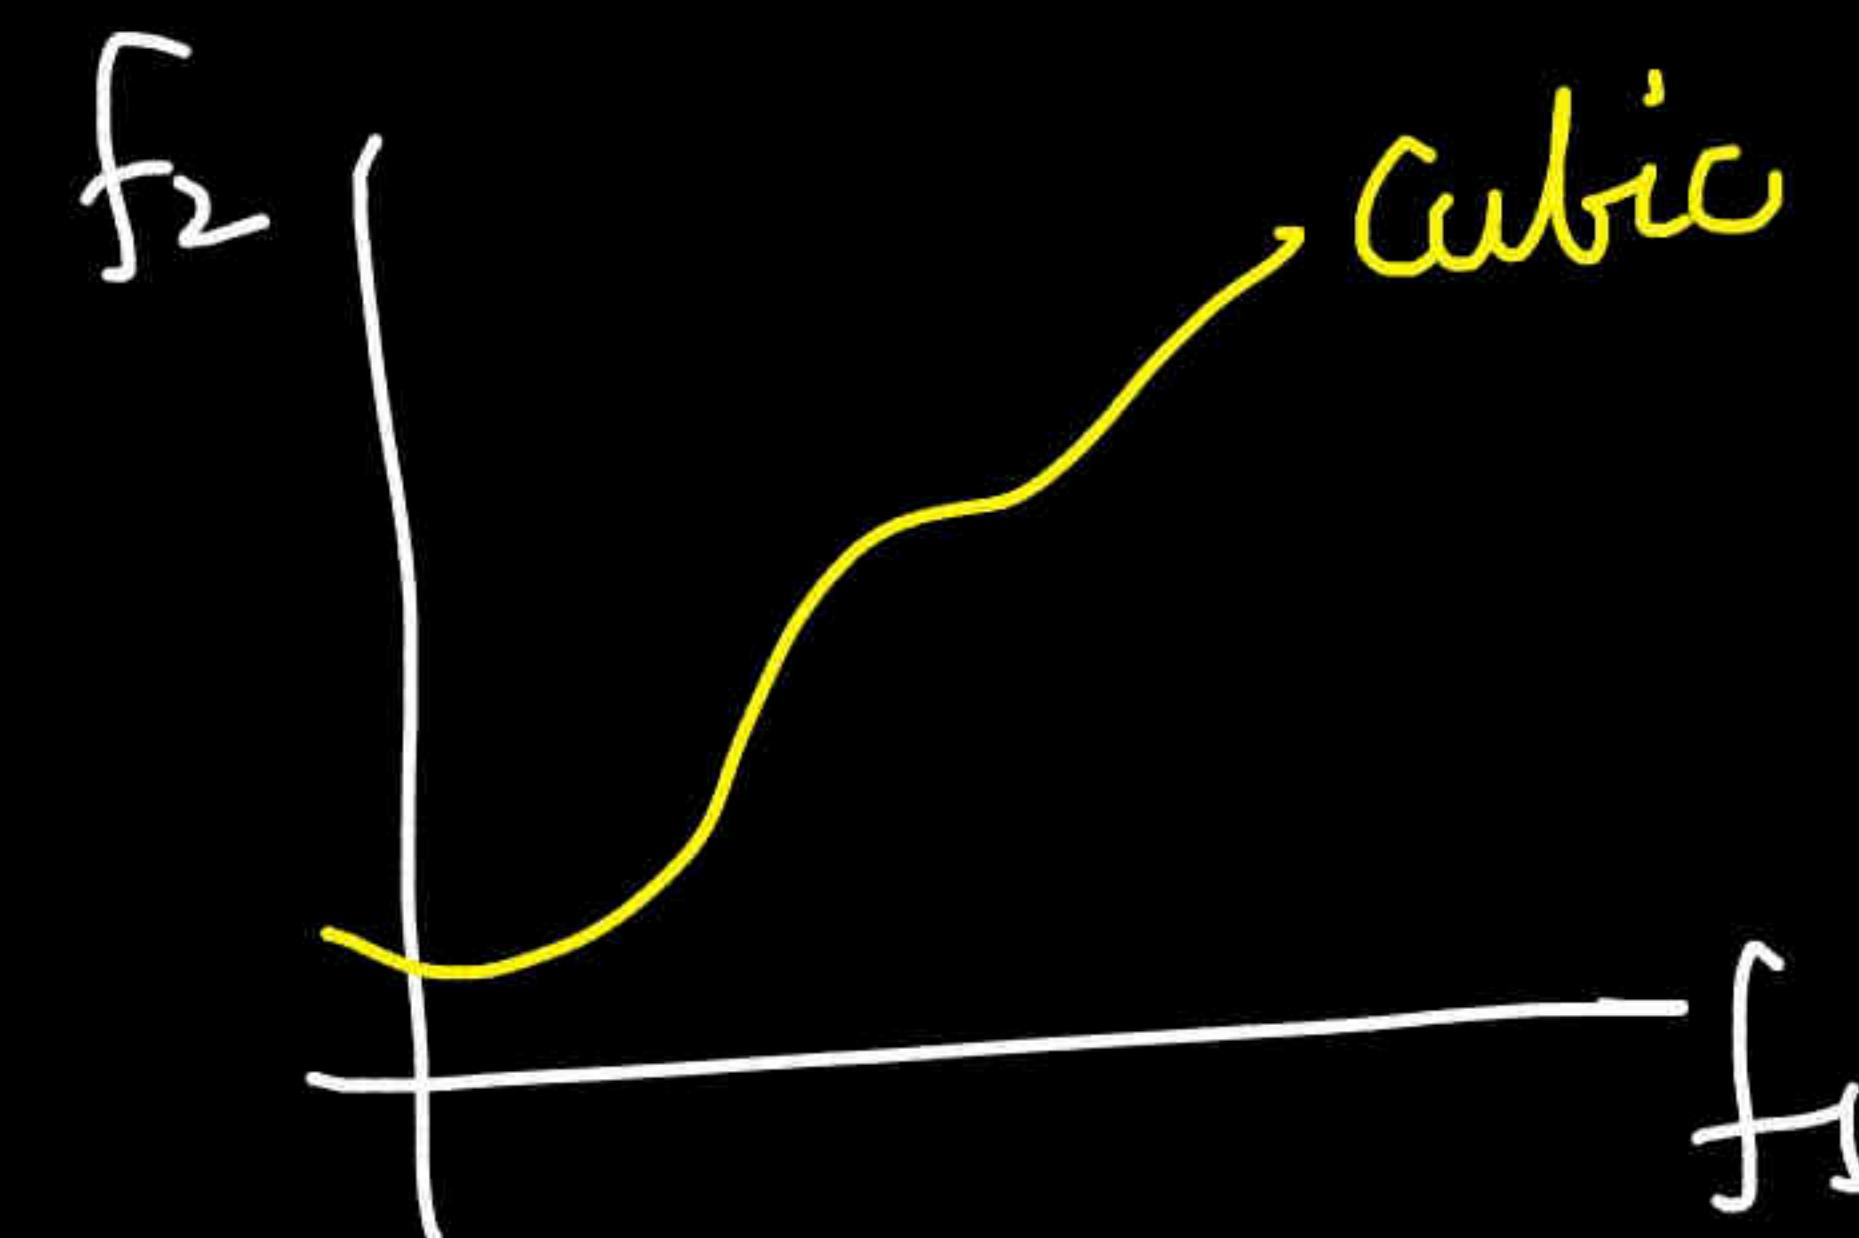
⑤ $f_1f_2, f_1^2, f_2^2, f_1f_2$

Quadratic curve



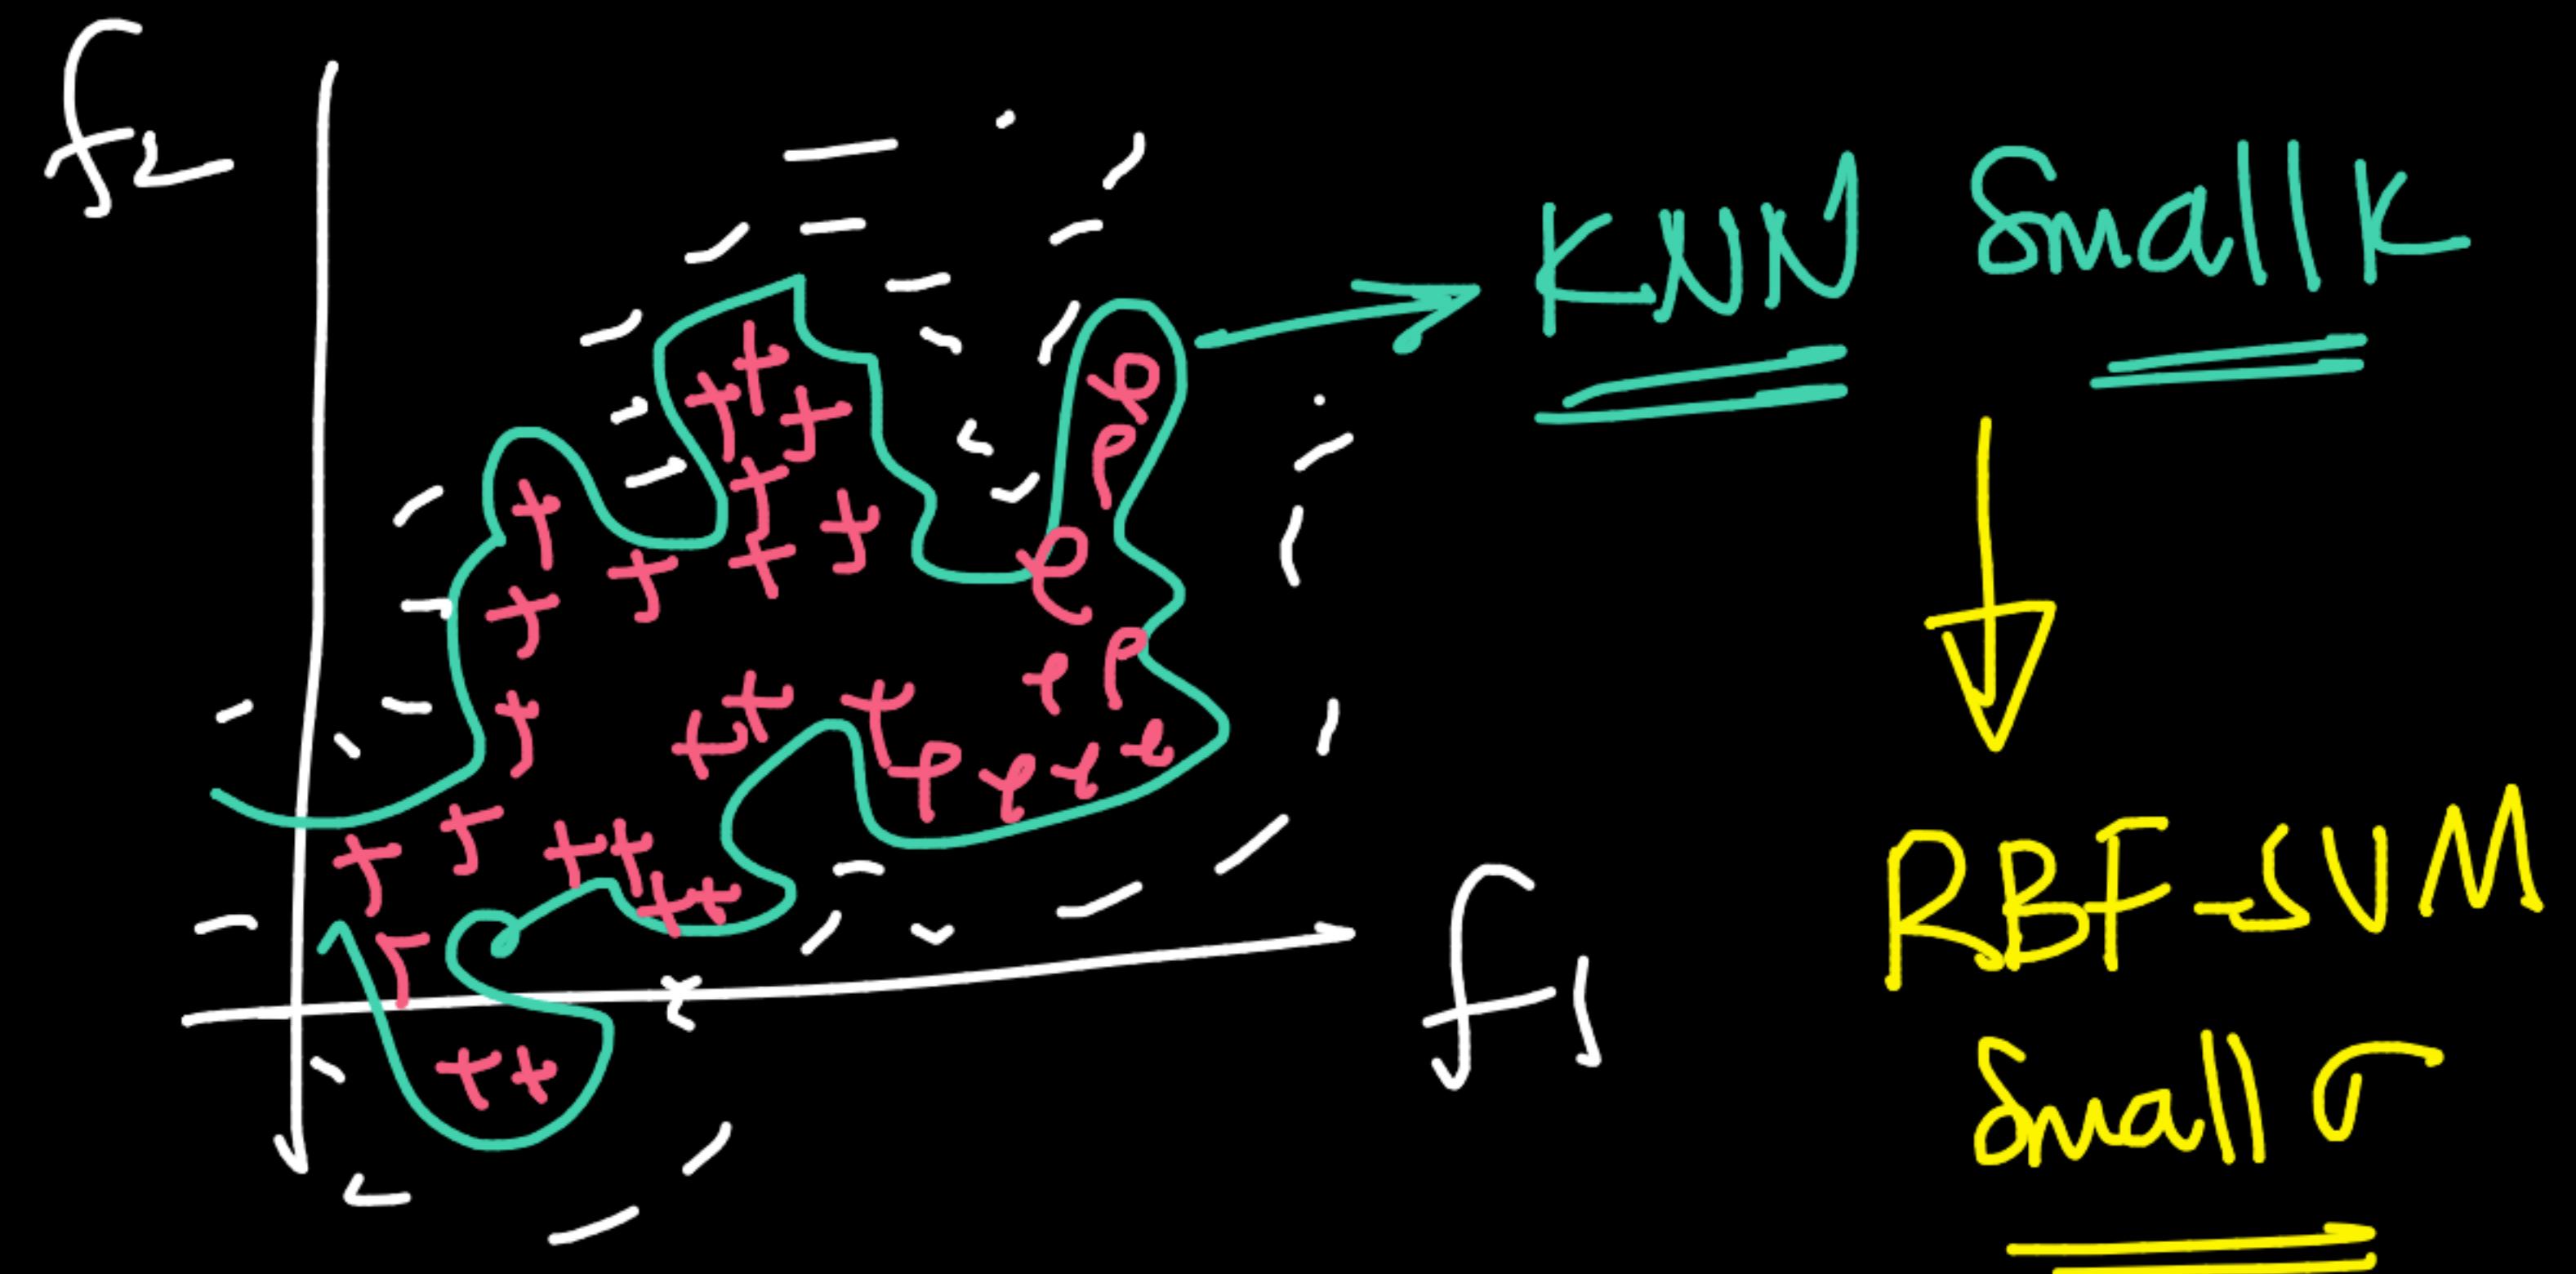
Linear model

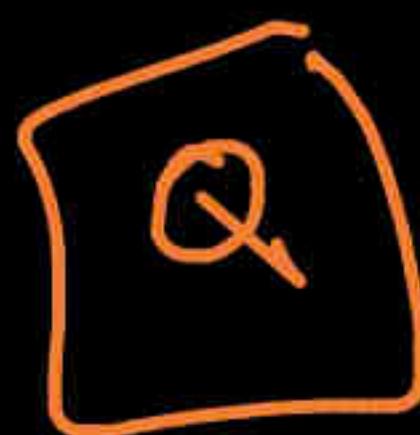
9 → $f_1, f_2, f_1^2, f_2^2, f_1 f_2, f_1^3, f_2^3, f_1^2 f_2, f_2 f_1^2$



many :-

$f_1, f_2, f_1^2, f_2^2, f_1^3, \dots$





! $\Rightarrow \approx$
+ - -
+ + - -
+ + - -

KNN with $k=1$

||

RBF-SVM small r

Pink: 2a)

Blue: train

+ + +



$K(x_1, x_2)$
RBF

d-dim

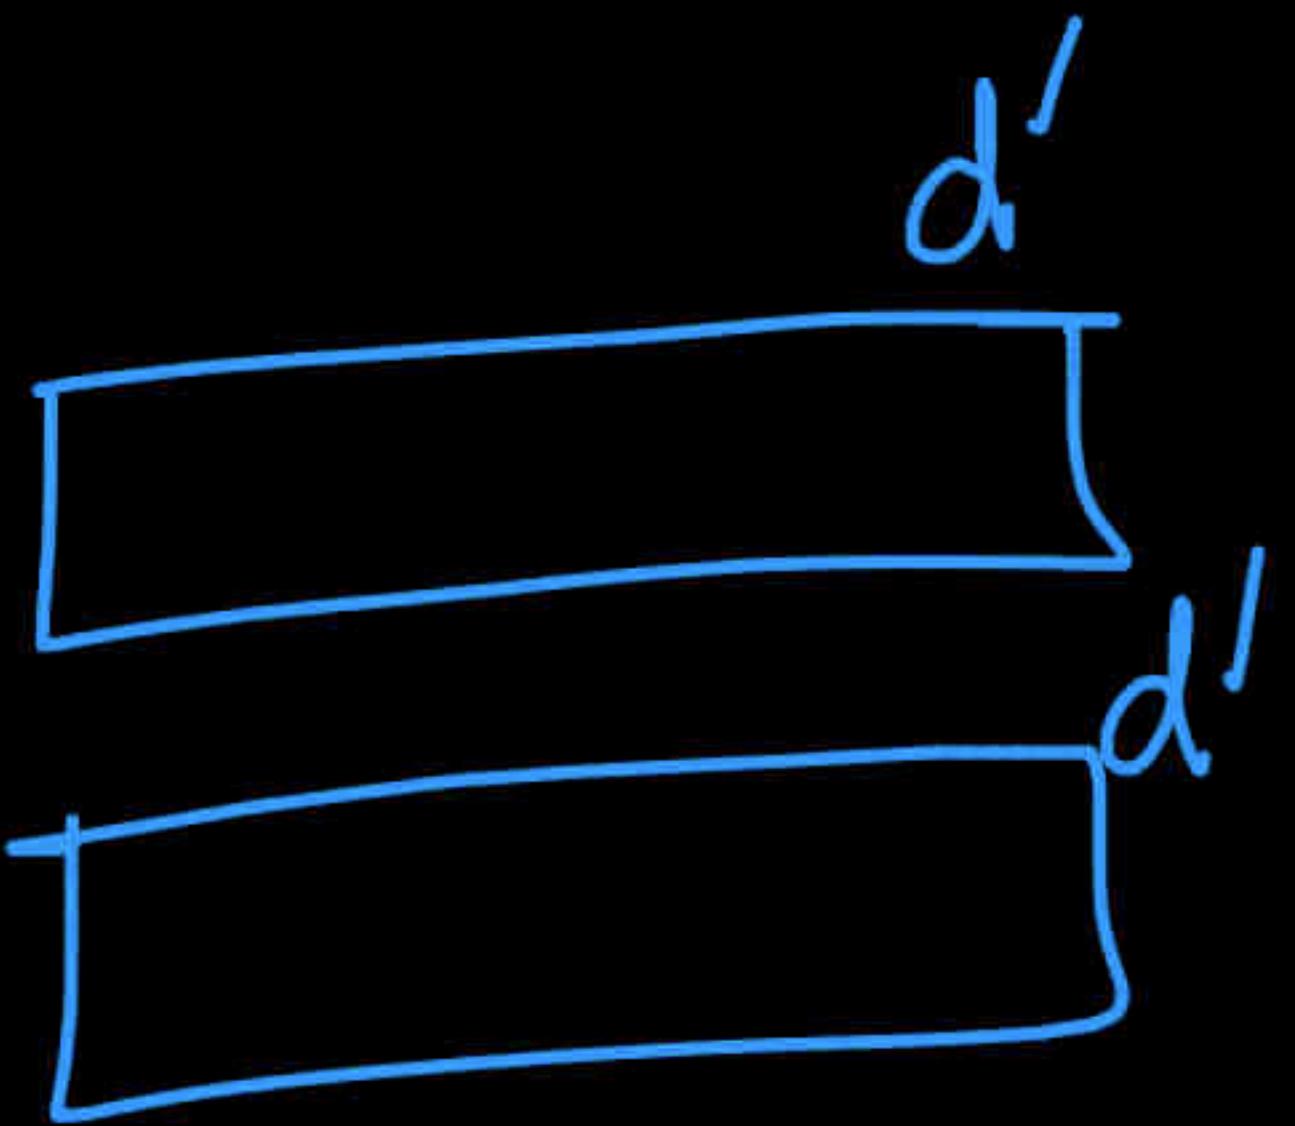
d

Sim

+ x_1
+ x_2

d'
+ x_1'

x_1' :



x_2' :

IMPOSSIBLE

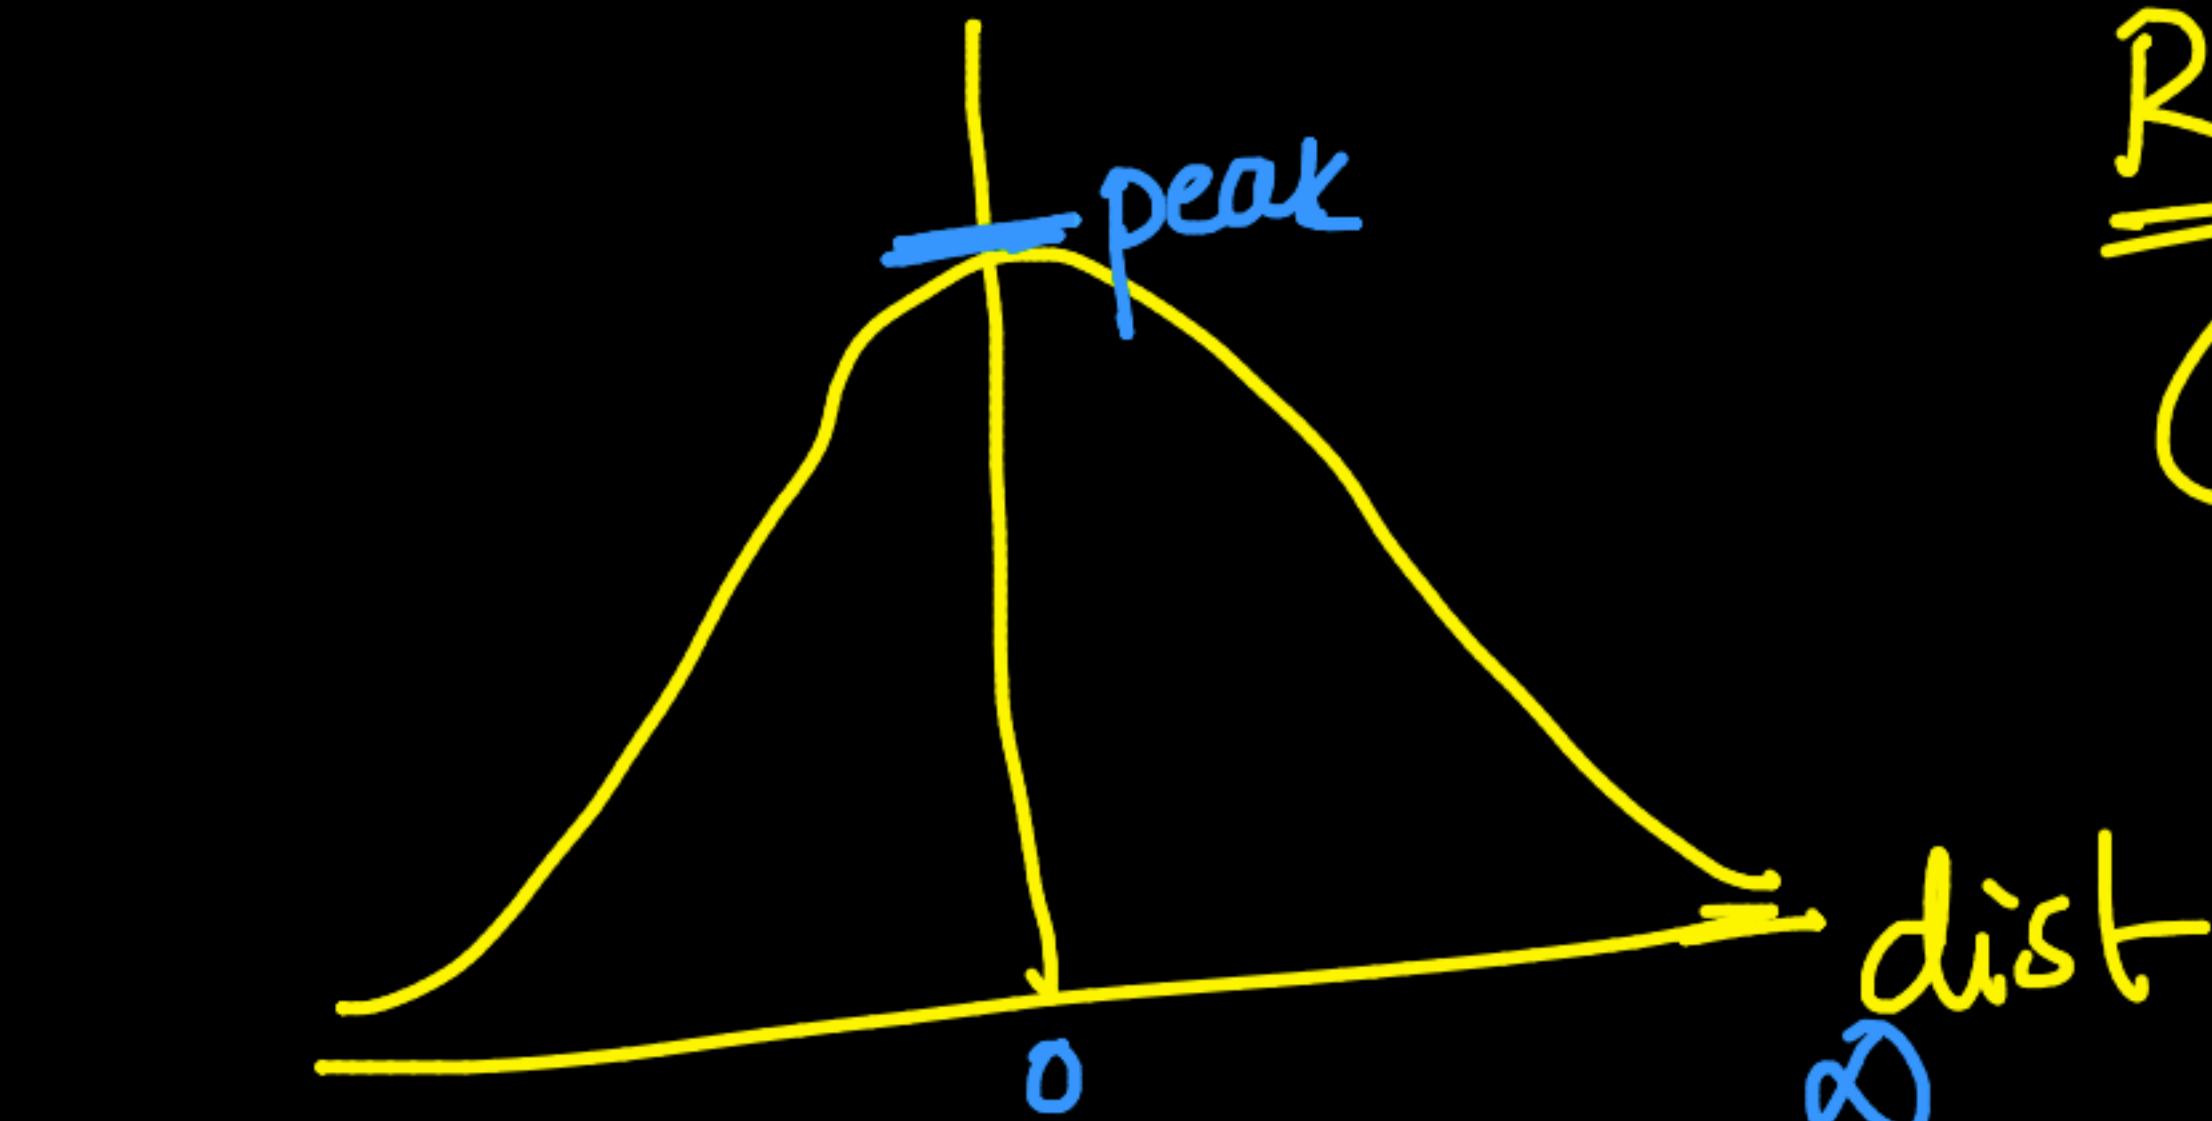
$d' \rightarrow \infty$ - large





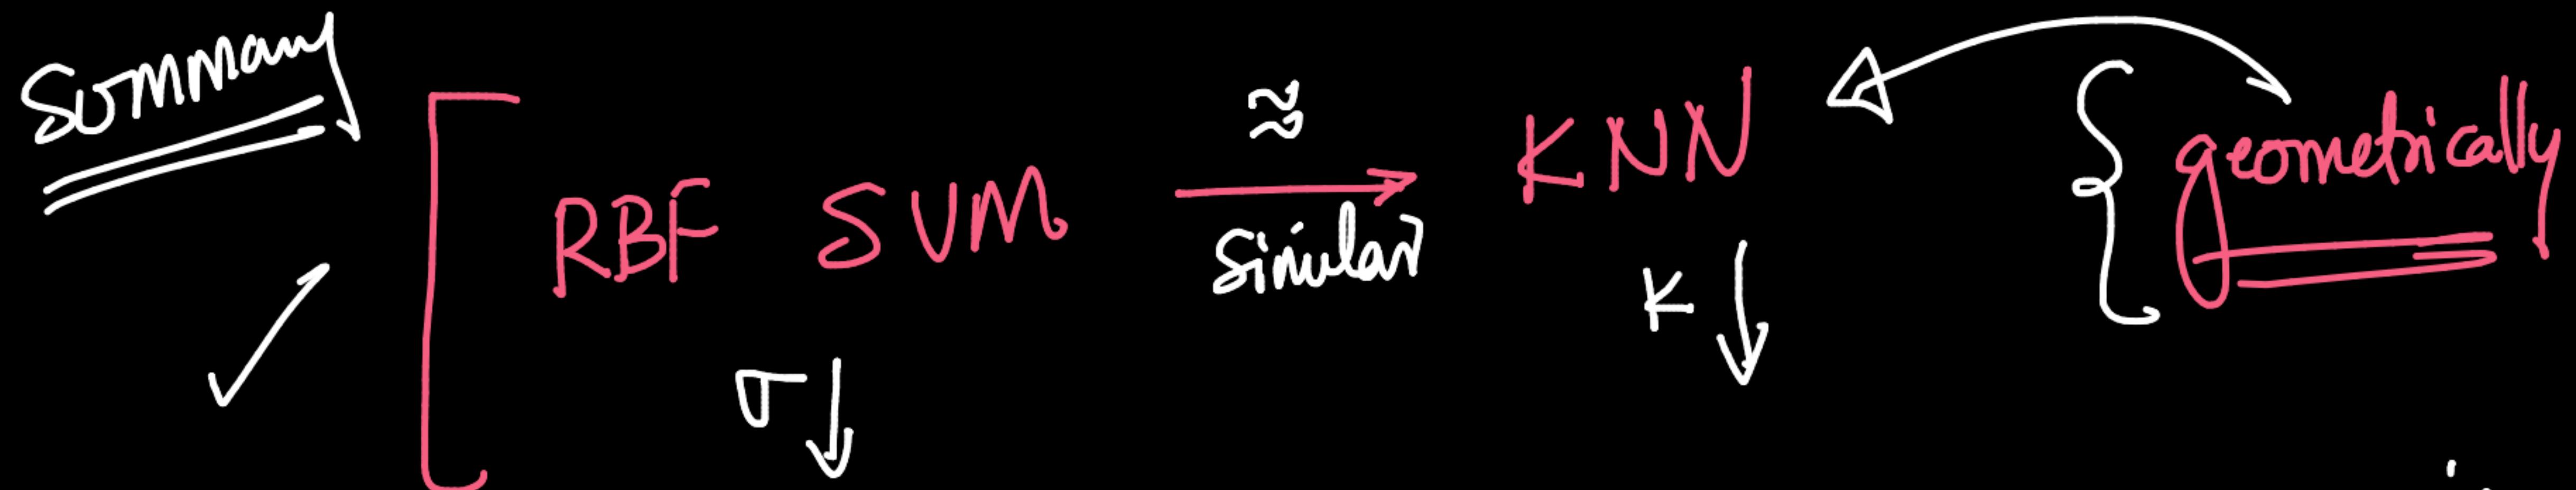
$$-\infty \leq \text{dist}(x_1, x_2) \leq +\infty$$

$d \rightarrow d'$
↓
loop
(let)
↓
euc.dist



$$\|x_1 - x_2\|^2 = d_{12}^2 \text{ (use adjm)}$$

RBF
 $\text{G dist} \rightarrow \text{RBF}(\text{dist})$



Run-time complex SVM:- $f(\underline{x}_a) = \sum_{i=1}^n \alpha_i y_i k(\underline{x}_i \underline{x}_a) + b$

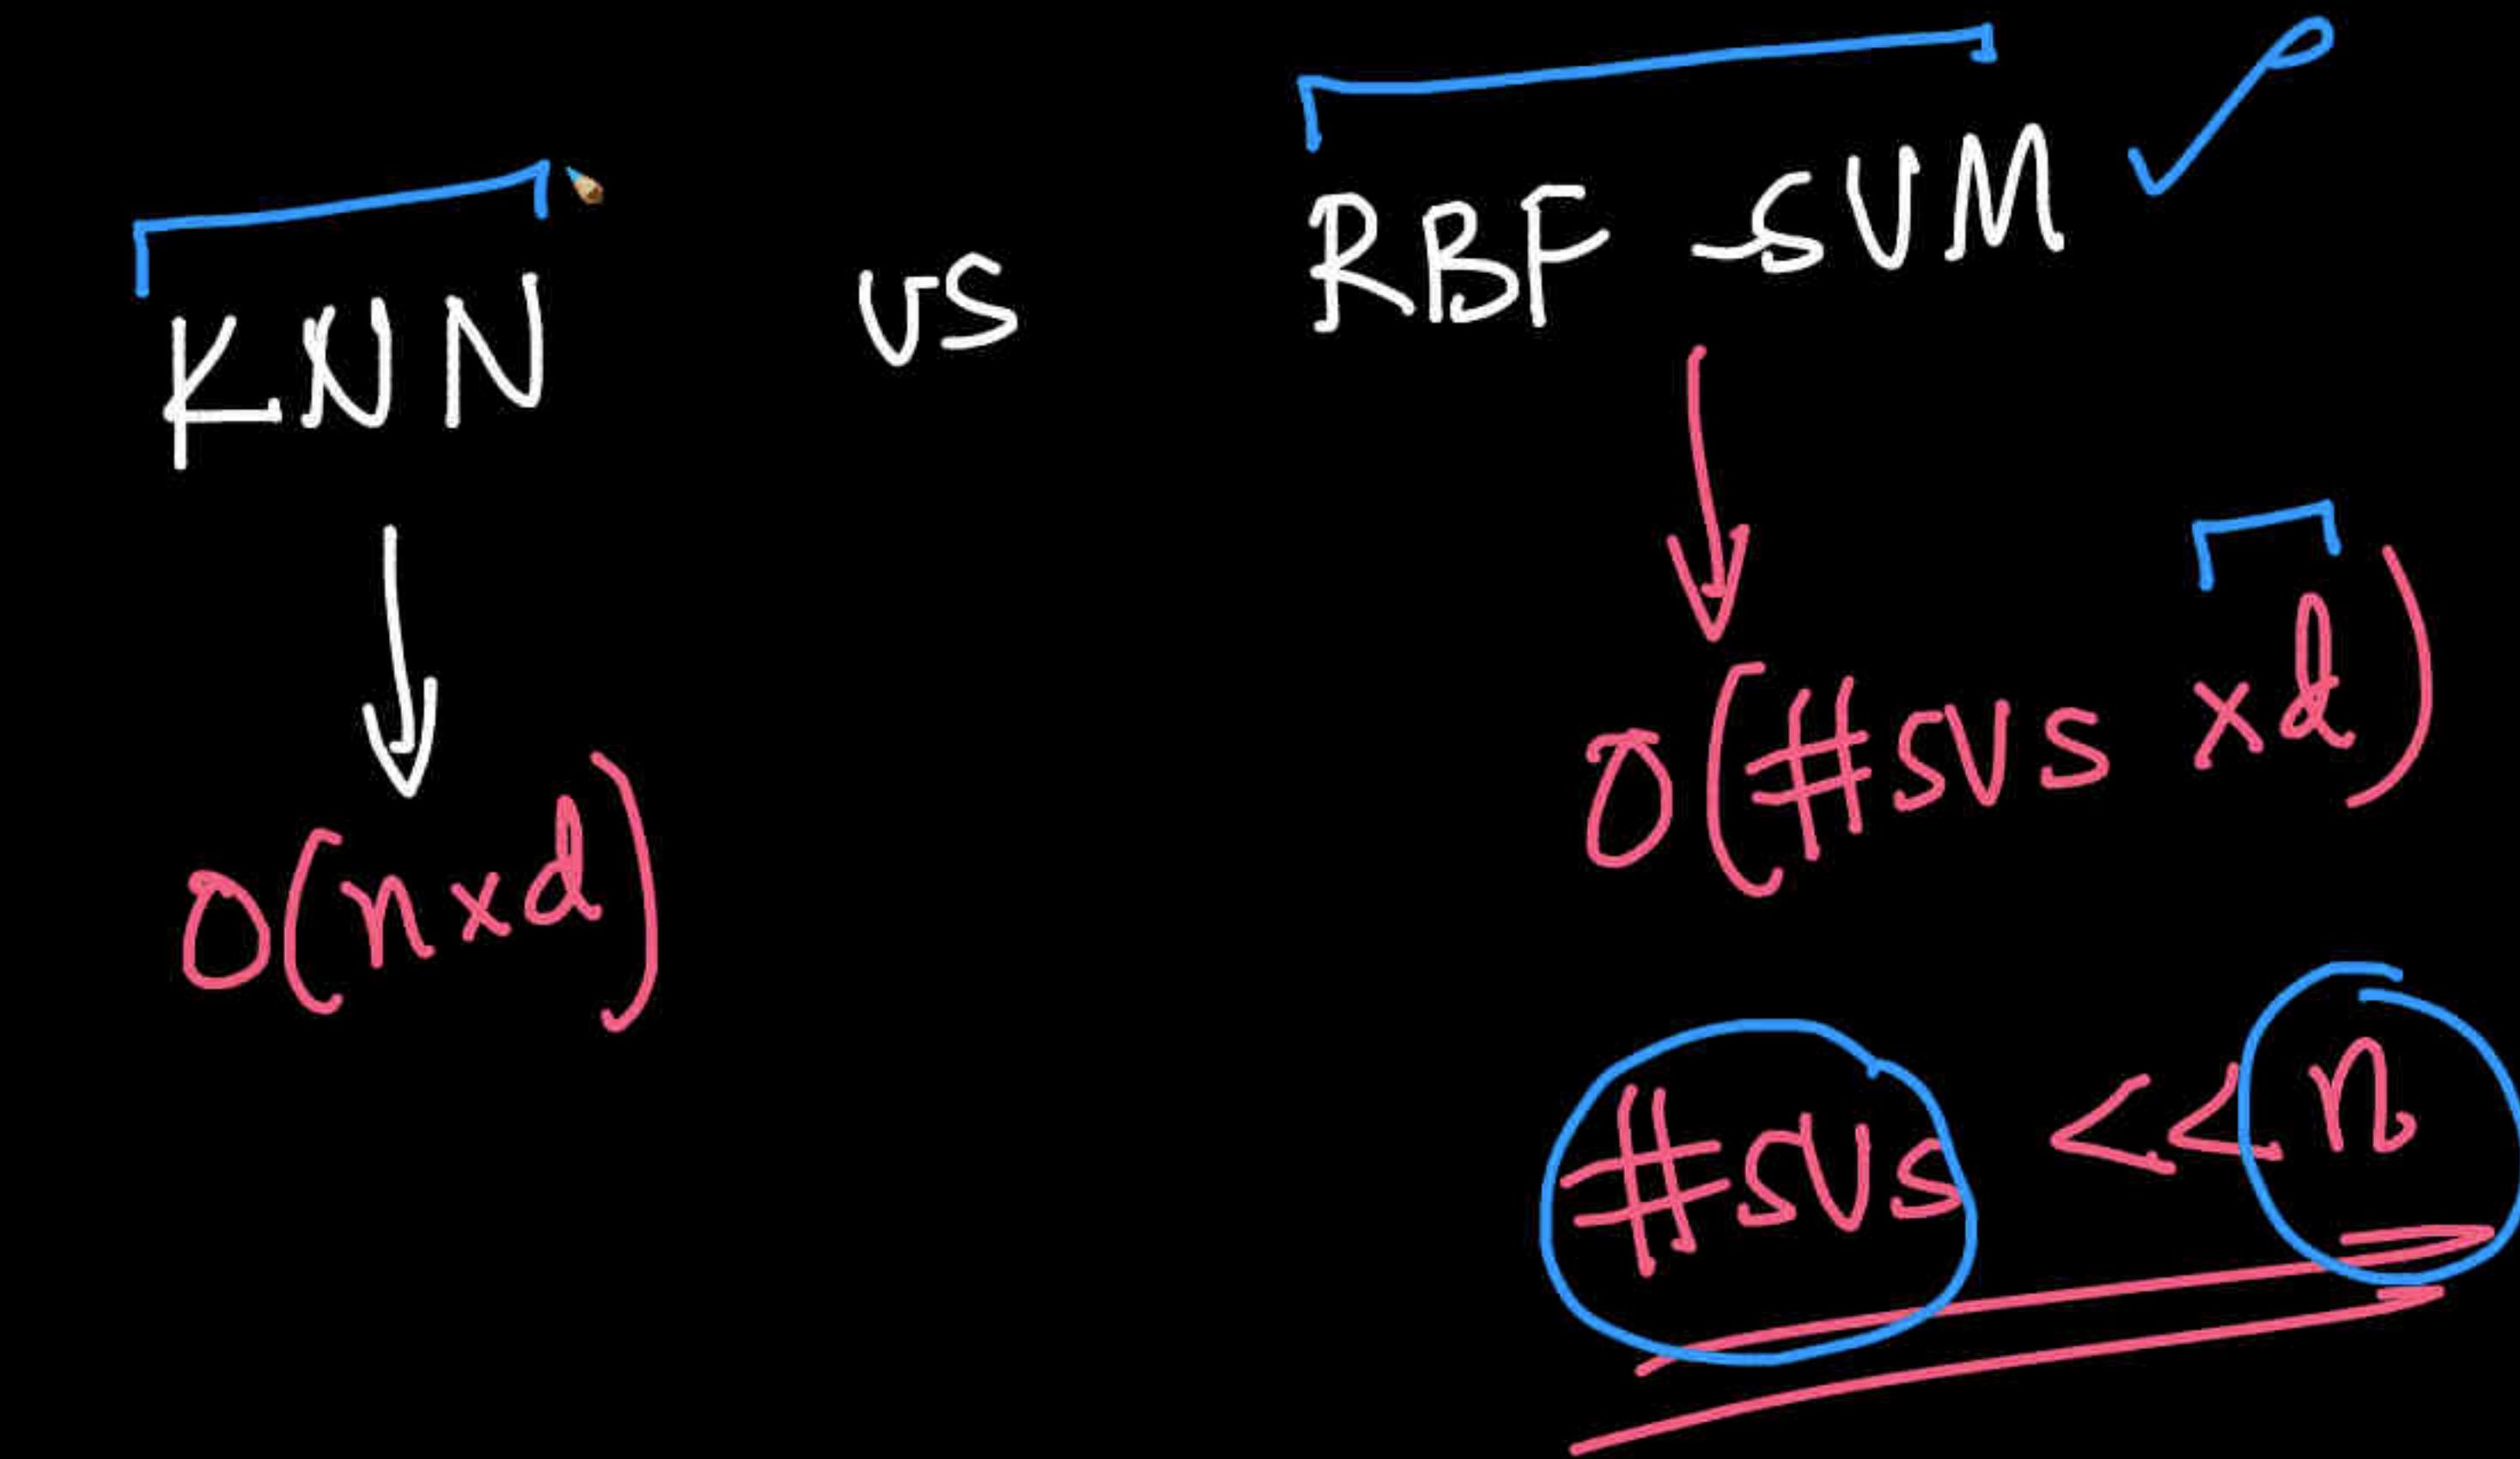
(Test) $\rightarrow O(\cancel{n}), O(\cancel{d})$

$$f(x_q) = \sum_{i=1}^n d_i y_i K(x_i, x_q) + b$$

$d_i = 0$ for non SVs
 $d_i > 0$ for SVs

Specializing
in \mathcal{O} -dim

Run-time
= $\mathcal{O}(\#SVs \times d)$



Summarize

LR-SVMs

i

Margin - Max Linear model
↳ regularization

hard-margin
soft-margin
almost lr.
sep

soft-margin LR-SVM
 \approx logistic reg

min hinge-loss

②

non-linear / kernel SUMS

- primal \rightarrow dual

$$K(x_i, x_j) \quad \text{Similarity}$$

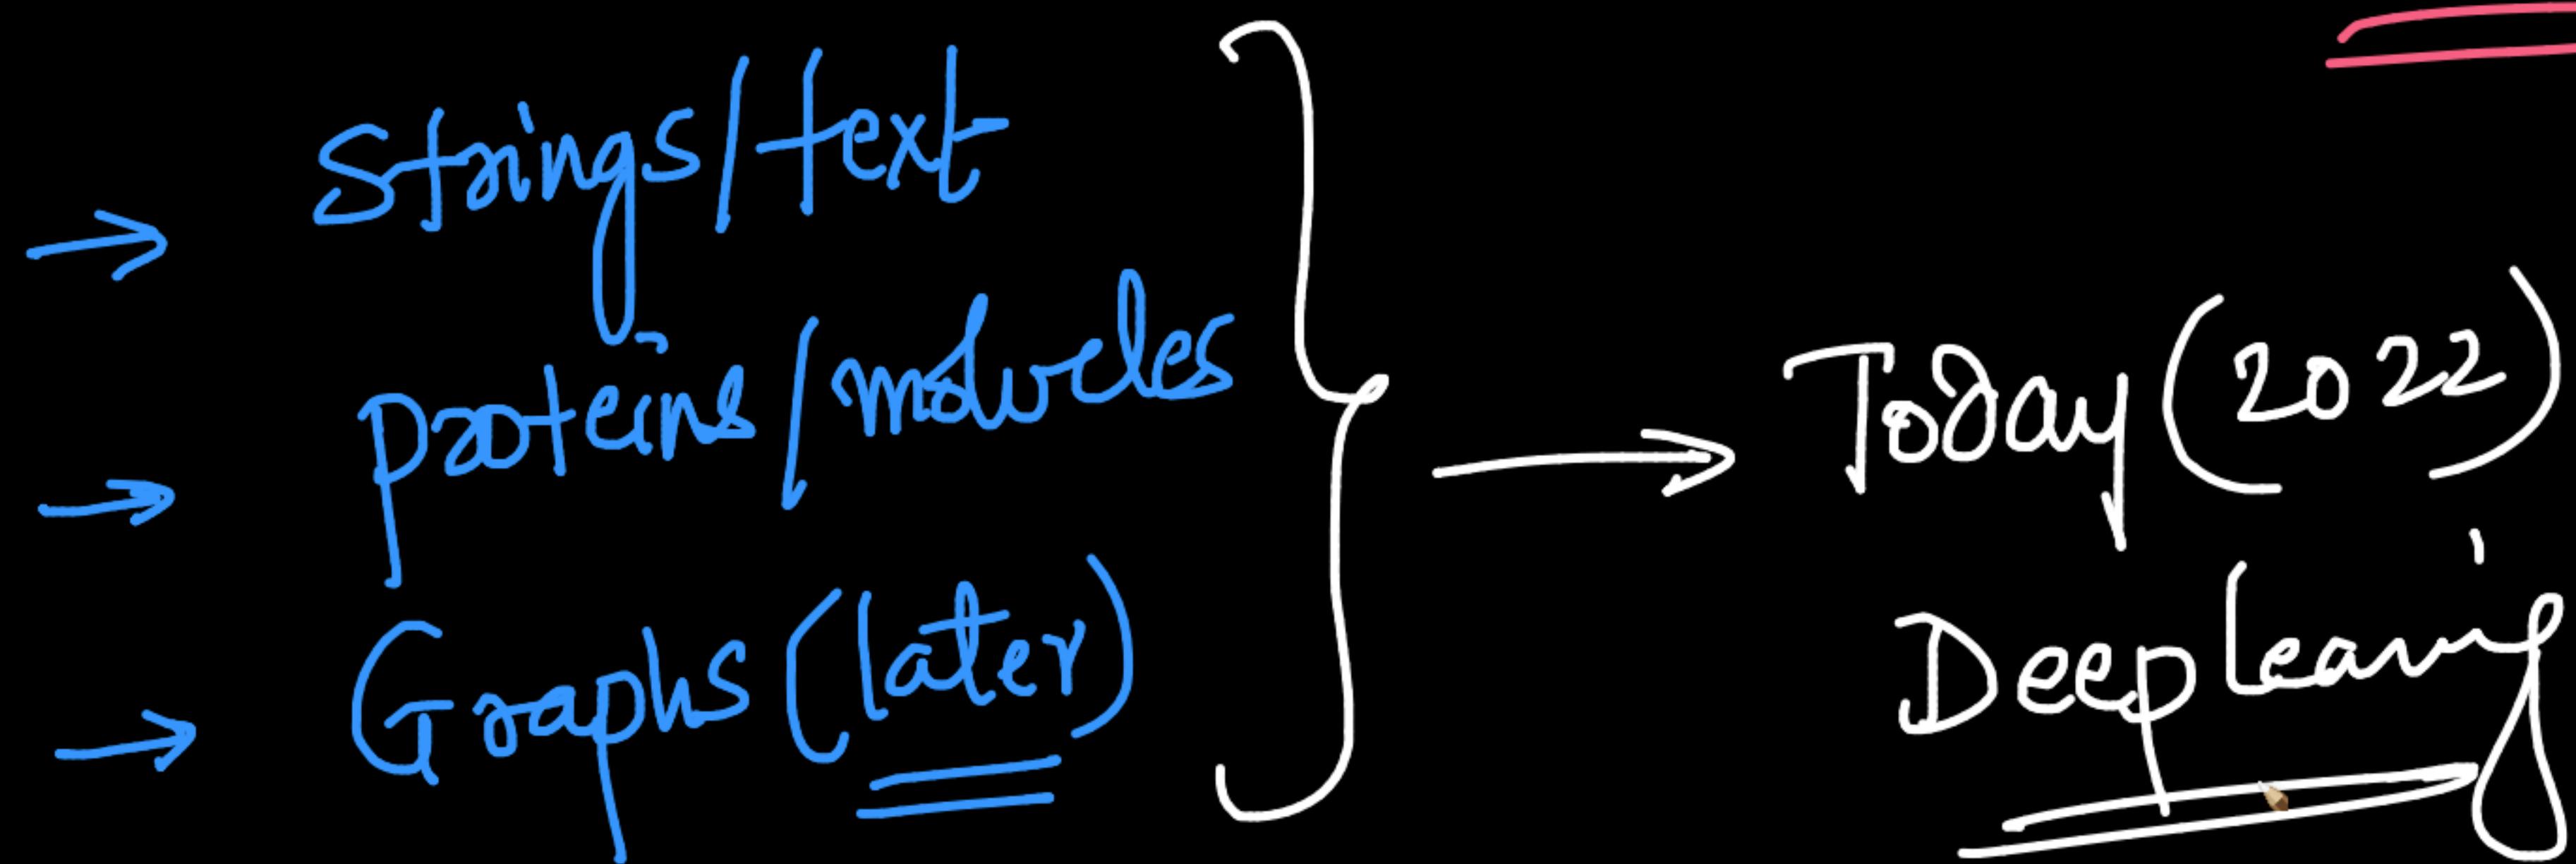
$$\underline{x_i^T x_j} : \text{Similarity}$$

- Quad. Kernel: $2d \rightarrow 6d$ (explicitly)

- RBF / Gaussian Kernel \rightarrow Geom! - $d \rightarrow d'$
 (costage)

$$\sum_{(r)} \text{RBF-SUM} \approx KNN(K)$$

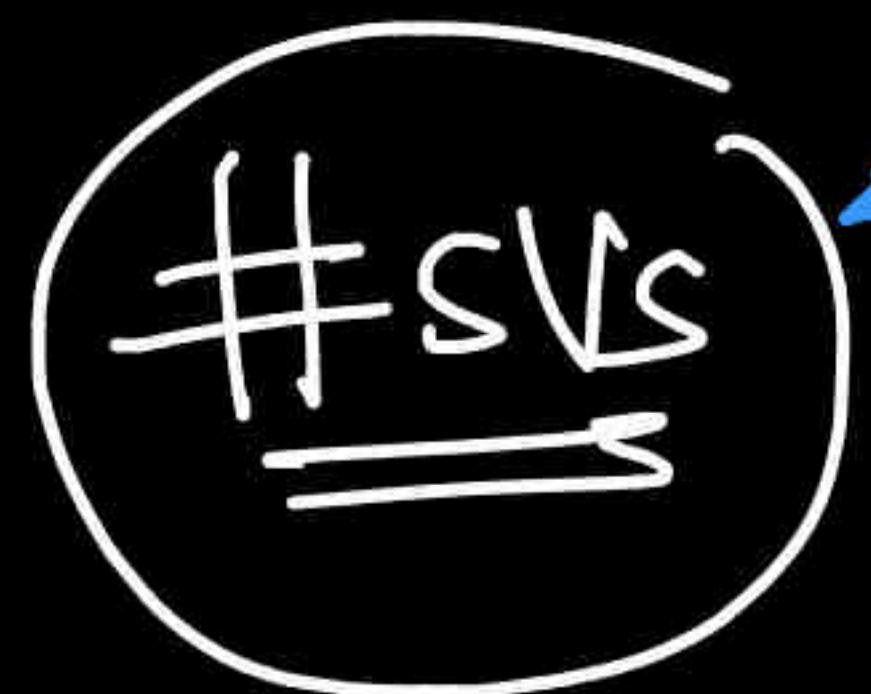
Domain Specific
Kernels



SVMs \rightarrow libSVM (C++) \rightarrow sklearn internally
 \hookrightarrow SGD ; alternative GD (SMO)

SVMs! Training time:- $\underline{\underline{O(n^2)}}$ (improved later...on)
 $n = LB$

Runtime Complex:-



Controlling it
is hard

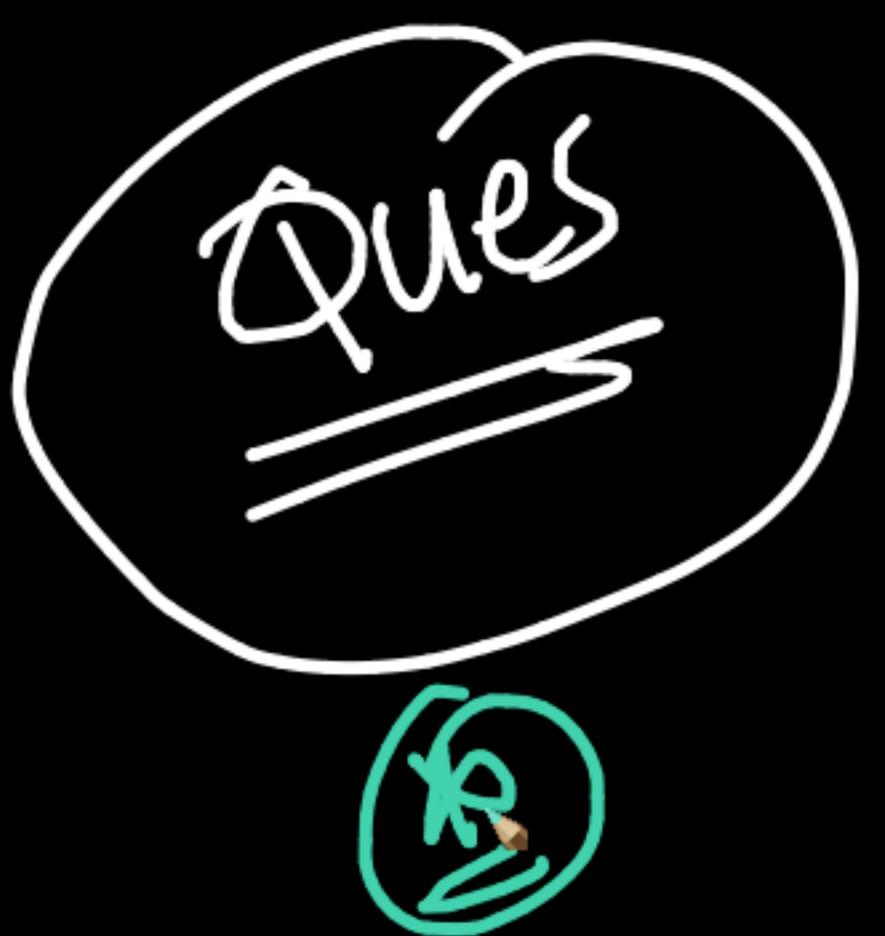
Concept:

intuition ?

Kernel-tick

$d \rightarrow d'$
(∞ -large)

SSVMS



$$K(x_i, x_j) \stackrel{?}{=} x_i^T x_j = \|x_i\| \|x_j\| \cos \theta$$

angular sim $\stackrel{?}{=}$

$$K_{\text{Quad}}(\underline{x}_1, \underline{x}_2) = (1 + \underline{x}_1^T \underline{x}_2)^2$$

✓ $K_{\text{RBF}}(\underline{x}_1, \underline{x}_2) = \exp\left\{-\frac{\|\tilde{\underline{x}}_1 - \tilde{\underline{x}}_2\|^2}{2\sigma^2}\right\}$

Simplest → formulation

hard-SVR

$\text{cl: } \xi_i > 0$

$$\min_{w, b} \frac{1}{2} \|w\|^2 + C \sum_i \xi_i$$

s.t.

$$y_i - (\bar{w}^T \bar{x}_i + \bar{b}) \leq \xi_i$$

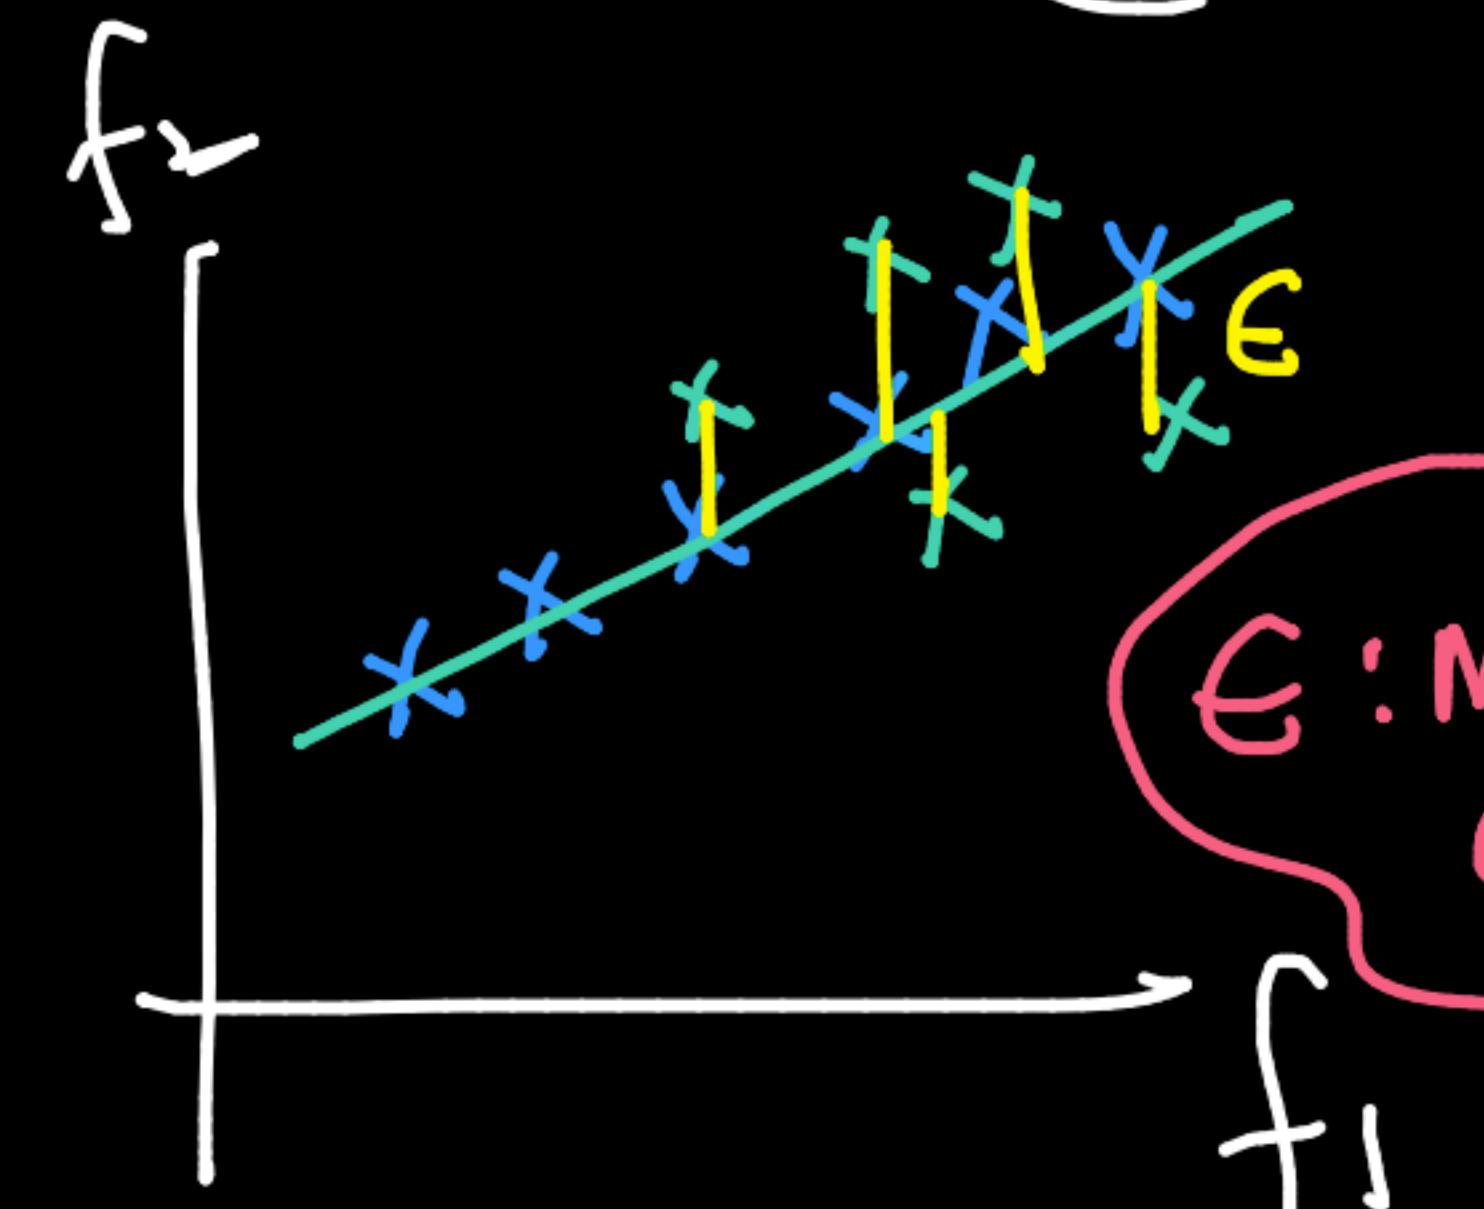
$$(\bar{w}^T \bar{x}_i + \bar{b}) - y_i \leq \xi_i$$

$$\xi_i \geq 0$$

not very popular

SV-R

(Kernelized → SVR)



ϵ : Max error

hyper-params:

RBF-SVM

① $\sigma \downarrow$

overfit

Primal:-

$$\min_{\omega, b} \frac{1}{2} \|\omega\|^2 + C \sum_{i=1}^n \xi_i$$

\uparrow reg
 \uparrow loss
 \uparrow C

② $C \uparrow$

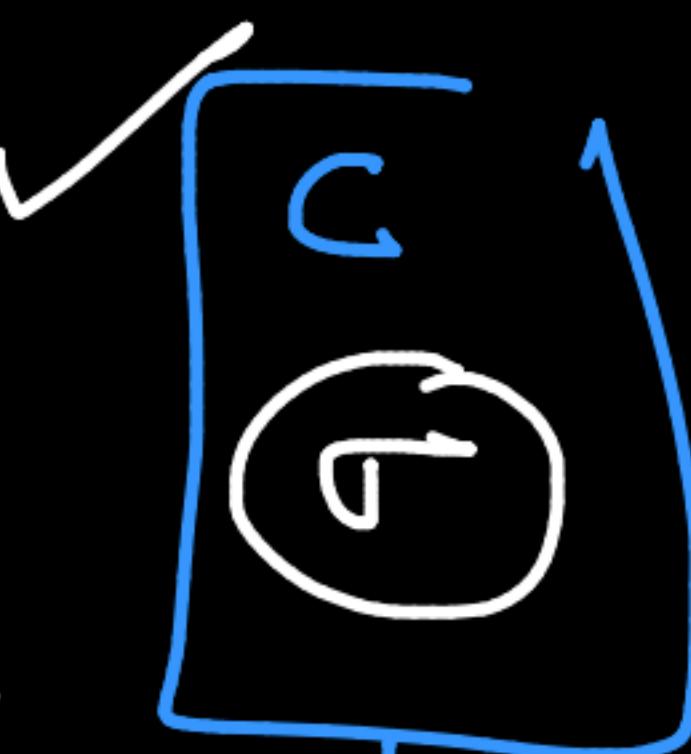
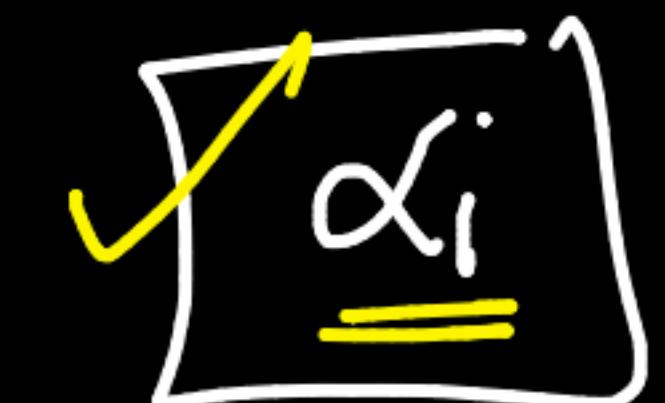
overfit

log-reg: \uparrow \uparrow underfit





Interpretability of SVMs



$$K_{RBF} =$$

No interpretation
drawbacks

hack

RBF-SVM \approx K-NN

$$x_{q1} = 1$$

nearest neighbors
nearest

plot(exp(-x^2/(2*1*1))) - GeoGebra x sklearn.svm.SVC — scikit-learn x +

scikit-learn.org/stable/modules/generated/sklearn.svm.SVC.html

scikit learn Install User Guide API Examples Community More ▾ Go

Prev Up Next

scikit-learn 1.1.1 Other versions

Please cite us if you use the software.

sklearn.svm.SVC

Examples using `sklearn.svm.SVC`

sklearn.svm.SVC

```
class sklearn.svm.SVC(*, C=1.0, kernel='rbf', degree=3, gamma='scale', coef0=0.0, shrinking=True, probability=False, tol=0.001, cache_size=200, class_weight=None, verbose=False, max_iter=-1, decision_function_shape='ovr', break_ties=False, random_state=None)
```

[source]

C-Support Vector Classification.

The implementation is based on libsvm. The fit time scales at least quadratically with the number of samples and may be impractical beyond tens of thousands of samples. For large datasets consider using `LinearSVC` or `SGDClassifier` instead, possibly after a `Nystroem` transformer.

The multiclass support is handled according to a one-vs-one scheme.

For details on the precise mathematical formulation of the provided kernel functions and how `gamma`, `coef0` and `degree` affect each other, see the corresponding section in the narrative documentation: [Kernel functions](#).

Read more in the [User Guide](#).

Toggle Menu

Page 40 / 40

Google plot(exp(-x^2/(2*1*1))) - Geogebra x sklearn.svm.SVC — scikit-learn x +

scikit-learn.org/stable/modules/generated/sklearn.svm.SVC.html

scikit learn

Prev Up Next

scikit-learn 1.1.1

Other versions

Please cite us if you use the software.

sklearn.svm.SVC

Examples using sklearn.svm.SVC

$\gamma = \frac{1}{\sigma^2}$

kernel : {‘linear’, ‘poly’, ‘rbf’, ‘sigmoid’, ‘precomputed’} or callable, default=‘rbf’

Specifies the kernel type to be used in the algorithm. If none is given, ‘rbf’ will be used. If a callable is given it is used to pre-compute the kernel matrix from data matrices; that matrix should be an array of shape (n_samples, n_samples).

degree : int, default=3

Degree of the polynomial kernel function (‘poly’). Ignored by all other kernels.

gamma : {‘scale’, ‘auto’} or float, default=‘scale’

Kernel coefficient for ‘rbf’, ‘poly’ and ‘sigmoid’.

- if gamma='scale' (default) is passed then it uses $1 / (\text{n_features} * \text{X.var()})$ as value of gamma,
- if ‘auto’, uses $1 / \text{n_features}$.

Changed in version 0.22: The default value of gamma changed from ‘auto’ to ‘scale’!

coef0 : float, default=0.0

Independent term in kernel function. It is only significant in ‘poly’ and ‘sigmoid’.

shrinking : bool, default=True

Whether to use the shrinking heuristic. See the [User Guide](#).

Toggle Menu

41/41

plot(exp(-x^2/(2*1*1))) - GeoGebra x sklearn.svm.SVC — scikit-learn x +

scikit-learn.org/stable/modules/generated/sklearn.svm.SVC.html

scikit learn Install User Guide API Examples Community More ▾ Go

Prev Up Next

scikit-learn 1.1.1 Other versions

Please cite us if you use the software.

sklearn.svm.SVC

Examples using `sklearn.svm.SVC`

sklearn.svm.SVC

$w^T x_i + b$

```
class sklearn.svm.SVC(*, C=1.0, kernel='rbf', degree=3, gamma='scale', coef0=0.0, shrinking=True, probability=False, tol=0.001, cache_size=200, class_weight=None, verbose=False, max_iter=-1, decision_function_shape='ovr', break_ties=False, random_state=None)
```

[source]

C-Support Vector Classification.

The implementation is based on libsvm. The fit time scales at least quadratically with the number of samples and may be impractical beyond tens of thousands of samples. For large datasets consider using `LinearSVC` or `SGDClassifier` instead, possibly after a `Nystroem` transformer.

The multiclass support is handled according to a one-vs-one scheme.

For details on the precise mathematical formulation of the provided kernel functions and how `gamma`, `coef0` and `degree` affect each other, see the corresponding section in the narrative documentation: [Kernel functions](#).

Read more in the [User Guide](#).

Toggle Menu

Page 42 / 42

regularization is inversely proportional

plot(exp(-x^2/(2*1*1))) - GeoGebra x sklearn.svm.SVC — scikit-learn x +

scikit-learn.org/stable/modules/generated/sklearn.svm.SVC.html

scikit learn Install User Guide API Examples Community More ▾ Go

Prev Up Next

scikit-learn 1.1.1 Other versions

Please cite us if you use the software.

[sklearn.svm.SVC](#)

Examples using `sklearn.svm.SVC`

sklearn.svm.SVC

```
class sklearn.svm.SVC(*, C=1.0, kernel='rbf', degree=3, gamma='scale', coef0=0.0, shrinking=True, probability=False, tol=0.001, cache_size=200, class_weight=None, verbose=False, max_iter=-1, decision_function_shape='ovr', break_ties=False, random_state=None)
```

[source]

C-Support Vector Classification.

The implementation is based on libsvm. The fit time scales at least quadratically with the number of samples and may be impractical beyond tens of thousands of samples. For large datasets consider using [LinearSVC](#) or [SGDClassifier](#) instead, possibly after a [Nystroem](#) transformer.

The multiclass support is handled according to a one-vs-one scheme.

For details on the precise mathematical formulation of the provided kernel functions and how `gamma`, `coef0` and `degree` affect each other, see the corresponding section in the narrative documentation: [Kernel functions](#).

Read more in the [User Guide](#).

Toggle Menu

Page 43 / 43

plot(exp(-x^2/(2*1*1))) - GeoGebra x sklearn.svm.SVC — scikit-learn x +

scikit-learn.org/stable/modules/generated/sklearn.svm.SVC.html

scikit learn Install User Guide API Examples Community More ▾

Prev Up Next

scikit-learn 1.1.1 Other versions

Please cite us if you use the software.

sklearn.svm.SVC

Examples using `sklearn.svm.SVC`

sklearn.svm.SVC

class `sklearn.svm.SVC(*, C=1.0, kernel='rbf', degree=3, gamma='scale', coef0=0.0, shrinking=True, probability=False, tol=0.001, cache_size=200, class_weight=None, verbose=False, max_iter=-1, decision_function_shape='ovr', break_ties=False, random_state=None)` [source]

C-Support Vector Classification.

The implementation is based on `libsvm`. The fit time scales at least quadratically with the number of samples and may be impractical beyond tens of thousands of samples. For large datasets consider using `LinearSVC` or `SGDClassifier` instead, possibly after a `Nystroem` transformer.

The multiclass support is handled according to a one-vs-one scheme.

For details on the precise mathematical formulation of the provided kernel functions and how `gamma`, `coef0` and `degree` affect each other, see the corresponding section in the narrative documentation: [Kernel functions](#).

Read more in the [User Guide](#).

Toggle Menu

Page 1/6 Go

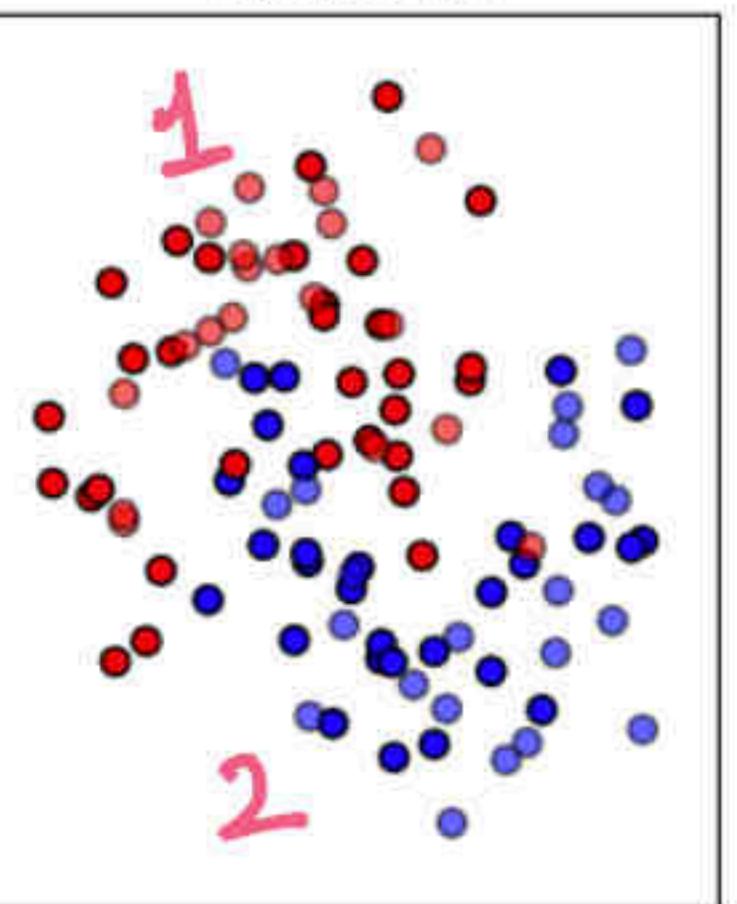
regularization is inversely proportional

Google plot(exp(-x^2/(2*1*1))) - Geogebra

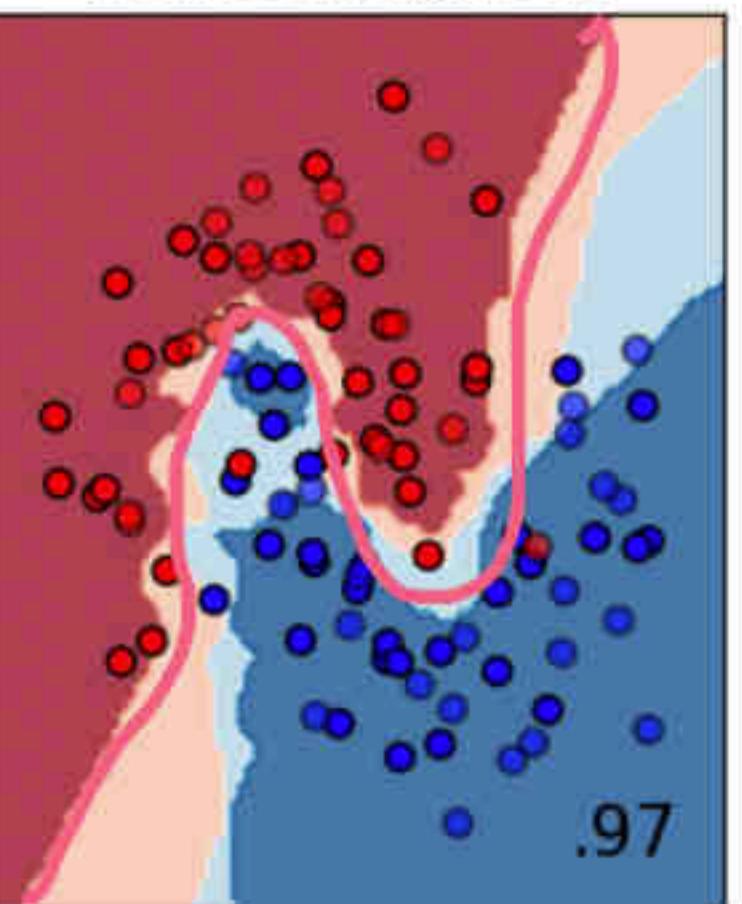
Classifier comparison — scikit-learn 0.20.2 documentation

sphx_glr_plot_classifier_comparison_001.png

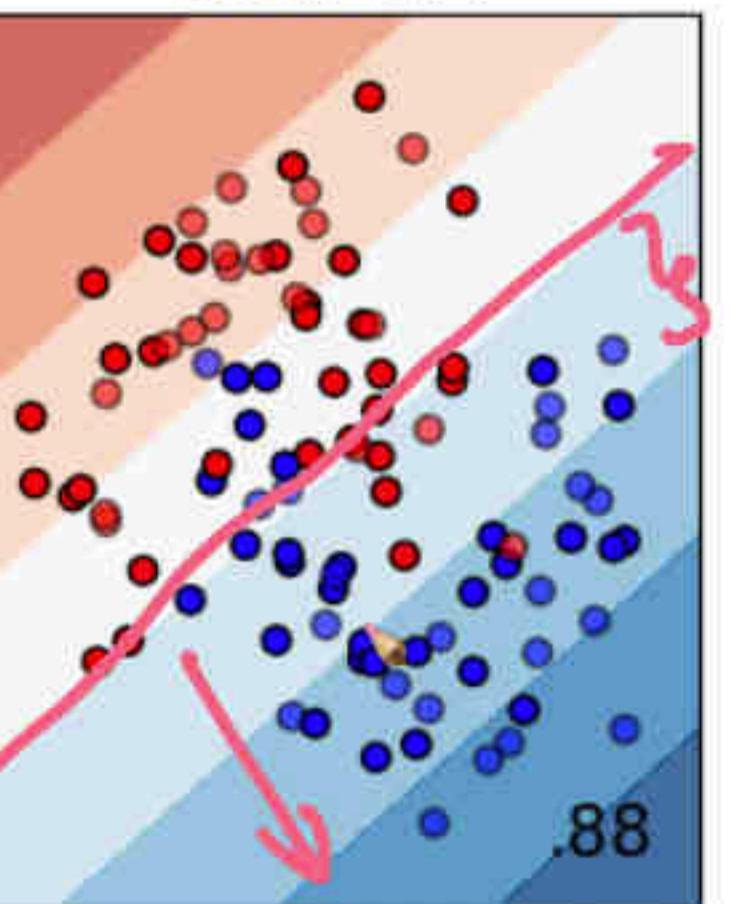
Input data



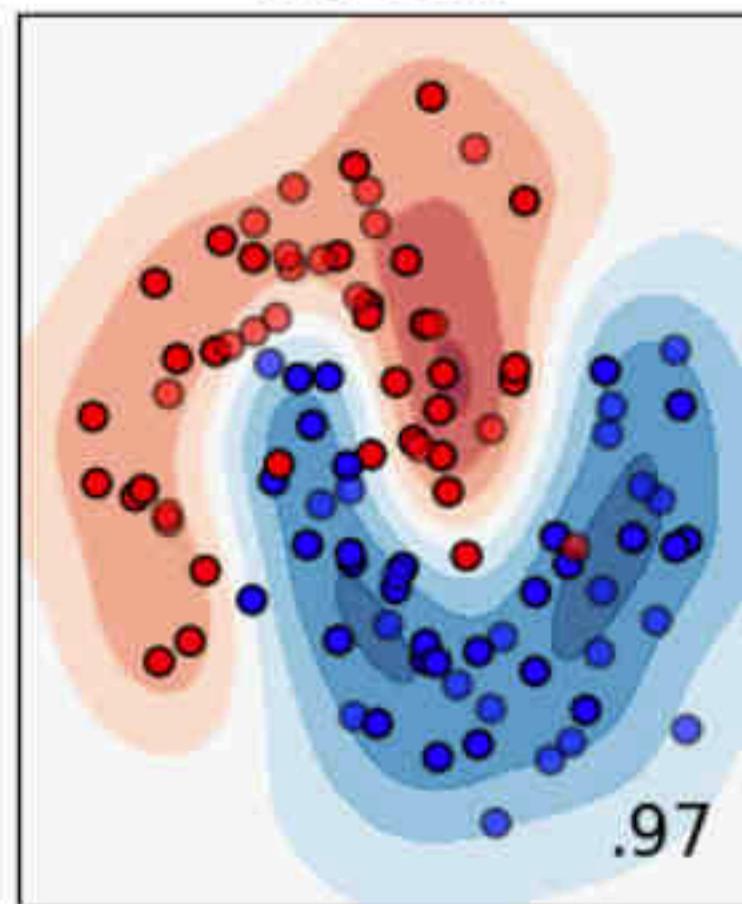
Nearest Neighbors



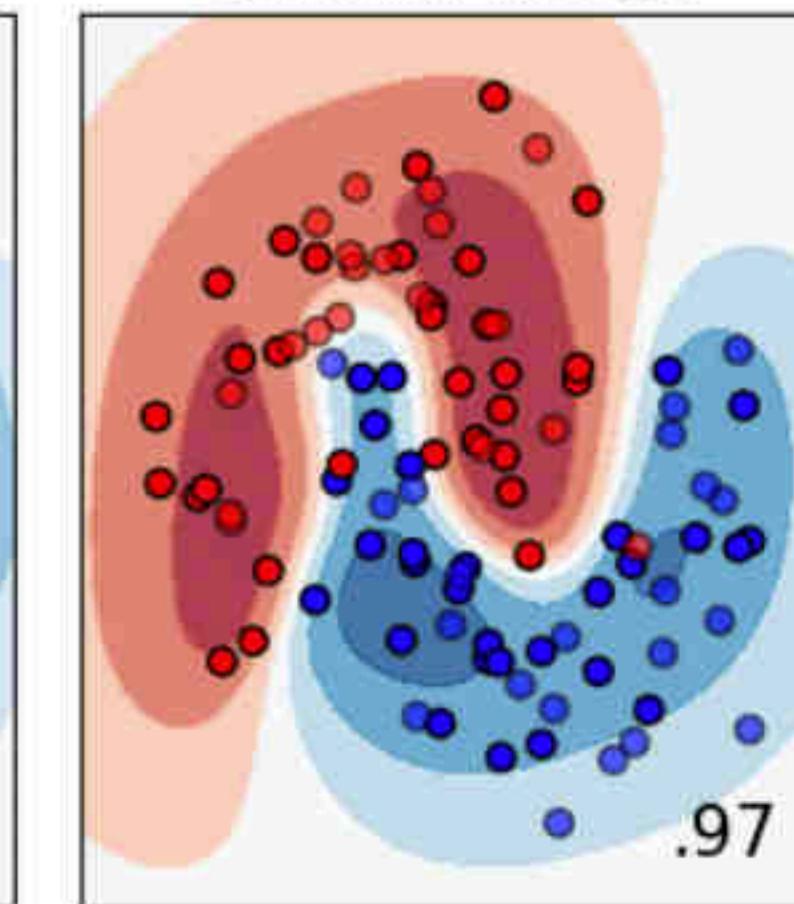
Linear SVM



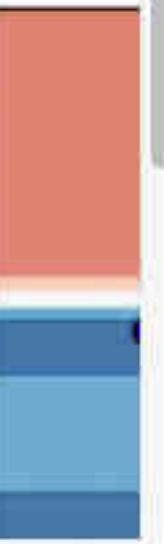
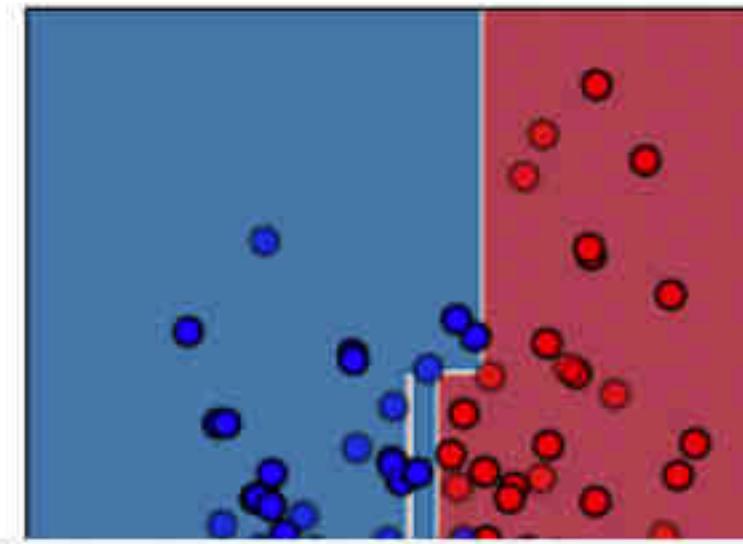
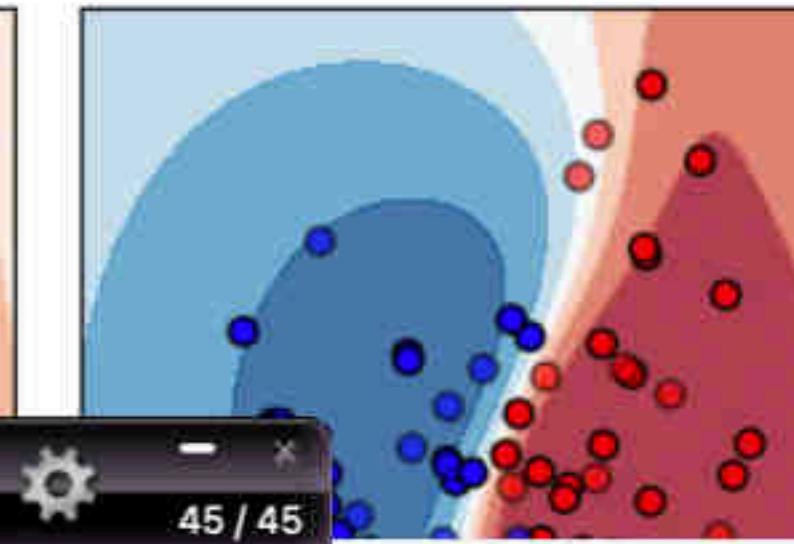
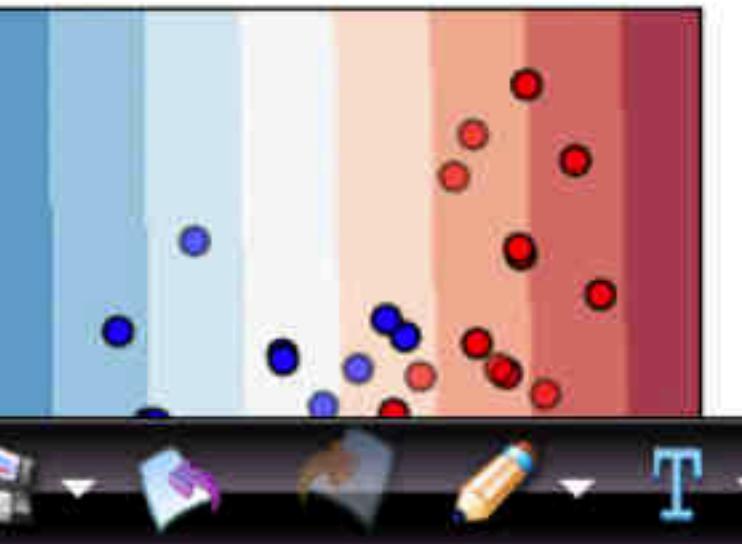
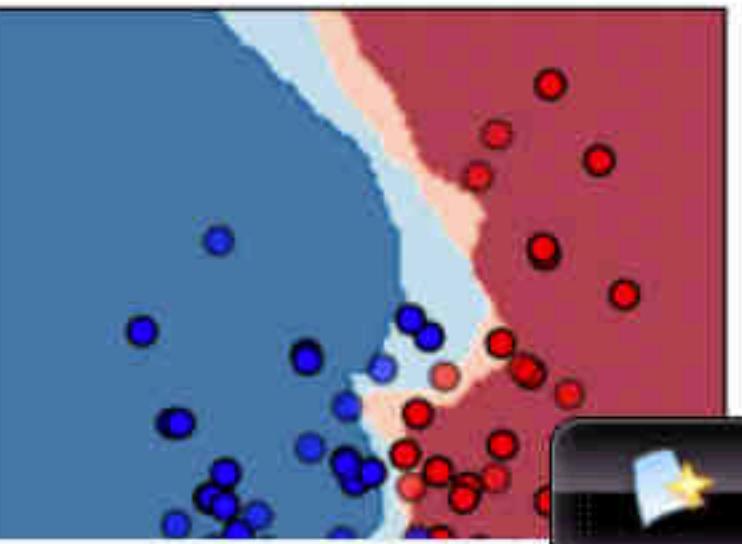
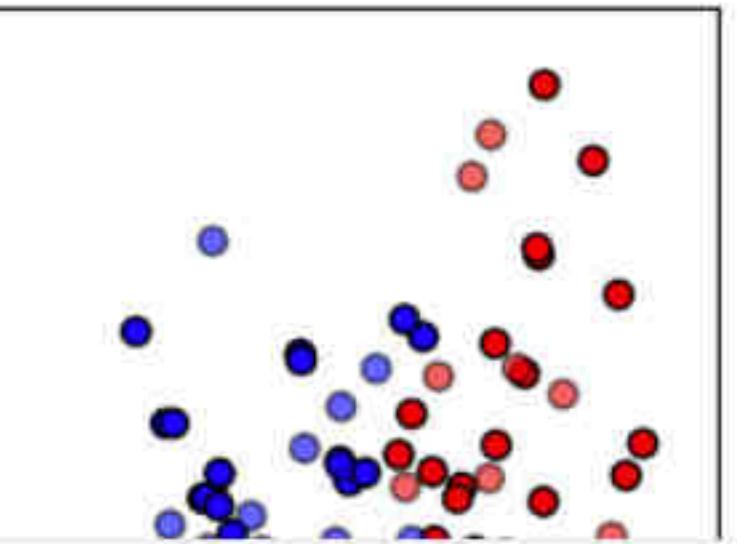
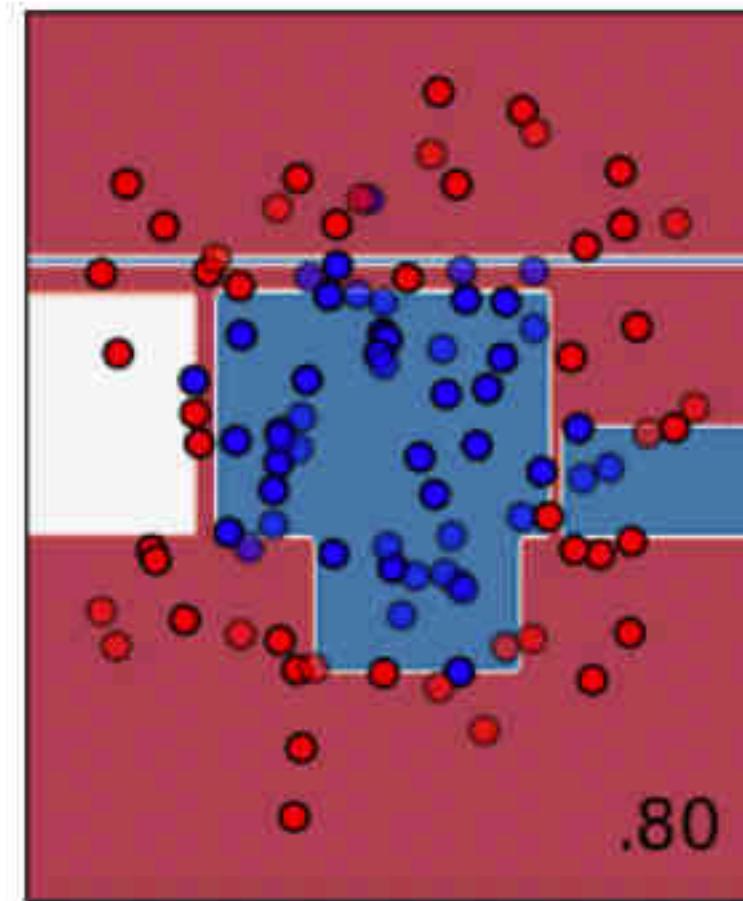
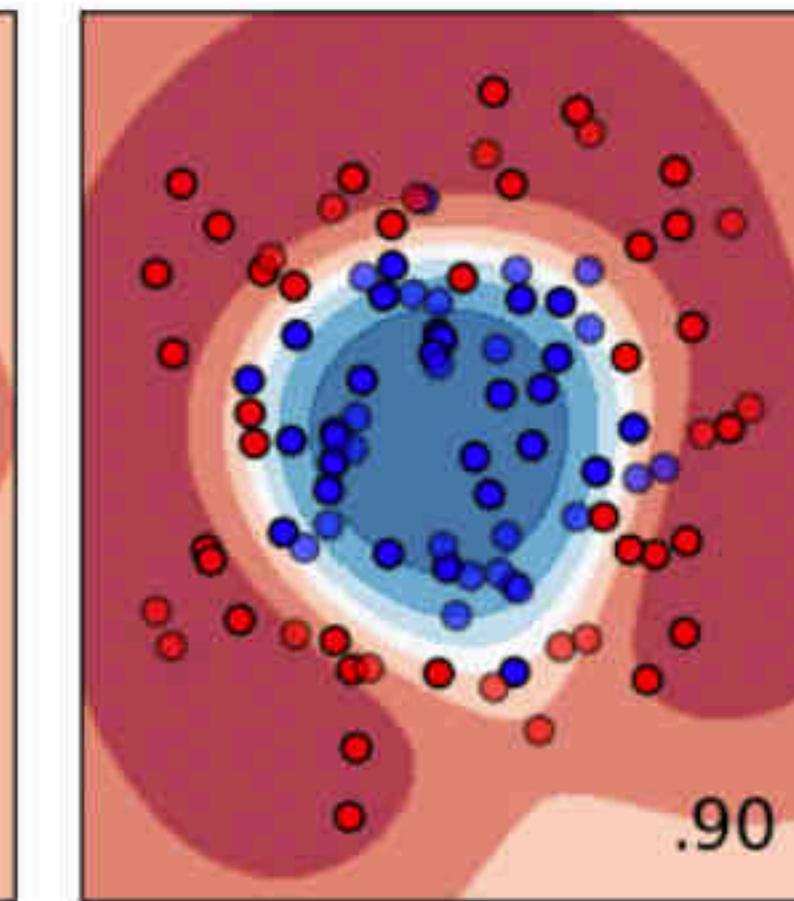
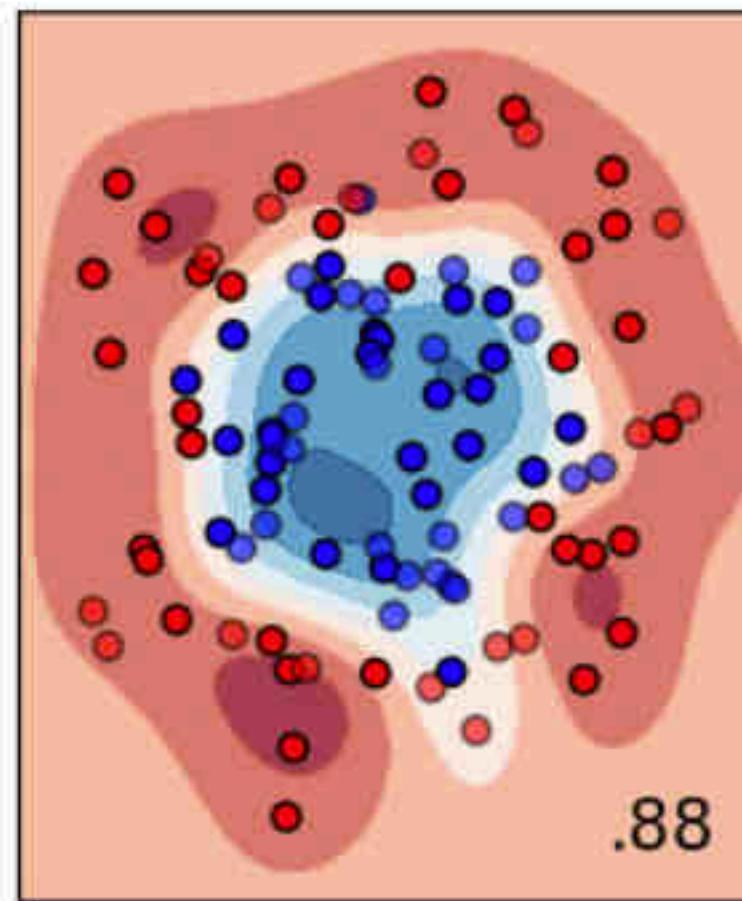
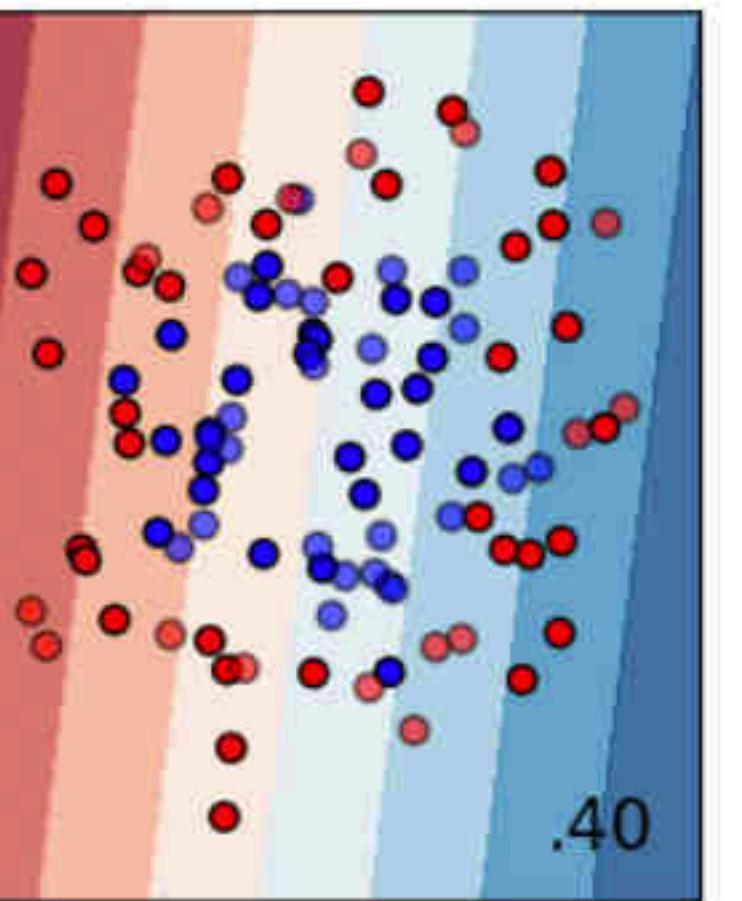
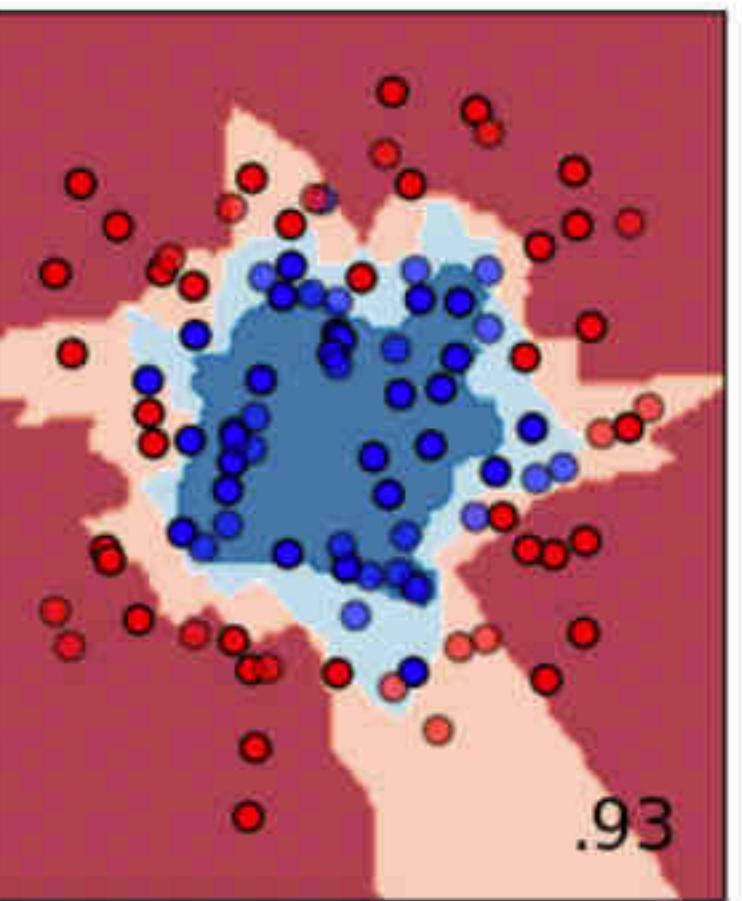
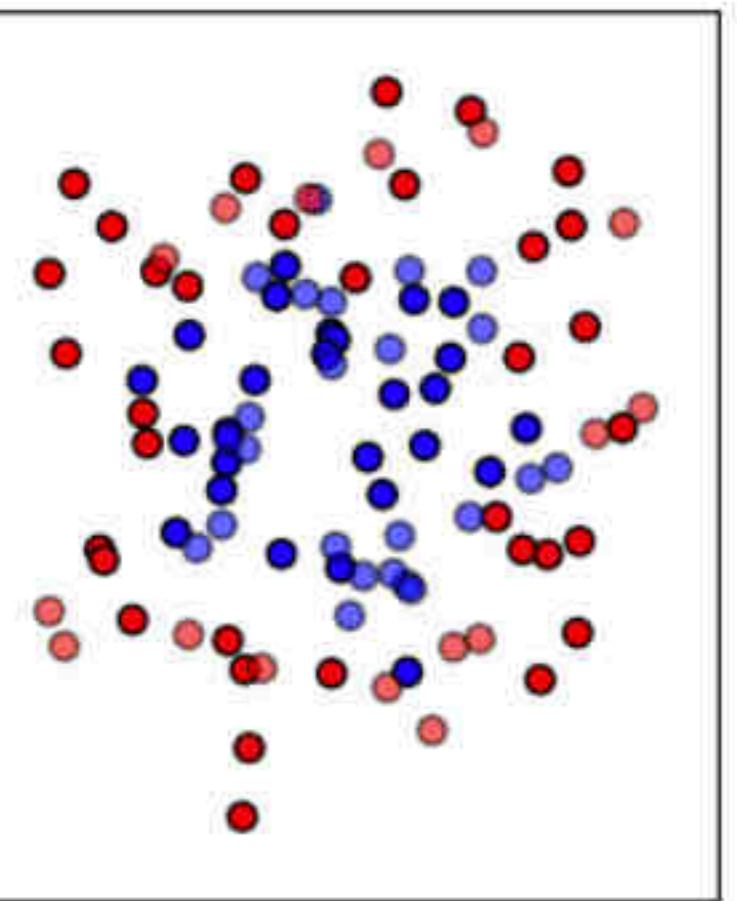
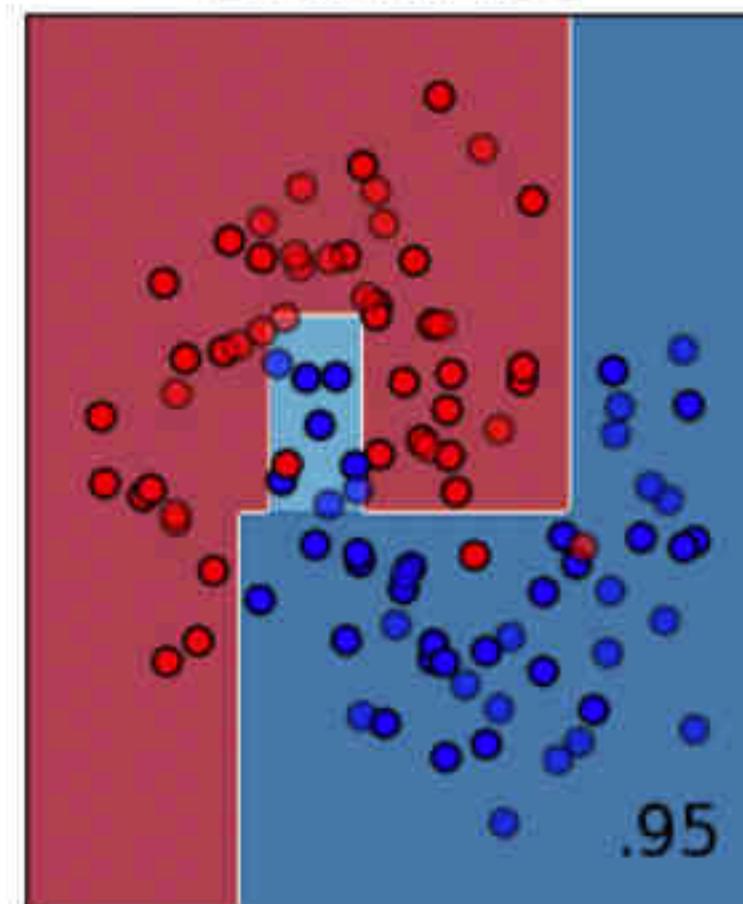
RBF SVM



Gaussian Process



Decision Tree

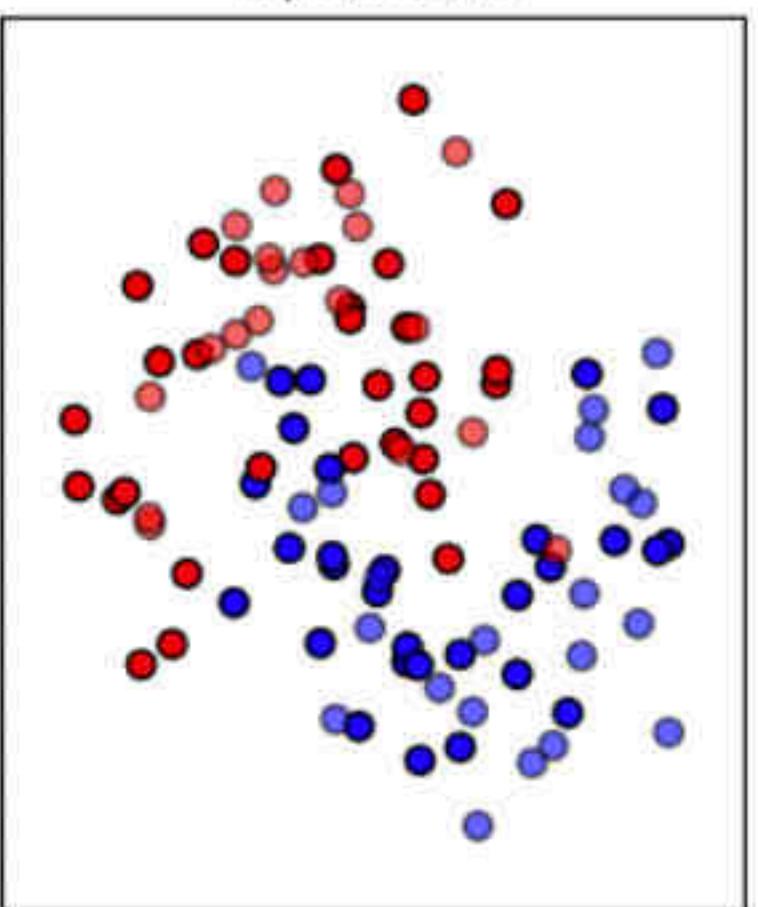


Google plot(exp(-x^2/(2*1*1))) - Geogebra

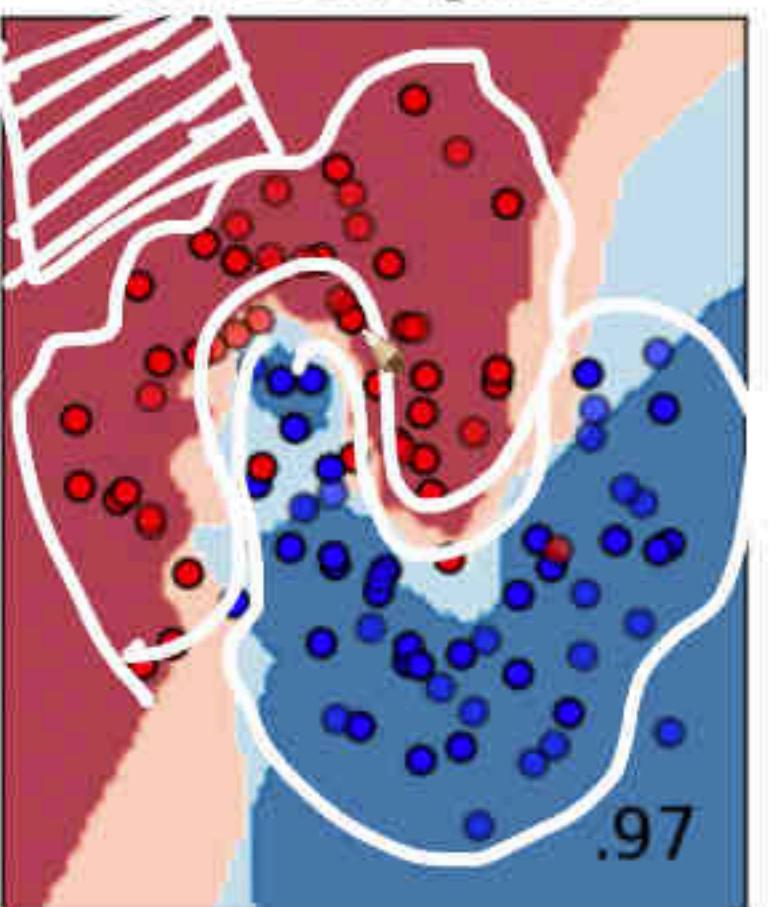
Classifier comparison — scikit-learn 0.23.2 documentation

sphx_glr_plot_classifier_comparison_001.png

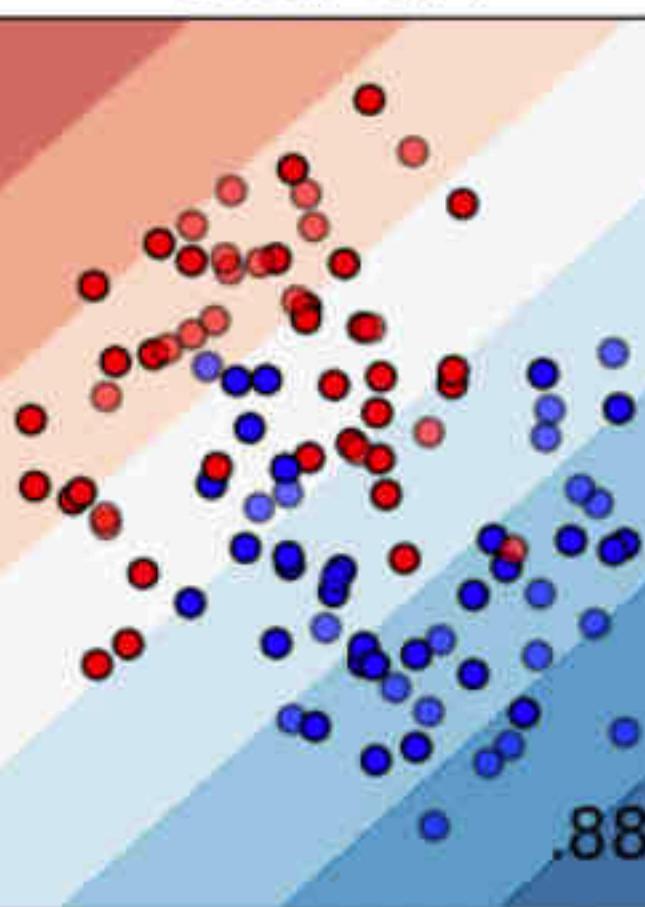
Input data



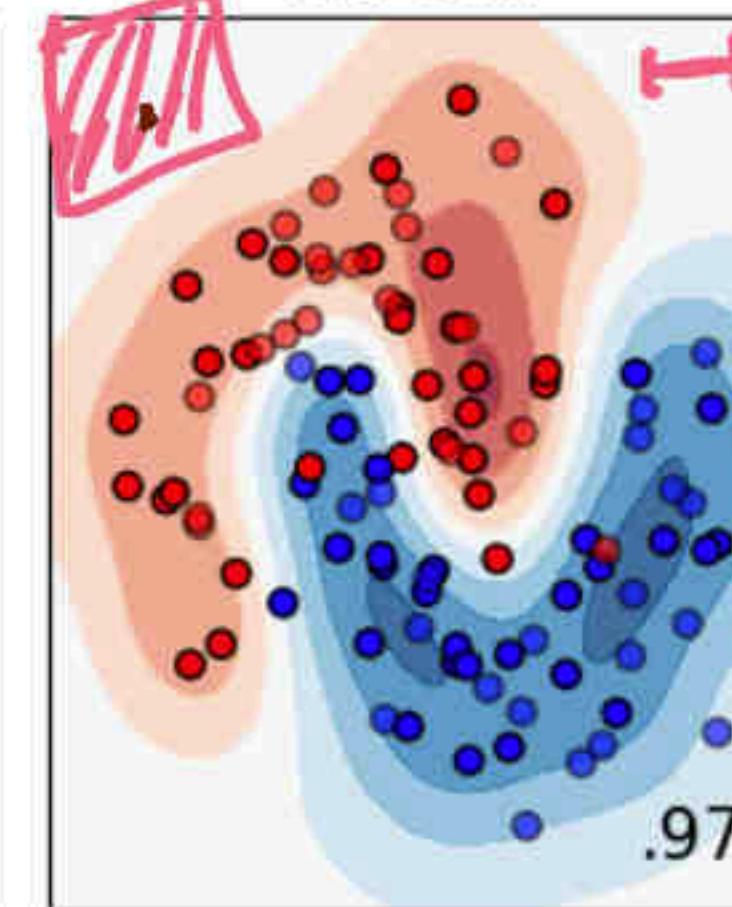
Nearest Neighbors



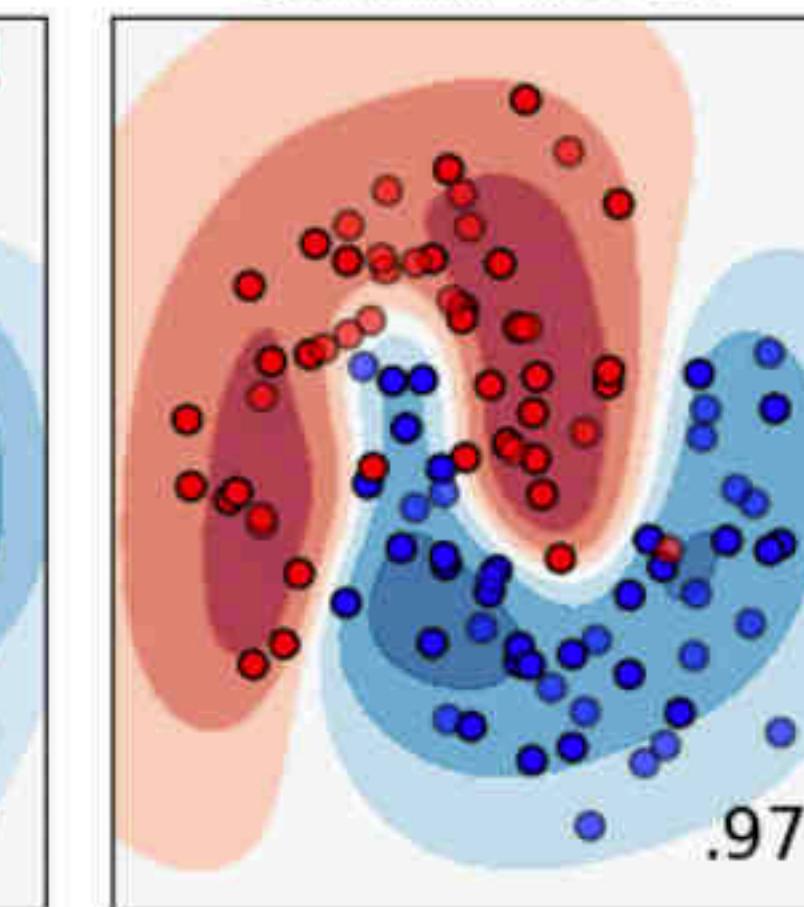
Linear SVM



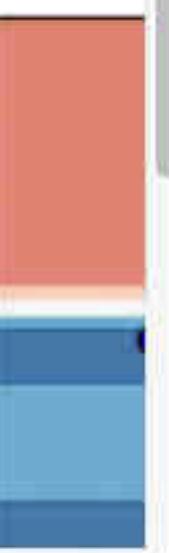
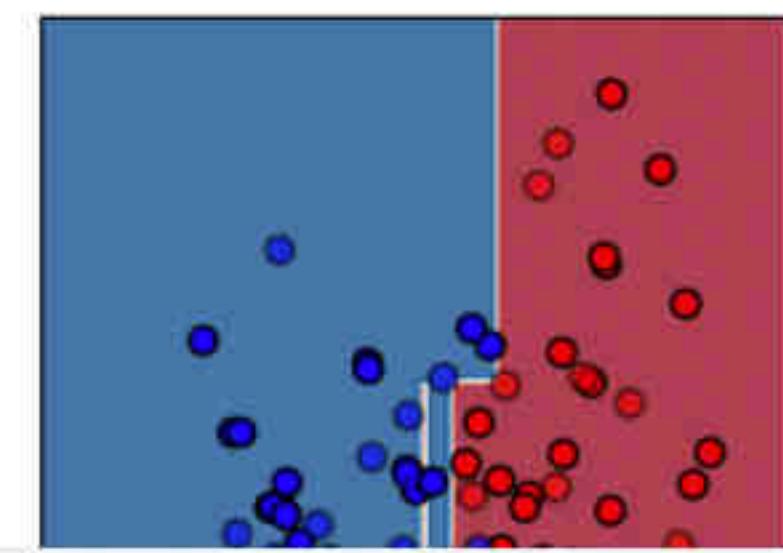
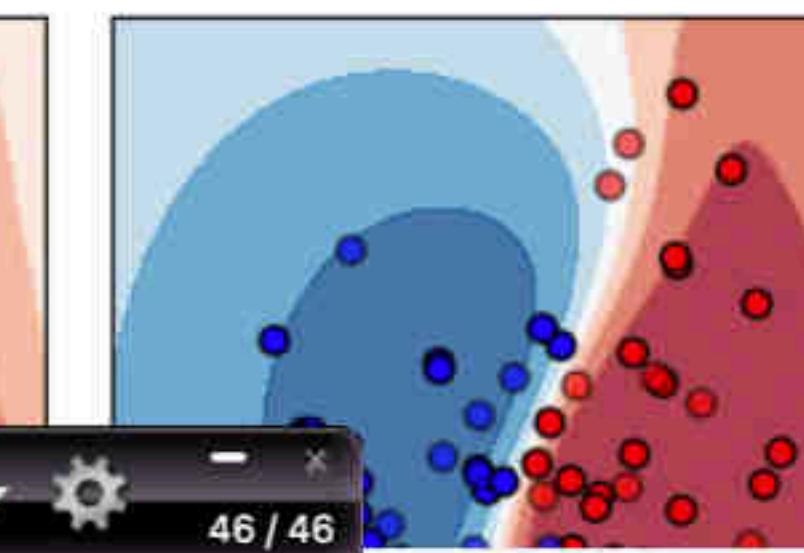
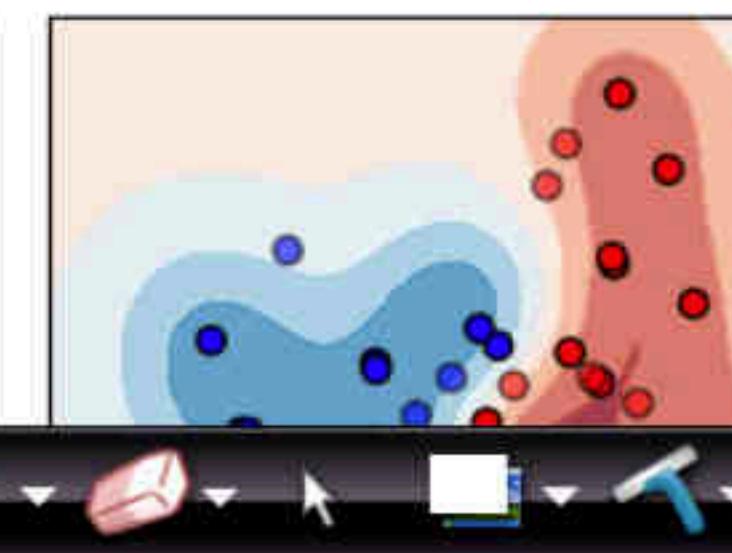
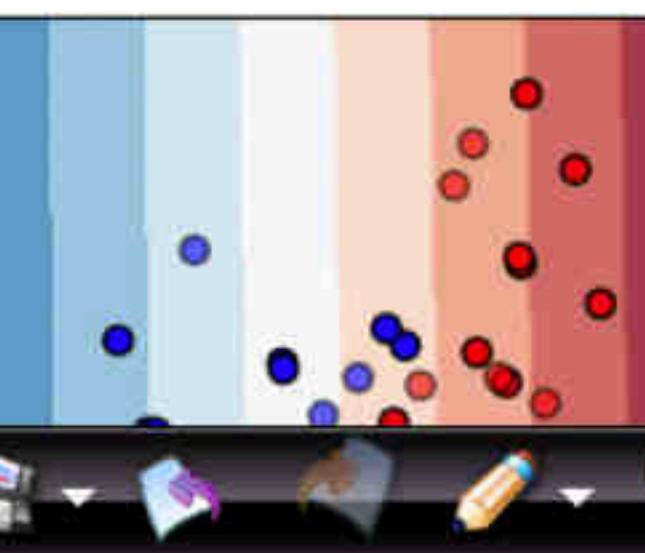
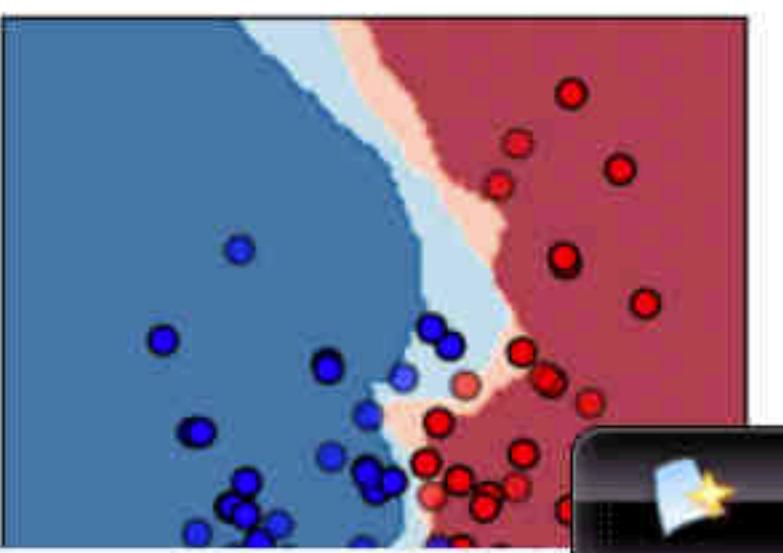
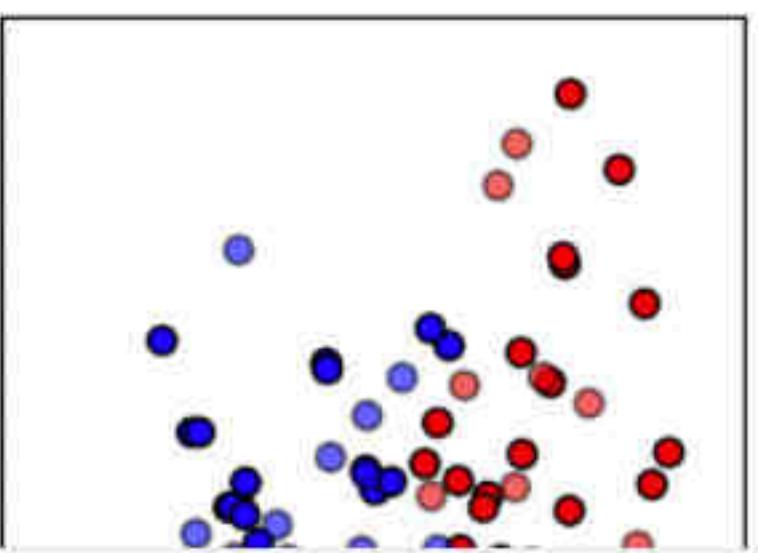
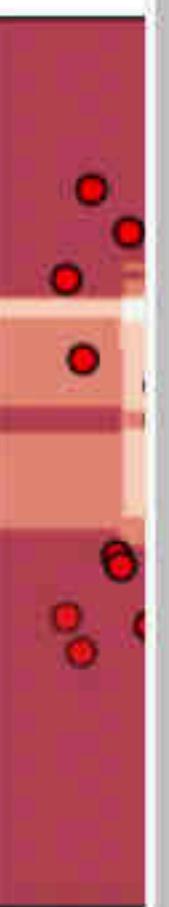
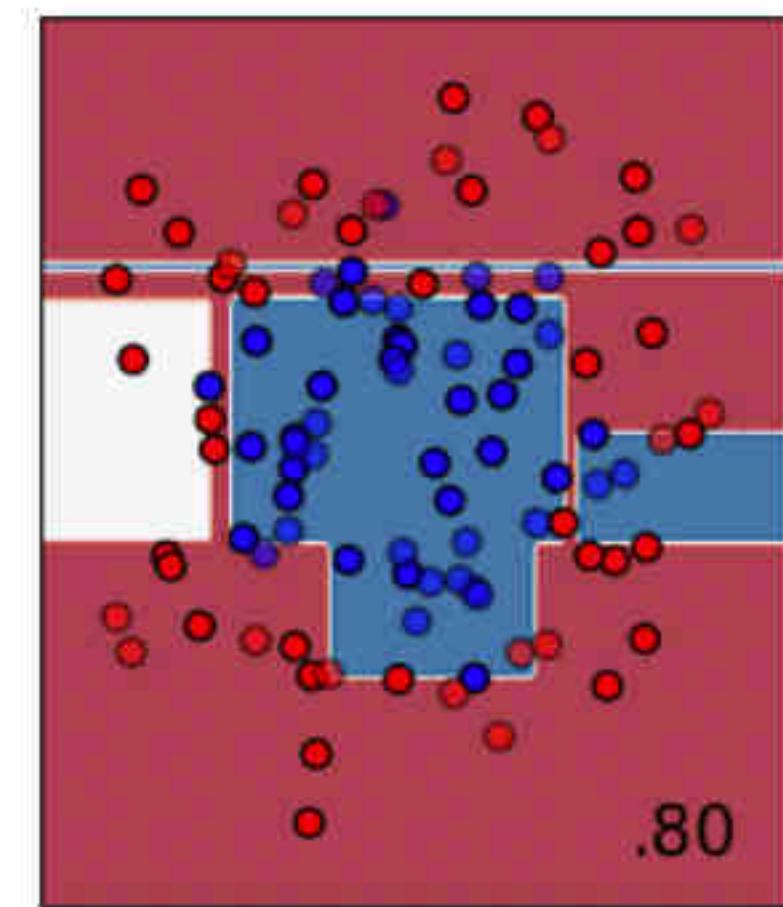
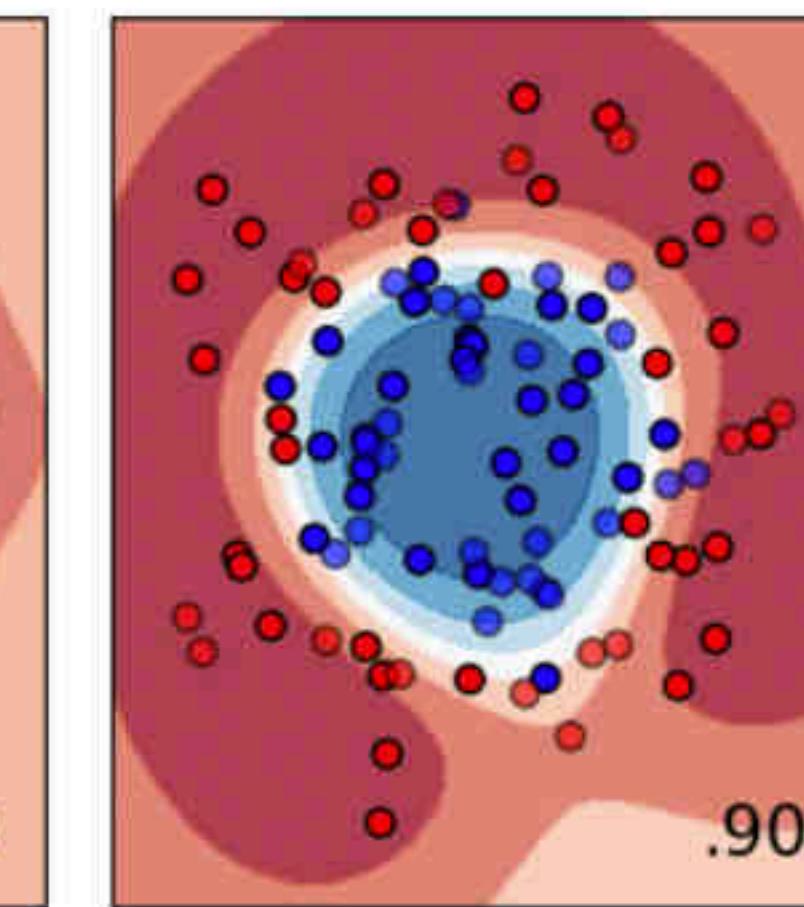
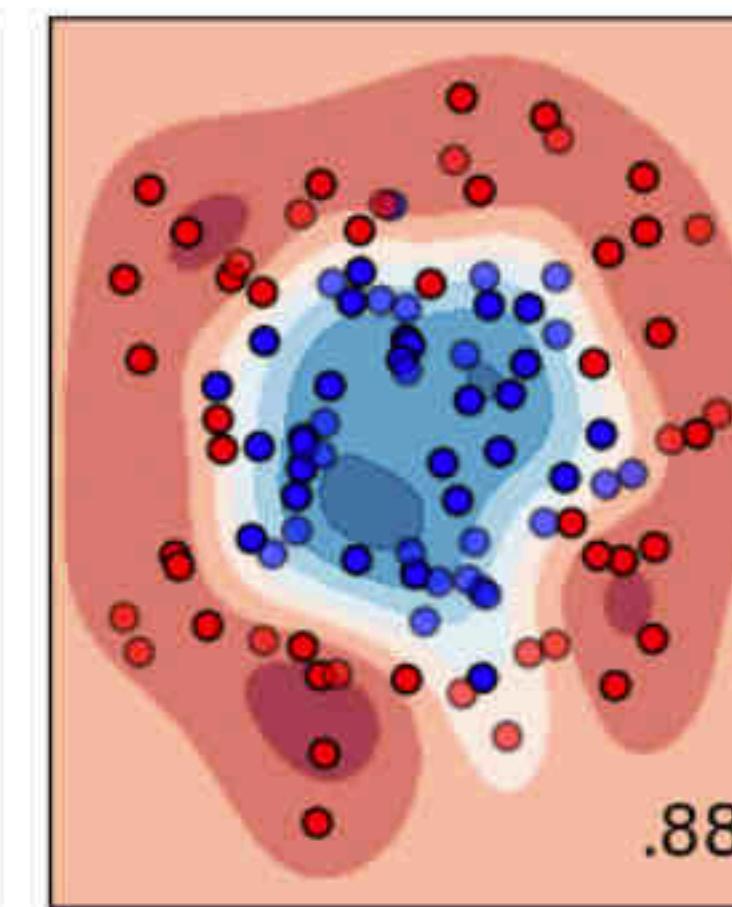
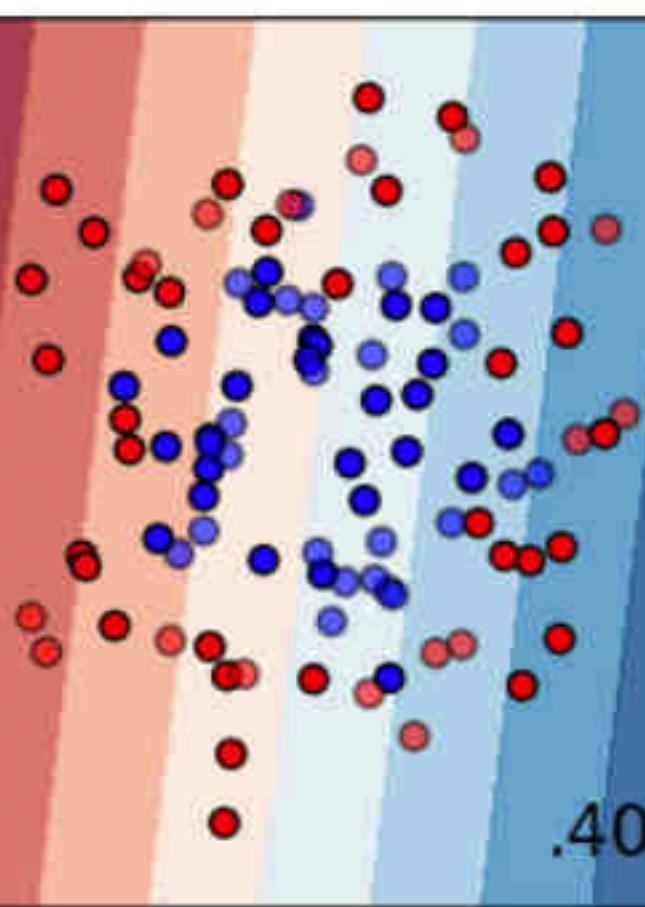
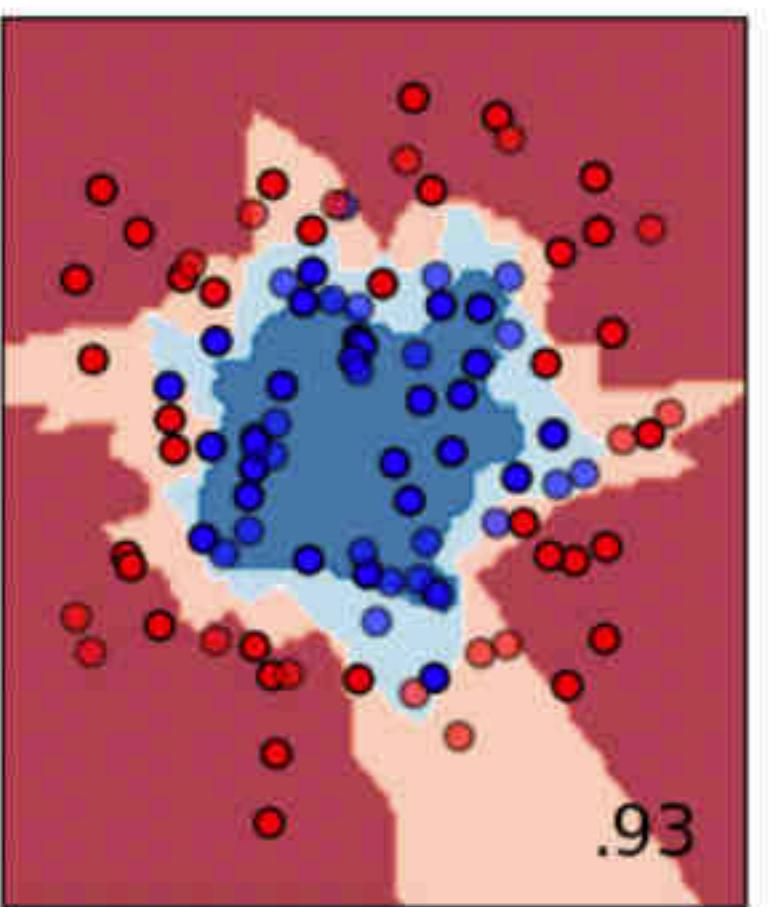
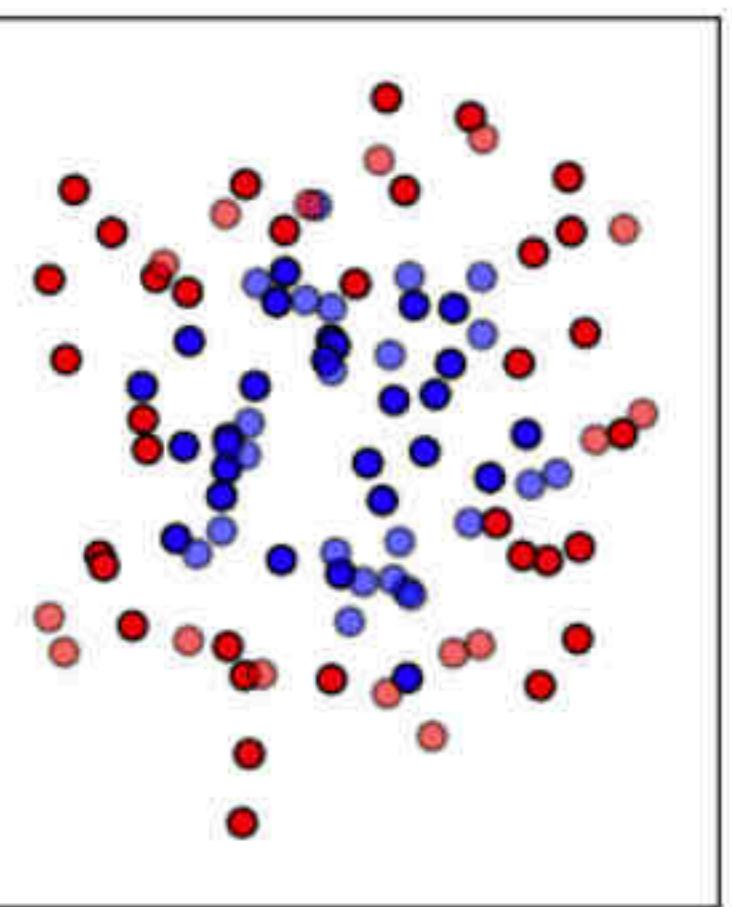
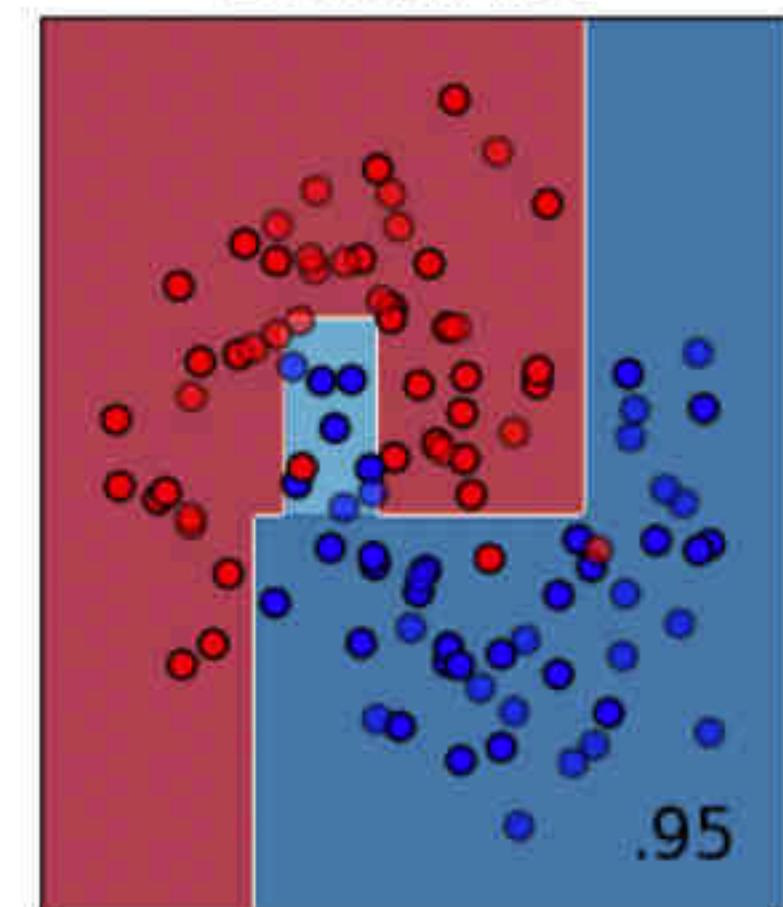
RBF SVM

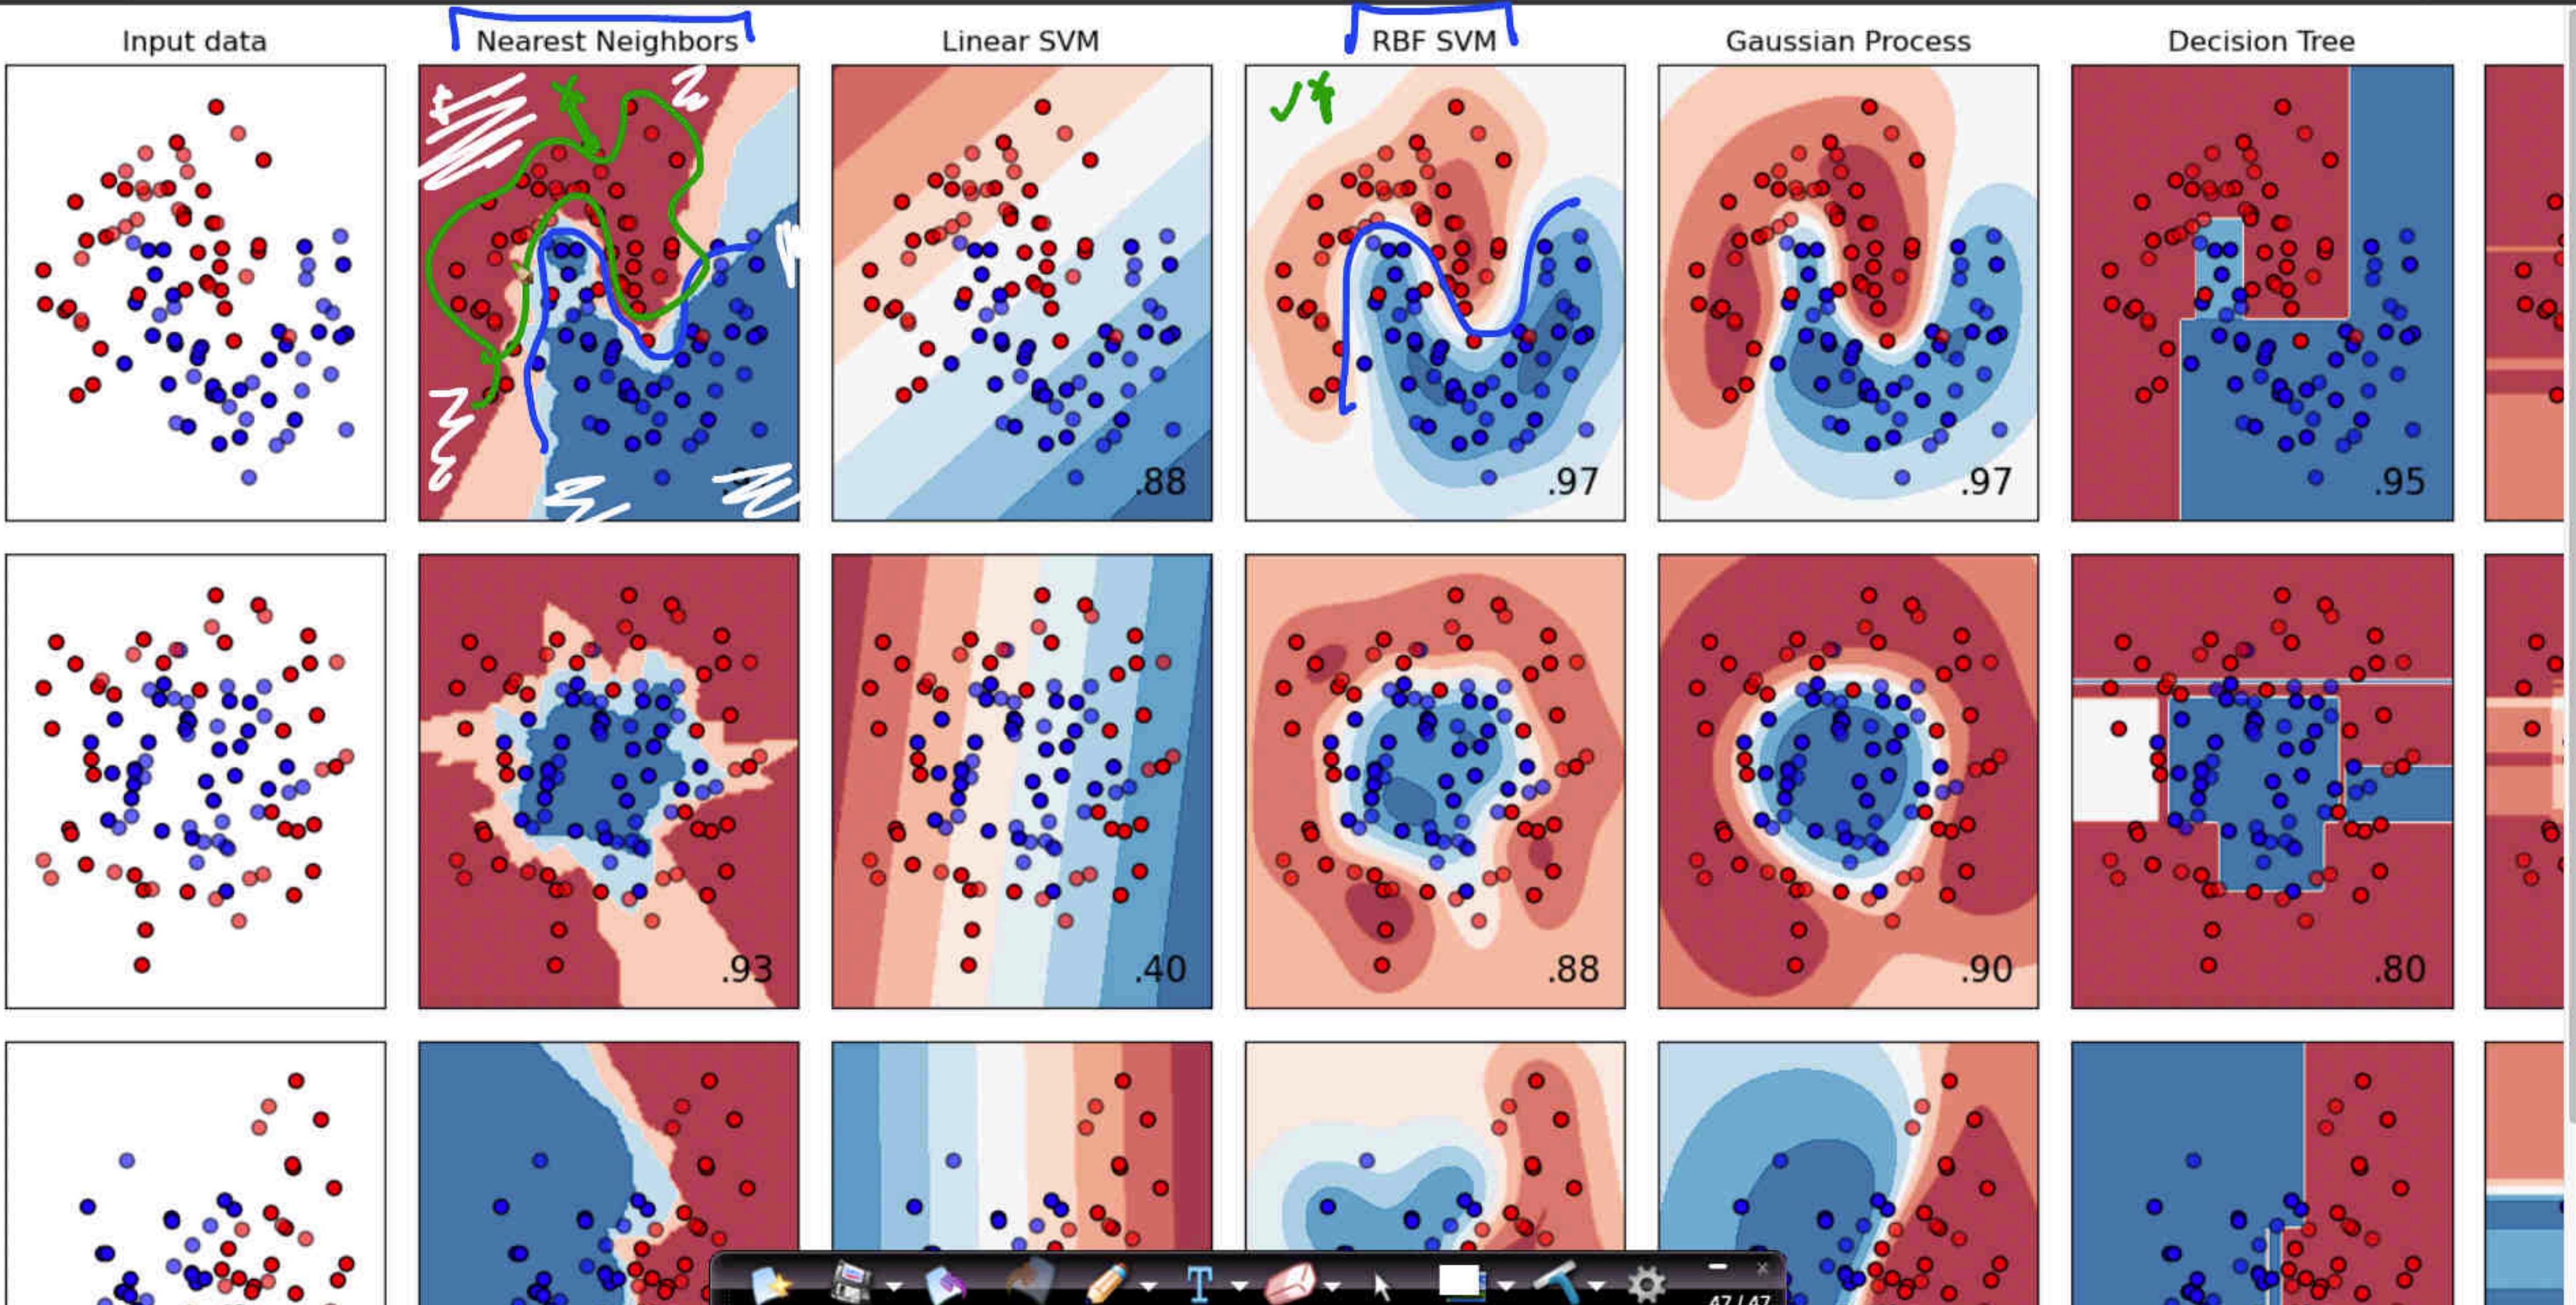


Gaussian Process

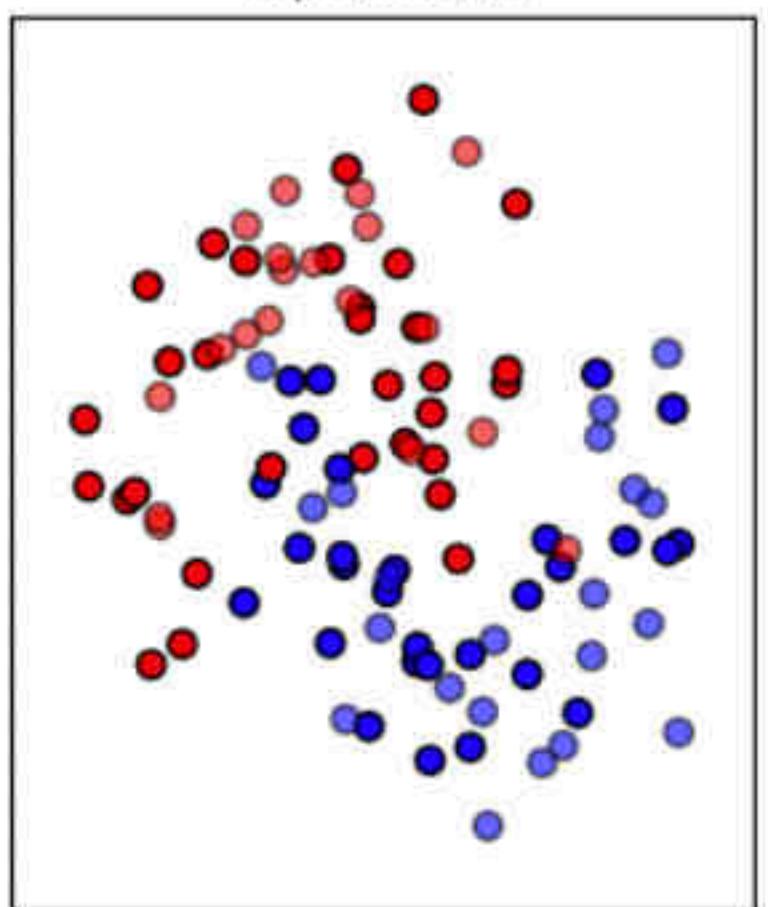


Decision Tree

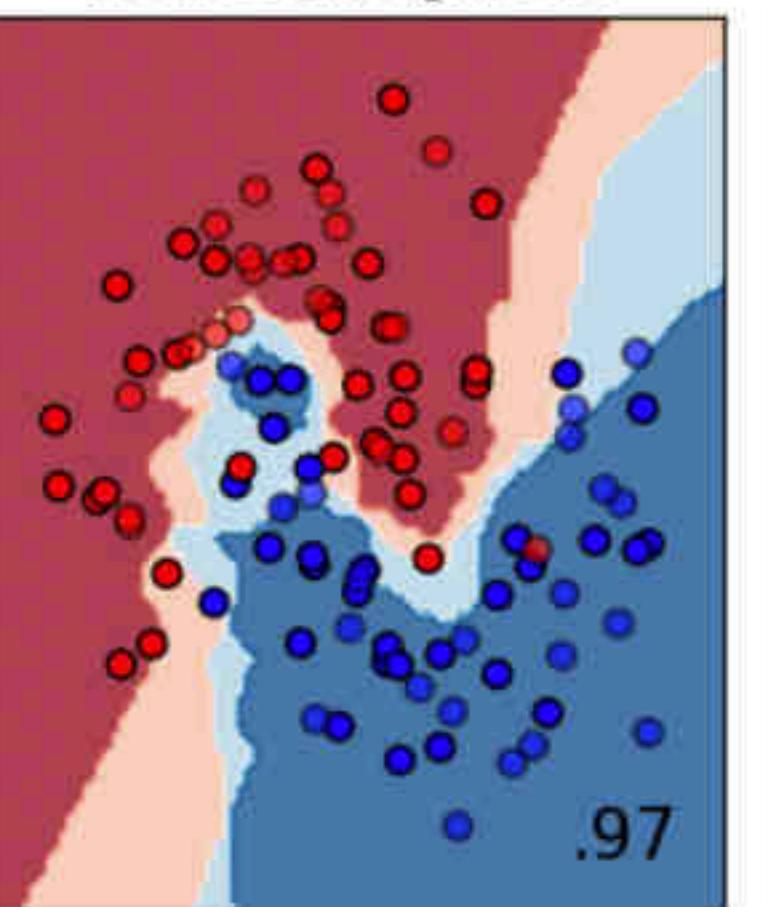




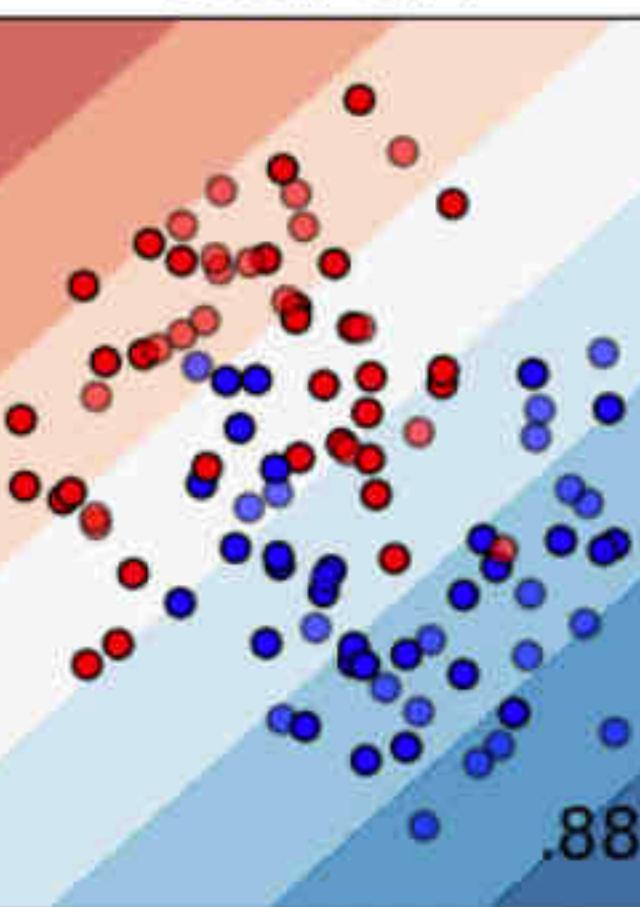
Input data



Nearest Neighbors

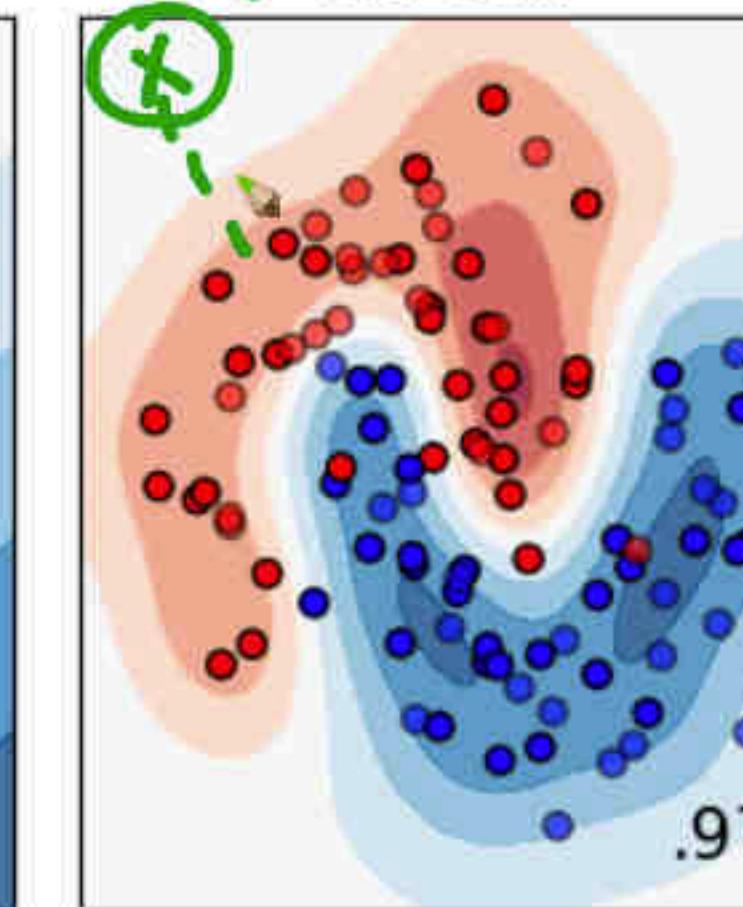


Linear SVM

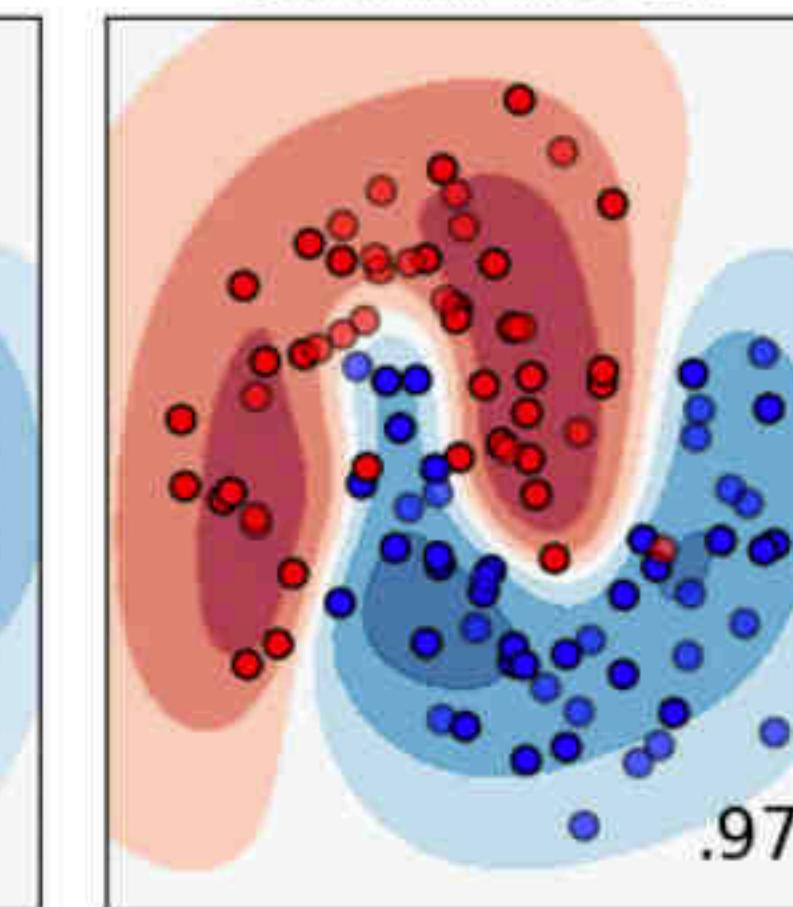


201

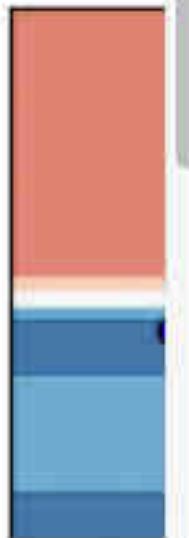
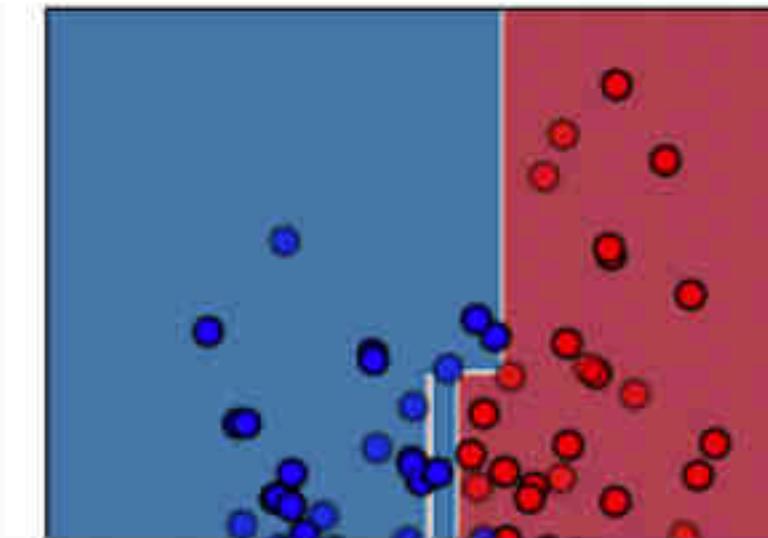
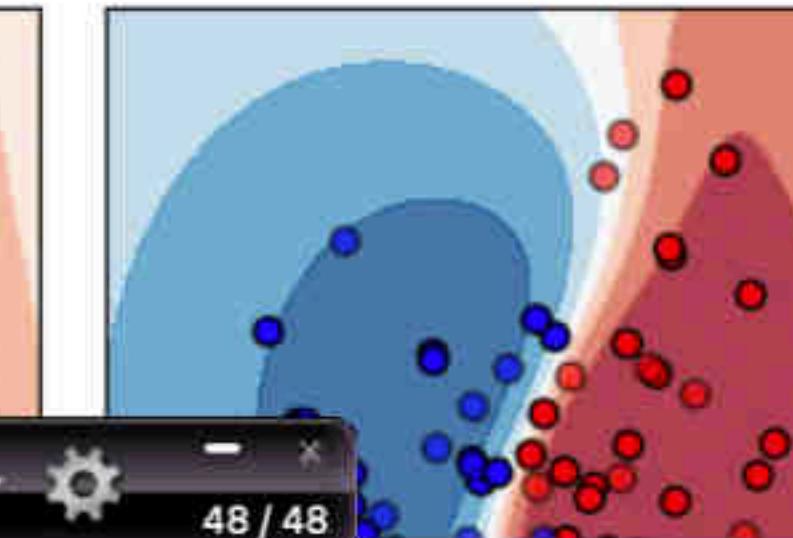
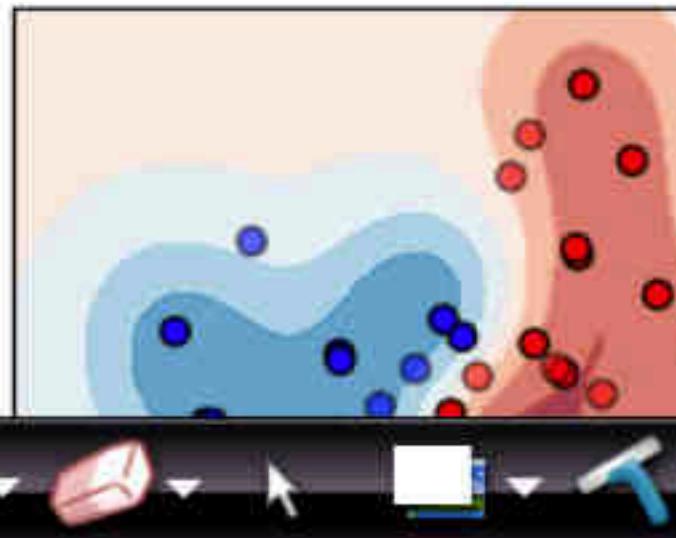
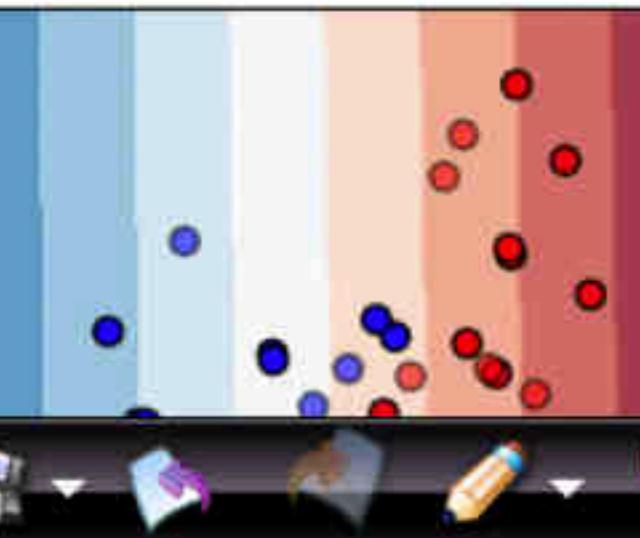
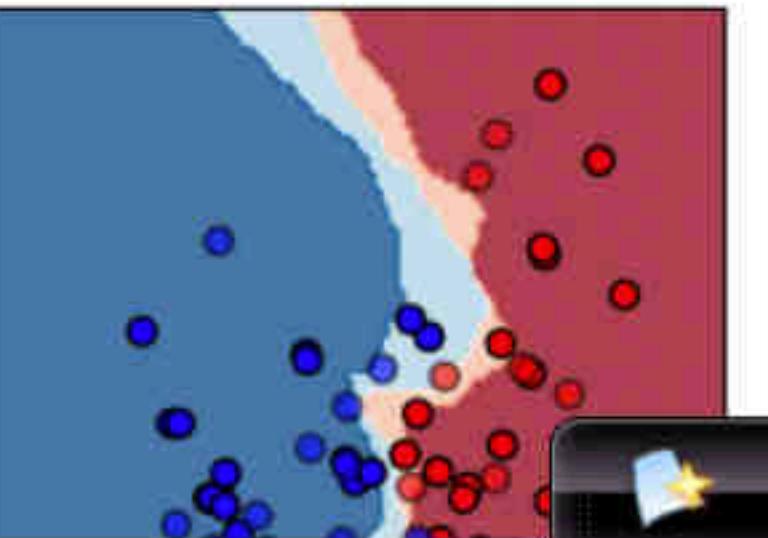
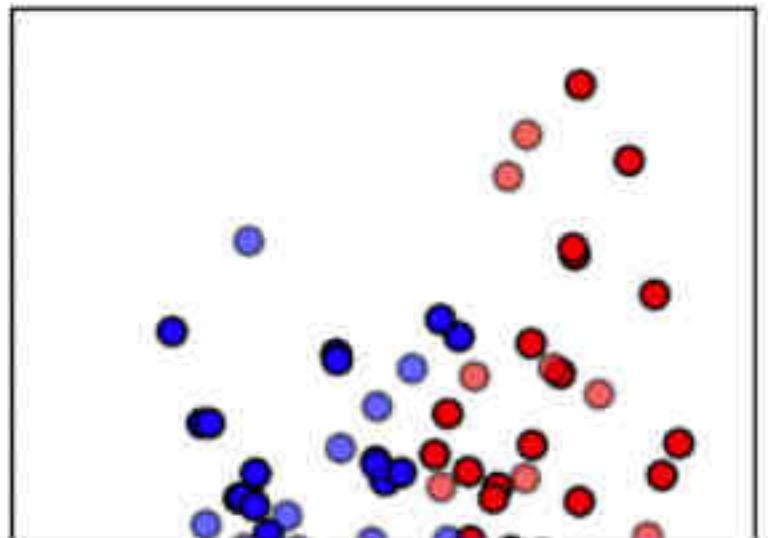
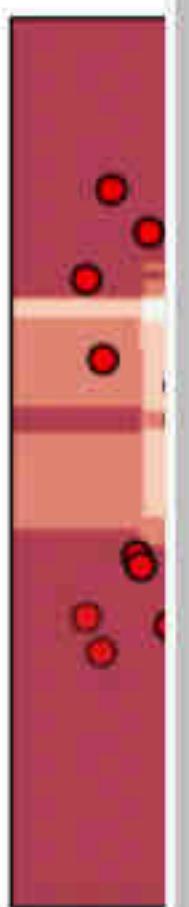
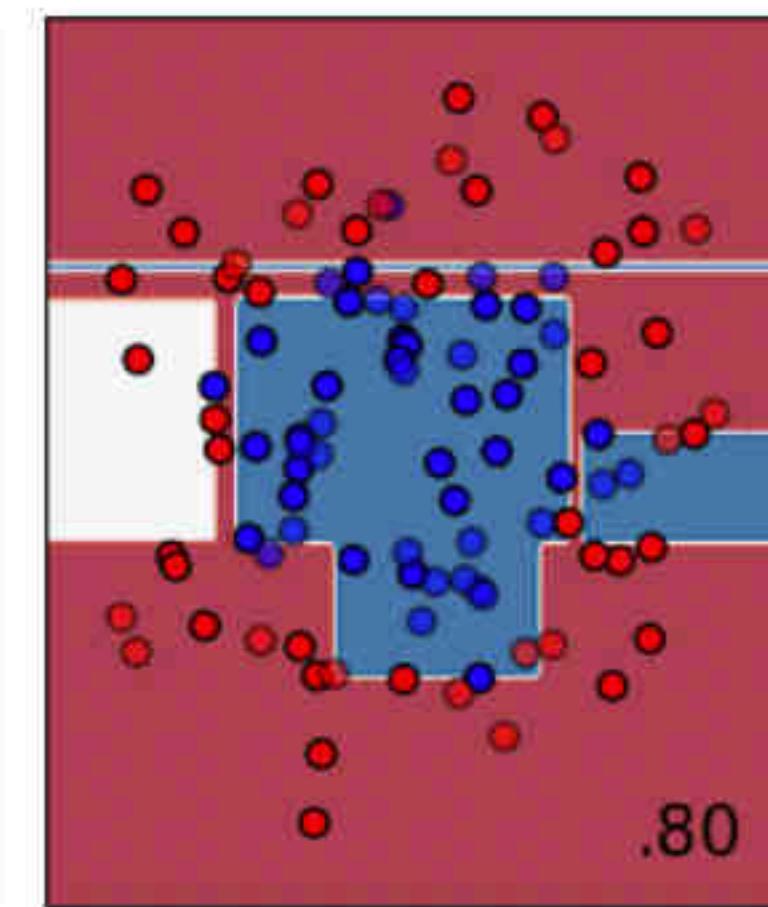
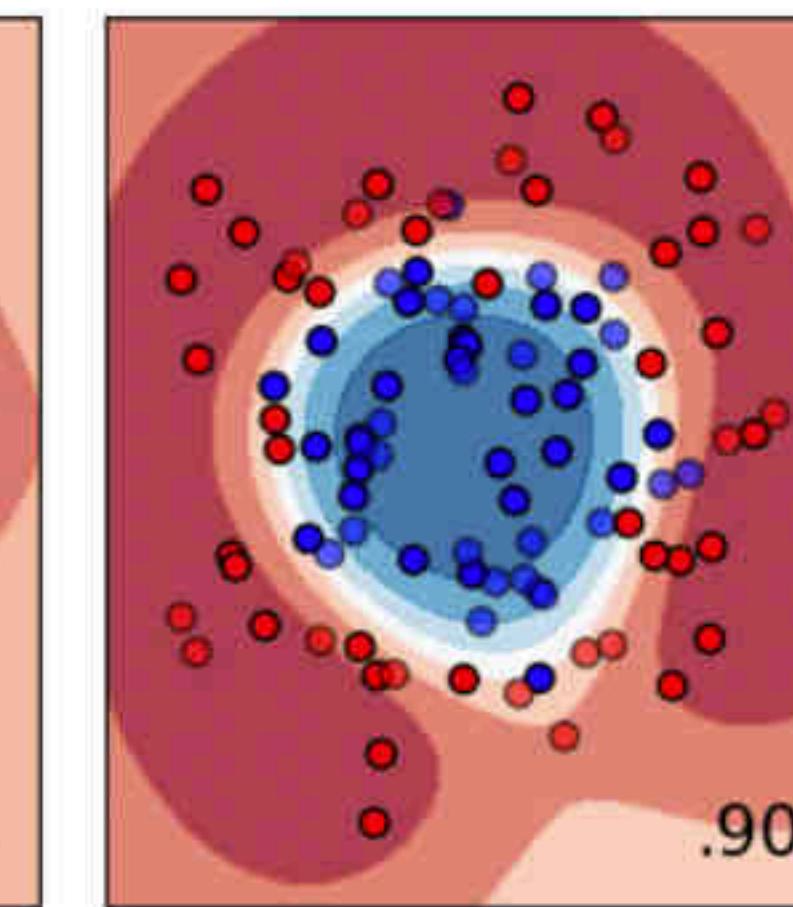
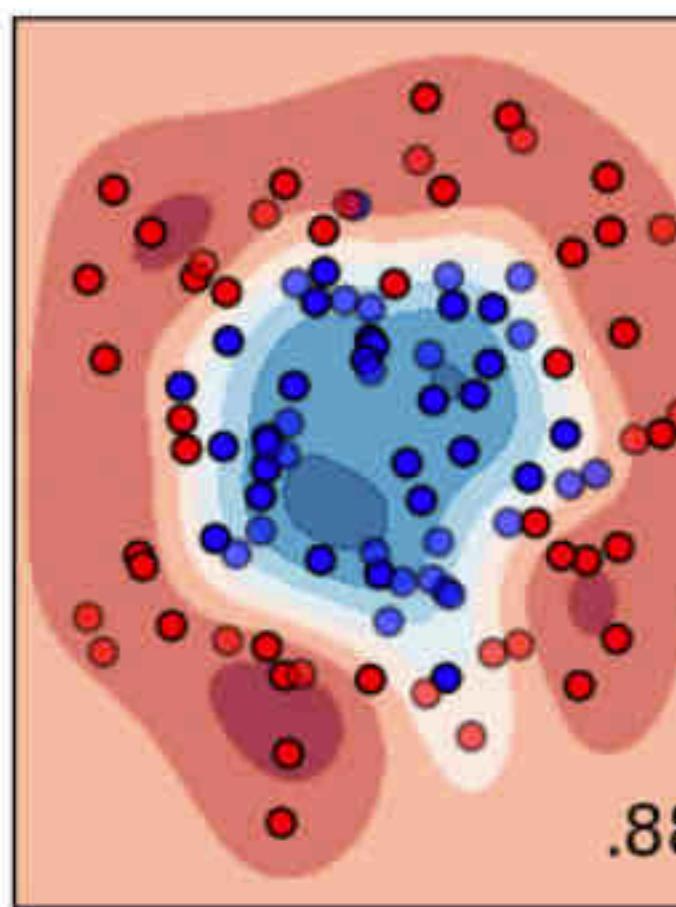
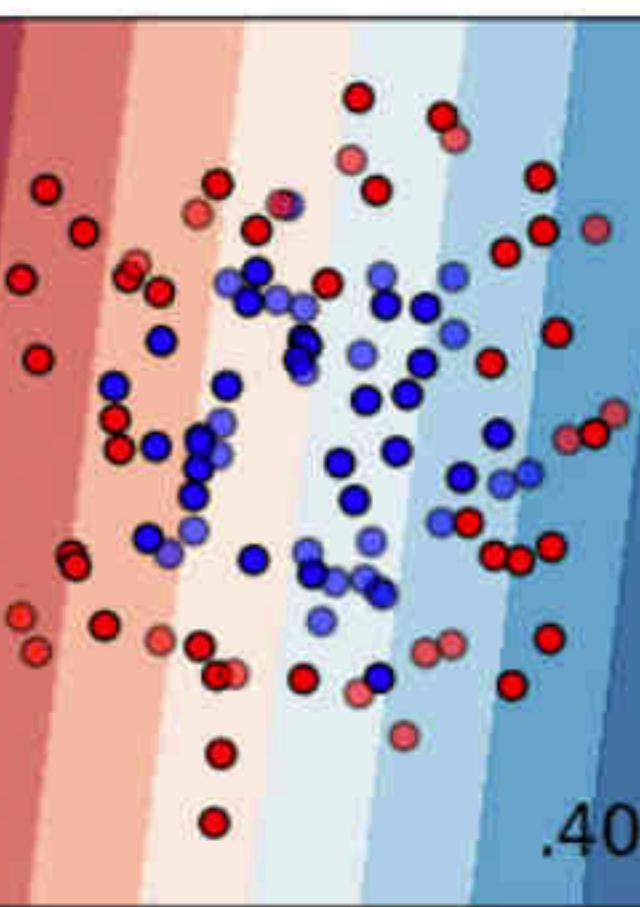
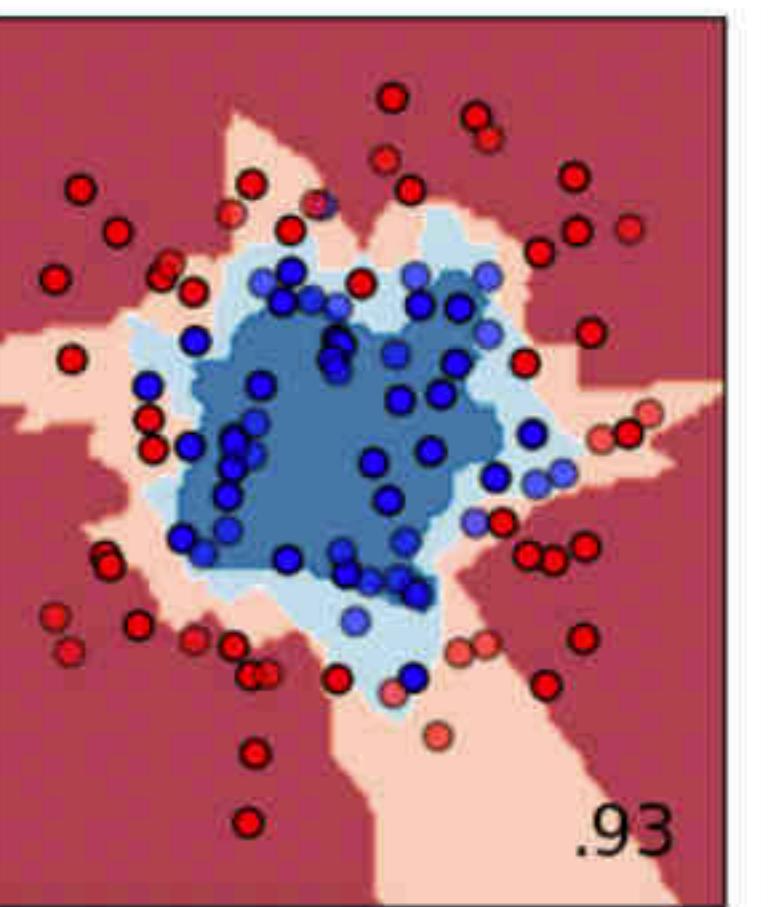
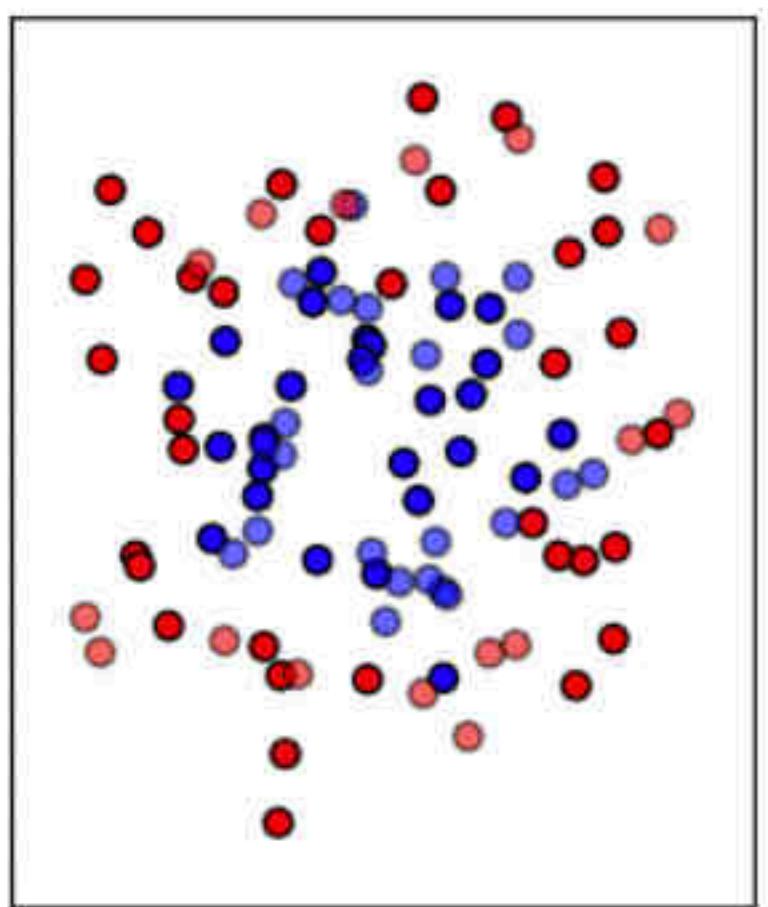
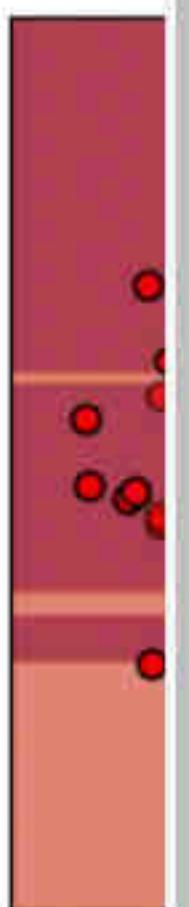
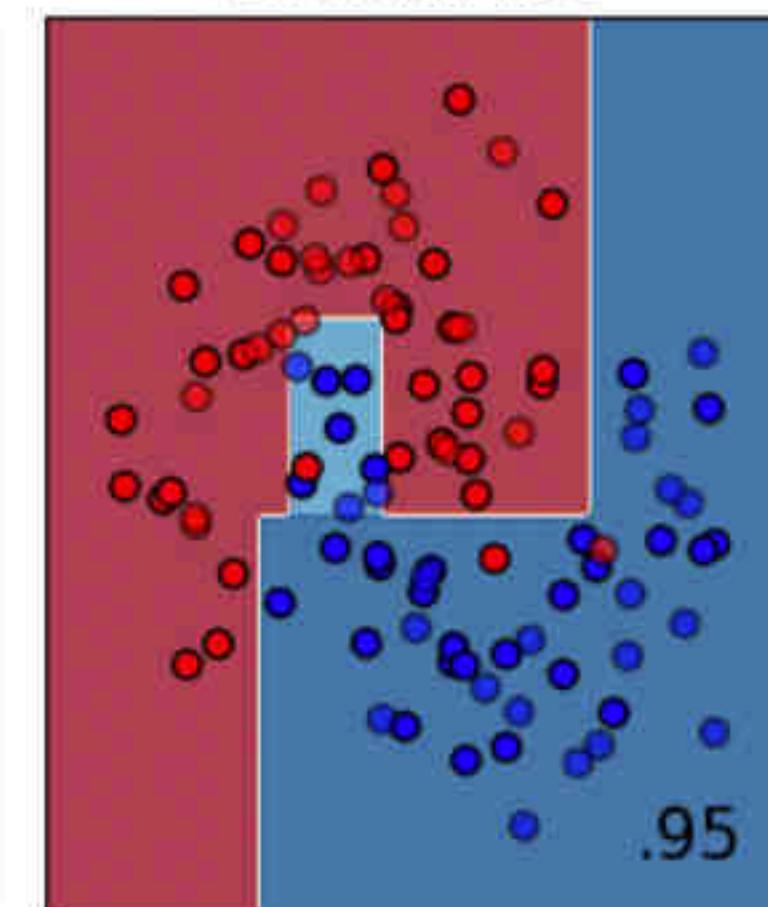
RBF SVM



Gaussian Process



Decision Tree



scikit-learn.org/stable/modules/generated/sklearn.svm.SVC.html

scikit-learn

Install User Guide API Examples Community More ▾

Prev Up Next

scikit-learn 1.1.1 Other versions

Please cite us if you use the software.

sklearn.svm.SVC Examples using sklearn.svm.SVC

1 2 - - j n
1 2
3
n
n x n

Parameters:

C : float, default=1.0
Regularization parameter. The strength of the regularization is inversely proportional to C. Must be strictly positive. The penalty is a squared l2 penalty.

kernel : {'linear', 'poly', 'rbf', 'sigmoid', 'precomputed'} or callable, default='rbf'
Specifies the kernel type to be used in the algorithm. If none is given, 'rbf' will be used. If a callable is given it is used to pre-compute the kernel matrix from data matrices; that matrix should be an array of shape (n_samples, n_samples).

degree : int, default=3
Degree of the polynomial kernel function ('poly'). Ignored by all other kernels.

gamma : {'scale', 'auto'} or float, default='scale'
Kernel coefficient for 'rbf', 'poly' and 'sigmoid'.

- if gamma='scale' (default) is passed then it uses $1 / (\text{n_features} * \text{X.var()})$ as value of gamma,
- if 'auto', uses $1 / \text{n_features}$.

Changed in version 0.22: The default value of gamma changed from 'auto' to 'scale'!

Toggle Menu

[Prev](#) [Up](#) [Next](#)

scikit-learn 1.1.1

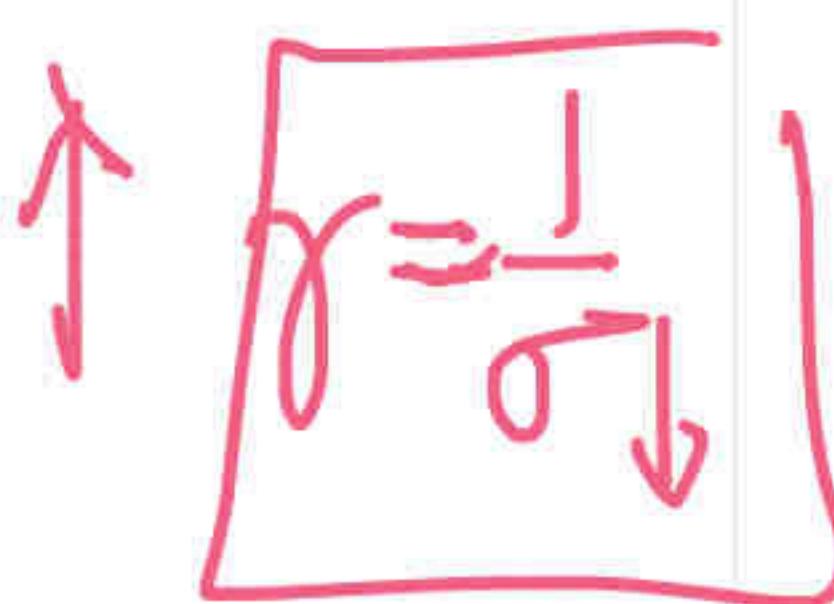
[Other versions](#)

Please [cite us](#) if you use the software.

RBF SVM parameters

[Load and prepare data set](#)[Train classifiers](#)[Visualization](#)

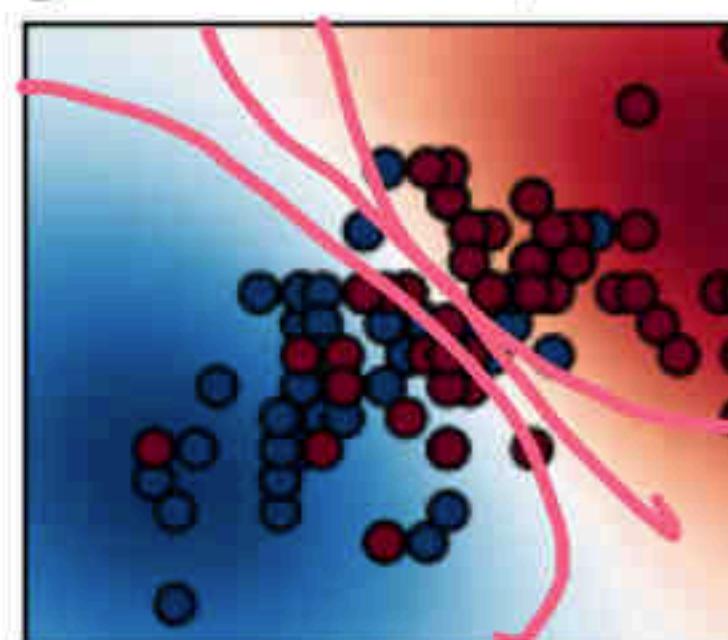
RBF-SVM



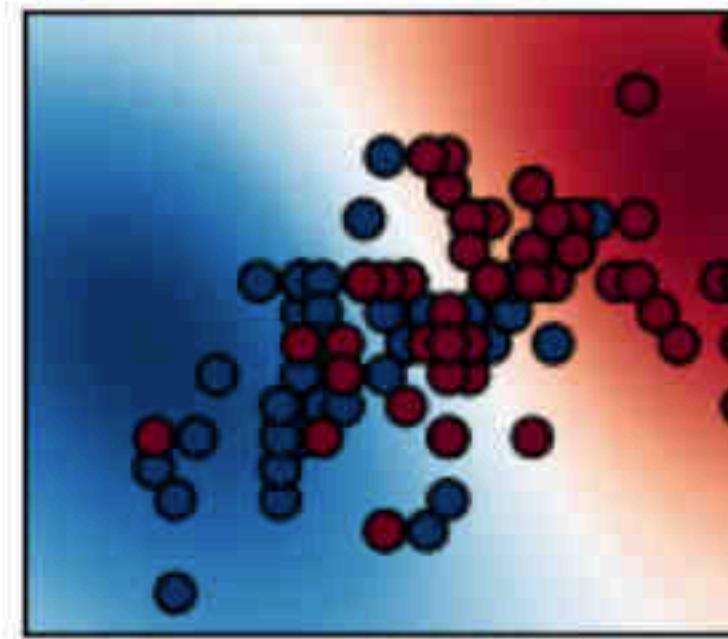
```
plt.yticks(())
plt.axis("tight")
```

```
scores = grid.cv_results_["mean_test_score"].reshape(len(C_range), len(gamma_range))
```

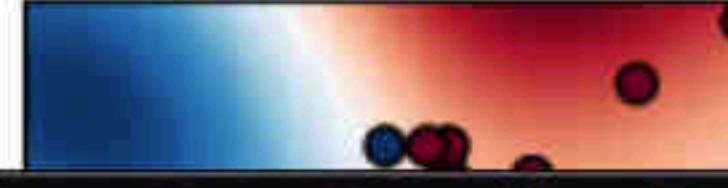
$\sigma \uparrow$ Small
gamma=10⁻¹, C=10⁻²



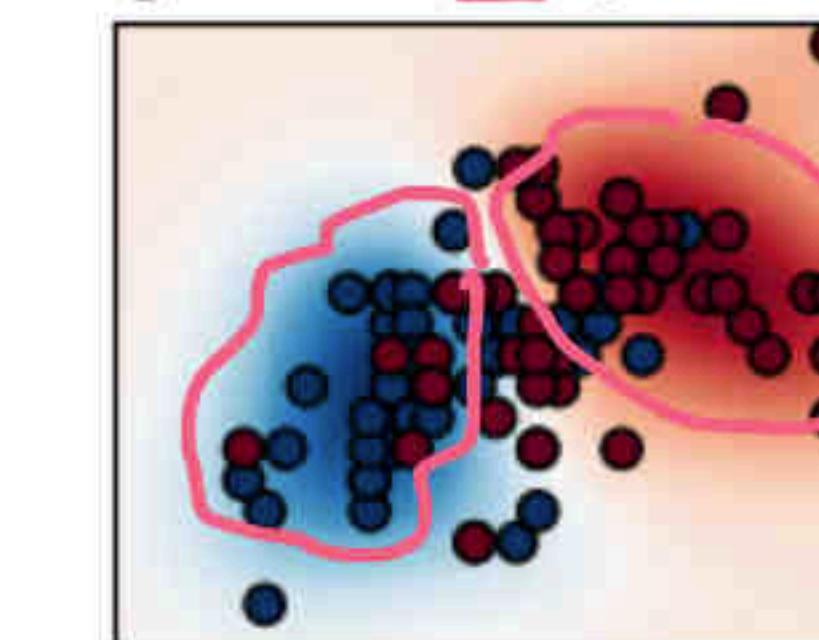
gamma=10⁻¹, C=10⁰



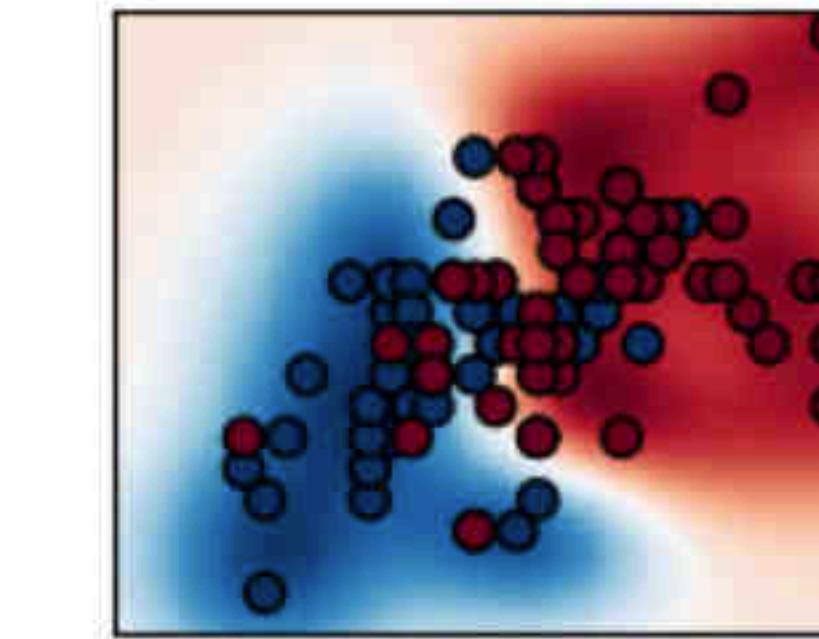
gamma=10⁻¹, C=10²



$\sigma \downarrow$
gamma=10⁰, C=10⁻²

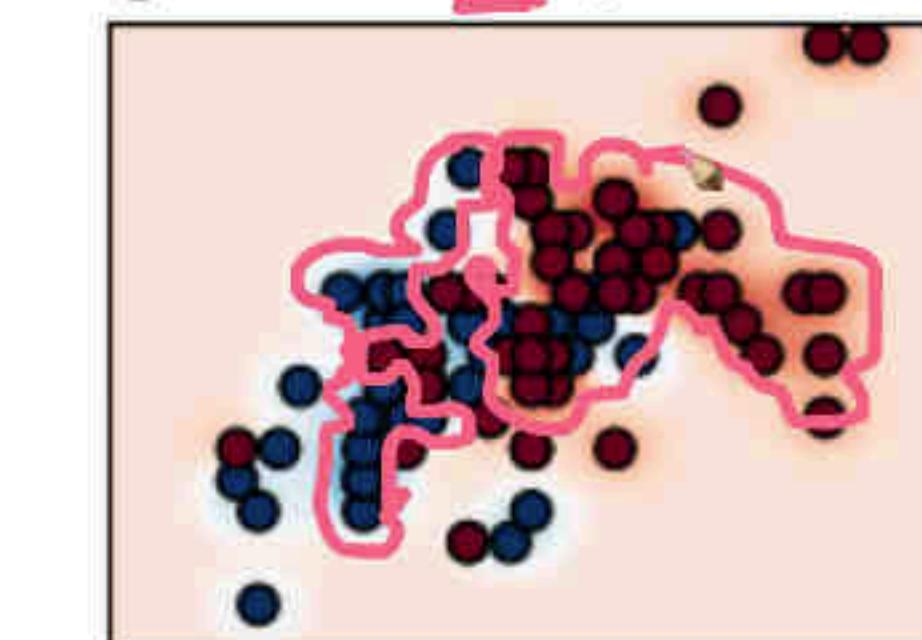


gamma=10⁰, C=10⁰

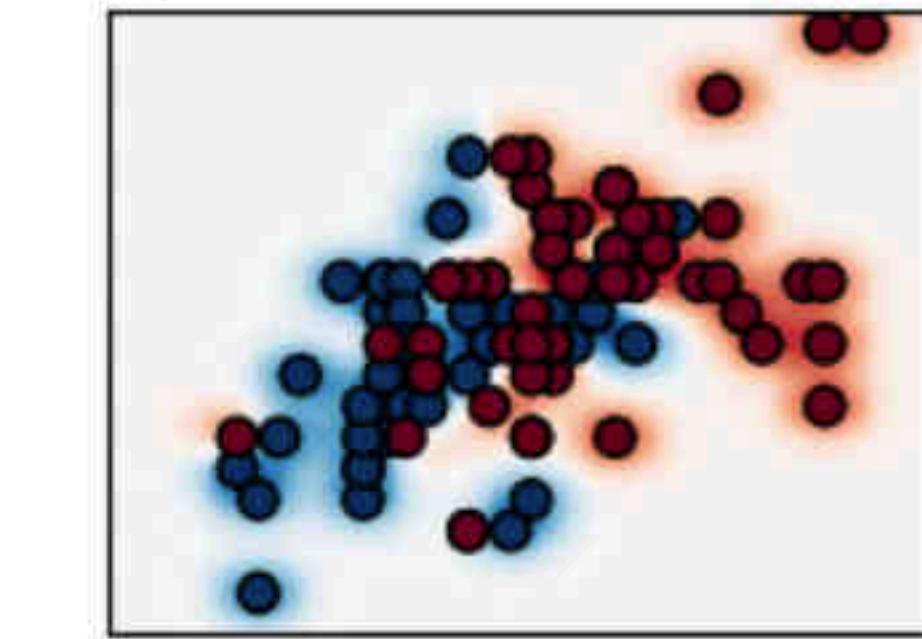


gamma=10⁰, C=10²

$\sigma \downarrow$
gamma=10¹, C=10⁻²

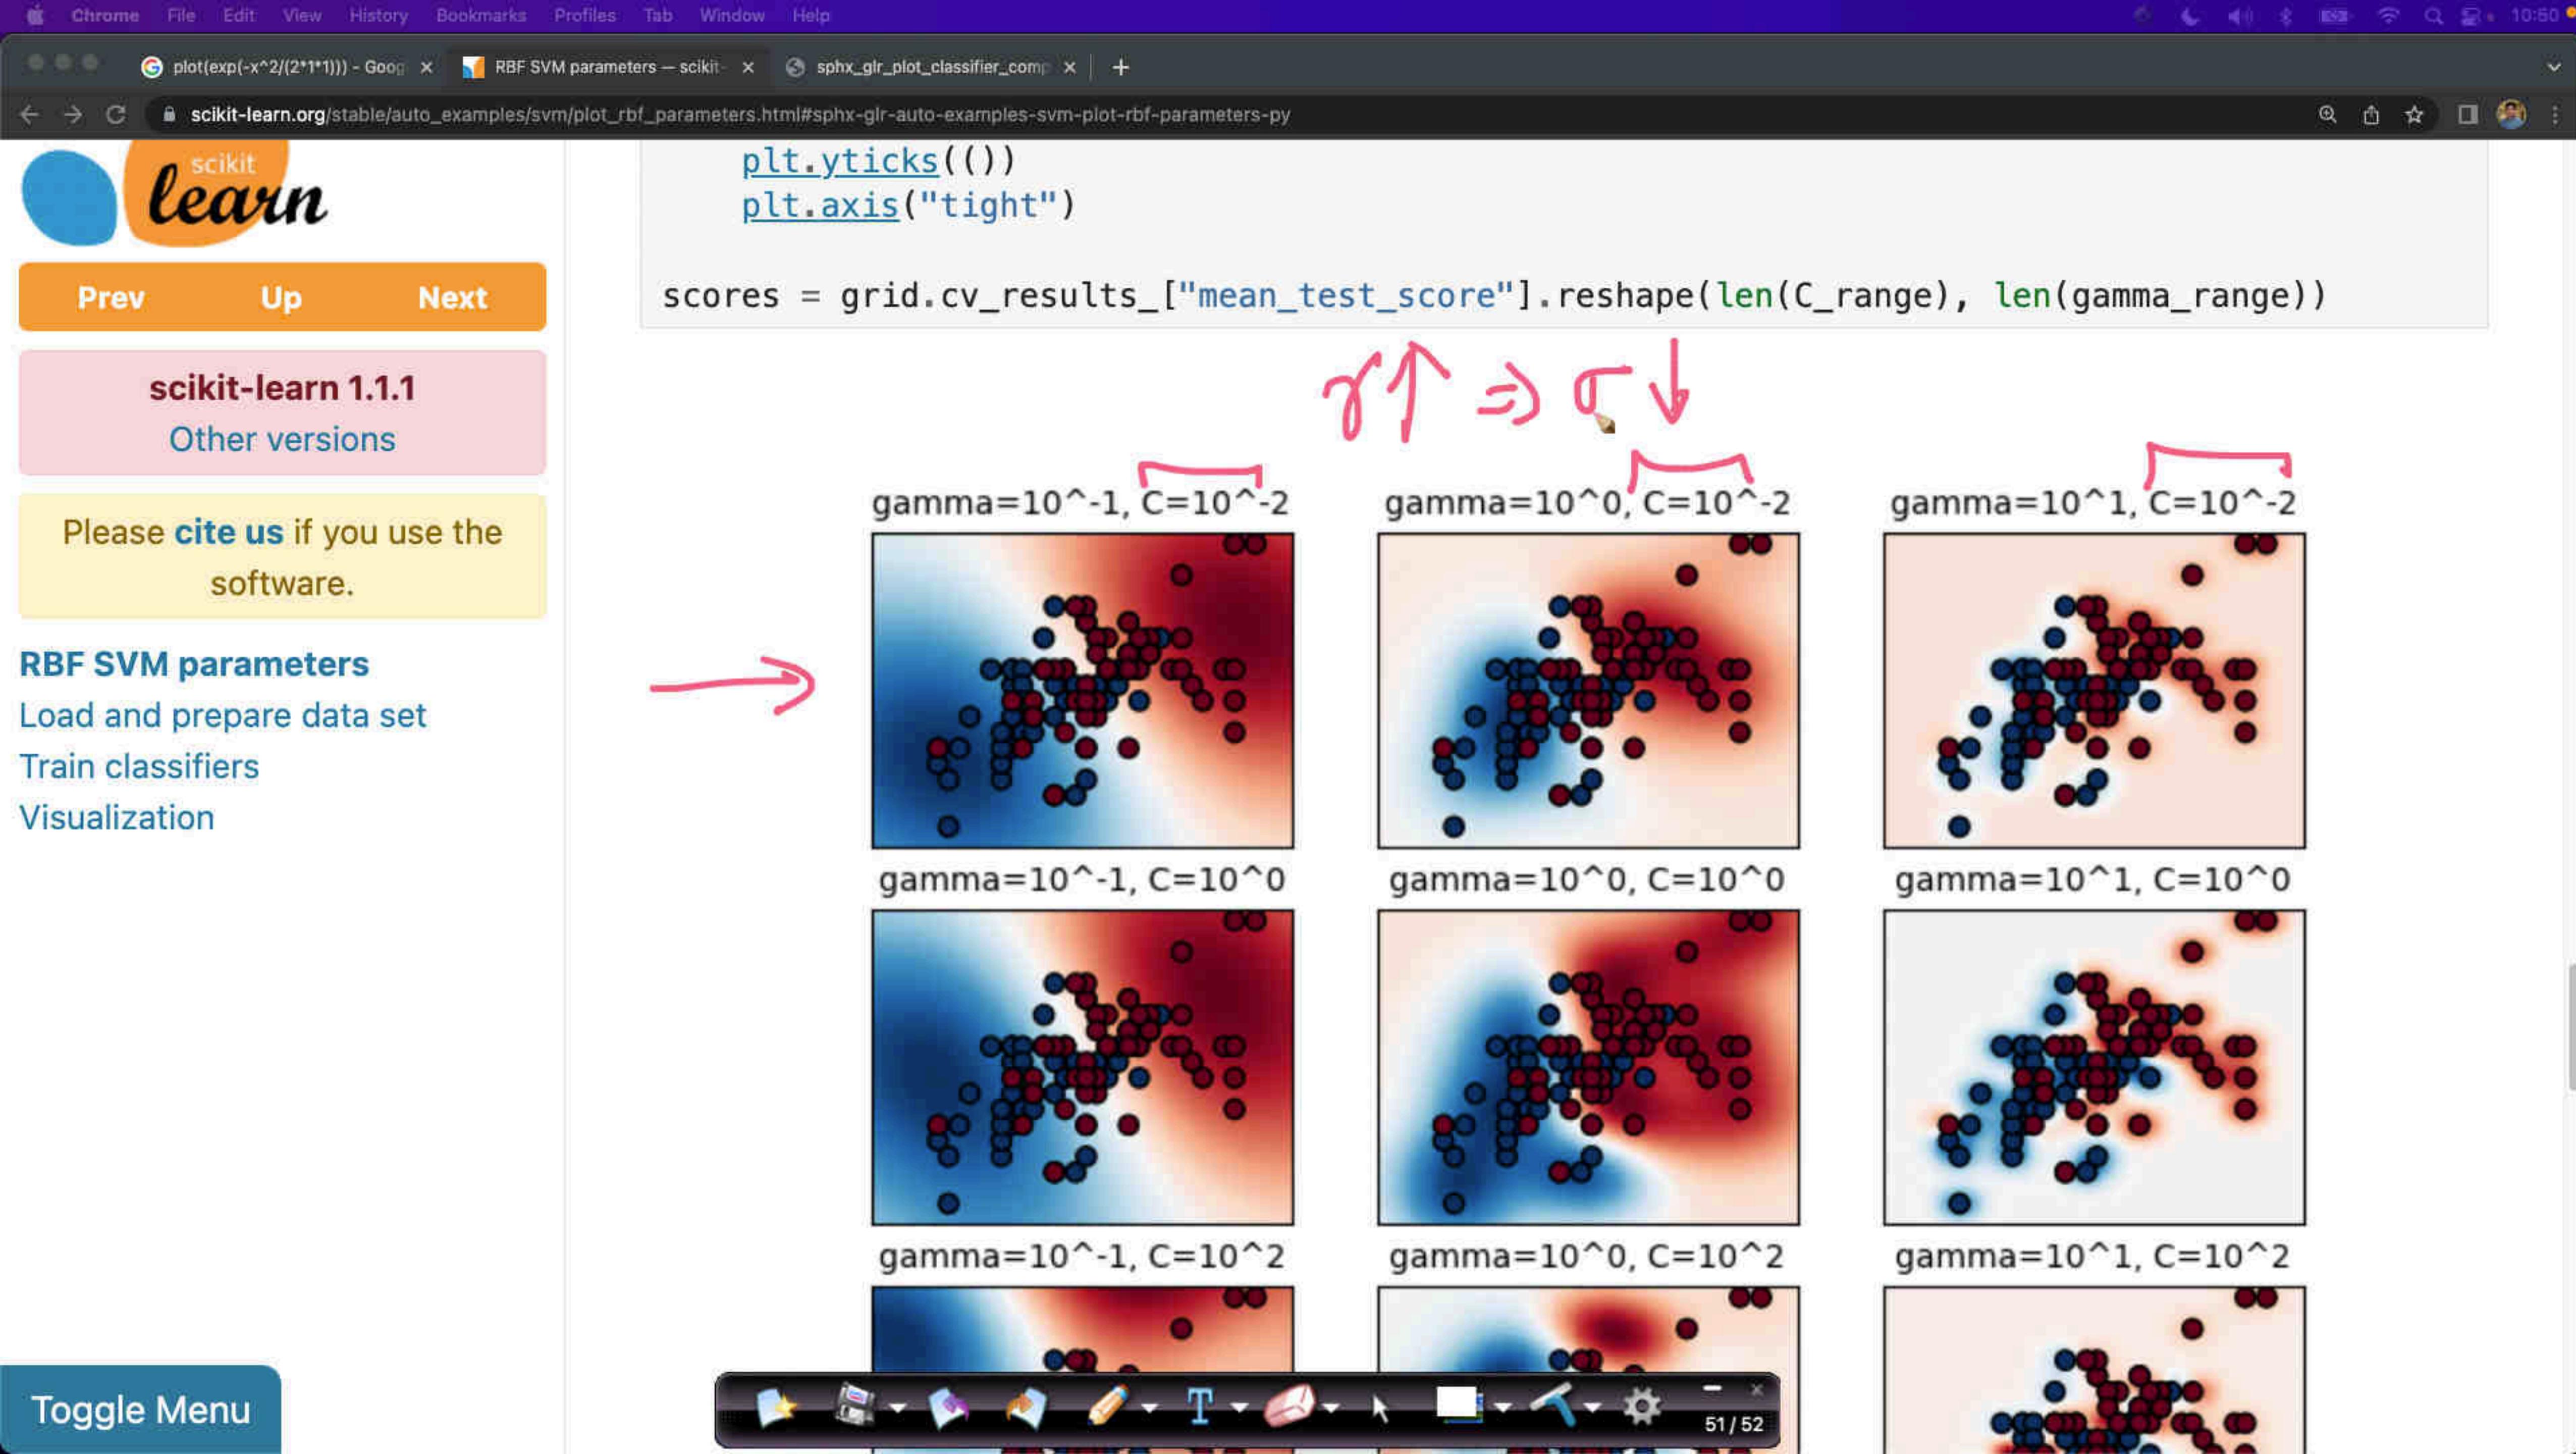


gamma=10¹, C=10⁰



gamma=10¹, C=10²

[Toggle Menu](#)



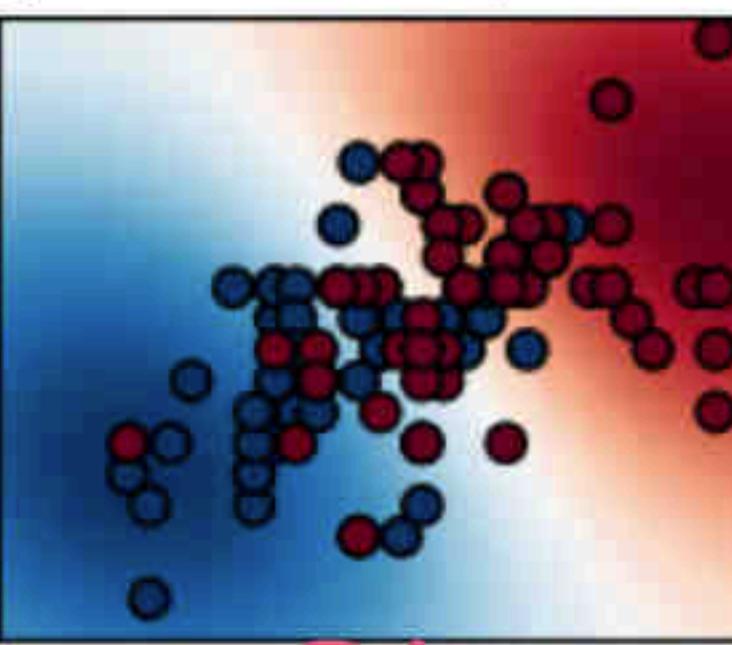
[Prev](#) [Up](#) [Next](#)**scikit-learn 1.1.1**[Other versions](#)

Please [cite us](#) if you use the software.

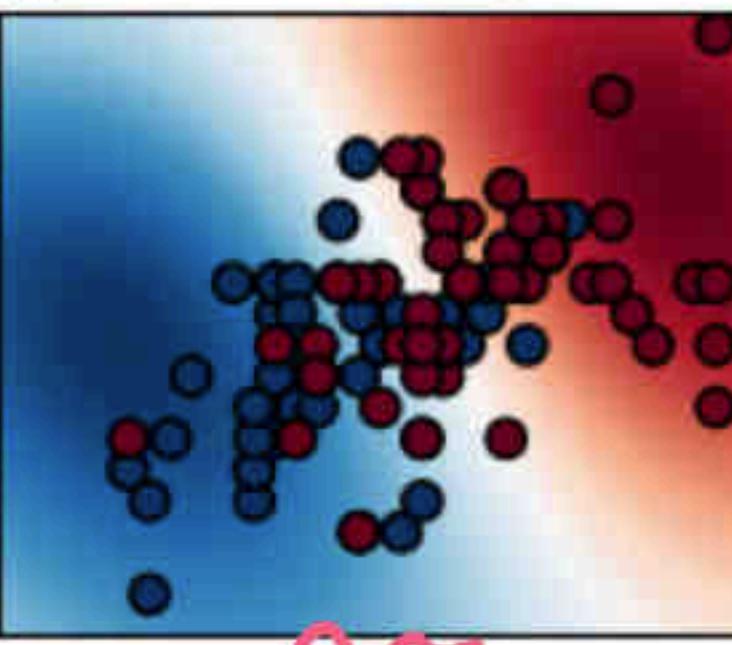
RBF SVM parameters[Load and prepare data set](#)[Train classifiers](#)[Visualization](#)[Toggle Menu](#)

inc
C

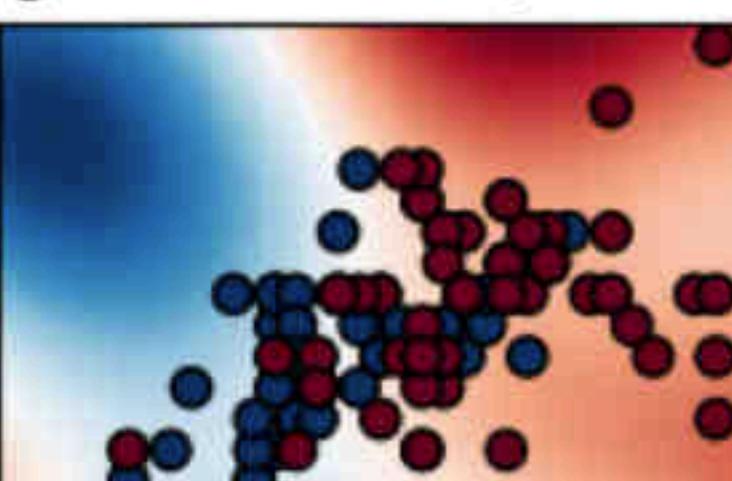
gamma=10⁻¹, C=10⁻²



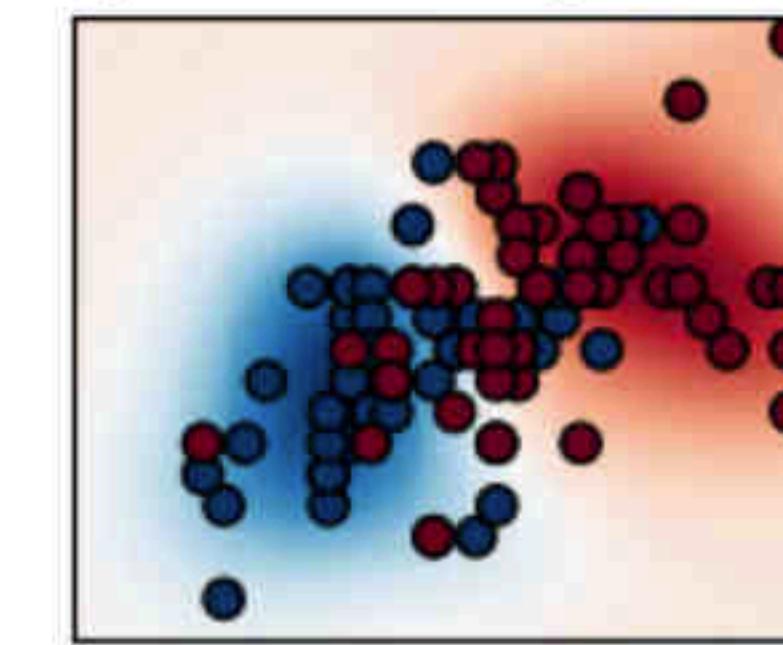
gamma=10⁻¹, C=10⁰



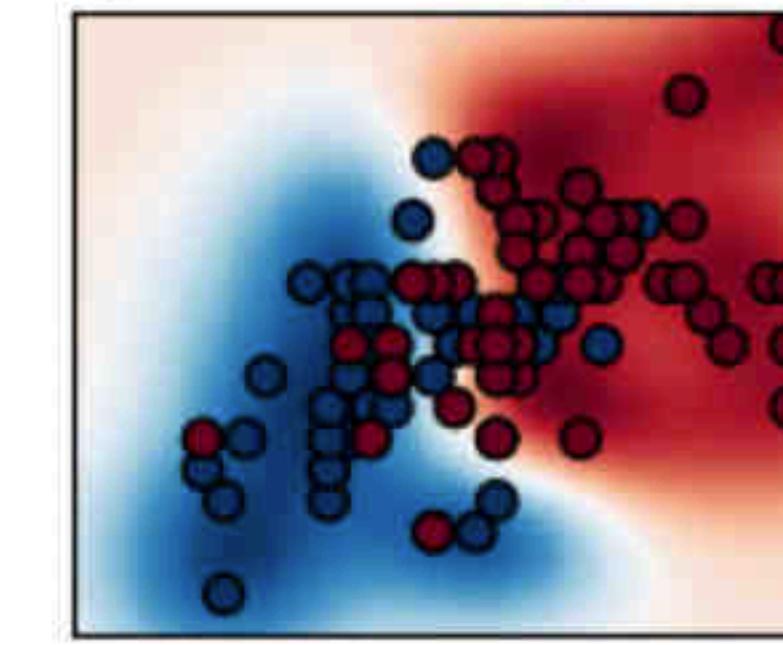
gamma=10⁻¹, C=10²



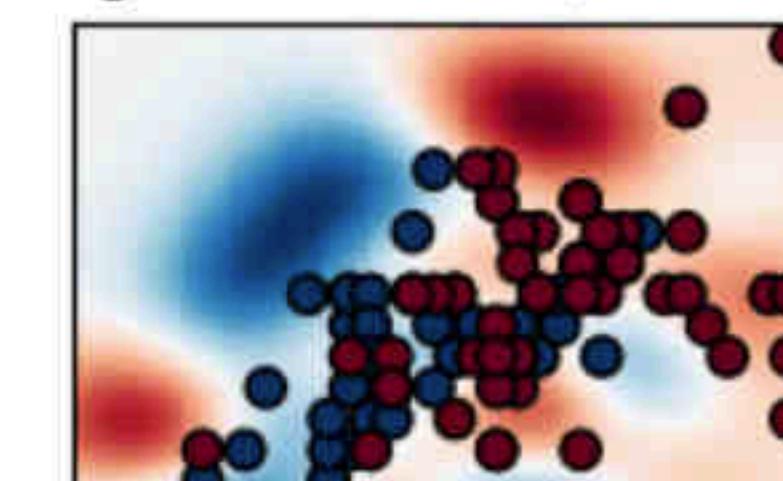
gamma=10⁰, C=10⁻²



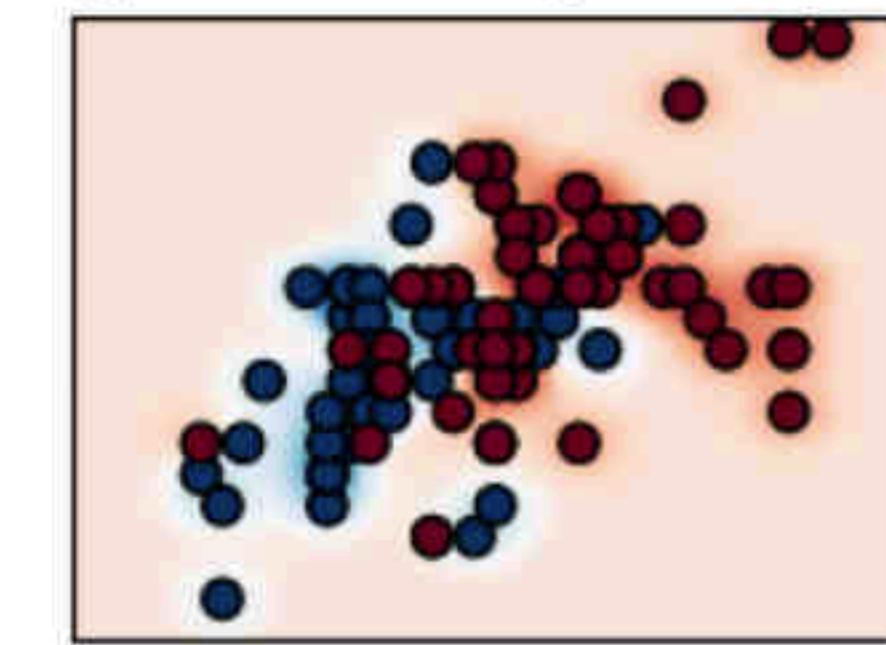
gamma=10⁰, C=10⁰



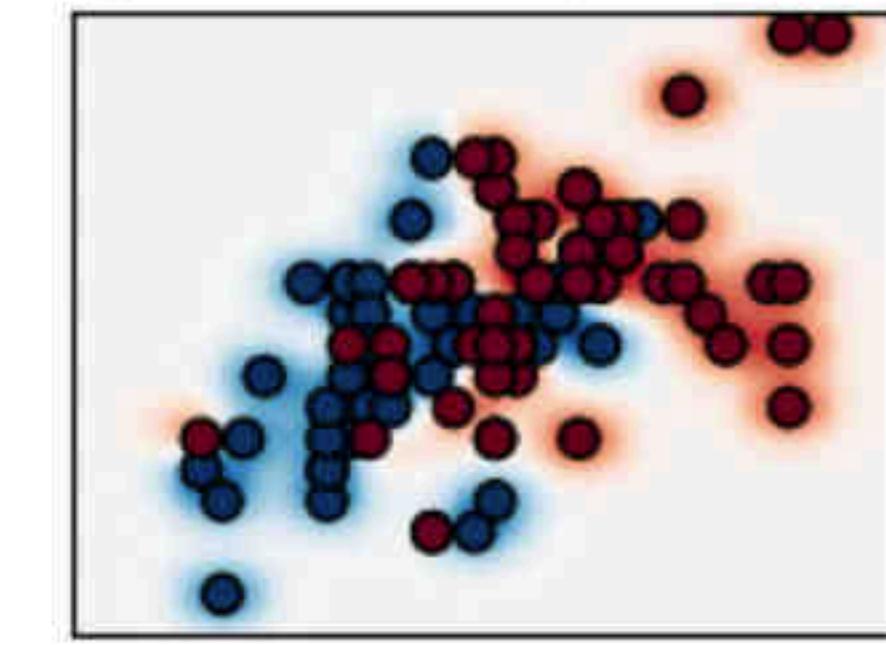
gamma=10⁰, C=10²



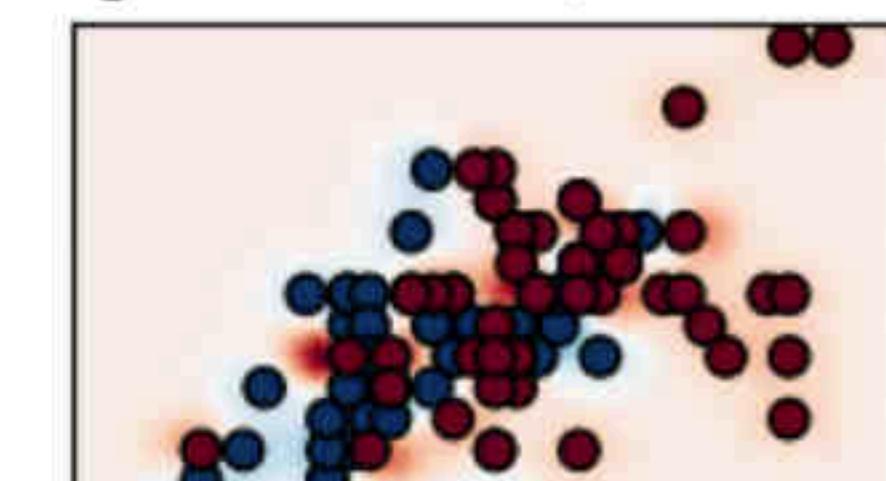
gamma=10¹, C=10⁻²



gamma=10¹, C=10⁰



gamma=10¹, C=10²



scikit-learn.org/stable/auto_examples/svm/plot_rbf_parameters.html#sphx-glr-auto-examples-svm-plot-rbf-parameters-py

scikit learn

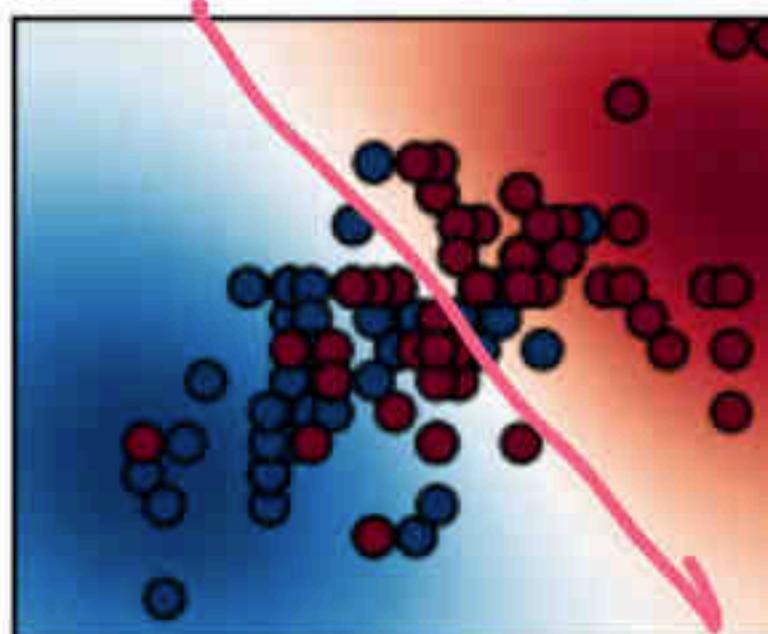
Prev Up Next

scikit-learn 1.1.1
Other versions

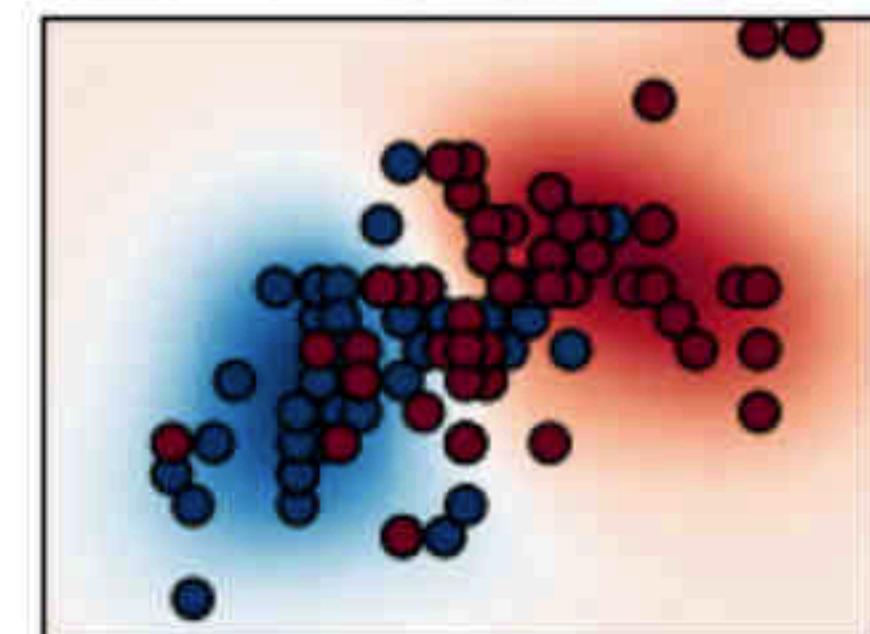
Please cite us if you use the software.

RBF SVM parameters
Load and prepare data set
Train classifiers
Visualization

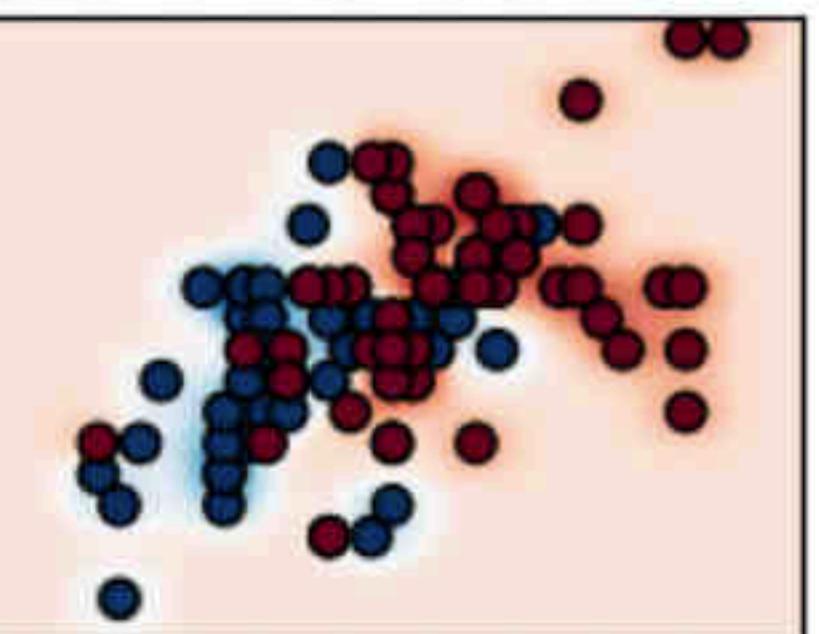
gamma=10^-1, C=10^-2



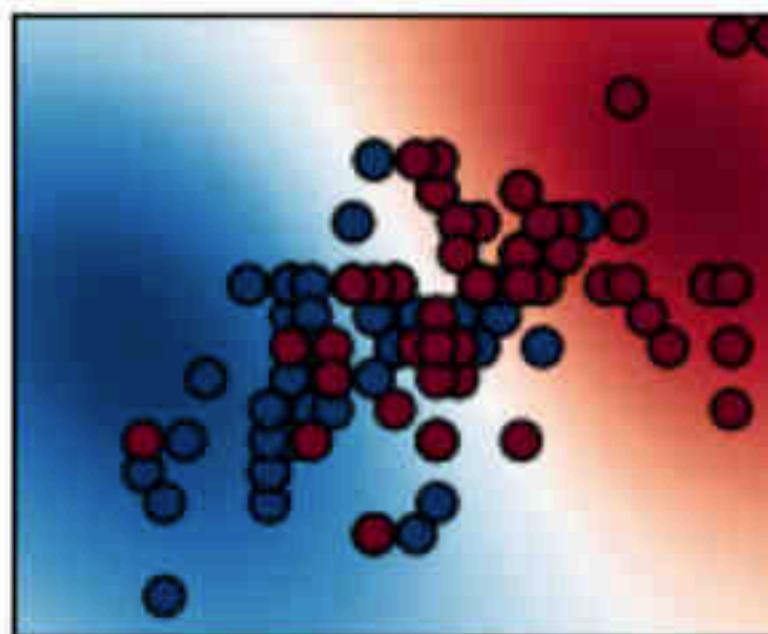
gamma=10^0, C=10^-2



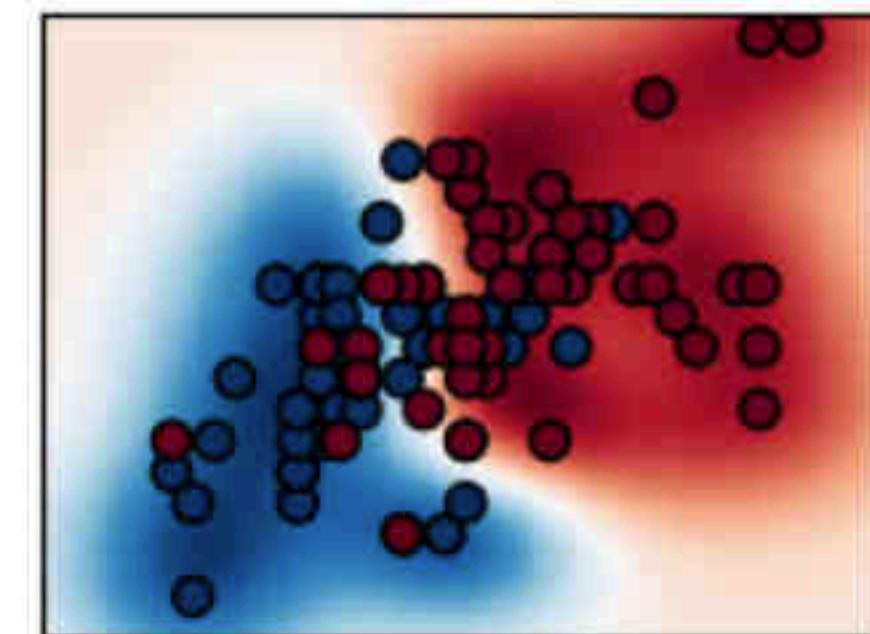
gamma=10^1, C=10^-2



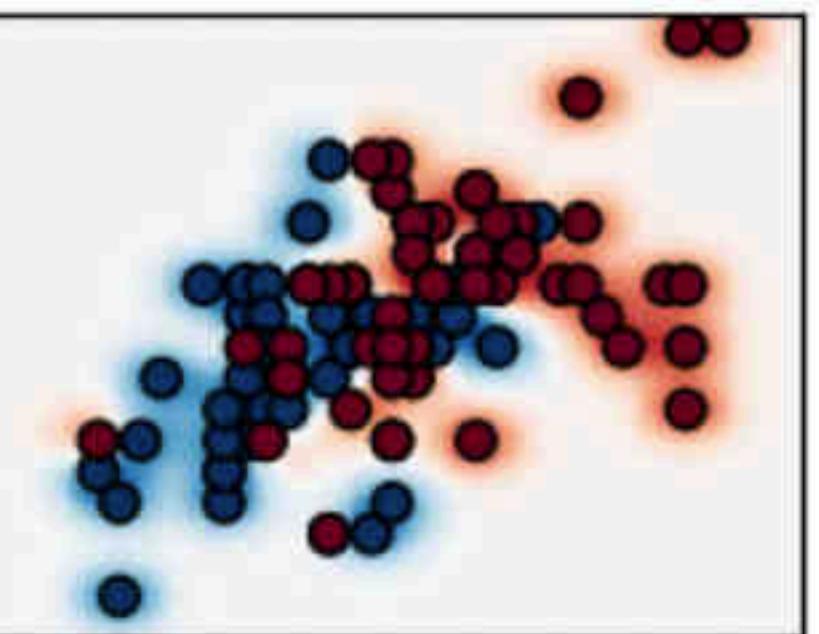
gamma=10^-1, C=10^0



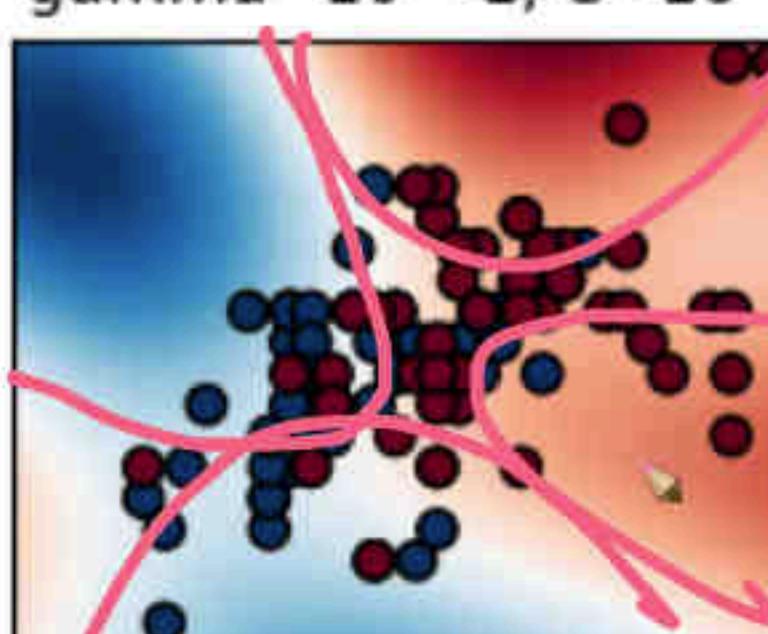
gamma=10^0, C=10^0



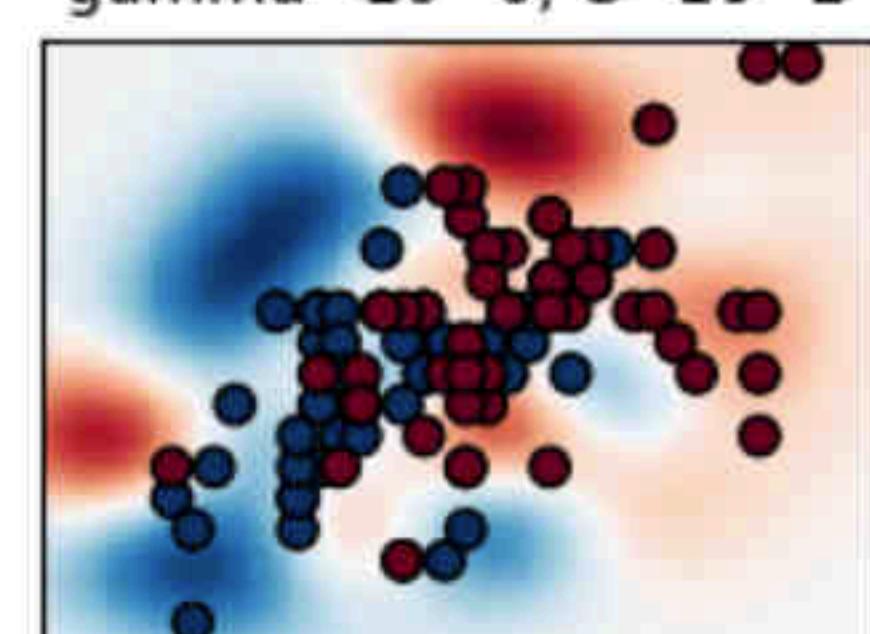
gamma=10^1, C=10^0



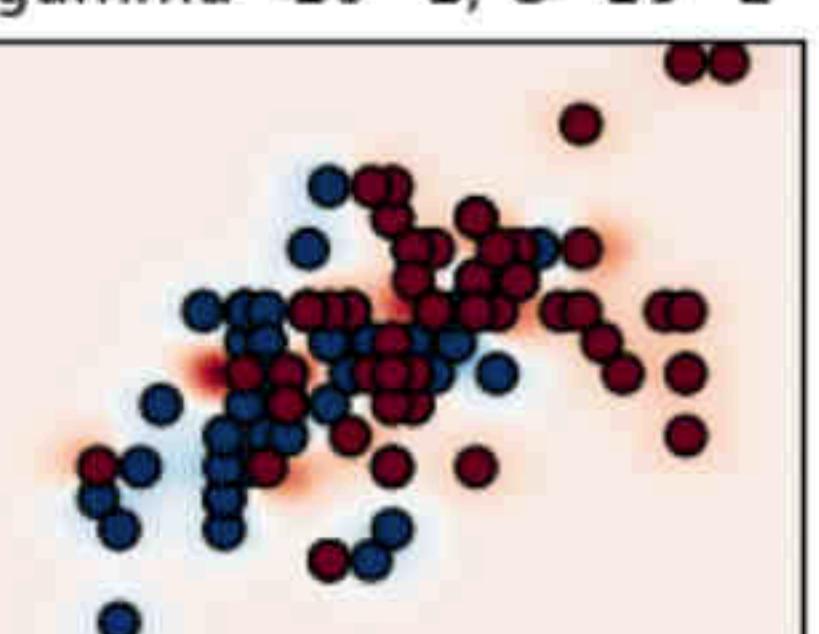
gamma=10^-1, C=10^2



gamma=10^0, C=10^2



gamma=10^1, C=10^2



scikit-learn.org/stable/auto_examples/svm/plot_rbf_parameters.html#sphx-glr-auto-examples-svm-plot-rbf-parameters-py

scikit learn

Prev Up Next

scikit-learn 1.1.1
Other versions

Please cite us if you use the software.

RBF SVM parameters
Load and prepare data set
Train classifiers
Visualization

$\gamma = \frac{1}{\text{C}}$

$C \uparrow$

overfit

gamma=10^-1, C=10^-2

gamma=10^0, C=10^-2

gamma=10^1, C=10^-2

gamma=10^-1, C=10^0

gamma=10^0, C=10^0

gamma=10^1, C=10^0

gamma=10^-1, C=10^2

gamma=10^0, C=10^2

gamma=10^1, C=10^2

Toggle Menu

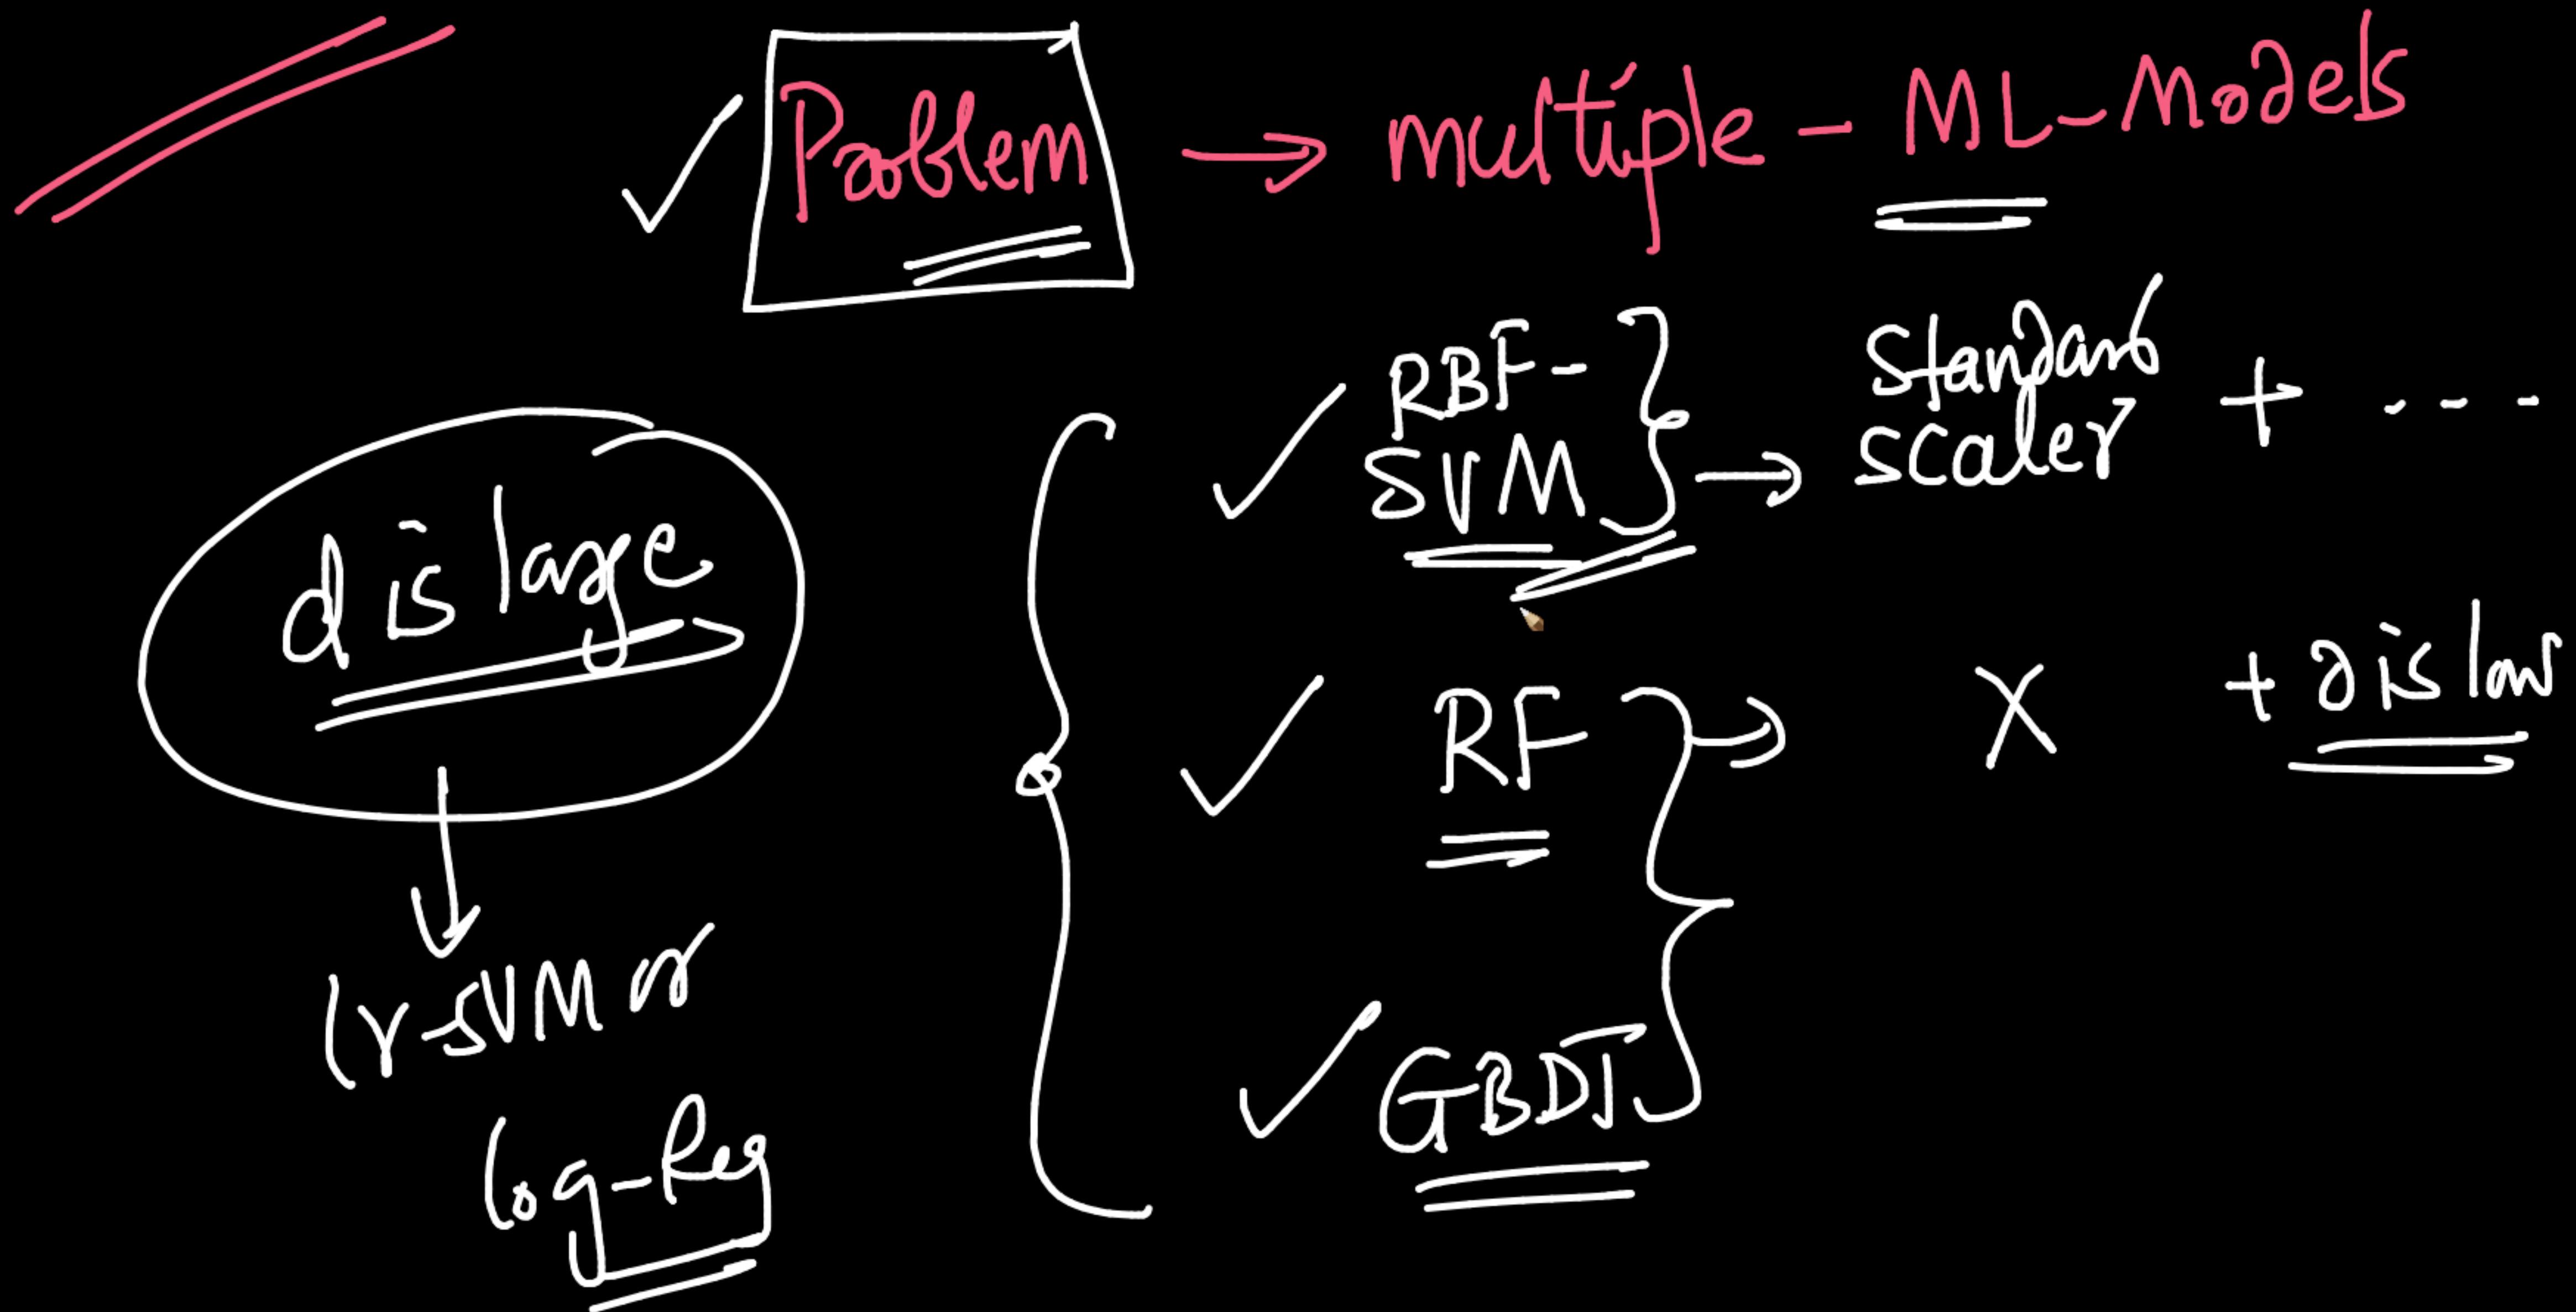
54 / 55

The figure displays a 3x3 grid of plots, each showing a 2D dataset with two classes (blue and red) and their corresponding RBF SVM decision boundaries. The plots are arranged in three rows and three columns. Each plot includes its parameter values in the title:

- Top row: $\gamma = 10^{-1}$, $C = 10^{-2}$; $\gamma = 10^0$, $C = 10^{-2}$; $\gamma = 10^1$, $C = 10^{-2}$
- Middle row: $\gamma = 10^{-1}$, $C = 10^0$; $\gamma = 10^0$, $C = 10^0$; $\gamma = 10^1$, $C = 10^0$
- Bottom row: $\gamma = 10^{-1}$, $C = 10^2$; $\gamma = 10^0$, $C = 10^2$; $\gamma = 10^1$, $C = 10^2$

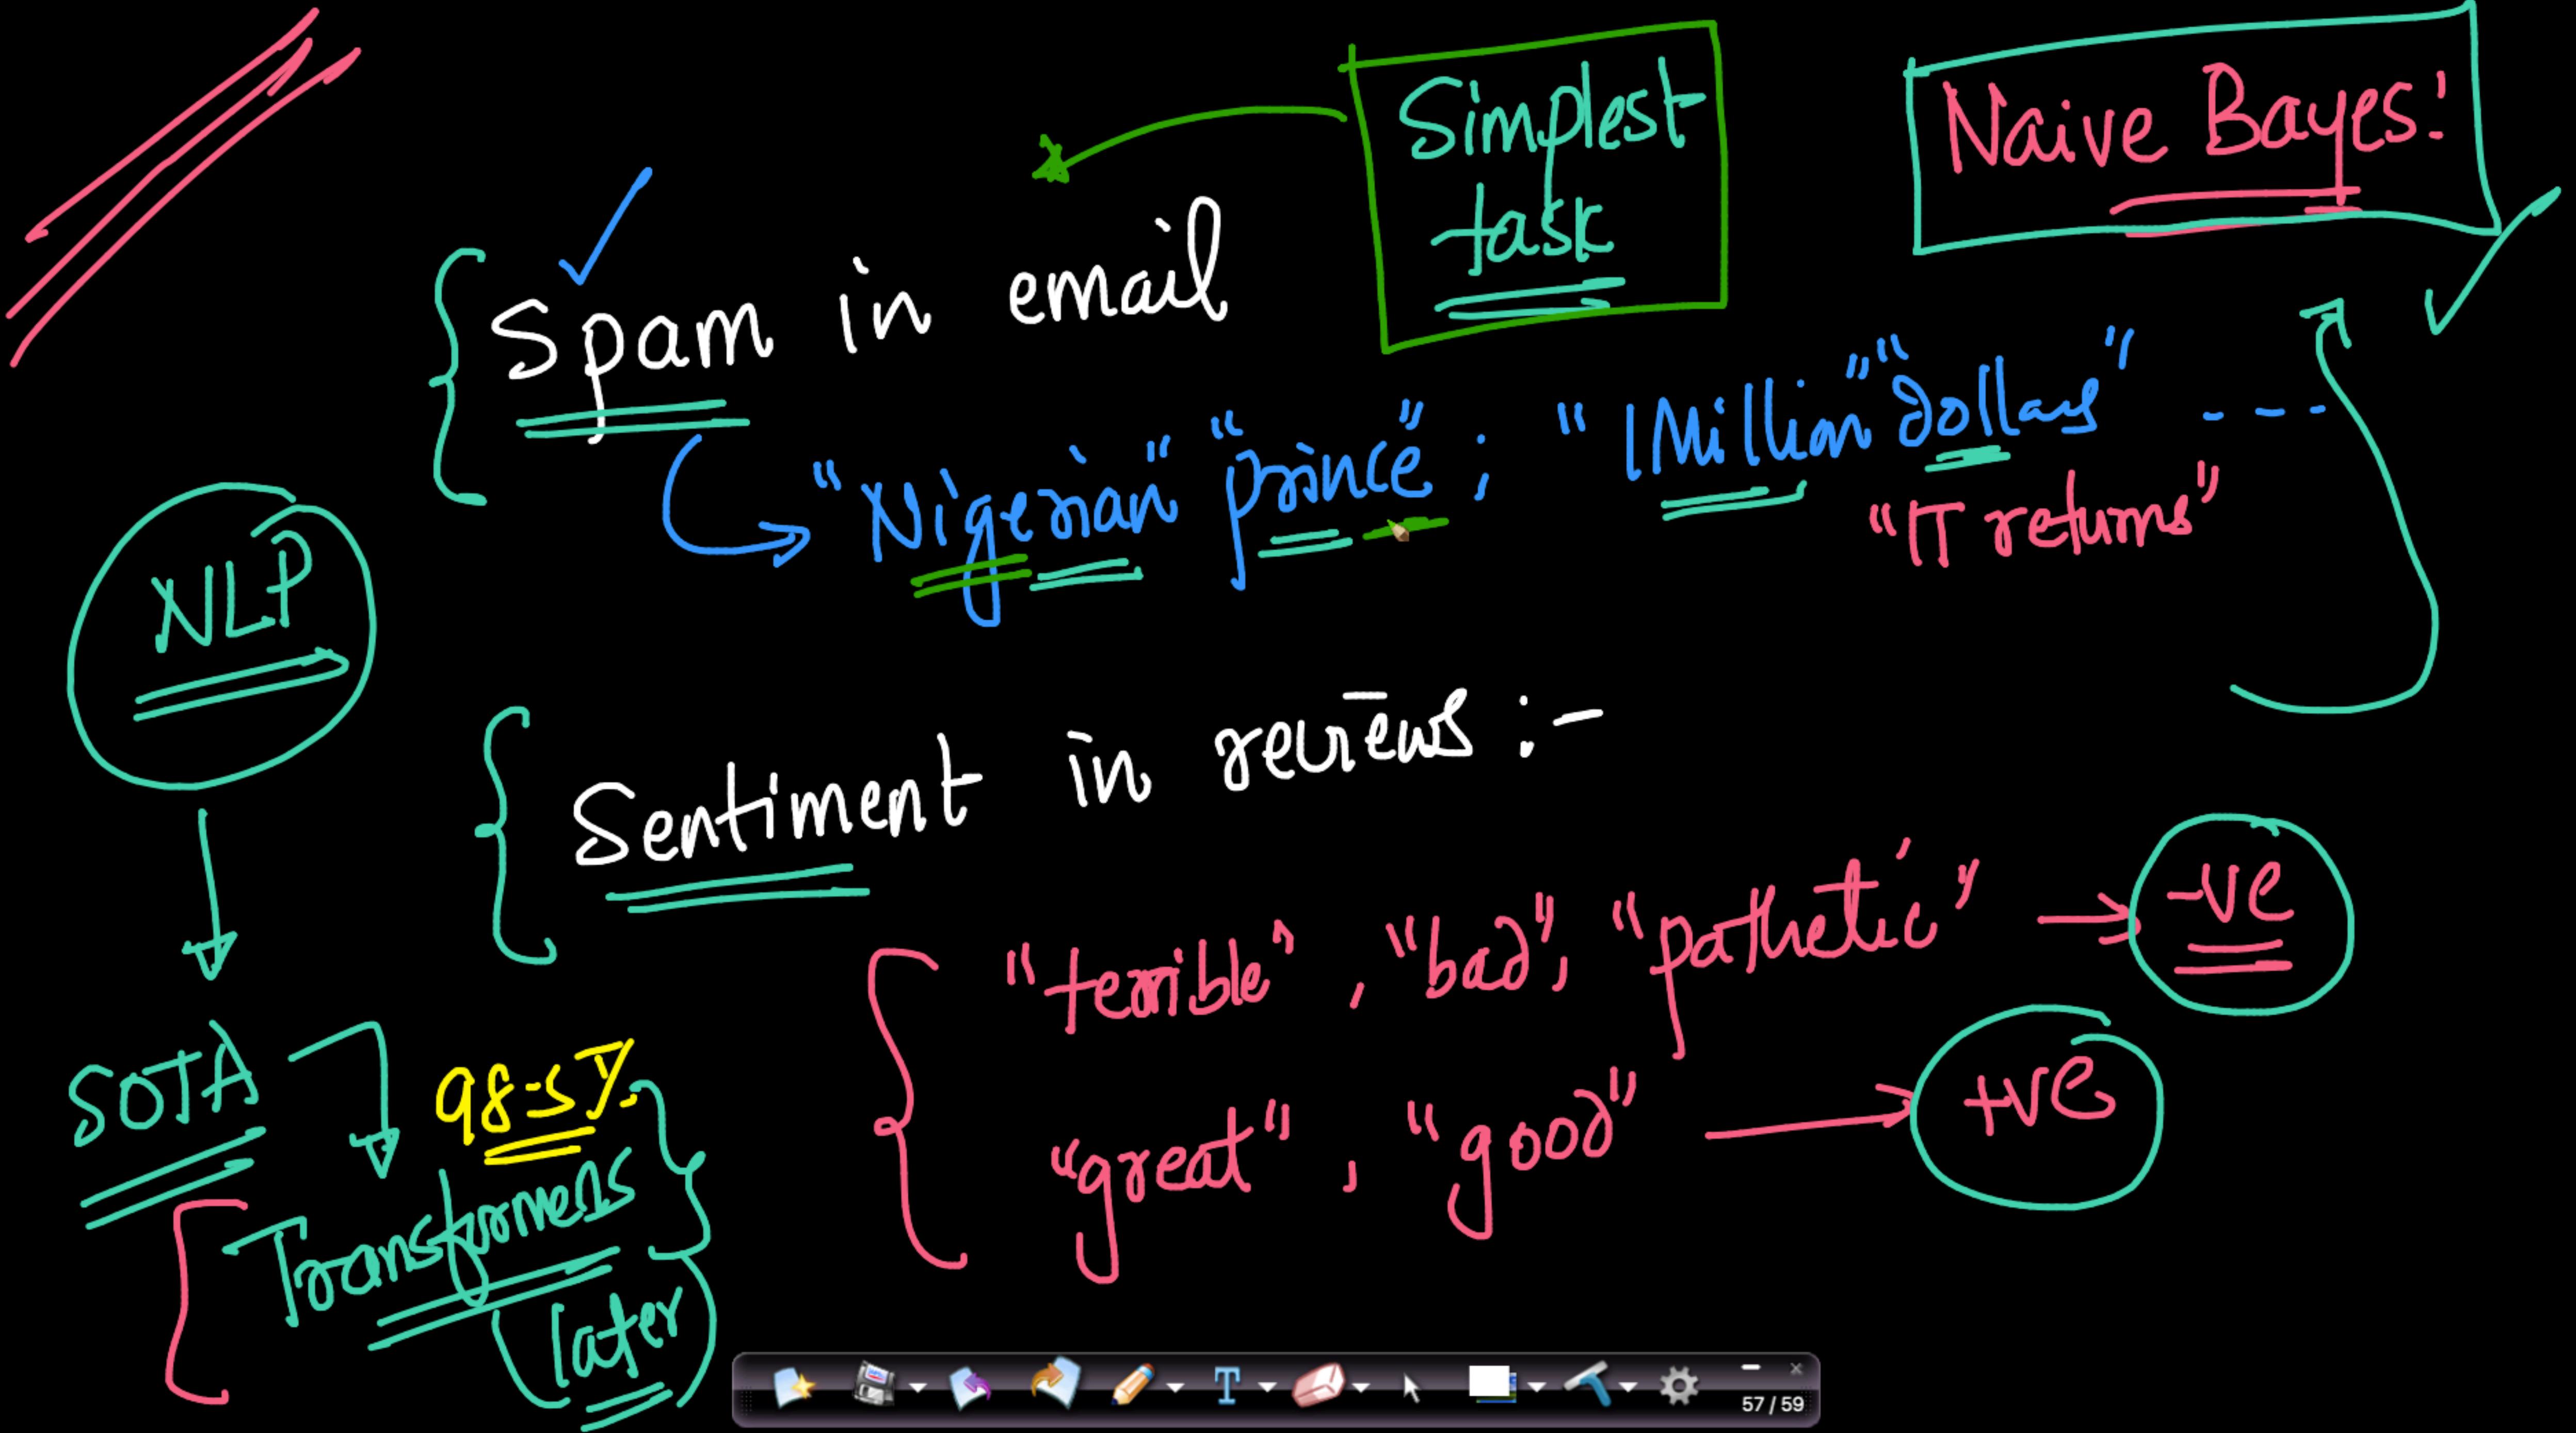
Handwritten annotations are present on the left side of the grid:

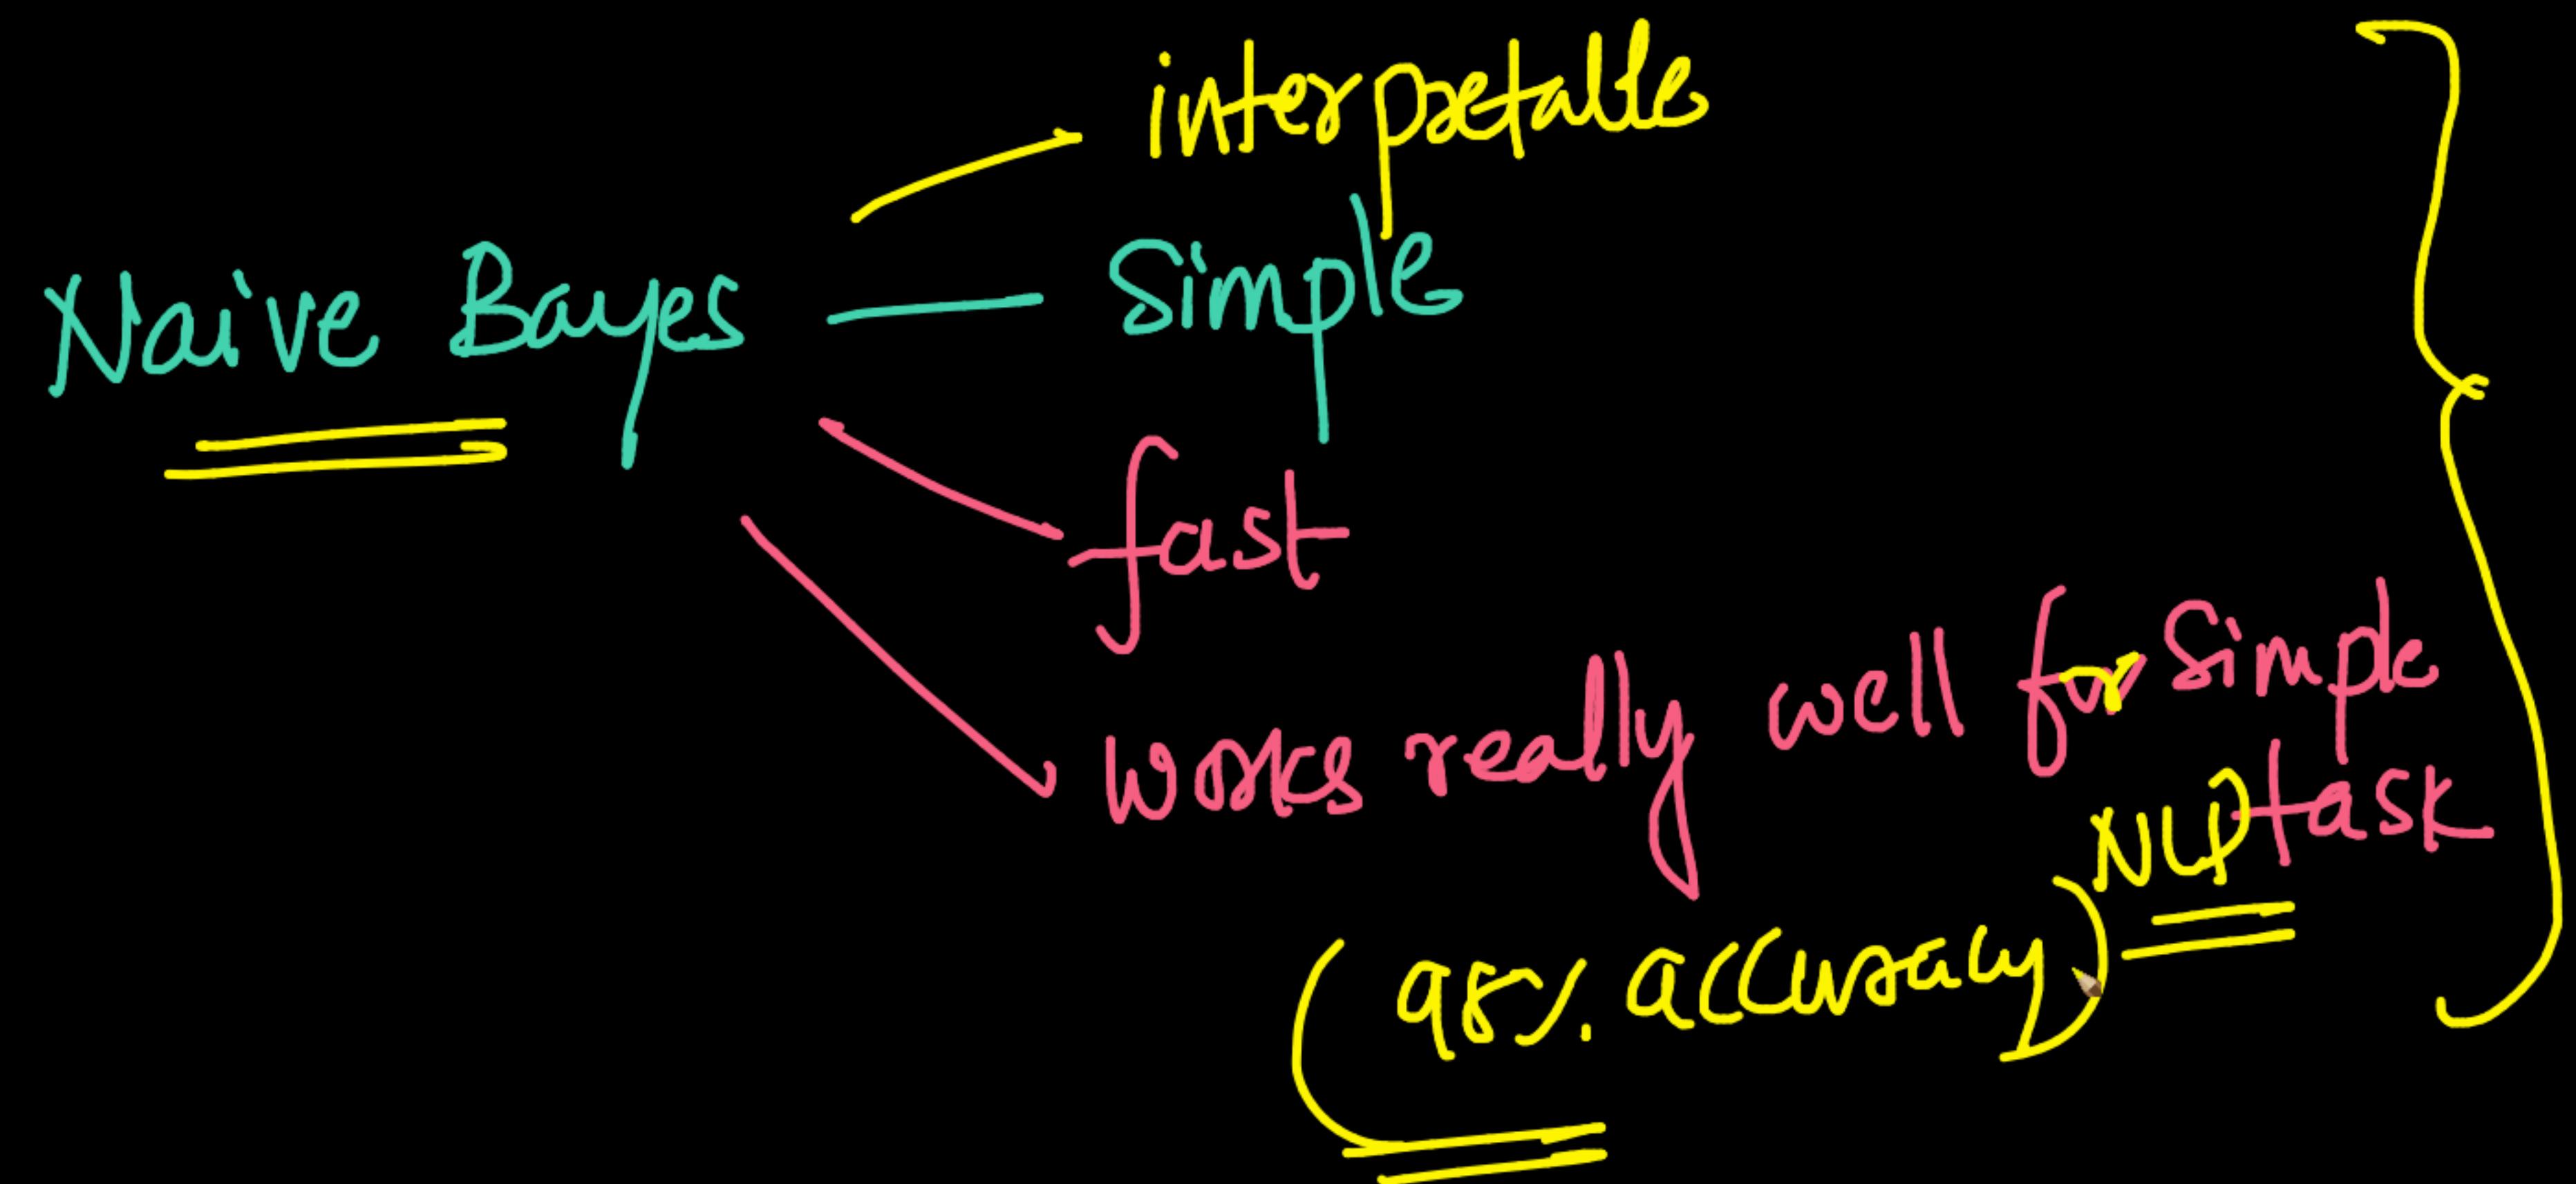
- A large red circle encloses the first column of plots ($\gamma = 10^{-1}$, $\gamma = 10^0$, $\gamma = 10^1$).
- A red arrow labeled "C ↑" points upwards between the first and second columns.
- The word "overfit" is written below the first column.
- A red circle contains the formula $\gamma = \frac{1}{C}$.



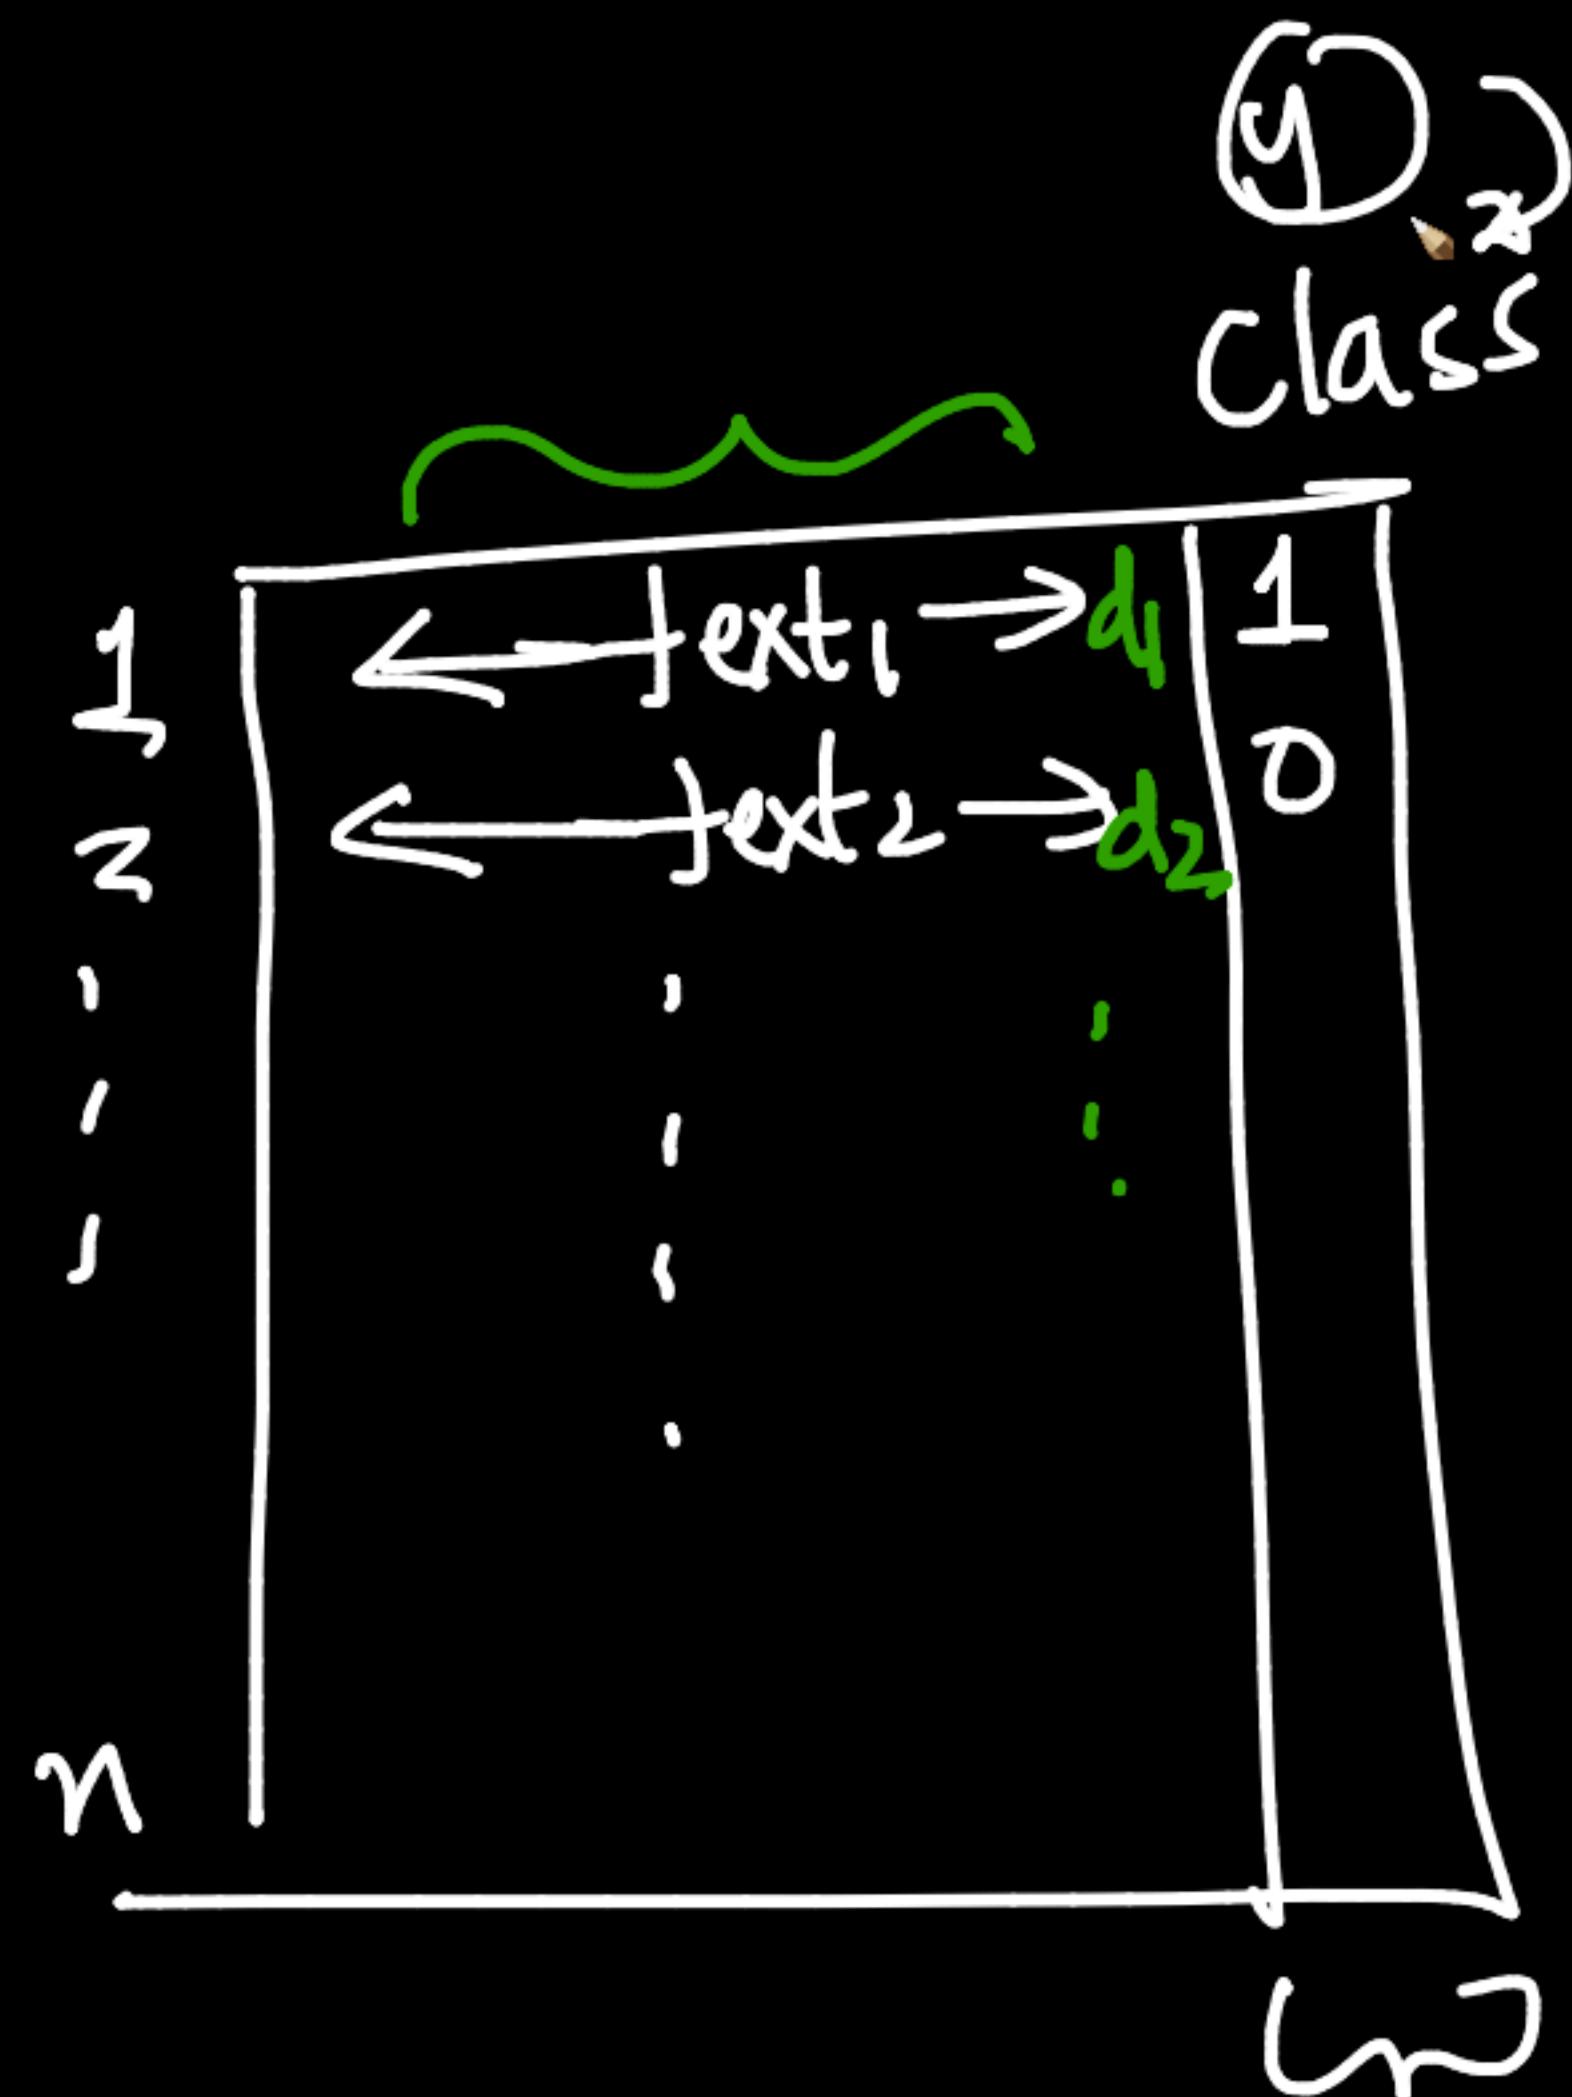
11:00

Naive Bayes





Data:-



text → remove

Stop-woods
of the in-t-.

Three pink hand-drawn diagonal lines.

$$P(y=1 \mid \text{text})$$

$$P(y=0 \mid \text{text})$$

$$P(y=1 \mid \text{text}) = P(y=1 \mid w_1, w_2, w_3, \dots, w_d)$$

$$= \frac{P(y=1, w_1, w_2, w_3, \dots, w_d)}{P(w_1, w_2, w_3, \dots, w_d)}$$

$$P(w_1, w_2, w_3, \dots, w_d)$$

$$= P(w_1, w_2, \dots, w_d \mid y=1) P(y=1)$$

~~$$P(w_1, w_2, \dots, w_d)$$~~ K

$$\begin{aligned} & P(A \mid B) \\ &= P(A \cap B) \\ &= P(A \text{ and } B) \end{aligned}$$

Bayes Thm

$y=1$

$$P(y=1 | w_1, \dots, w_d) = P(w_1, \dots, w_d | y=1) P(y=1)$$

Compare

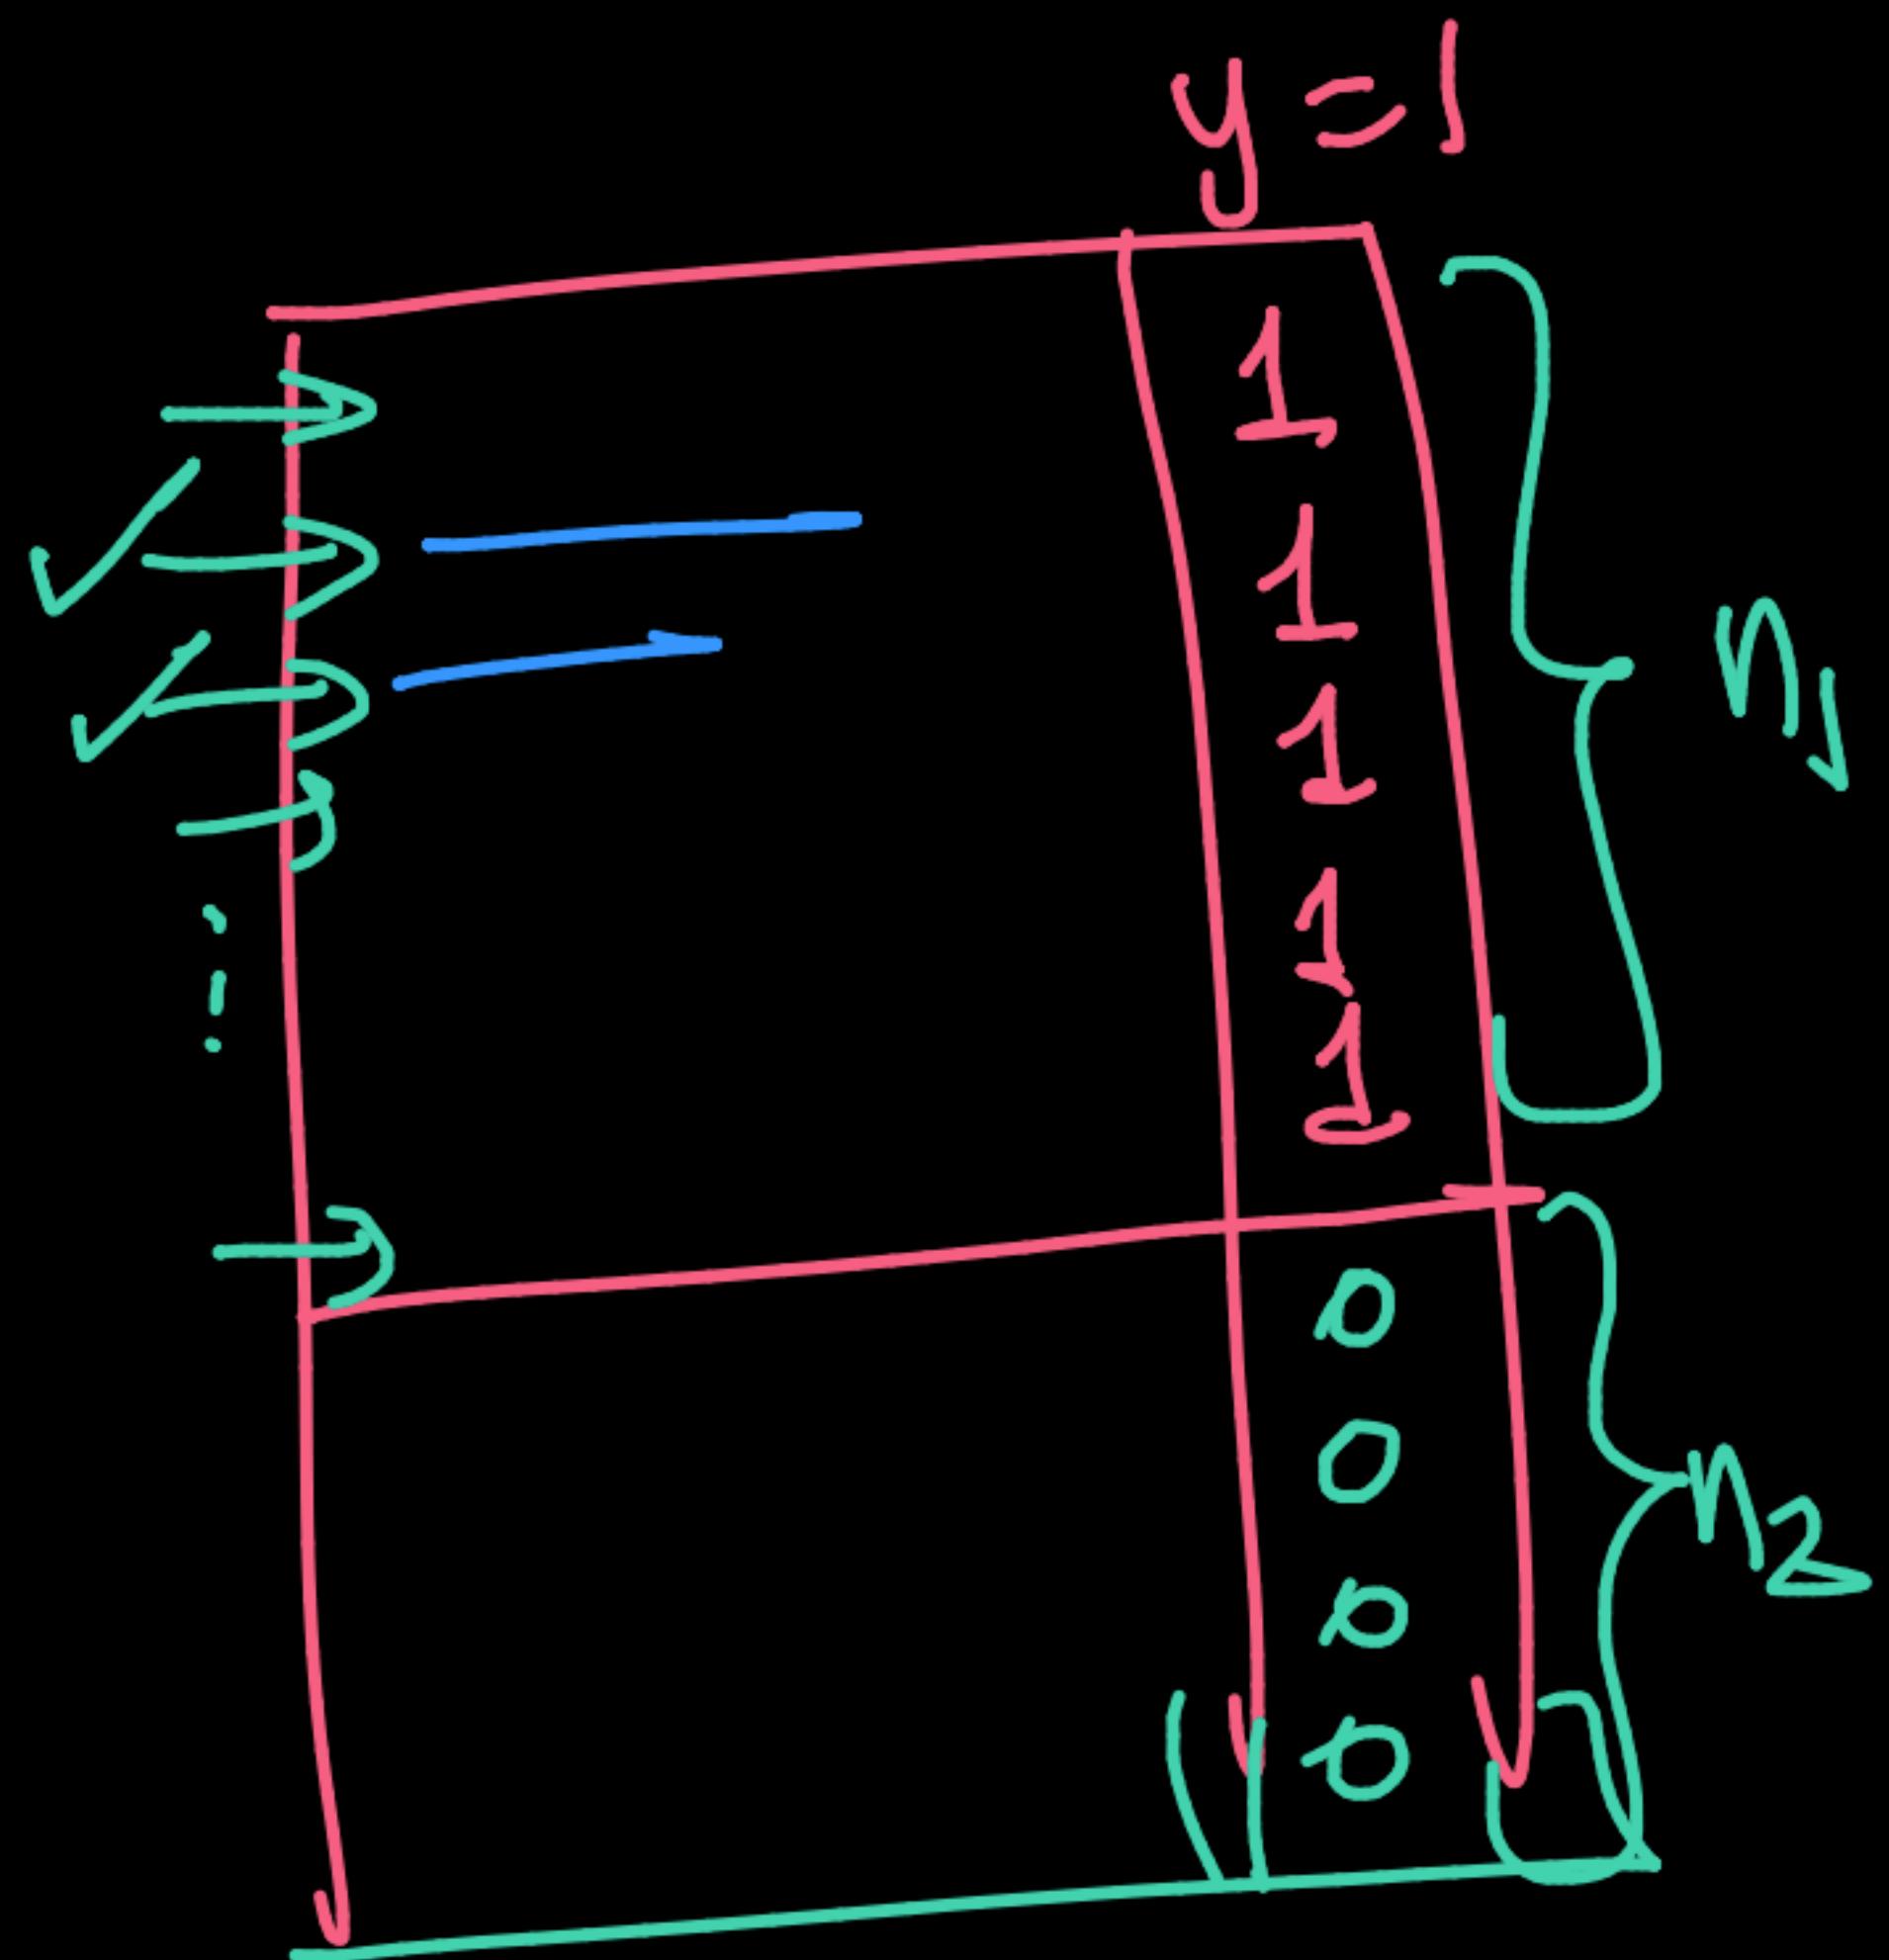
$$P(y=0 | w_1, \dots, w_d) = P(w_1, \dots, w_d | y=0) P(y=0)$$

$$P(Y=1) \rightarrow \frac{\text{# train pts with } Y_i=1}{\text{total # of train pts}}$$

$$P(Y=0) \rightarrow \frac{\text{# train pts with } Y_i=0}{\text{total # of train pts}}$$

$$P(w_1, w_2, \dots, w_d | y=1)$$

Naive



✓ [If A & B are indep \rightarrow [Indep \neq mutually exc] $\leftarrow \dots$] } defn

$$P(A, B) = P(A) \cdot P(B)$$

P(A, B)

A & B are indep

✓ { w₁ & w₂ are indep. conditioned on the class

Naive, assumption } P(w₁, w₂, | y = i) → spam
Nigerian p_{spam} = P(w₁ | y₁) · P(w₂ | y₁)

$$P(w_1, w_2, w_3, \dots, w_d | y=1)$$

Naive

$$= \underbrace{P(w_1 | y=1)}_{\text{Naive}} \cdot P(w_2 | y=1) \cdot \dots \cdot P(w_d | y=1)$$

$$\begin{aligned} P(w_1 | y=1) \\ = \frac{n_1}{n_1 + n_2} \end{aligned}$$



$$P(y=1 \mid \text{text}) = \frac{P(w_1 w_2 \dots w_d \mid y=1) p(y=1)}{k}$$

Σ : sum

\prod : product

$$= p(w_1 | y=1) p(w_2 | y=1) \dots \\ p(w_d | y=1) \cdot p(y=1) / k$$

$$= \prod_{i=1}^d p(w_i | y=1) p(y=1) / k$$

$$P(y=1 \mid \text{text}) = \prod_{i=1}^d P(w_i \mid y=1)$$

↗ likelihoods

↙ Class priors

↖ Compose

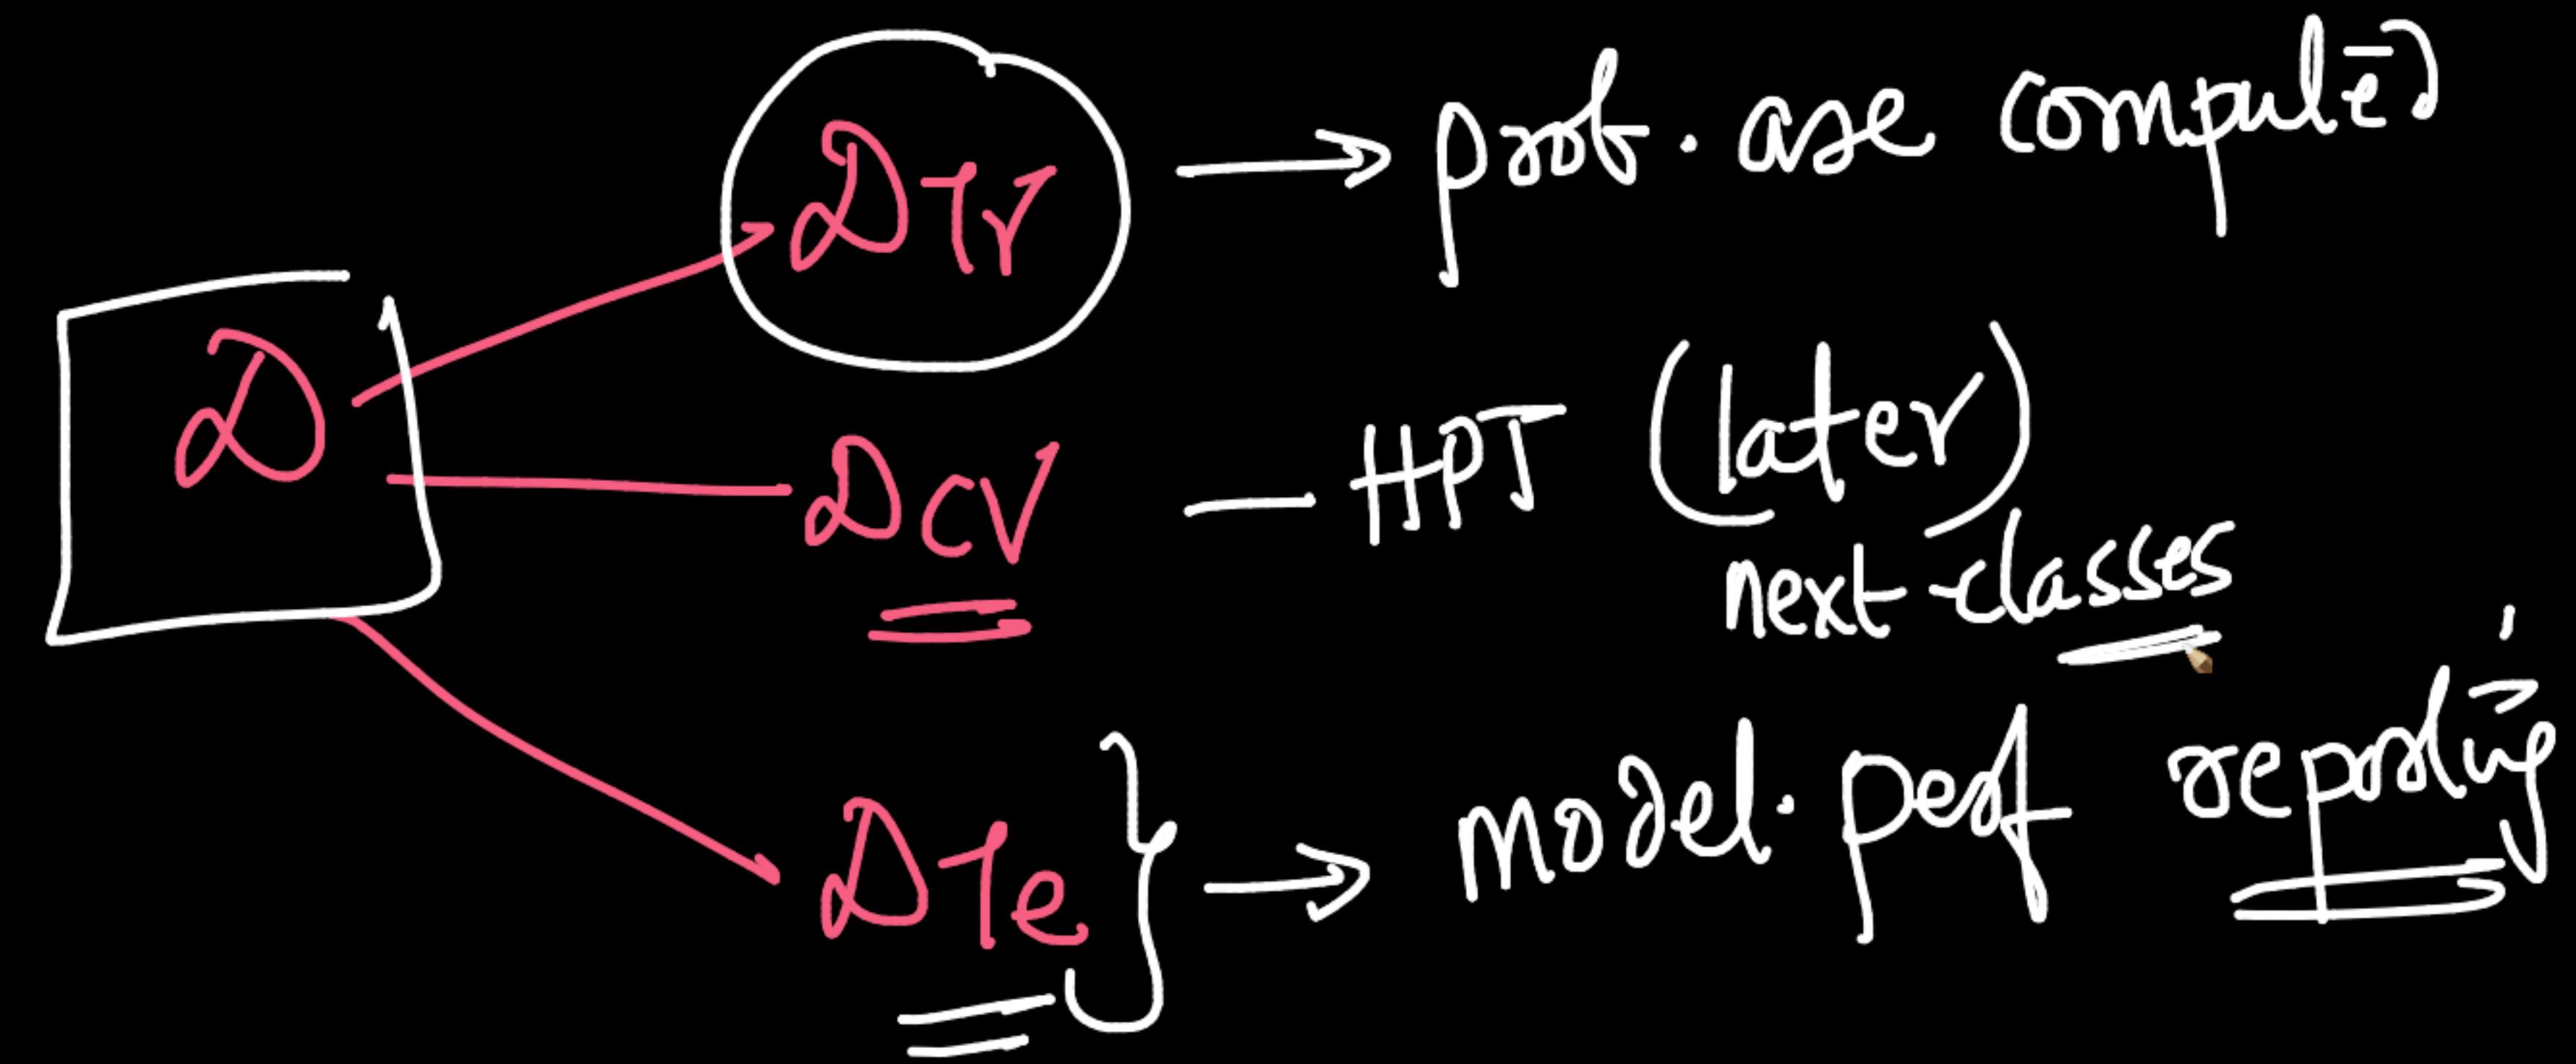
$$P(y=0 \mid \text{text}) = \prod_{i=1}^d P(w_i \mid y=0)$$

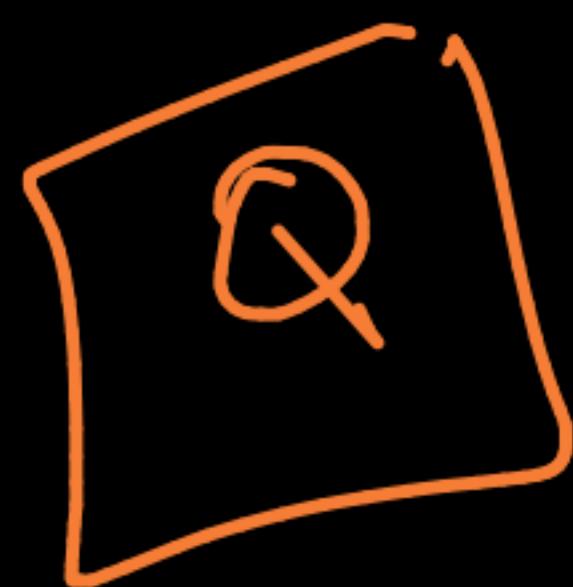
↗ likelihoods

↙ Class priors

↖ Compose

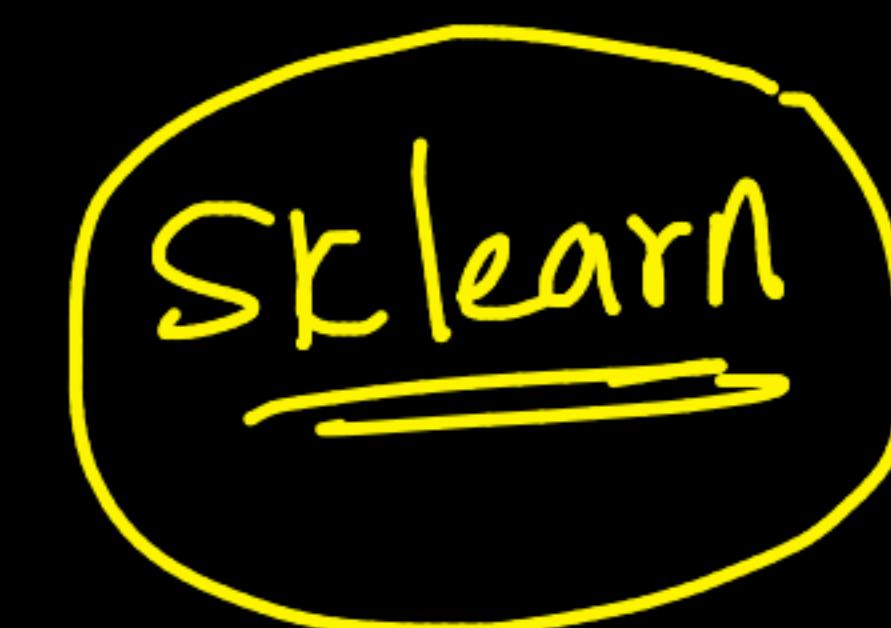
Naive Bayes





log-reg: - $\hat{y} = \frac{1}{1 + e^{-x}}$

SVM \rightarrow $+1, -1$



Label-encoding

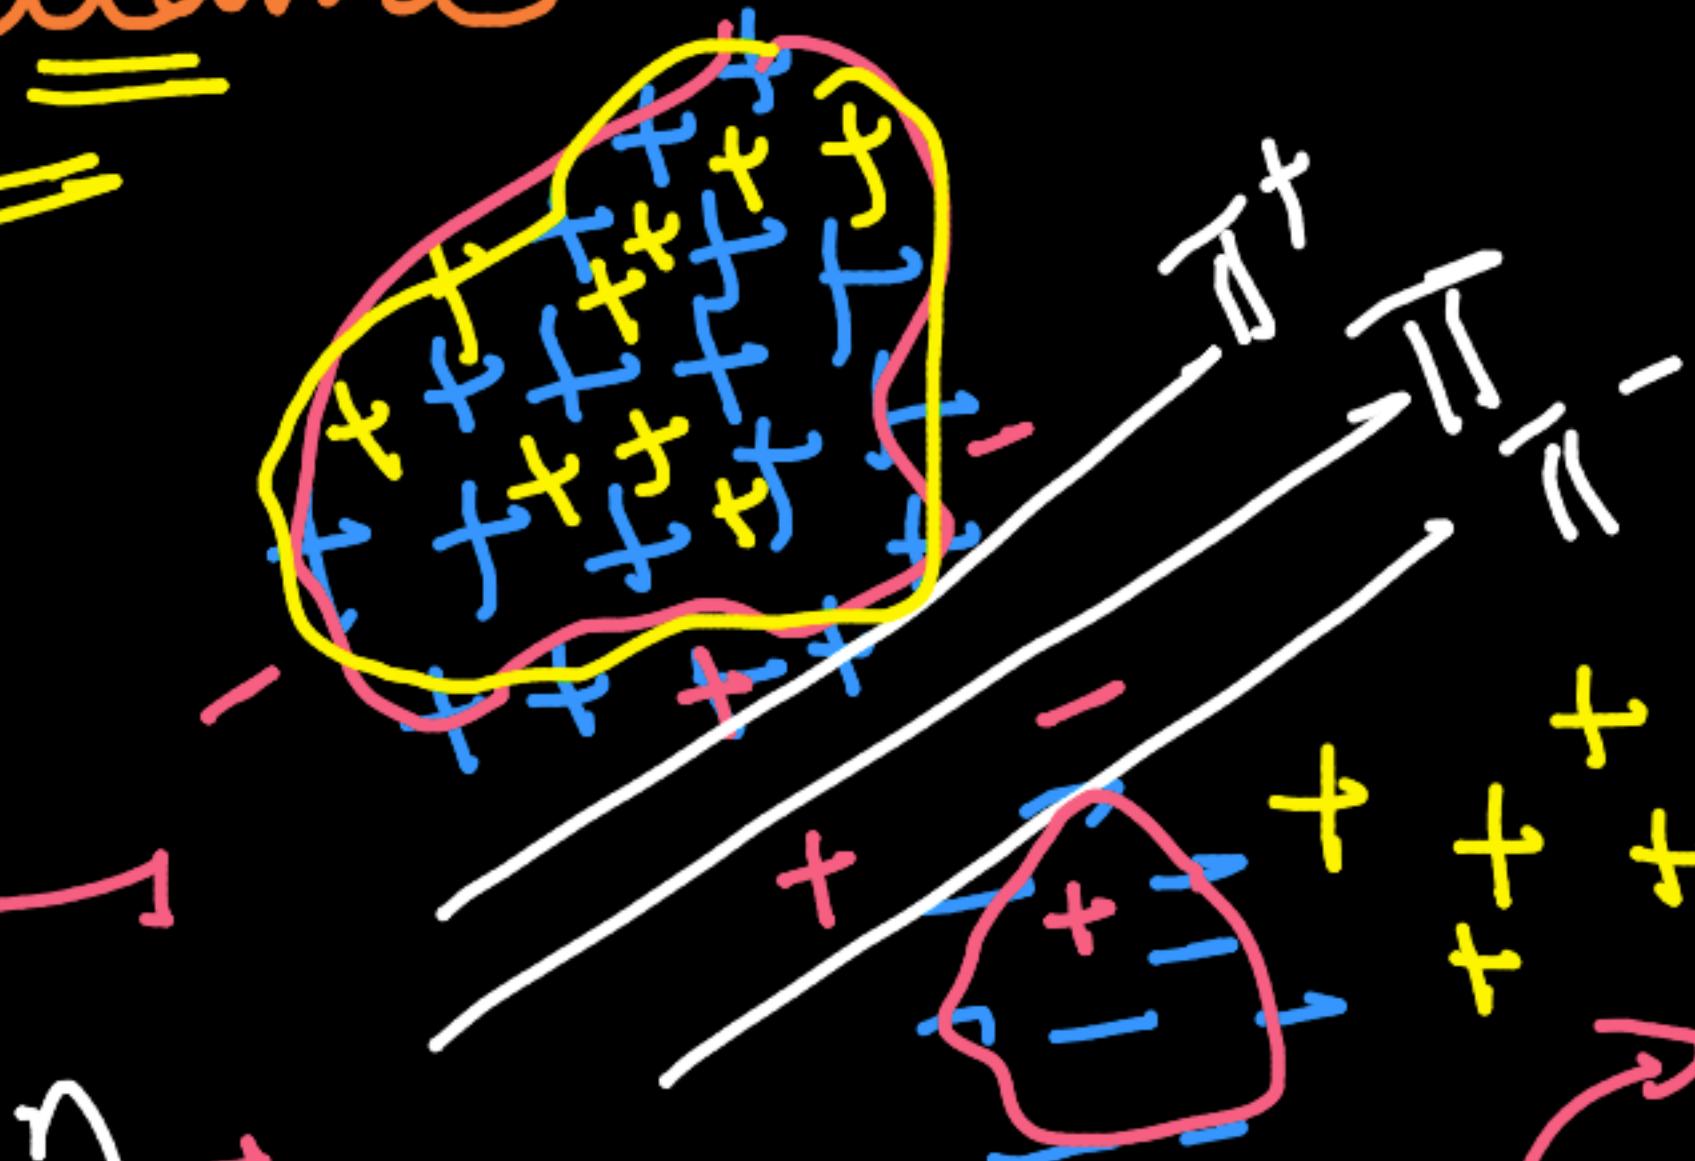
Q

SVM: $\left\{ \begin{array}{l} \text{dow} \\ \text{imbalance} \end{array} \right.$

Ly-SVM:

$$\min_{w,b} \frac{\|w\|}{2} + C \sum_{i=1}^n \xi_i$$

loss



hard Margin
soft Margin

May | may not
#SVs from each class
are same

imbalance → class weights ...
↳ rebalance ...

en.wikipedia.org/wiki/Radial_basis_function#~text=A%20radial%20basis%20function%20(RBF,property%20is%20a%20radial%20function.

Not logged in Talk Contributions Create account Log in

Article Talk Read Edit View history Search Wikipedia

Radial basis function

From Wikipedia, the free encyclopedia

A **radial basis function (RBF)** is a **real-valued function** φ whose value depends only on the distance between the input and some fixed point, either the **origin**, so that $\varphi(\mathbf{x}) = \hat{\varphi}(\|\mathbf{x}\|)$, or some other fixed point \mathbf{c} , called a **center**, so that $\varphi(\mathbf{x}) = \hat{\varphi}(\|\mathbf{x} - \mathbf{c}\|)$. Any function φ that satisfies the property $\varphi(\mathbf{x}) = \hat{\varphi}(\|\mathbf{x}\|)$ is a **radial function**. The distance is usually **Euclidean distance**, although other **metrics** are sometimes used. They are often used as a collection $\{\varphi_k\}_k$ which forms a **basis** for some **function space** of interest, hence the name.

Sums of radial basis functions are typically used to **approximate given functions**. This approximation process can also be interpreted as a simple kind of **neural network**; this was the context in which they were originally applied to machine learning, in work by **David Broomhead** and **David Lowe** in 1988,^{[1][2]} which stemmed from **Michael J. D. Powell**'s seminal research from 1977.^{[3][4][5]} RBFs are also used as a **kernel** in **support vector classification**.^[6] The technique has proven effective and flexible enough that radial basis functions are now applied in a variety of engineering applications.^{[7][8]}

Contents [hide]

- 1 Definition
 - 1.1 Examples
- 2 Approximation
- 3 RBF Network
- 4 RBFs for PDEs
- 5 See also
- 6 References
- 7 Further reading

Definition [edit]



SMOTE
D_{train} → Model

{ metric → F1-Score
D_{test} → Model → f1-score

WIKIPEDIA
The Free EncyclopediaArticle  Talk

Read

Edit

View history

Search Wikipedia



Radial basis function

From Wikipedia, the free encyclopedia

A **radial basis function (RBF)** is a [real-valued function](#) φ whose value depends only on the distance between the input and some fixed point, either the [origin](#), so that $\varphi(\mathbf{x}) = \hat{\varphi}(\|\mathbf{x}\|)$, or some other fixed point \mathbf{c} , called a *center*, so that $\varphi(\mathbf{x}) = \hat{\varphi}(\|\mathbf{x} - \mathbf{c}\|)$. Any function φ that satisfies the property $\varphi(\mathbf{x}) = \hat{\varphi}(\|\mathbf{x}\|)$ is a [radial function](#). The distance is usually [Euclidean distance](#), although other [metrics](#) are sometimes used. They are often used as a collection $\{\varphi_k\}_k$ which forms a [basis](#) for some [function space](#) of interest, hence the name.

Sums of radial basis functions are typically used to [approximate given functions](#). This approximation process can also be interpreted as a simple kind of [neural network](#); this was the context in which they were originally applied to machine learning, in work by [David Broomhead](#) and [David Lowe](#) in 1988,^{[1][2]} which stemmed from [Michael J. D. Powell's](#) seminal research from 1977.^{[3][4][5]} RBFs are also used as a [kernel](#) in [support vector classification](#).^[6] The technique has proven effective and flexible enough that radial basis functions are now applied in a variety of engineering applications.^{[7][8]}

Contents [hide]

1 Definition

1.1 Examples

2 Approximation

3 RBF Network

4 RBFs for PDEs

5 See also



77/77