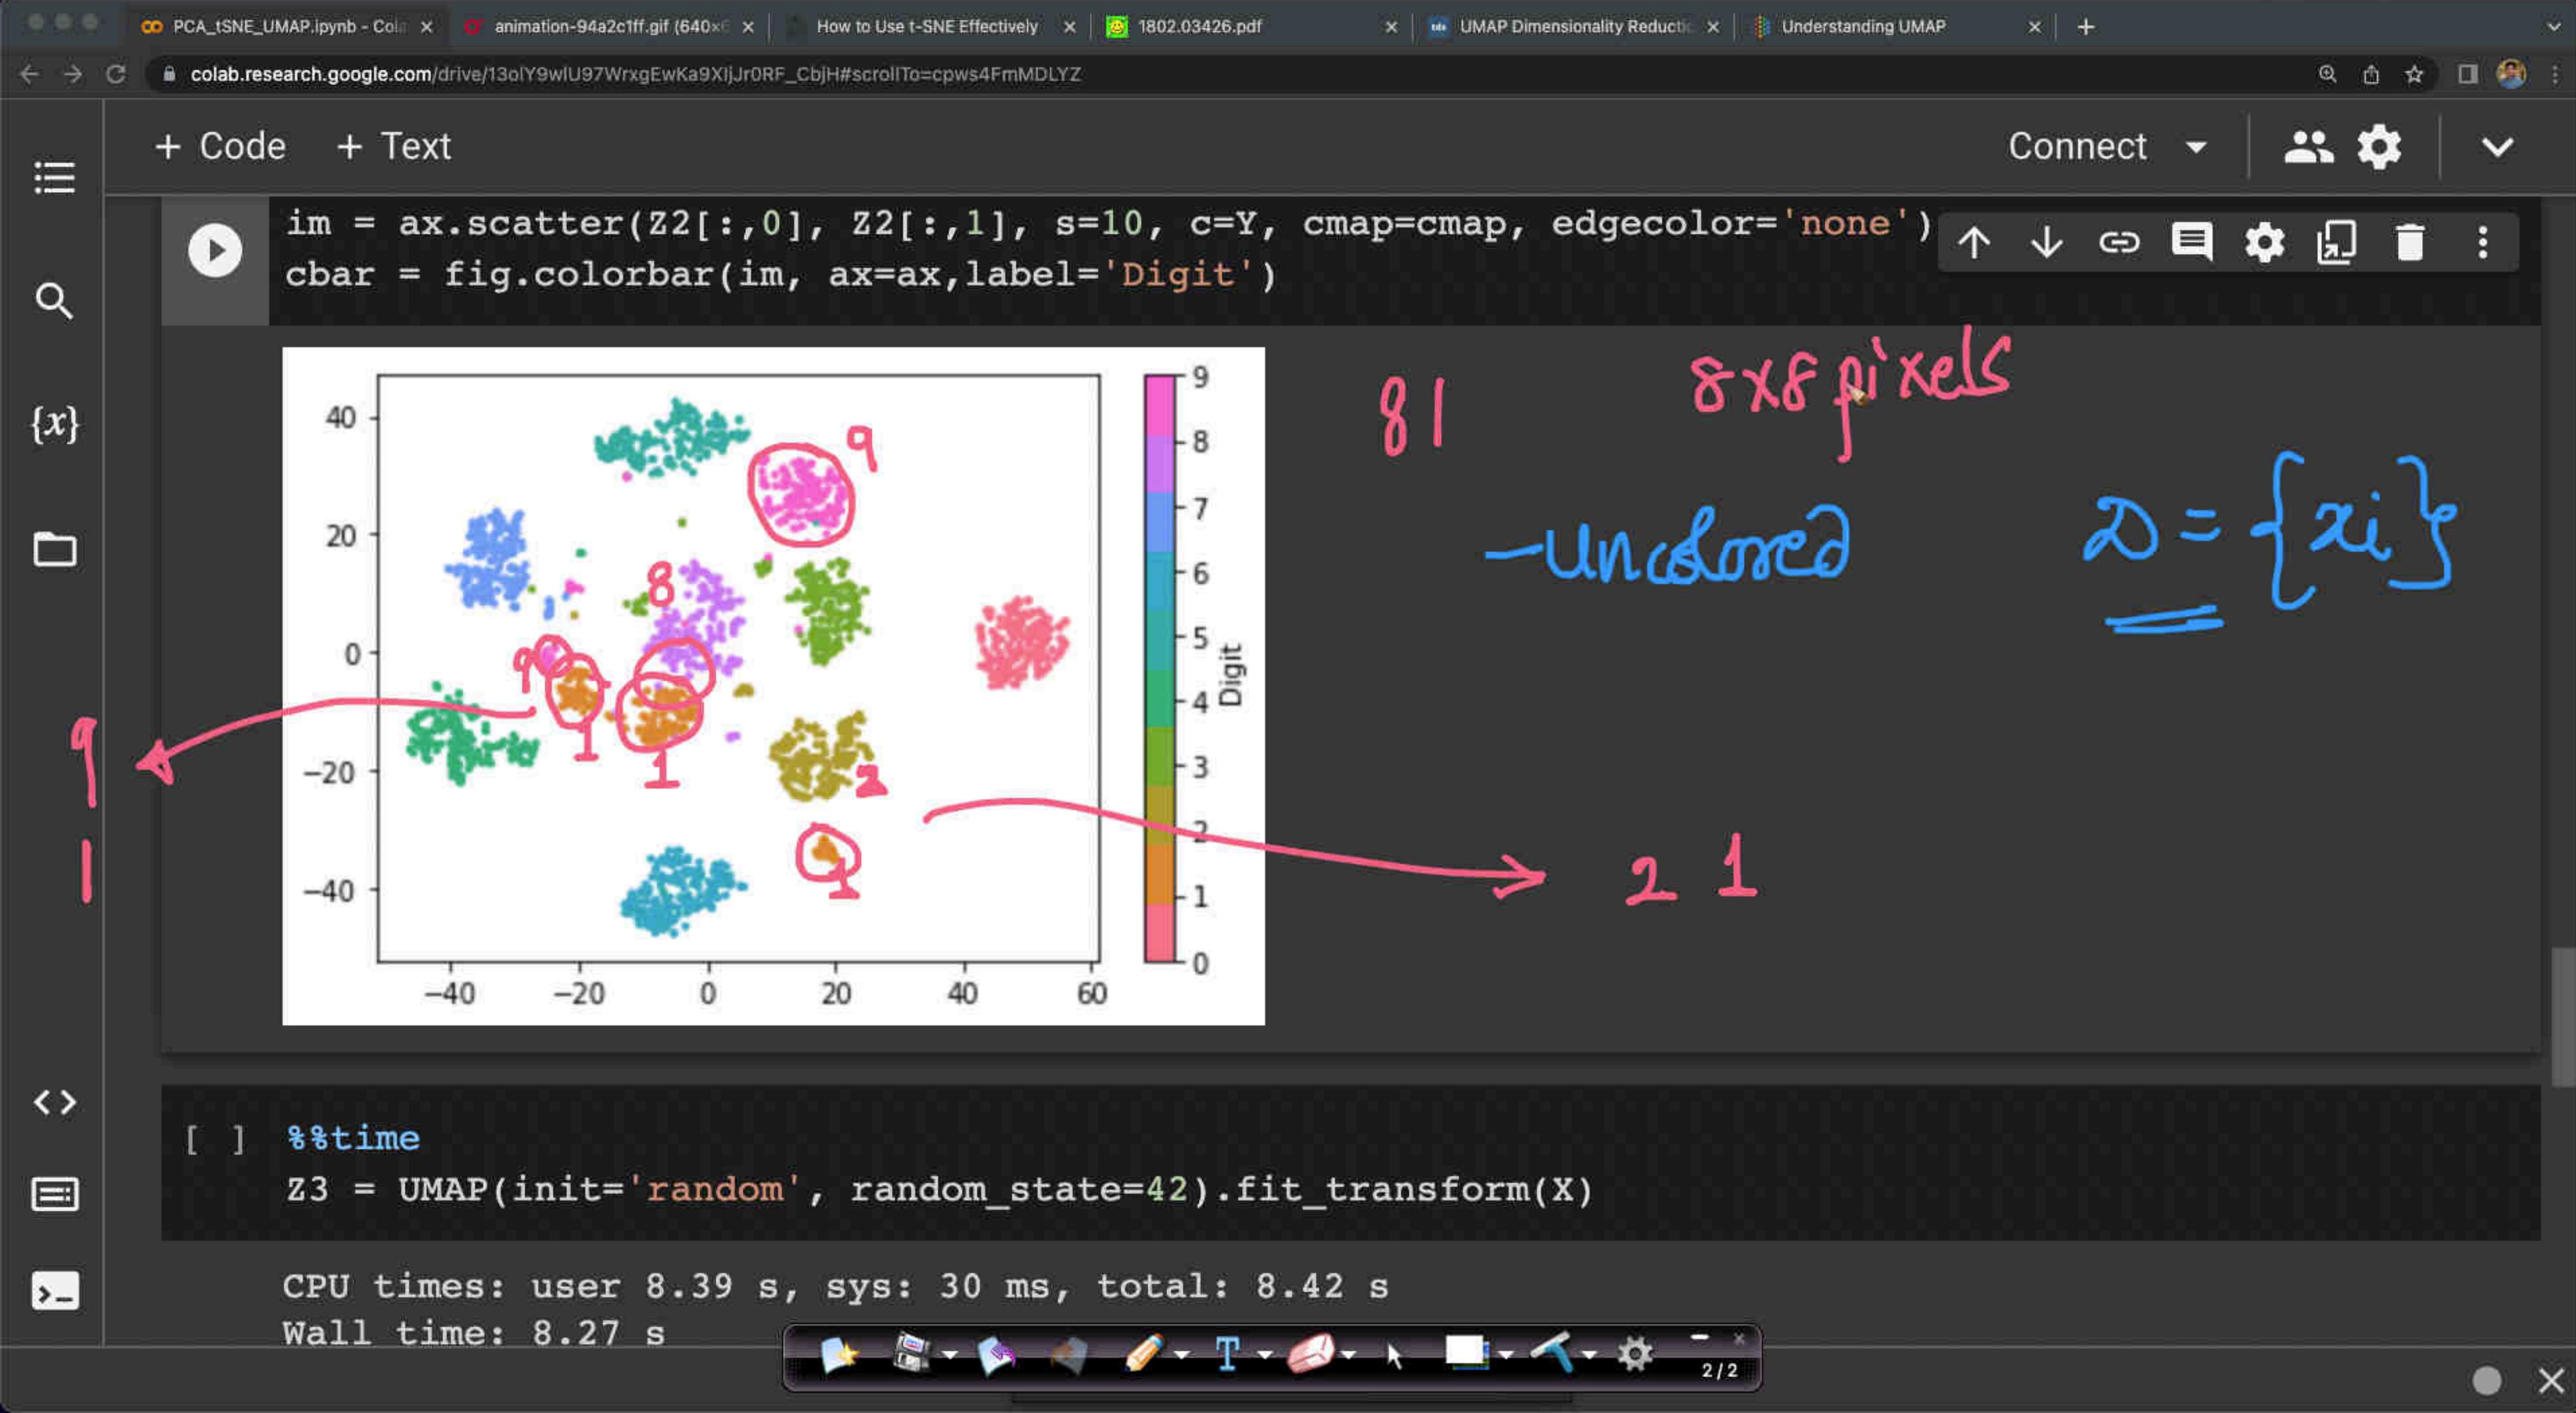
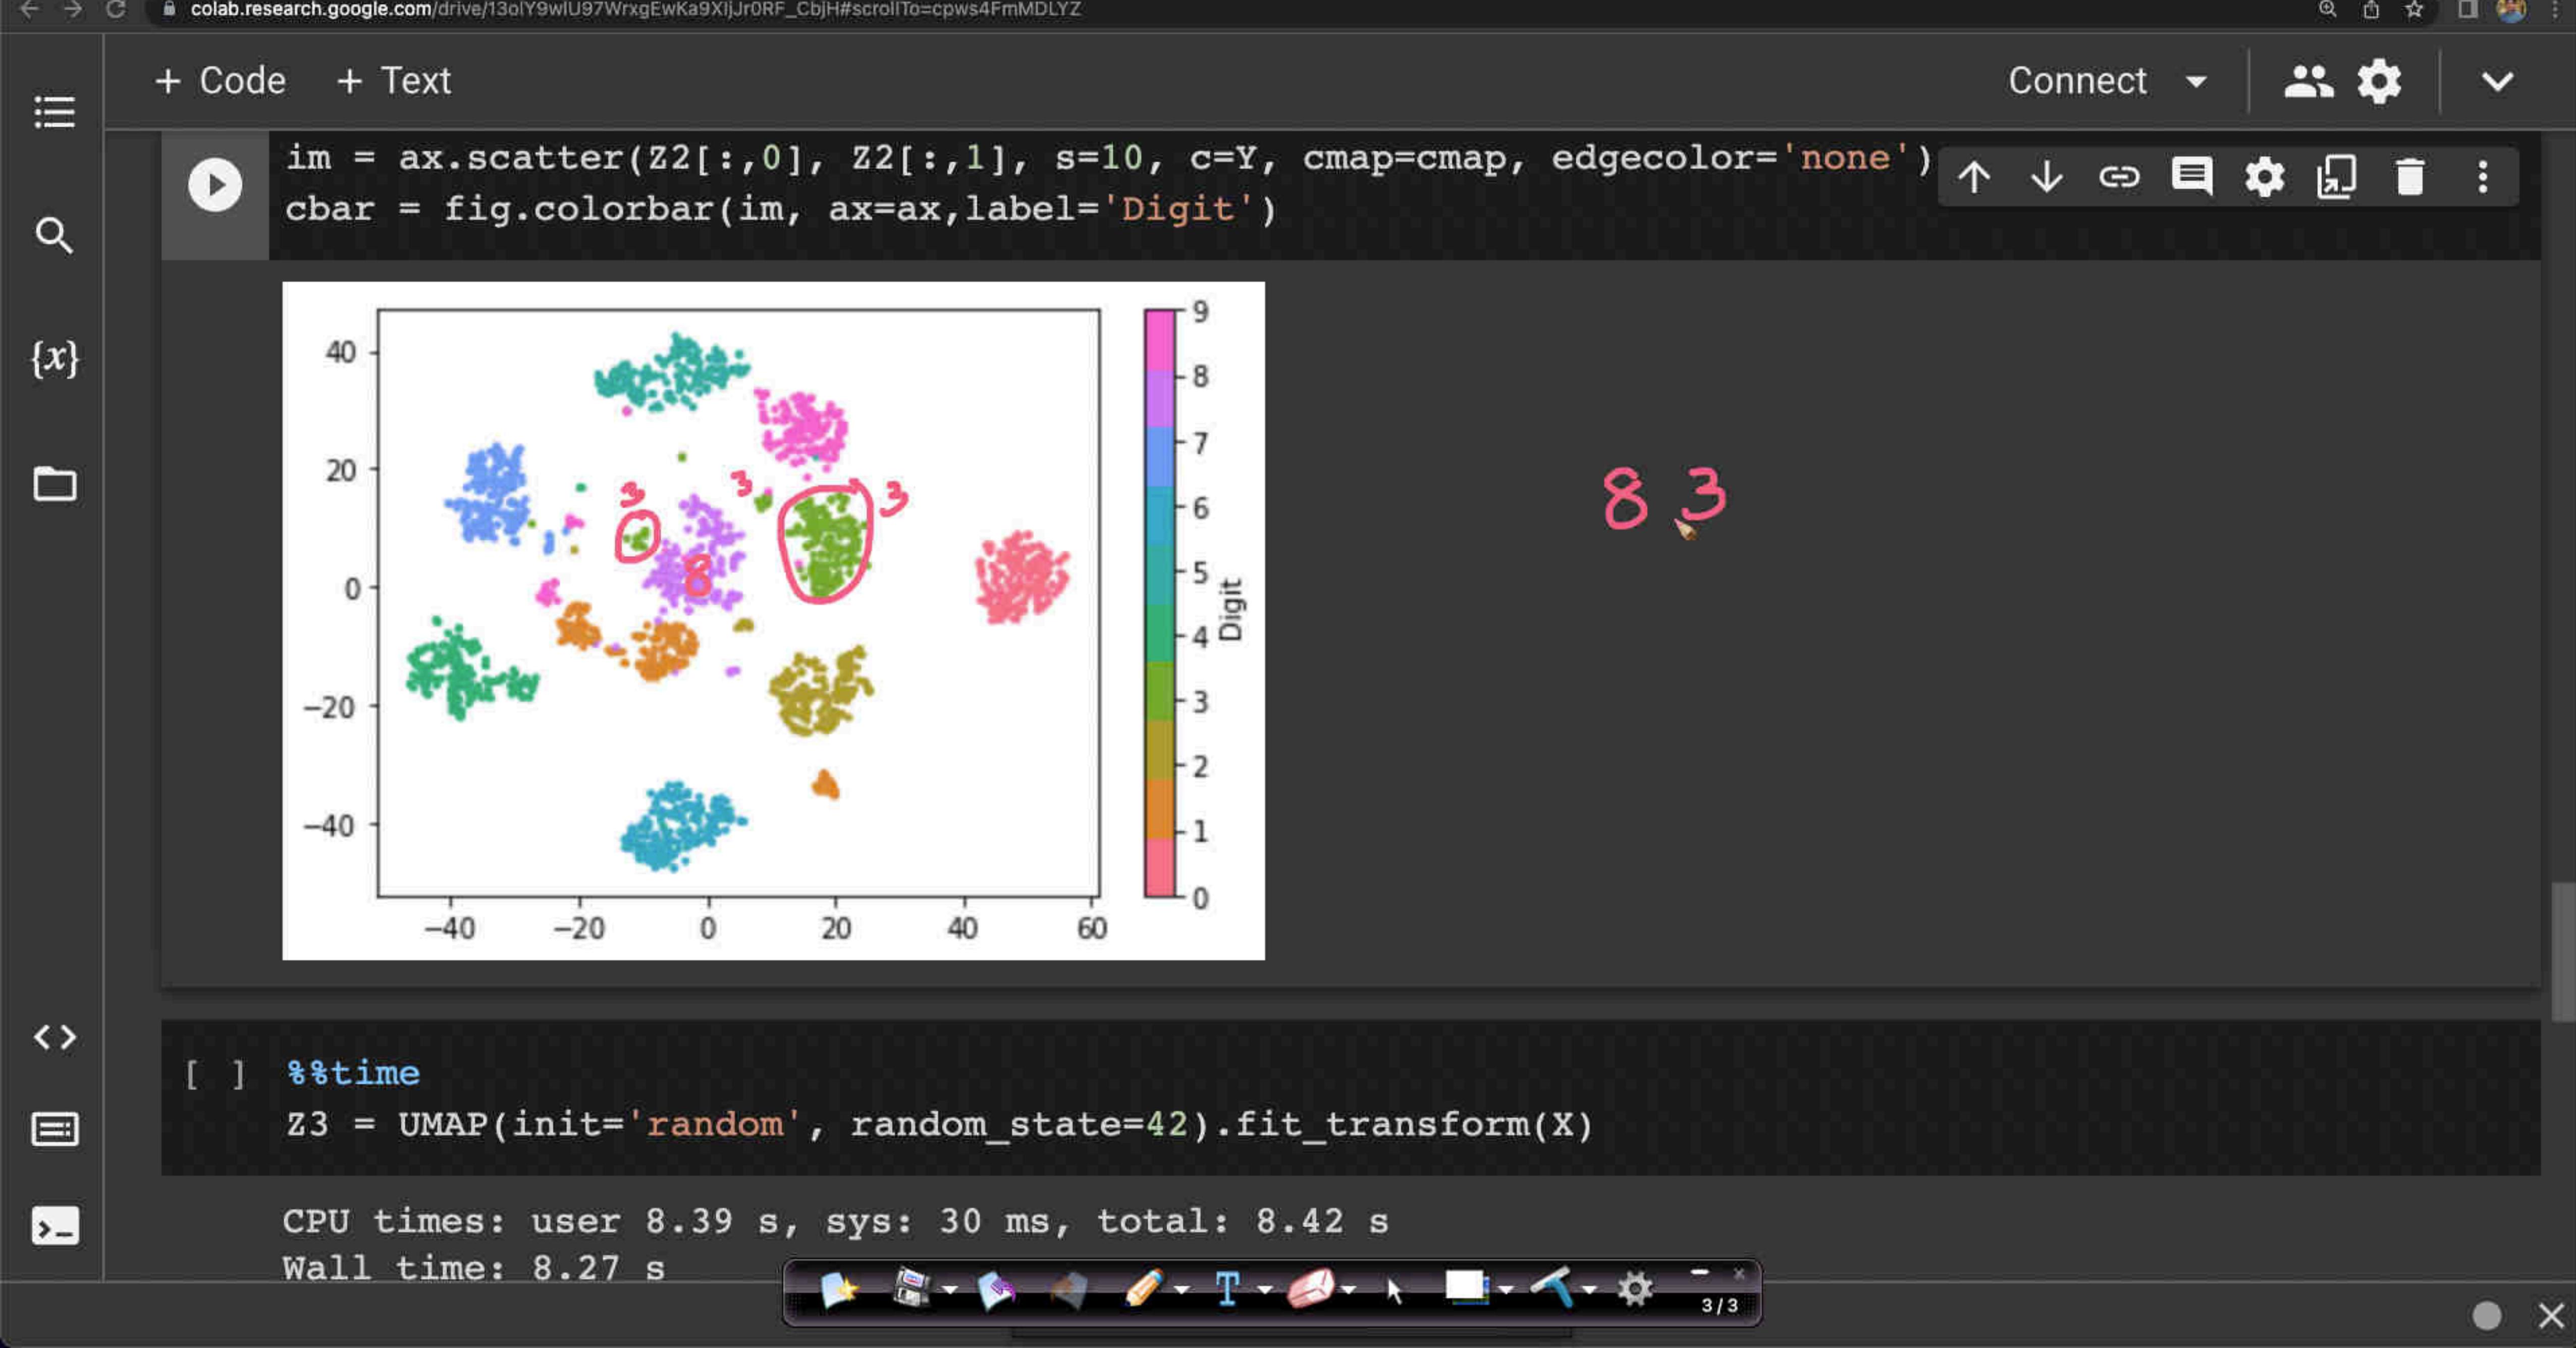
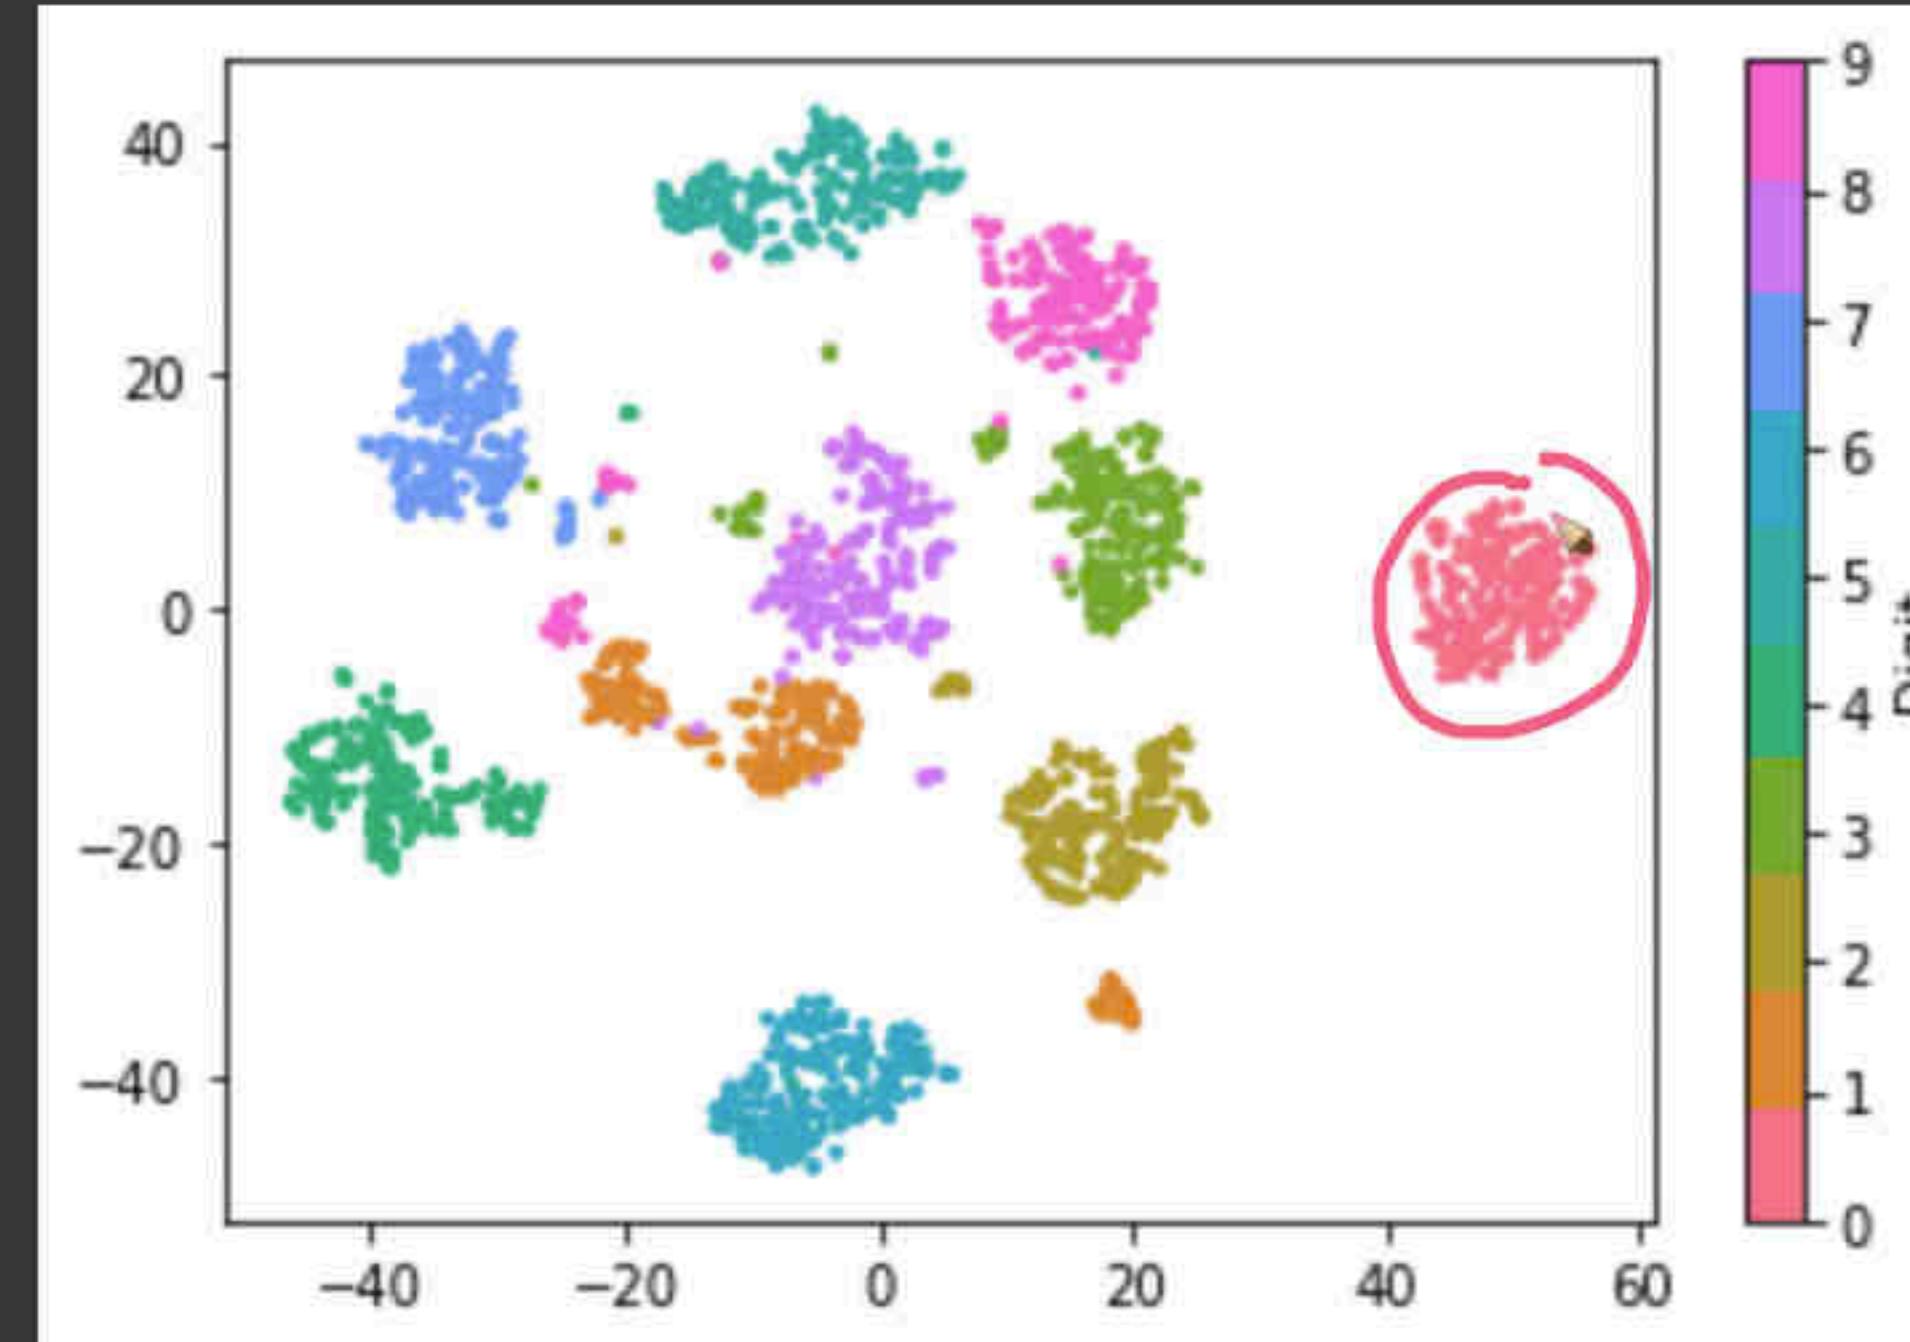
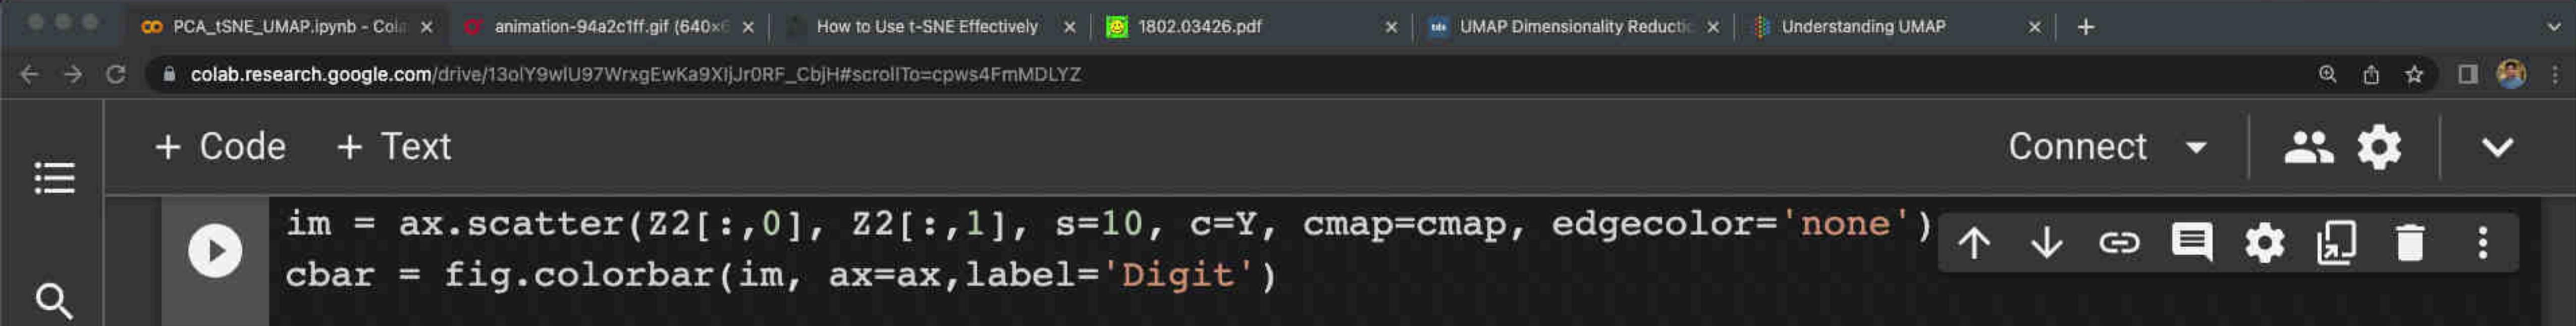


# Topics

- ✓ ① tSNE → practical aspects ] + code
- [ ② UMAP → Math + practical aspects
- { ③ Recommender-Systems : Content - based ✓
- ④ RecSys : Collaborative ✓
- ⑤ RecSys : Matrix factorization = ✓



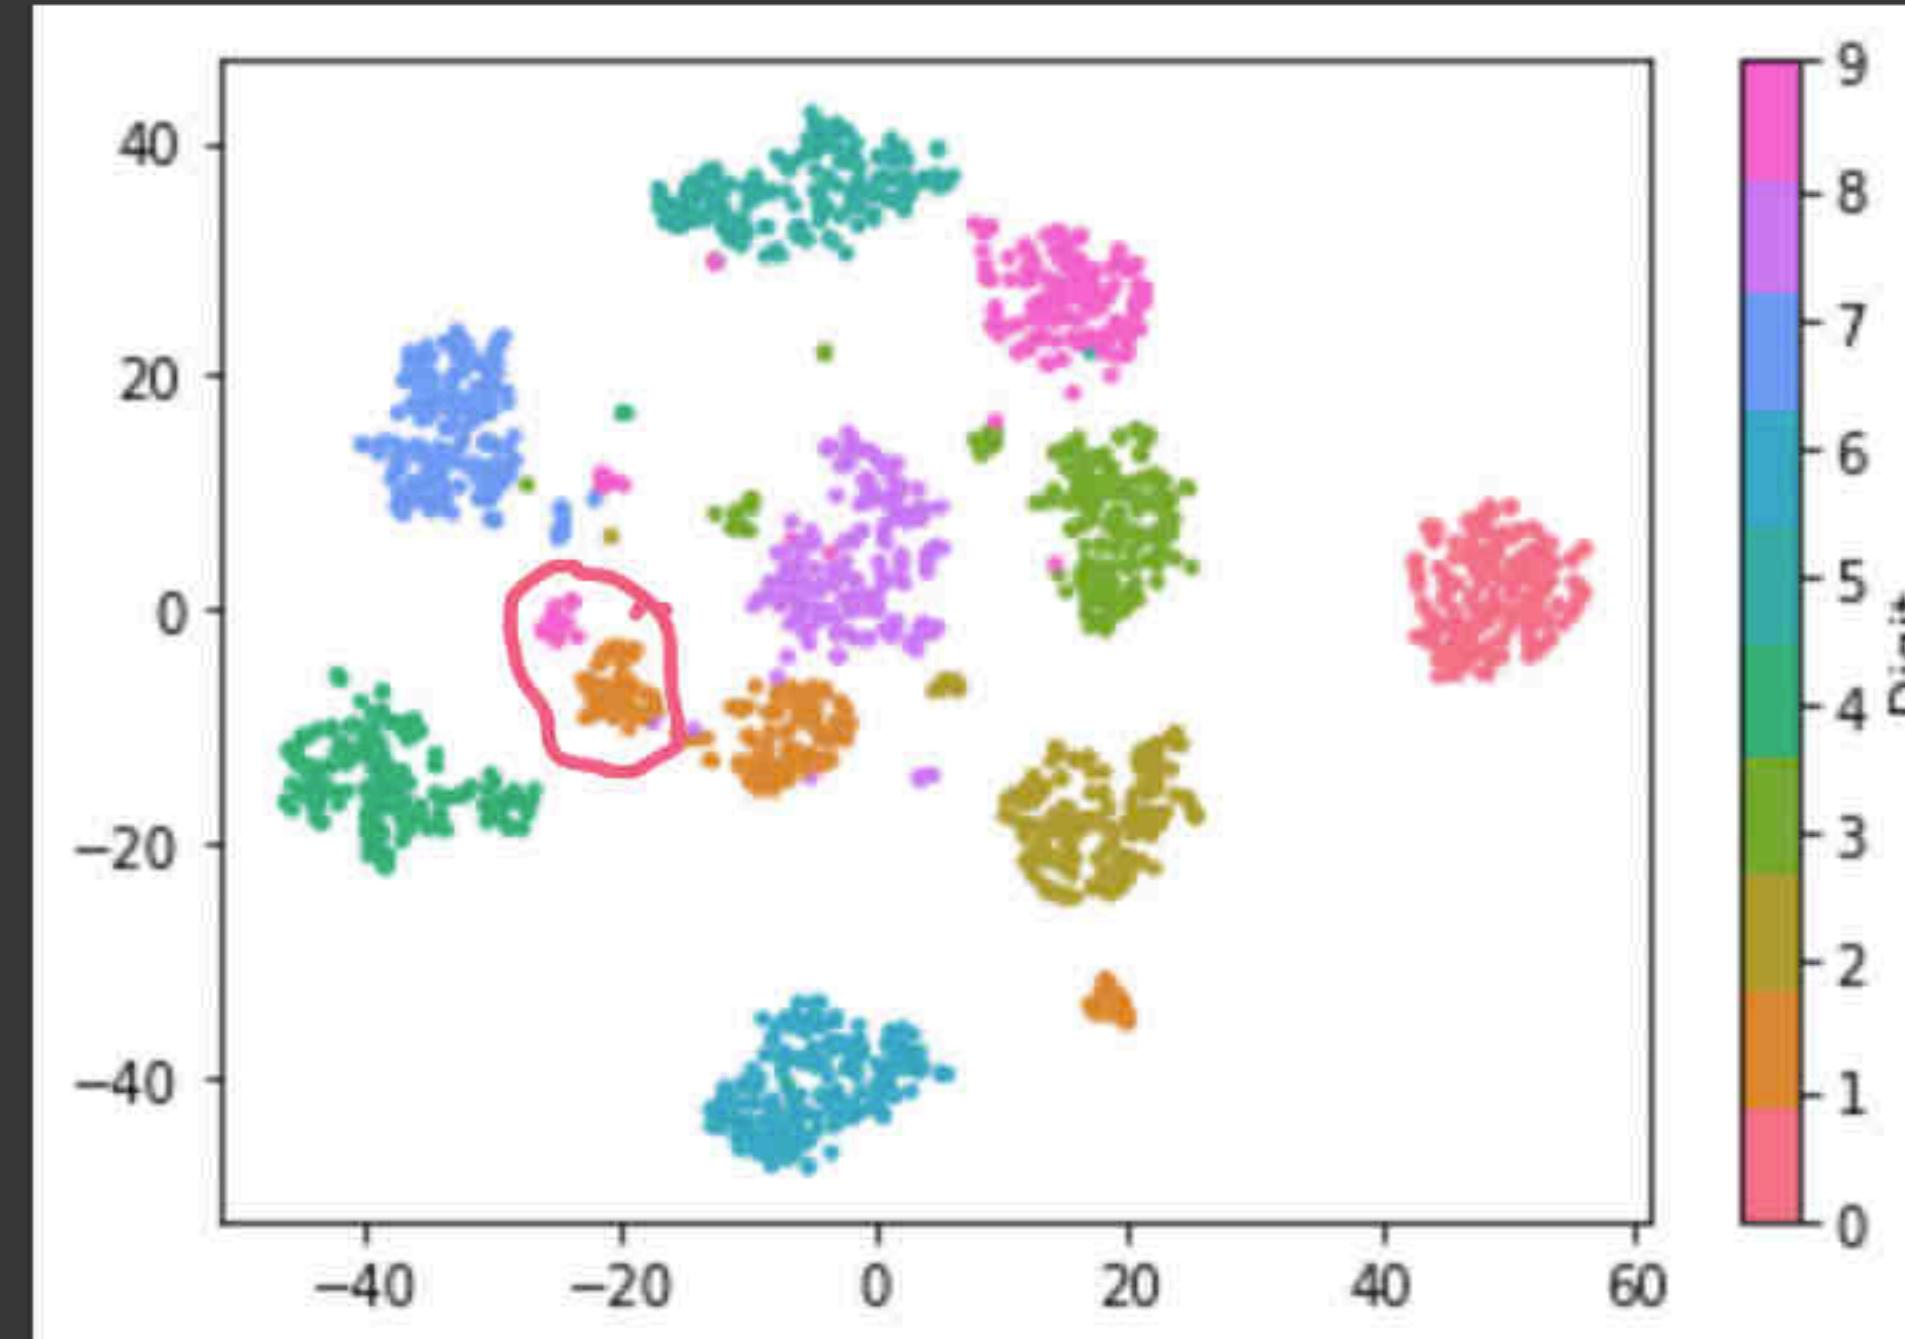
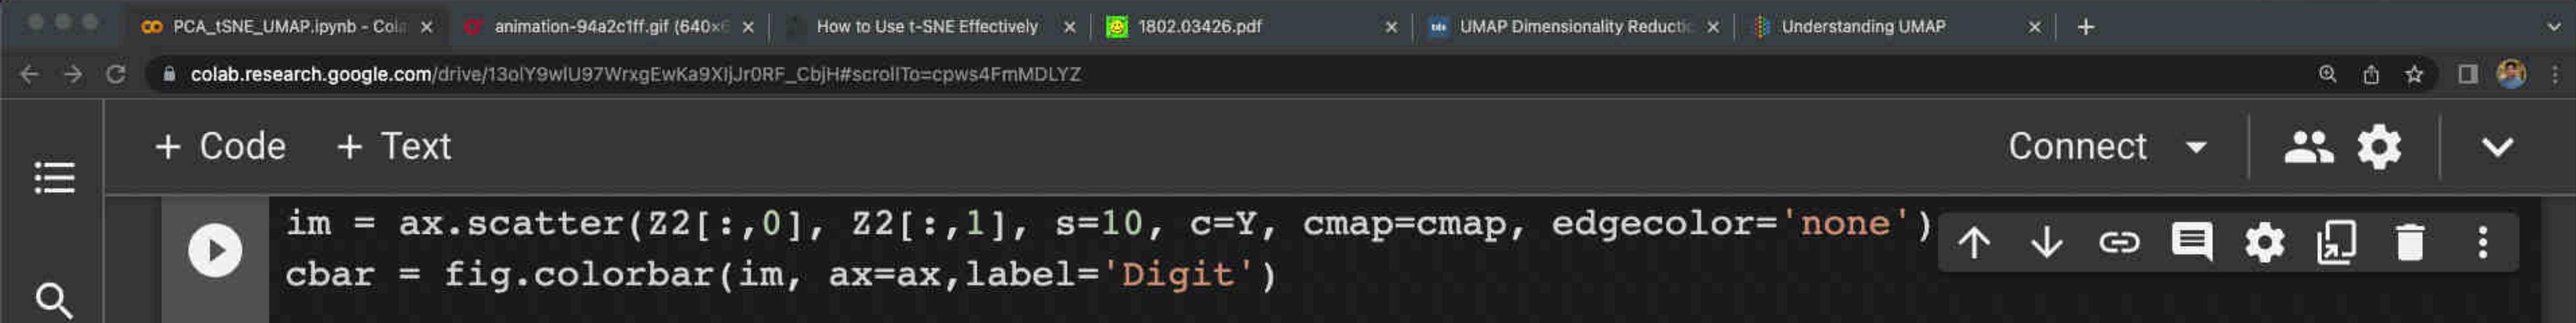




```
[ ] %%time
Z3 = UMAP(init='random', random_state=42).fit_transform(X)
```

CPU times: user 8.39 s, sys: 30 ms, total: 8.42 s  
Wall time: 8.27 s



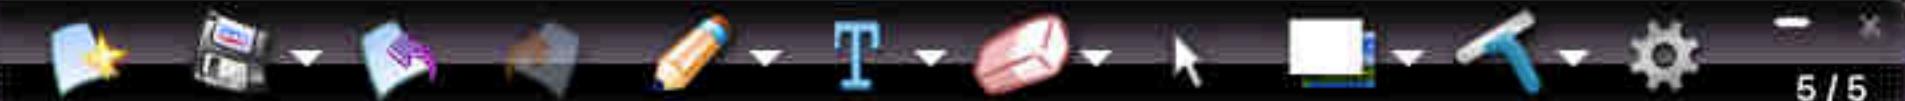


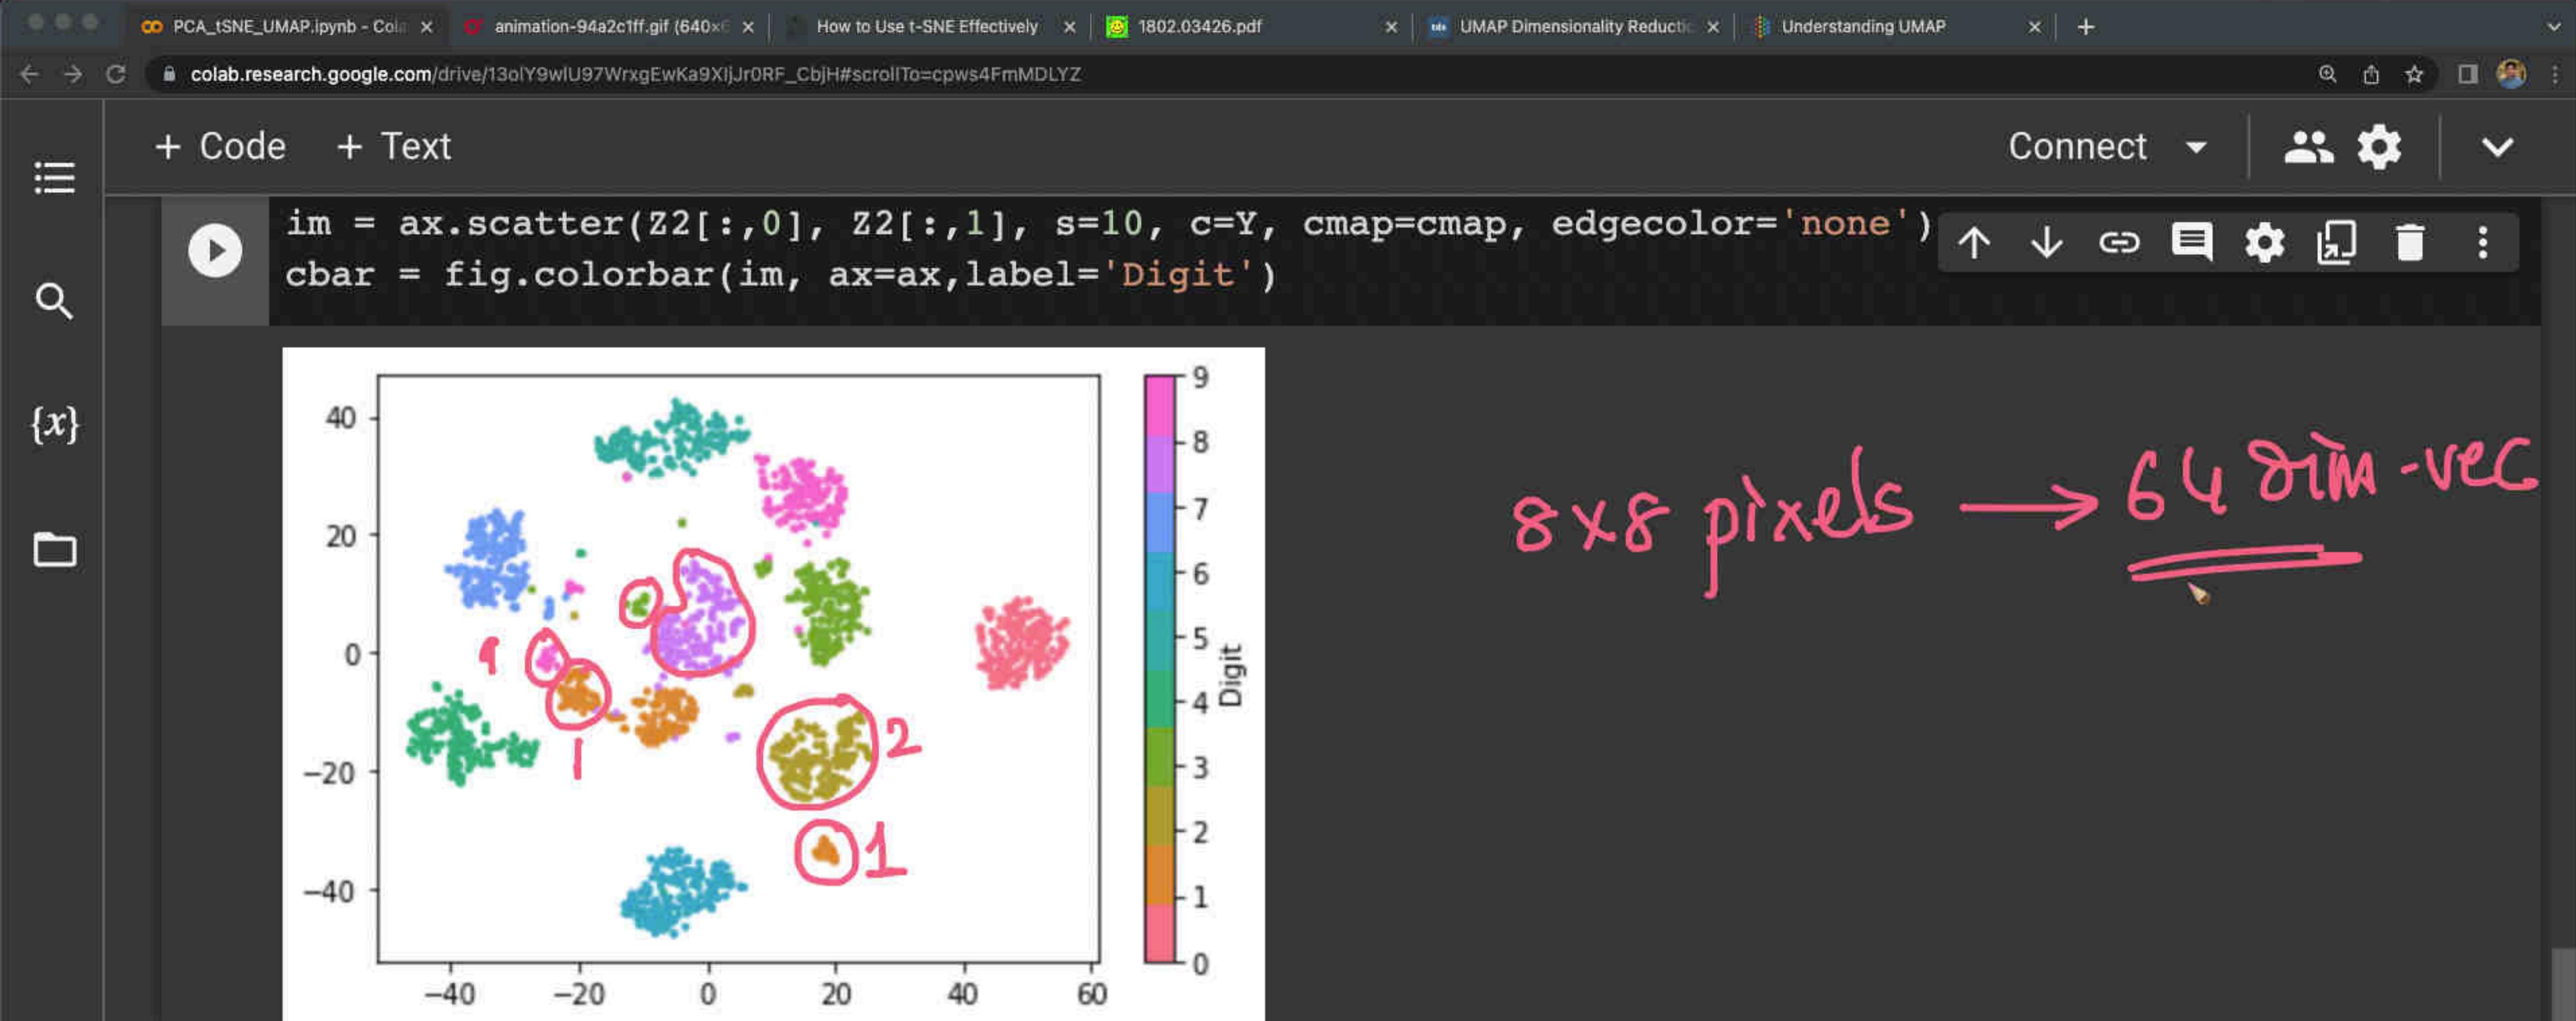
→ not colored & no labels

→ closed

```
[ ] %time
Z3 = UMAP(init='random', random_state=42).fit_transform(X)
```

CPU times: user 8.39 s, sys: 30 ms, total: 8.42 s  
Wall time: 8.27 s





```
[ ] %time  
Z3 = UMAP(init='random', random_state=42).fit_transform(X)
```

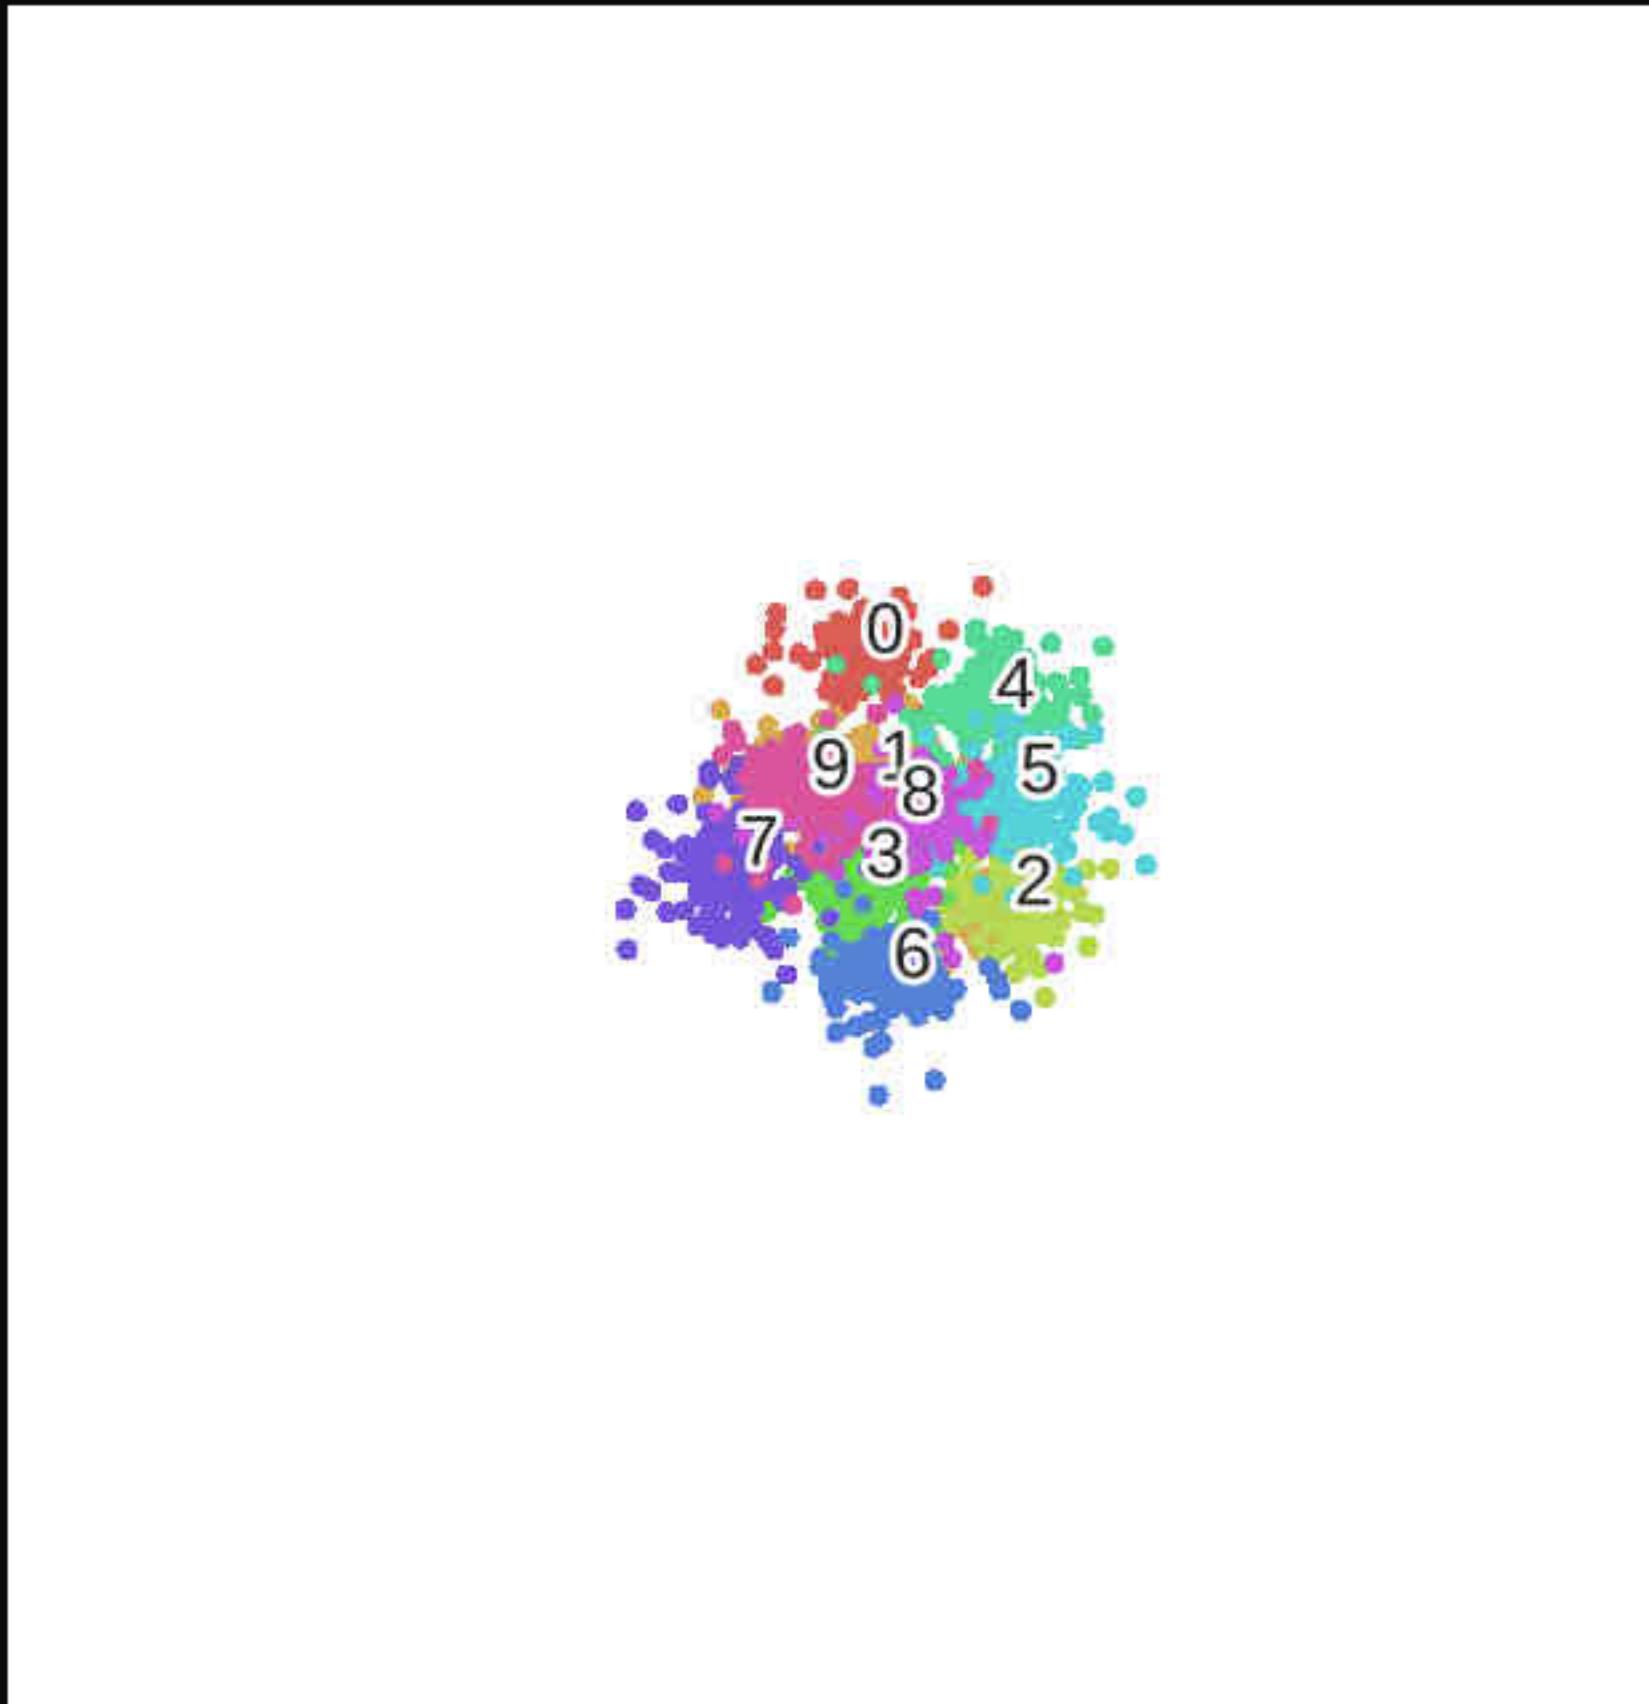
CPU times: user 8.39 s, sys: 30 ms, total: 8.42 s  
Wall time: 8.27 s



$t\text{SNE} \rightarrow \text{opt pool}$

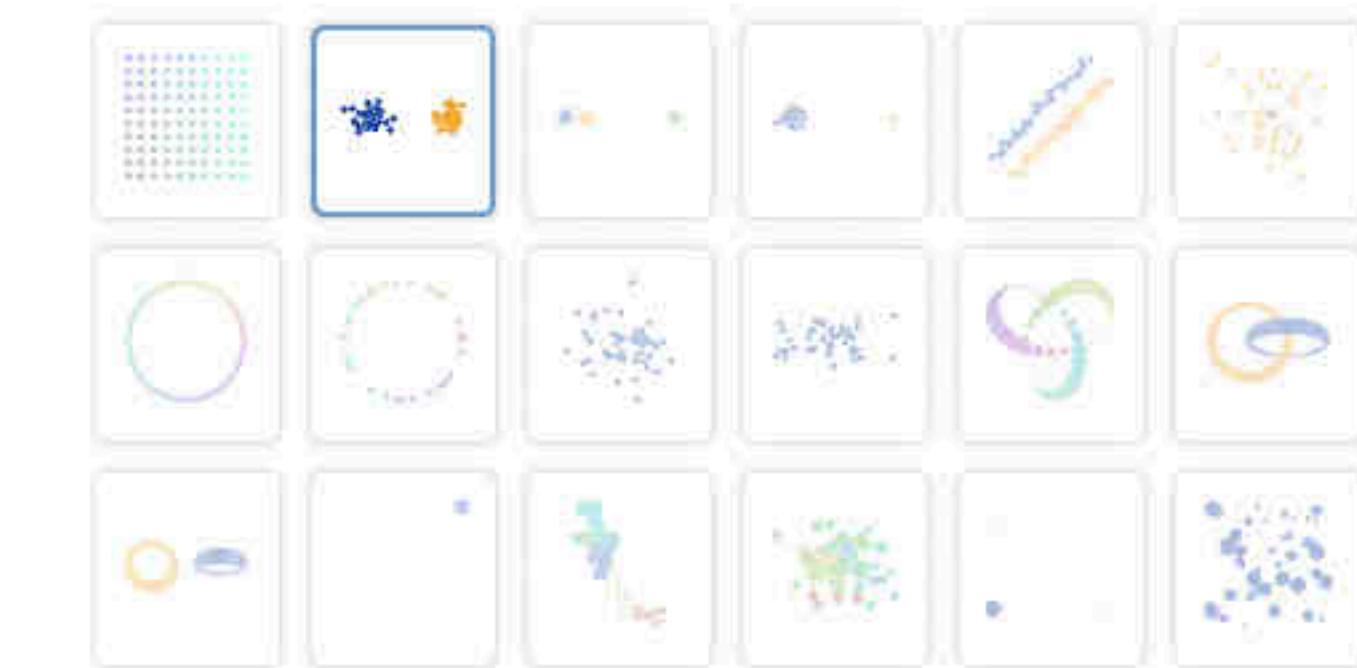
(SGD)

$\min KLDiv(P, Q)$



animation  
of 1st iter n

Sometimes be mysterious or misleading. By exploring how it behaves in simple cases, we can learn to use it more effectively.



Step  
5,000

Points Per Cluster 50

Dimensions 3

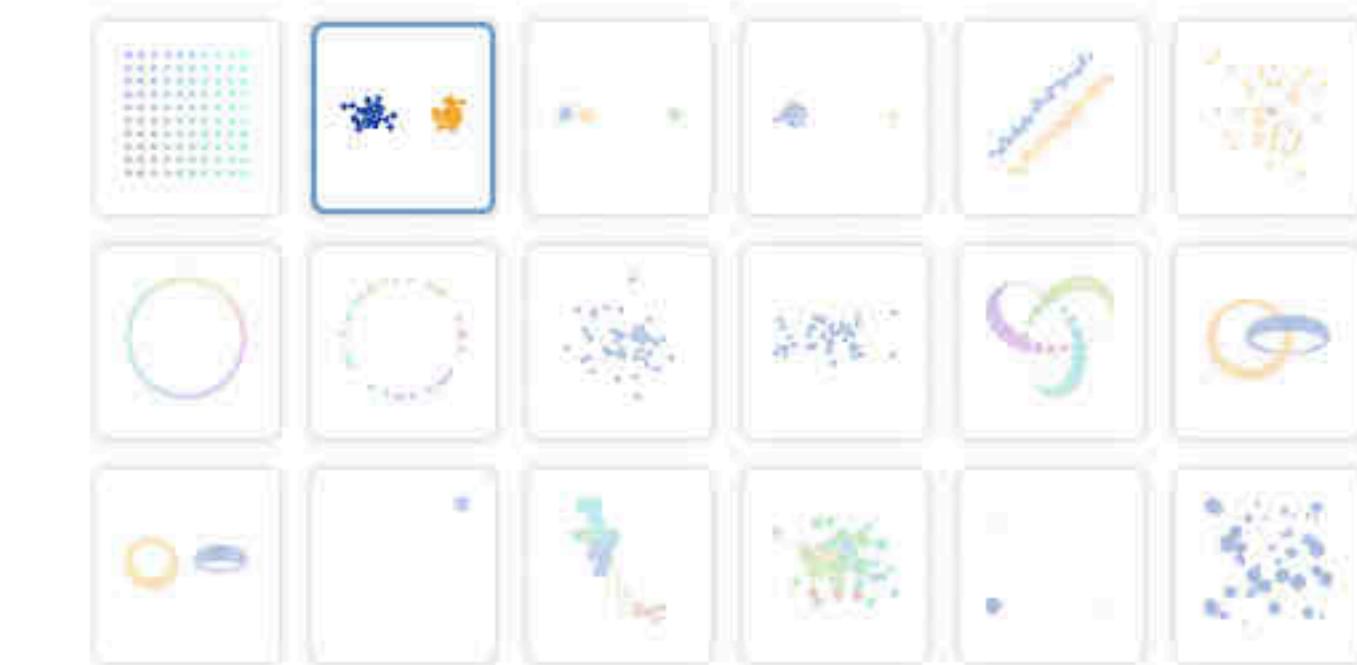
Perplexity 10

Epsilon 5

Two clusters with equal numbers of points.

Share this view

Sometimes be mysterious or misleading. By exploring how it behaves in simple cases, we can learn to use it more effectively.



Step  
5,000

Points Per Cluster 50

Dimensions 3

Perplexity 10 ✓

Epsilon 5 ↗

Two clusters with equal numbers of points.

Share this view

sometimes be mysterious or misleading. By exploring how it behaves in simple cases, we can learn to use it more effectively.

Global  
{ local structure}

Perp = 90

Coexisting

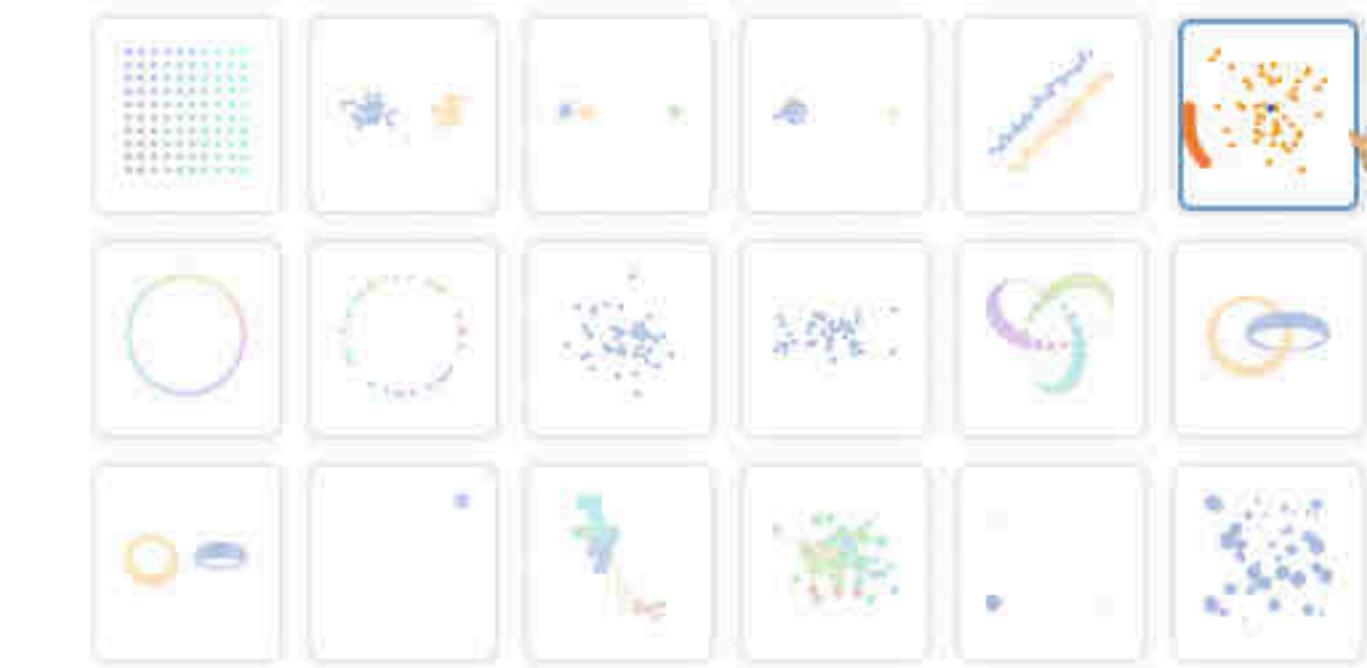
Perp ↑  
preserve global structure

(5-50)

≈ effective # NN  
whose dist we want to preserve

MARTIN WATTENBERG FERNANDA VIÉGAS IAN JOHNSON Oct. 13 Citation:  
Google Brain Google Brain Google Cloud 2016 Wattenberg, et al., 2016

Sometimes be mysterious or misleading. By exploring how it behaves in simple cases, we can learn to use it more effectively.

Step  
5,000

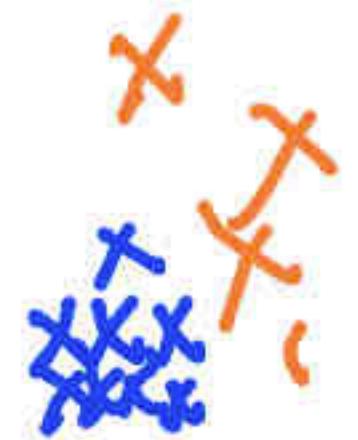
Points Per Cluster 50

Dimensions 2

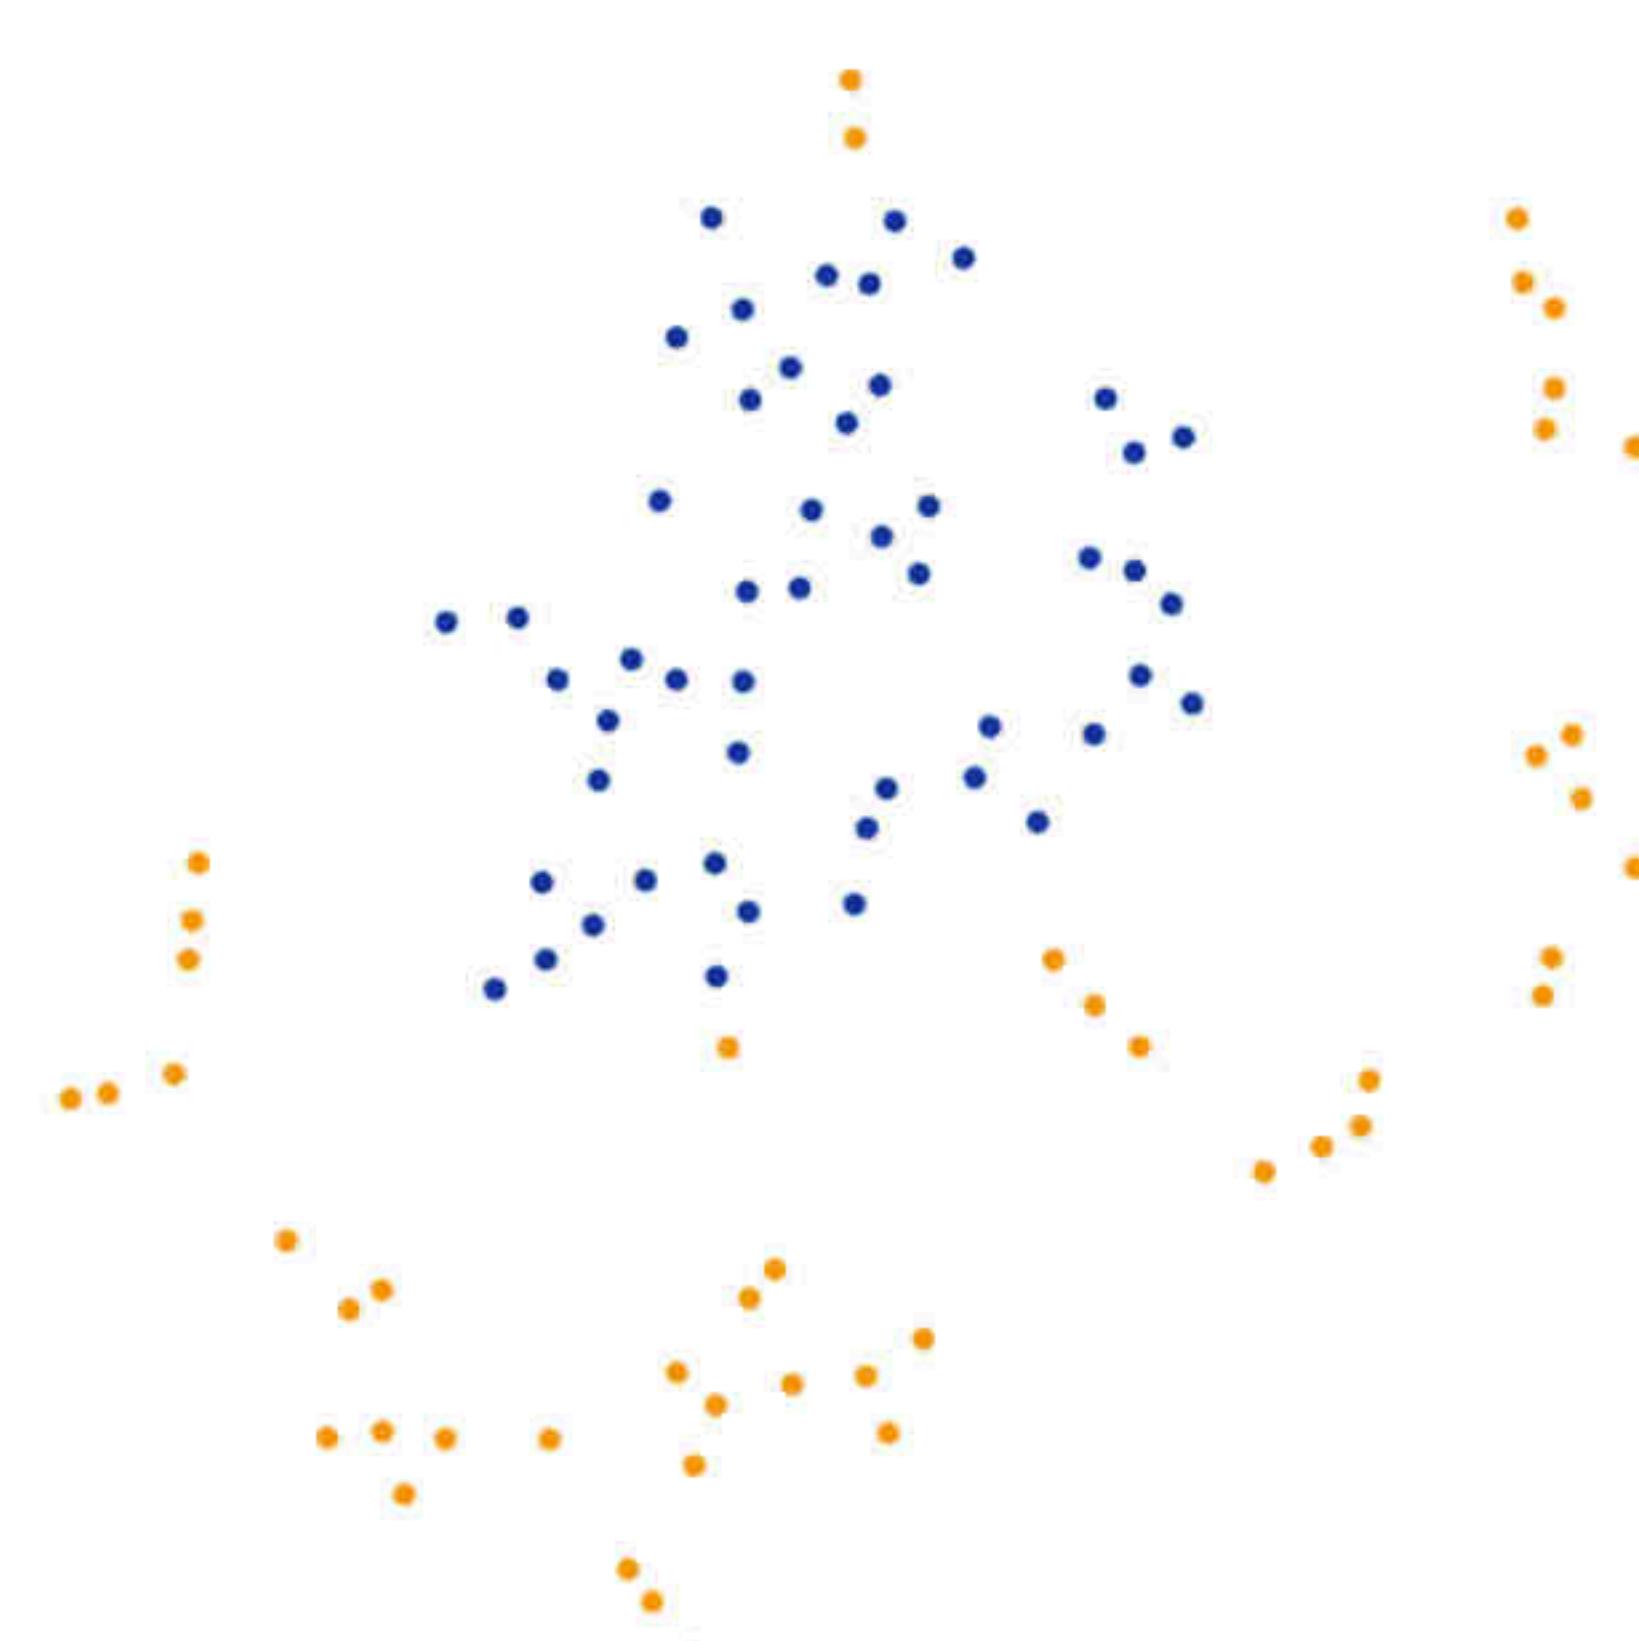
Perplexity 10

Epsilon 5

A dense, tight cluster inside of a wide, sparse cluster. Perplexity makes a big difference here.

[Share this view](#)

sometimes be mysterious or misleading. By exploring how it behaves in simple cases, we can learn to use it more effectively.



Step  
5,000

Points Per Cluster 50

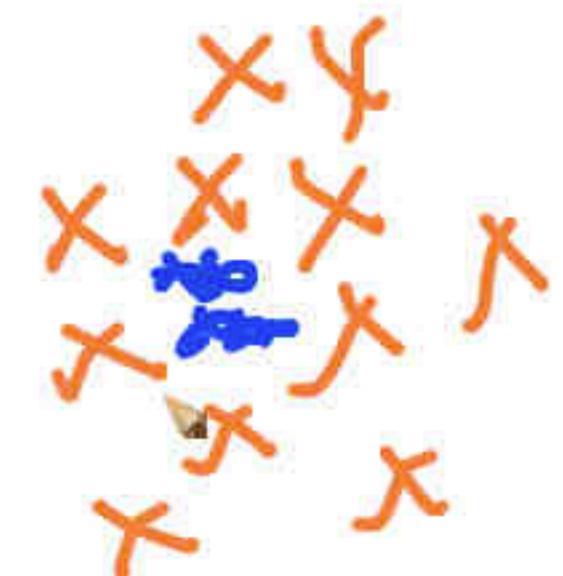
Dimensions 4

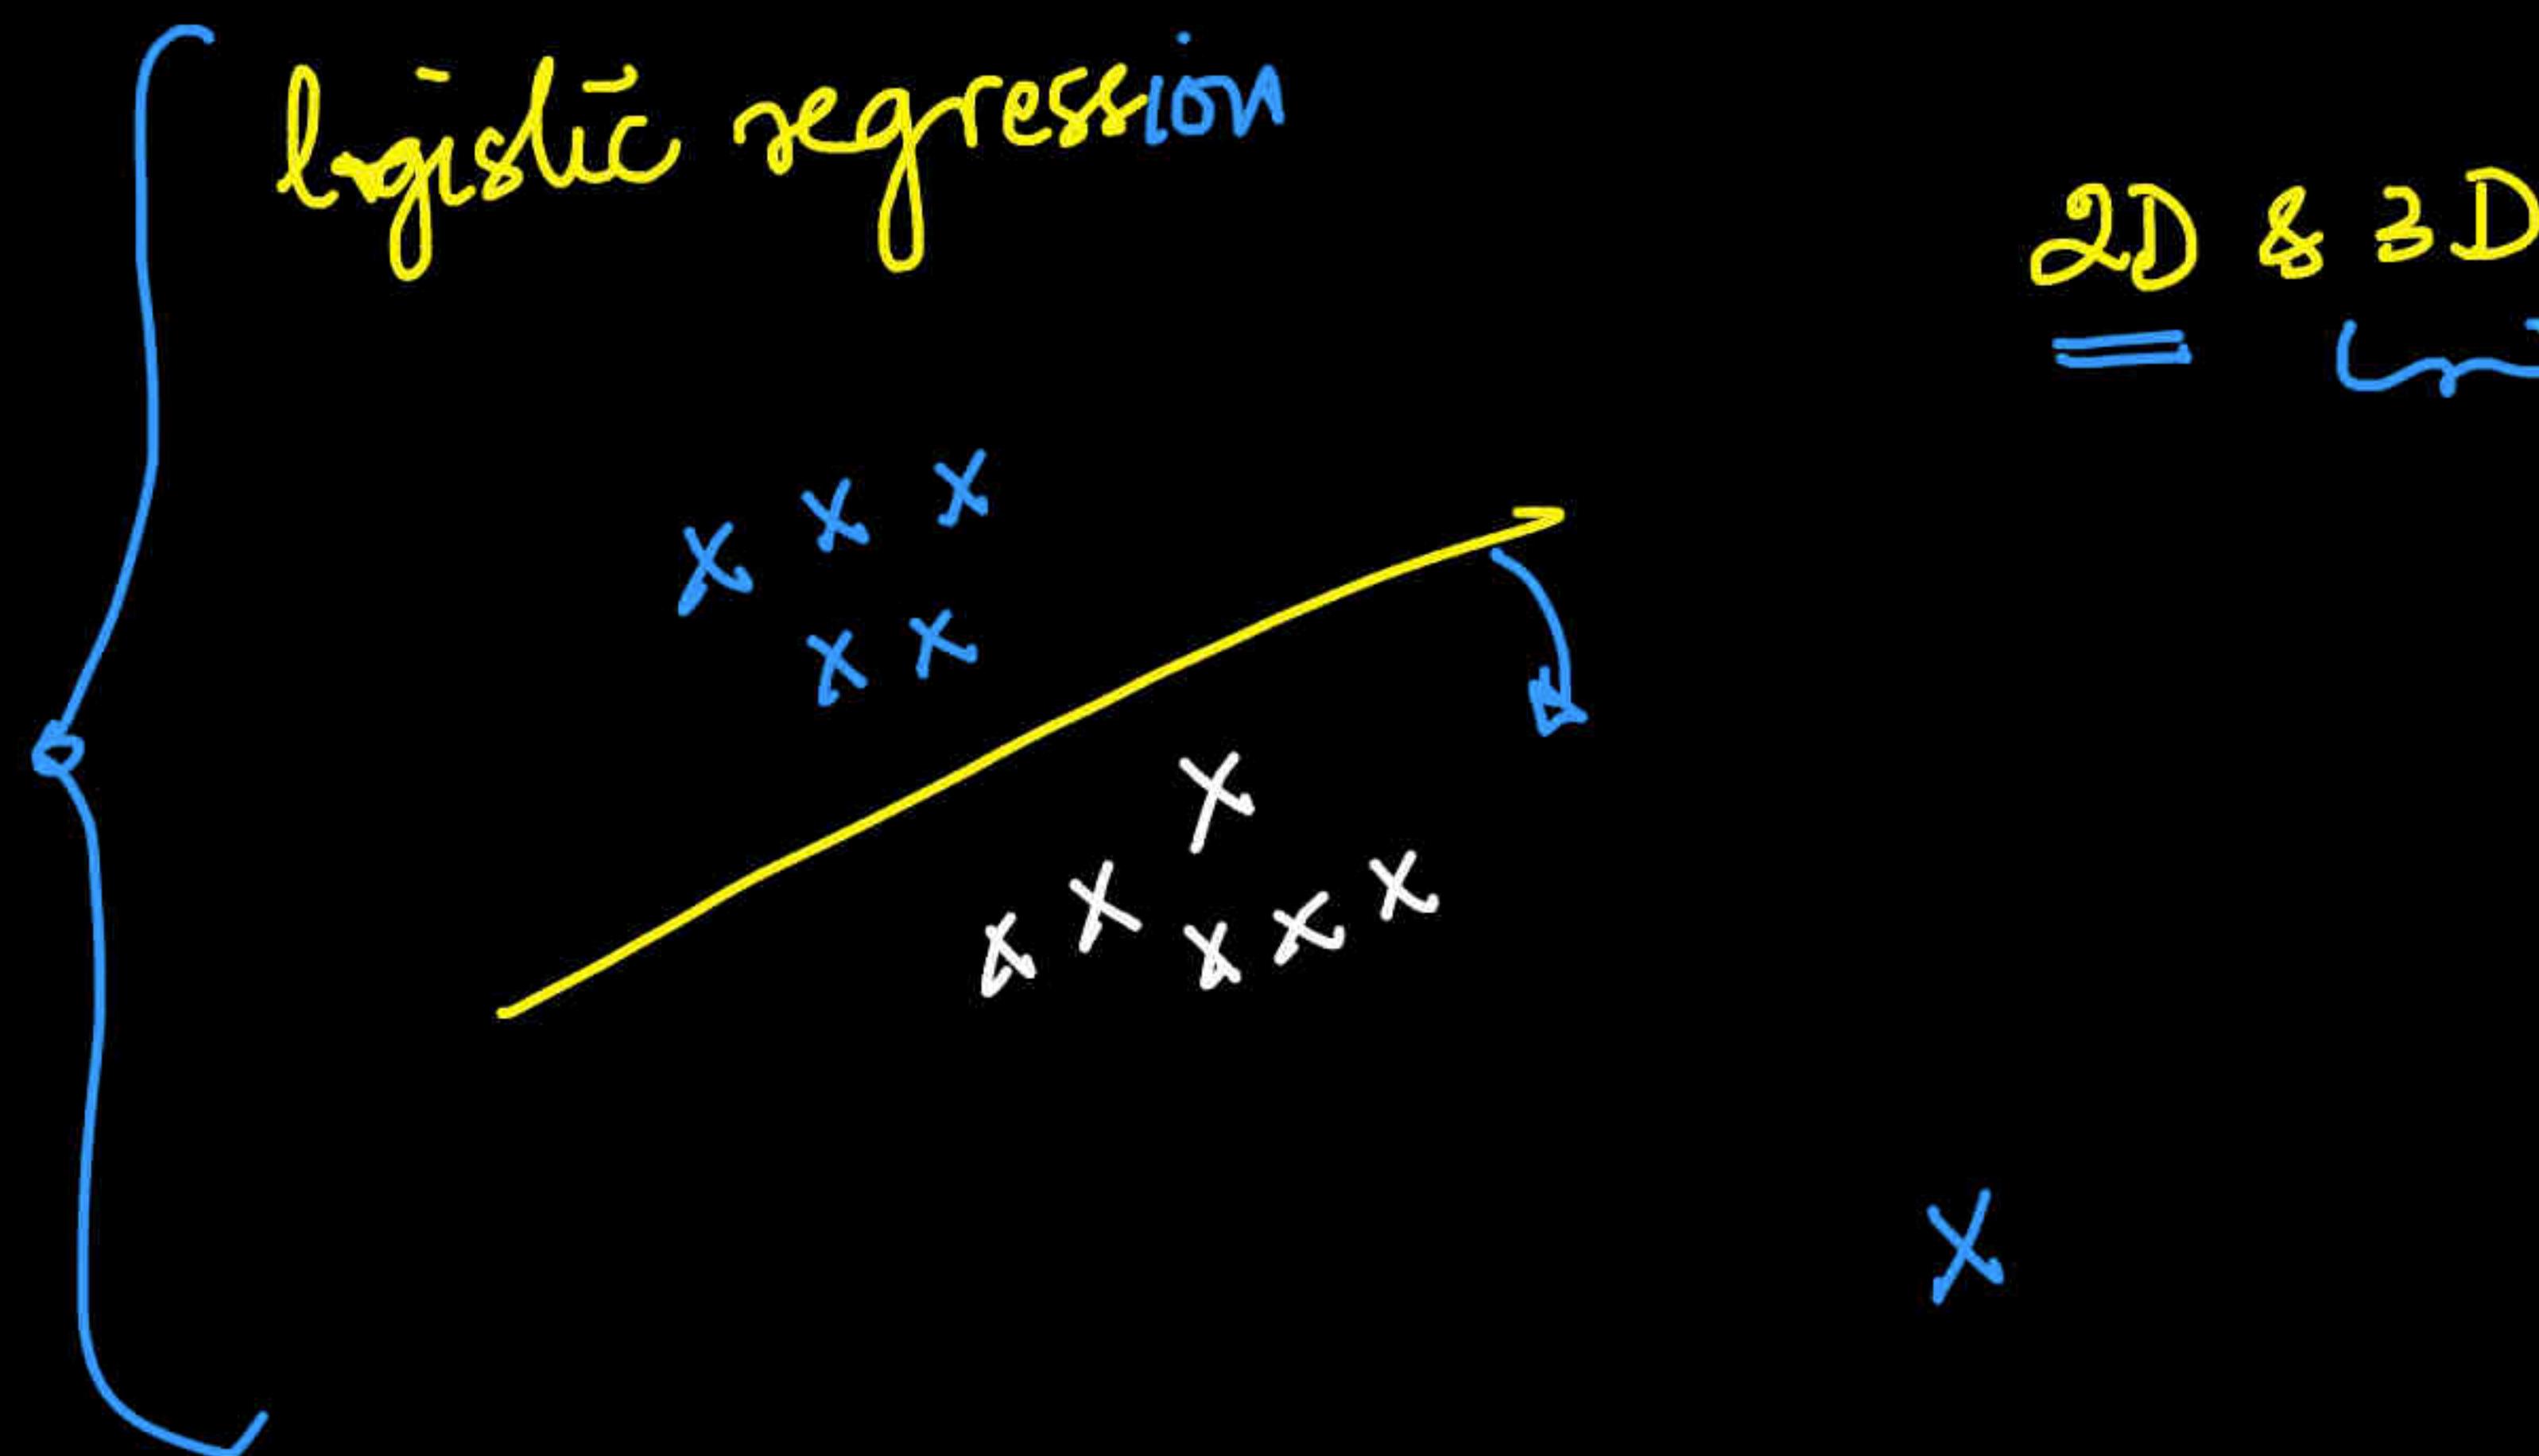
Perplexity 10

Epsilon 5

A dense, tight cluster inside of a wide, sparse cluster. Perplexity makes a big difference here.

Share this view

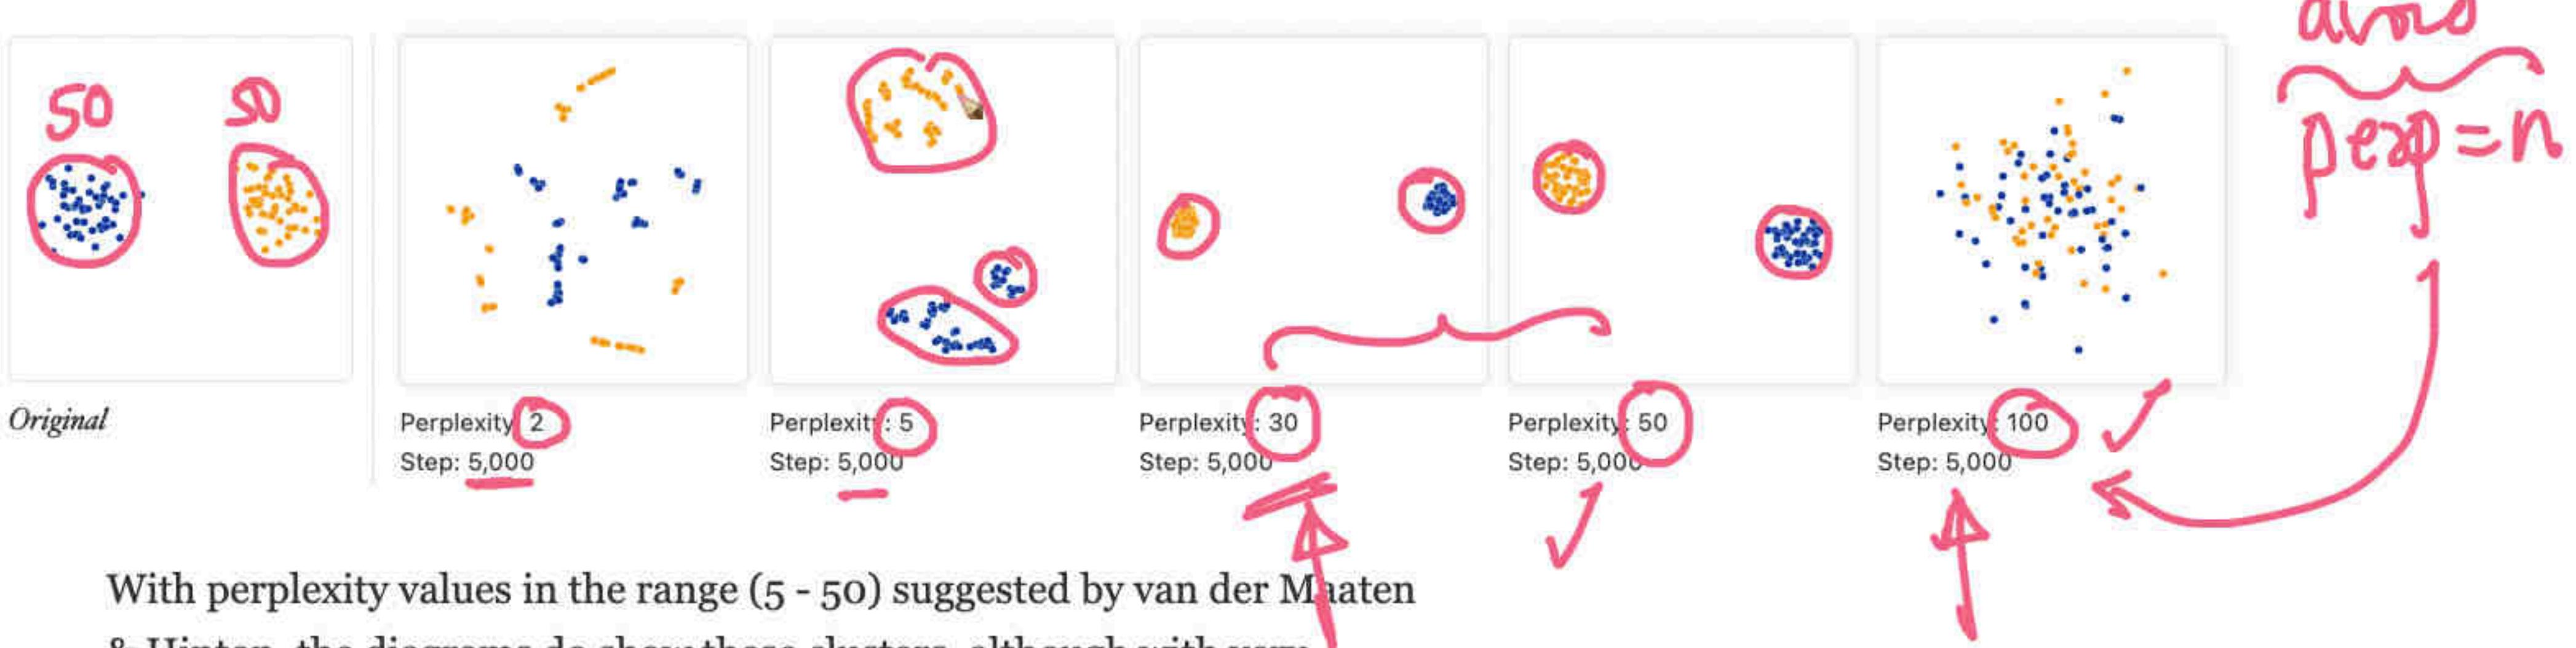




## Lesson 1. Those hyperparameters really matter

Let's start with the "hello world" of t-SNE: a data set of two widely separated clusters. To make things as simple as possible, we'll consider clusters in a 2D plane, as shown in the lefthand diagram. (For clarity, the two clusters are color coded.) The diagrams at right show t-SNE plots for five different perplexity values.

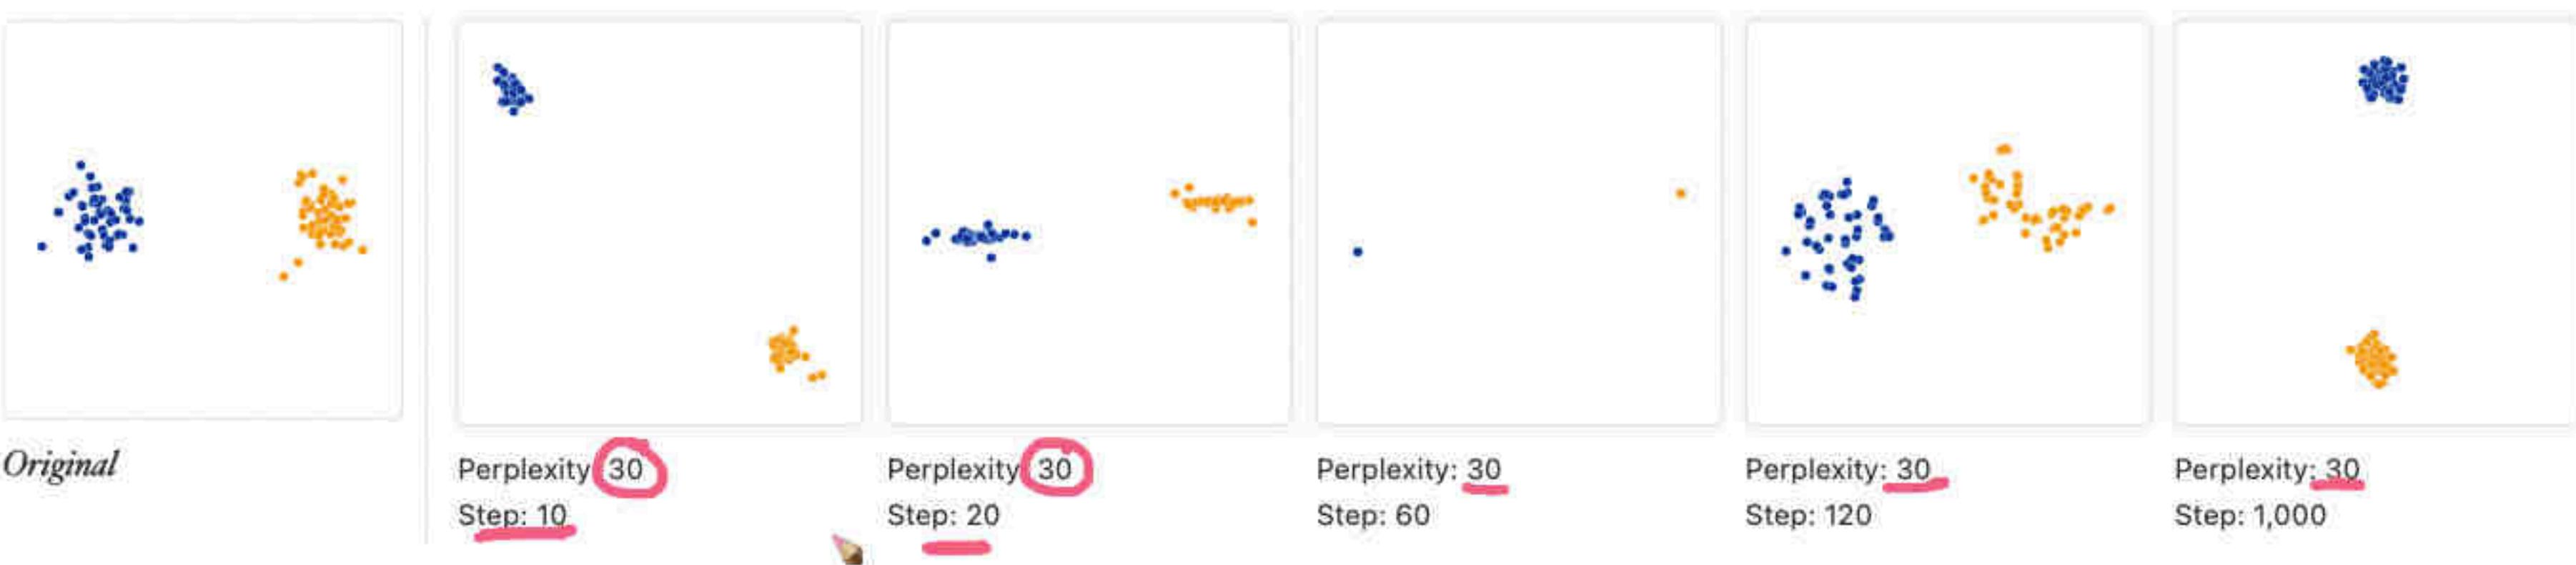
5 - 50



With perplexity values in the range (5 - 50) suggested by van der Maaten & Hinton, the diagrams do show these clusters, although with very different shapes. Outside that range, things get a little weird. With perplexity 2, local variations dominate. The image for perplexity 100, with merged clusters, illustrates a pitfall: for the algorithm to operate properly, the perplexity really should be smaller than the number of points. Implementation

points. Implementations can give unexpected behavior otherwise.

Each of the plots above was made with 5,000 iterations with a learning rate (often called “epsilon”) of 10, and had reached a point of stability by step 5,000. How much of a difference do those values make? In our experience, the most important thing is to iterate until reaching a stable configuration.

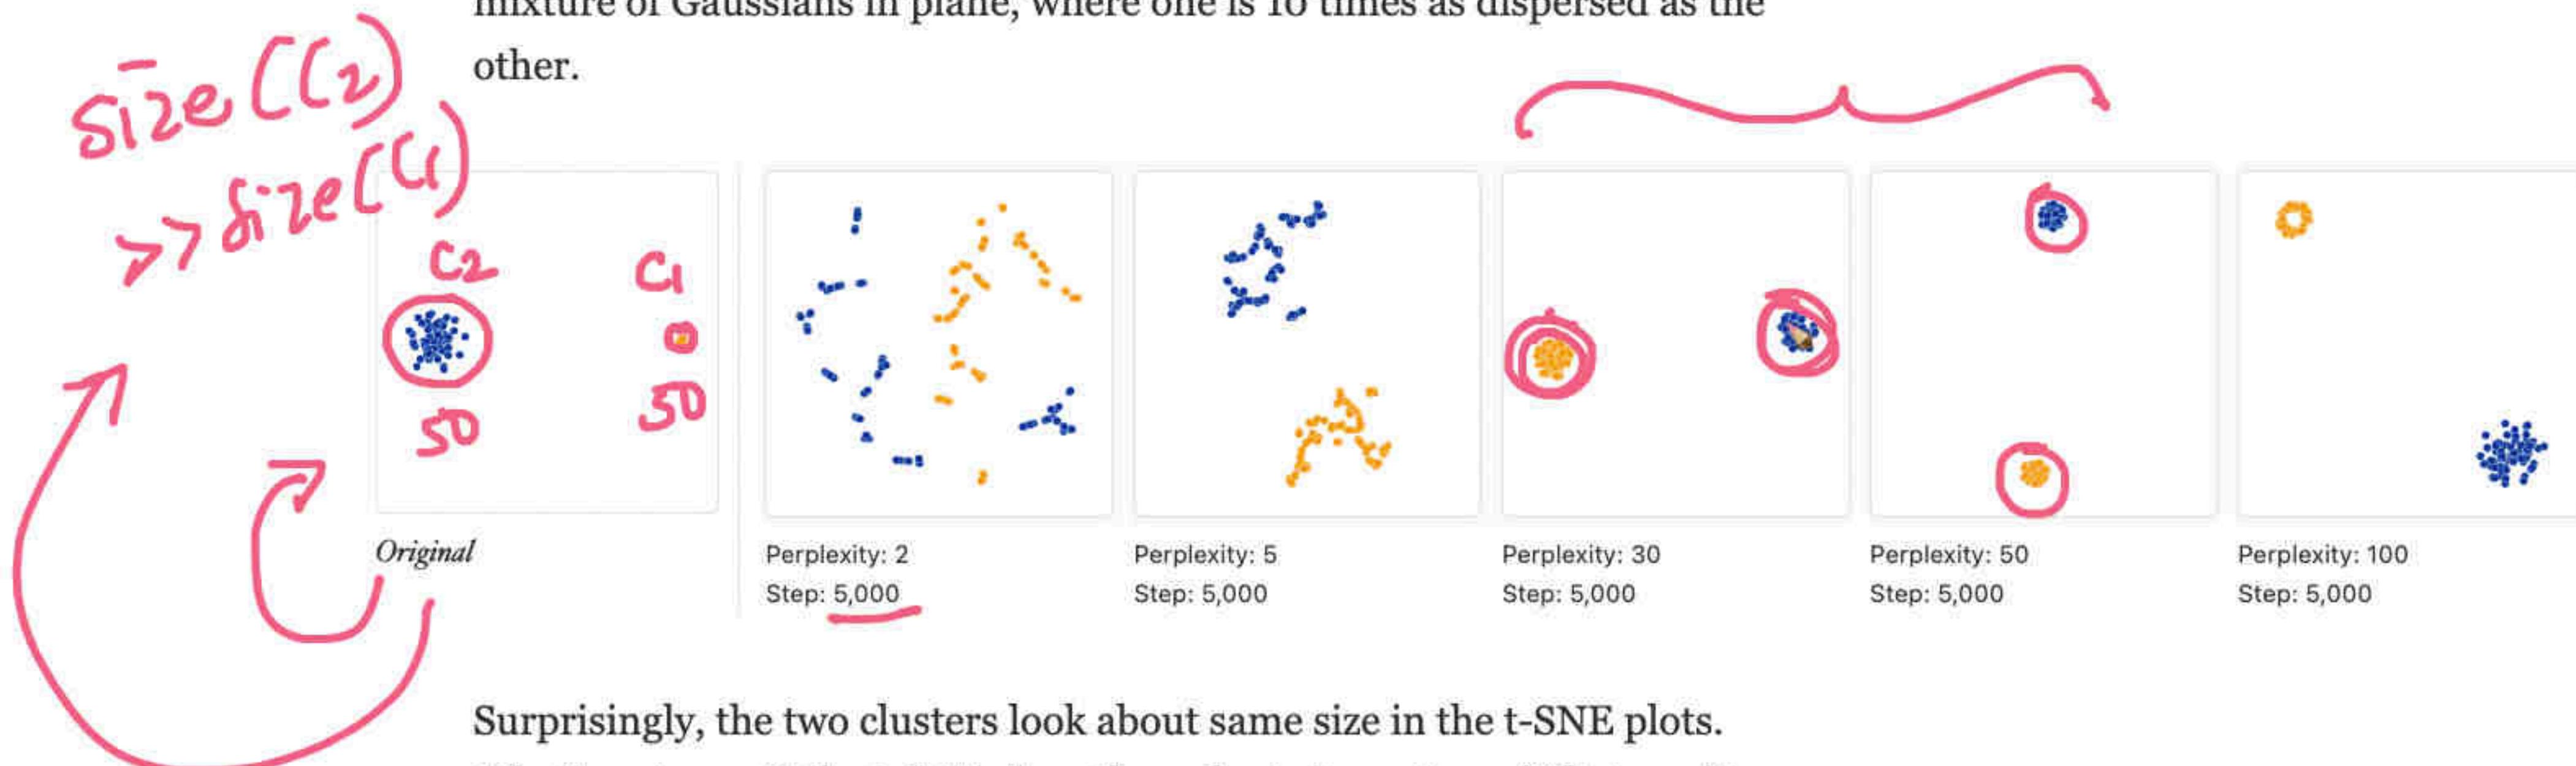


The images above show five different runs at perplexity 30. The first four were stopped before stability. After 10, 20, 60, and 120 steps you can see layouts with seeming 1-dimensional and even pointlike images of the clusters. If you see a t-SNE plot with strange “pinched” shapes, chances are the process was stopped too early. Unfortunately, there’s no fixed number of steps that yields a stable result. Different data sets can require different numbers of iterations to converge.

Another natural question is whether different choices of hyperparameters produce the same results. In this simple two-cluster

## Lesson 2. Cluster sizes in a t-SNE plot mean nothing

So far, so good. But what if the two clusters have different standard deviations, and so different sizes? (By size we mean bounding box measurements, not number of points.) Below are t-SNE plots for a mixture of Gaussians in plane, where one is 10 times as dispersed as the other.

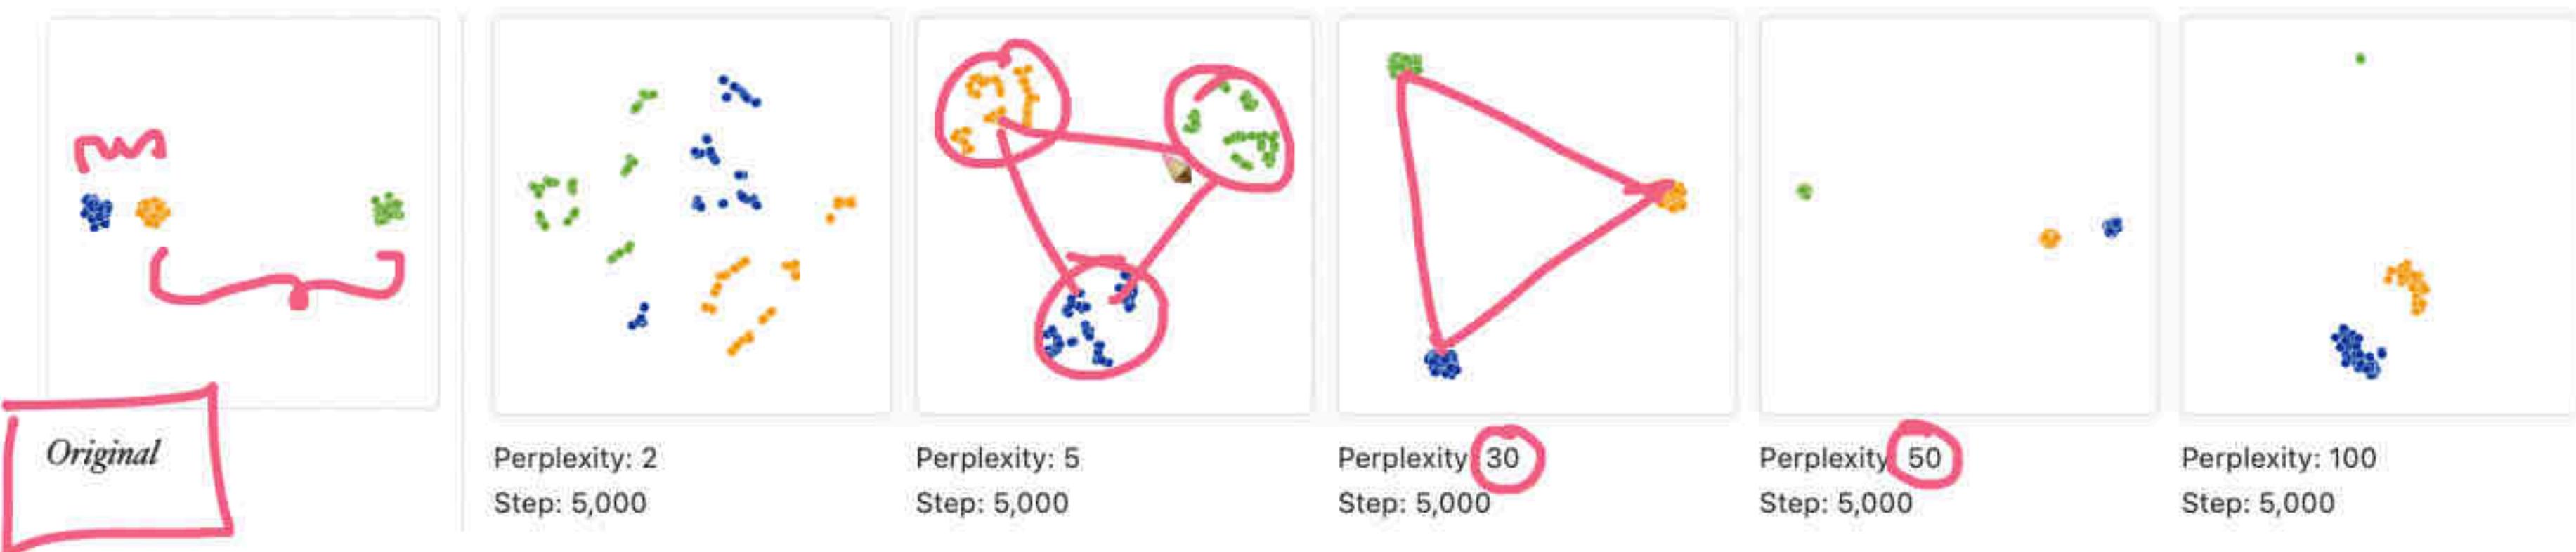


Surprisingly, the two clusters look about same size in the t-SNE plots. What's going on? The t-SNE algorithm adapts its notion of "distance" to regional density variations in the data set. As a result, it naturally expands dense clusters, and contracts sparse ones, evening out cluster sizes. To be clear, this is a different effect than the run-of-the-mill fact that any dimensionality reduction technique will distort distances. (After all, in this example all

# Lesson

## 3. Distances between clusters might not mean anything

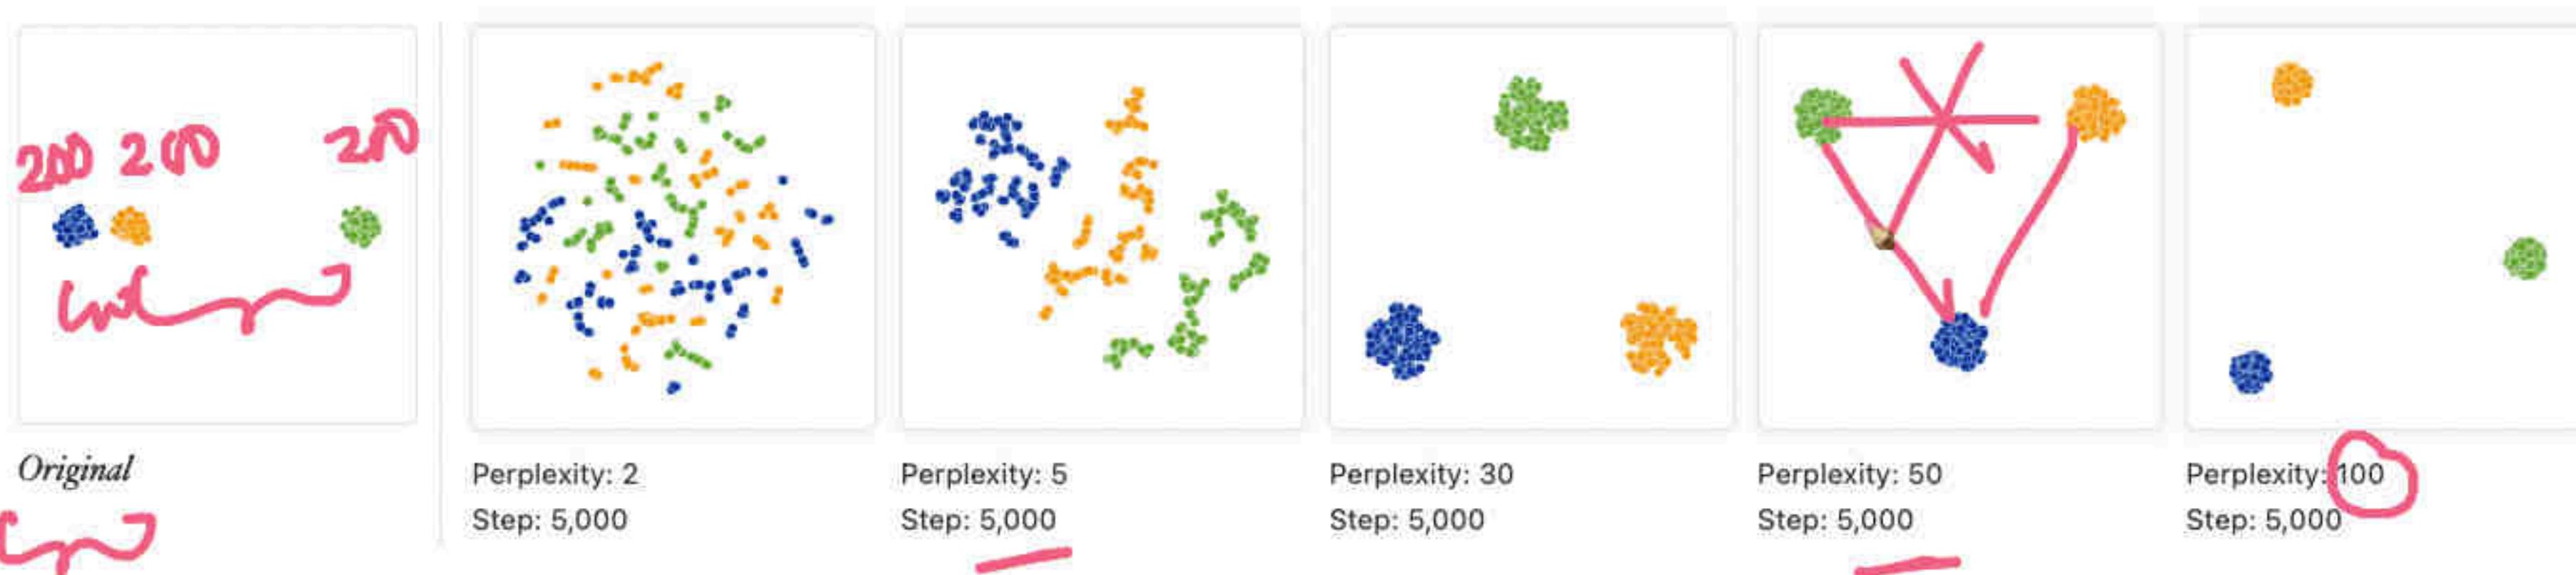
What about distances *between* clusters? The next diagrams show three Gaussians of 50 points each, one pair being 5 times as far apart as another pair.



At perplexity 50, the diagram gives a good sense of the global geometry. For lower perplexity values the clusters look equidistant. When the perplexity is 100, we see the global geometry fine, but one of the cluster appears, falsely, much smaller than the others. Since perplexity 50 gave us a good picture in this example, can we always set perplexity to 50 if we want to see global geometry?

appears, falsely, much smaller than the others. Since perplexity 50 gave us a good picture in this example, can we always set perplexity to 50 if we want to see global geometry?

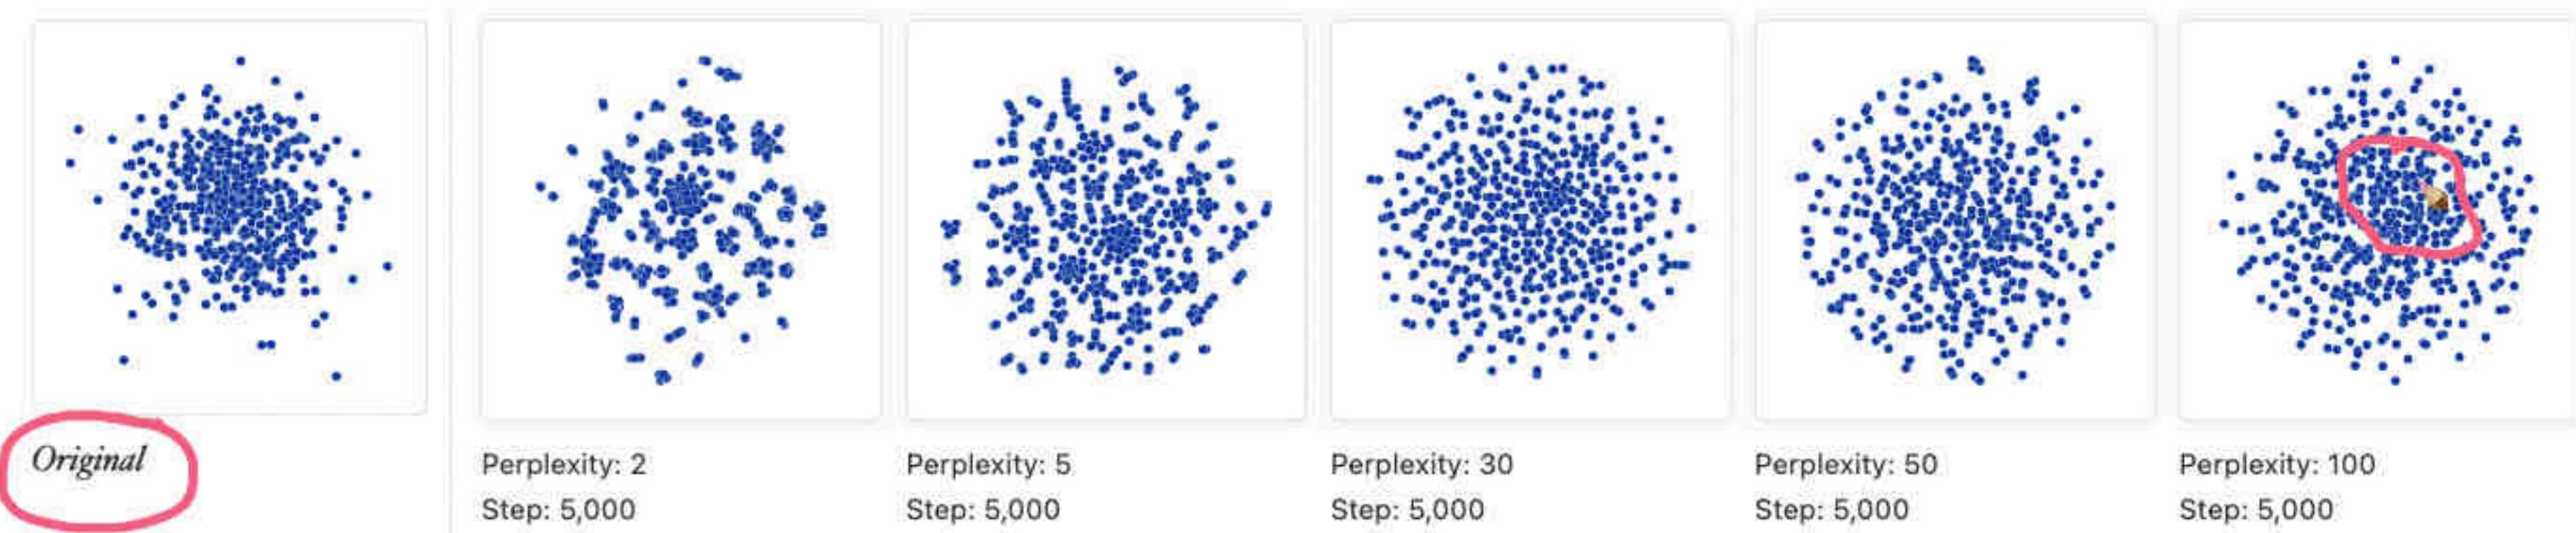
Sadly, no. If we add more points to each cluster, the perplexity has to increase to compensate. Here are the t-SNE diagrams for three Gaussian clusters with 200 points each, instead of 50. Now none of the trial perplexity values gives a good result.



It's bad news that seeing global geometry requires fine-tuning perplexity. Real-world data would probably have multiple clusters with different numbers of elements. There may not be one perplexity value that will capture distances across all clusters—and sadly perplexity is a global parameter. Fixing this problem might be an interesting area for future research.

## Lesson 4. Random noise doesn't always look random.

A classic pitfall is thinking you see patterns in what is really just random data. Recognizing noise when you see it is a critical skill, but it takes time to build up the right intuitions. A tricky thing about t-SNE is that it throws a lot of existing intuition out the window. The next diagrams show genuinely random data, 500 points drawn from a unit Gaussian distribution in 100 dimensions. The left image is a projection onto the first two coordinates.



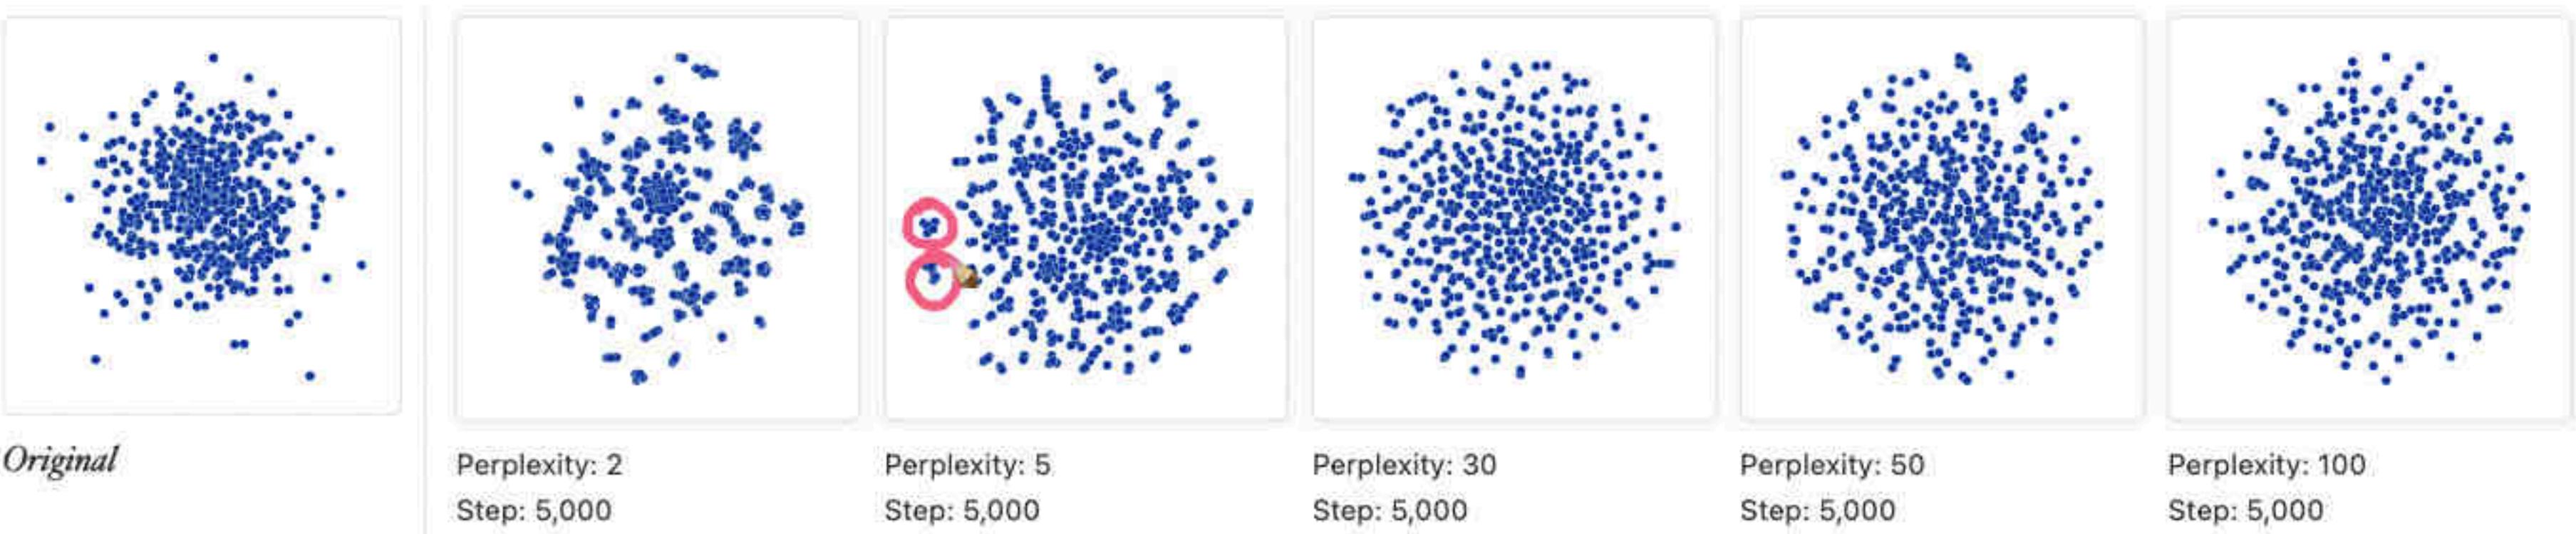
The plot with perplexity 2 seems to show dramatic clusters. If you were tuning perplexity to bring out structure in the data, you might think you'd hit the jackpot.

Of course, since we kn



## 4. Random noise doesn't always look random.

A classic pitfall is thinking you see patterns in what is really just random data. Recognizing noise when you see it is a critical skill, but it takes time to build up the right intuitions. A tricky thing about t-SNE is that it throws a lot of existing intuition out the window. The next diagrams show genuinely random data, 500 points drawn from a unit Gaussian distribution in 100 dimensions. The left image is a projection onto the first two coordinates.



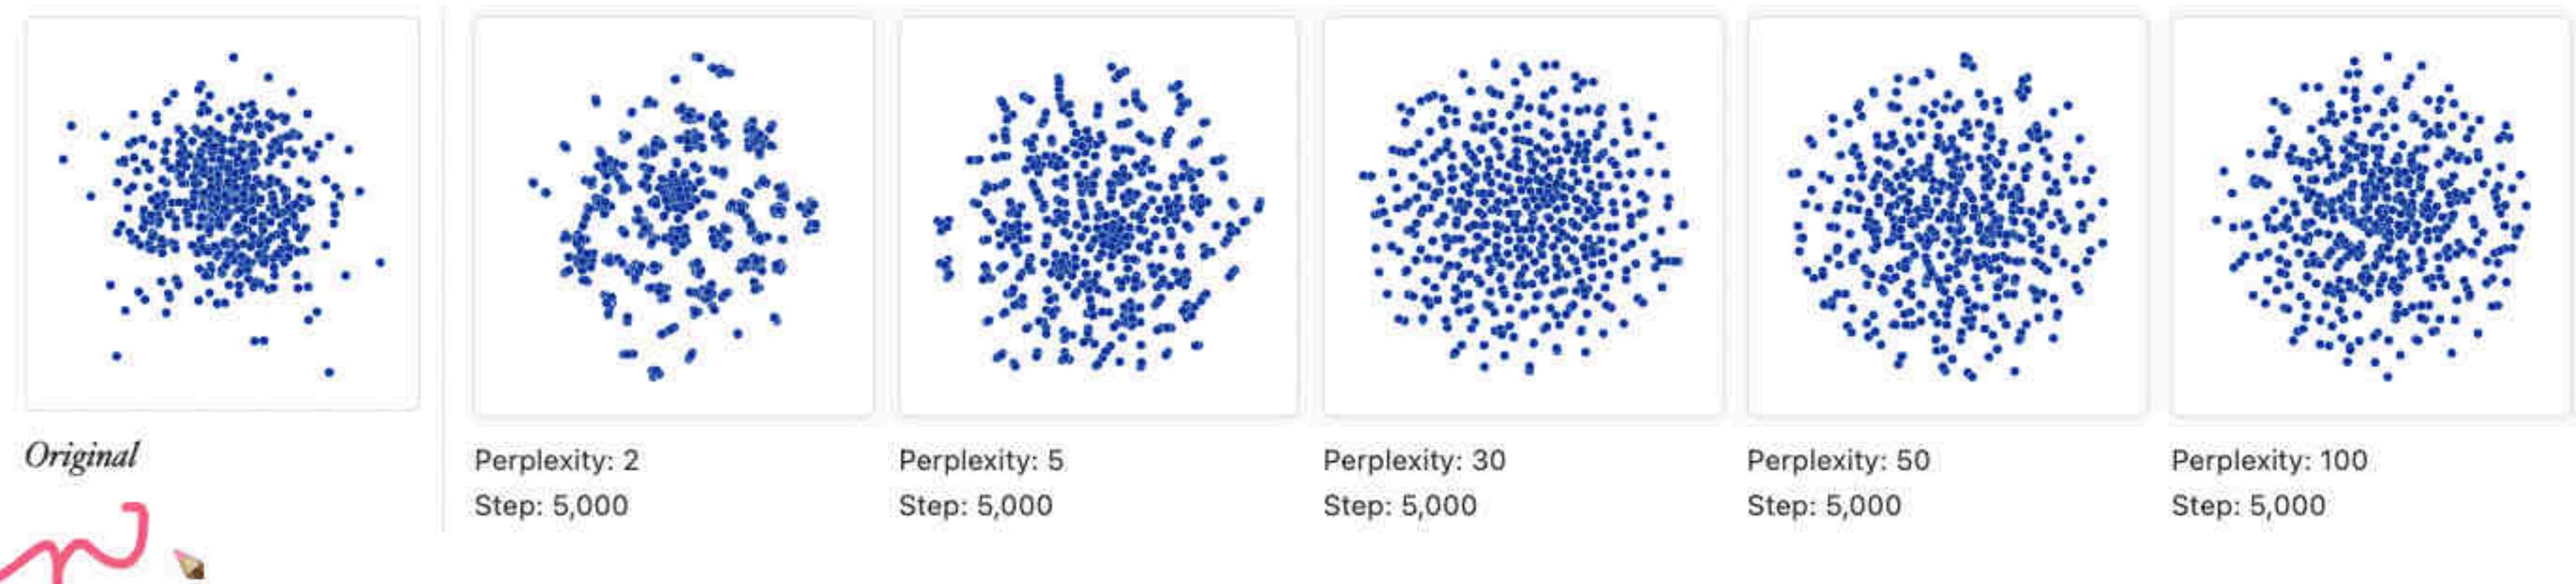
The plot with perplexity 2 seems to show dramatic clusters. If you were tuning perplexity to bring out structure in the data, you might think you'd hit the jackpot.

Of course, since we kn



## 4. Random noise doesn't always look random.

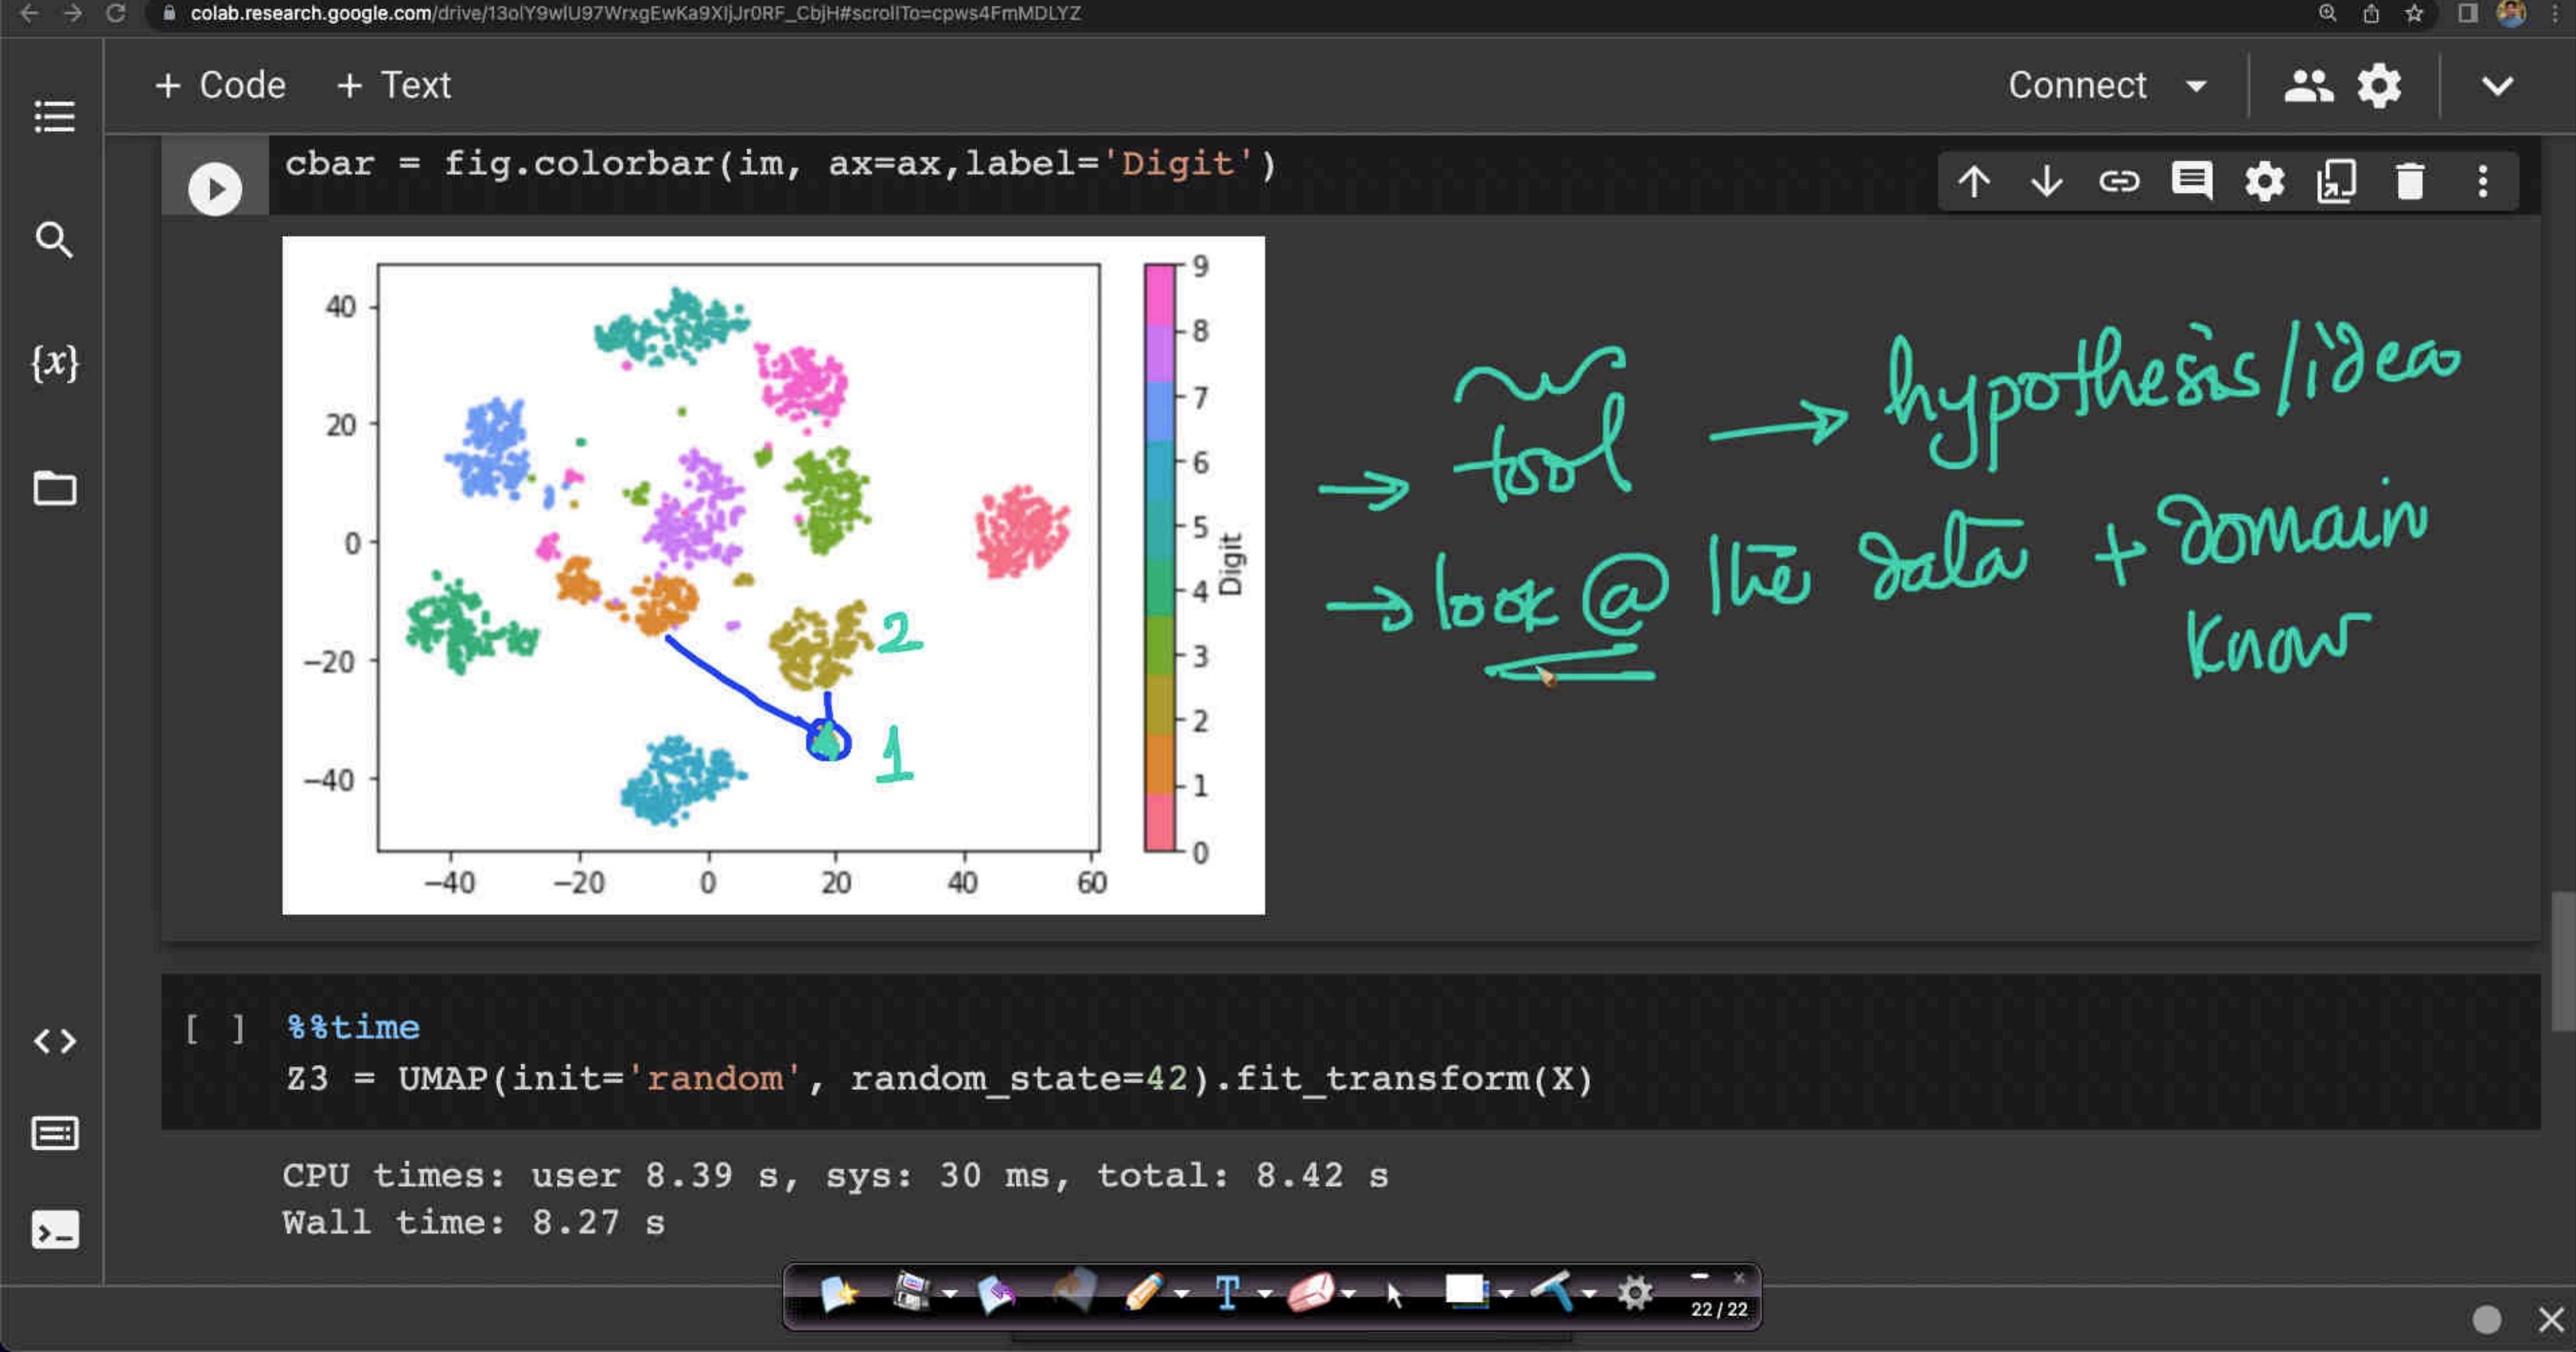
A classic pitfall is thinking you see patterns in what is really just random data. Recognizing noise when you see it is a critical skill, but it takes time to build up the right intuitions. A tricky thing about t-SNE is that it throws a lot of existing intuition out the window. The next diagrams show genuinely random data, 500 points drawn from a unit Gaussian distribution in 100 dimensions. The left image is a projection onto the first two coordinates.



The plot with perplexity 2 seems to show dramatic clusters. If you were tuning perplexity to bring out structure in the data, you might think you'd hit the jackpot.

Of course, since we kn





tSNE in ML

$$\mathcal{D} = \{(x_i, \hat{y}_i)\}_{i=1}^n$$

$\downarrow$   
 $y_i$

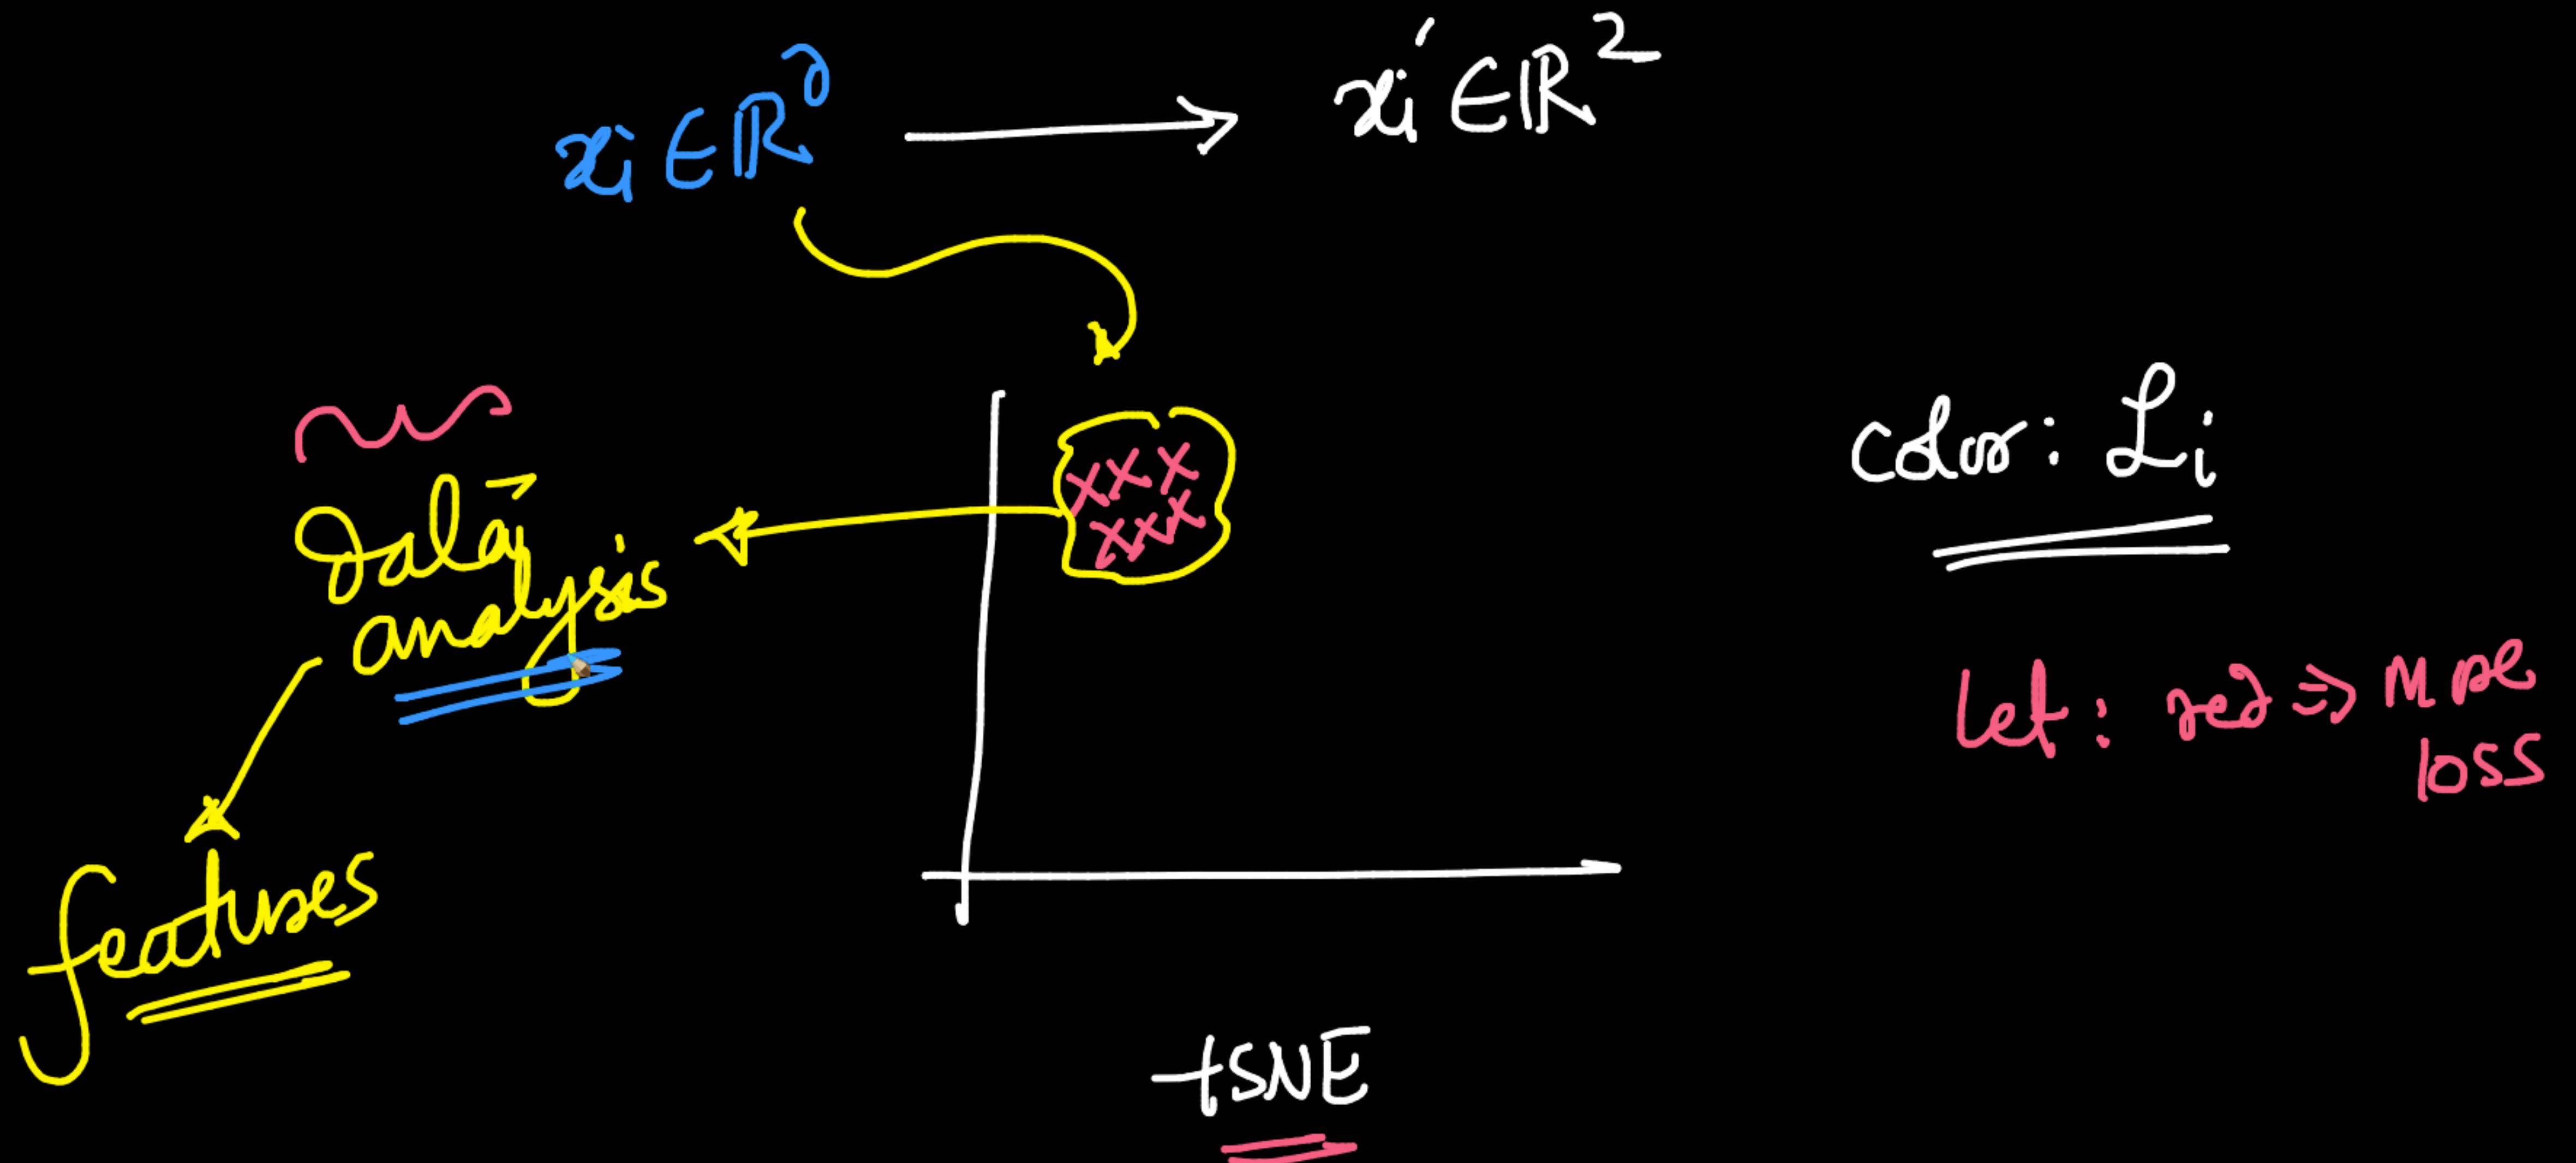
Classification

→ outliers  
→ clustering



$$\tilde{l}_i = \text{diff}(y_i, \hat{y}_i)$$

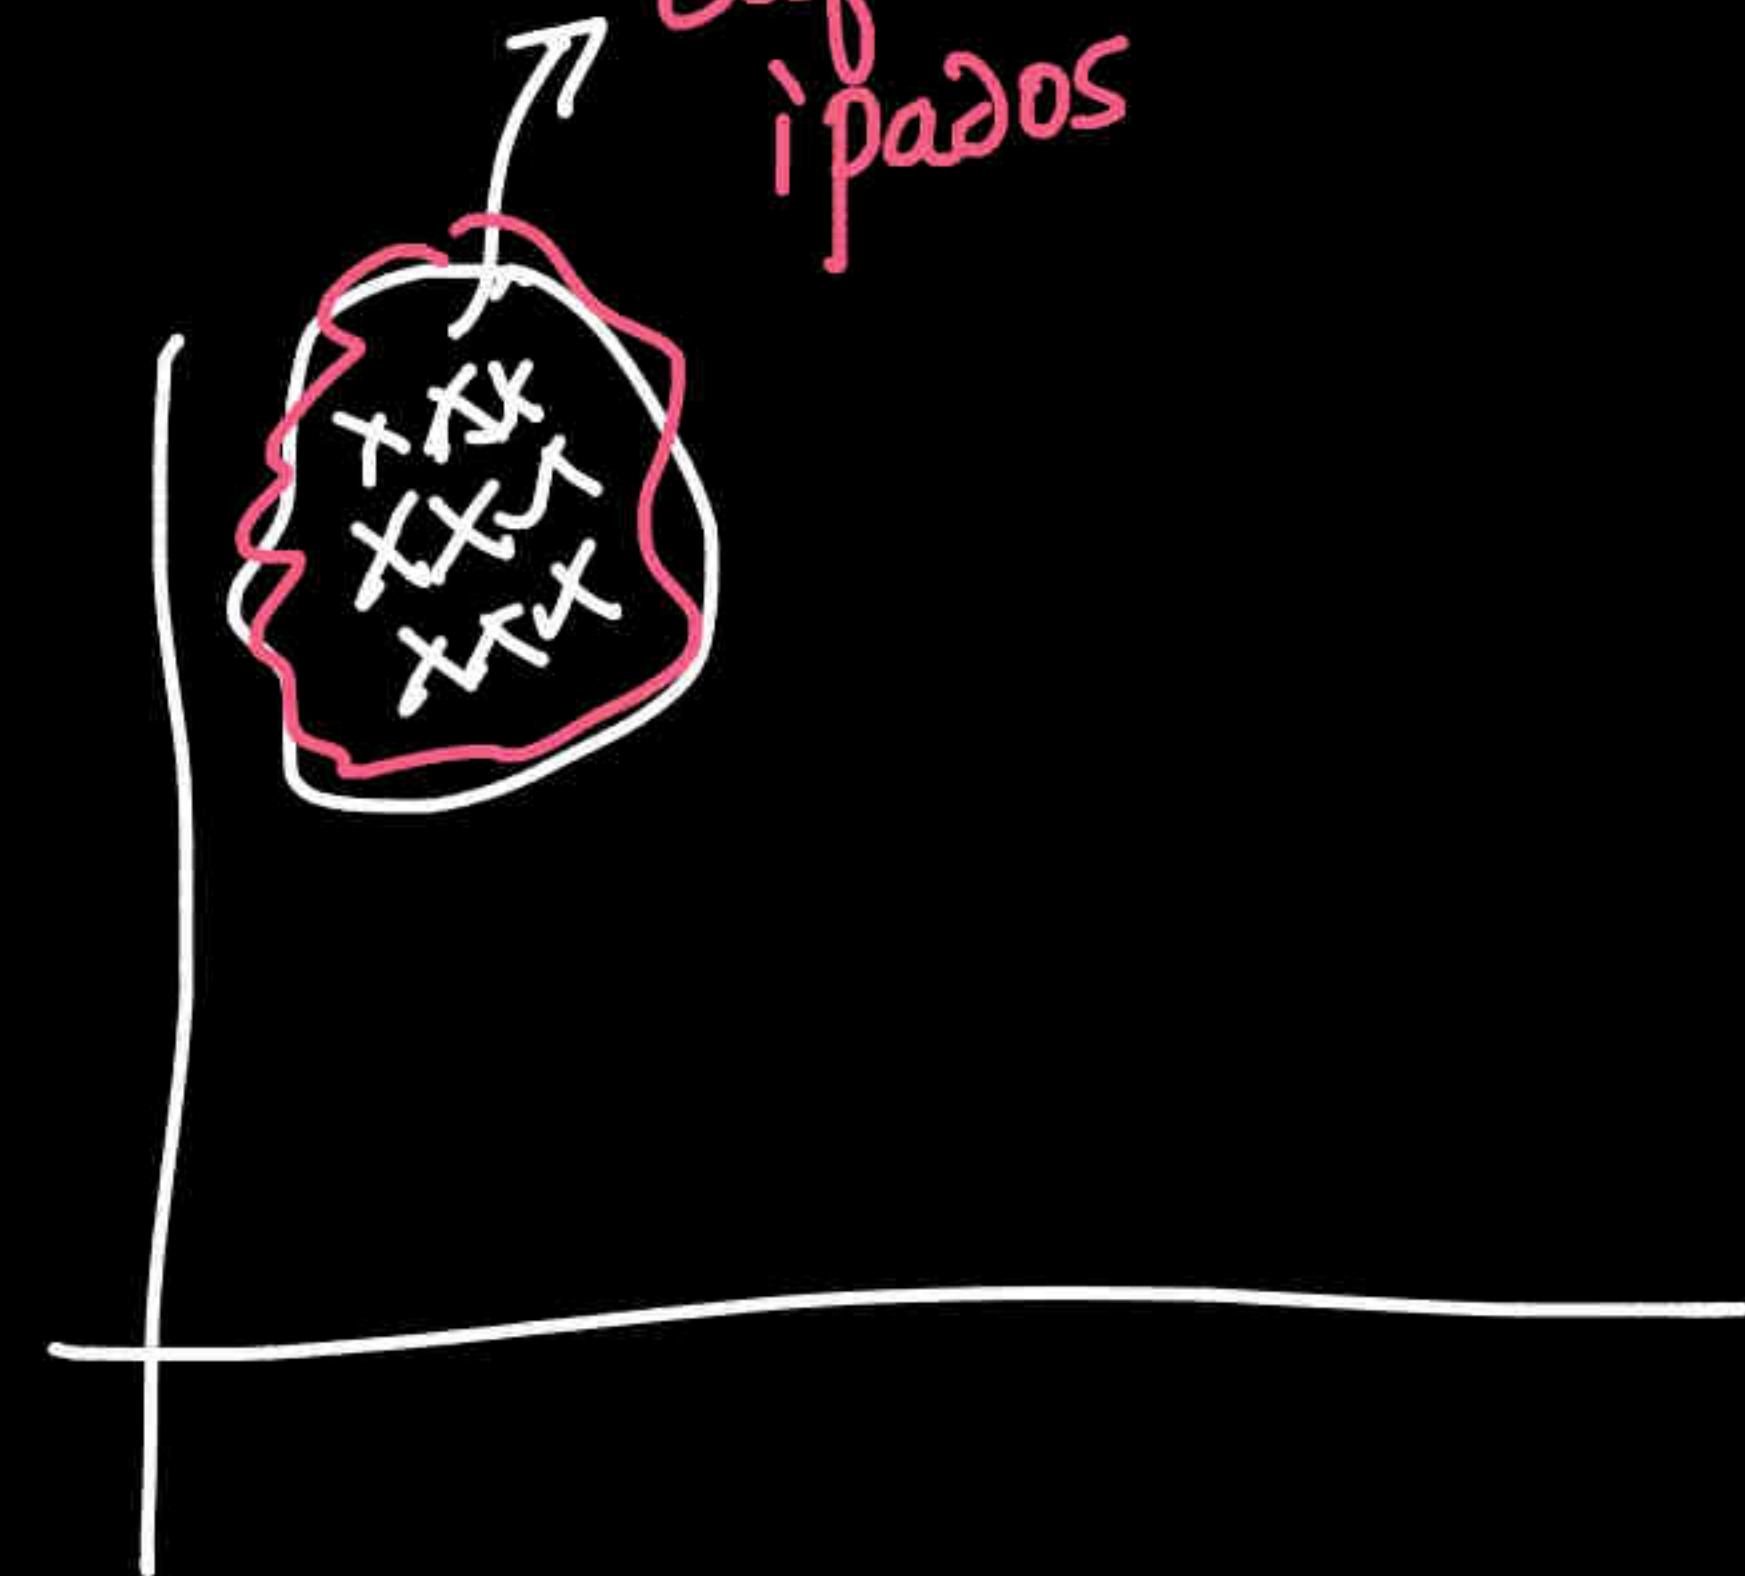
↑ classfn. loss

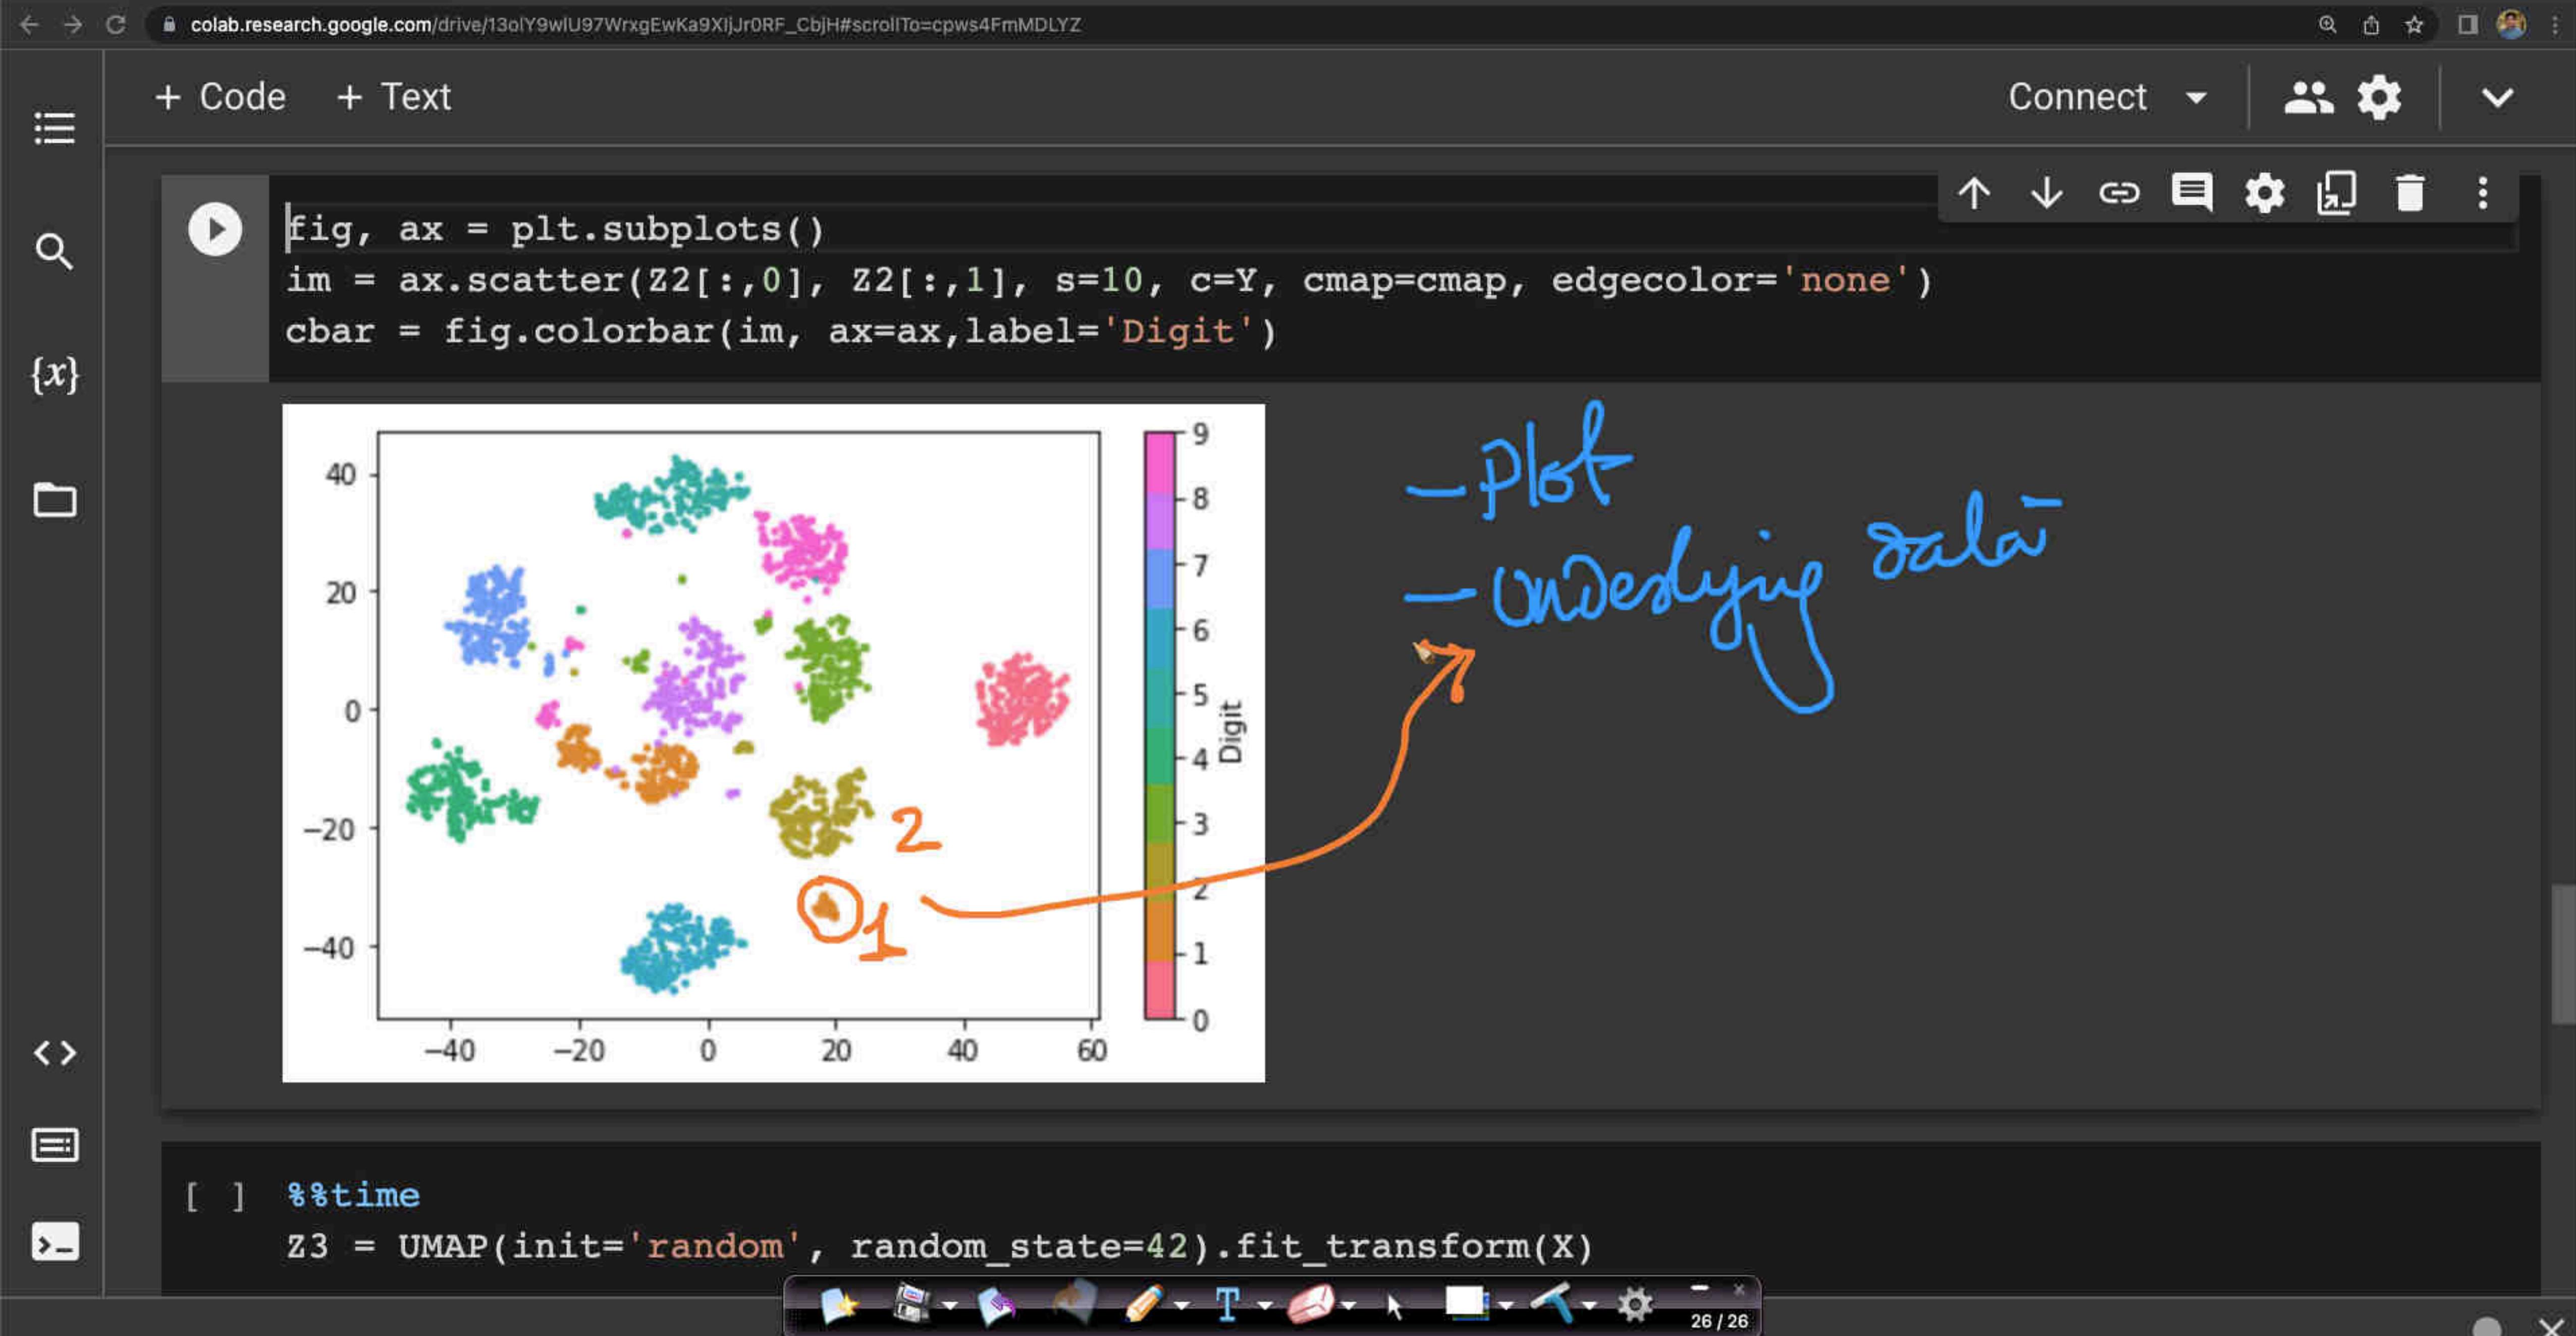


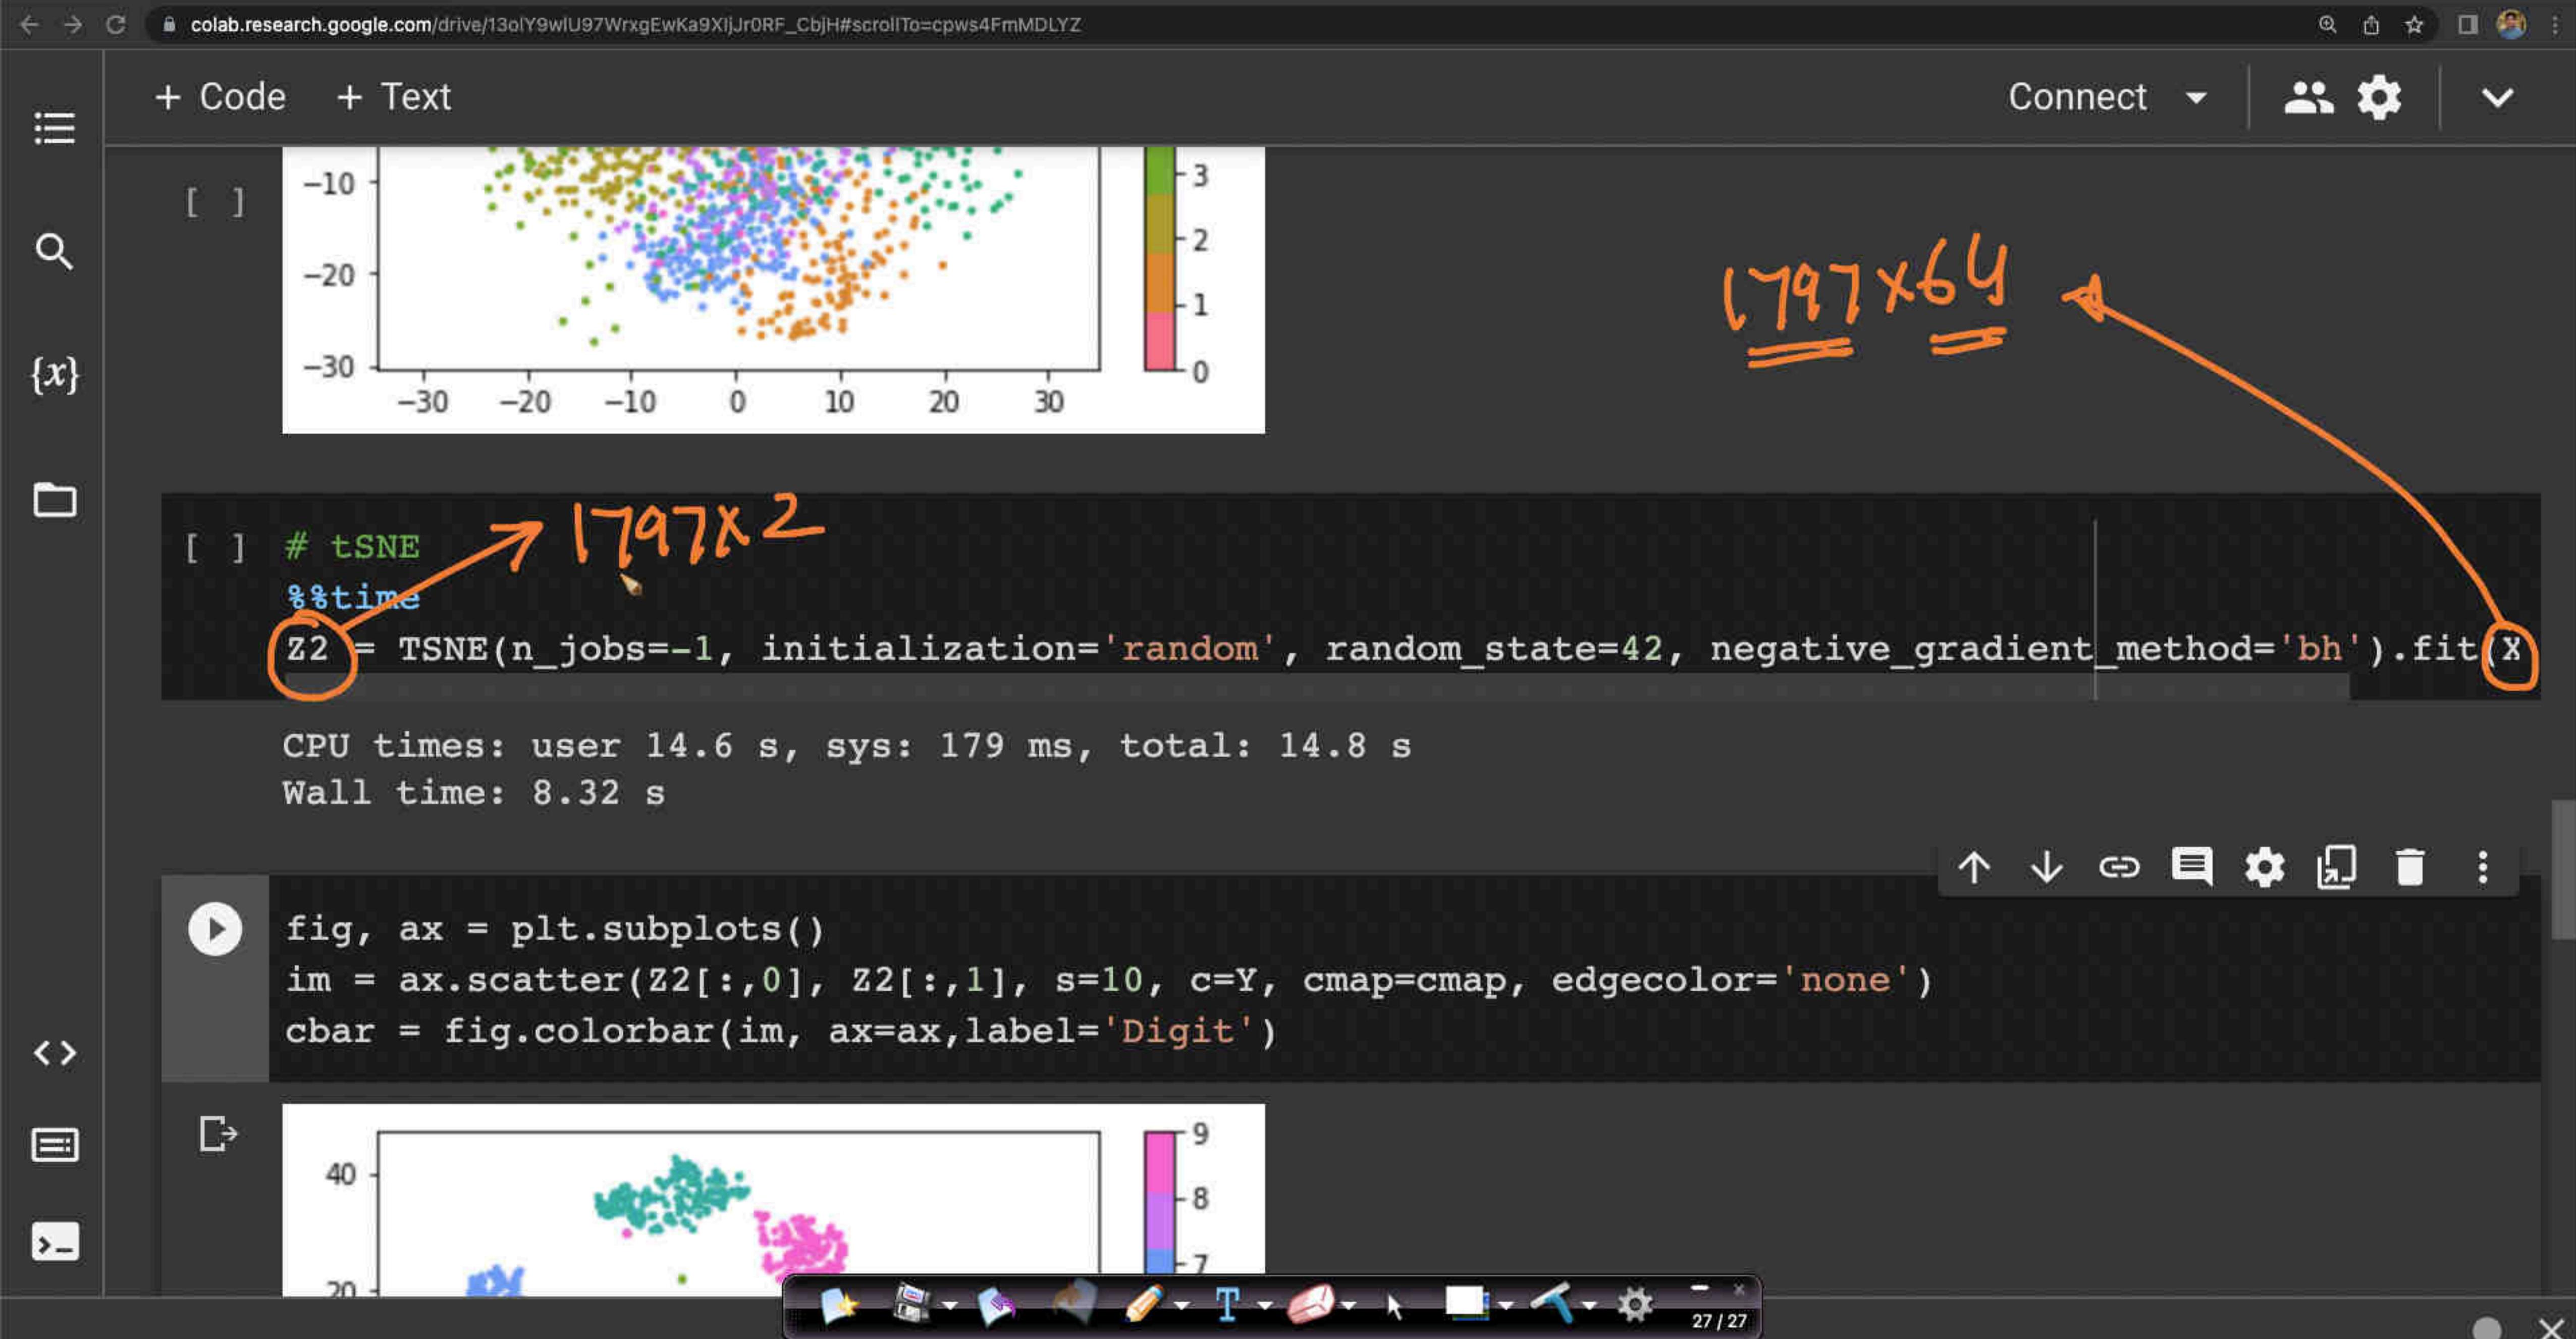
concretely:

Safari on  
ipados

ads on web-browser/  
app/mobile/lap







+ Code + Text

Connect

Q

{x}

b2

[ ] # tSNE  
%%time  
**b2** Z2 = TSNE(n\_jobs=-1, initialization='random', random\_state=42, negative\_gradient\_method='bh').fit(X)

CPU times: user 14.6 s, sys: 179 ms, total: 14.8 s  
Wall time: 8.32 s

[ ] fig, ax = plt.subplots()  
im = ax.scatter(Z2[:,0], Z2[:,1], s=10, c=Y, cmap=cmap, edgecolor='none')  
cbar = fig.colorbar(im, ax=ax, label='Digit')

40  
20  
0

9  
8  
7  
6  
5  
4  
3  
2  
1  
0

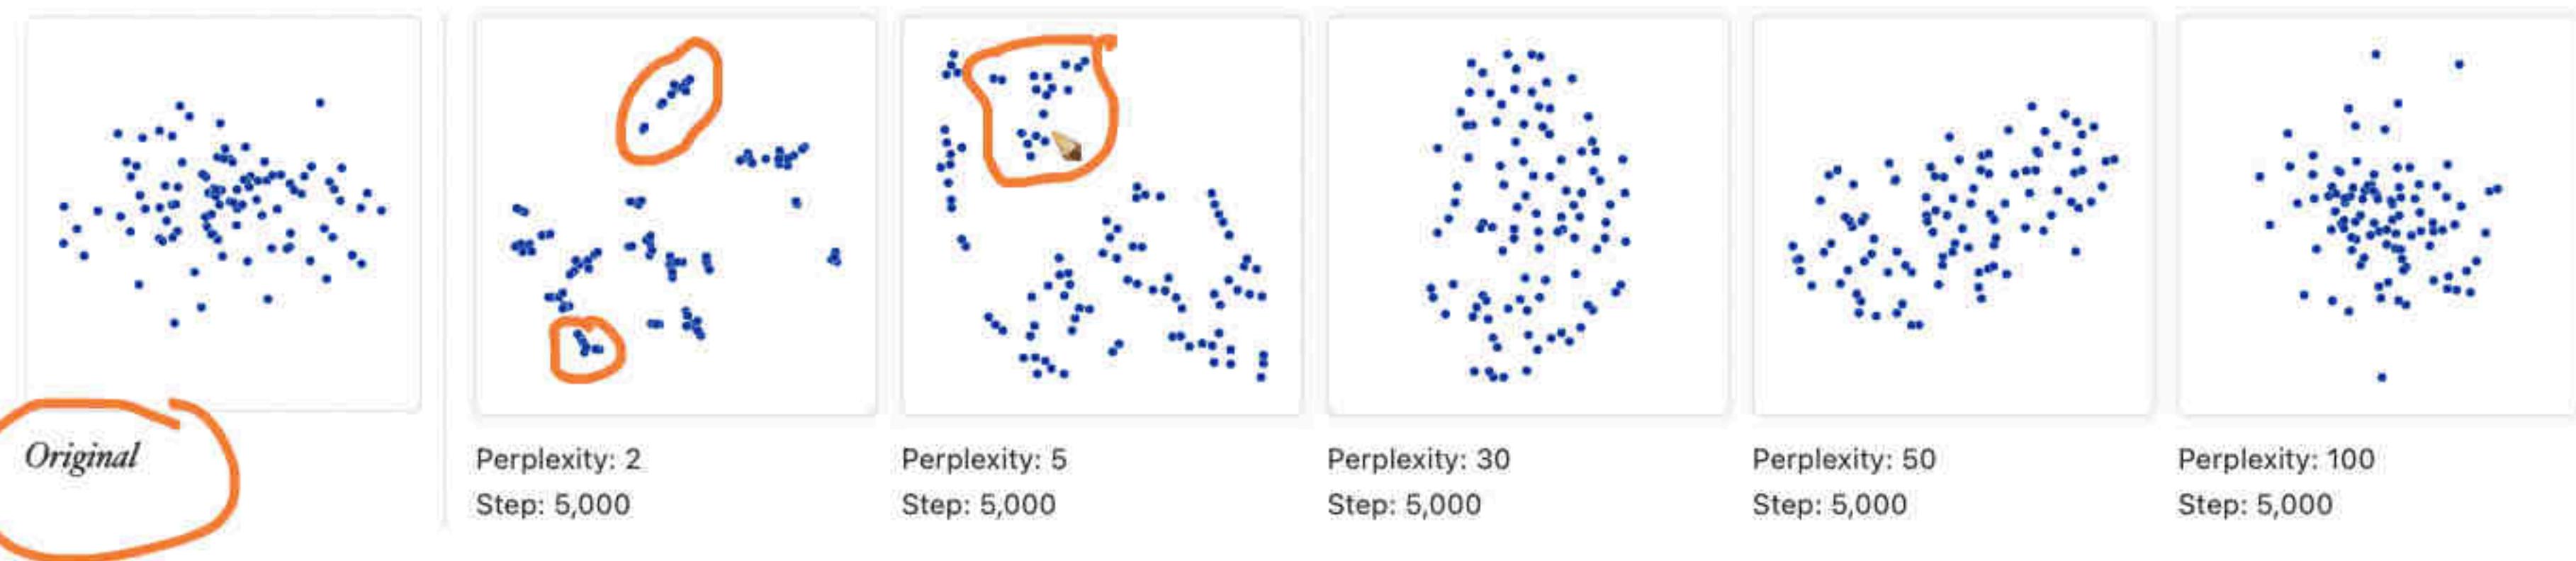
Digit

28 / 28

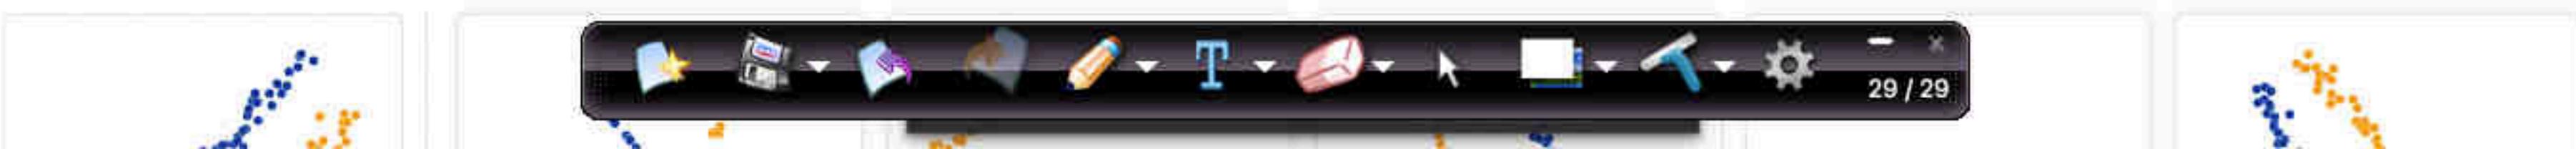
# Lesson

## 5. You can see some shapes, sometimes

It's rare for data to be distributed in a perfectly symmetric way. Let's take a look at an axis-aligned Gaussian distribution in 50 dimensions, where the standard deviation in coordinate  $i$  is  $1/i$ . That is, we're looking at a long-ish ellipsoidal cloud of points.

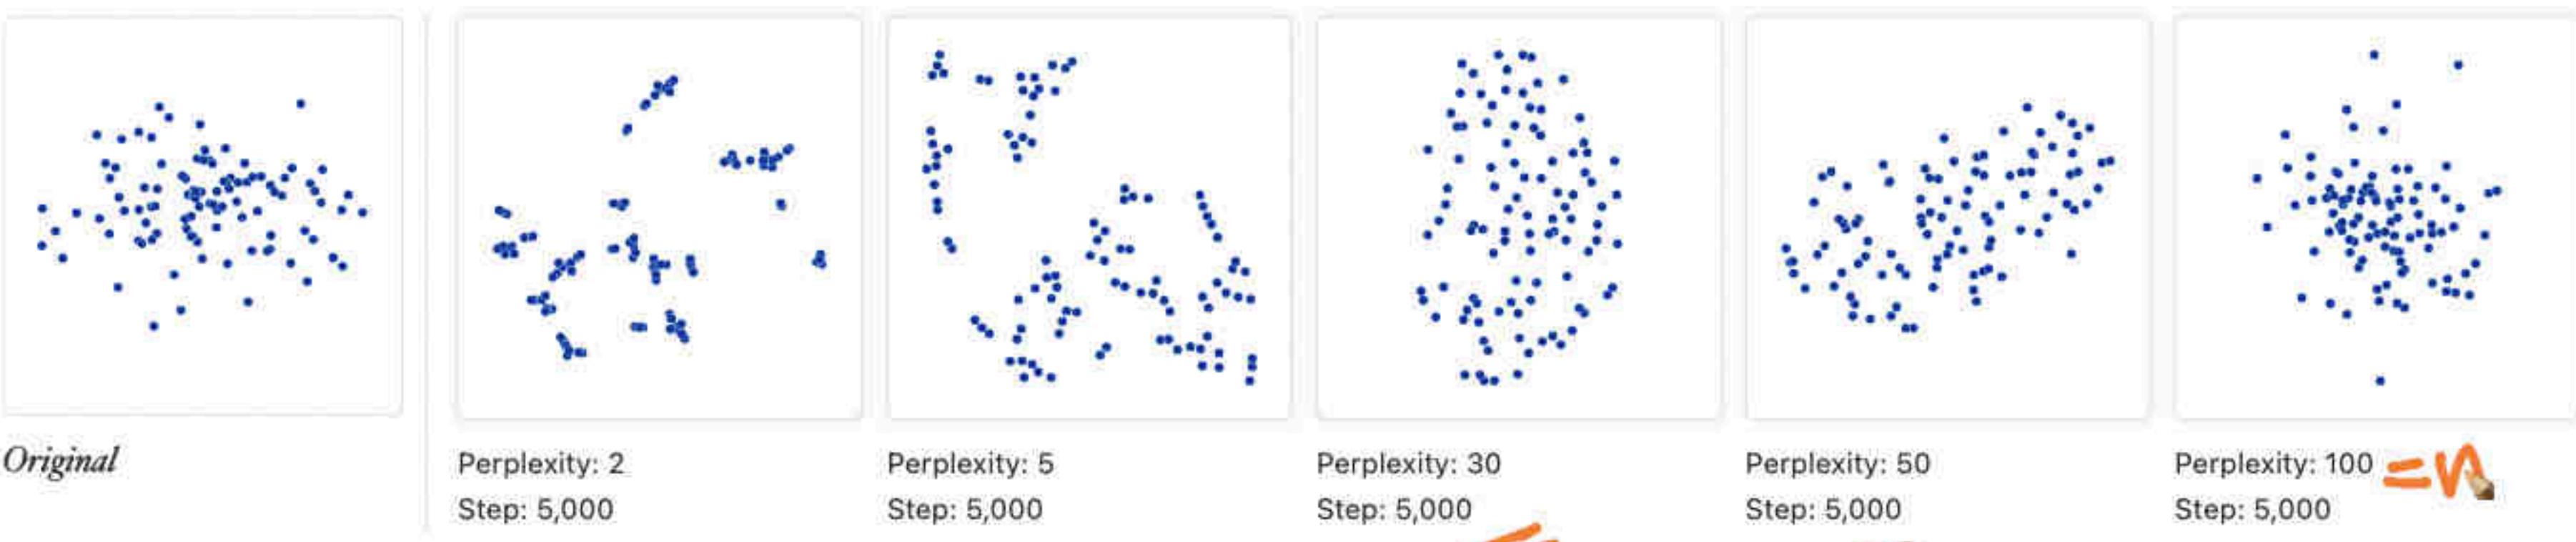


For high enough perplexity values, the elongated shapes are easy to read. On the other hand, at low perplexity, local effects and meaningless “clumping” take center stage. More extreme shapes also come through, but again only at the right perplexity. For example, here are two clusters of 75 points each in 2D, arranged in parallel lines with a bit of noise.



## 5. You can see some shapes, sometimes

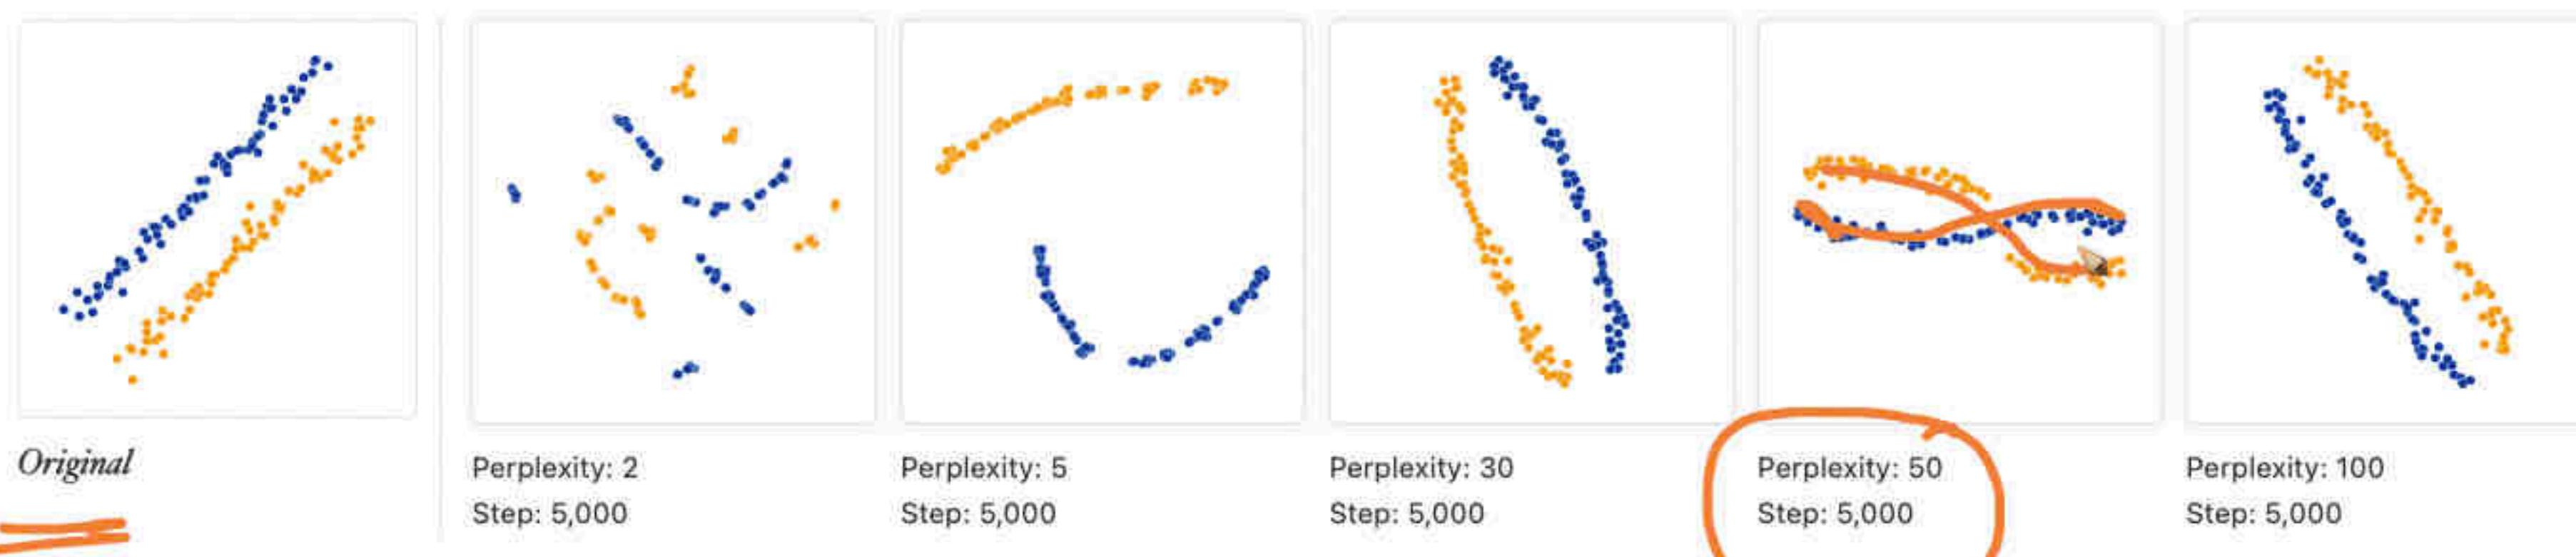
It's rare for data to be distributed in a perfectly symmetric way. Let's take a look at an axis-aligned Gaussian distribution in 50 dimensions, where the standard deviation in coordinate  $i$  is  $1/i$ . That is, we're looking at a long-ish ellipsoidal cloud of points.



For high enough perplexity values, the elongated shapes are easy to read. On the other hand, at low perplexity, local effects and meaningless “clumping” take center stage. More extreme shapes also come through, but again only at the right perplexity. For example, here are two clusters of 75 points each in 2D, arranged in parallel lines with a bit of noise.



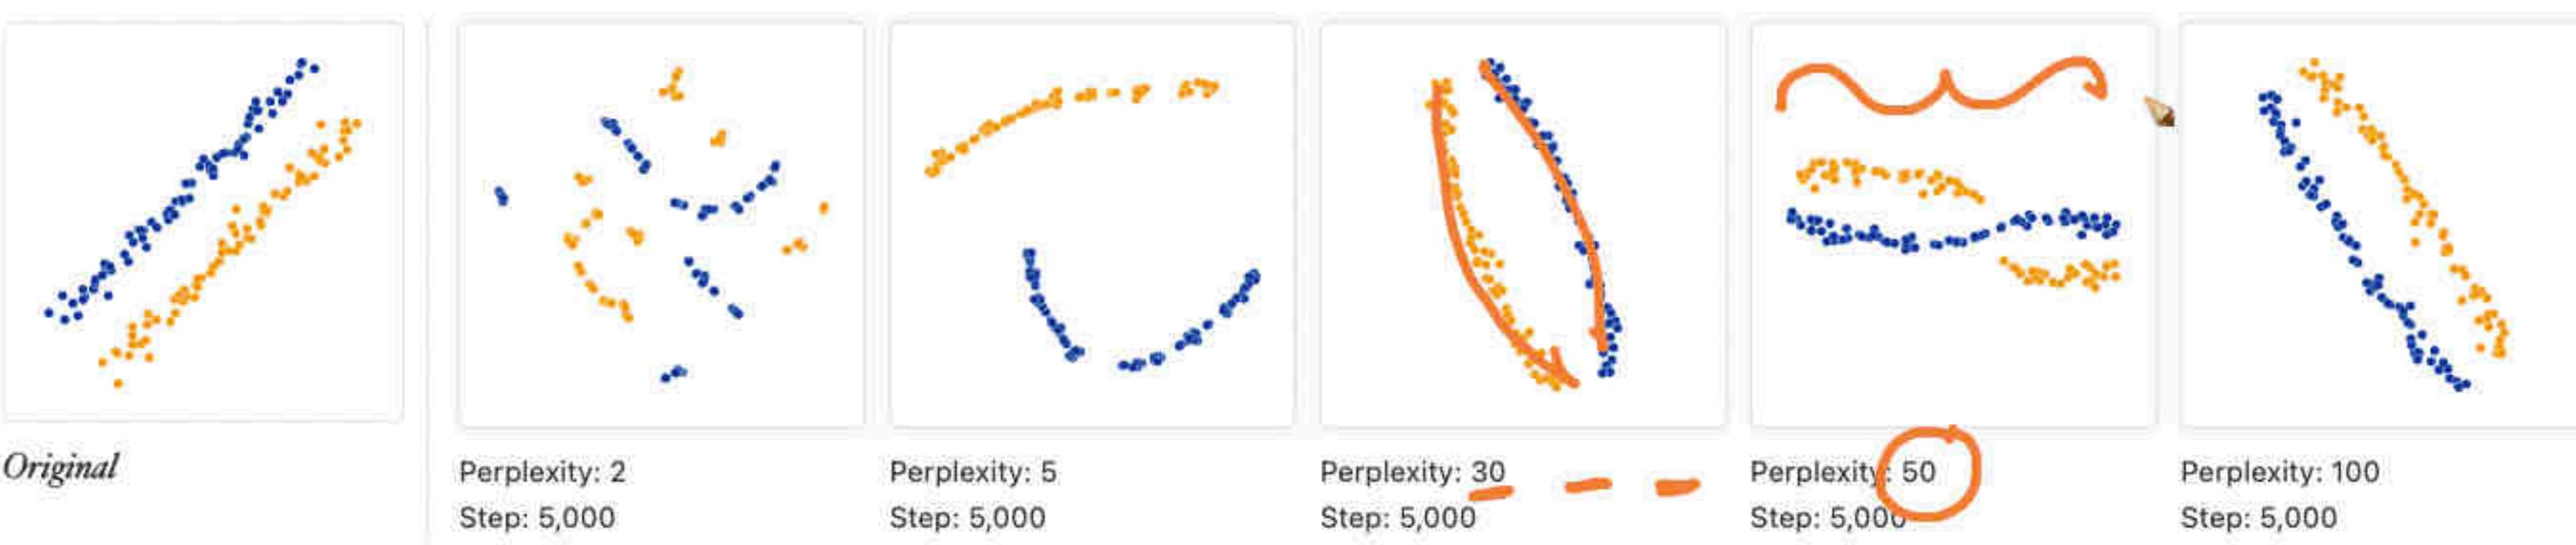
On the other hand, at low perplexity, local effects and meaningless “clumping” take center stage. More extreme shapes also come through, but again only at the right perplexity. For example, here are two clusters of 75 points each in 2D, arranged in parallel lines with a bit of noise.



For a certain range of perplexity the long clusters look close to correct, which is reassuring.

Even in the best cases, though, there's a subtle distortion: the lines are slightly curved outwards in the t-SNE diagram. The reason is that, as usual, t-SNE tends to expand denser regions of data. Since the middles of the clusters have less empty space around them than the ends, the algorithm magnifies them.

On the other hand, at low perplexity, local effects and meaningless “clumping” take center stage. More extreme shapes also come through, but again only at the right perplexity. For example, here are two clusters of 75 points each in 2D, arranged in parallel lines with a bit of noise.



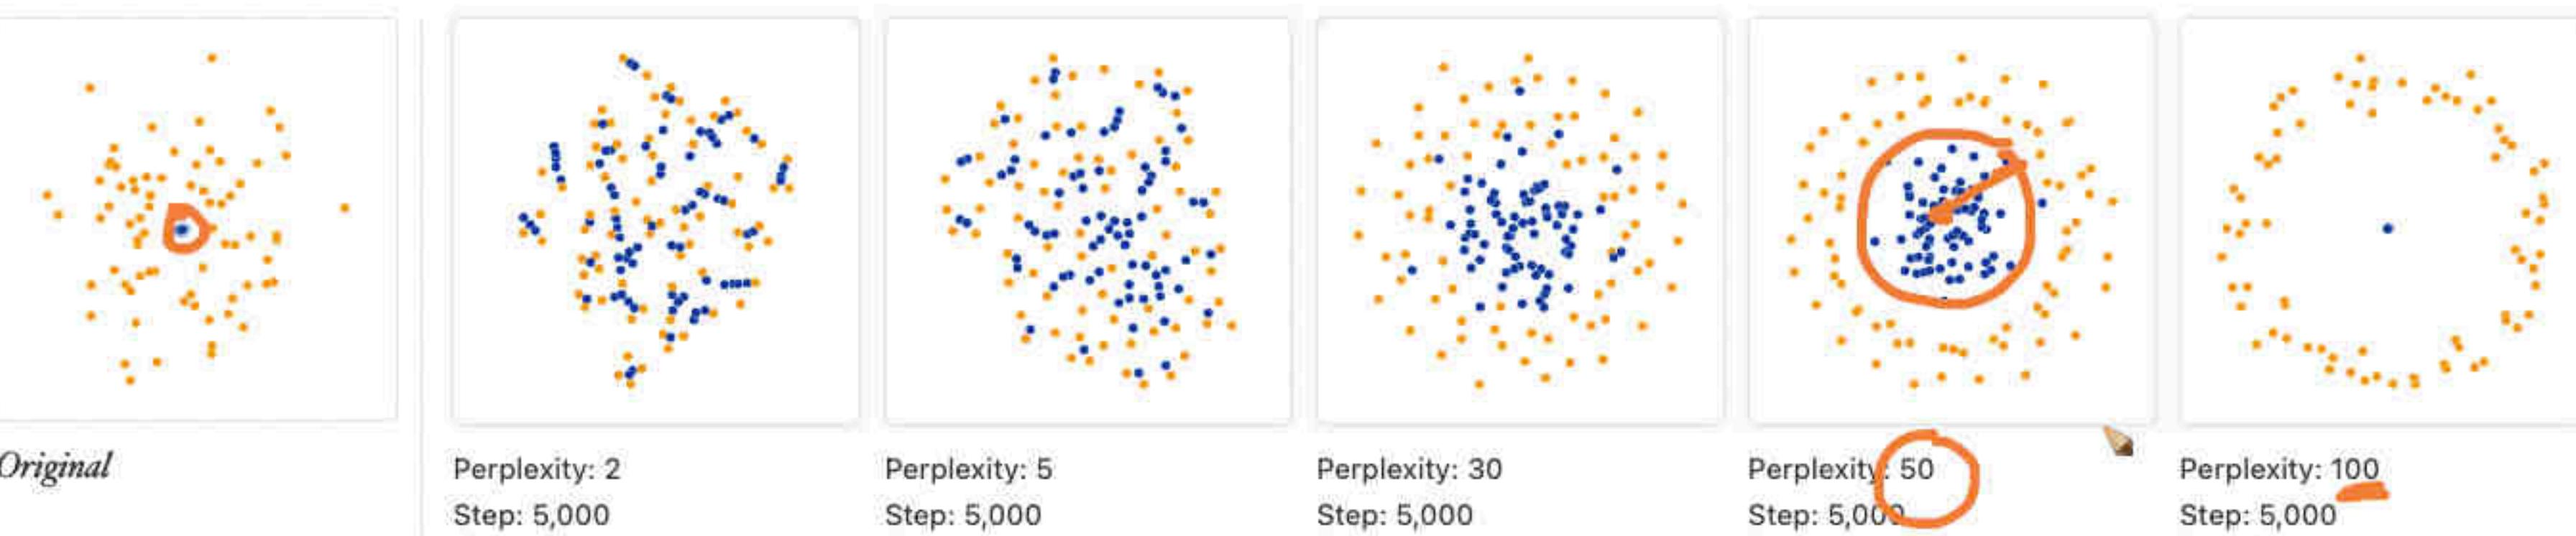
For a certain range of perplexity the long clusters look close to correct, which is reassuring.

Even in the best cases, though, there's a subtle distortion: the lines are slightly curved outwards in the t-SNE diagram. The reason is that, as usual, t-SNE tends to expand denser regions of data. Since the middles of the clusters have less empty space around them than the ends, the algorithm magnifies them.

[distill.pub/2016/misread-tsne/](https://distill.pub/2016/misread-tsne/)

For topology, you may need more than one plot

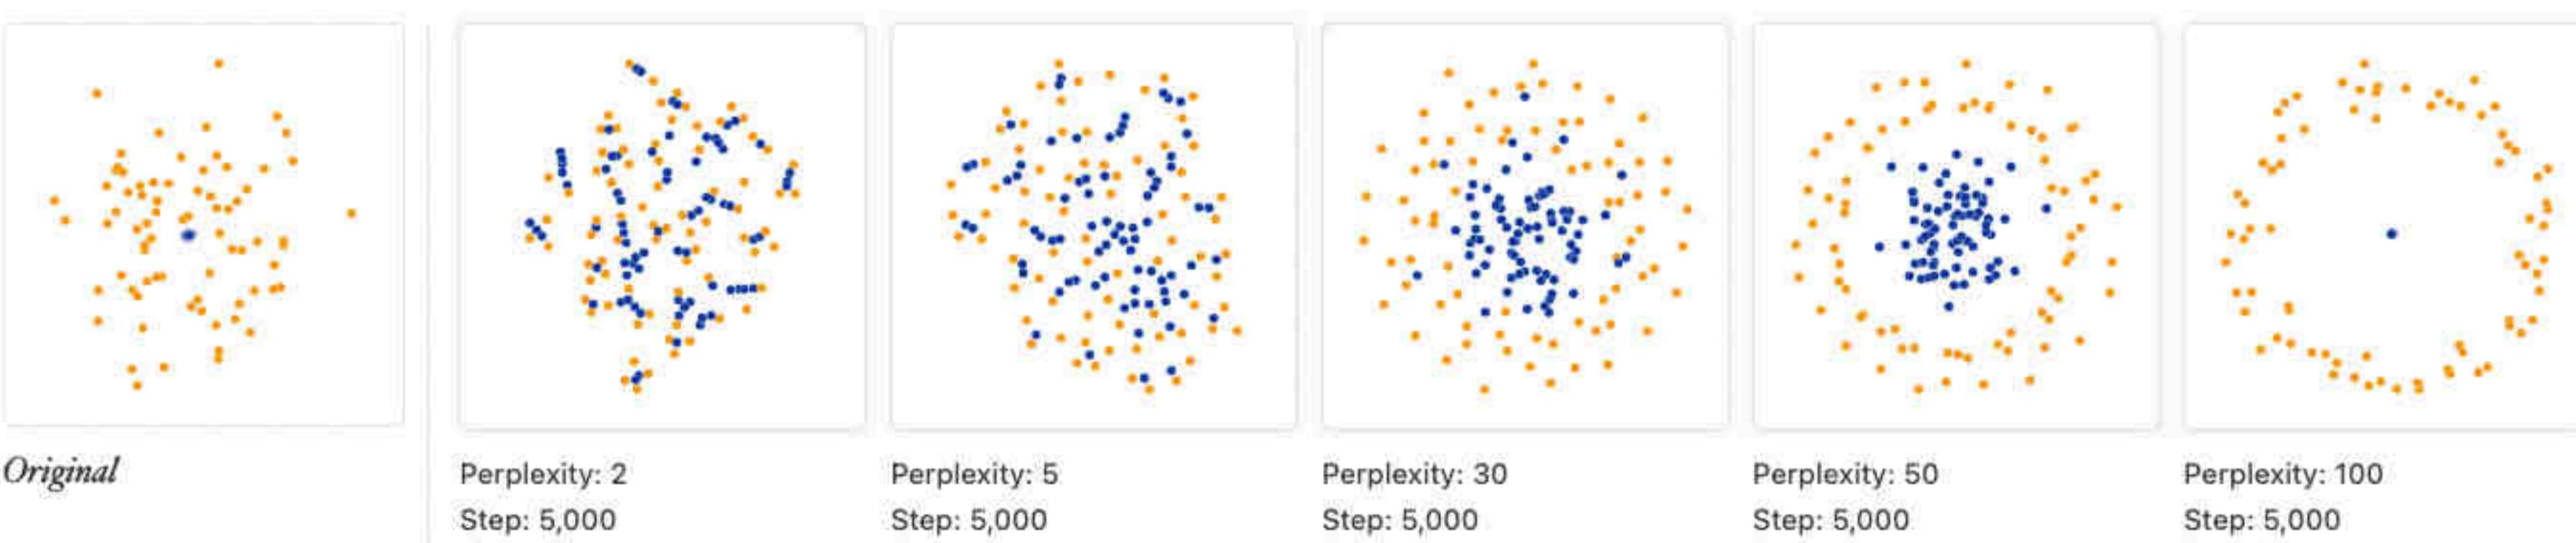
Sometimes you can read topological information off a t-SNE plot, but that typically requires views at multiple perplexities. One of the simplest topological properties is containment. The plots below show two groups of 75 points in 50 dimensional space. Both are sampled from symmetric Gaussian distributions centered at the origin, but one is 50 times more tightly dispersed than the other. The “small” distribution is in effect contained in the large one.



The perplexity 30 view shows the basic topology correctly, but again t-SNE greatly exaggerates the size of the smaller group of points. At perplexity 50, there's a new phenomenon: the outer group becomes a circle, as the plot tries to depict the fact that all its points are about the same distance from the inner group. If you looked at this image alone, it would be easy to misread these outer points as a one-dimensional structure.

For topology, you may need more than one plot

Sometimes you can read topological information off a t-SNE plot, but that typically requires views at multiple perplexities. One of the simplest topological properties is containment. The plots below show two groups of 75 points in 50 dimensional space. Both are sampled from symmetric Gaussian distributions centered at the origin, but one is 50 times more tightly dispersed than the other. The “small” distribution is in effect contained in the large one.



The perplexity 30 view shows the basic topology correctly, but again t-SNE greatly exaggerates the size of the smaller group of points. At perplexity 50, there's a new phenomenon: the outer group becomes a circle, as the plot tries to depict the fact that all its points are about the same distance from the inner group. If you looked at this image alone, it would be easy to misread these outer points as a one-dimensional structure.

→ various perplexity values  
+ t-SNE + data-analysis  
look @ the data  
↓  
Conclusion

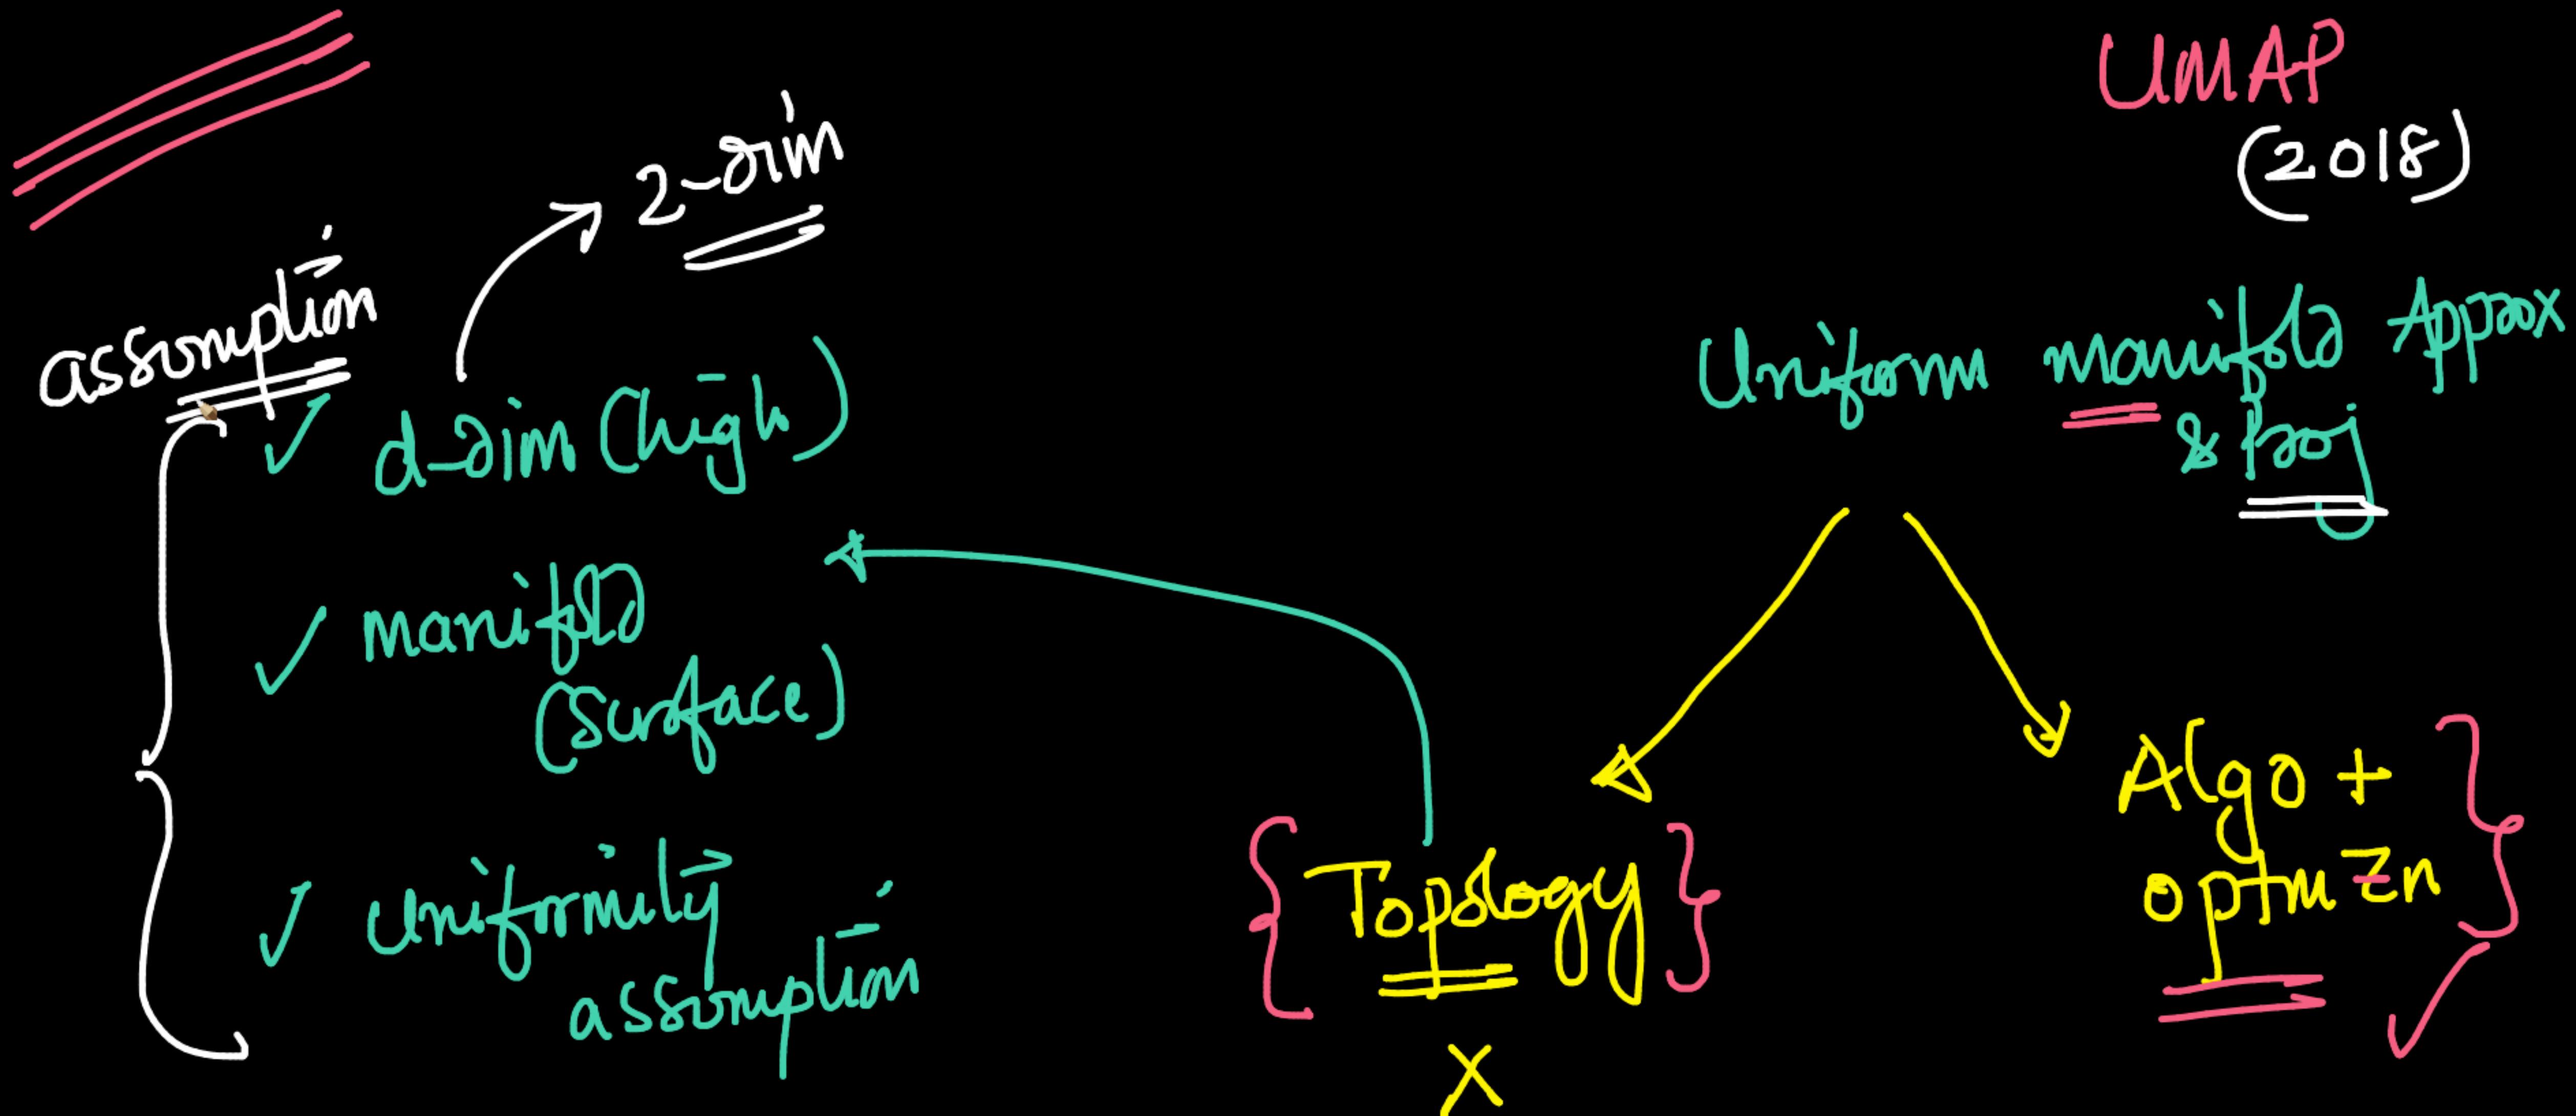
## Conclusion

There's a reason that t-SNE has become so popular: it's incredibly flexible, and can often find structure where other dimensionality-reduction algorithms cannot. Unfortunately, that very flexibility makes it tricky to interpret. Out of sight from the user, the algorithm makes all sorts of adjustments that tidy up its visualizations. Don't let the hidden "magic" scare you away from the whole technique, though. The good news is that by studying how t-SNE behaves in simple cases, it's possible to develop an intuition for what's going on.

## Acknowledgments

We are grateful to Chris Olah and Shan Carter for creating this platform, and for excellent design and editorial help from Shan Carter. Daniel Smilkov, James Wexler, and Chi Zeng provided many helpful comments. We also thank Andrej Karpathy for creating the `tsnejs` library used in the interactive diagrams.

UMAP  
(2018)



~~Algo + Opt~~

①

Finds

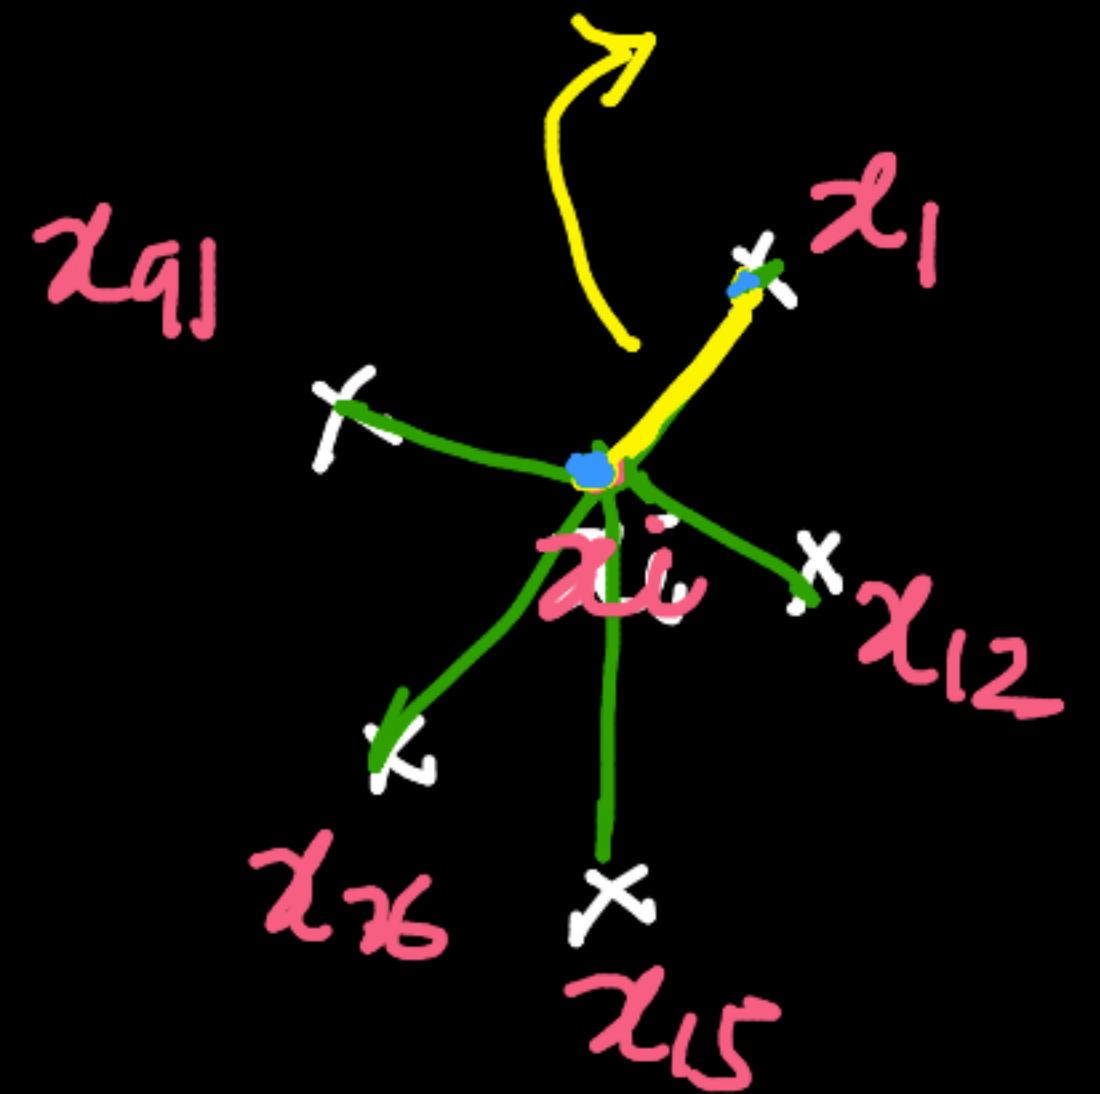
nearest neighbors

n-neighbors = 5

② Builds a graph (d-dim) (weighted)

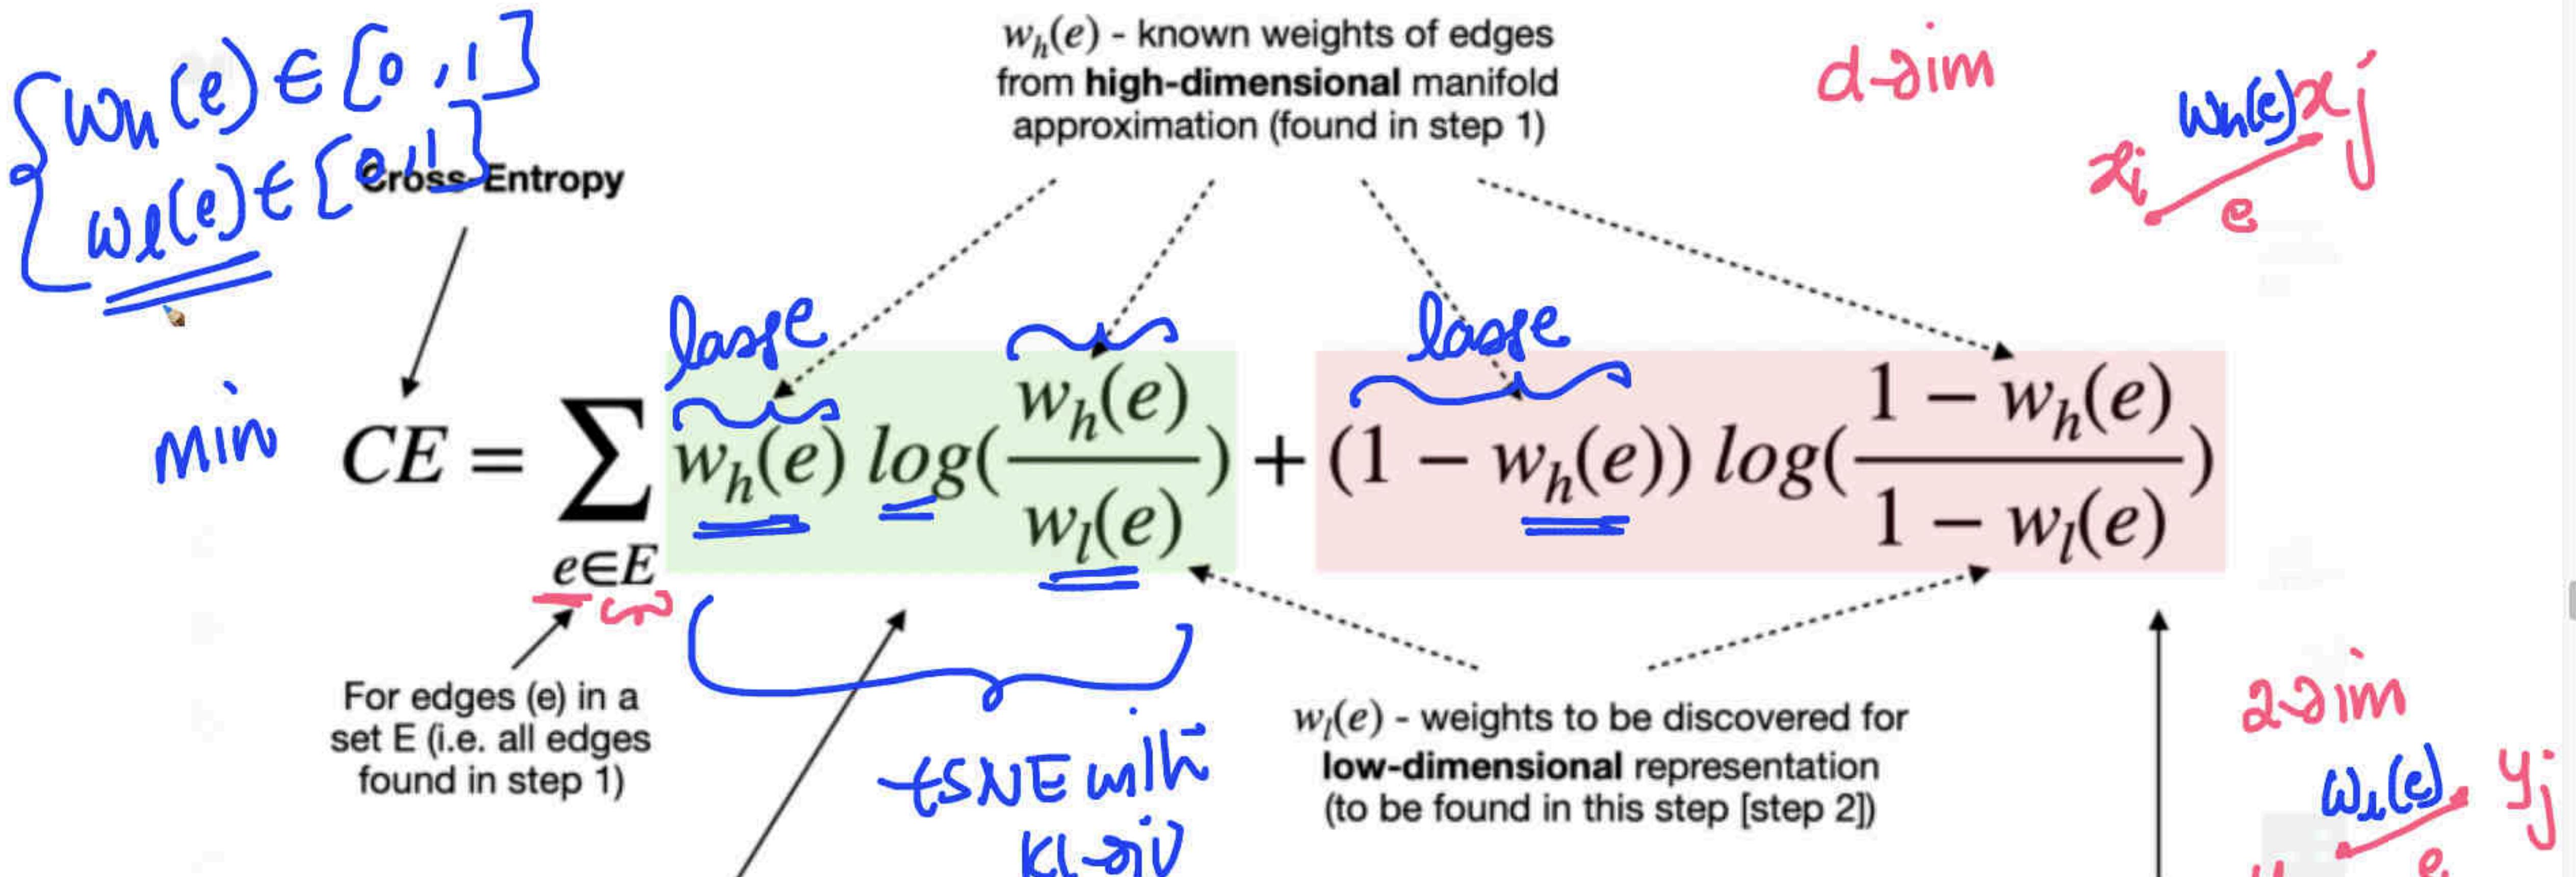
$$G \xrightarrow{G^{(d)}} W_{ij}^d \xrightarrow{W_{ij}^d} \text{dist}^d(x_i, x_j)$$

d-dim



- nodes: places
- edges: roads
- weights: distance

③ for each  $x_i$   $\xrightarrow{\text{Part}}$   $y_i$   
 $d$ -dim  $\quad \quad \quad 2$ -dim



The first term acts as an "**attractive force**" whenever there is a large weight associated to the high-dimensional case. This is because this term will be minimized when  $w_l(e)$  is as large as possible, which will occur when the distance between the points is as small as possible.

The second term acts as a "**repulsive force**" whenever high-dimensional weight  $w_h(e)$  is small. This is because the term will be minimized by making  $w_l(e)$  as small as possible.

On balance the push and pull of the two terms will allow the low-dimensional representation settle into a state that matches the original data.

**Cross-Entropy**

$$CE = \sum_{e \in E} w_h(e) \log\left(\frac{w_h(e)}{w_l(e)}\right) + (1 - w_h(e)) \log\left(\frac{1 - w_h(e)}{1 - w_l(e)}\right)$$

*w<sub>h</sub>(e)* - known weights of edges from **high-dimensional** manifold approximation (found in step 1)

*w<sub>l</sub>(e)* - weights to be discovered for **low-dimensional** representation (to be found in this step [step 2])

*w<sub>h</sub>(e)* small

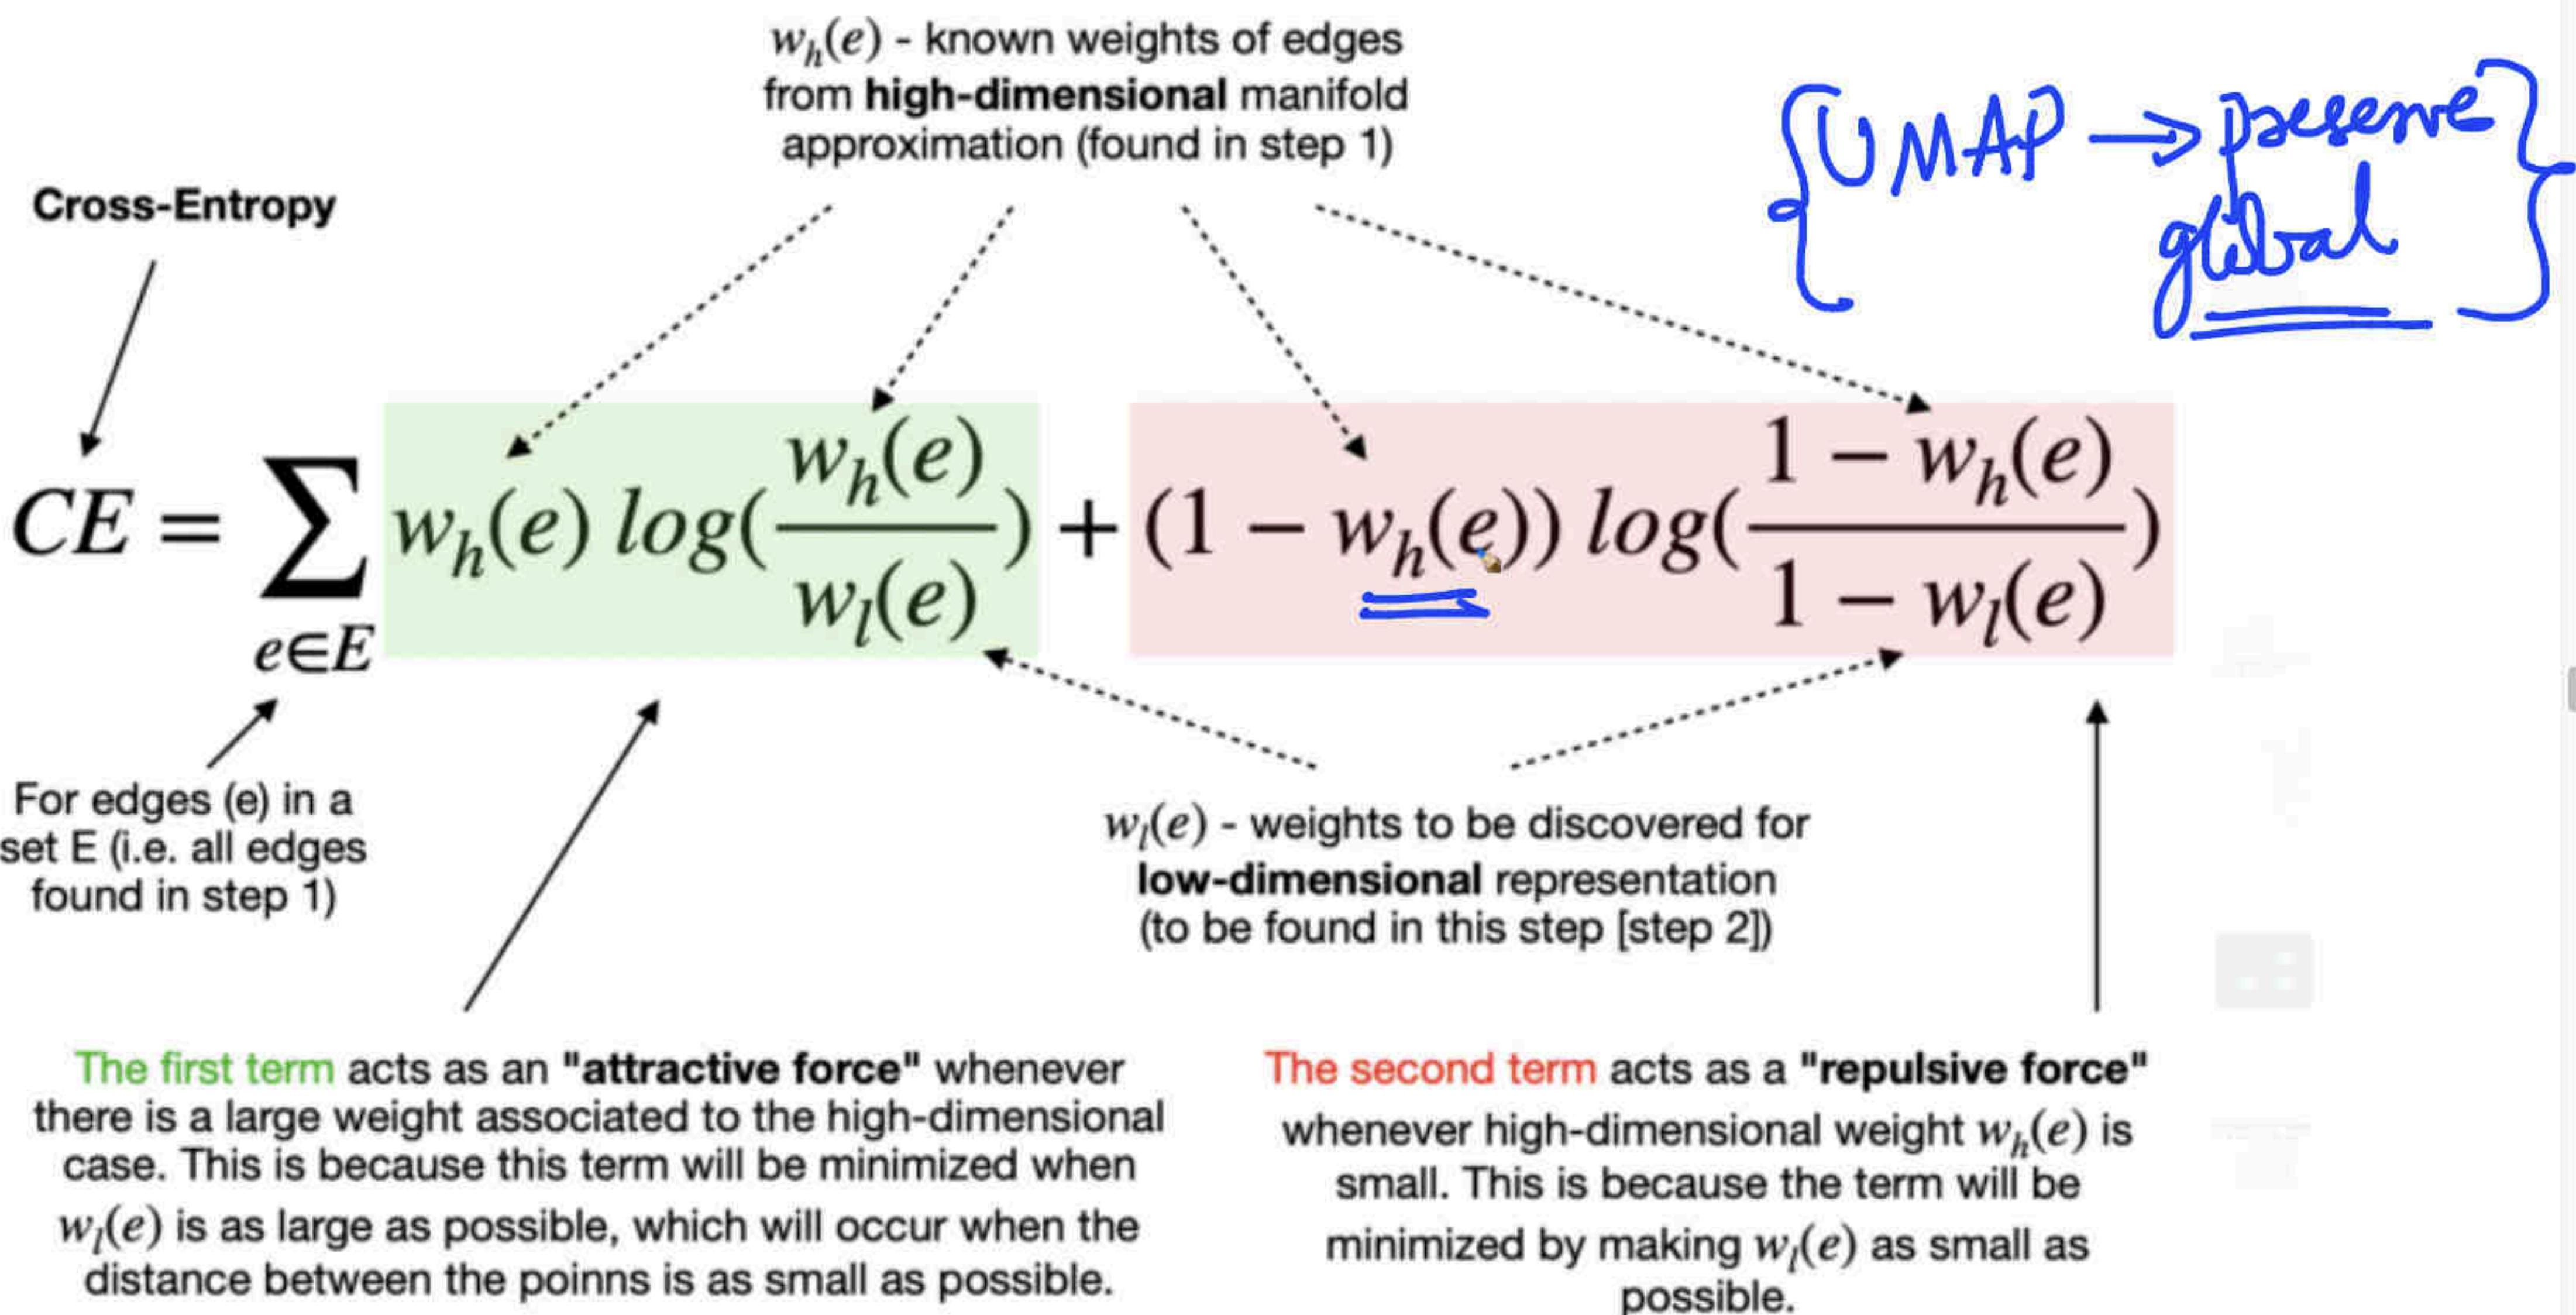
For edges ( $e$ ) in a set  $E$  (i.e. all edges found in step 1)

The first term acts as an "**attractive force**" whenever there is a large weight associated to the high-dimensional case. This is because this term will be minimized when  $w_l(e)$  is as large as possible, which will occur when the distance between the points is as small as possible.

$w_l(e)$  - weights to be discovered for **low-dimensional** representation (to be found in this step [step 2])

The second term acts as a "**repulsive force**" whenever high-dimensional weight  $w_h(e)$  is small. This is because the term will be minimized by making  $w_l(e)$  as small as possible.

On balance the push and pull of the two terms will allow the low-dimensional representation settle into a state that matches the original data.



On balance the push and pull of the two terms in the loss function cause the low-dimensional representation settle into a state that preserves the global structure of the original data.

$w_h(e)$  - known weights of edges from **high-dimensional** manifold approximation (found in step 1)

**Cross-Entropy**

$$CE = \sum_{e \in E} w_h(e) \log\left(\frac{w_h(e)}{w_l(e)}\right) + (1 - w_h(e)) \log\left(\frac{1 - w_h(e)}{1 - w_l(e)}\right)$$

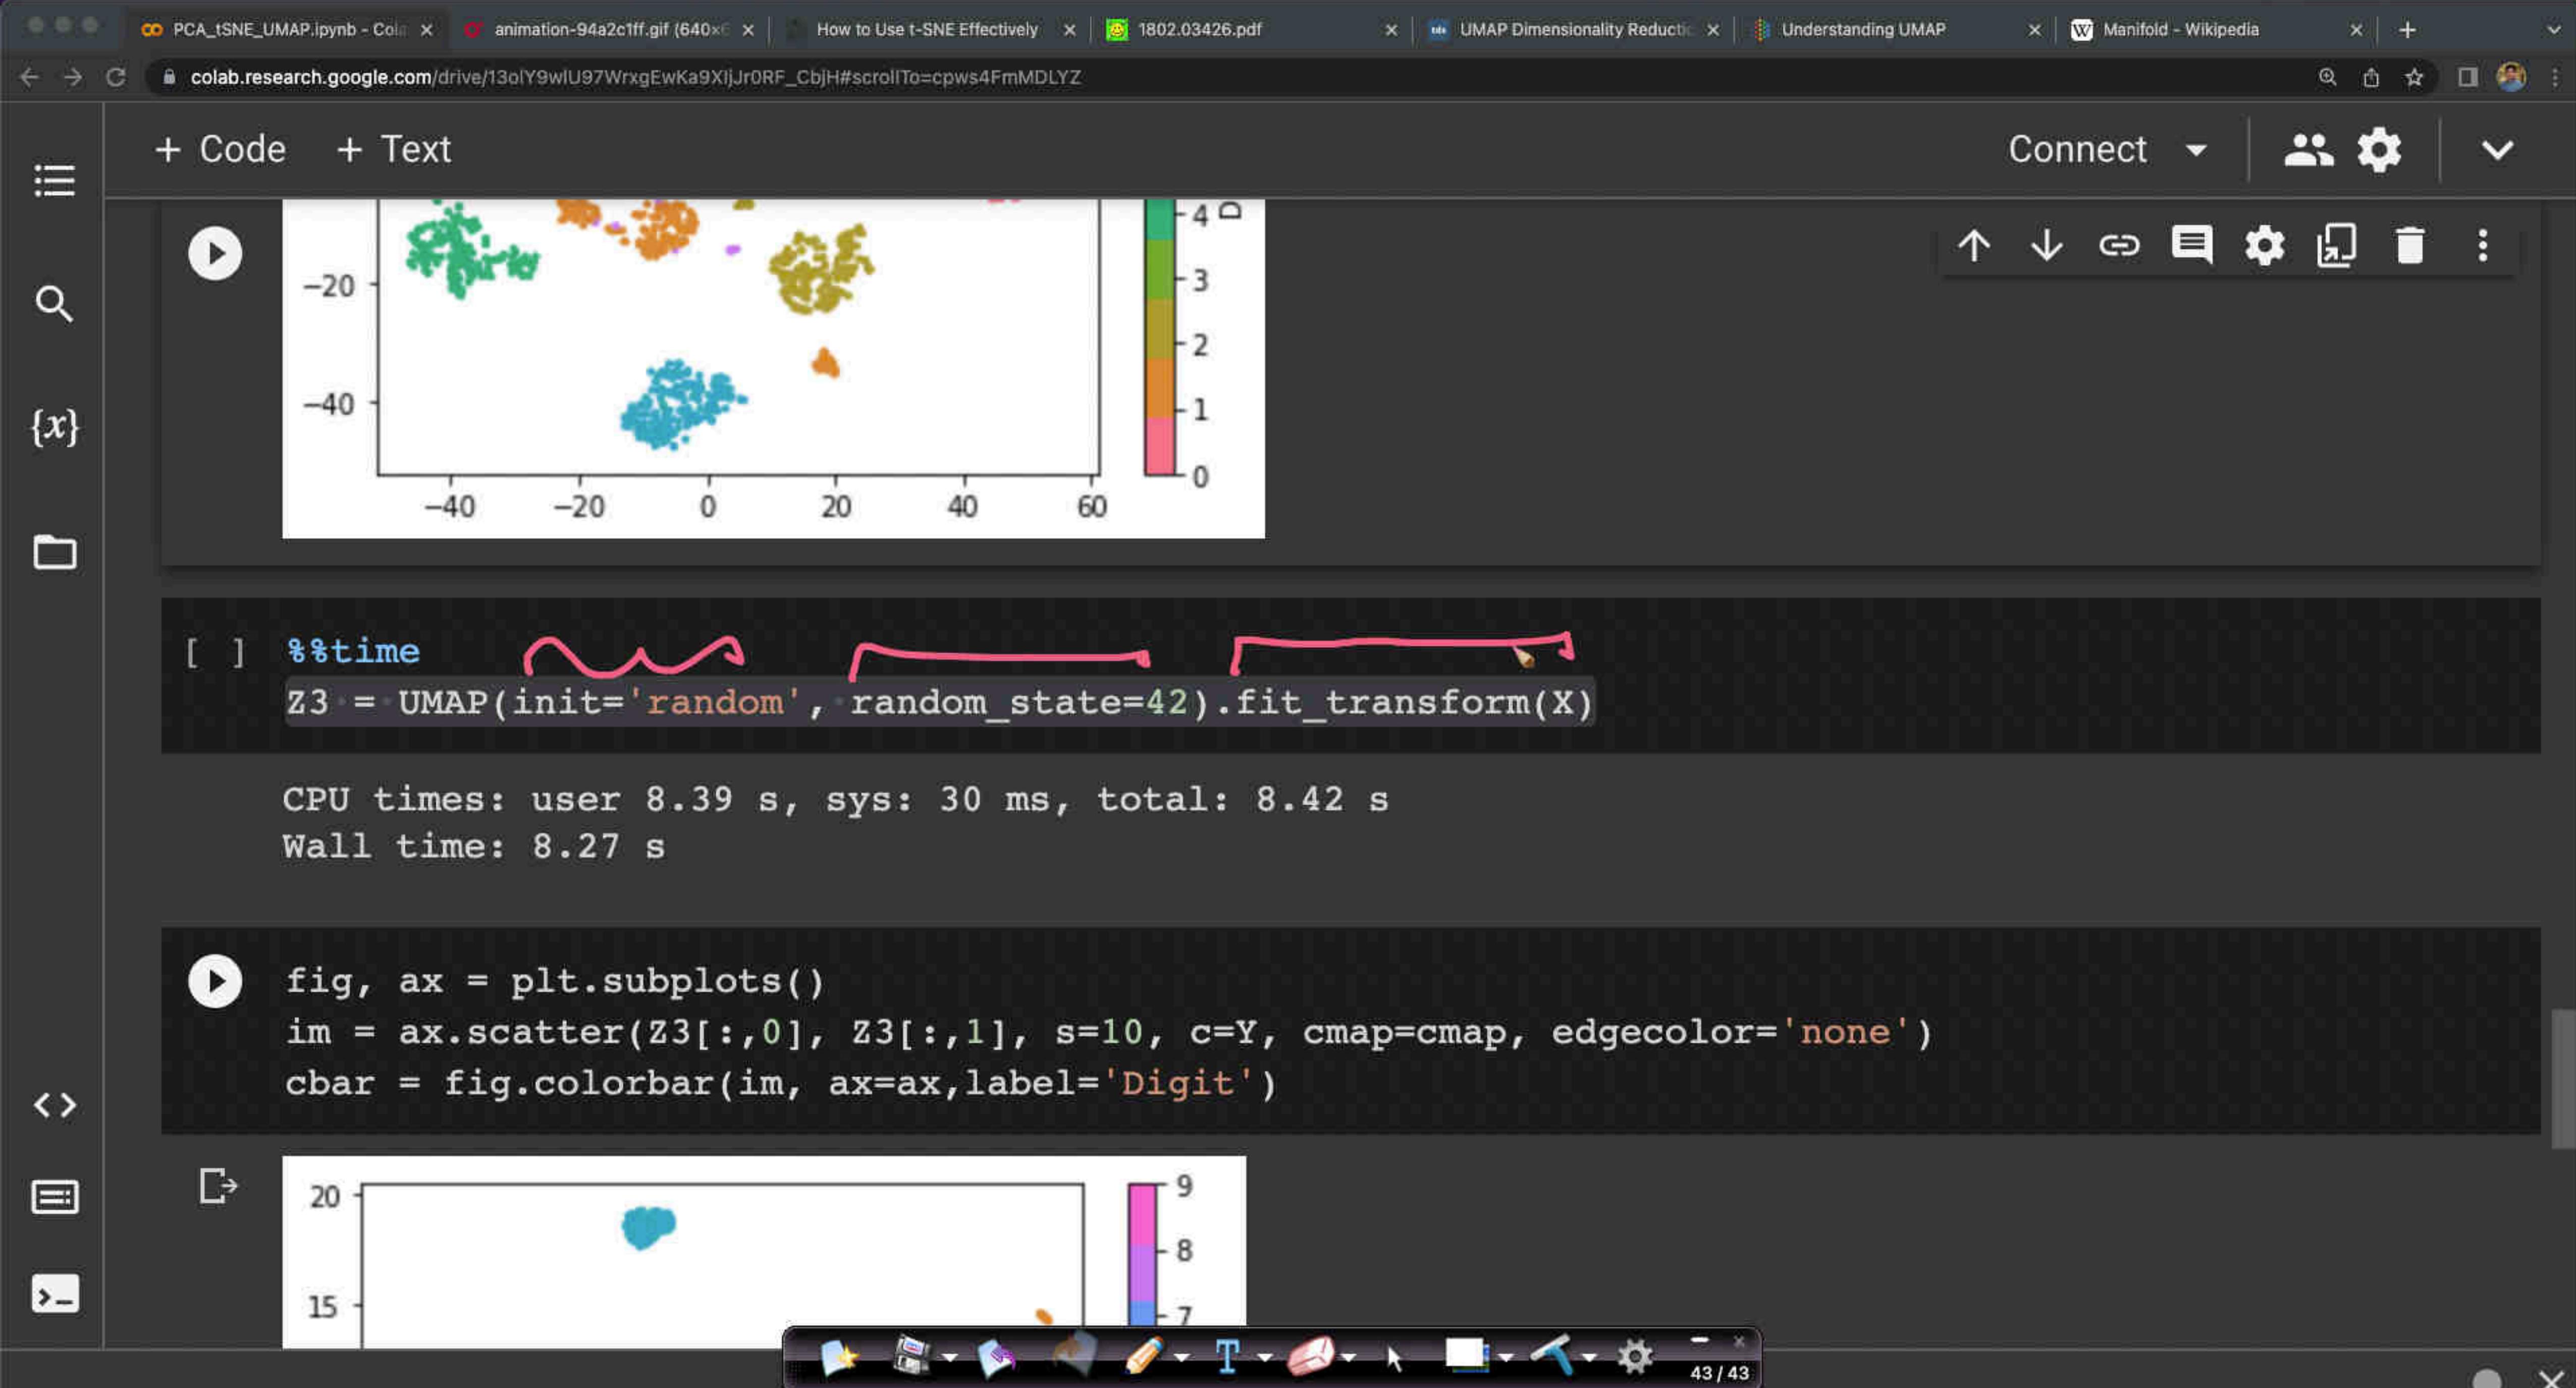
For edges ( $e$ ) in a set  $E$  (i.e. all edges found in step 1)

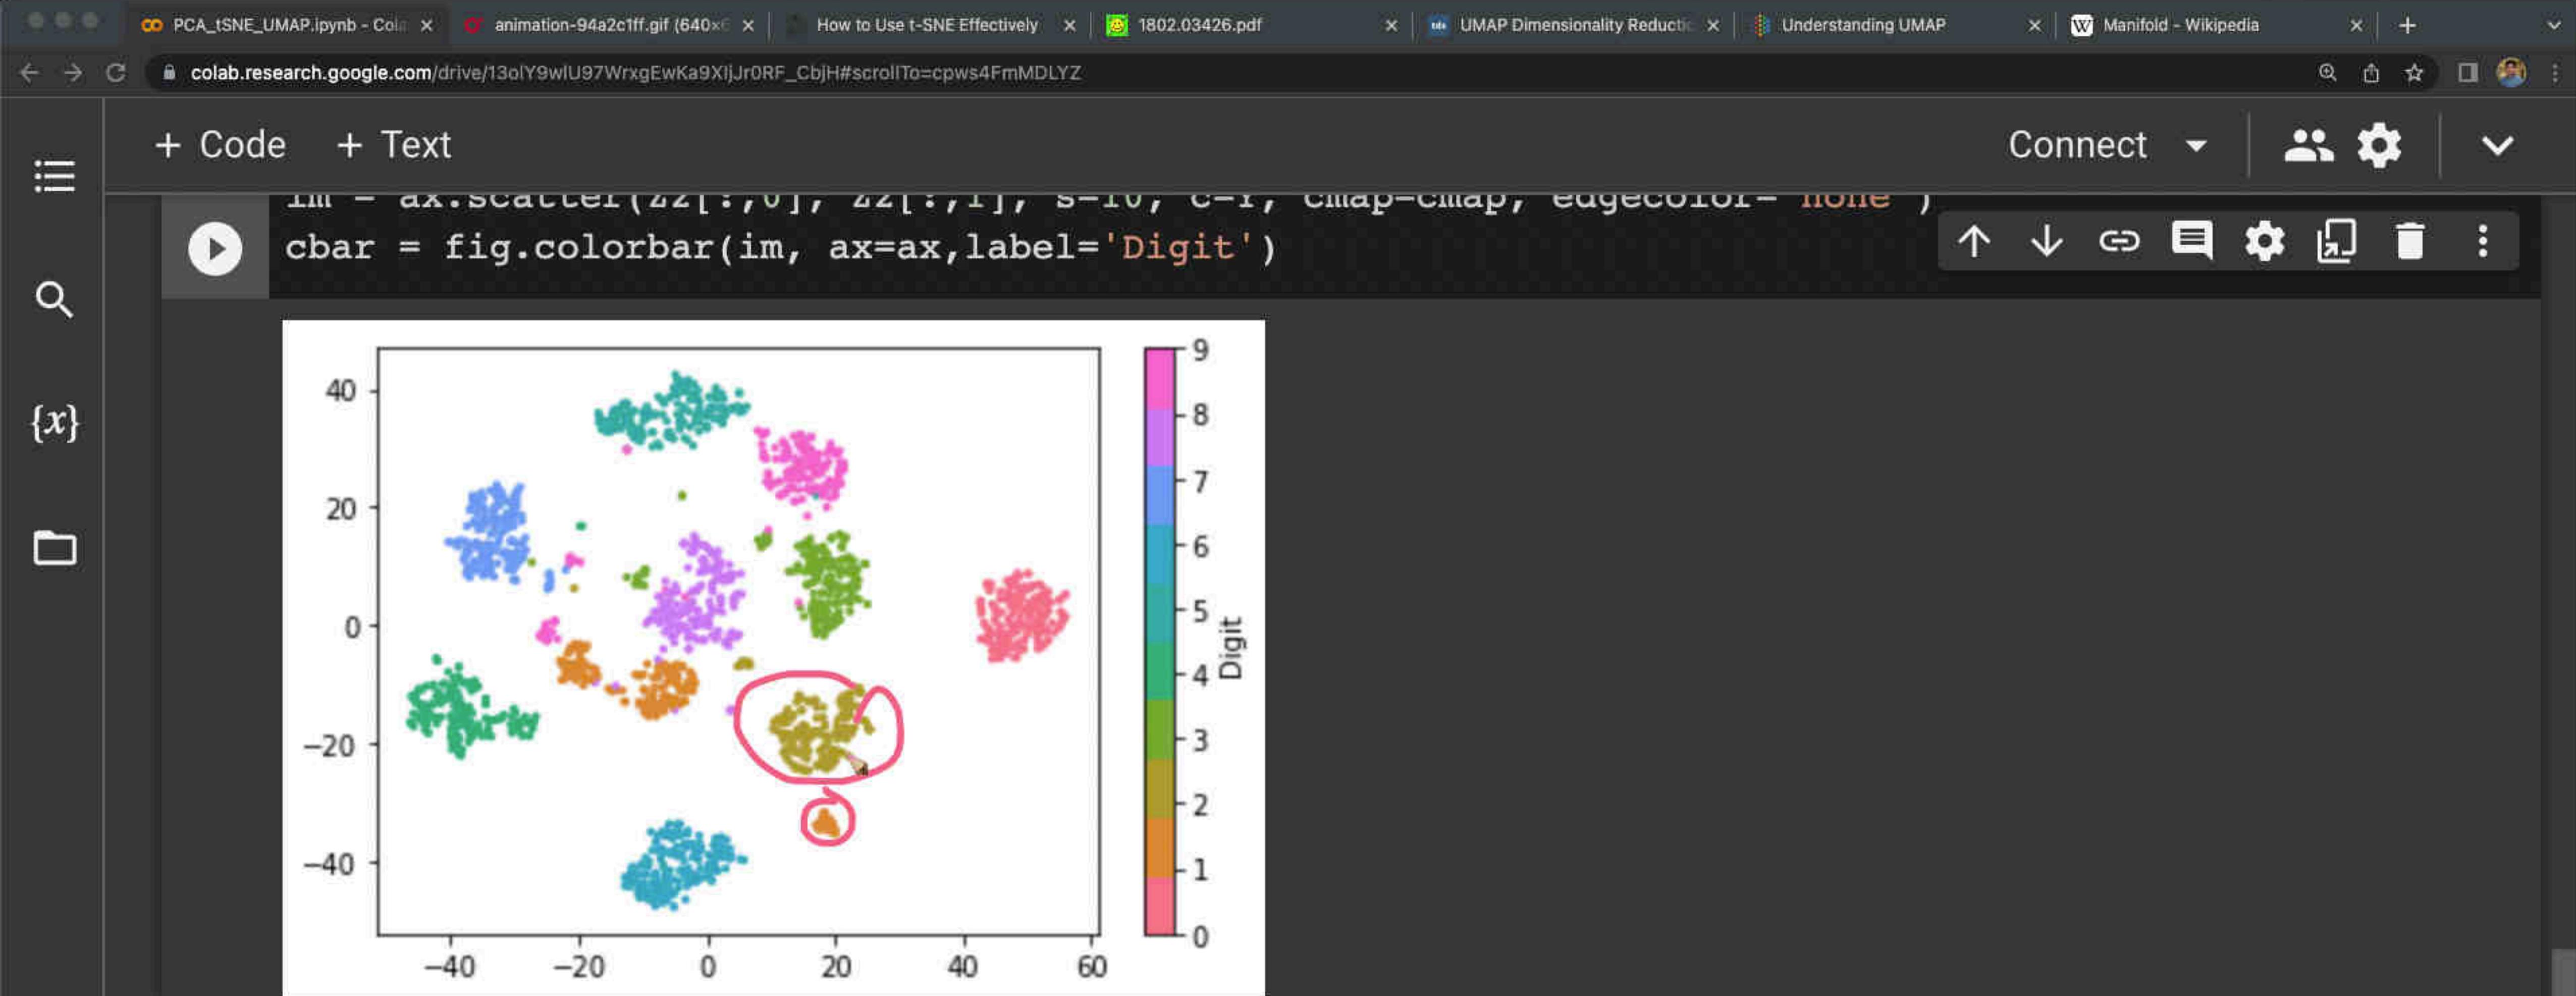
$w_l(e)$  - weights to be discovered for **low-dimensional** representation (to be found in this step [step 2])

The first term acts as an "**attractive force**" whenever there is a large weight associated to the high-dimensional case. This is because this term will be minimized when  $w_l(e)$  is as large as possible, which will occur when the distance between the points is as small as possible.

The second term acts as a "**repulsive force**" whenever high-dimensional weight  $w_h(e)$  is small. This is because the term will be minimized by making  $w_l(e)$  as small as possible.

On balance the push and pull of the two terms will allow the low-dimensional representation settle into a state that matches the original data.





[ ] %time

```
Z3 = UMAP(init='random', random_state=42).fit_transform(X)
```

CPU times: user 8.39 s, sys: 30 ms, total: 8.42 s

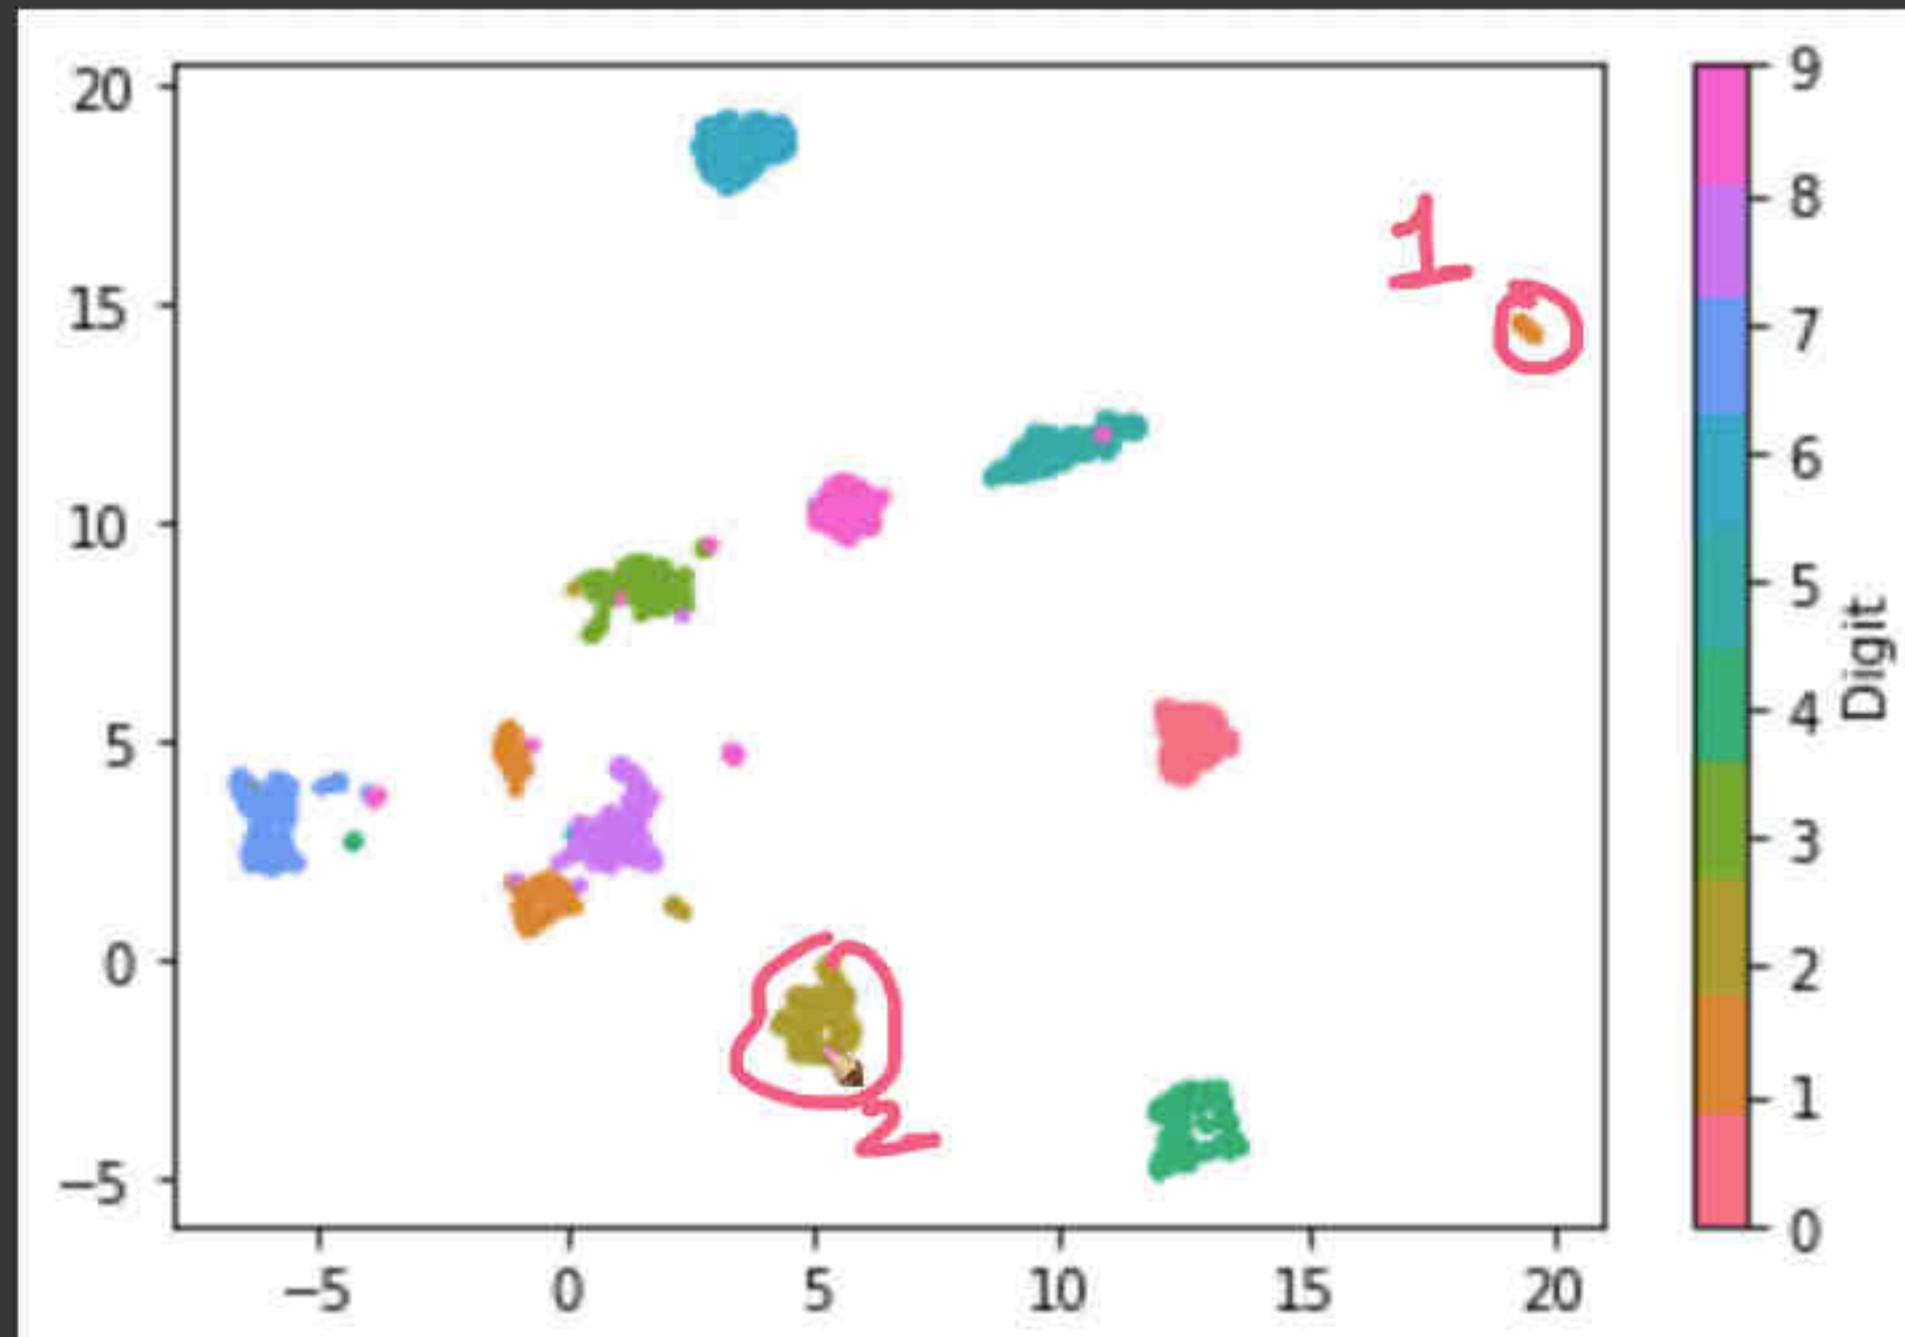
Wall time: 8.27 s

+ Code + Text

Connect



```
[ ] cbar = fig.colorbar(im, ax=ax, label='Digit')
```



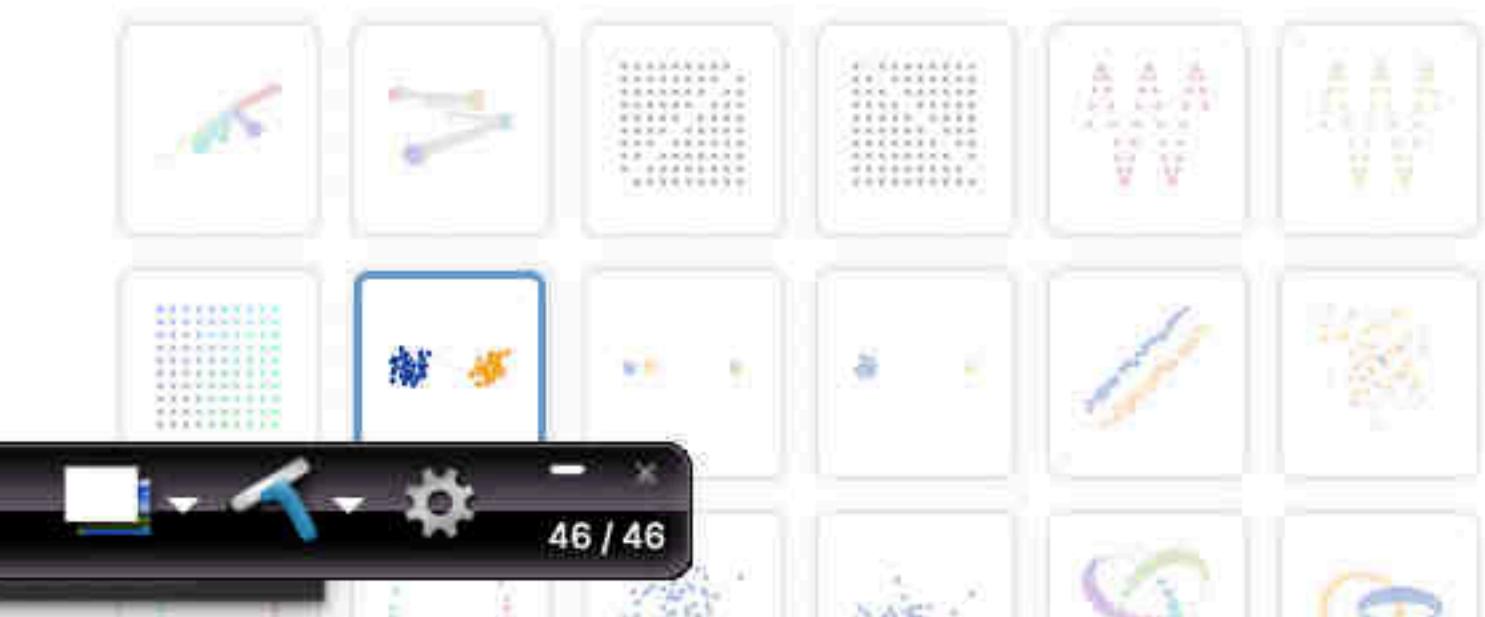
# Understanding UMAP

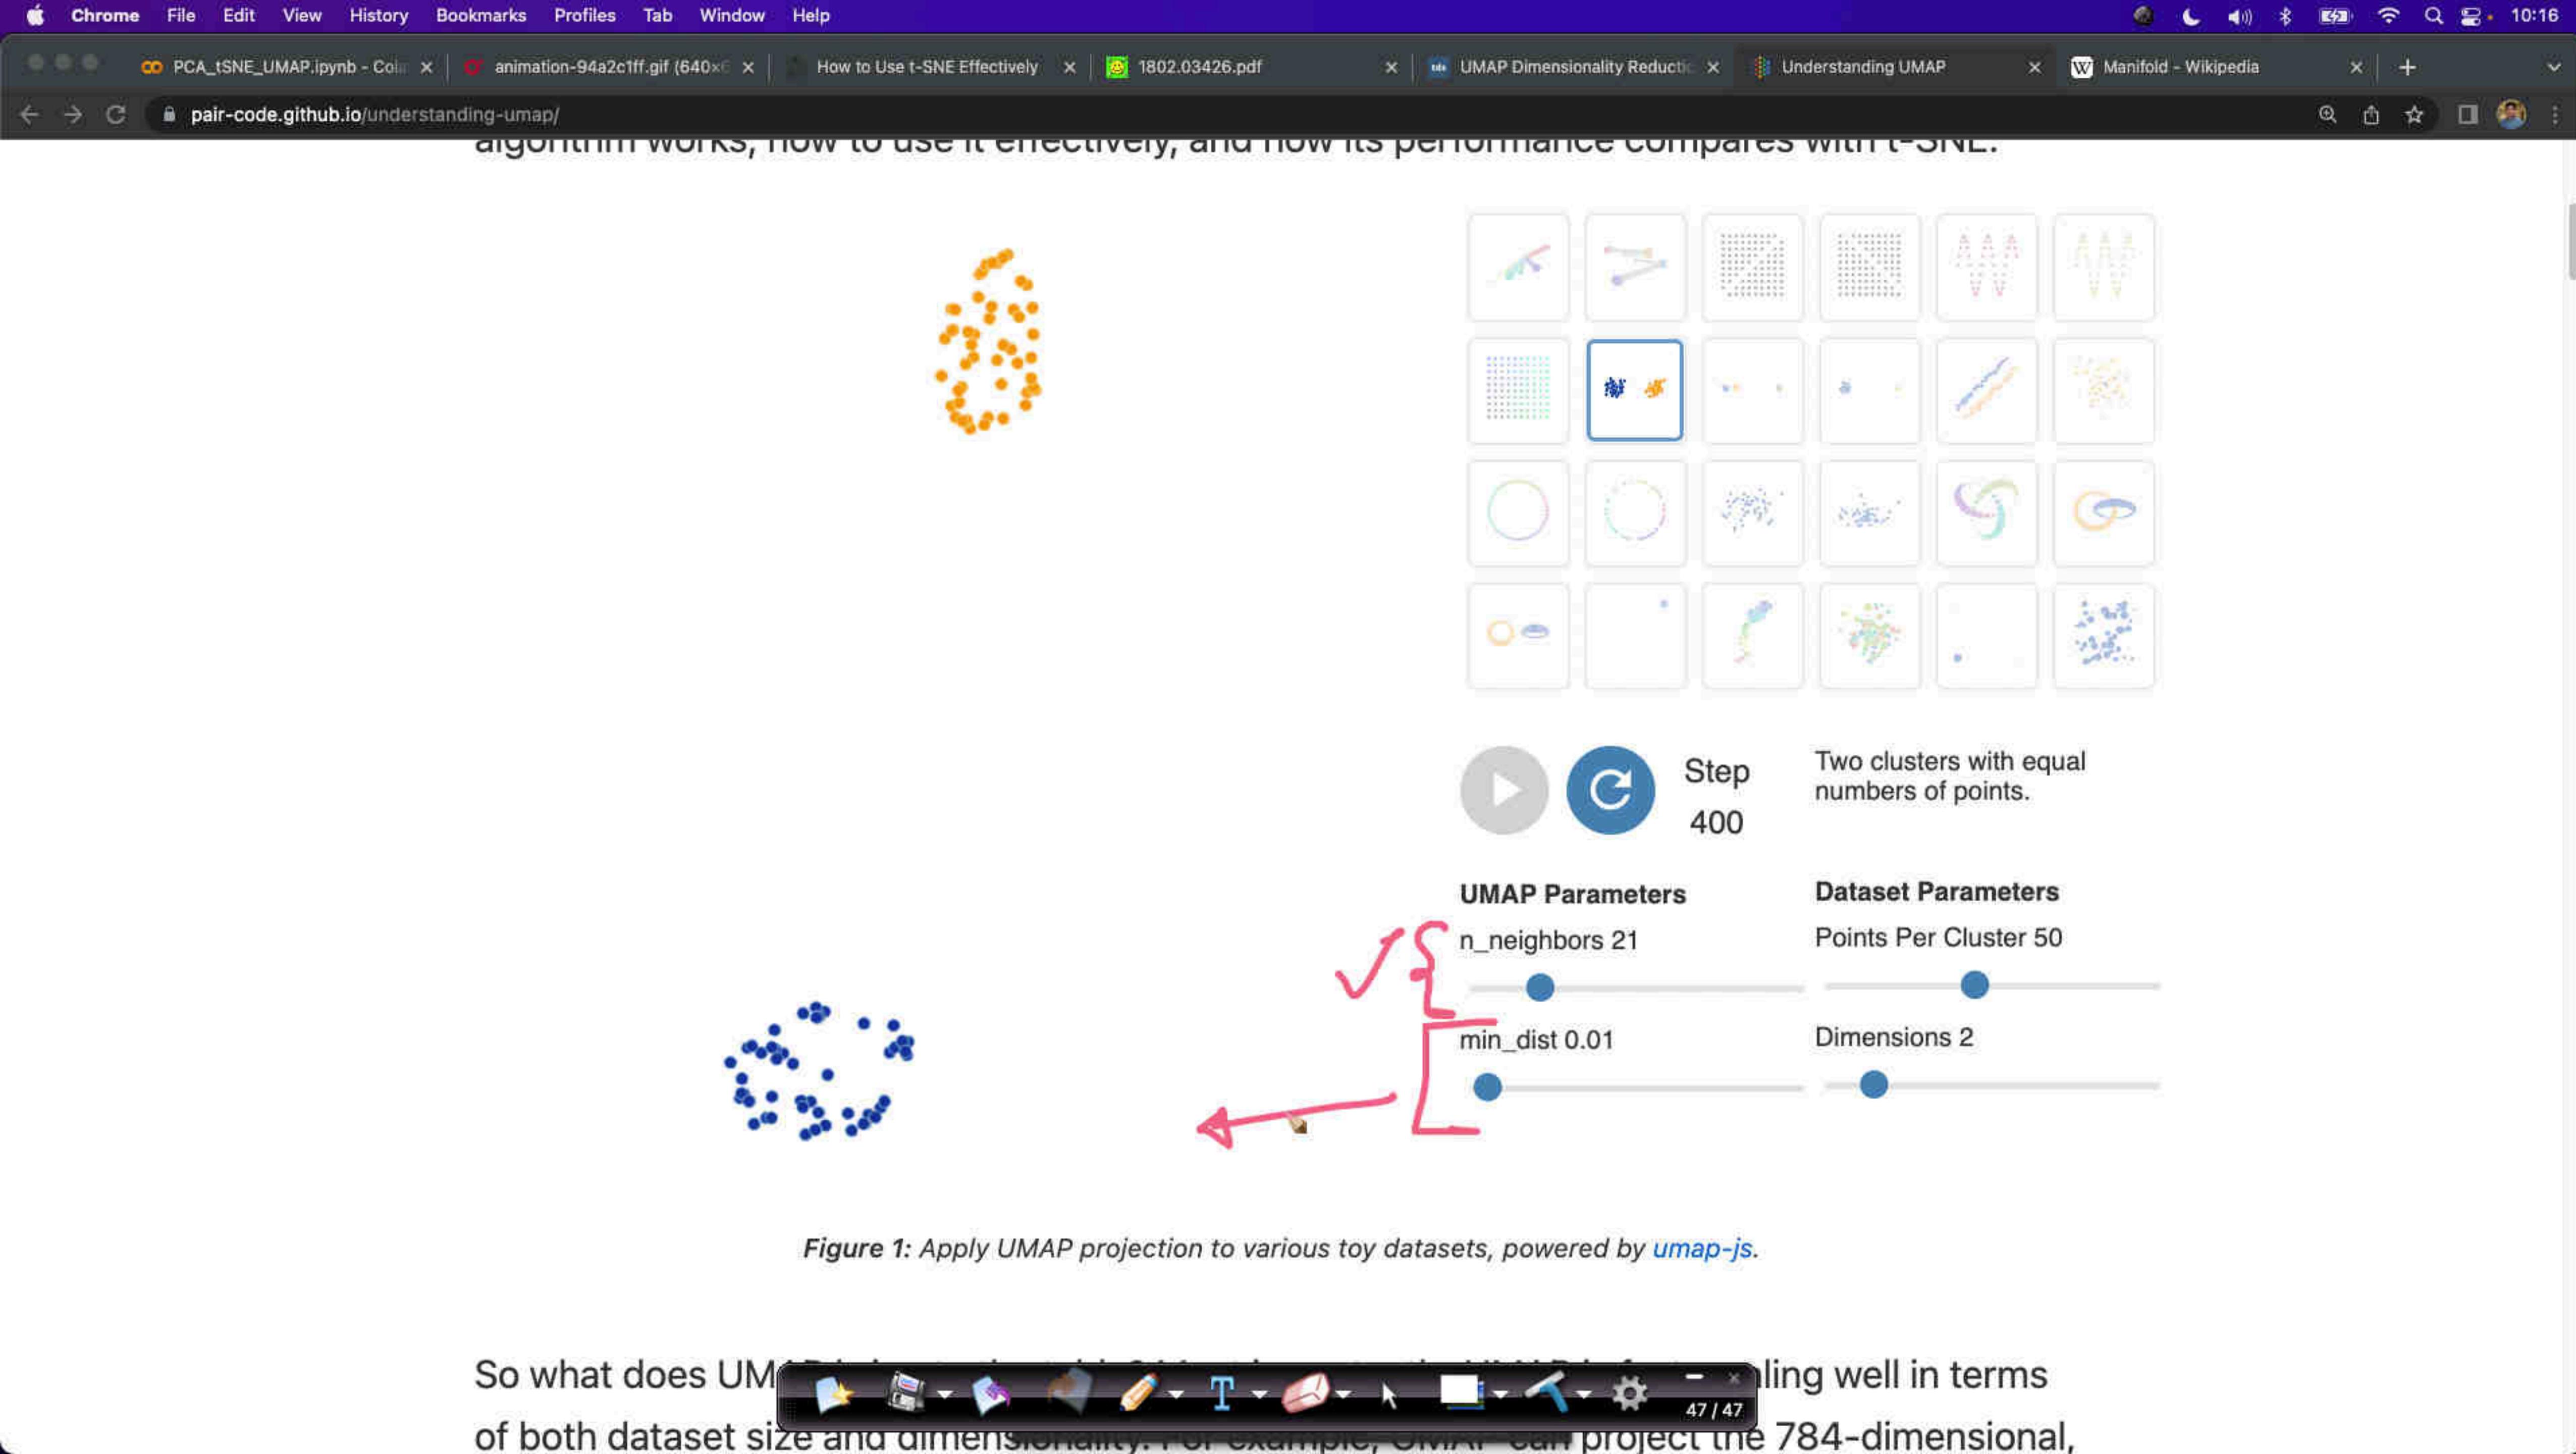
Andy Coenen, Adam Pearce | [Google PAIR](#)



Dimensionality reduction is a powerful tool for machine learning practitioners to visualize and understand large, high dimensional datasets. One of the most widely used techniques for visualization is [t-SNE](#), but its performance suffers with large datasets and using it correctly can be [challenging](#).

[UMAP](#) is a new technique by McInnes et al. that offers a number of advantages over t-SNE, most notably increased speed and better preservation of the data's global structure. In this article, we'll take a look at the theory behind UMAP in order to better understand how the algorithm works, how to use it effectively, and how its performance compares with t-SNE.





So what does UMAP do? It's a dimensionality reduction algorithm that's scaling well in terms of both dataset size and dimensionality. For example, UMAP can project the 784-dimensional,

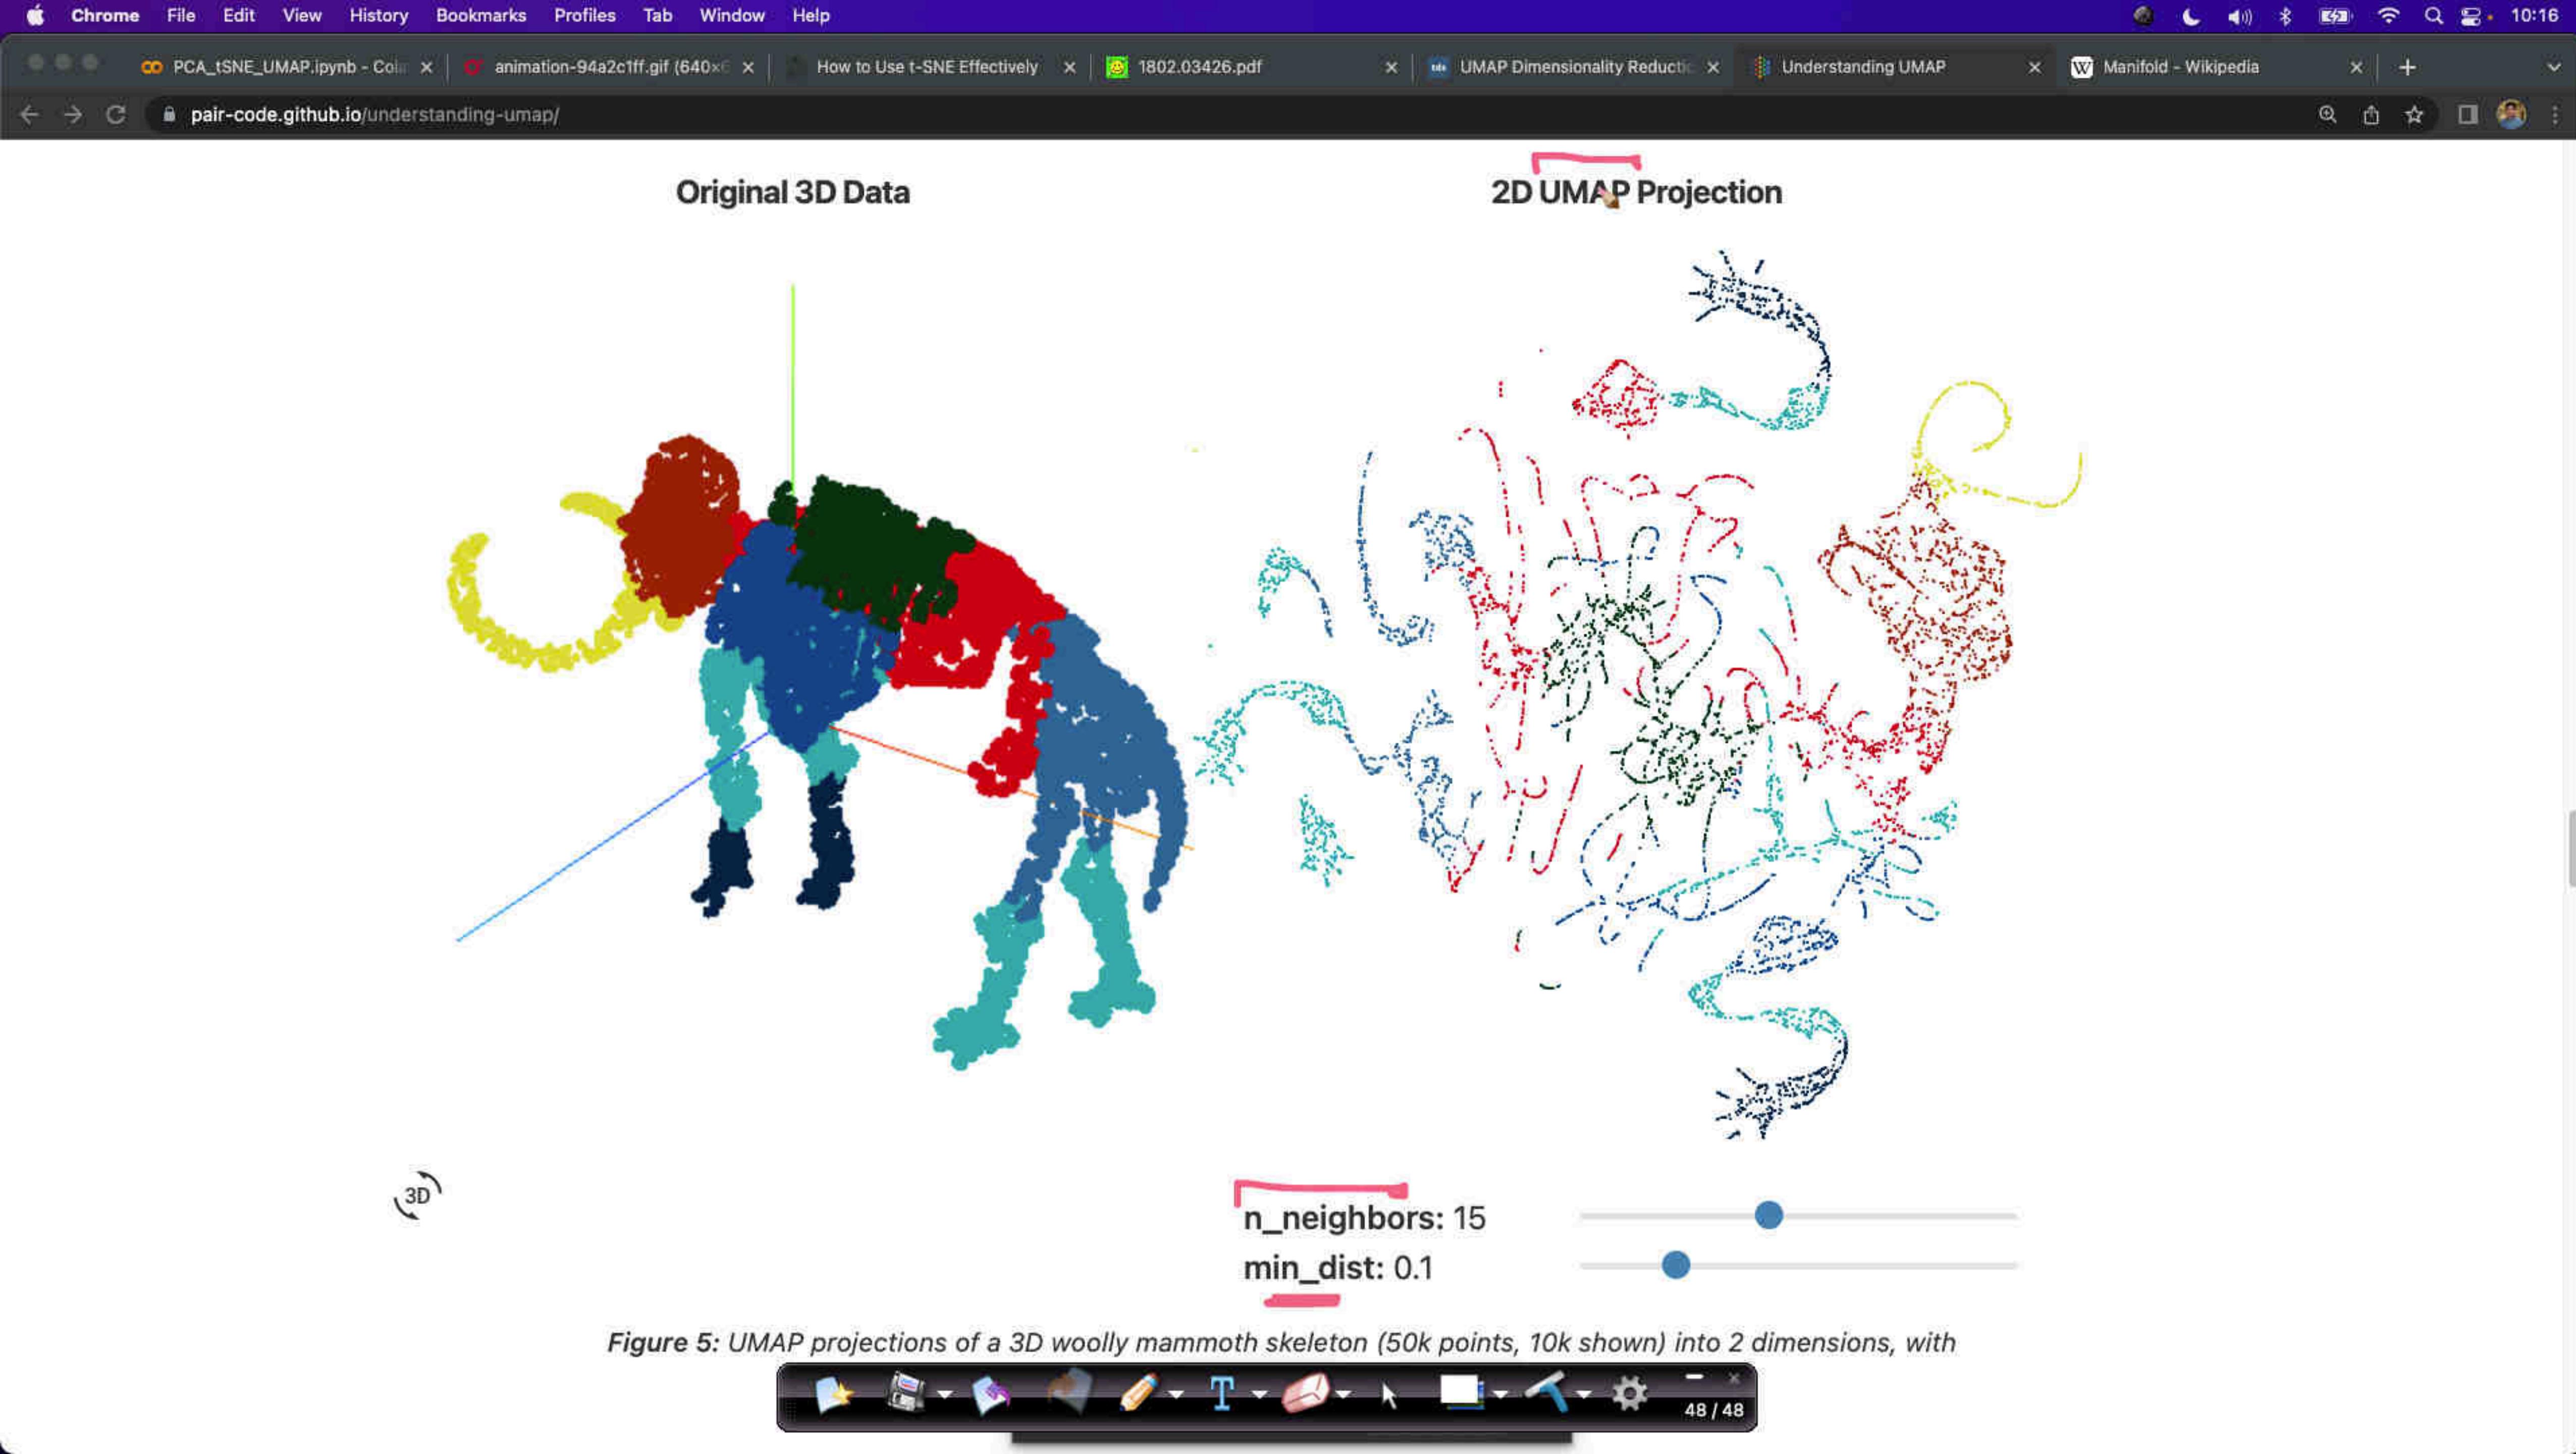


Figure 5: UMAP projections of a 3D woolly mammoth skeleton (50k points, 10k shown) into 2 dimensions, with

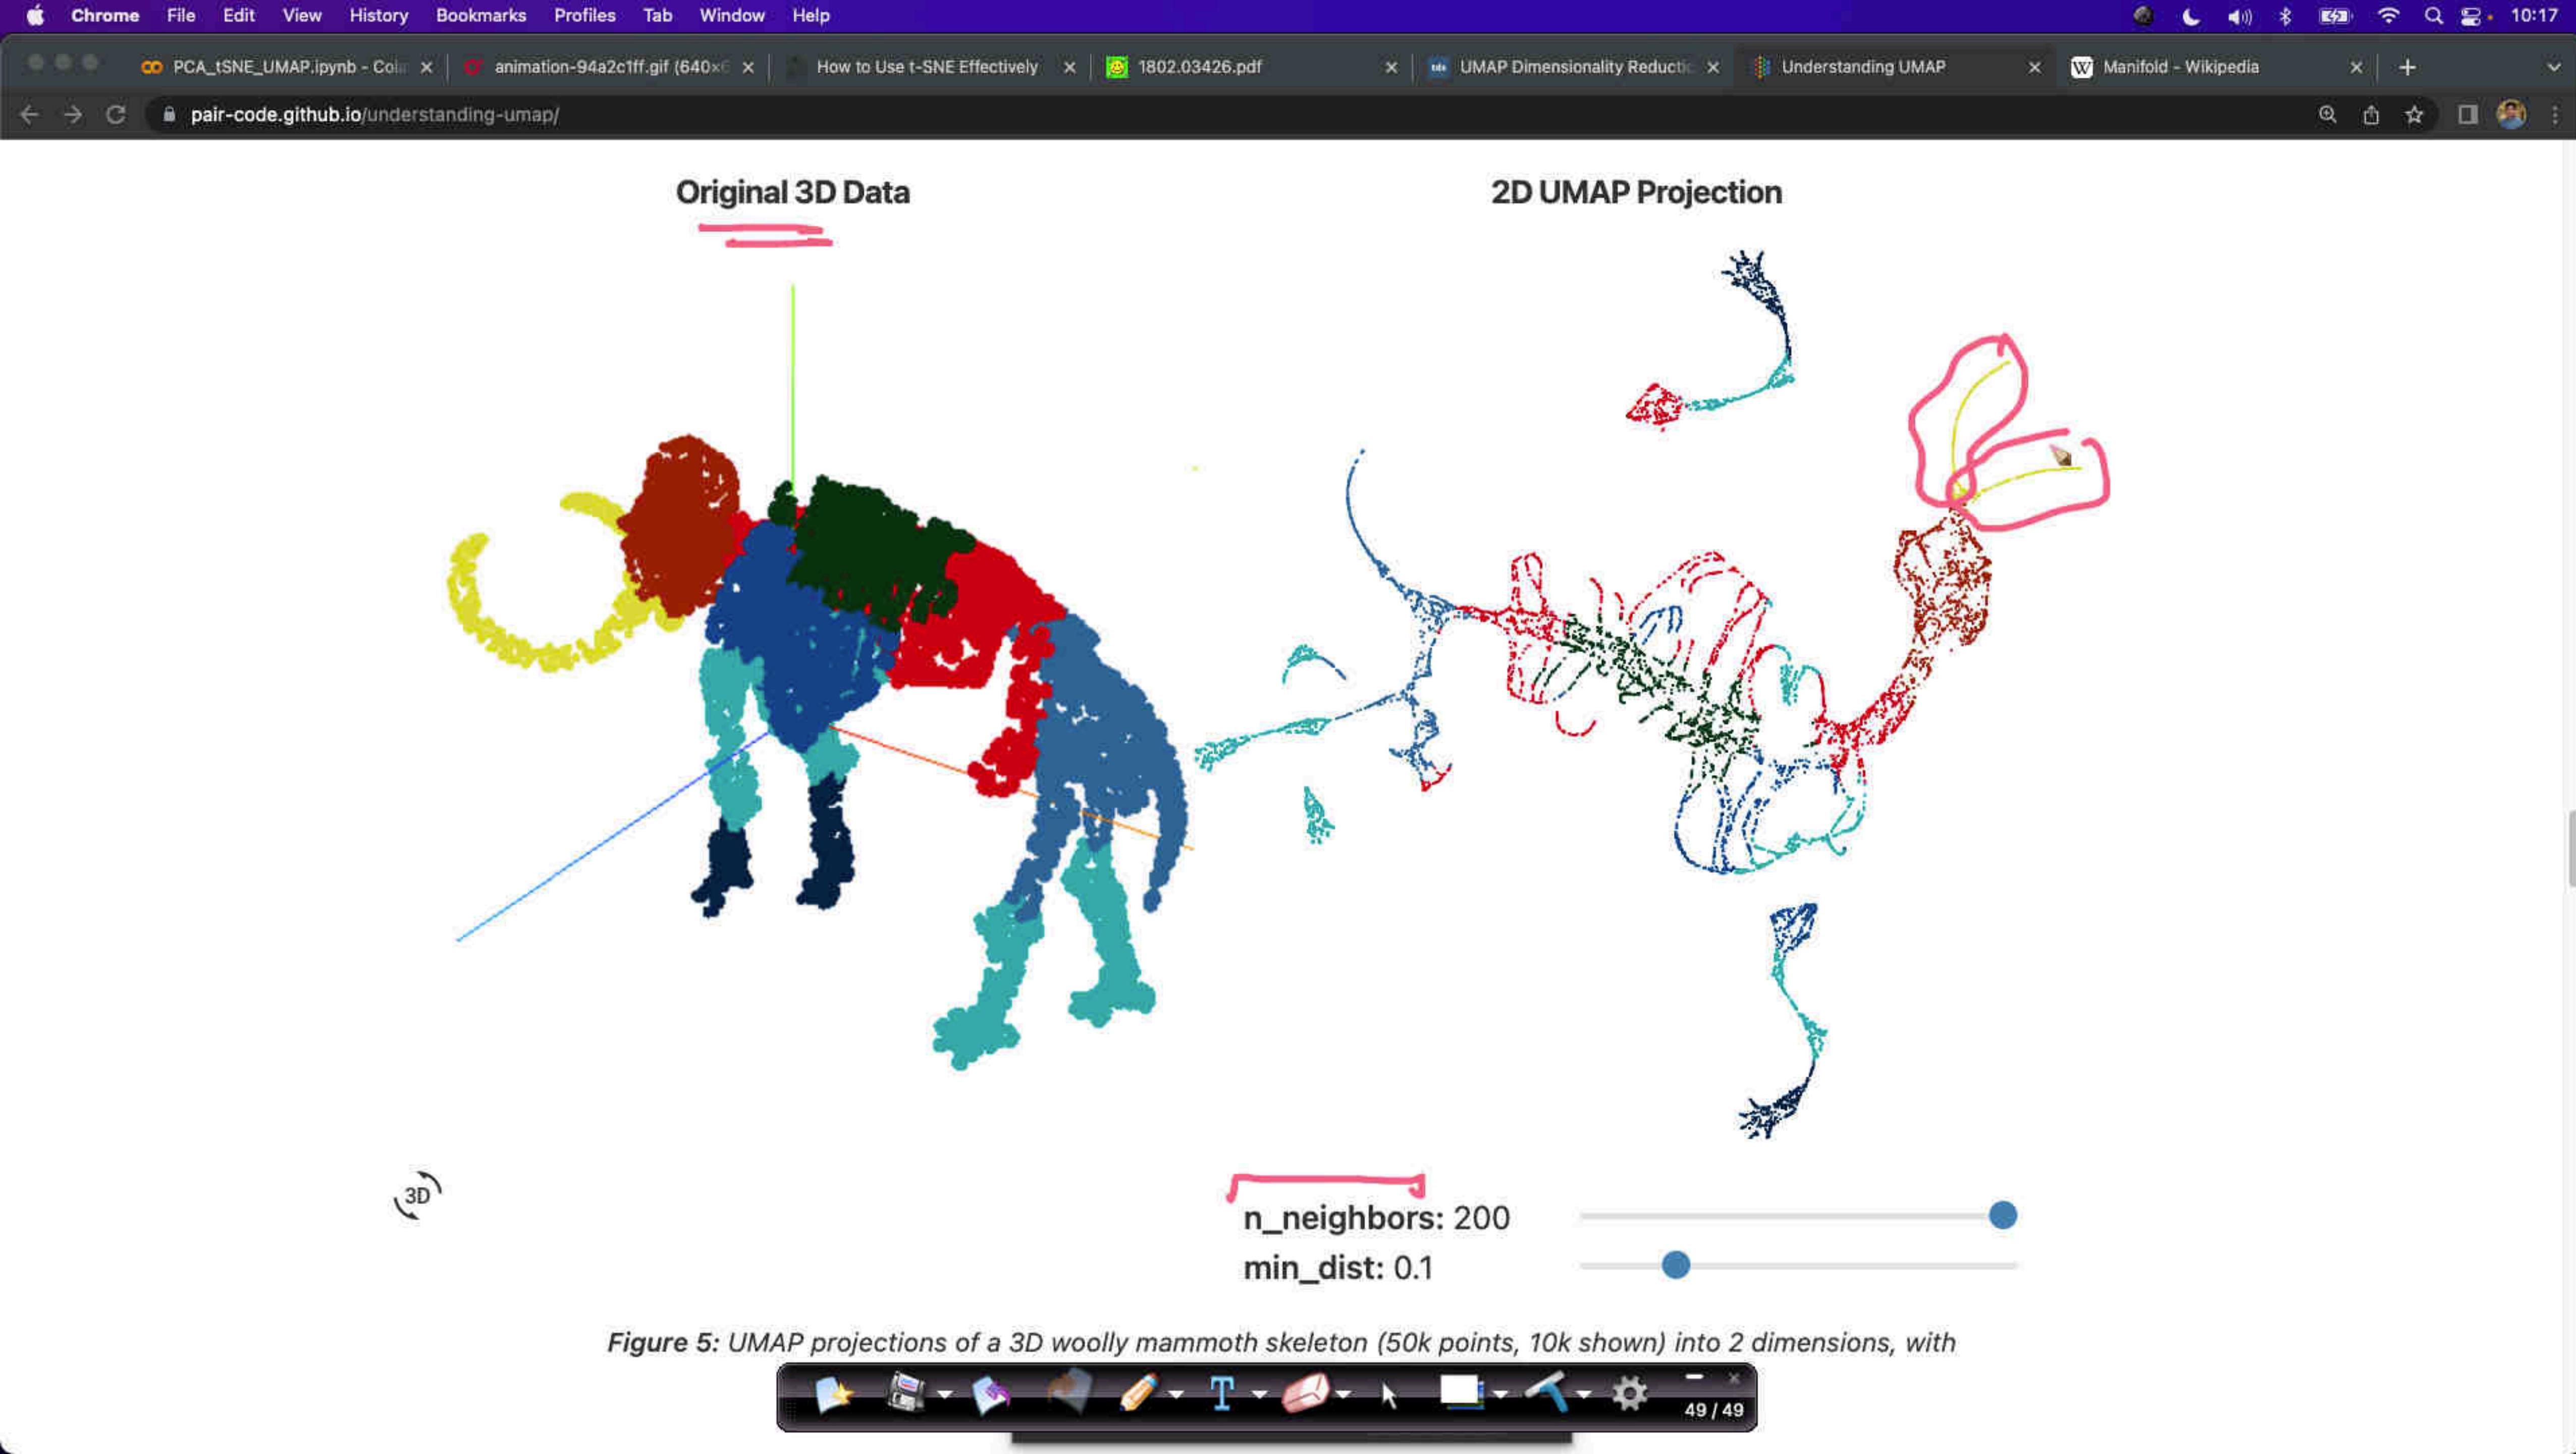


Figure 5: UMAP projections of a 3D woolly mammoth skeleton (50k points, 10k shown) into 2 dimensions, with

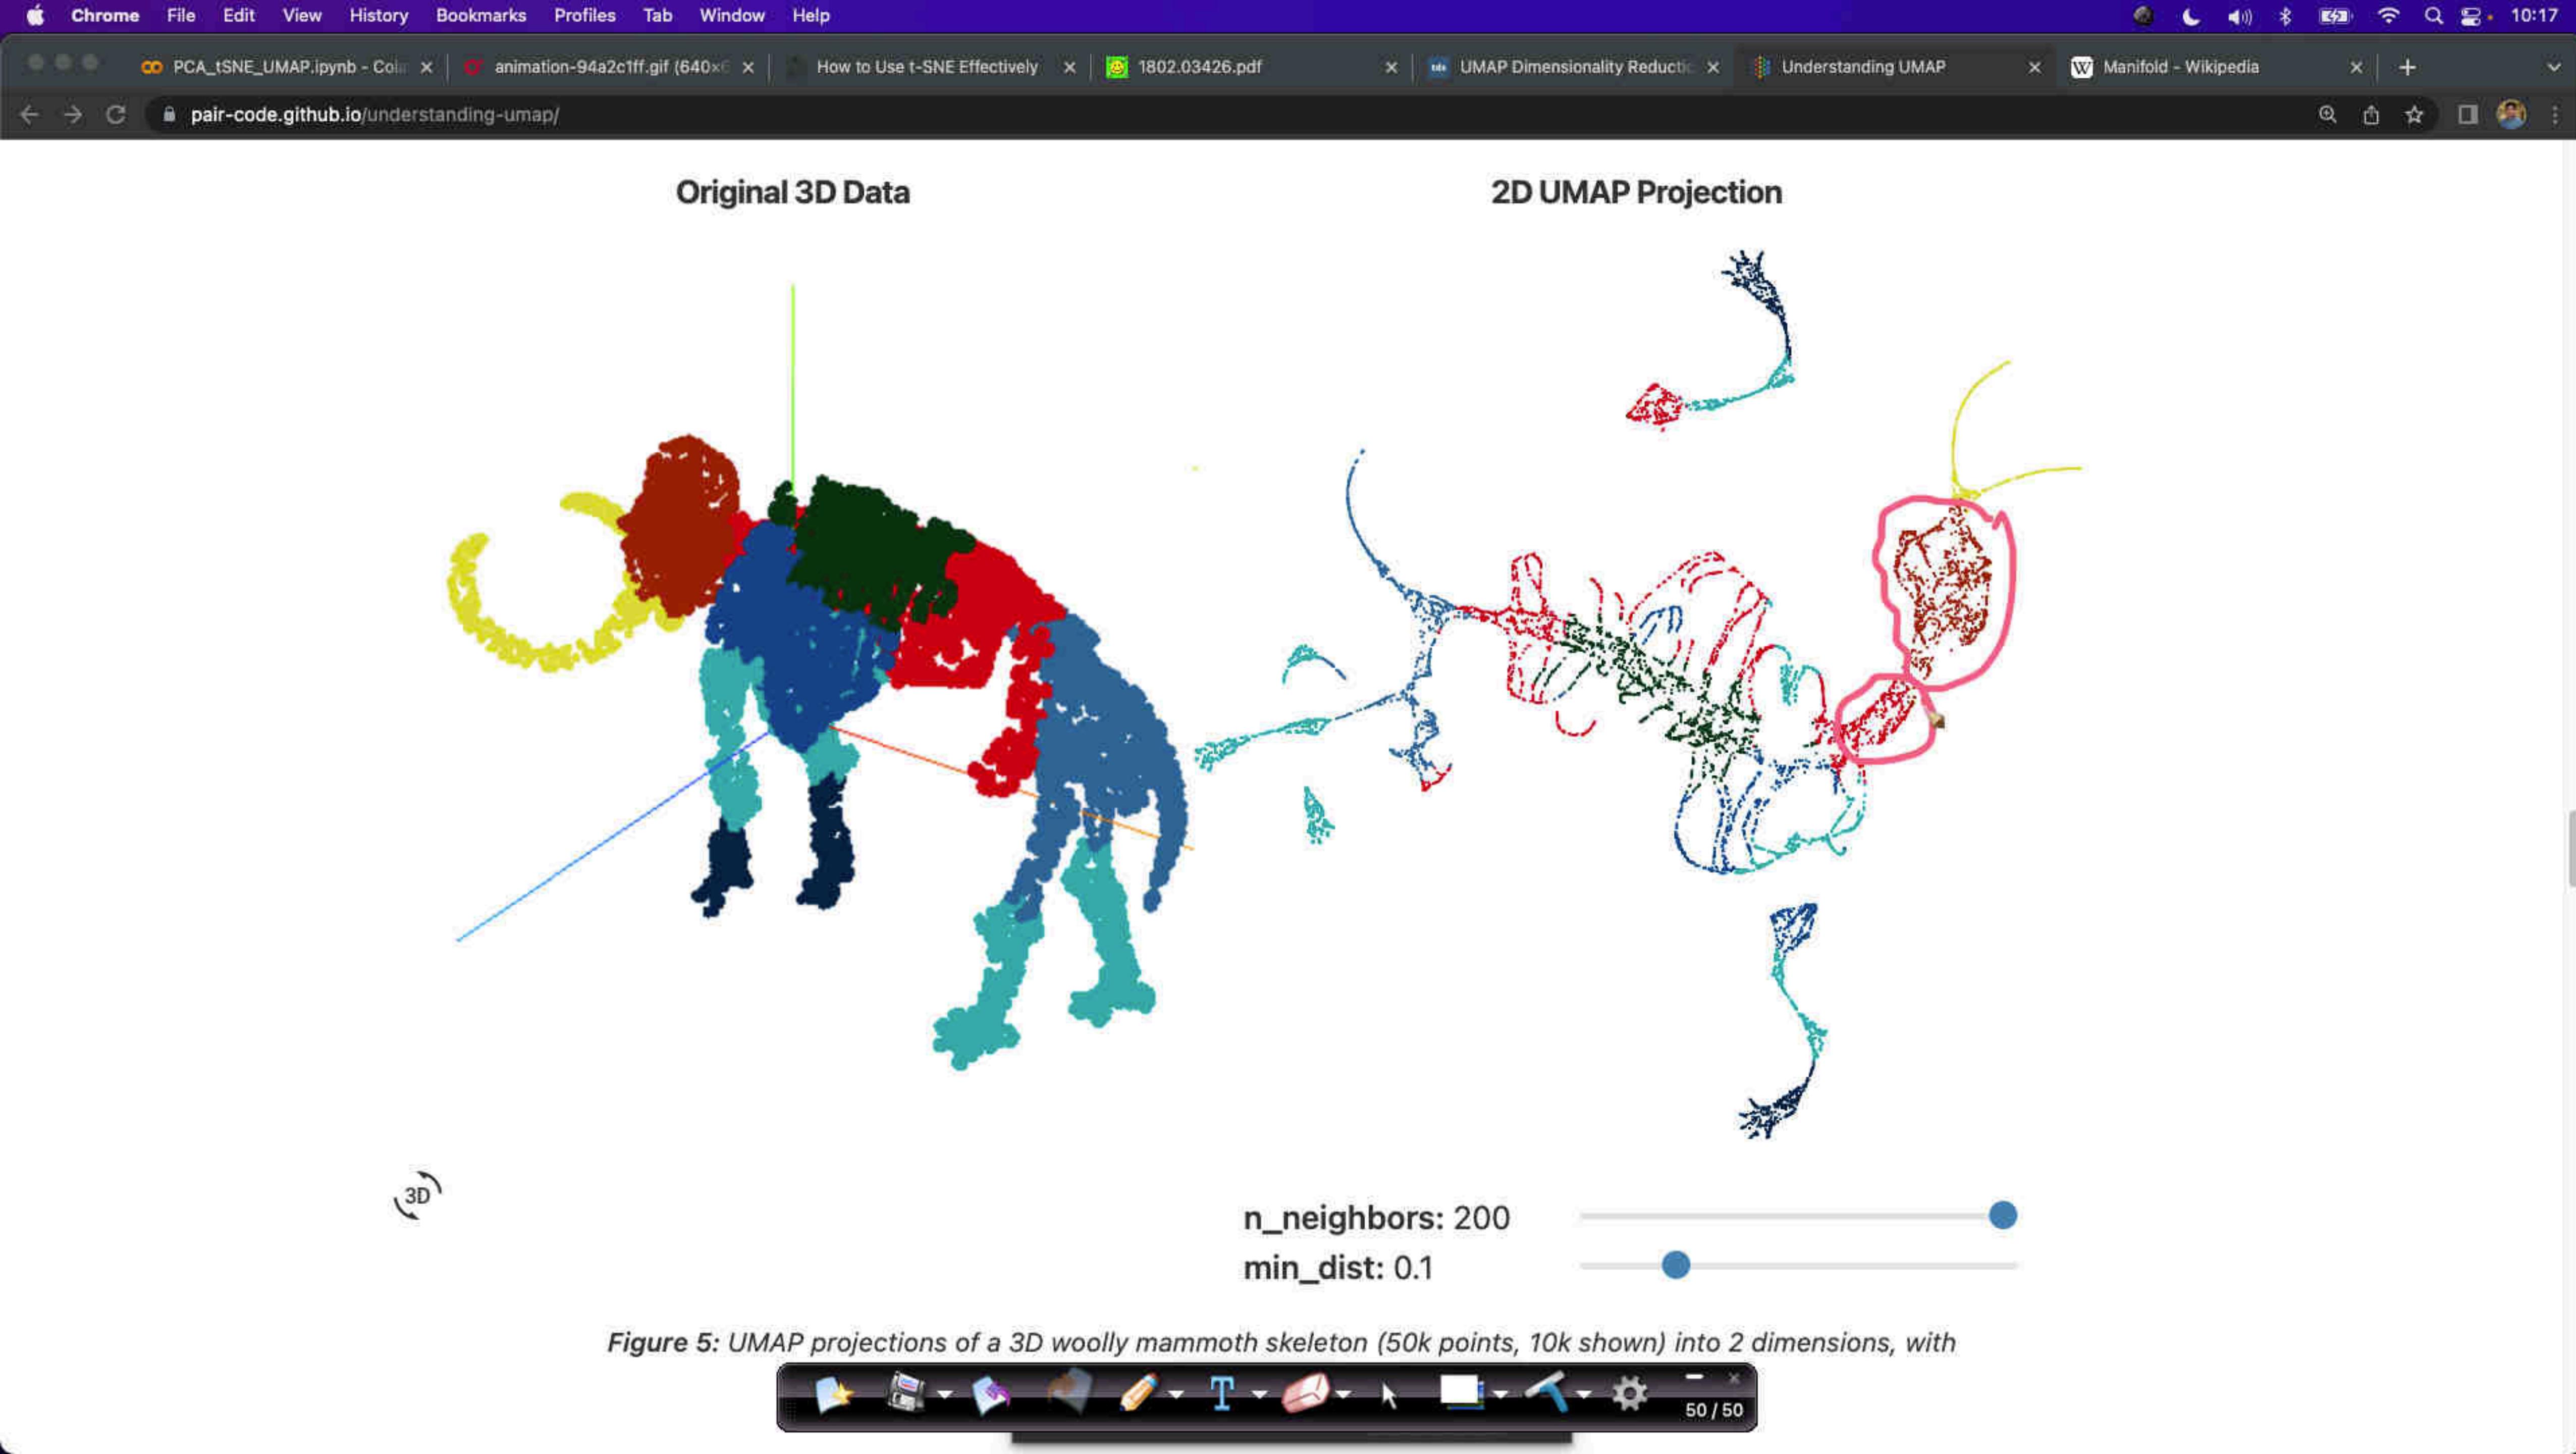


Figure 5: UMAP projections of a 3D woolly mammoth skeleton (50k points, 10k shown) into 2 dimensions, with

would distract from our intent.

### 4.3 Hyper-parameters

As described in Algorithm 1, the UMAP algorithm takes four hyper-parameters:

- { 1.  $n$ , the number of neighbors to consider when approximating the local metric;
- 2.  $d$ , the target embedding dimension;
- 3. min-dist, the desired separation between close points in the embedding space; and

100 → ② Uiz ✓

22

arxiv.org/pdf/1802.03426.pdf

22 / 63 | - 200% + | ☰

1802.03426.pdf

51 / 51

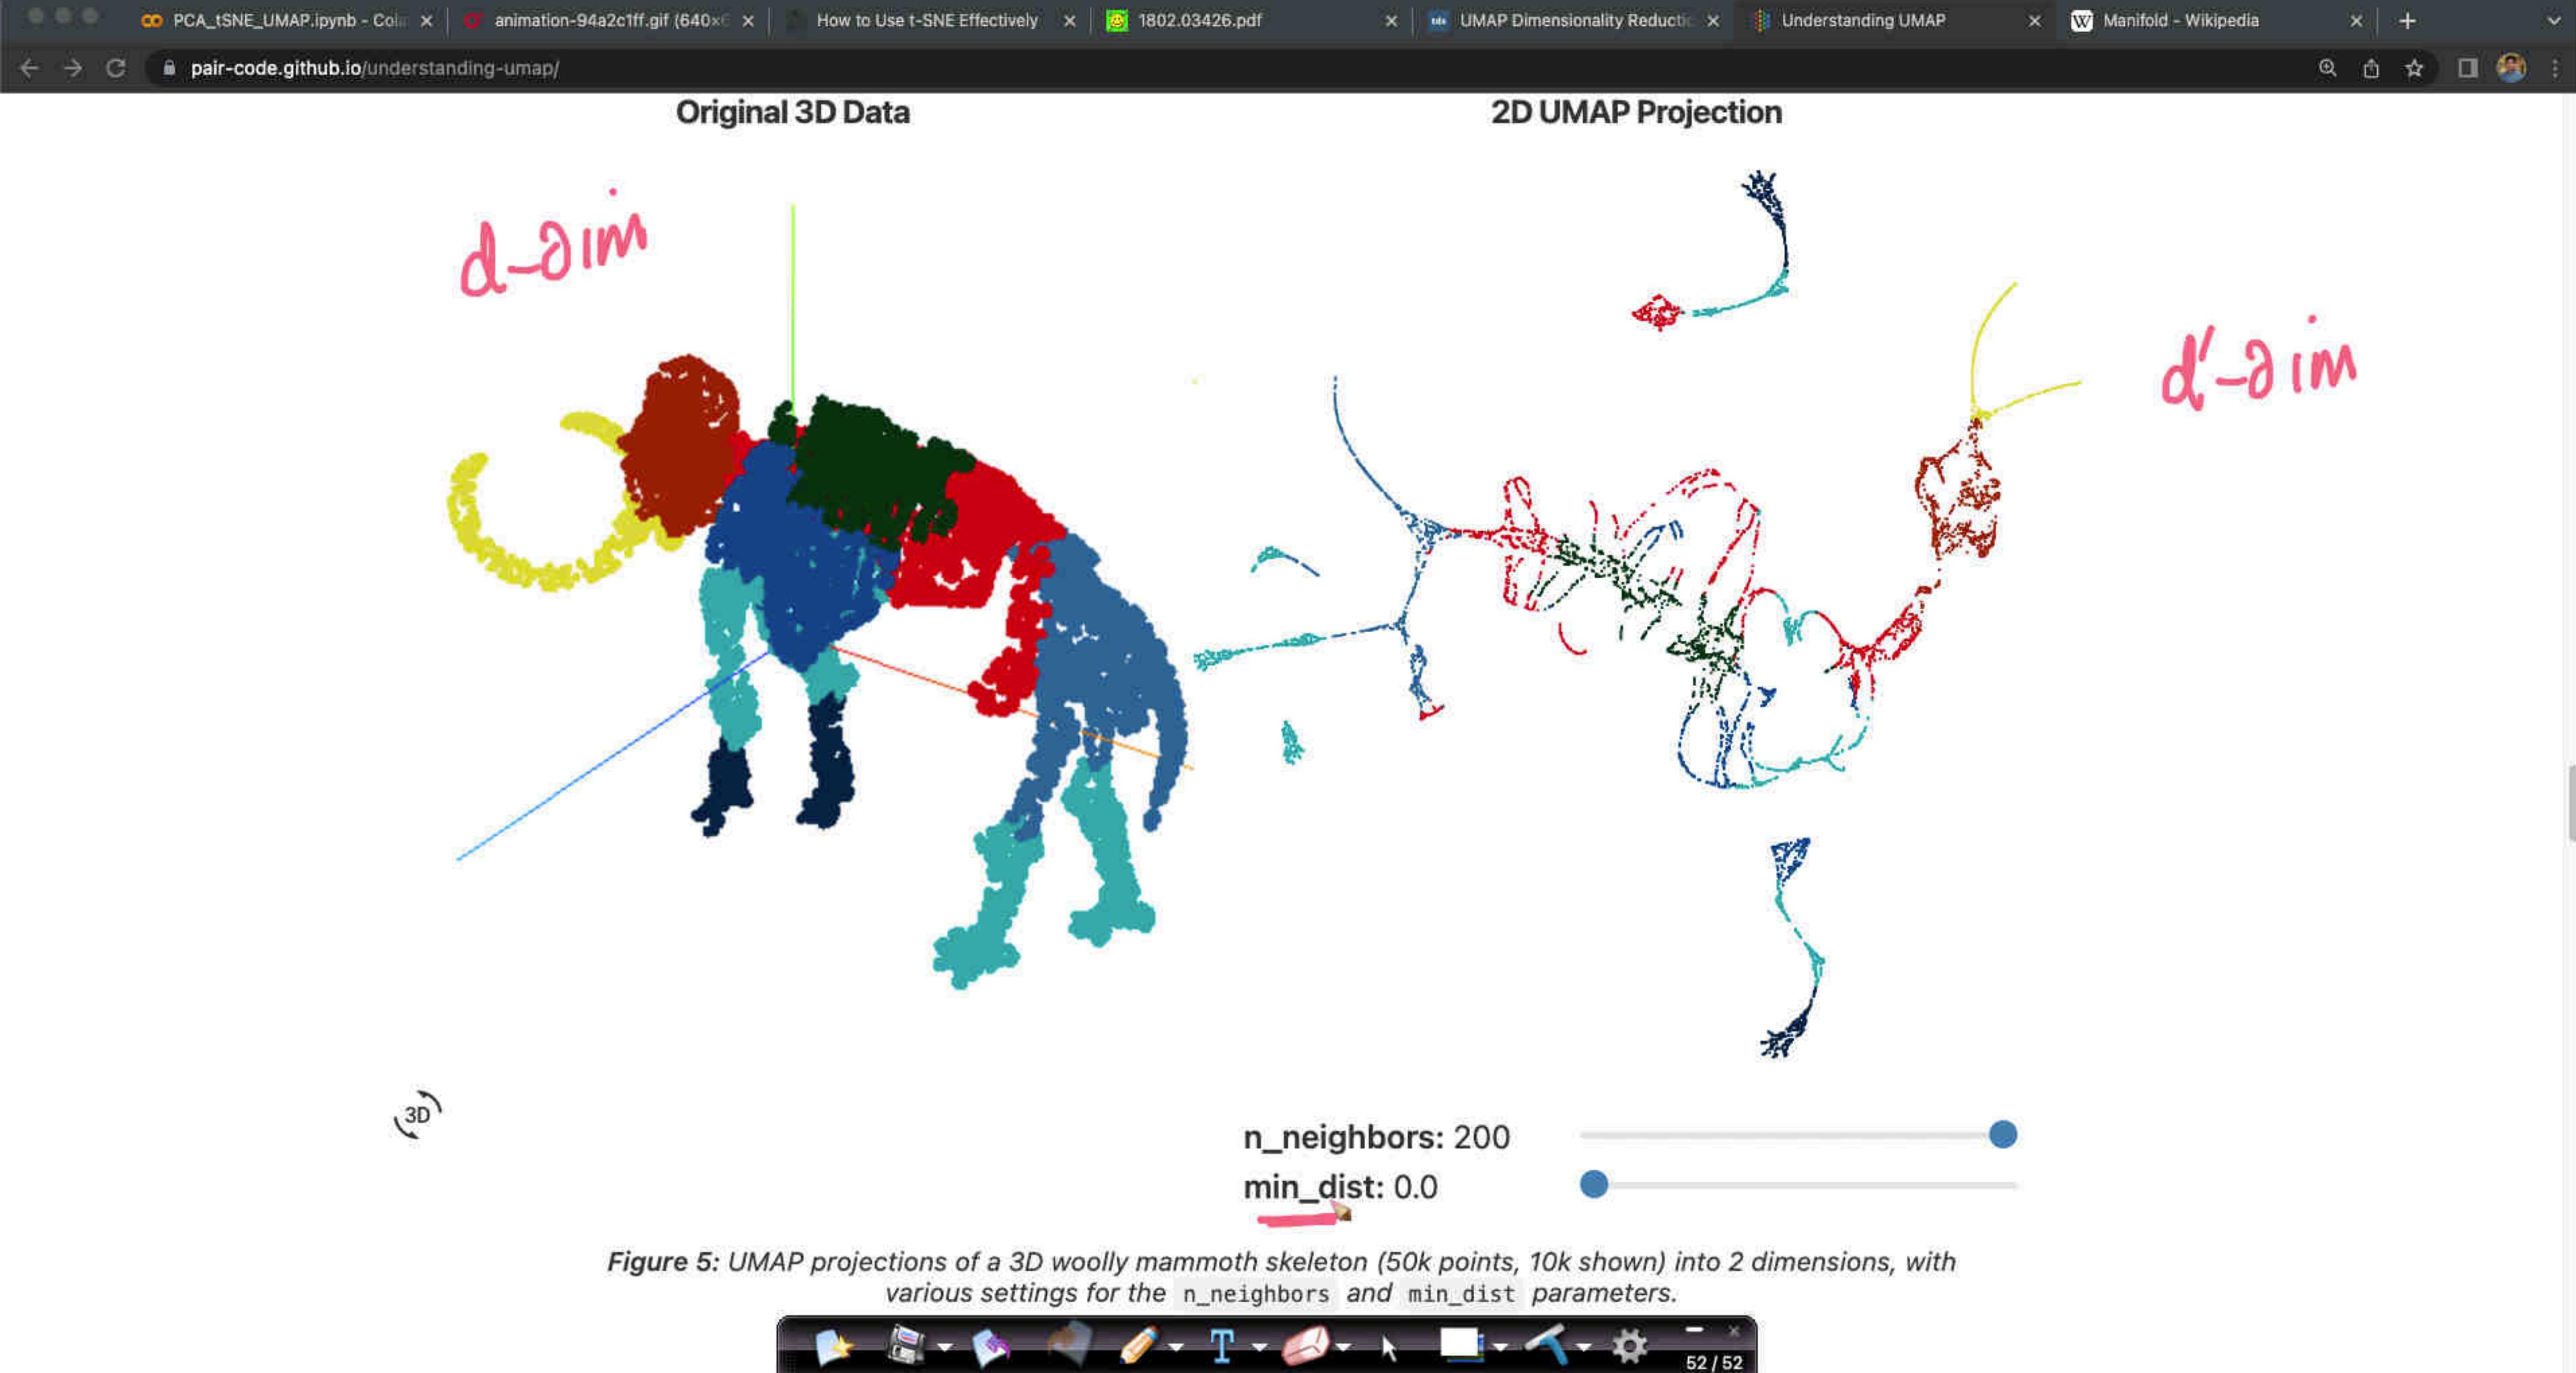


Figure 5: UMAP projections of a 3D woolly mammoth skeleton (50k points, 10k shown) into 2 dimensions, with various settings for the `n_neighbors` and `min_dist` parameters.

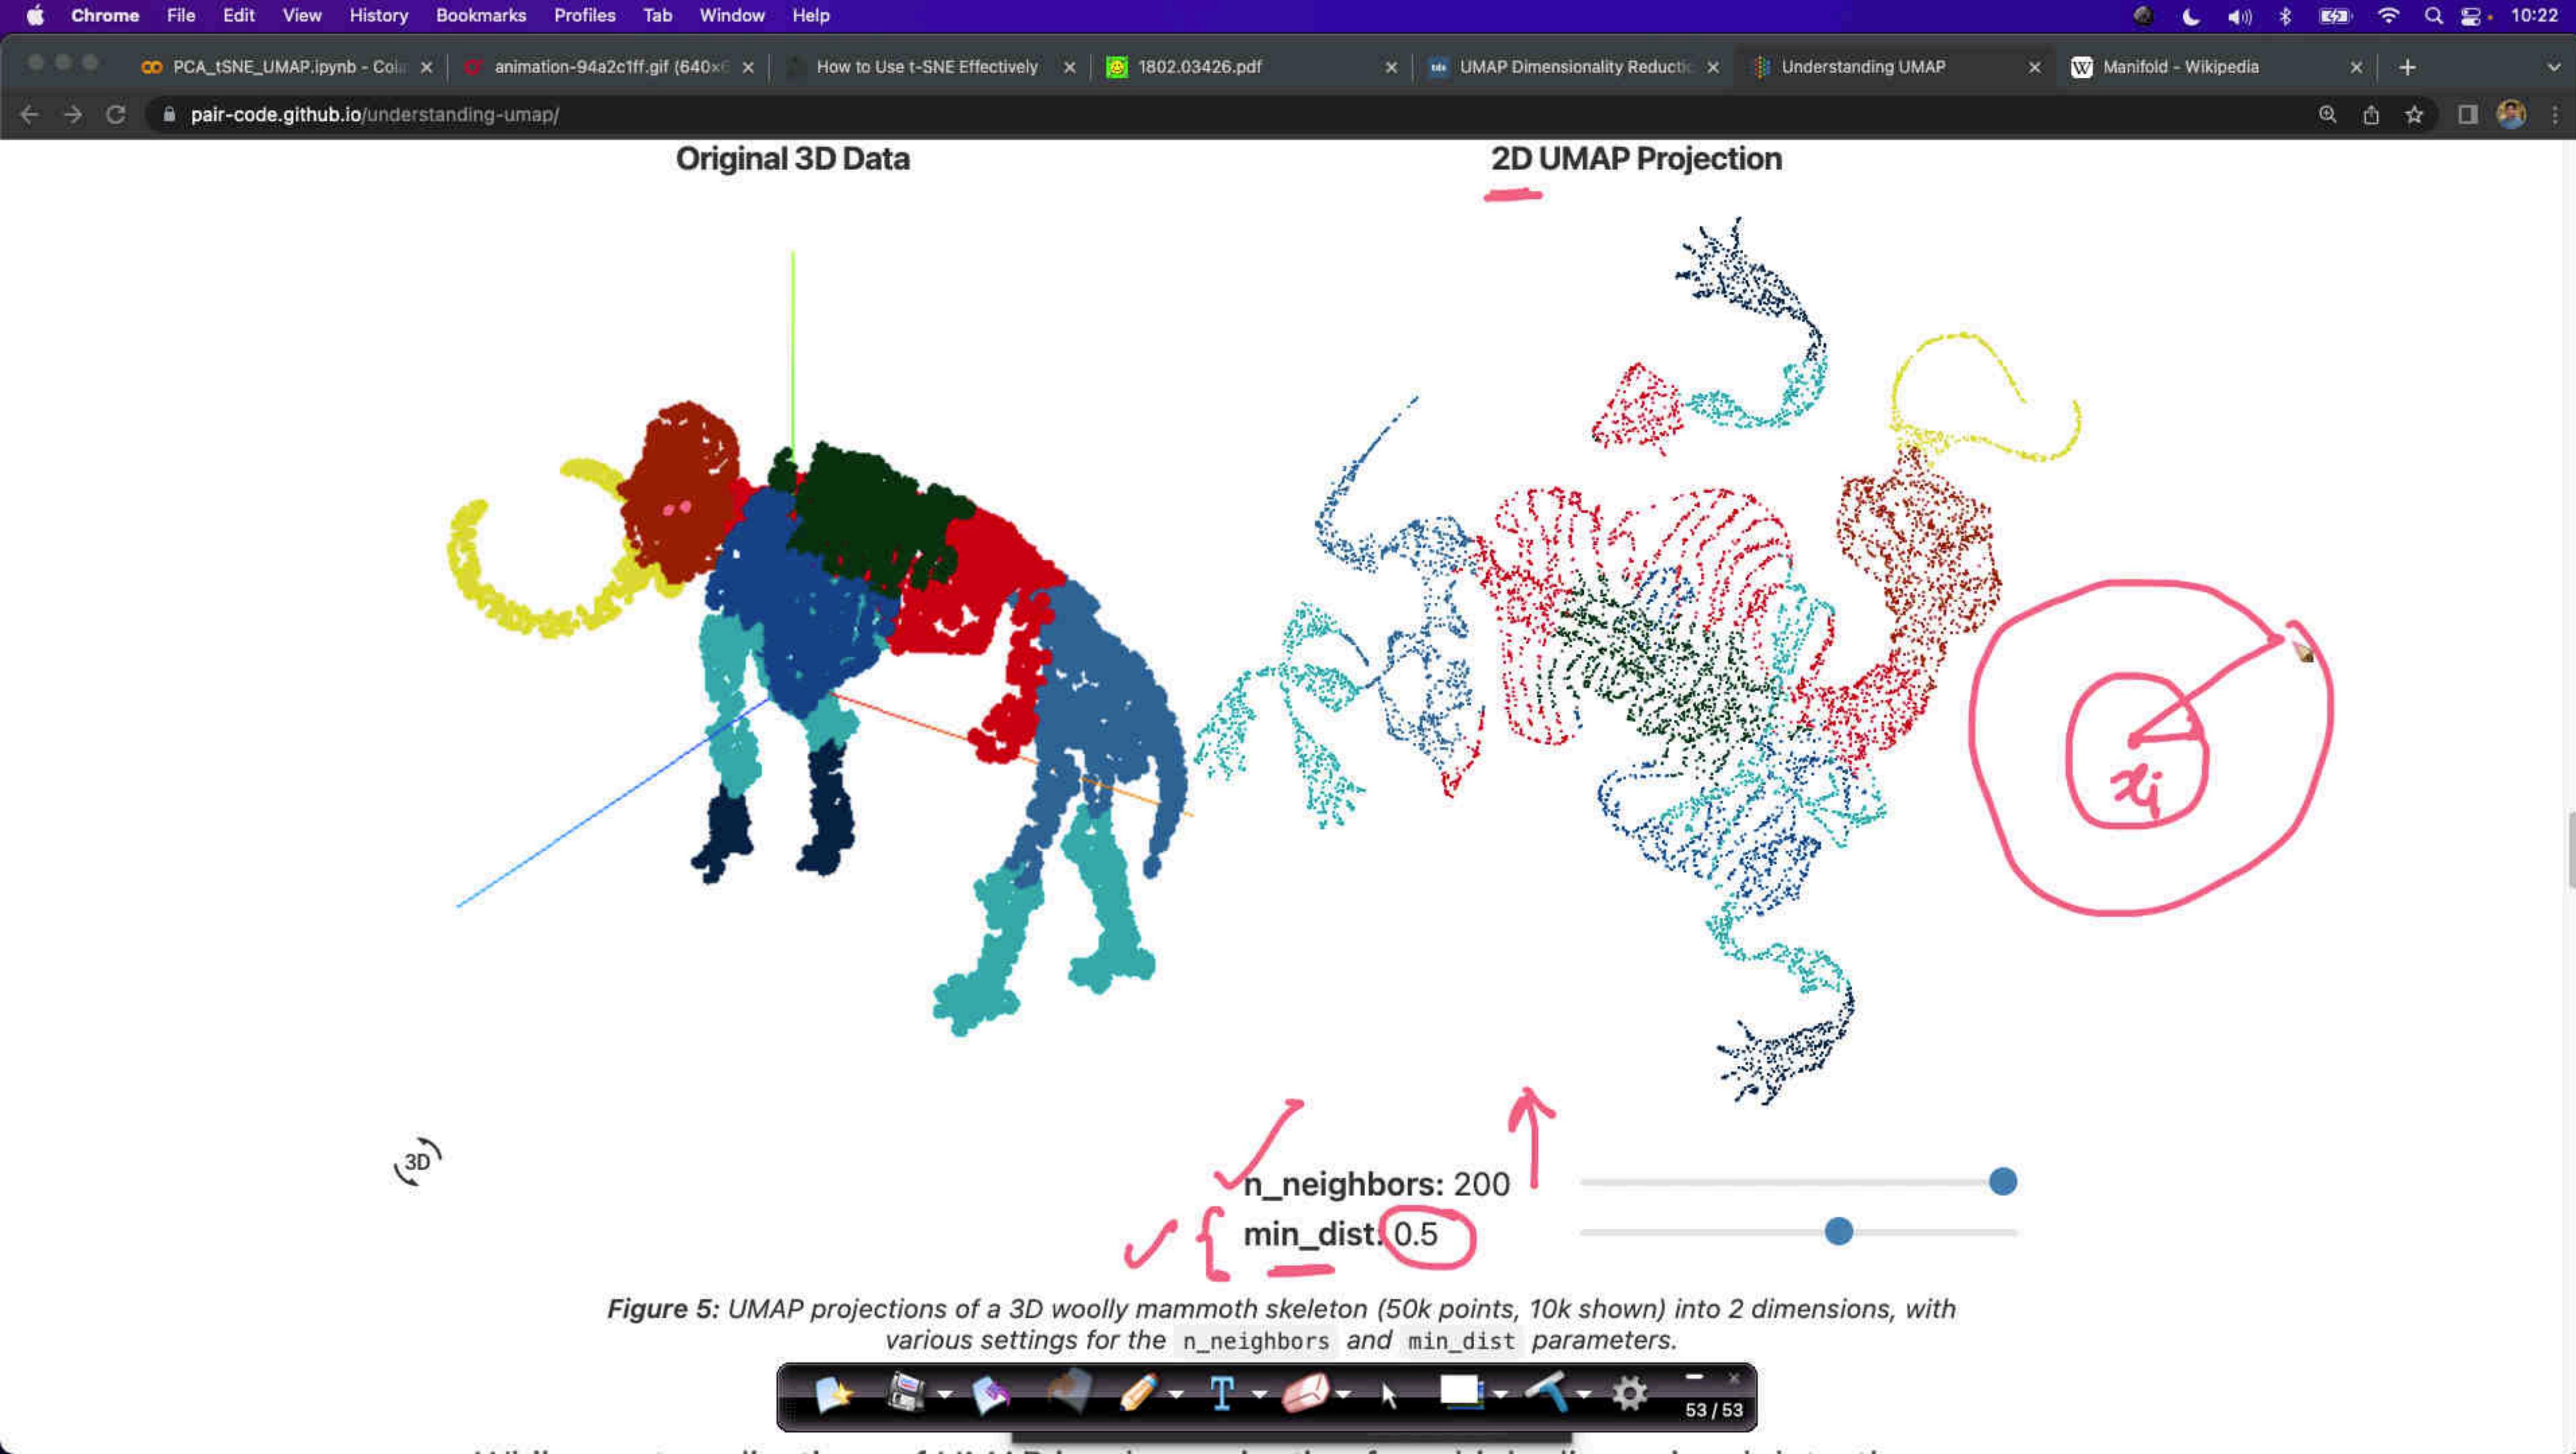


Figure 5: UMAP projections of a 3D woolly mammoth skeleton (50k points, 10k shown) into 2 dimensions, with various settings for the `n_neighbors` and `min_dist` parameters.



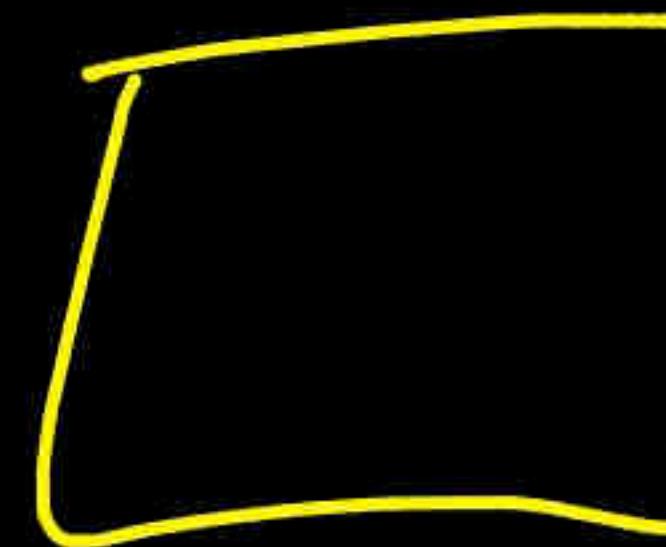
Obj: Viz

d-dim

(d')

→ 2dim

→ 3dim



reading and understanding its results requires some care. It's worth revisiting our [previous work on \(mis\)reading t-SNE](#), since many of the same takeaways apply to UMAP:

**1. Hyperparameters really matter**

Choosing good values isn't easy, and depends on both the data and your goals (eg, how tightly packed the projection ought to be). This is where UMAP's speed is a big advantage - By running UMAP multiple times with a variety of hyperparameters, you can get a better sense of how the projection is affected by its parameters.

**2. Cluster sizes in a UMAP plot mean nothing**

Just as in t-SNE, the size of clusters relative to each other is essentially meaningless. This is because UMAP uses local notions of distance to construct its high-dimensional graph representation.

**3. Distances between clusters might not mean anything**

Likewise, the distances between clusters is likely to be meaningless. While it's true that the global positions of clusters are better preserved in UMAP, the distances between them are not meaningful. Again, this is due to using local distances when constructing the graph.

**4. Random noise doesn't always look random.**

Especially at low noise levels, the points in a UMAP plot can appear very regular and organized, which can be misleading. Handwritten notes on the right side of the slide include "UMAP" with a downward arrow, "→ Speed", and "→ Global structure".

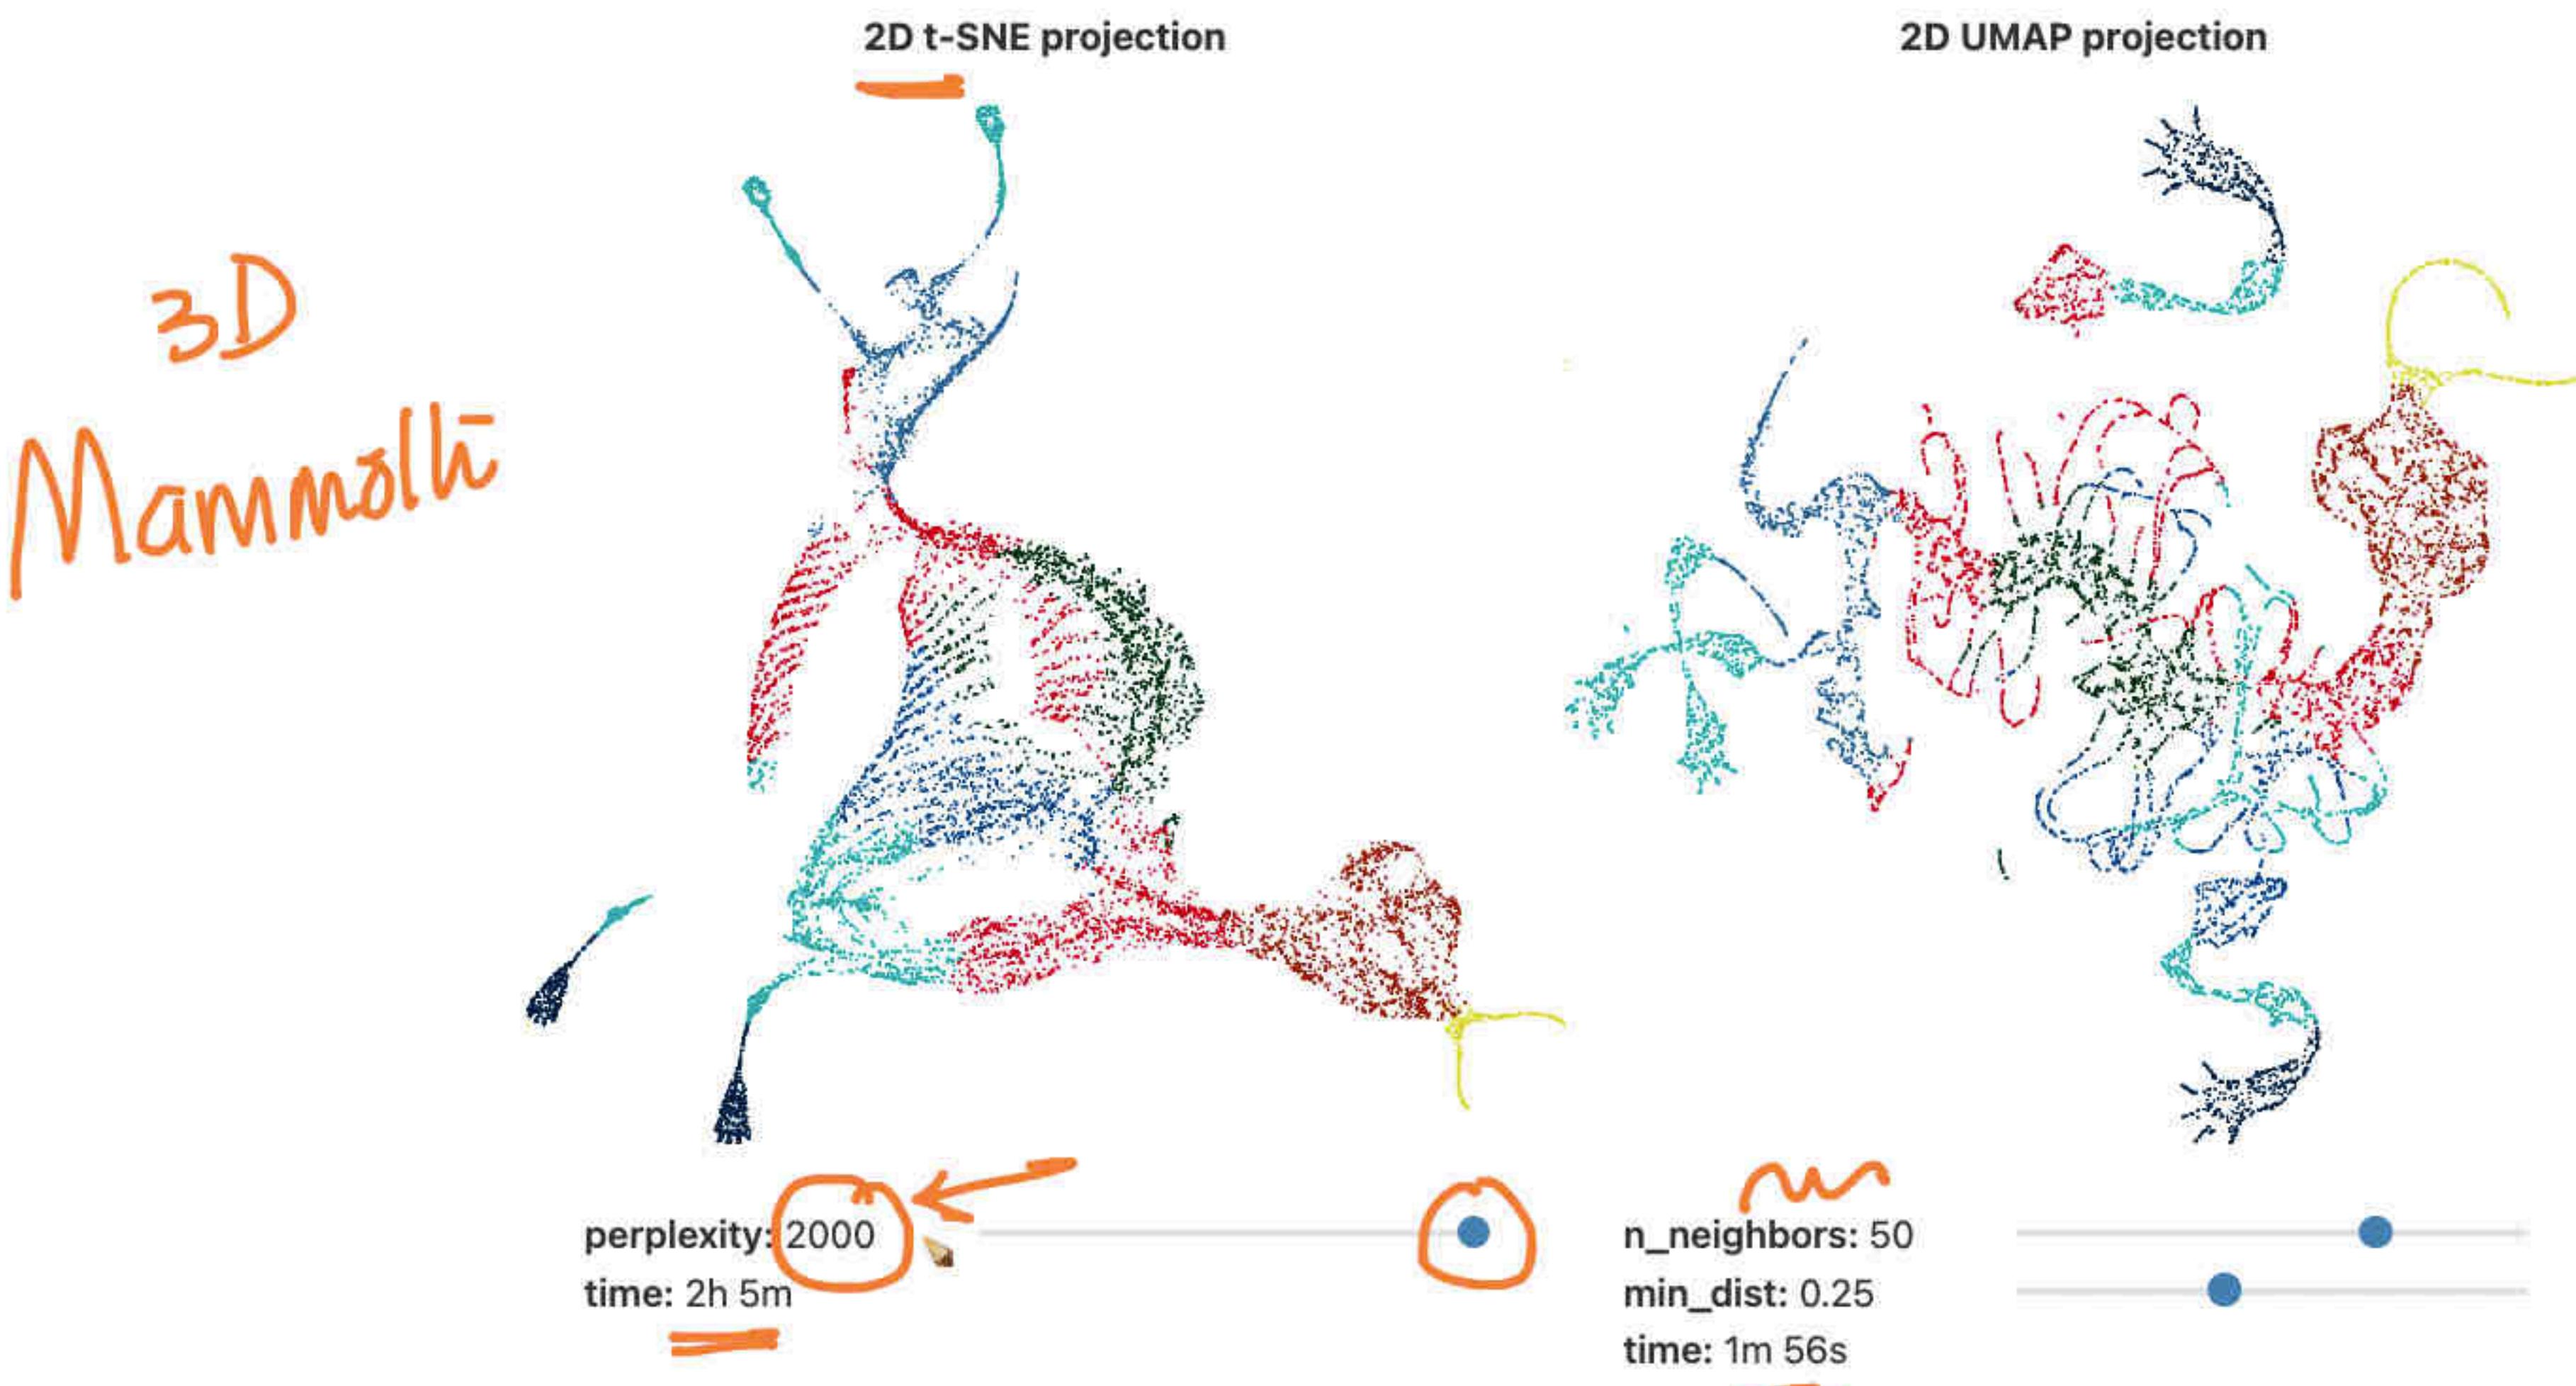
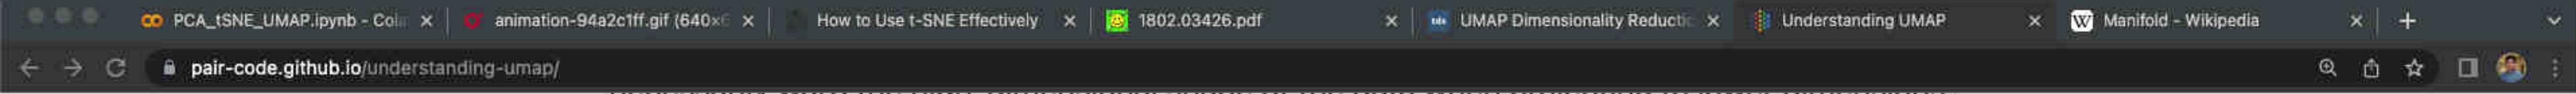


Figure 6: A comparison between UMAP and t-SNE projections of a 3D woolly mammoth skeleton (50,000 points) into 2 dimensions, with various settings for parameters. Notice how much more global structure is preserved with UMAP particularly with larger values of n\_neighbors.

pair-code.github.io/understanding-umap/

## HOW TO (MIS)READ UMAP

While UMAP offers a number of advantages over t-SNE, it's by no means a silver bullet - and reading and understanding its results requires some care. It's worth revisiting our [previous work on \(mis\)reading t-SNE](#), since many of the same takeaways apply to UMAP:

✓ **1. Hyperparameters really matter**

Choosing good values isn't easy, and depends on both the data and your goals (eg, how tightly packed the projection ought to be). This is where UMAP's speed is a big advantage - By running UMAP multiple times with a variety of hyperparameters, you can get a better sense of how the projection is affected by its parameters.

**2. Cluster sizes in a UMAP plot mean nothing**

Just as in t-SNE, the size of clusters relative to each other is essentially meaningless. This is because UMAP uses local notions of distance to construct its high-dimensional graph representation.

**3. Distances between clusters might not mean anything**

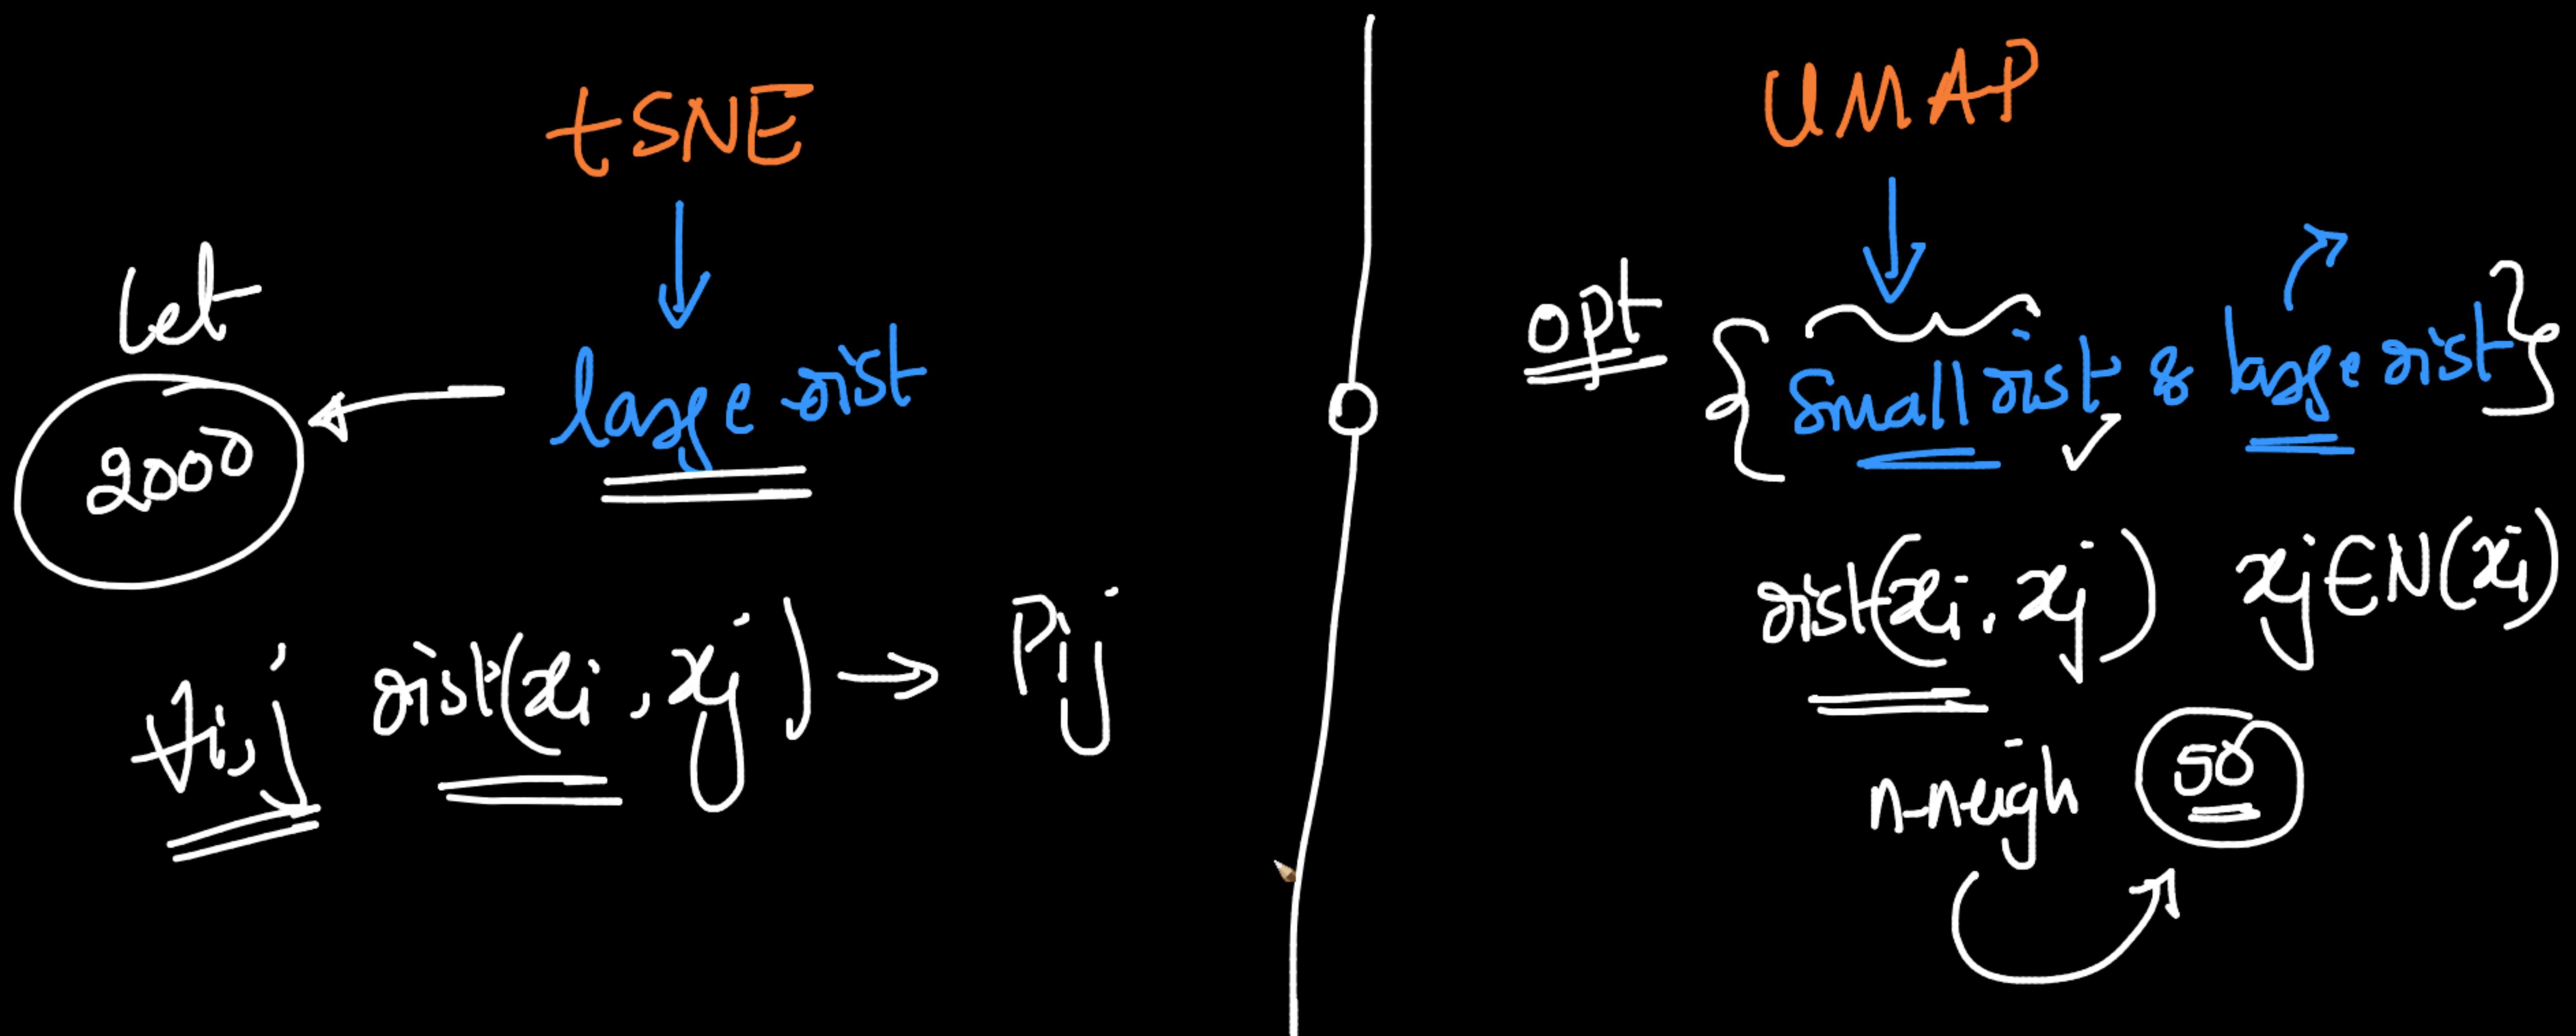
Likewise, the distances between clusters is likely to be meaningless. While it's true that the global positions of clusters are better preserved in UMAP, the distances between them are not meaningful. Again, this is due to using local distances when constructing the graph.

**4. Random noise doesn't always look random.**

Especially at low values of `n_neighbors`, spurious clustering can be observed.

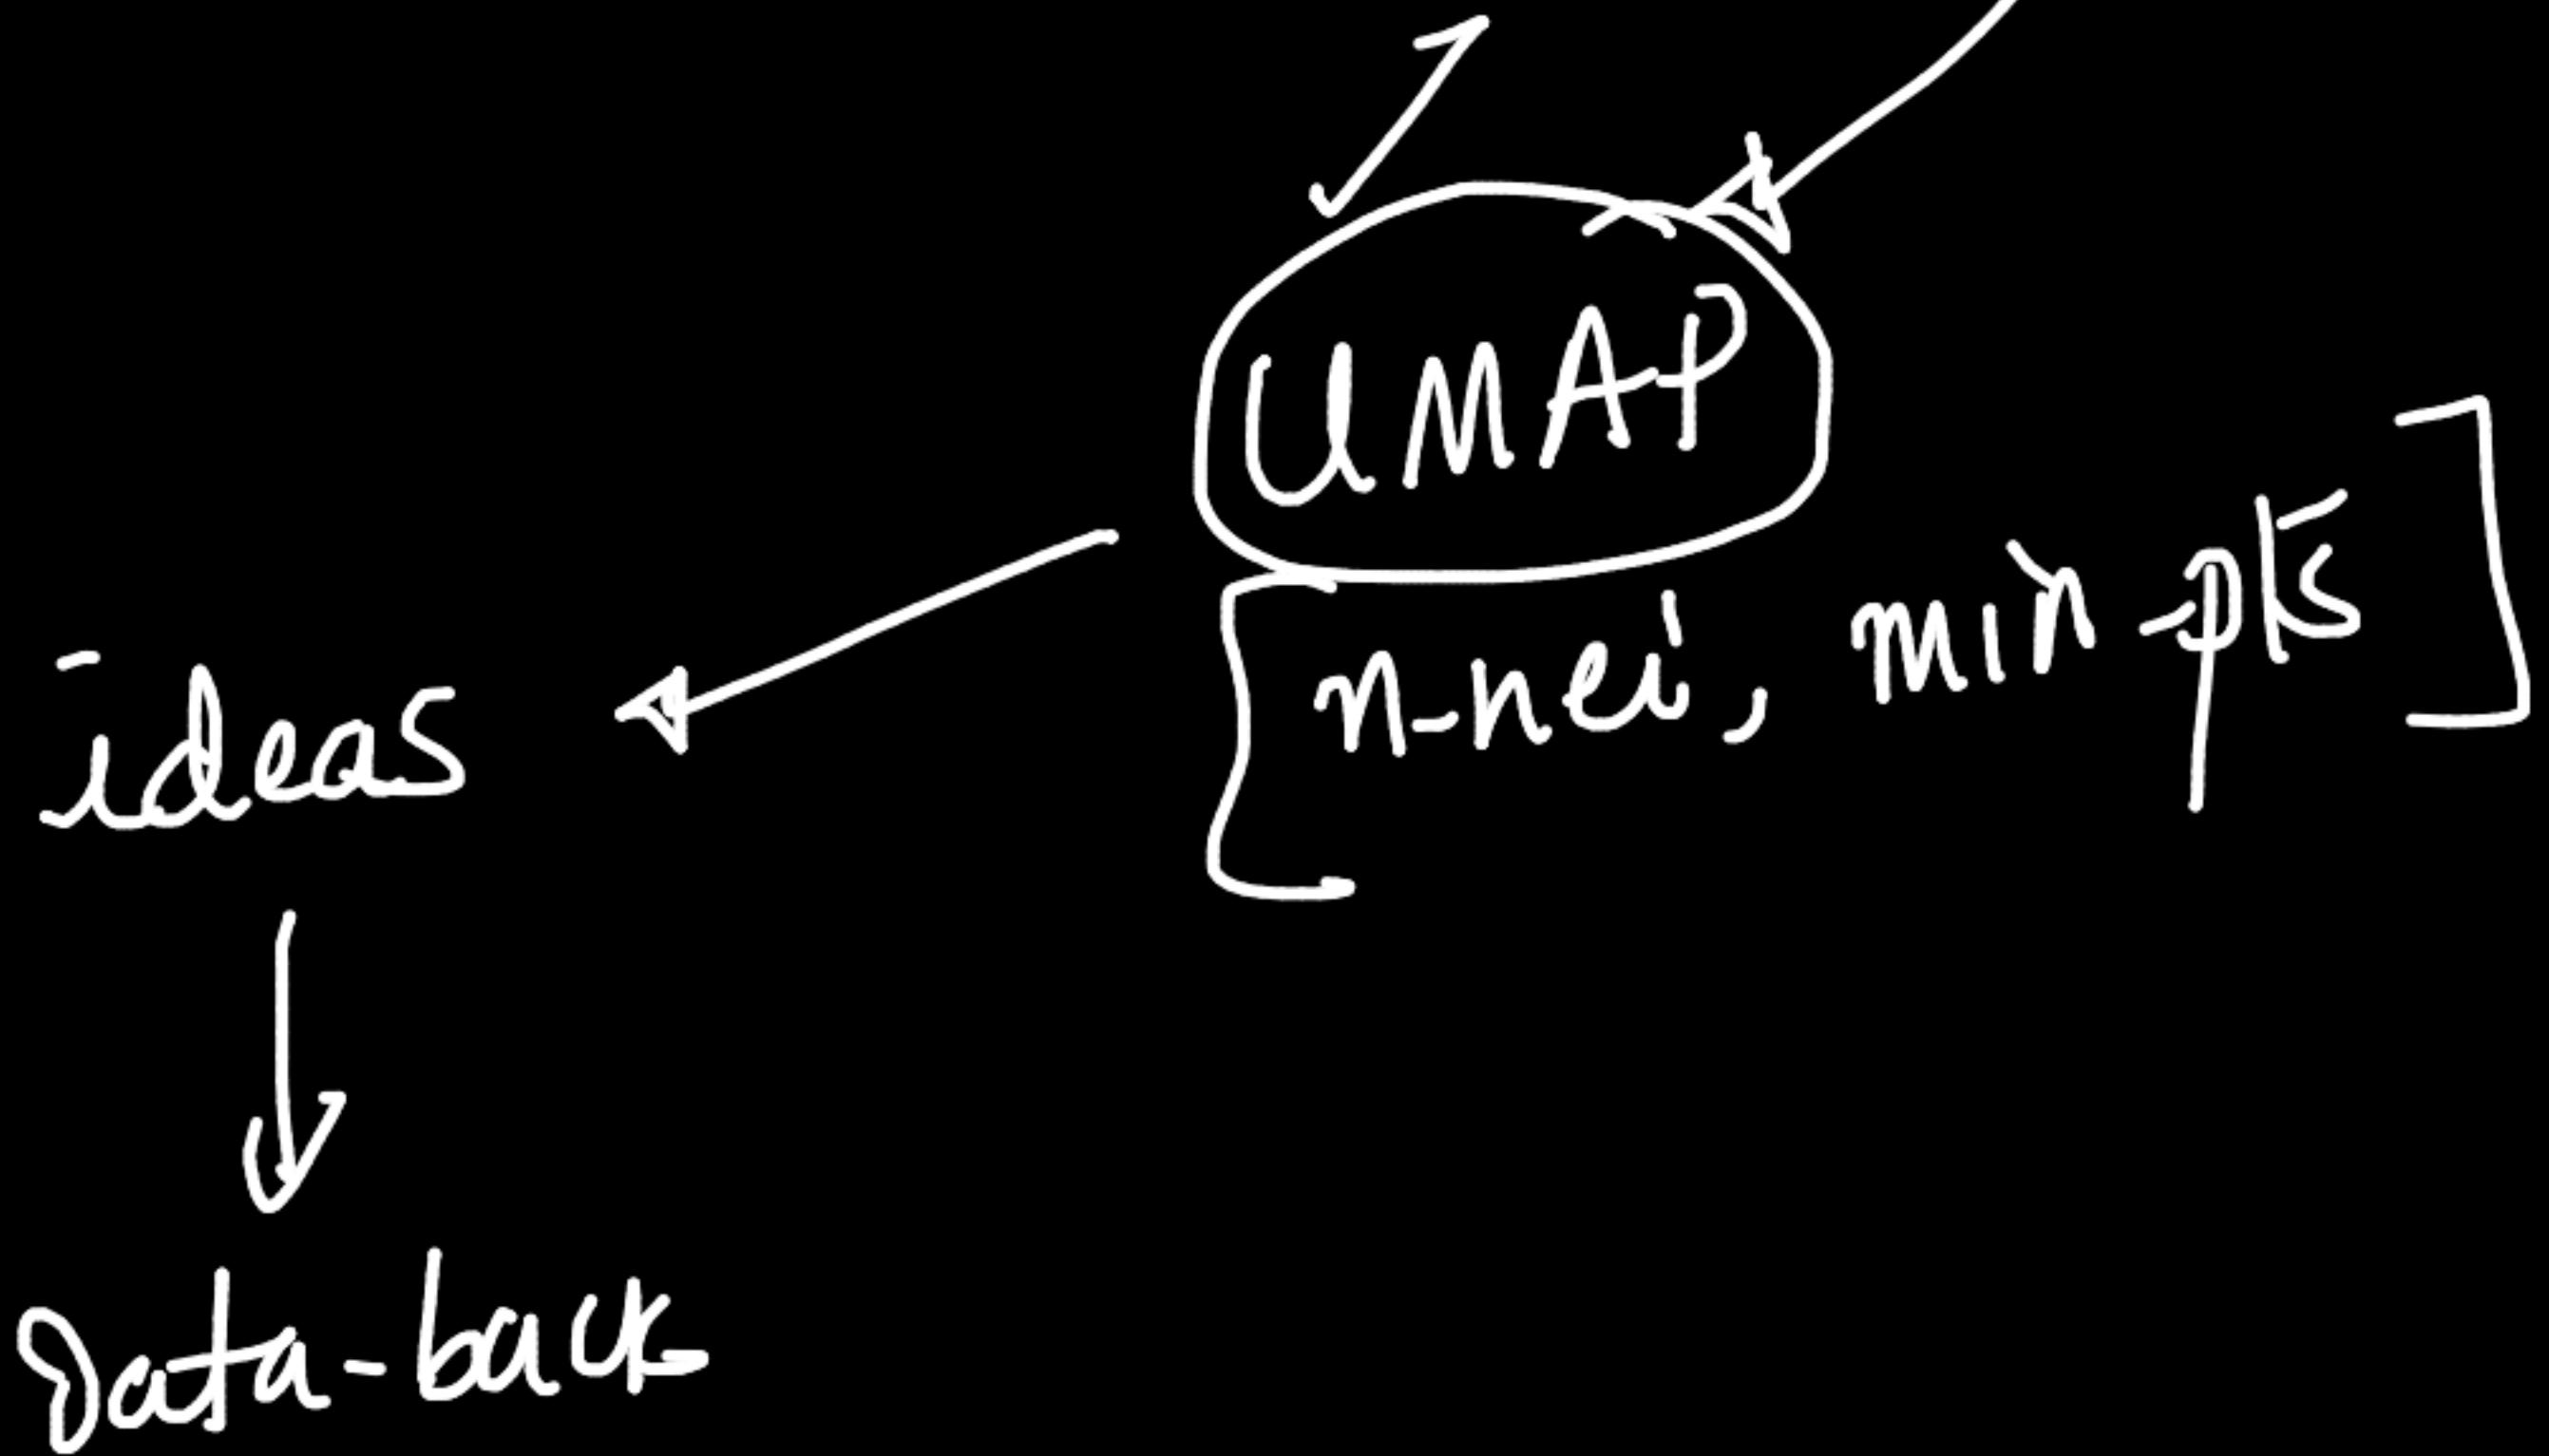
5. You may n

Speed &  
global-stre~~ss~~



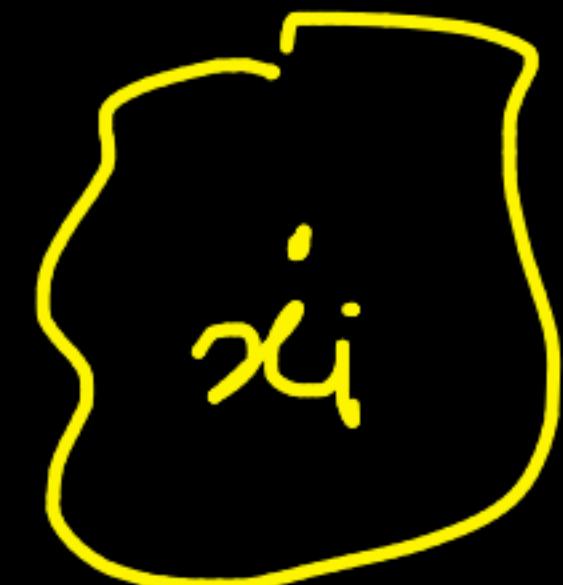
~~Takeaways:~~

Viz. high-dim  
data





local-str



small-dist  
↳ large Wule

global str

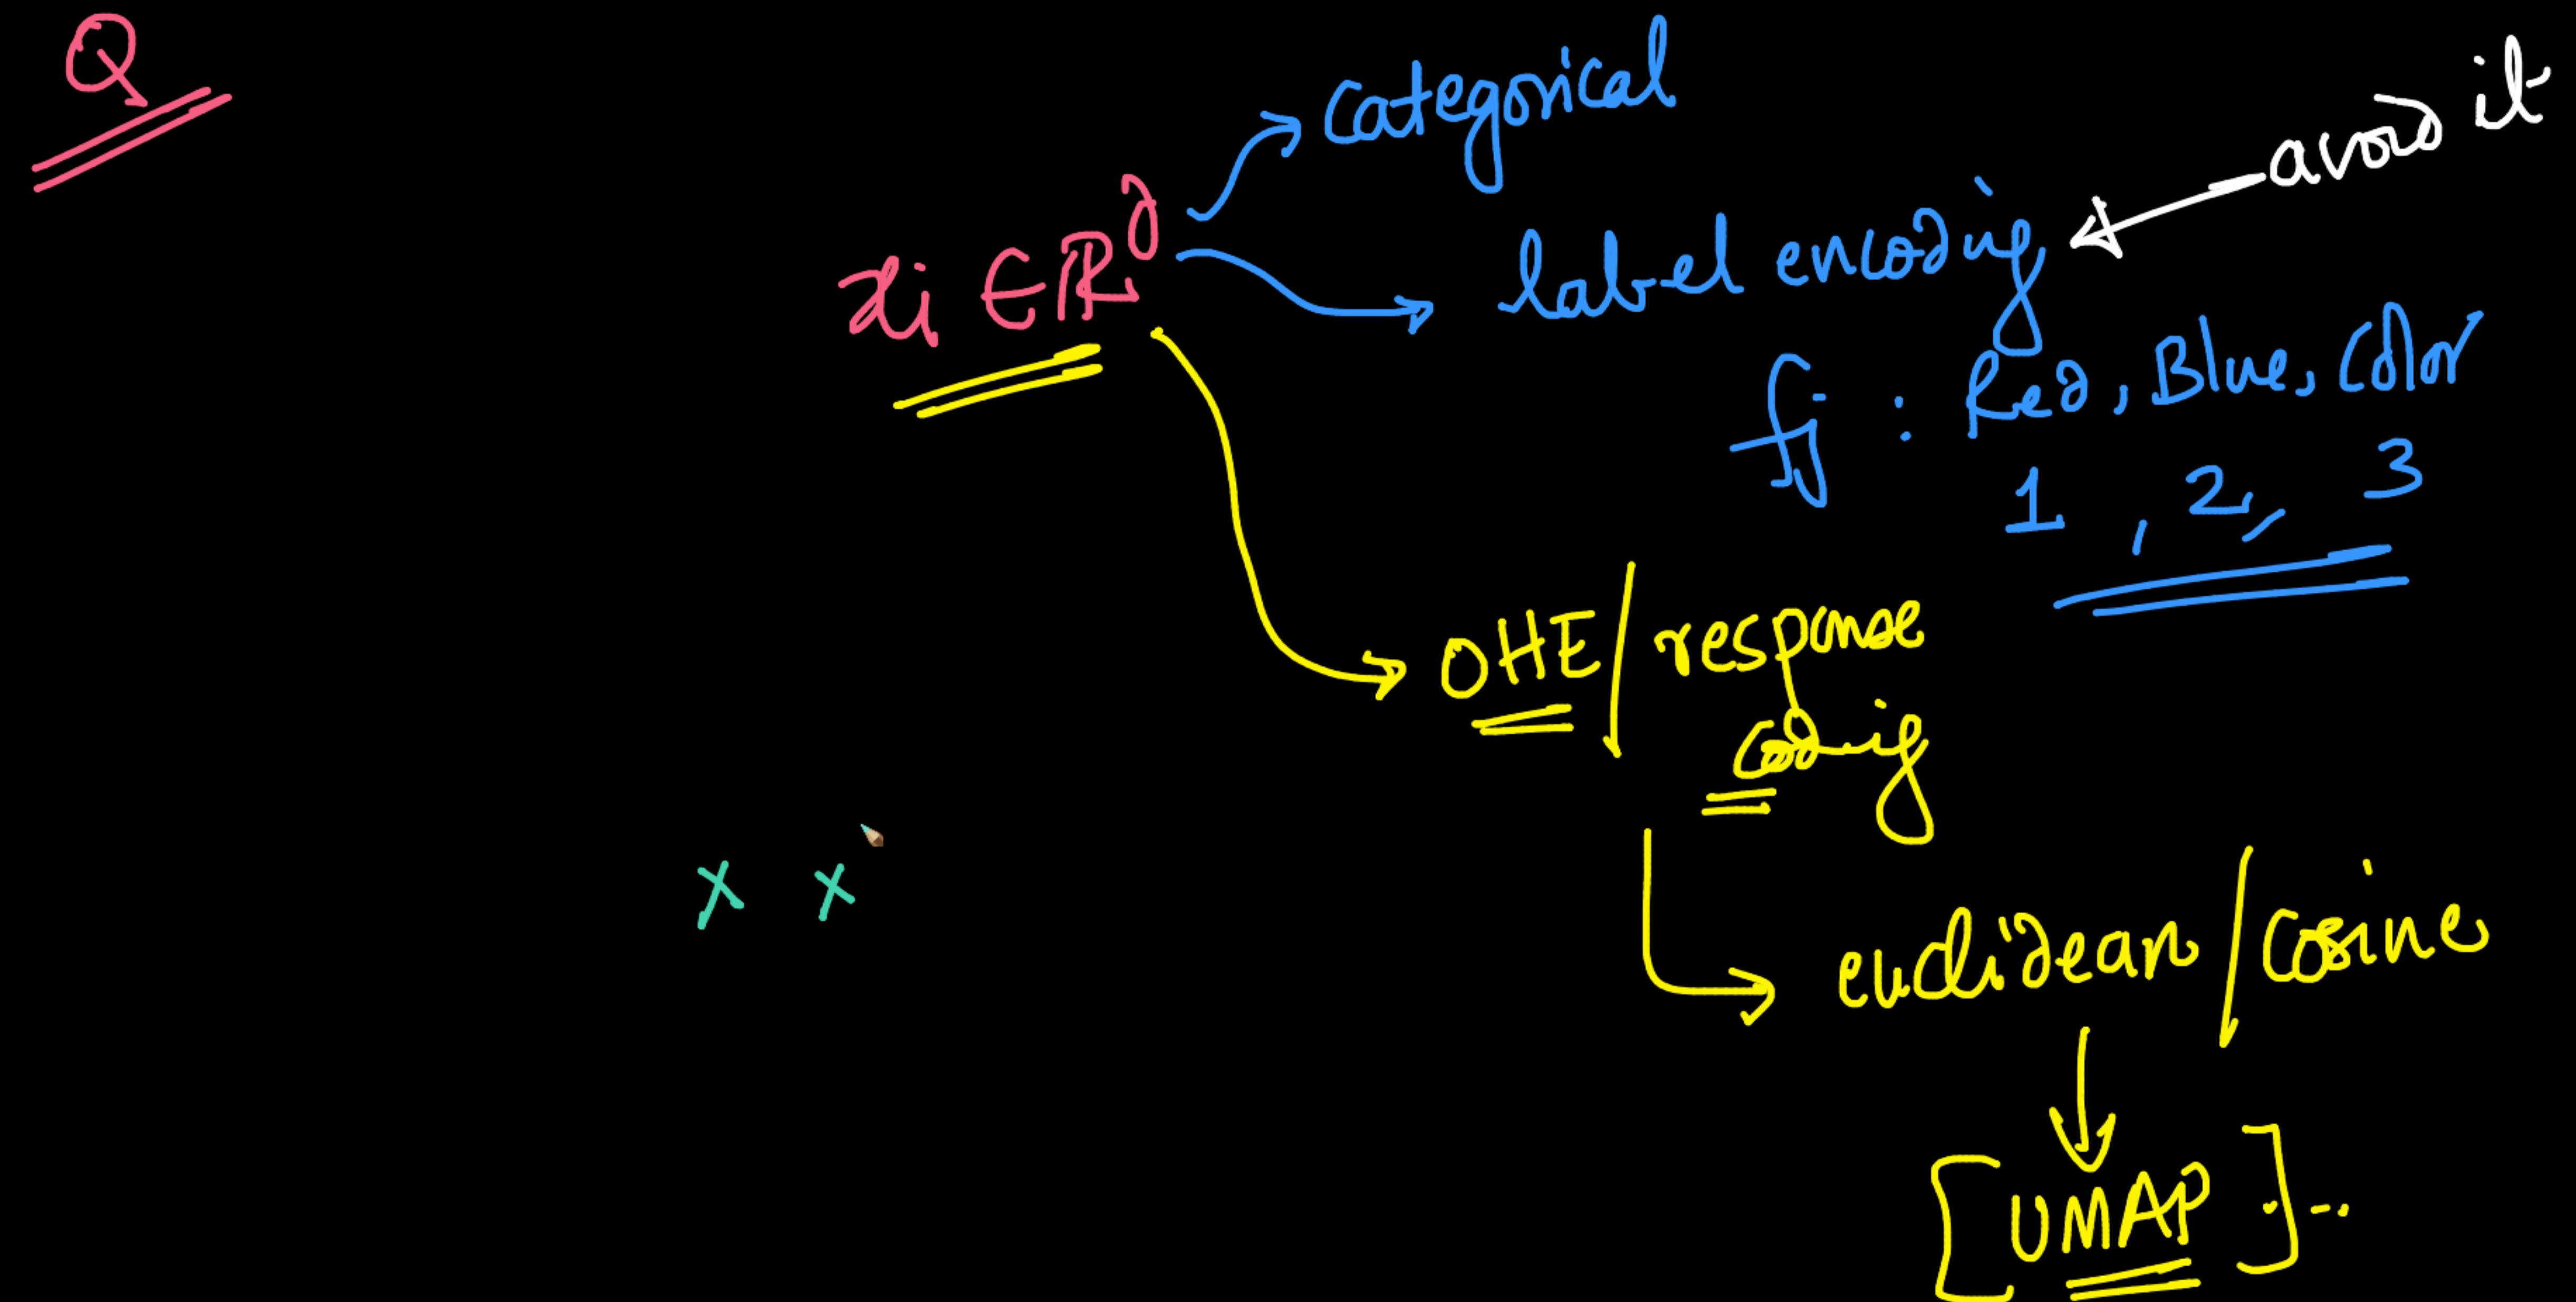


$x_i$

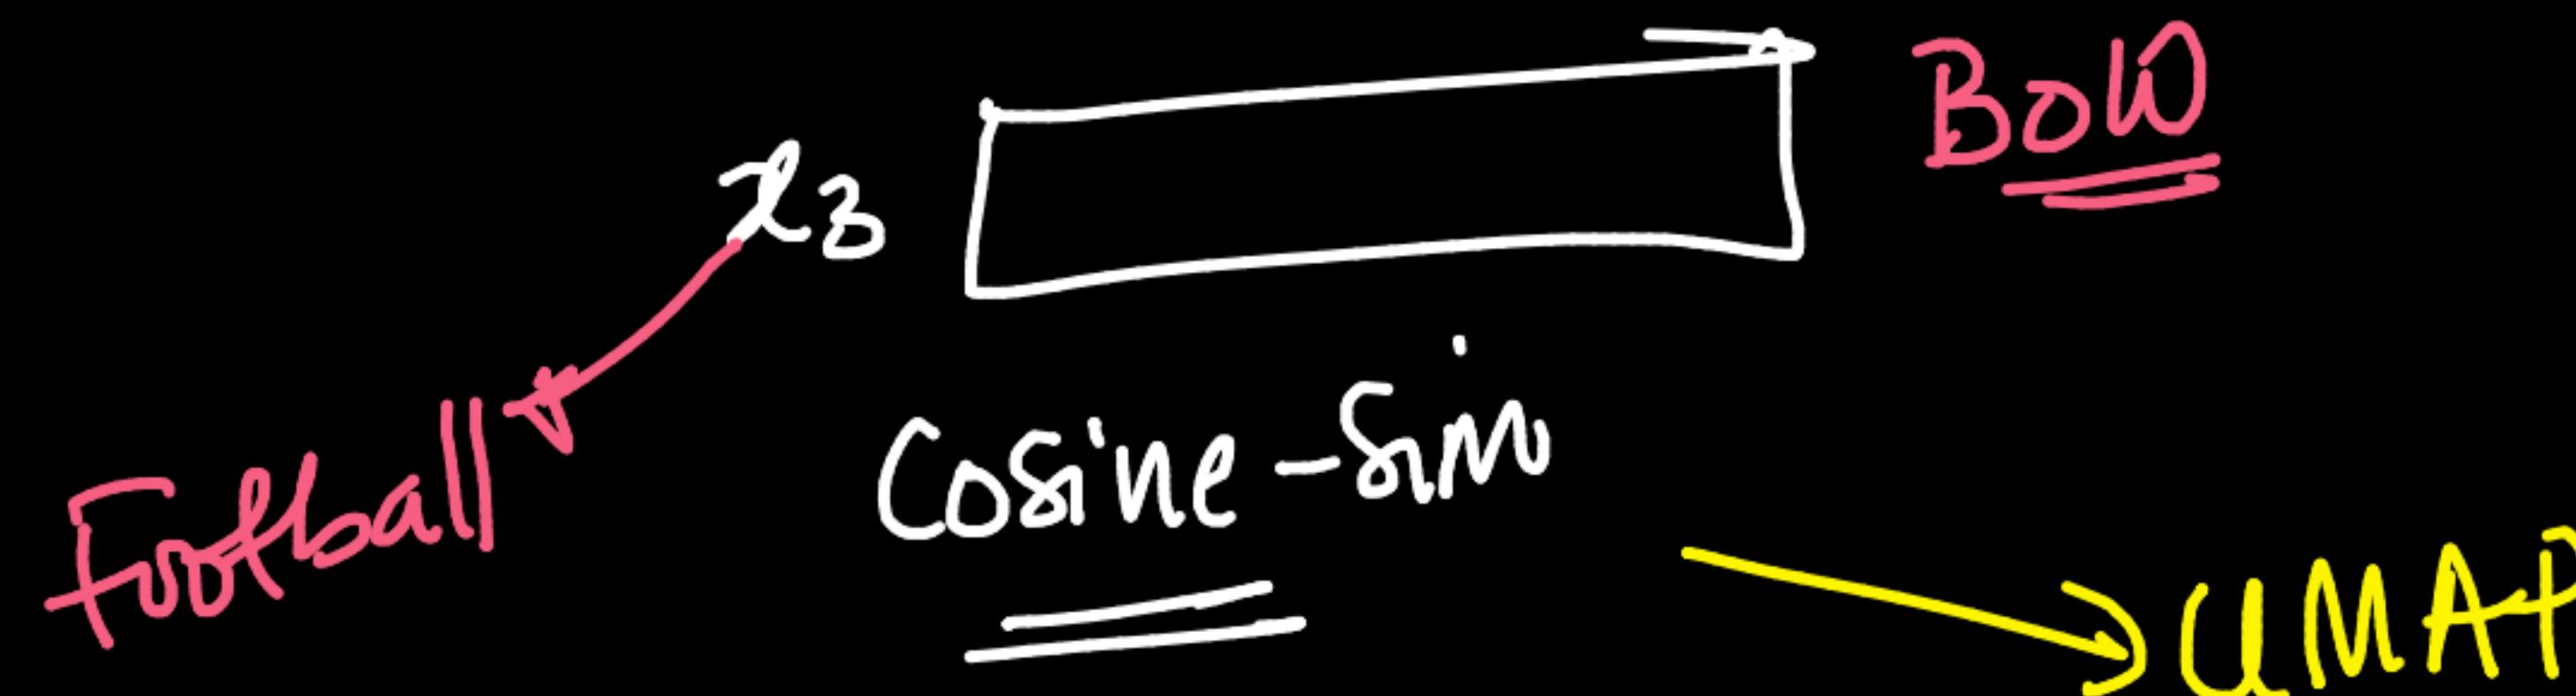
large dist /  
small Wule

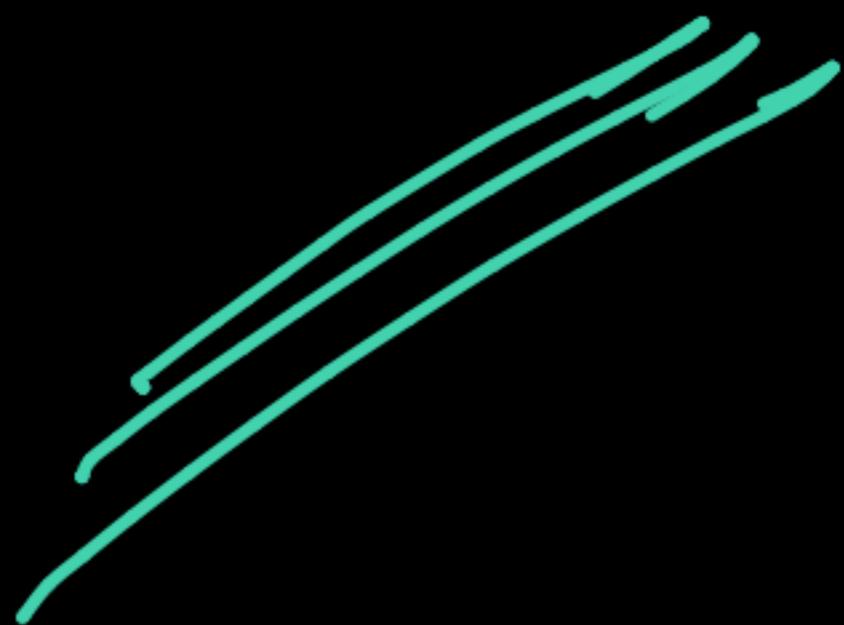
$y_j$

10:40



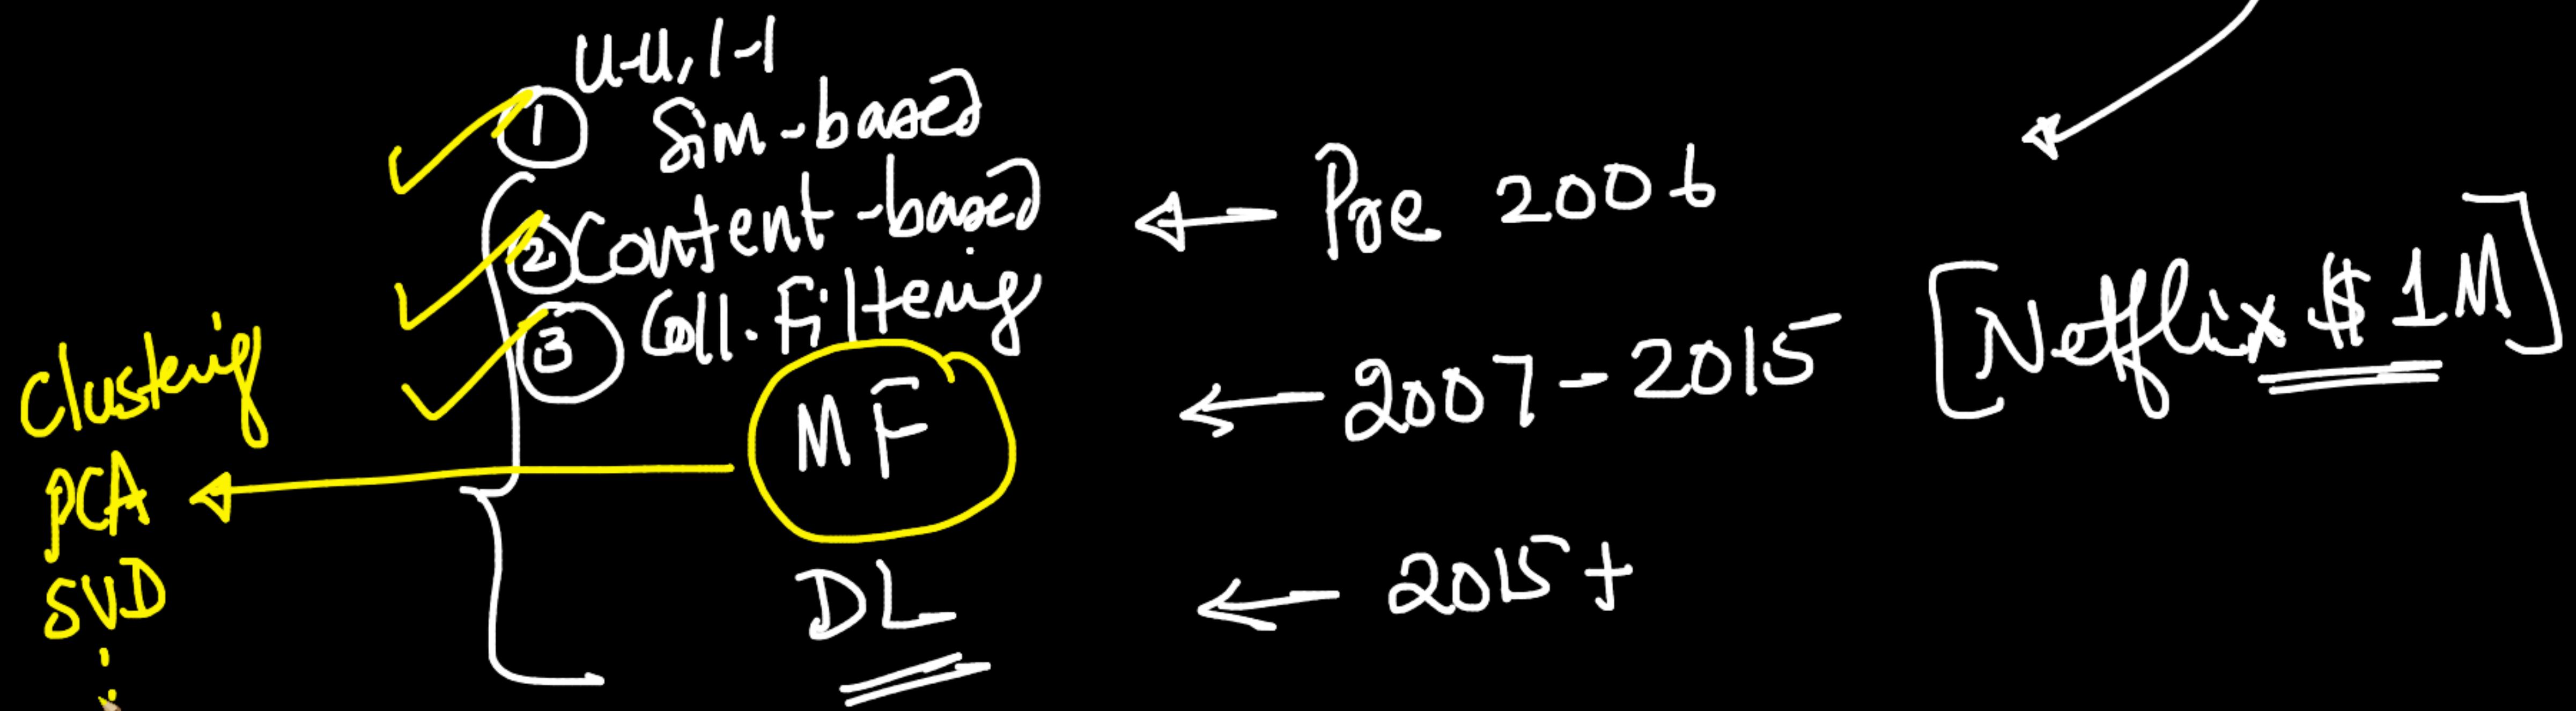
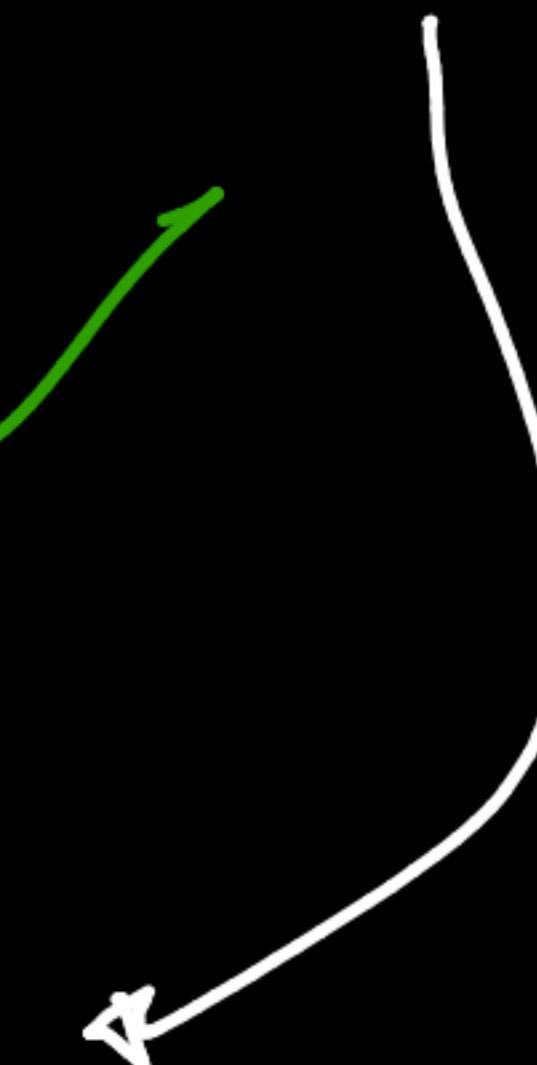
exercise





RecSys

Netflix; TikTok; Amzn; YouTube



# Formulation

User:  $u_i \quad i=1 \rightarrow n$  }  $\xrightarrow{\text{Billions}}$   $\rightarrow \text{V-V-large}$

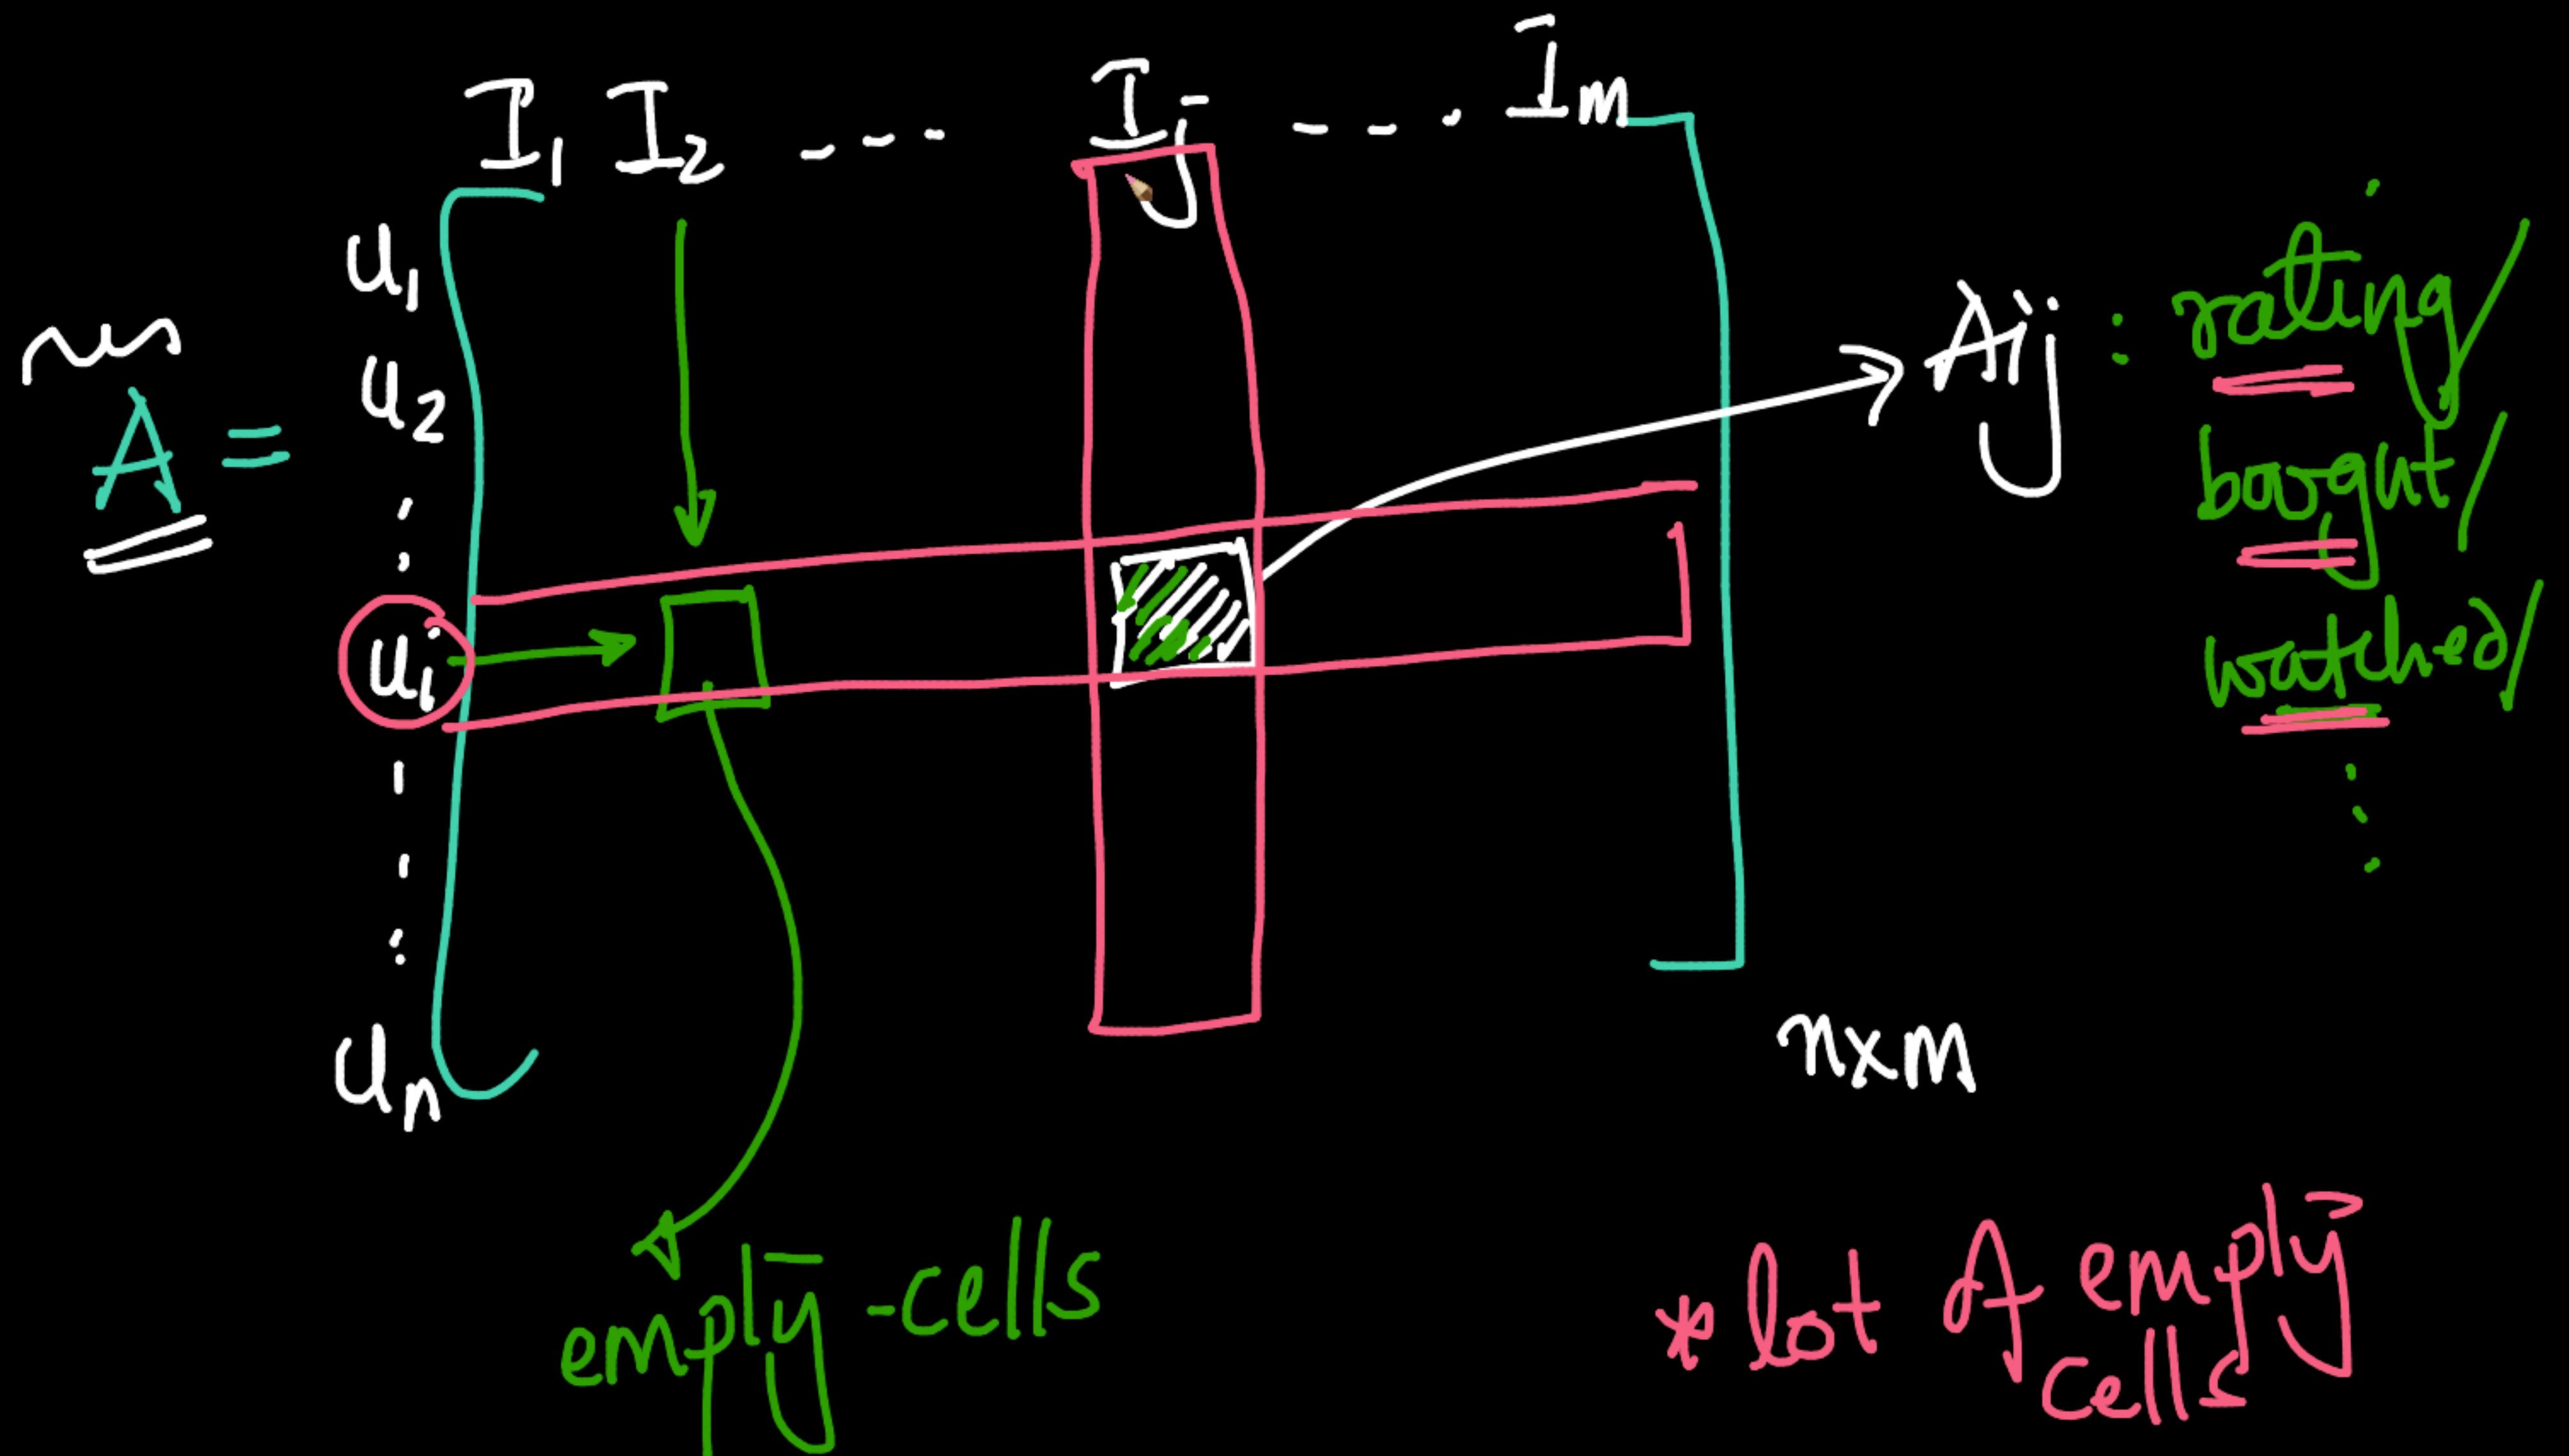
items:  $I_j \quad j=1 \rightarrow M$  }  $\xrightarrow{\text{loos of million}}$   
(Movie/product)

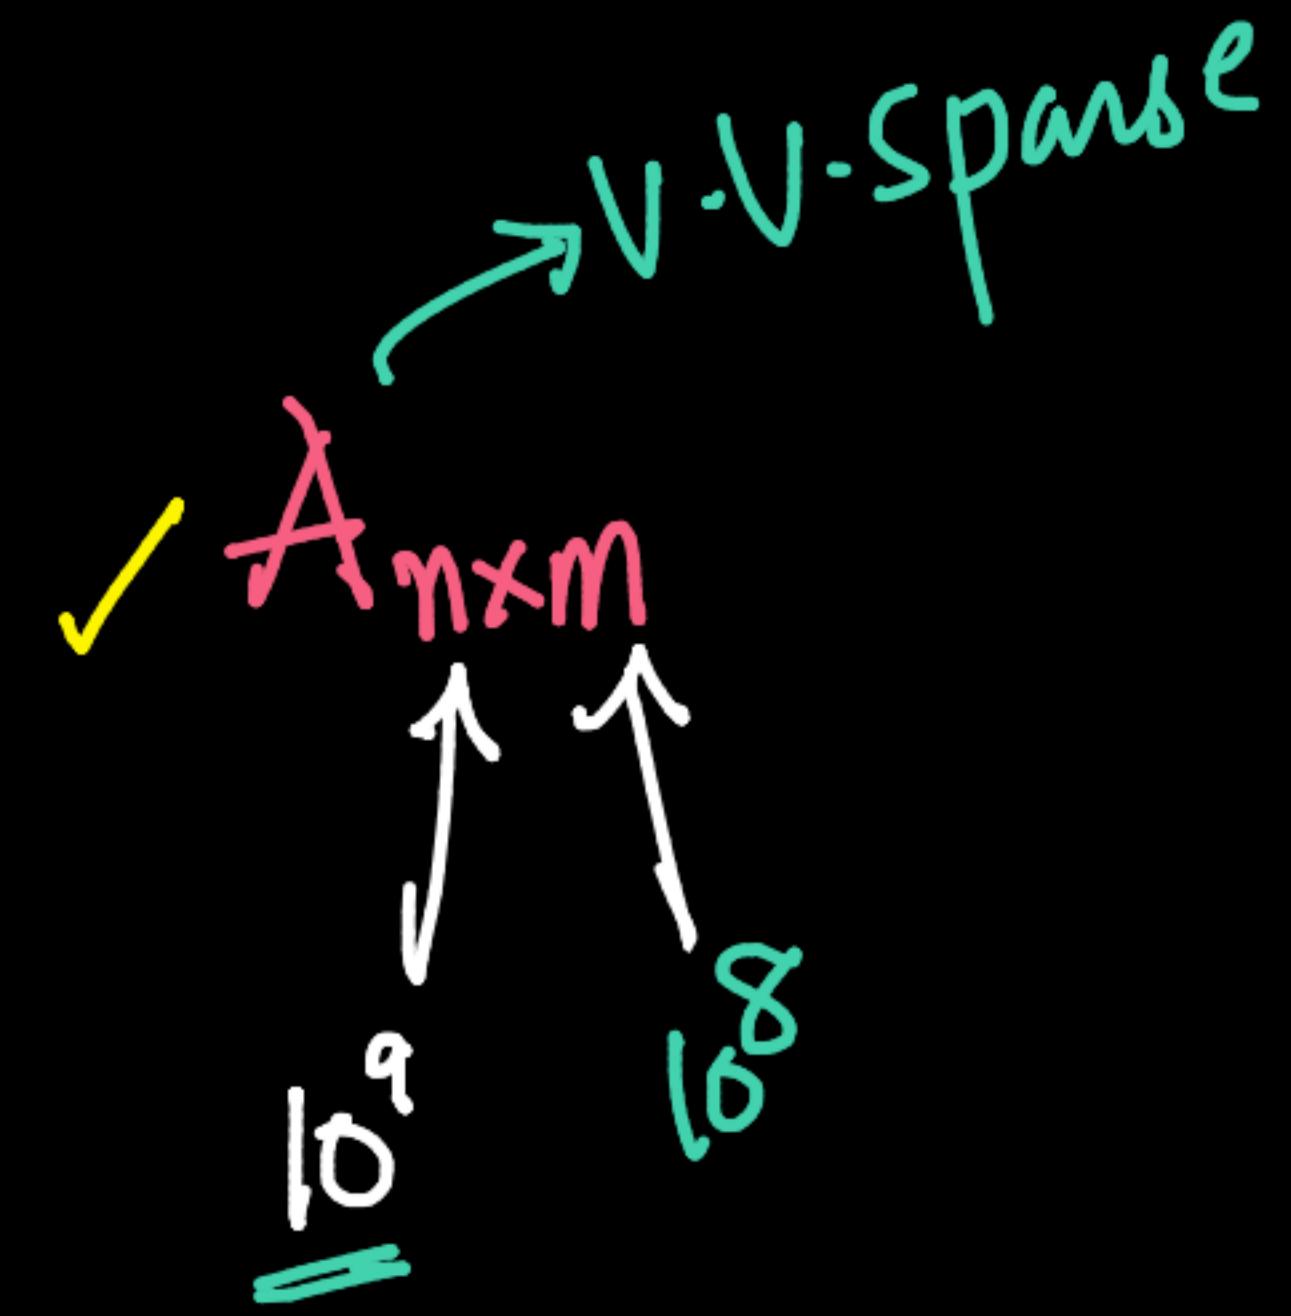
Task:  $u_i \rightarrow I_{10}, I_{12}, I_{16}, I_{18}$

based on historical data

Dataset

historical  
data





# cells: V.V.lage

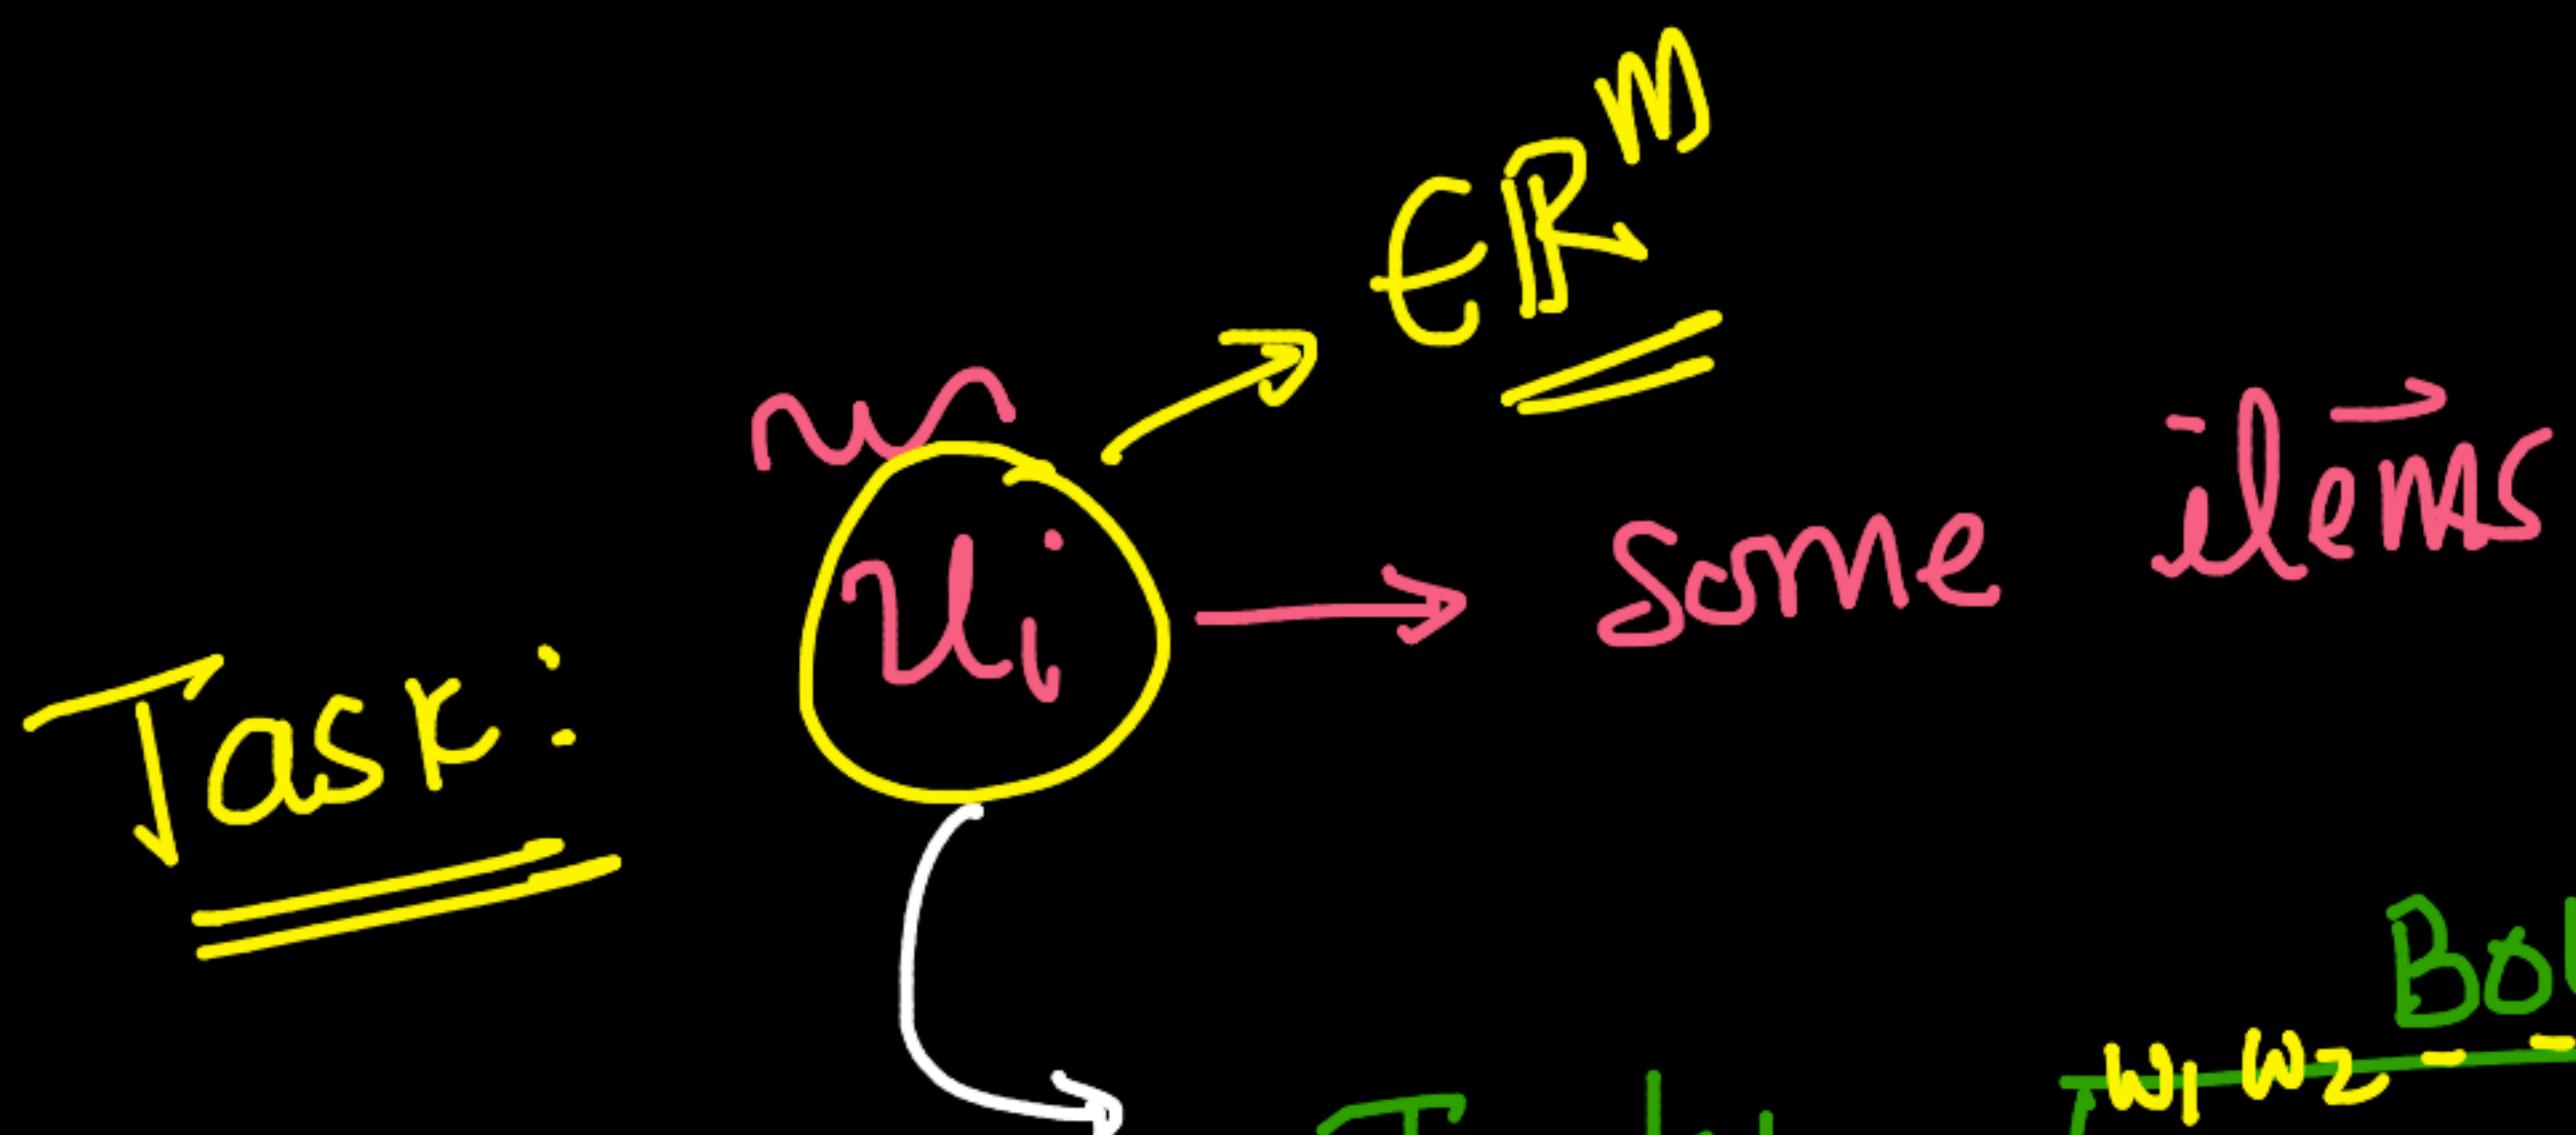
# non empty cells / Total

Sparcity:

$$\frac{10^9 \times 1000}{10^{17}} = \frac{10^{12}}{10^{17}}$$

$$= \boxed{\frac{1}{10^5}} = 10^{-5}$$

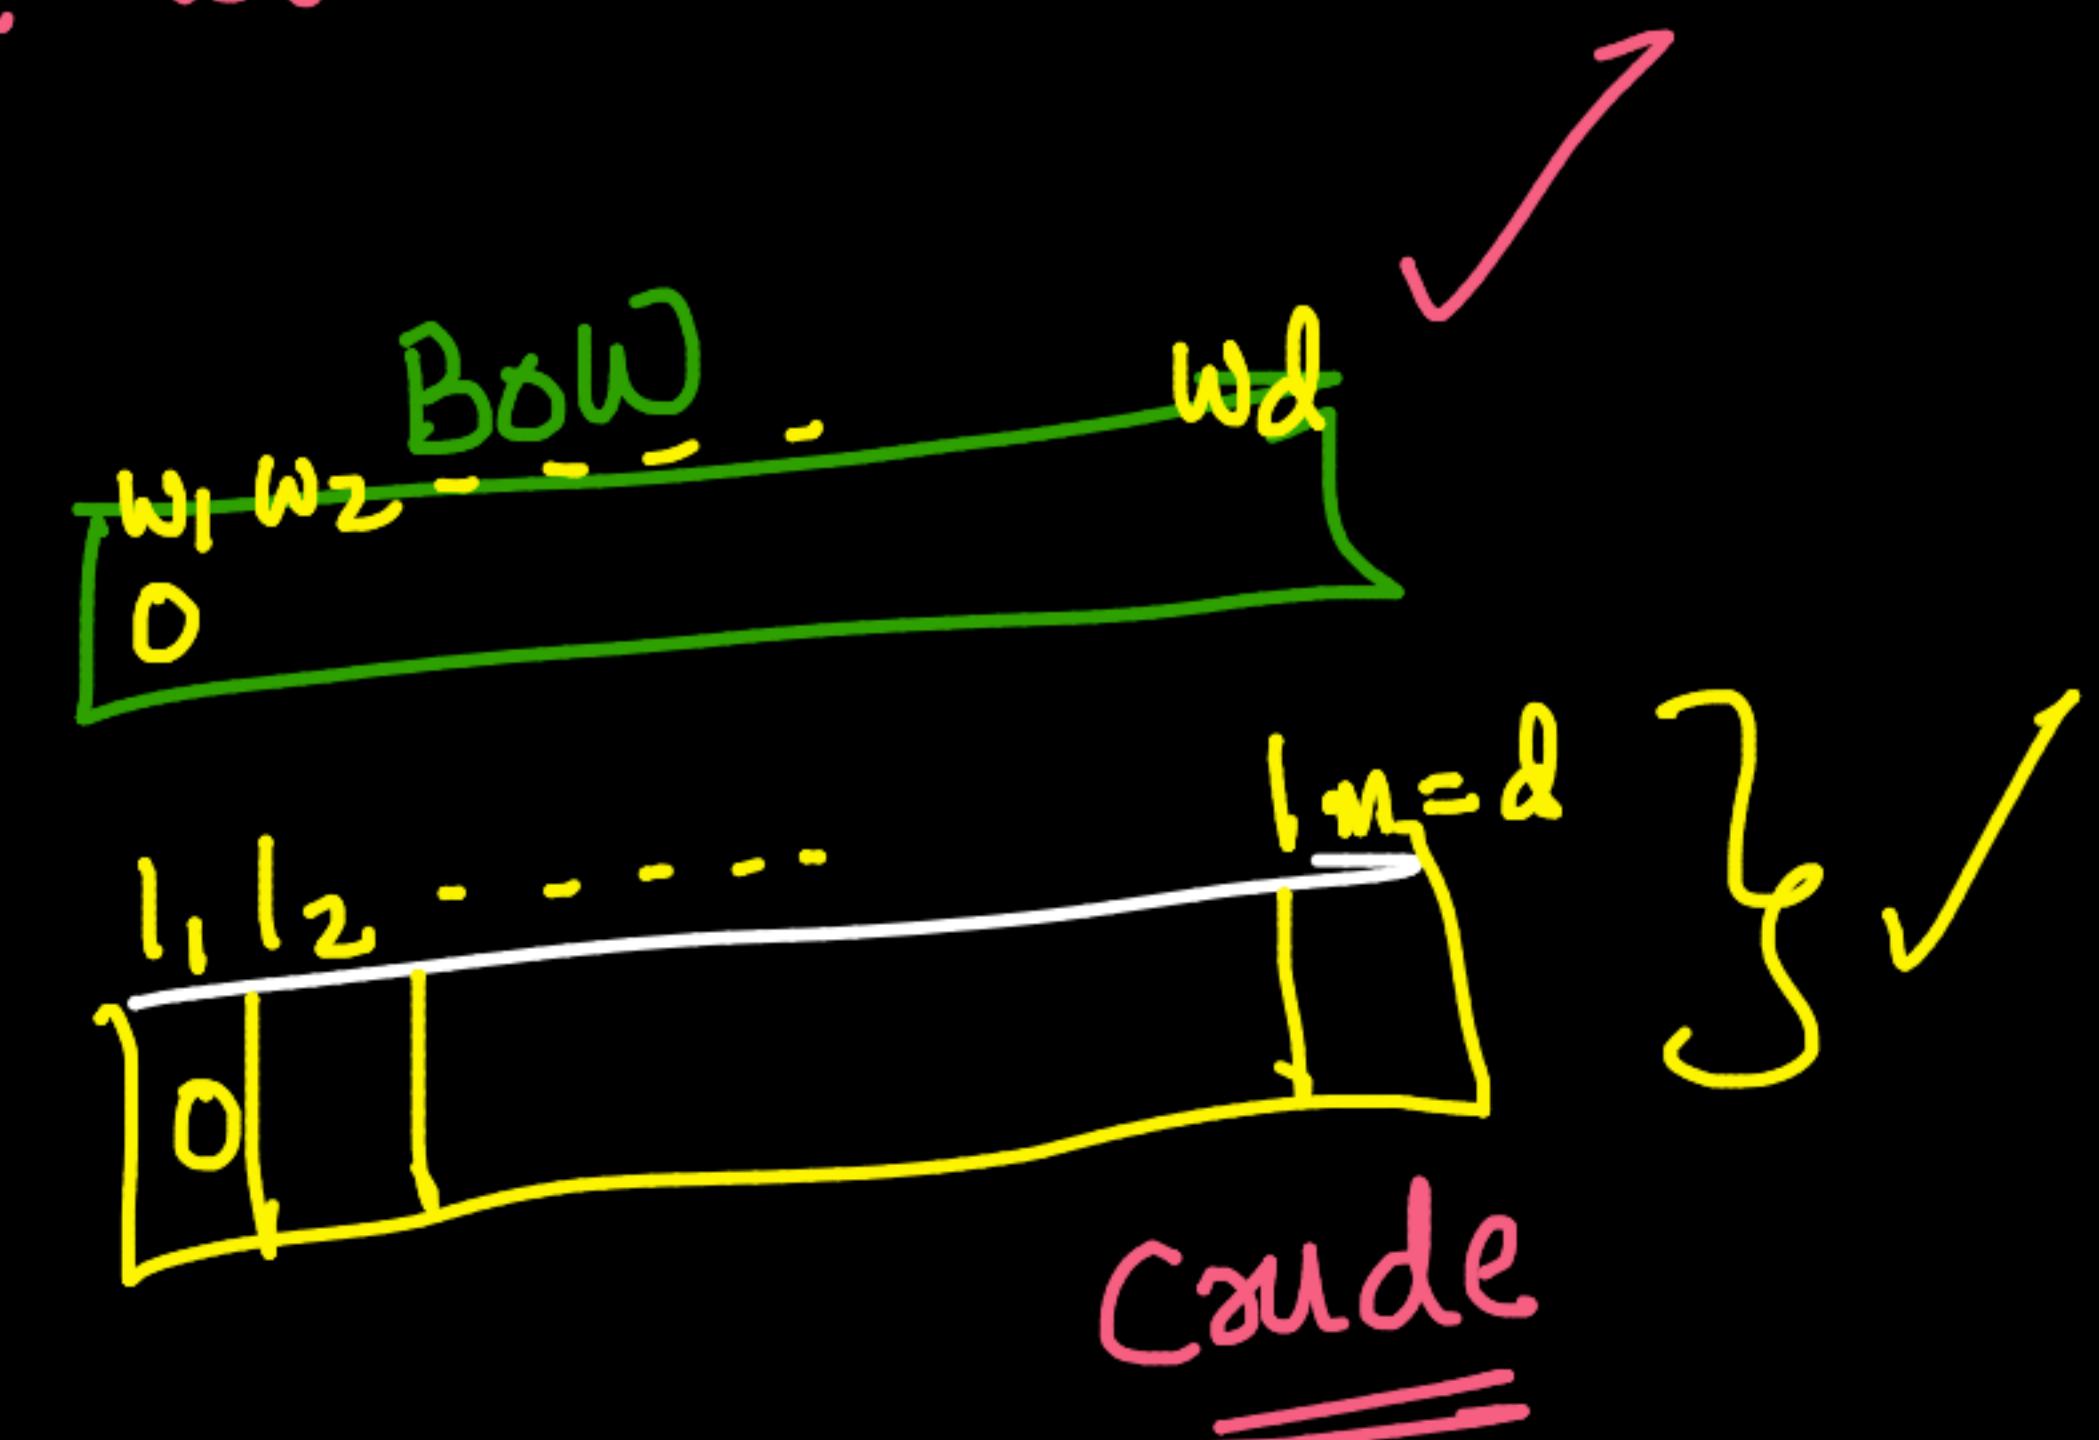
Already learnt  
technique



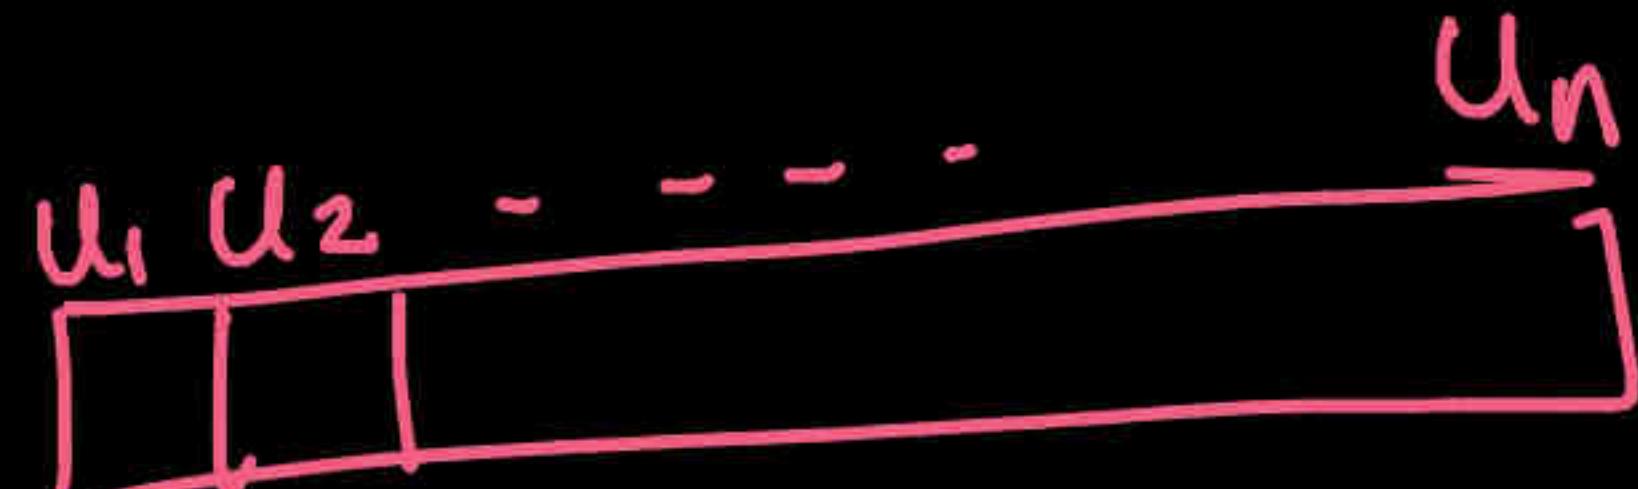
Represent  $u_i$  →

Text:

useY:



Represent  
 $I_j$

$I_j :$    
 $ER^n$

Very  
coarse

Given  $A^{n \times m}$

Task:

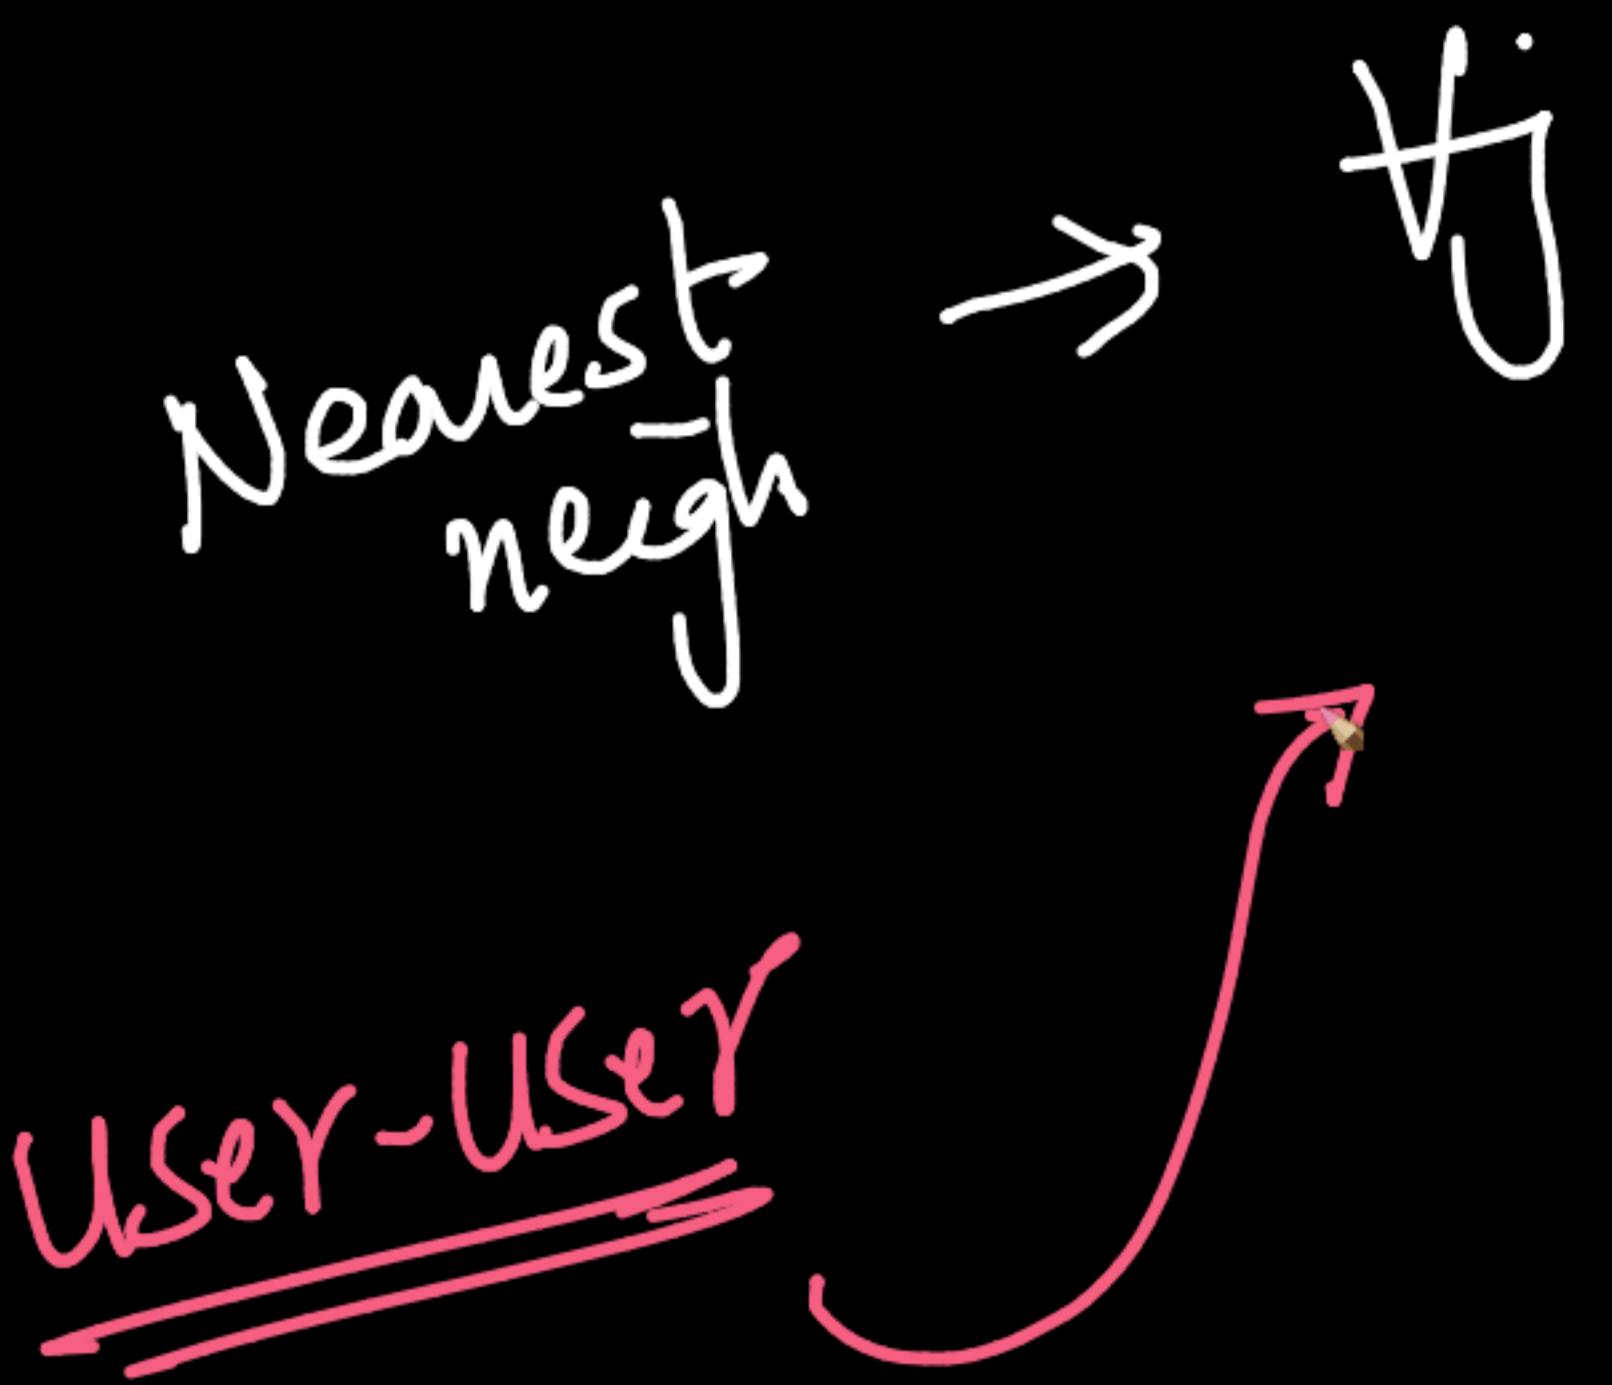
$u_i \rightarrow$  recommend some items

✓ idea1:

Item-Item Sim  
Nearest Neigh

$u_i \rightarrow I_{10}, I_{12}$  (already bought)  
 $\text{sim}(I_j, I_{10})$ ;  $\text{sim}(I_j, I_{12})$  is large  
cosine sim

idea2:

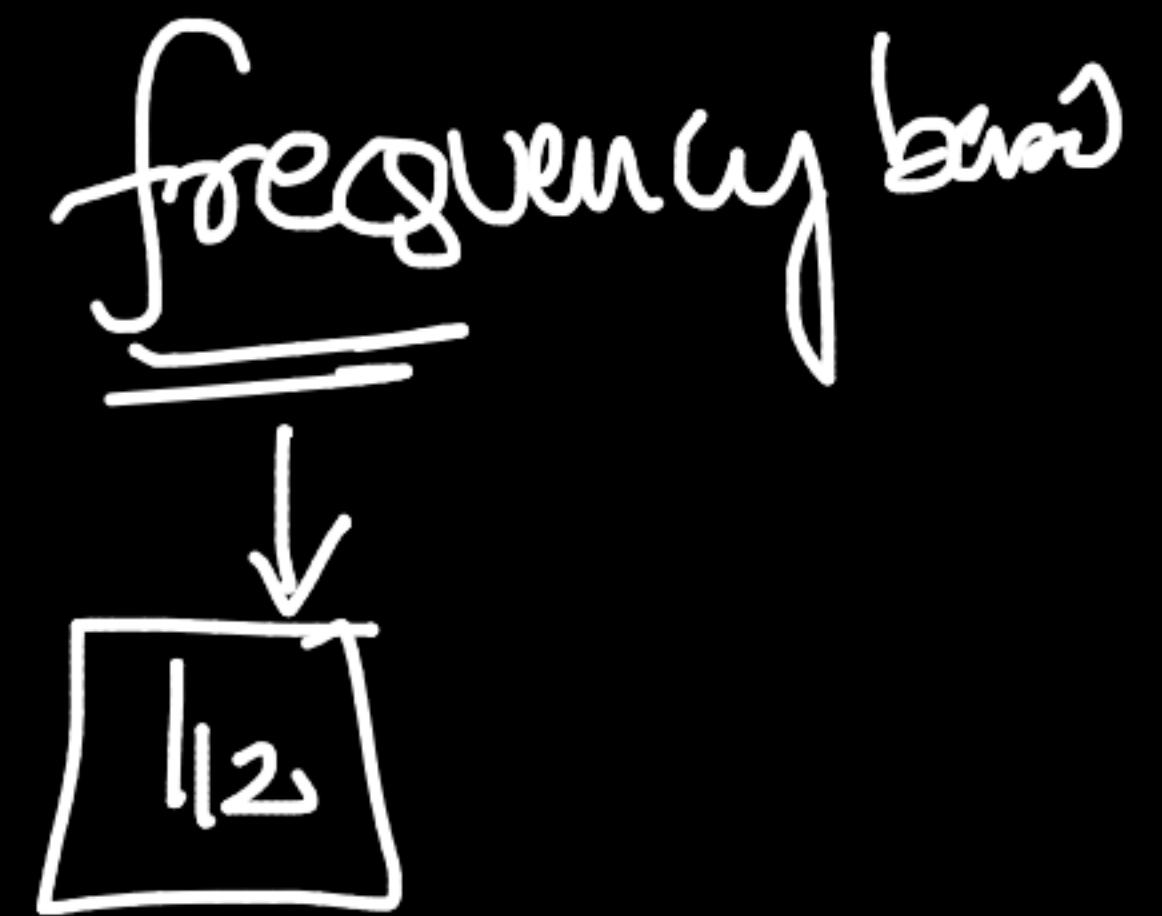


$u_i \rightarrow$  already bought  $|_{10}, |_{18}$

$\text{Sim}(u_j, u_i)$  is large

(let)  $u_{10}, u_{26}, u_{58}$  (most similar)

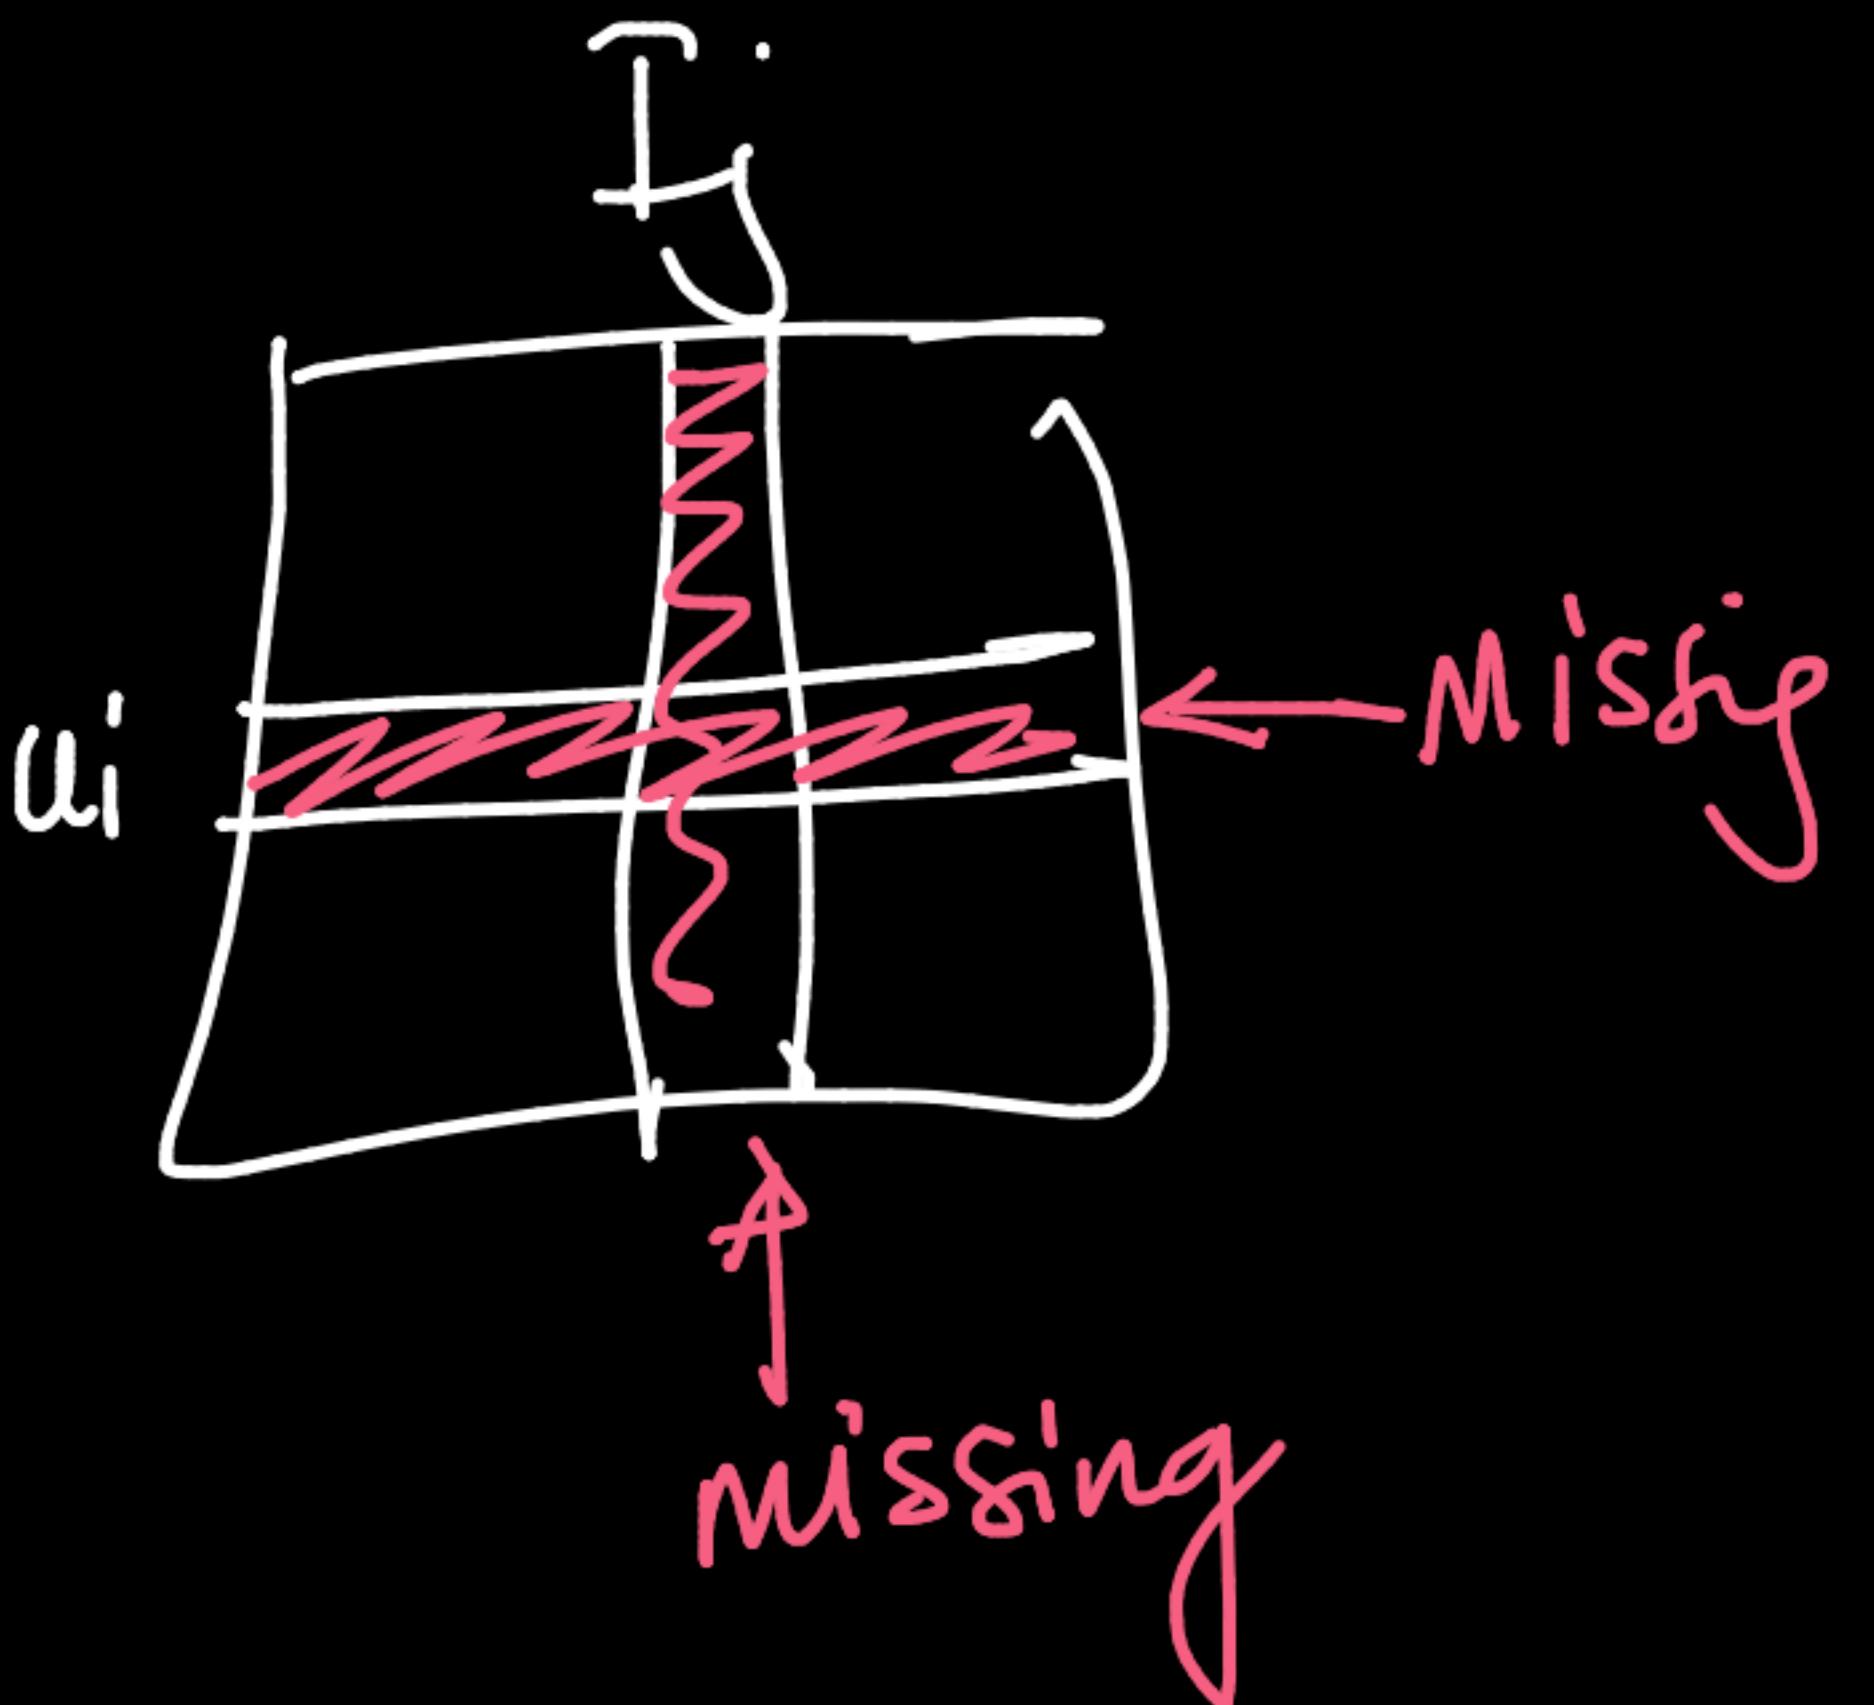
~~$|_{10}$~~        ~~$|_{10}$~~   
 $|_{12}$        $|_{18}$   
 ~~$|_{18}$~~        $|_{26}$   
 $|_{20}$        $|_{12}$



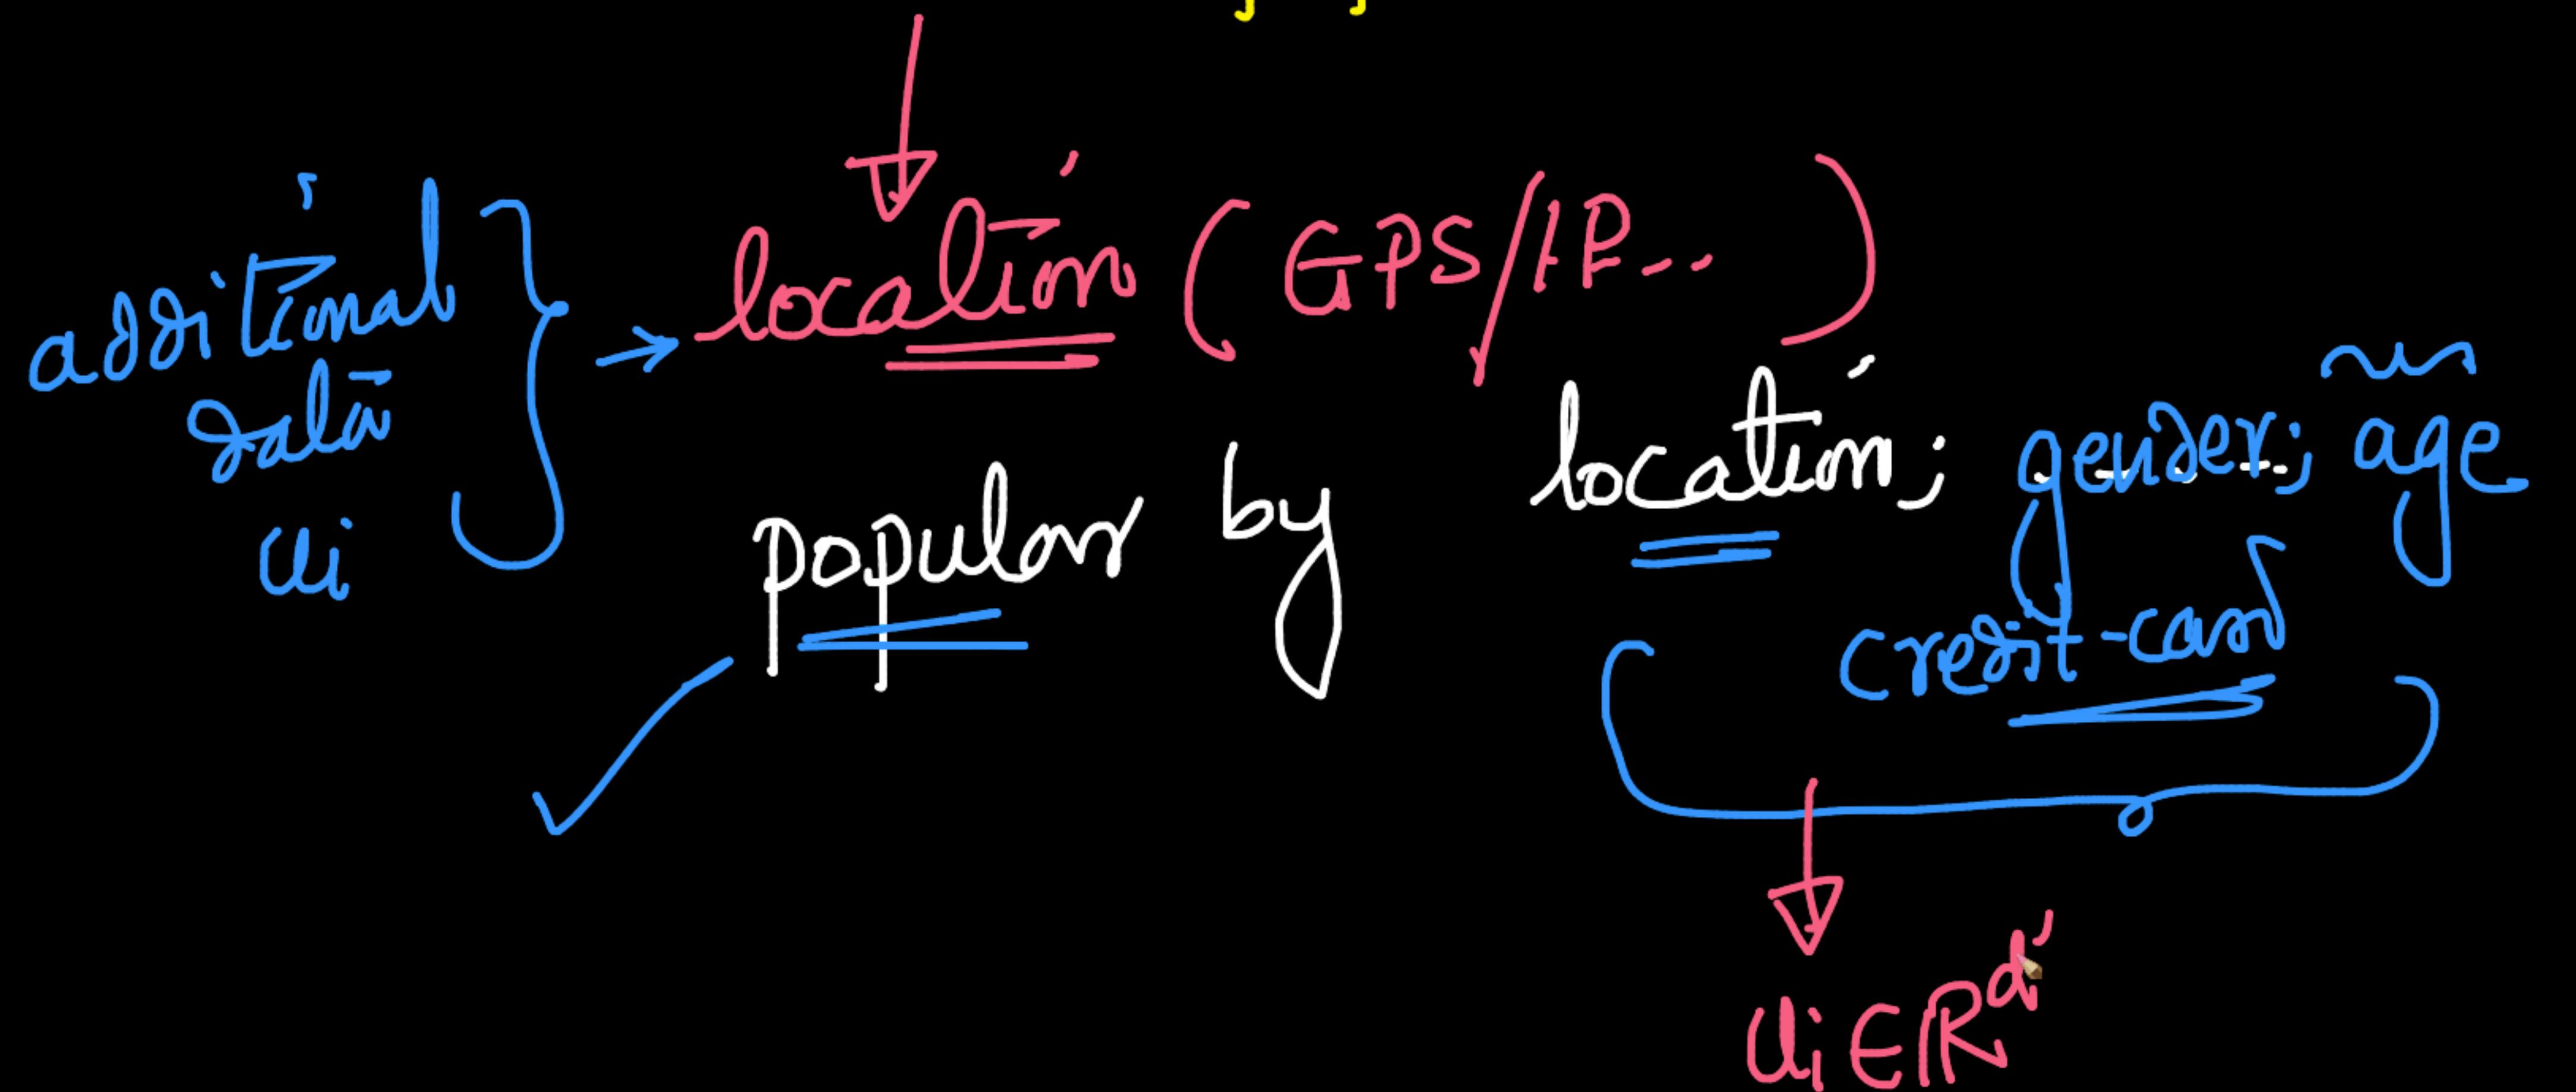
→ recommended items

$u_i$ : new-user

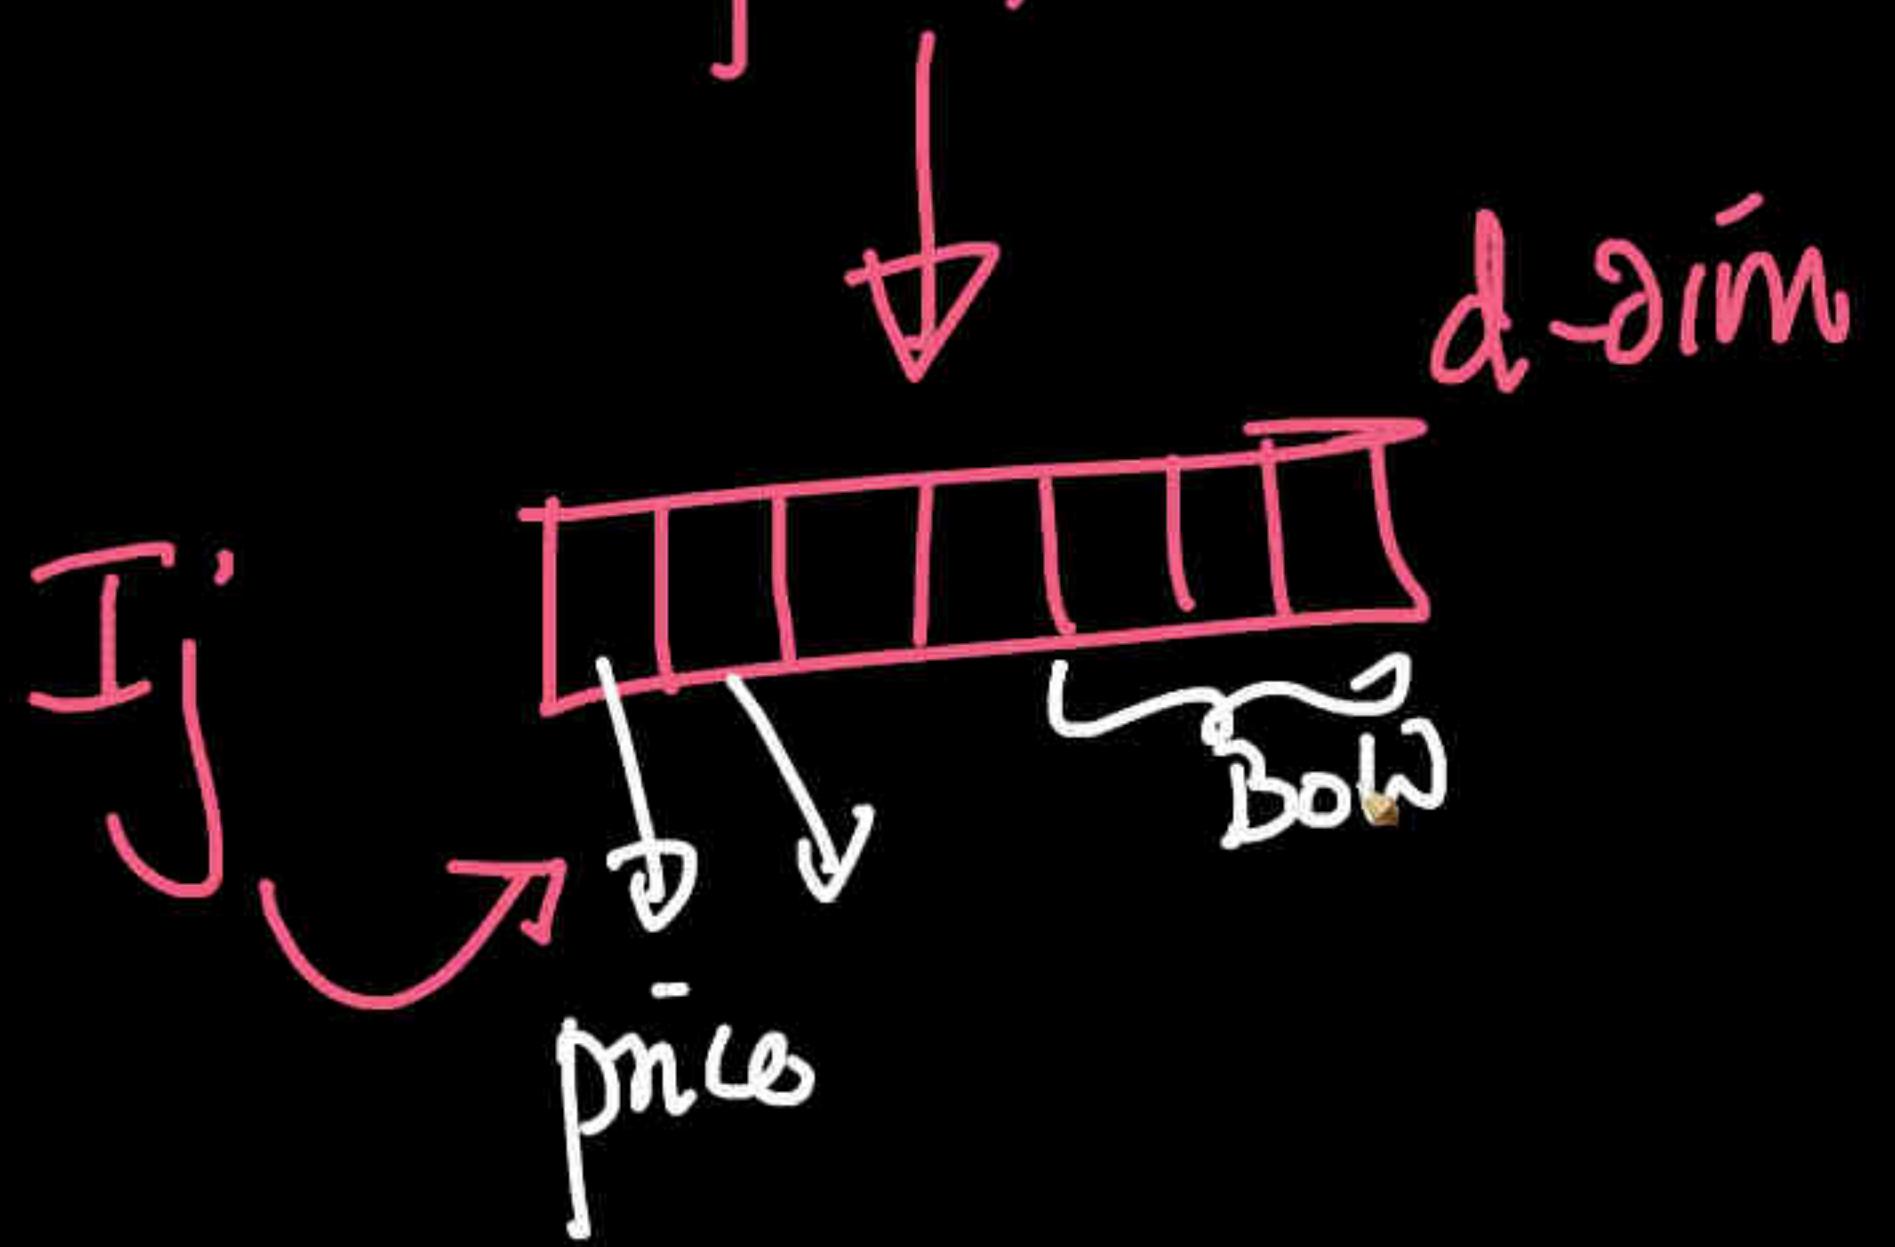
$I_j$ : new-item

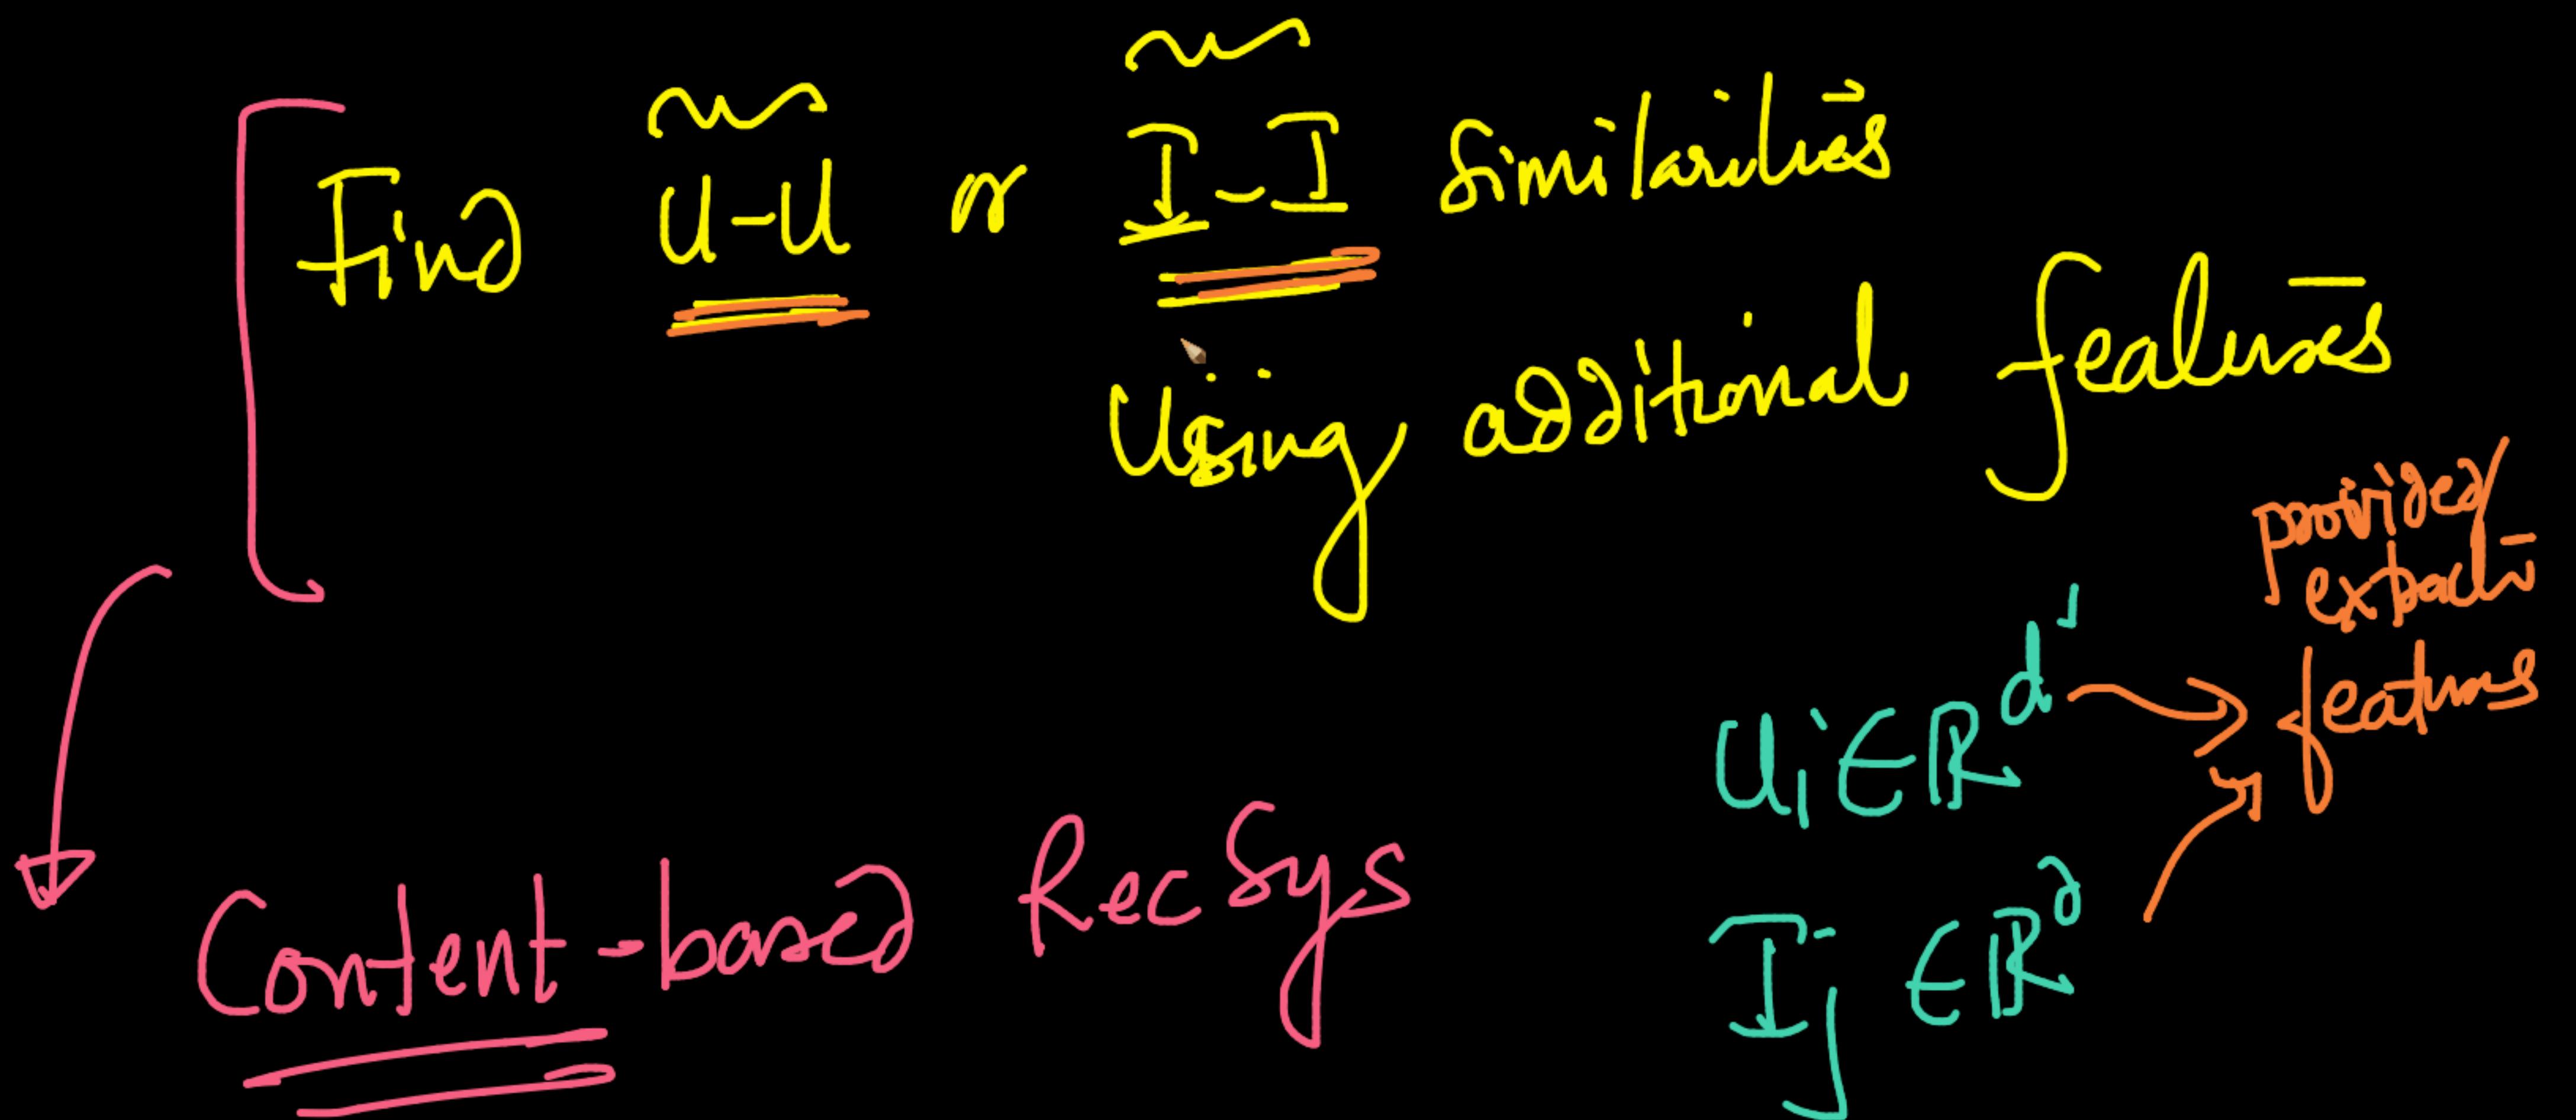


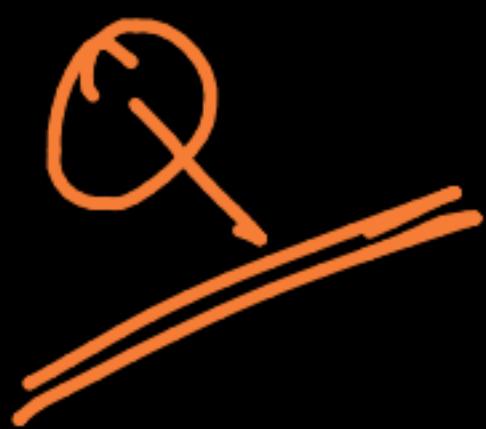
idea 1: frequently bought Items  $\xleftarrow{\text{(n) popular}}$



Item<sub>j</sub> is new → seller provided info

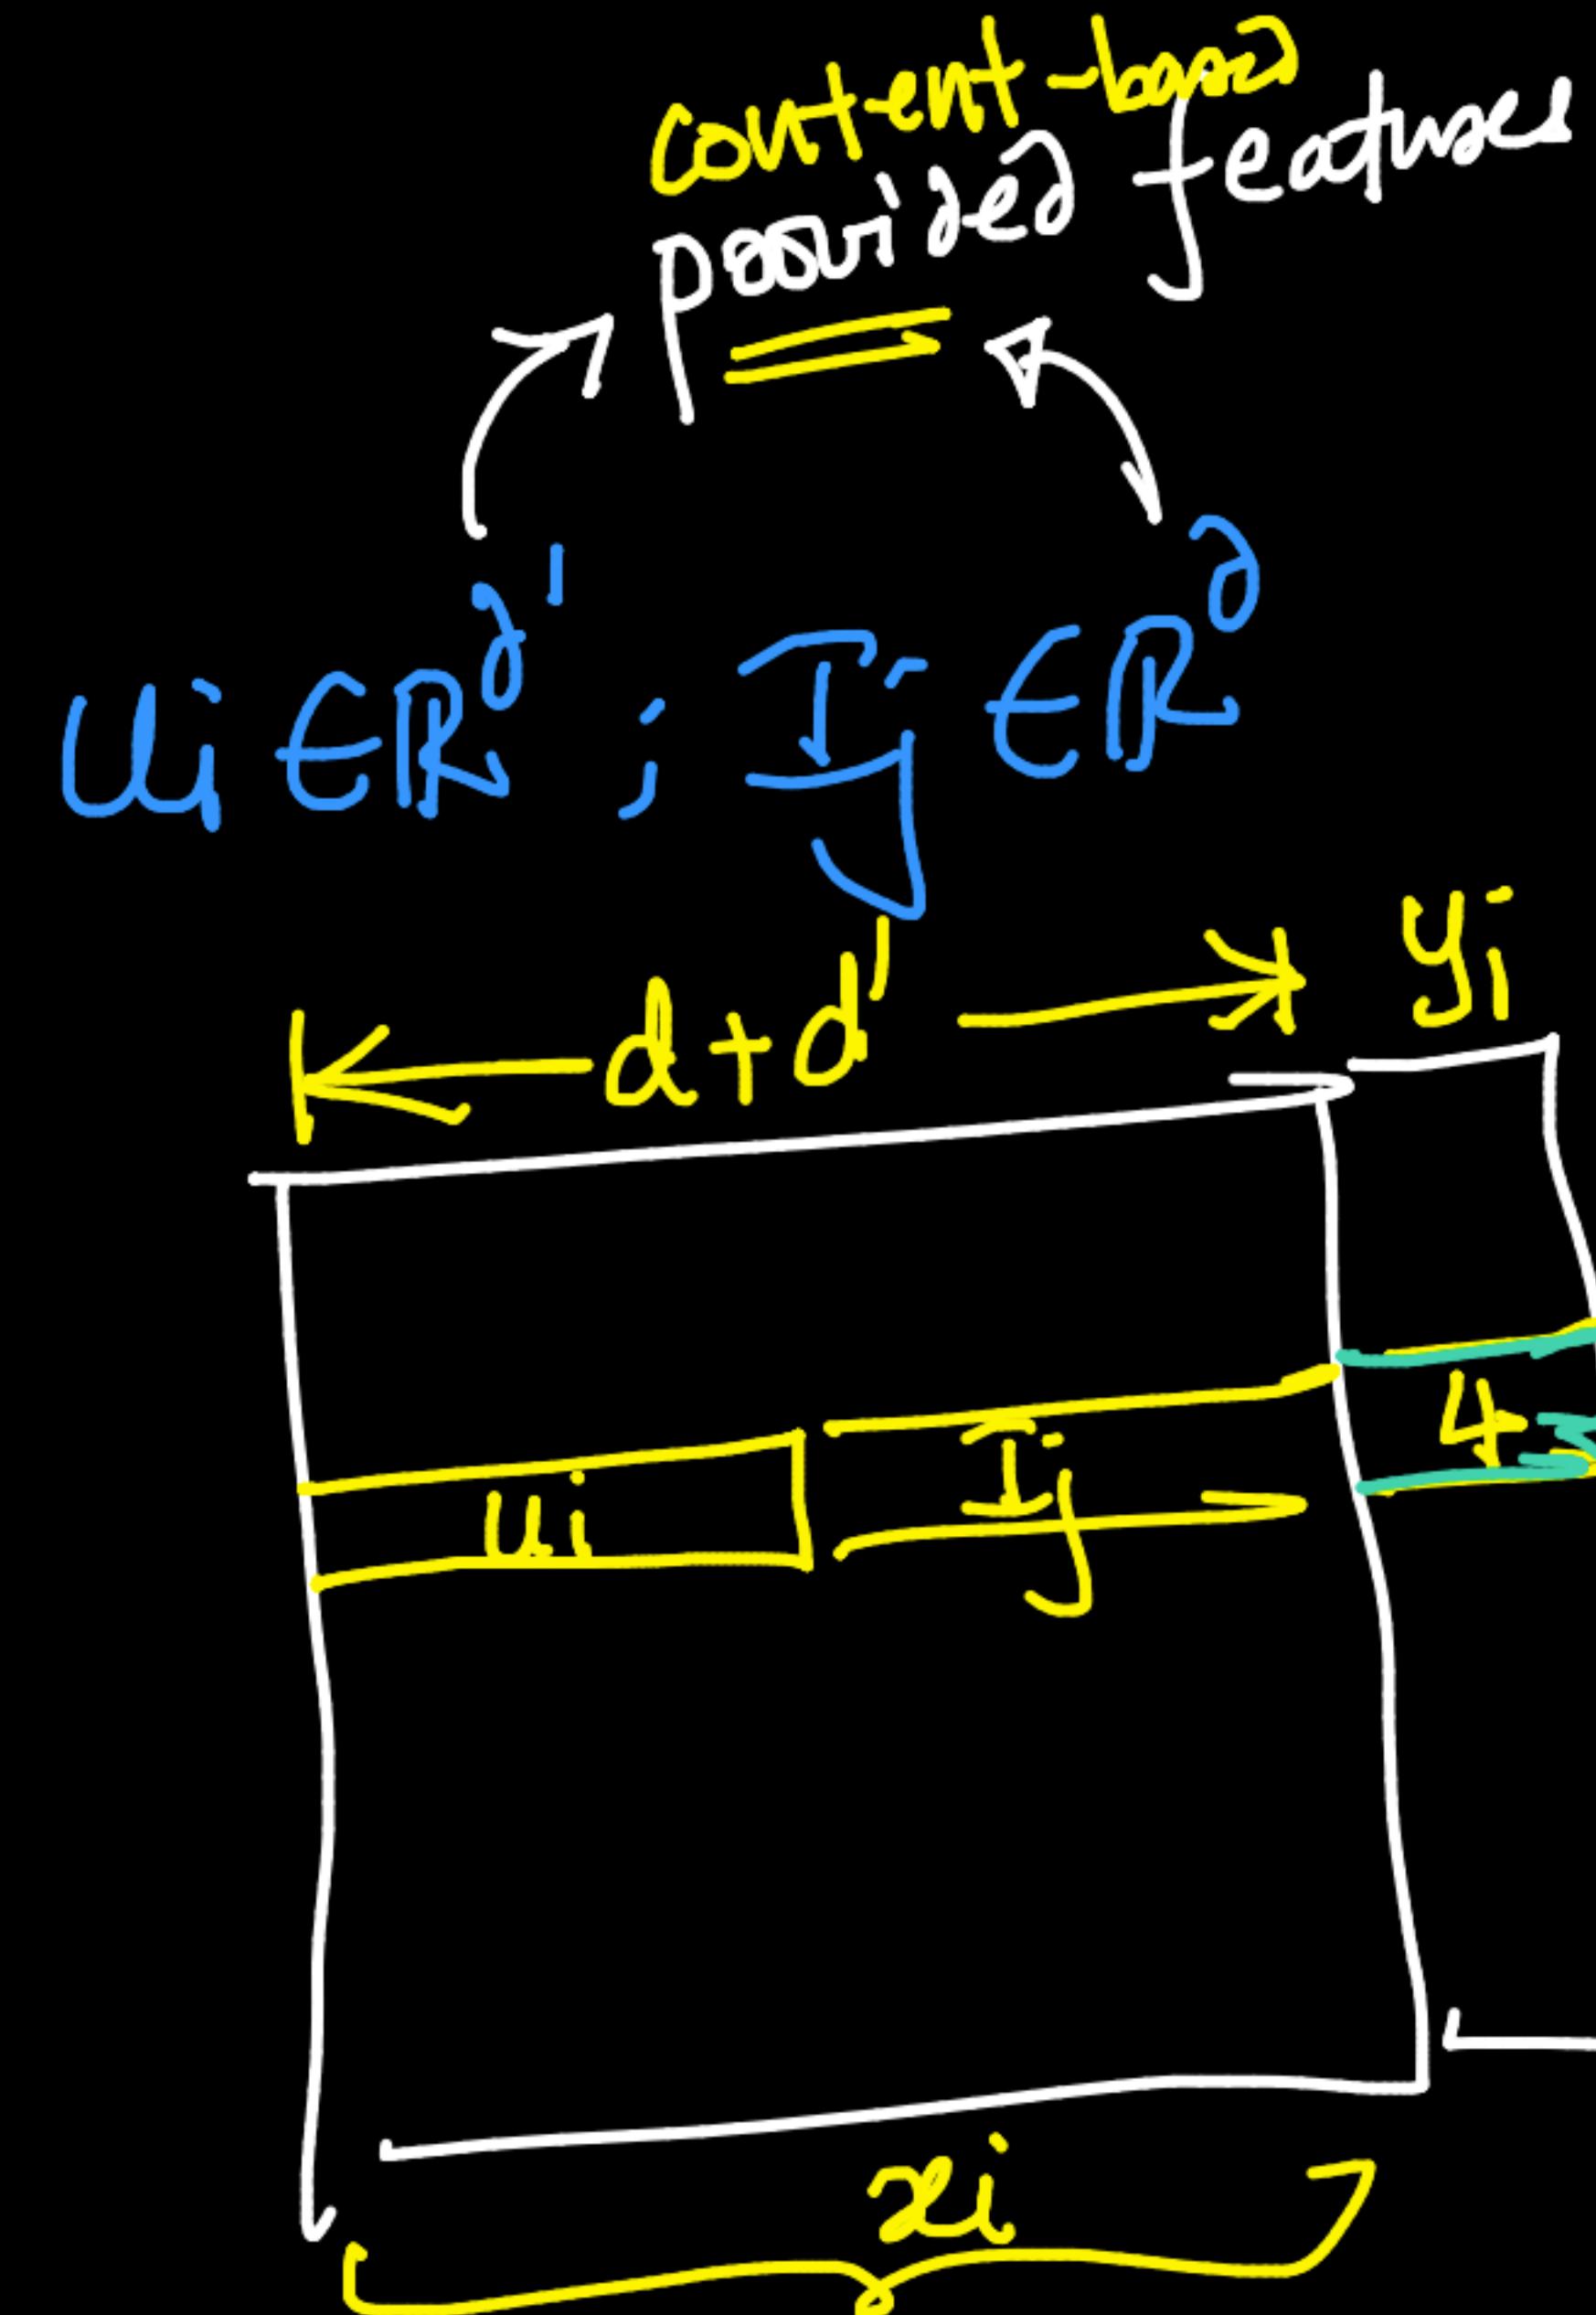






Idea'

$u_{1000}, I_{111}, y_i$

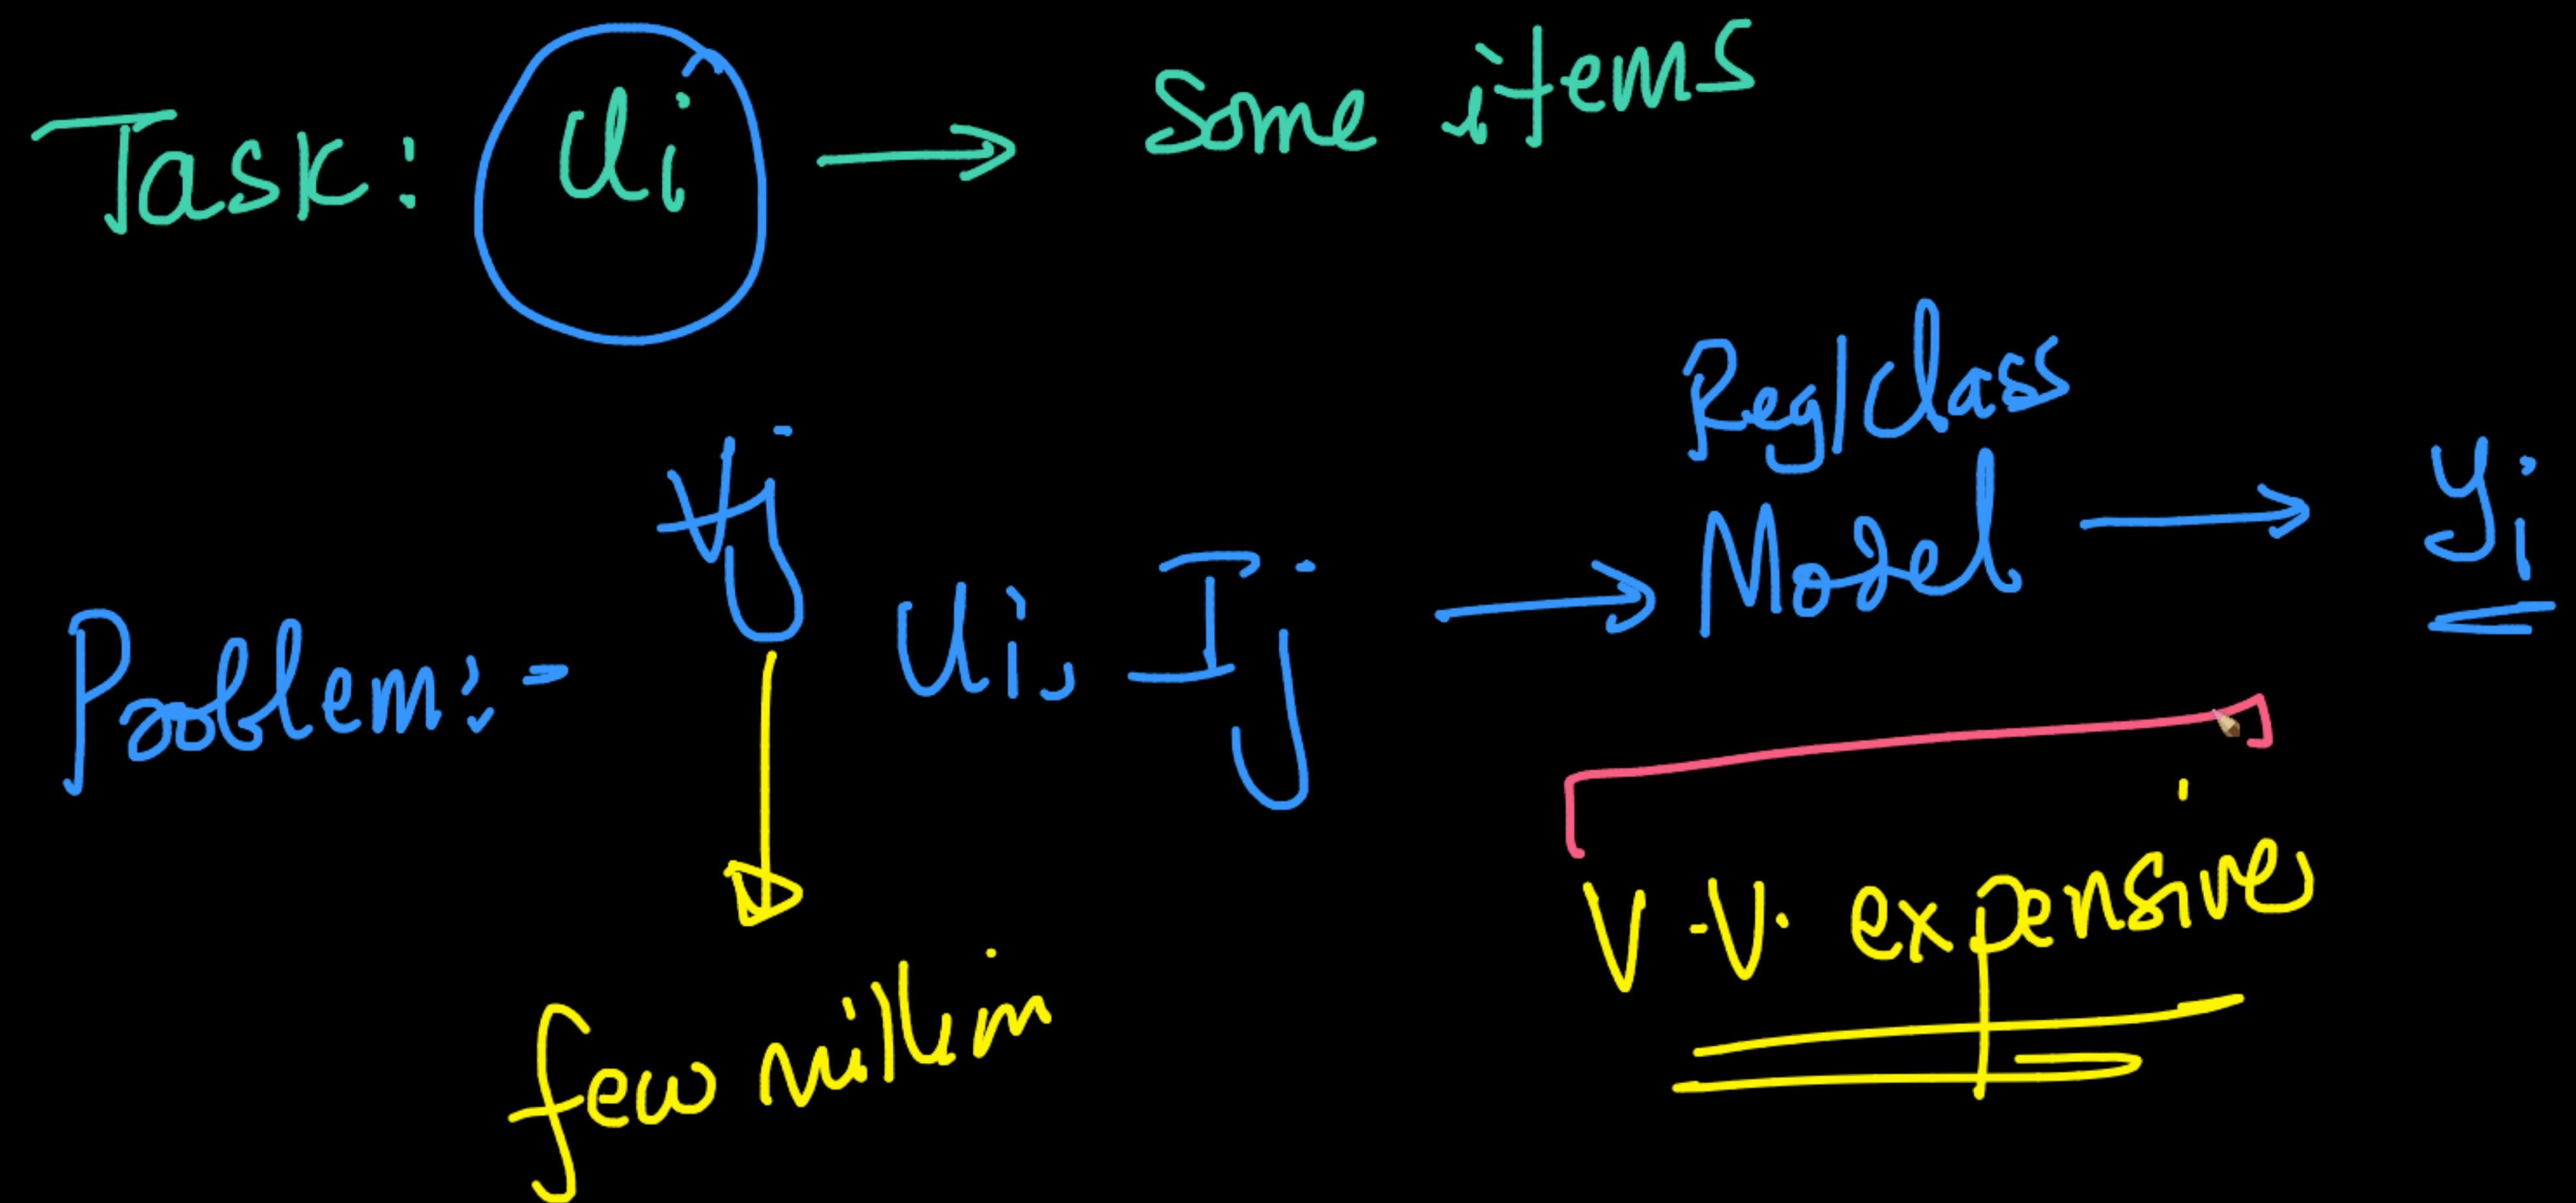


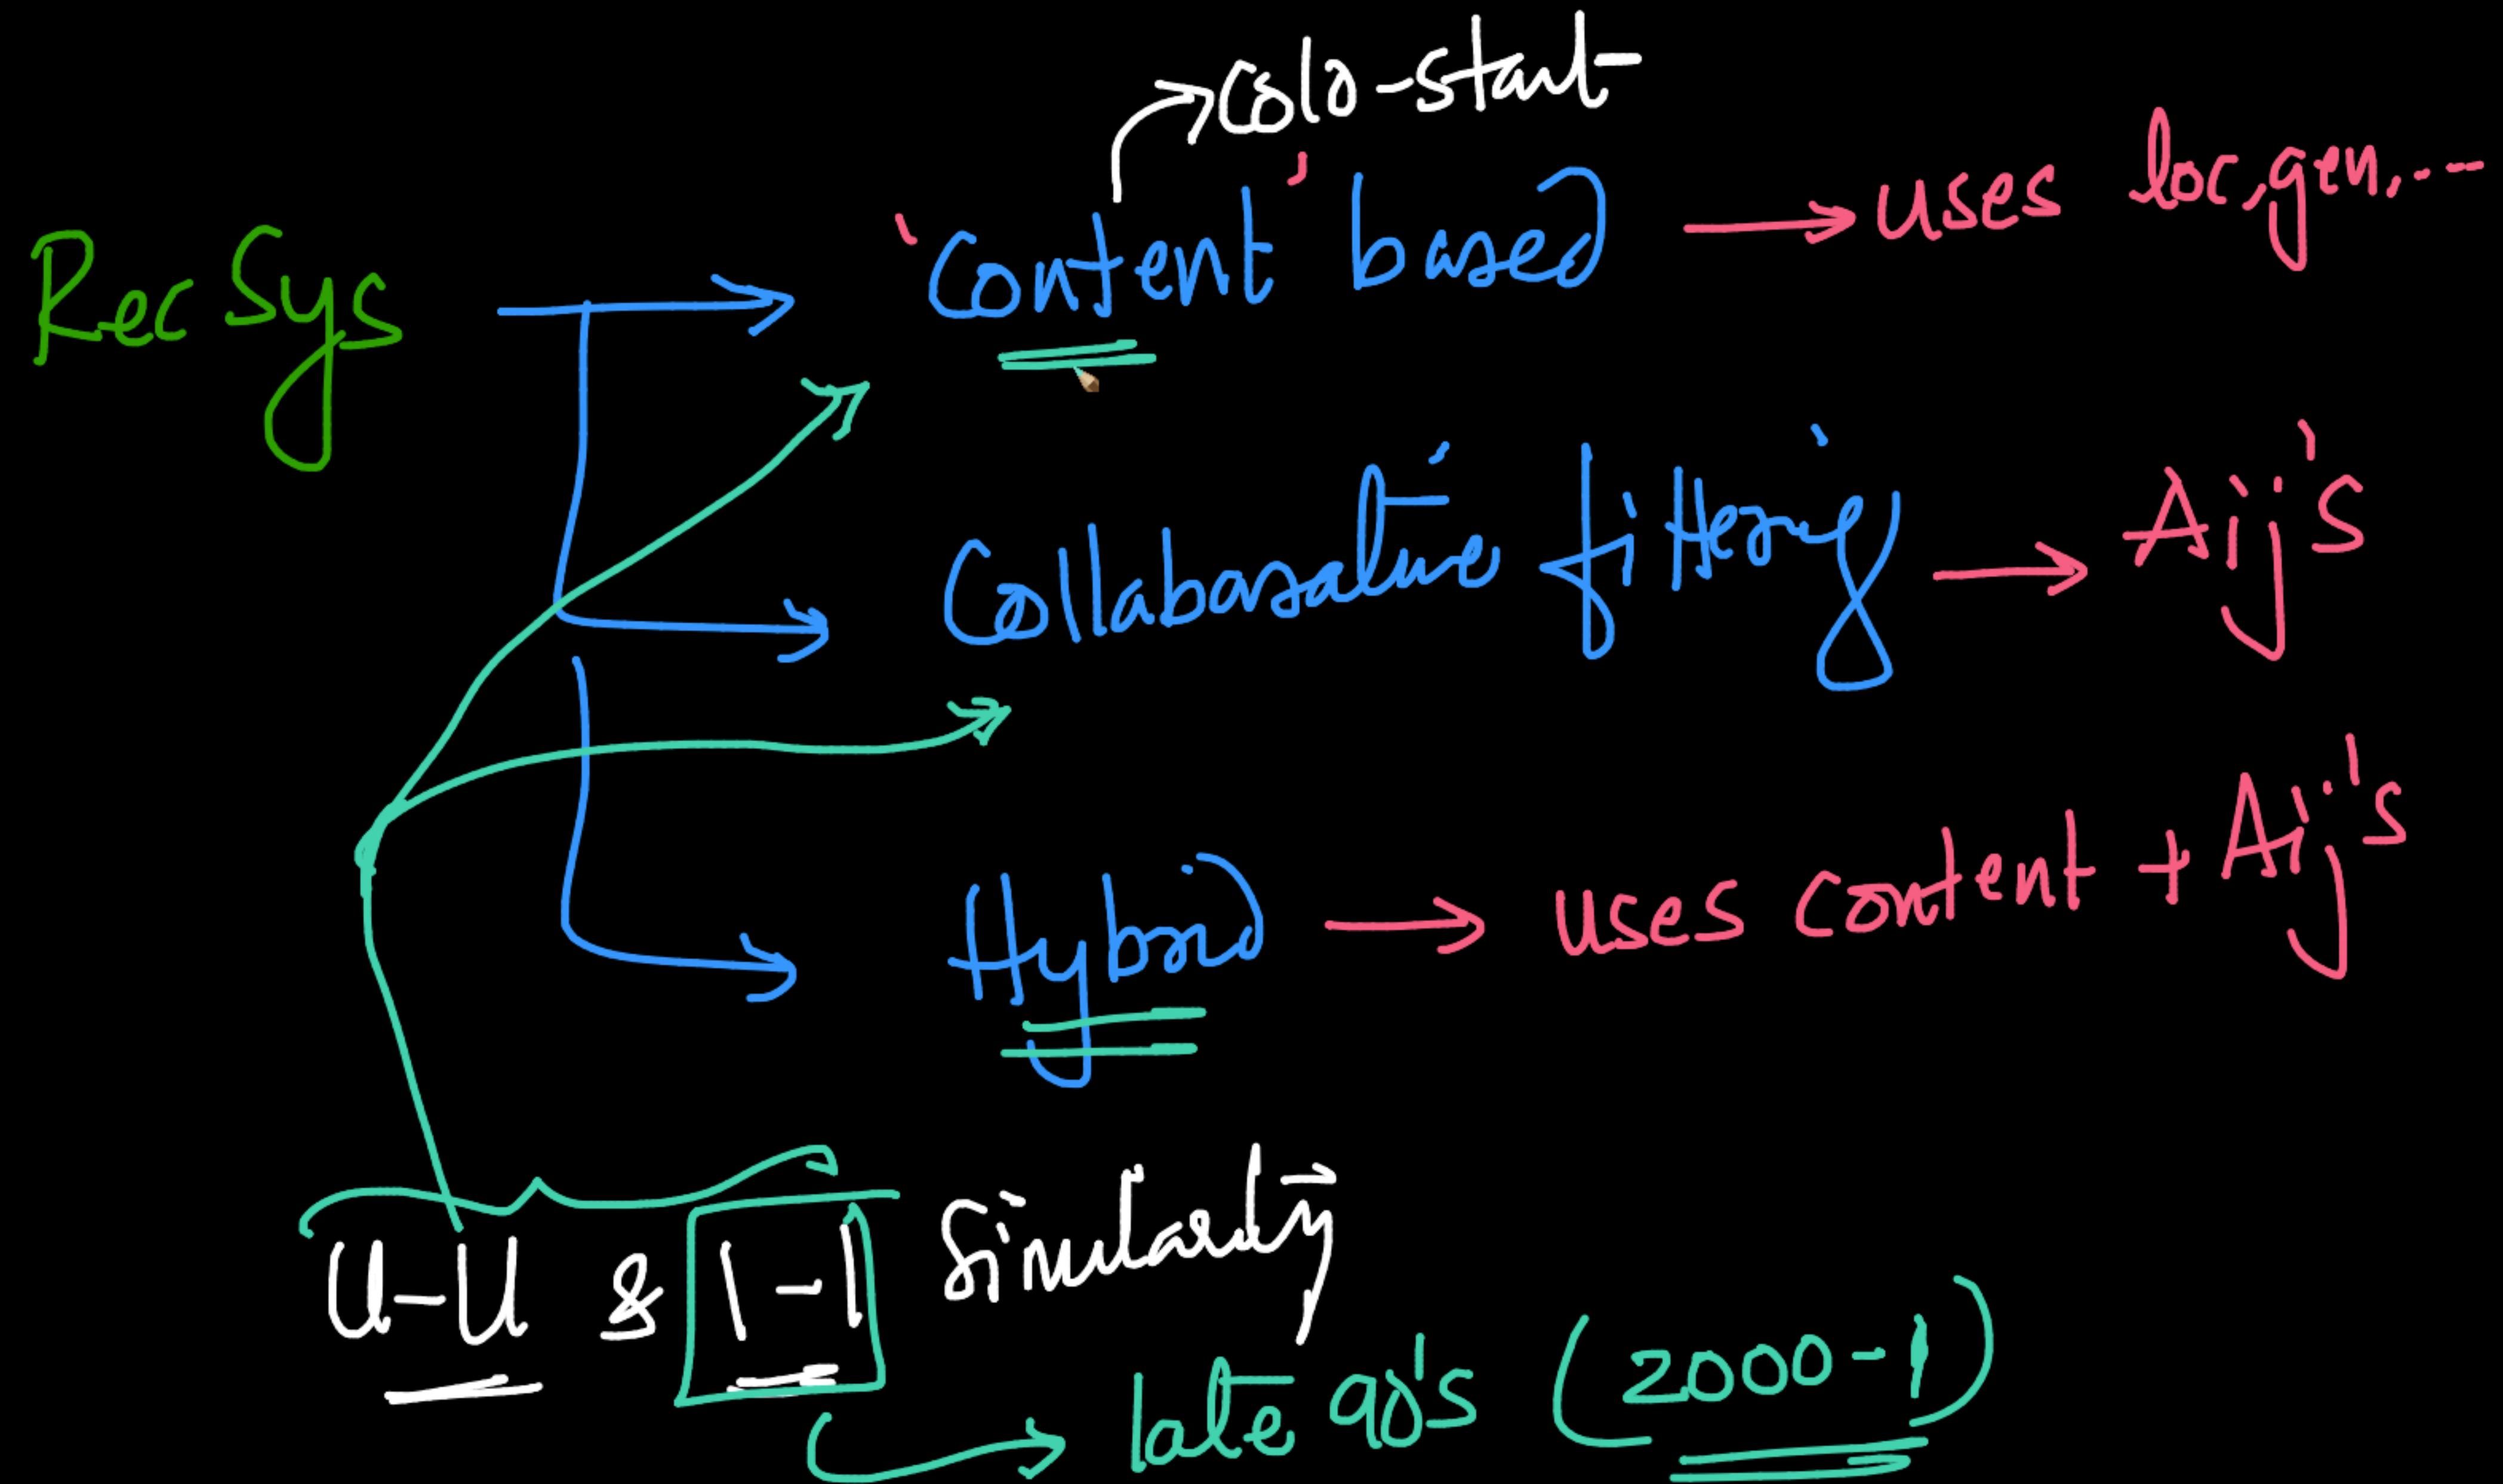
Reqd / classfn

list-data

$A_{ij} = 4$

$u_i, I_j \rightarrow 4$







item-item CF



I<sub>i</sub>, I<sub>j</sub>

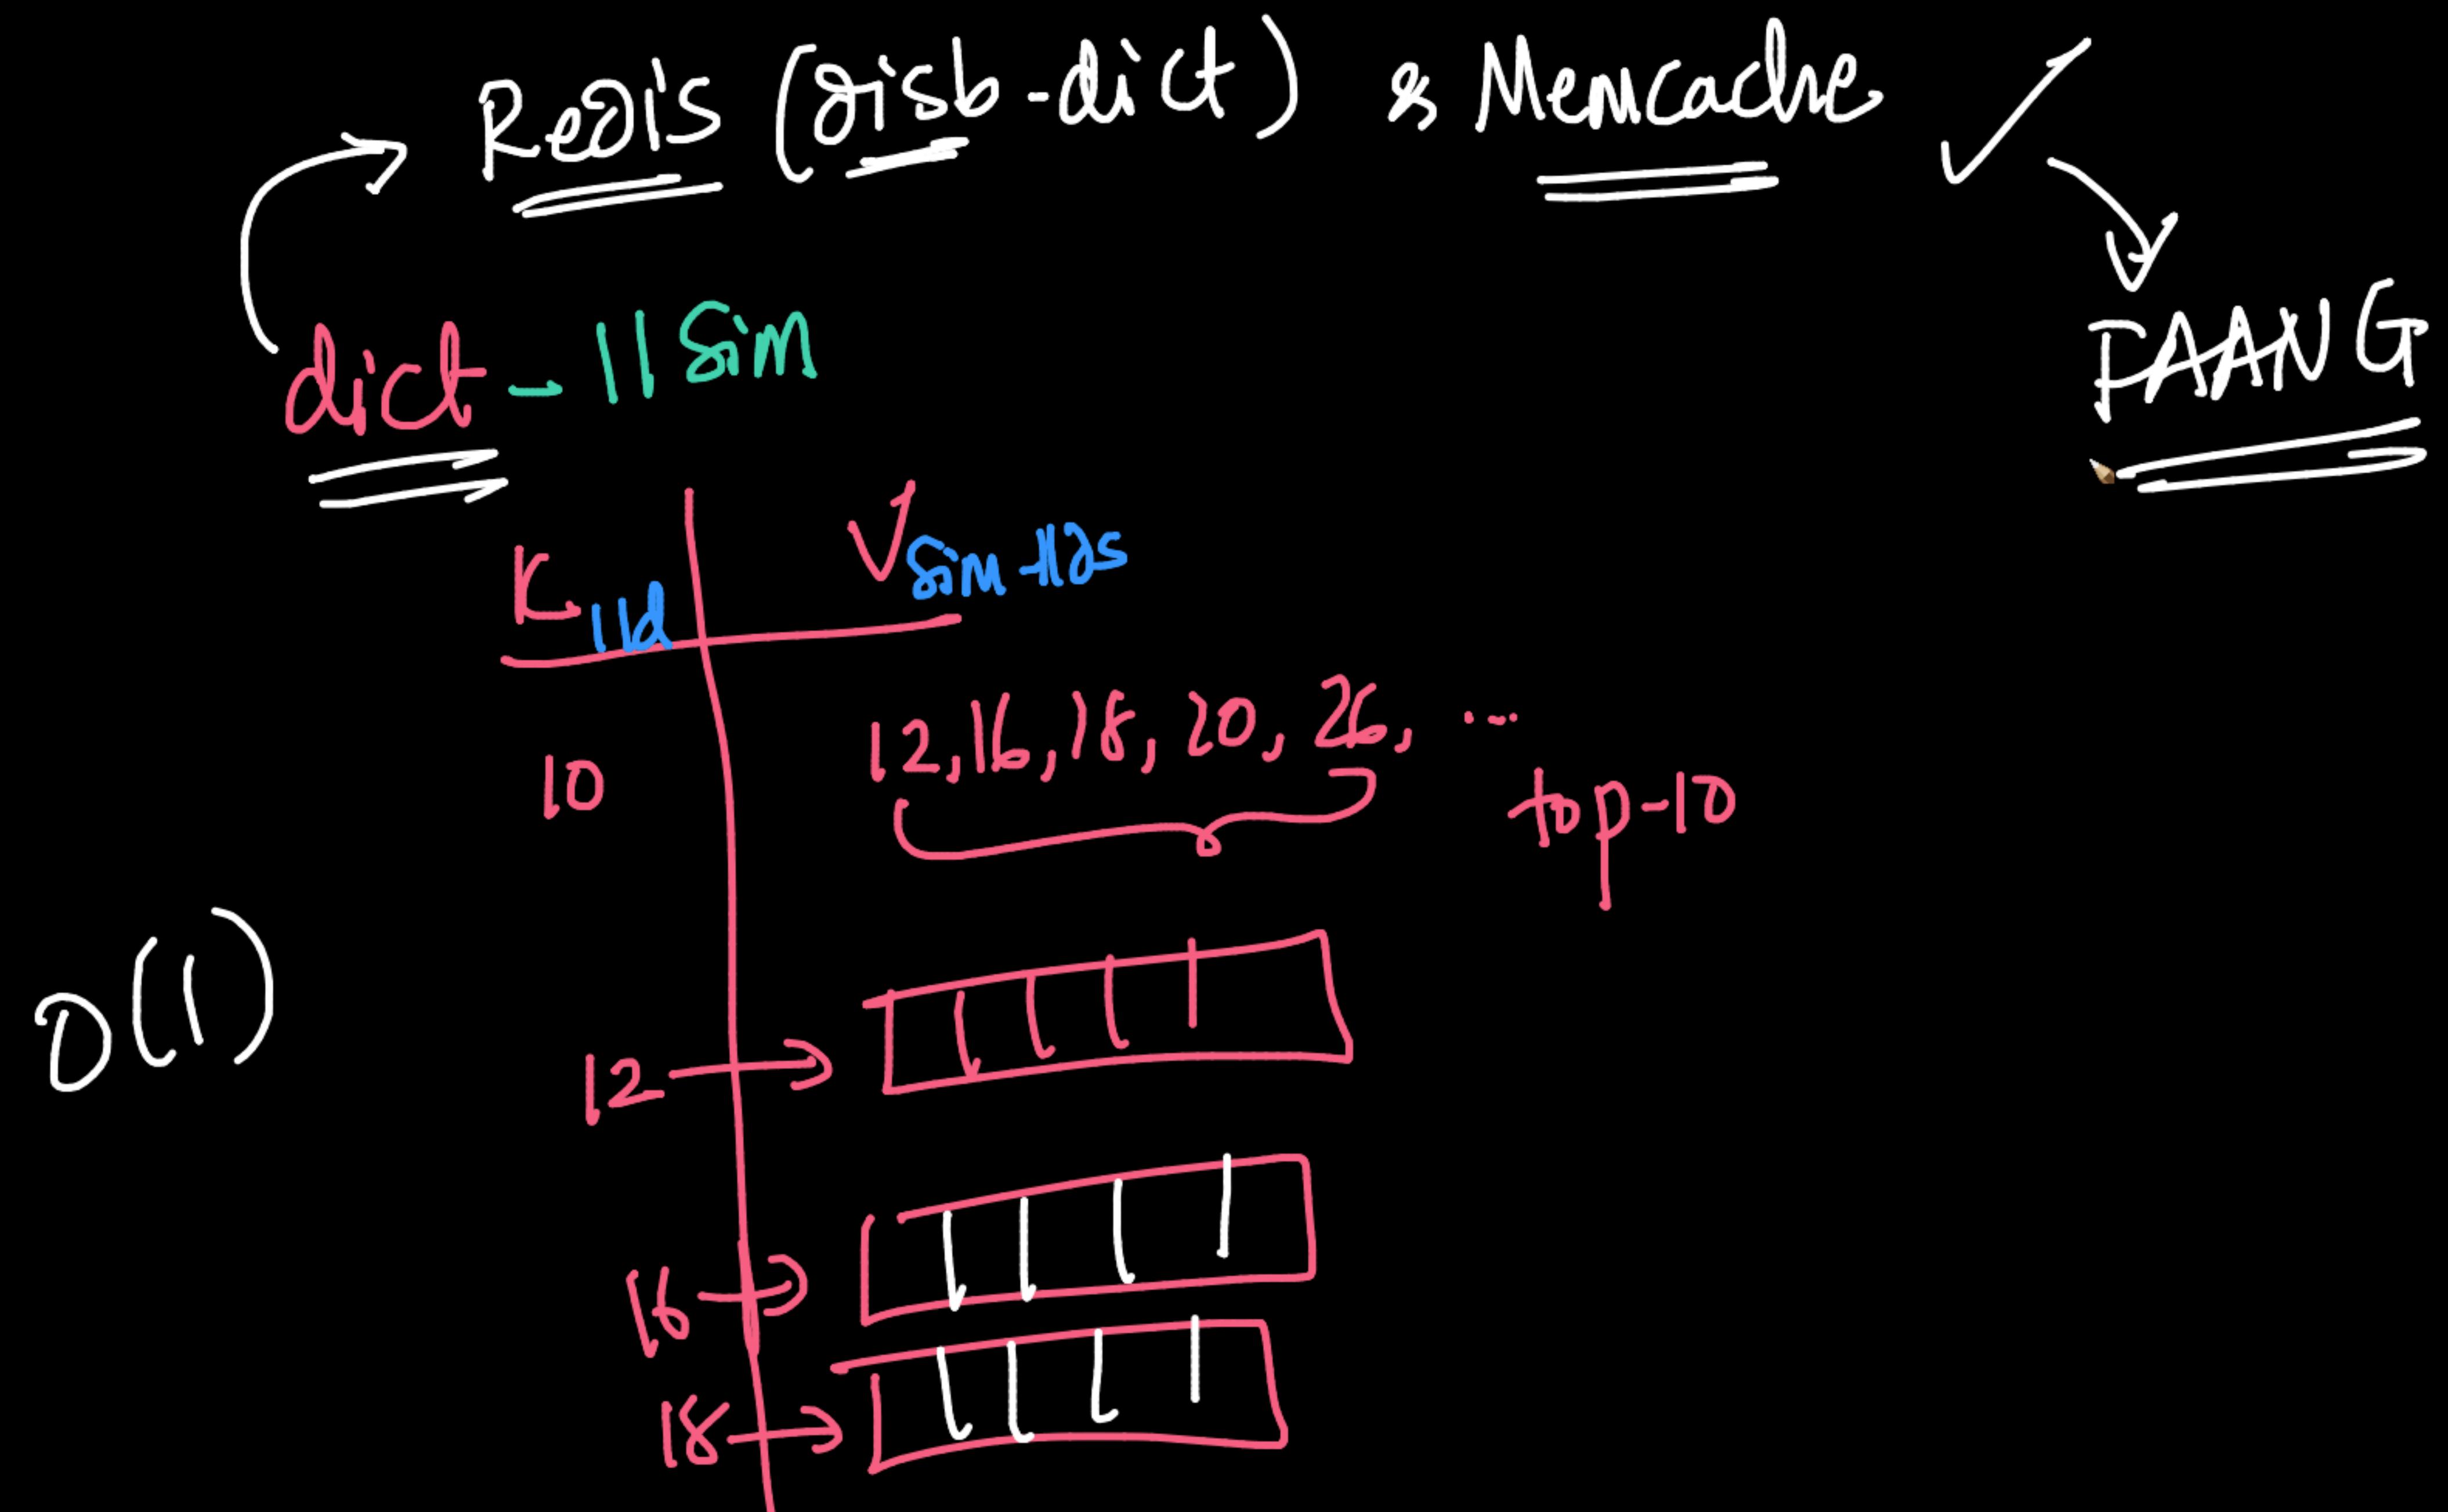
... - j ->

U<sub>i</sub> → ...

<100 ms

Sim - II =  $\begin{bmatrix} & \\ & \end{bmatrix}$

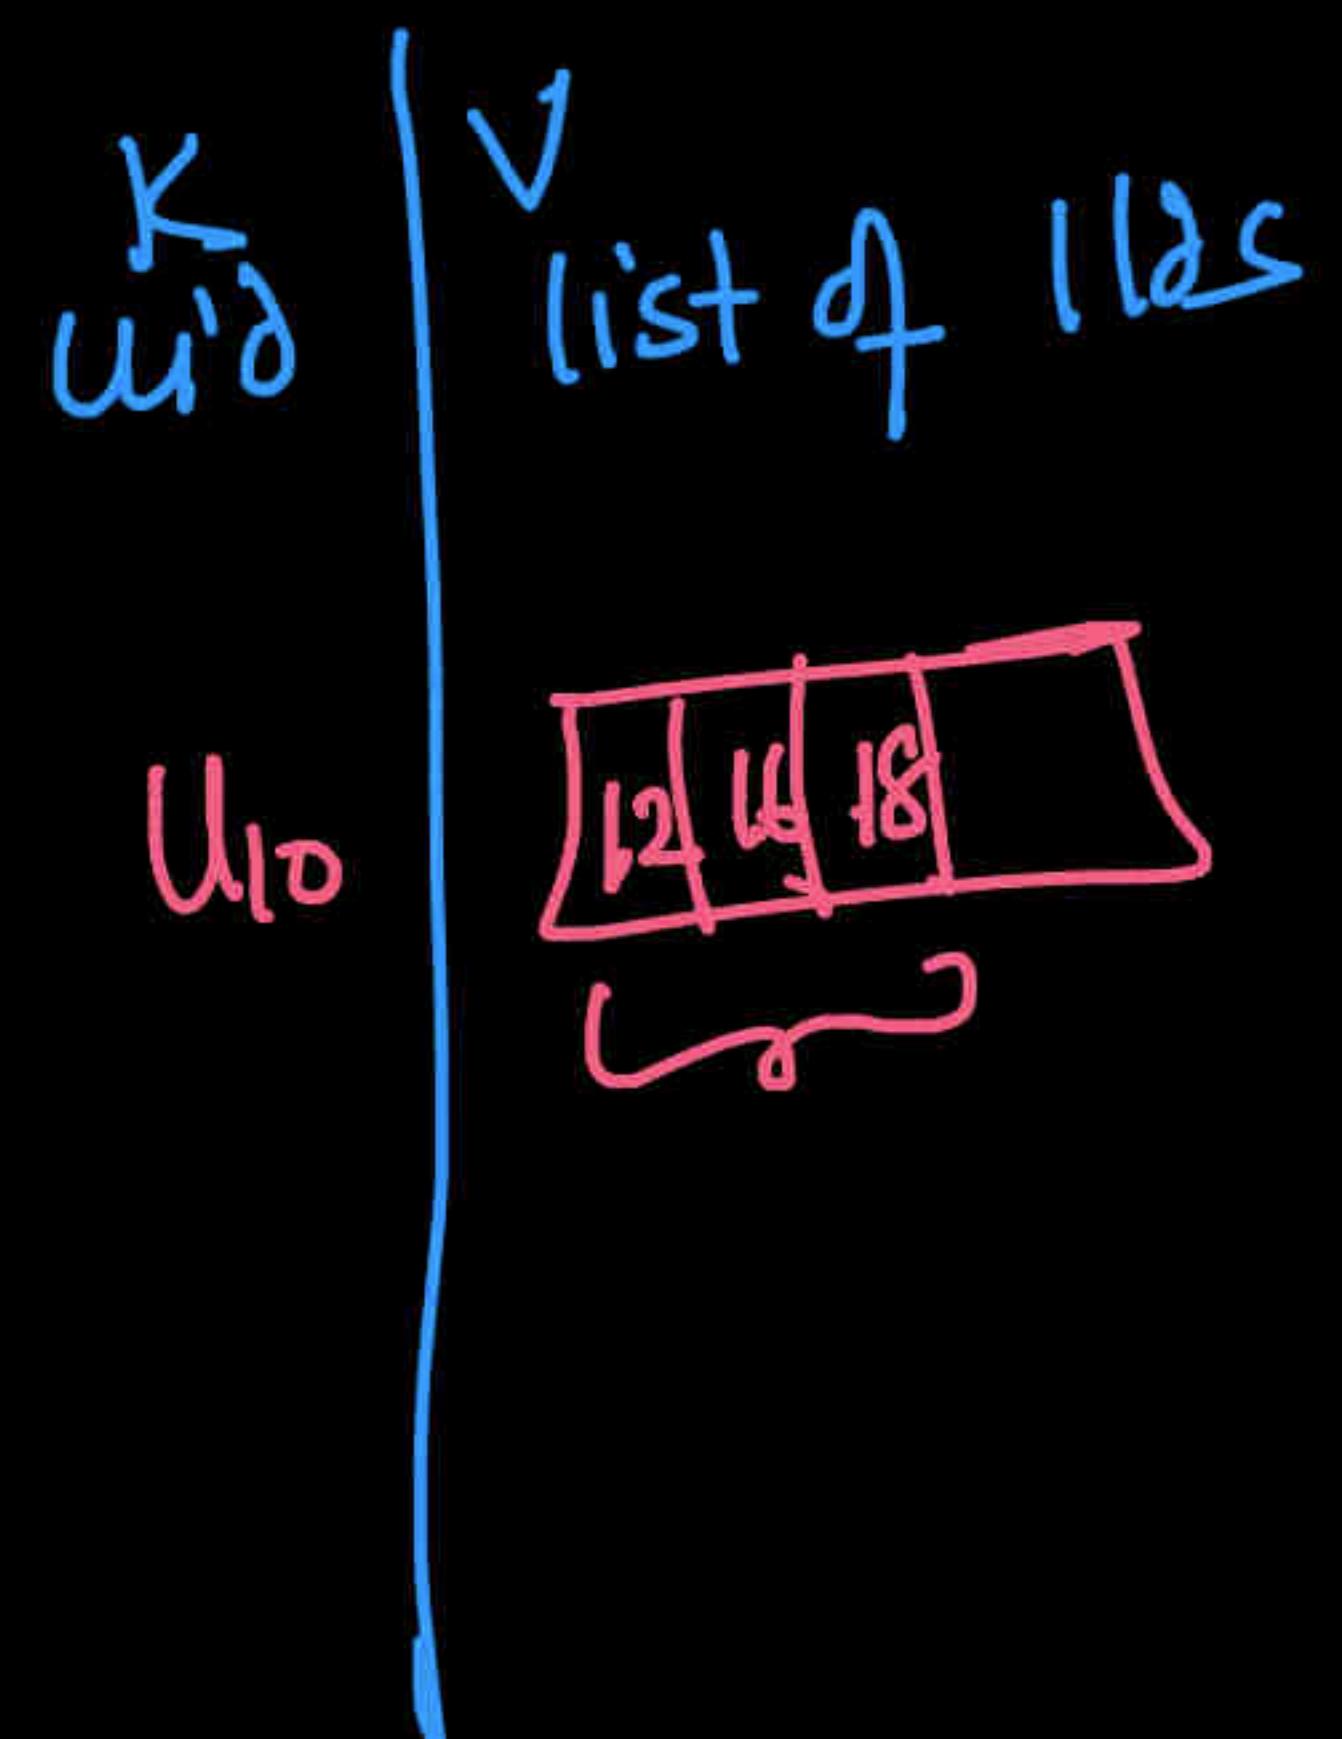
nightly / weekly

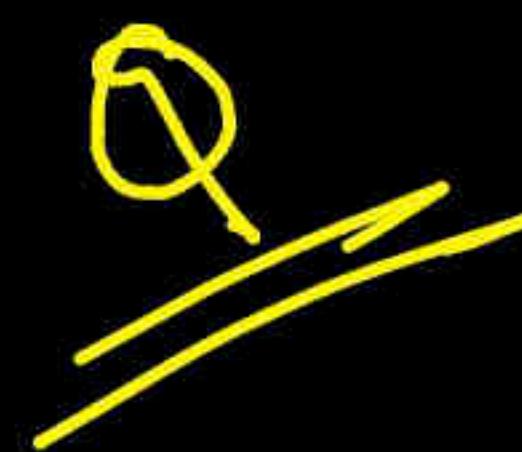


dict - U-1

$u_{10}$

$v(i)$





Netfli**x**

YouTube IN

Search

Live chat replay was turned off for this video.

All Republic Day Indian Army Troops >

Indira Gandhi & Treasure of Jaigarh Fort What is Indira Gandhi's connection with Jaigarh Fort's... StudyIQ IAS 2.9M views • 1 year ago By Aadesh Singh 15:25

जलवायु परिवर्तन अब वाद-विवाद का विषय नहीं रह गया है President of India 12K views • 13 days ago 0:46

Maa Tujhe Salaam | Daredevil Bike Stunts | Stunning Air... Music India Entertainment 11M views • 2 years ago 6:36

#EkThumkeSe #LoveYouLoktantra ... Zee Music Company 28K views • 13 hours ago New 0:16

Kargil: Valour & Victory Hindi... HISTORY TV18 468K views • 10 days ago 24:16

Namo Namo - Lyrical | Kedarnath | Sushant Rajput ... Zee Music Company 358M views • 2 years ago 5:29

The Pursuit Of Happyness: Job Interview Binge Society 10M views • 1 year ago 4:15

President of India 2.18M subscribers

#RepublicDay #RepublicDayParade2022 #RepublicDay2022

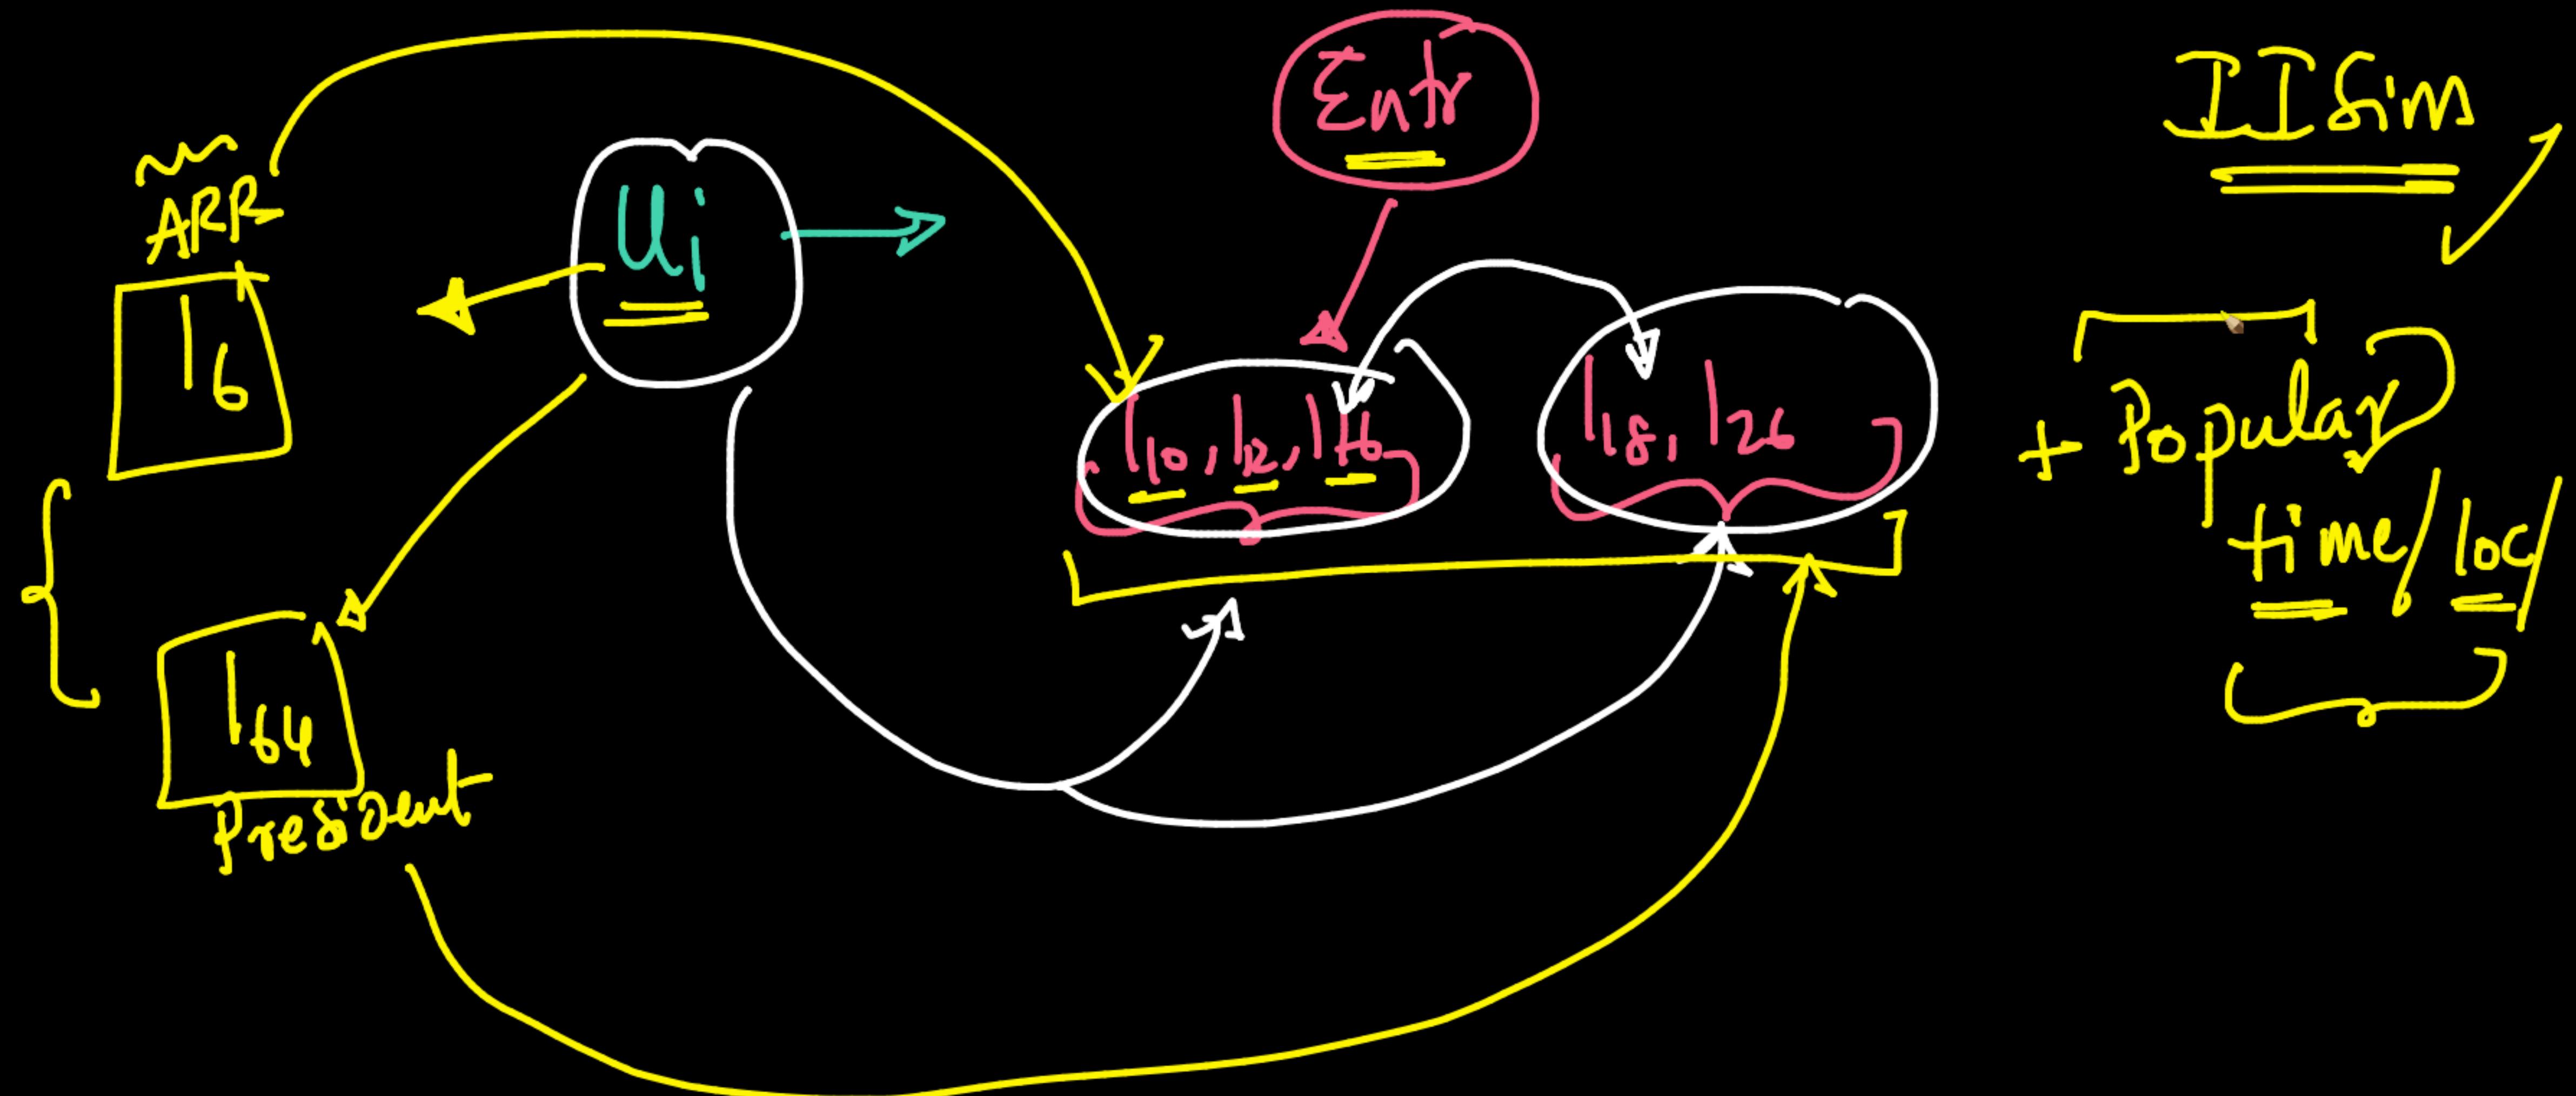
LIVE: Republic Day Parade - 2022

6,303,585 views • Streamed live on Jan 26, 2022

94K DISLIKE SHARE CLIP SAVE ...

83 / 83 SUBSCRIBE

0:00 / 2:15:55



PCA\_tSNE\_UMAP.ipynb · Amazon.com recommend · LIVE: Republic Day Parade · Amazon.in : tshirt · Buy CHKOKKO Men's Gy · Amazon.in : gift · Amazon.in : rakhi gift · Buy Collectible India Rakhi · +

amazon.in/CHKOKKO-Regular-Sleeves-T-Shirts-Seagreen/dp/B0856V6J1G/ref=sr\_1\_9?crid=2G0FEI6TC20N1&keywords=t-shirt&qid=1659723434&sprefix=tshir%2Caps%2C226&sr=8-9

amazon.in Hello, Sign in Account & Lists Returns & Orders Cart

All t-shirt

All Best Sellers Mobiles Today's Deals Customer Service Books Electronics Prime Fashion New Releases Home & Kitchen Amazon Pay Computers Coupons Toys & Games Sell PRIME EARLY ACCESS LIVE NOW

## Amazon Fashion

Women

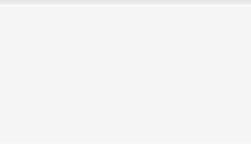
Men

Kids

Bags &amp; Luggage

Sportswear

Sales &amp; Deals

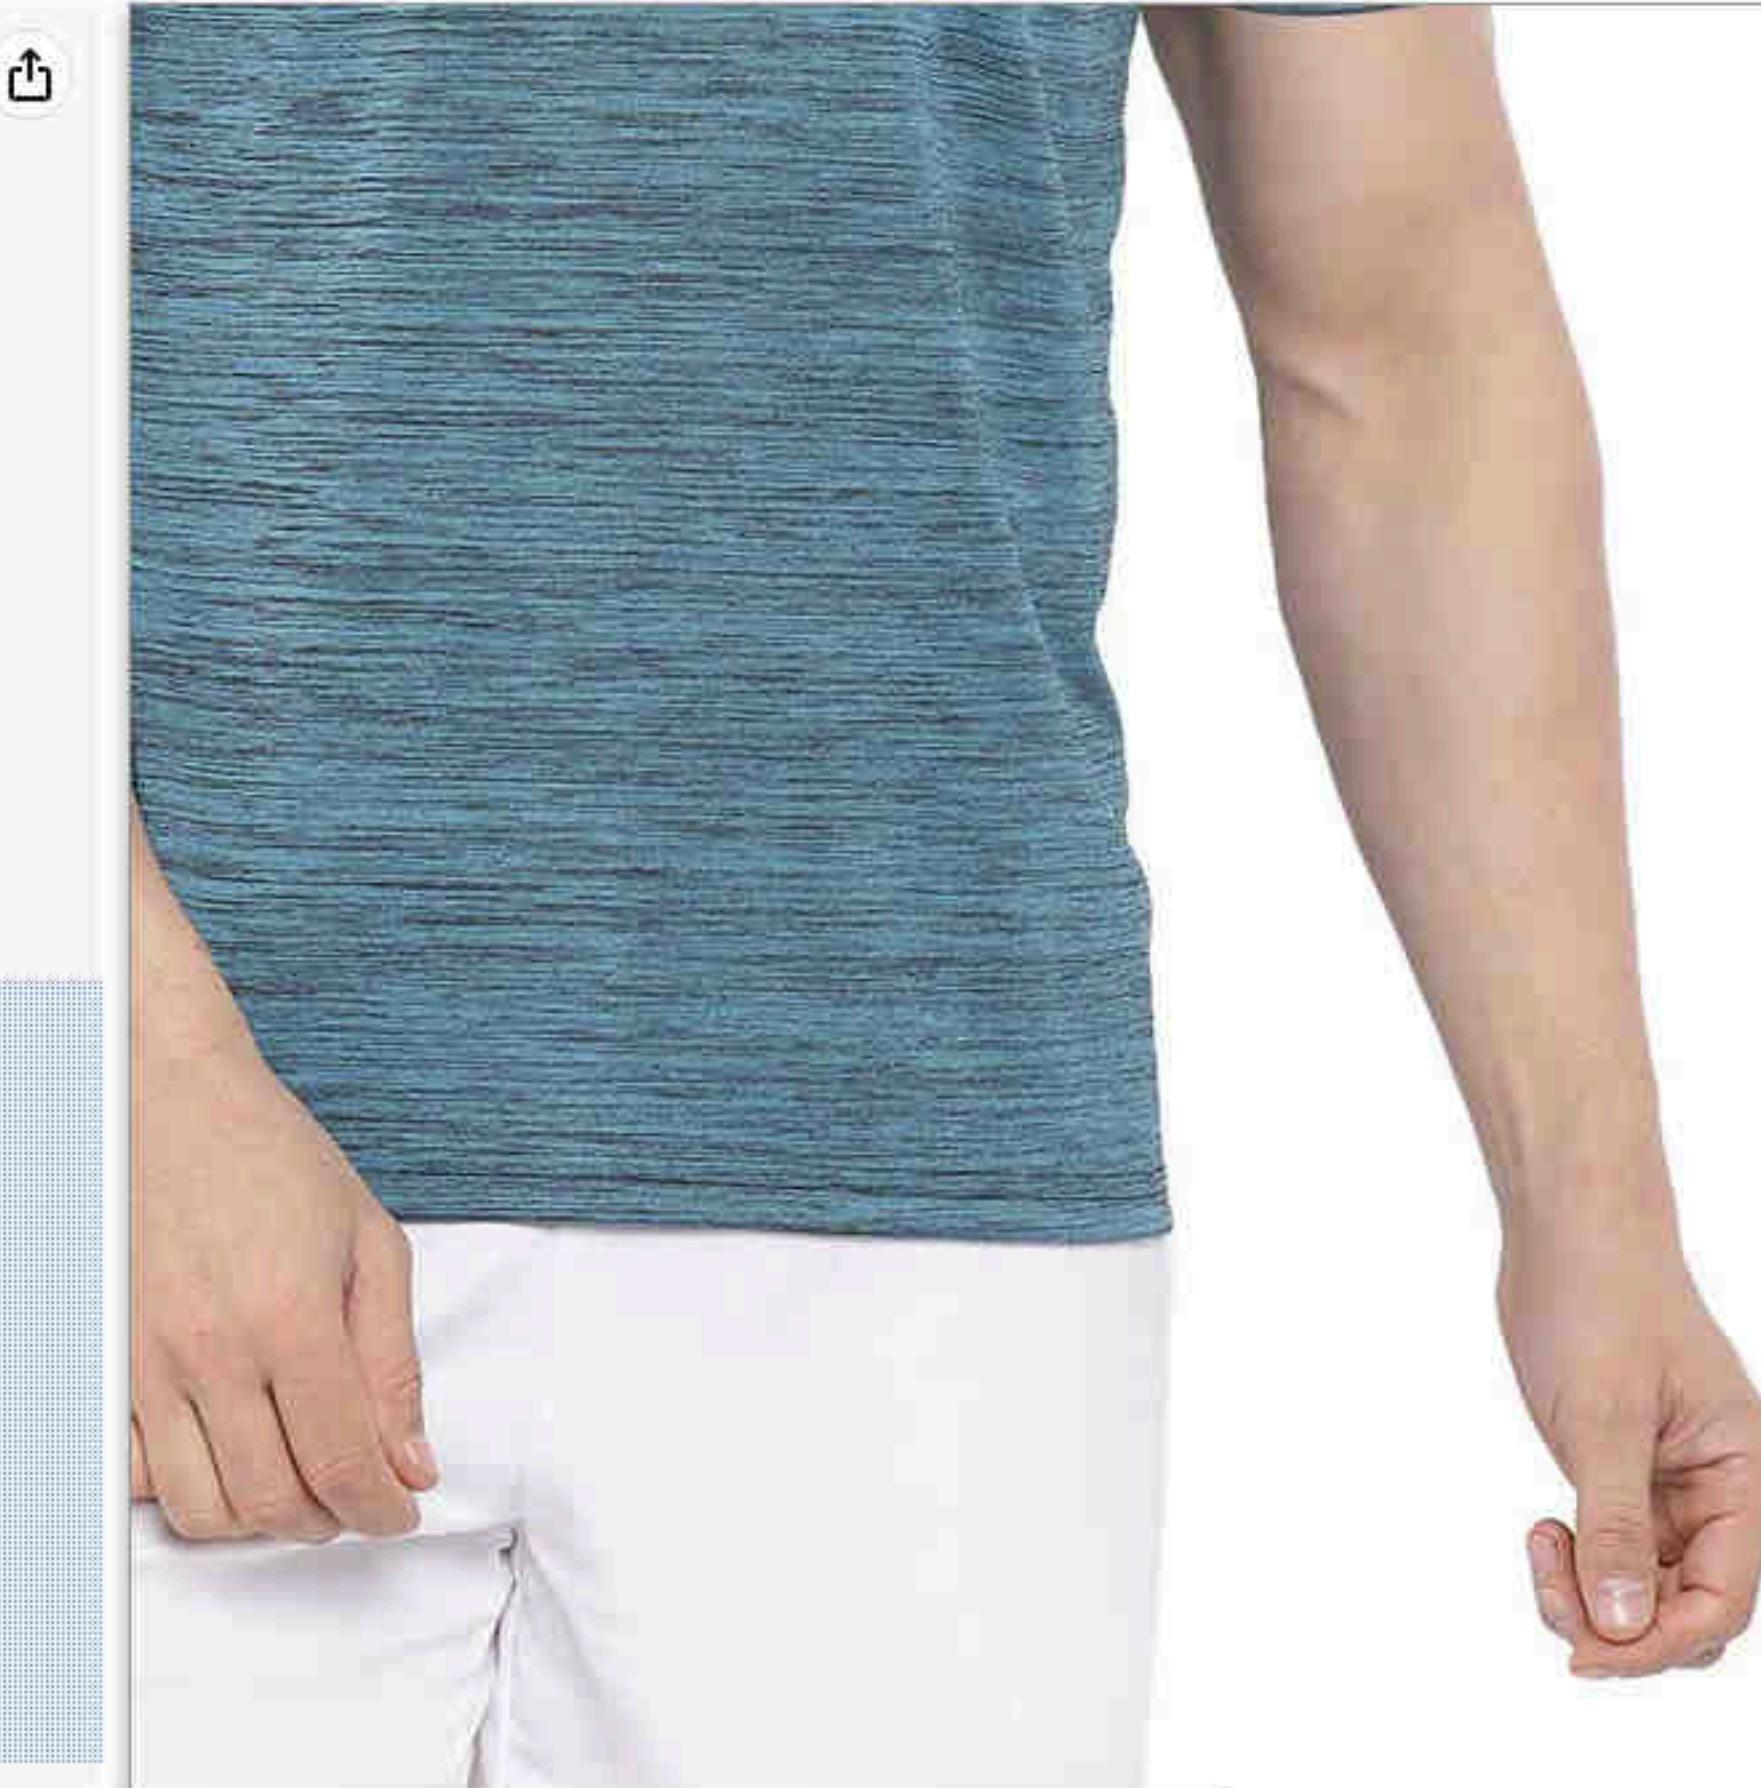
30 DAY RETURNS  
Restrictions Apply

Amazon Brand - Symbol Men's Regular Polo Shirt

★★★★★ 6,536  
₹299.00

Back to results

Sponsored



85 / 85

Click to open ext