

Suggestion/Homework:

To help improve your programming skills and your depth of understanding,

1. Implement Lloyd's algorithm from scratch
2. Implement KMeans++ from scratch

▼ Retail Customer Segmentation: Acquire Data

```
# https://drive.google.com/file/d/1lEccW5Y5\_2z00VRtLGOAJOU6YA9fl6W/view?usp=sharing
id = "1lEccW5Y5_2z00VRtLGOAJOU6YA9fl6W"
print("https://drive.google.com/uc?export=download&id=" + id)
```

https://drive.google.com/uc?export=download&id=1lEccW5Y5_2z00VRtLGOAJOU6YA9fl6W

```
!wget "https://drive.google.com/uc?export=download&id=1lEccW5Y5_2z00VRtLGOAJOU6YA9fl6W"
```

```

[+] --2022-06-01 13:27:48-- https://drive.google.com/uc?export=download&id=1lEccW5Y5\_2z00VRtLGOAJOU6YA9fl6W
Resolving drive.google.com (drive.google.com)... 172.253.62.100, 172.253.62.100
Connecting to drive.google.com (drive.google.com)|172.253.62.100|:443... connec
HTTP request sent, awaiting response... 303 See Other
Location: https://doc-10-64-docs.googleusercontent.com/docs/securesc/ha0ro937c
Warning: wildcards not supported in HTTP.
--2022-06-01 13:27:48-- https://doc-10-64-docs.googleusercontent.com/docs/securesc/ha0ro937c
Resolving doc-10-64-docs.googleusercontent.com (doc-10-64-docs.googleusercontent.com)... 172.253.62.100, 172.253.62.100
Connecting to doc-10-64-docs.googleusercontent.com (doc-10-64-docs.googleusercontent.com)|172.253.62.100|:443... connec
HTTP request sent, awaiting response... 200 OK
Length: 139827 (137K) [text/csv]
Saving to: 'E-commerce.csv'

```

```
E-commerce.csv      100%[=====>] 136.55K  --.-KB/s    in 0.001s
```

```
2022-06-01 13:27:48 (99.9 MB/s) - 'E-commerce.csv' saved [139827/139827]
```

```
!head ./E-commerce.csv
```

```

ID,n_clicks,n_visits,amount_spent,amount_discount,days_since_registration,prof
1476,130,65,213.90583071577163,31.600750627904915,233,235
1535,543,46,639.2230037736391,5.6891747173479414,228,170
1807,520,102,1157.4027626541078,844.3216058194998,247,409
1727,702,83,1195.903633609631,850.0417570033645,148,200
1324,221,84,180.75461615086704,64.2833000293408,243,259
1793,971,167,1700.9096451005262,1257.4171176204811,205,229
646,345,77,1314.029384121076,12.095727427593667,230,217
416,222,61,3869.409085656883,117.49933081401785,257,296
232,451,74,2598.1462935566806,103.64066379042382,65,102

```

```
import pandas as pd
import numpy as np
from matplotlib import pyplot as plt
import seaborn as sns
from sklearn.cluster import KMeans
```

```
df = pd.read_csv('./E-commerce.csv')
```

```
df.head()
```

	ID	n_clicks	n_visits	amount_spent	amount_discount	days_since_registra
0	1476	130	65	213.905831	31.600751	
1	1535	543	46	639.223004	5.689175	
2	1807	520	102	1157.402763	844.321606	
3	1727	702	83	1195.903634	850.041757	
4	1324	221	84	180.754616	64.283300	

```
## We do not need 5 decimal places, so we round it to 2 places instead.
```

```
df['amount_spent'].round(decimals=2)
```

```
df['amount_discount'].round(decimals=2)
```

```
0      31.60
1       5.69
2     844.32
3     850.04
4      64.28
...
2495    373.41
2496    122.64
2497      0.00
2498     78.13
2499   1065.42
Name: amount_discount, Length: 2500, dtype: float64
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2500 entries, 0 to 2499
Data columns (total 7 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   ID                                    2500 non-null   int64
1   n_clicks                             2500 non-null   int64
2   n_visits                             2500 non-null   int64
3   amount_spent                         2500 non-null   float64
4   amount_discount                      2500 non-null   float64
5   days_since_registration              2500 non-null   int64
6   profile_information                  2500 non-null   int64
dtypes: float64(2), int64(5)
memory usage: 136.8 KB
```

▼ Cleaning + Preprocessing

```
X=df.drop("ID",axis=1)
X
```

	n_clicks	n_visits	amount_spent	amount_discount	days_since_registration
0	130	65	213.905831	31.600751	200.9
1	543	46	639.223004	5.689175	200.9
2	520	102	1157.402763	844.321606	200.9
3	702	83	1195.903634	850.041757	130.0
4	221	84	180.754616	64.283300	200.9
...
2495	804	120	502.643798	373.413462	300.0
2496	482	60	530.014805	122.639755	100.0
2497	375	111	0.000000	0.000000	100.0
2498	271	32	3190.499018	78.133067	100.0
2499	814	123	1394.589041	1065.415902	200.9

2500 rows x 6 columns

```
X.describe()
```

	n_clicks	n_visits	amount_spent	amount_discount	days_since_registration
count	2500.00000	2500.000000	2500.000000	2500.000000	2500.000000
mean	408.68000	94.475600	1445.090745	388.508637	200.900000
std	186.41409	38.866356	1167.663473	487.143968	99.100000
min	50.00000	10.000000	0.000000	0.000000	0.000000
25%	274.75000	67.000000	609.618538	56.298615	130.000000
50%	378.00000	92.000000	1036.189112	137.454623	200.900000
75%	522.00000	119.000000	1949.270949	679.540536	268.000000
max	1246.00000	259.000000	6567.402267	2428.406527	514.000000

```
# Do we need scaling at all?
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
```

```
scaler.fit(X)
```

```
X
```

```
array([[ -1.49525046, -0.75853514, -1.05461141, -0.73280039,  0.32311781,
         0.33919174],
       [  0.72069055, -1.24748762, -0.69029218, -0.78600183,  0.27267227,
        -0.31003449],
       [  0.59728463,  0.19363547, -0.24642848,  0.93587151,  0.46436533,
         2.07712043],
       ...,
       [-0.18070918,  0.42524454, -1.2378394 , -0.7976828 , -1.18015931,
         2.52658475],
       [-0.73871854, -1.60776839,  1.49508613, -0.63726061, -0.90775339,
         0.13942982],
       [ 2.17473416,  0.73405663, -0.04325887,  1.38982053,  0.47445444,
        -0.51978451]])
```

```
from sklearn.cluster import KMeans
```

```
k = 4 ## arbitrary value
```

```
kmeans = KMeans(n_clusters=k)
```

```
y_pred = kmeans.fit_predict(X)
```

```
## what are learned labels(cluster #)
```

```
y_pred
```

```
array([3, 3, 0, ..., 1, 2, 0], dtype=int32)
```

```
##coordinates of the cluster centers
```

```
kmeans.cluster_centers_
```

```
array([[ 1.0387607 ,  0.87726845, -0.0308178 ,  1.51719781, -0.01890781,
        -0.02259216],
       [-0.10975713, -0.22690132, -0.62750954, -0.4811005 , -0.79076605,
        -0.0765108 ],
       [-0.84848342, -0.69485355,  1.62272465, -0.62820588, -0.03079898,
         0.13822599],
       [-0.18505778, -0.02992418, -0.6241805 , -0.42574709,  0.89260056,
        -0.00923737]])
```

```
y_pred is kmeans.labels_
```

```
True
```

```
#Visualize clusters
```

```
clusters = pd.DataFrame(X, columns=df.drop("ID",axis=1).columns)
```


```
clusters['label'] = kmeans.labels_
```

```
clusters
```

	n_clicks	n_visits	amount_spent	amount_discount	days_since_registration
0	-1.495250	-0.758535	-1.054611	-0.732800	0.3231
1	0.720691	-1.247488	-0.690292	-0.786002	0.2726
2	0.597285	0.193635	-0.246428	0.935872	0.4643
3	1.573801	-0.295317	-0.213449	0.947616	-0.5344
4	-1.006992	-0.269583	-1.083008	-0.665697	0.4240
...
2495	2.121079	0.656854	-0.807284	-0.030993	1.6347
2496	0.393397	-0.887207	-0.783838	-0.545880	-0.8976
2497	-0.180709	0.425245	-1.237839	-0.797683	-1.1801
2498	-0.738719	-1.607768	1.495086	-0.637261	-0.9077


```
def viz_clusters(kmeans):
    plt.scatter(clusters['n_clicks'], clusters['n_visits'], c=clusters['label'])
    plt.scatter(kmeans.cluster_centers_[0], kmeans.cluster_centers_[1],
                color="red",
                marker="x",
                s=100)
```

```
viz_clusters(kmeans)
```




```
#polar plot
```

```
polar = clusters.groupby("label").mean().reset_index()
polar = pd.melt(polar, id_vars=["label"])
polar
```

	label	variable	value	
0	0	n_clicks	1.038761	
1	1	n_clicks	-0.109757	
2	2	n_clicks	-0.848483	
3	3	n_clicks	-0.185058	
4	0	n_visits	0.877268	
5	1	n_visits	-0.226901	
6	2	n_visits	-0.694854	
7	3	n_visits	-0.029924	
8	0	amount_spent	-0.030818	
9	1	amount_spent	-0.627510	
10	2	amount_spent	1.622725	
11	3	amount_spent	-0.624181	
12	0	amount_discount	1.517198	
13	1	amount_discount	-0.481101	
14	2	amount_discount	-0.628206	
15	3	amount_discount	-0.425747	
16	0	days_since_registration	-0.018908	
17	1	days_since_registration	-0.790766	
18	2	days_since_registration	-0.030799	
19	3	days_since_registration	0.892601	
20	0	profile_information	-0.022592	

```
import plotly.express as px
```

```
fig = px.line_polar(polar, r="value", theta="variable", color="label", line_close=True)
fig.show()
```



Insights

1. The plot is read and interpreted radially - values increase as we move away from the center showing the influence of a feature on that label.
2. Many overlapping lines - green(2) & red(1) and blue and overlap on all the features except one. Looking at this plot, we have different customer segments:
3. Bargain shoppers(label 0) - people who buy heavily discounted items. Action: show them more discounted items.
4. Inactive old users (label 1) - people who have been a long time user of the app but have not shown much interest now. Action: start sending notifications, emails, etc - get them back on the platform.
5. New and inactive users - Users who have recently joined but haven't bought much and are not that actively looking for items. Somewhat similar to label 1.
6. Premium shopper (label 3) - Heavy spenders who like to buy items. Action: Show them more quality

▼ Finding the best K

```
# Inertia = Within Cluster Sum of Squares
kmeans_per_k = [KMeans(n_clusters=k, random_state=42).fit(X)
```

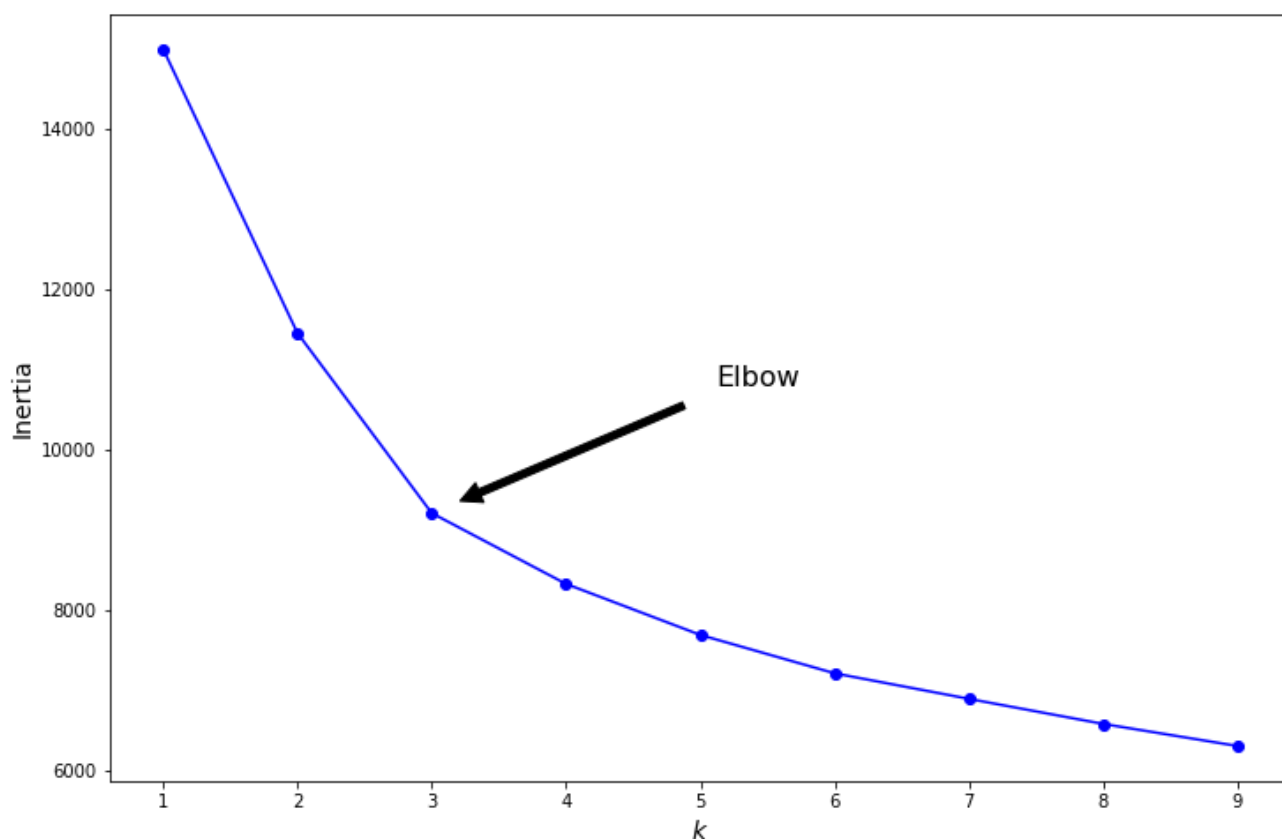
```

for k in range(1, 10)]

inertias = [model.inertia_ for model in kmeans_per_k]

plt.figure(figsize=(12, 8))
plt.plot(range(1, 10), inertias, "bo-")
plt.xlabel("$k$", fontsize=14)
plt.ylabel("Inertia", fontsize=14)
plt.annotate('Elbow',
             xy=(3, inertias[2]),
             xytext=(0.55, 0.55),
             textcoords='figure fraction',
             fontsize=16,
             arrowprops=dict(facecolor='black', shrink=0.1)
            )
plt.show()

```



▼ Silhouette score


1. Silhouette Coefficient of $X_i = (b - a) / \max(a, b)$
2. Diagram: https://www.researchgate.net/figure/Derivation-of-the-Overall-Silhouette-Coefficient-OverallSil_fig1_221570710
3. Range: Worst (-1) to Best(+1)

4. Average Silhouette Coefficients of all points.

```
from sklearn.metrics import silhouette_score

silhouette_scores = [silhouette_score(X, model.labels_)
                     for model in kmeans_per_k[1:]]

plt.figure(figsize=(10, 6))
plt.plot(range(2, 10), silhouette_scores, "bo-")
plt.xlabel("$k$", fontsize=14)
plt.ylabel("Silhouette score", fontsize=14)
plt.show()
```



✓ 0s completed at 19:28

