Introduction to Gen AI

Generative AI refers to artificial intelligence algorithms that can generate new content or data that is similar to the data they were trained on.The generative part of the name indicates that these AI systems are capable of creating something new, rather than just processing or analyzing existing data. traditional AI algorithms may be used to identify patterns within a training data set and make predictions, generative AI uses machine learning algorithms to create outputs based on a training data set.

Mostly Gen AI models are trained to generate new original content based on natural language input. In other words, you can describe a desired output in normal everyday language, and the model can respond by creating appropriate text, image, or even code output.

There are several reasons for the Gen AI frenzy:

- **Unprecedented Capabilities:** Generative AI models have shown remarkable abilities in tasks that were once thought to be uniquely human, such as creative writing, art creation, and problem-solving.
- **Rapid Advancement:** The field is progressing at an astonishing pace, with new breakthroughs and improved models emerging frequently.
- **Wide-ranging Applications:** Generative AI has potential applications across numerous industries, from healthcare and finance to entertainment and education.
- **Economic Impact:** According to McKinsey, generative AI could add trillions of dollars to the global economy in the coming years.
- **Accessibility:** Many generative AI tools are becoming increasingly user-friendly, allowing non-experts to leverage their power.

Overview of this module

After this module you would be able to master most of the gen AI capabilities

- You'll have an understanding of how computers understand text and images and work with them
- You would be able to create simple systems to understand text and then
    - summarise them (like in shorts), Sentiment analysis and customer service chatbots (https://www.chatbot.com/solutions/chatbot-for-support/)

- Work not with just text but also images, Know what all tasks can be done using what kind of sevices and models
    - Create object detection, image captioning (https://ai.cloudinary.com/demos/captions) questions and answers based on images (https://vision-explorer.allenai.org/visual_question) and much more.
    - Understand the basic of how computers generate images based on text prompts https://magicstudio.com/ai-art-generator/
- Would be able to utilise chat gpt and other llm apis to build powerful systems on top of that
    - https://analyticsindiamag.com/7-chatgpt-api-use-cases-that-caught-our-eyes/ (this article has great examples)
    - Youn will also understand how to classify, moderate input and outputs and also train apis for your specific needs.
- Understand how LLMs can work on your own data using RAGs
    - Creating chatbots and LLMS on your own data
    - How to use LLMs with data in csvs and other formats.
- Understand the peak of generative ai developments with Agents to build the most advanced and capable systems combining all the tools and knowledge
    - Making llms use your own tools, search information from internet, run code and most advance capabiliteis

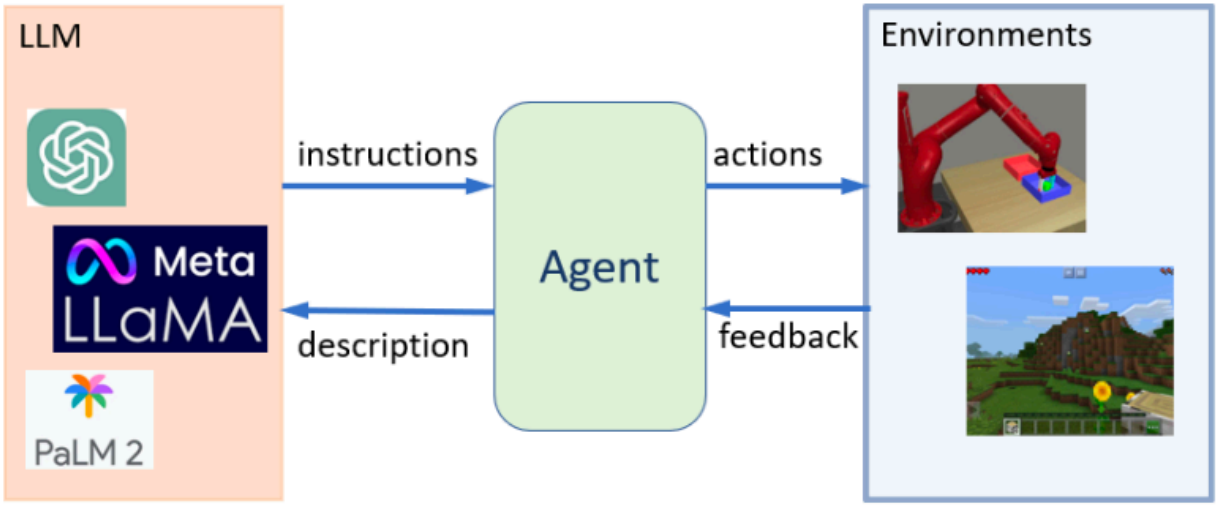

Here's a breakdown of the core concept of GEN AI:

- **Data Feeding**: The generative model is fed a vast amount of data relevant to the kind of content it should create – text for writing, images for creating visuals, and so on.
- **Content Generation**: Once it grasps the data, the model can begin generating new content. It does this by taking a prompt or starting point from the user and using its knowledge to extrapolate and create something original, adhering to the learned patterns.

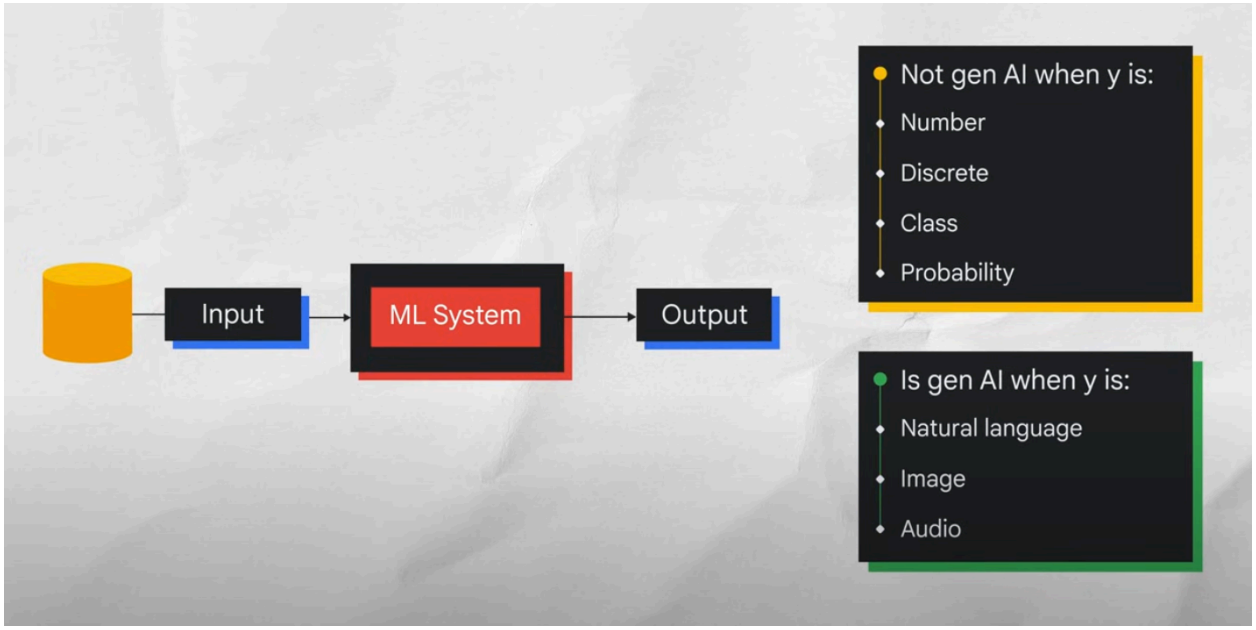

From ML to Gen AI

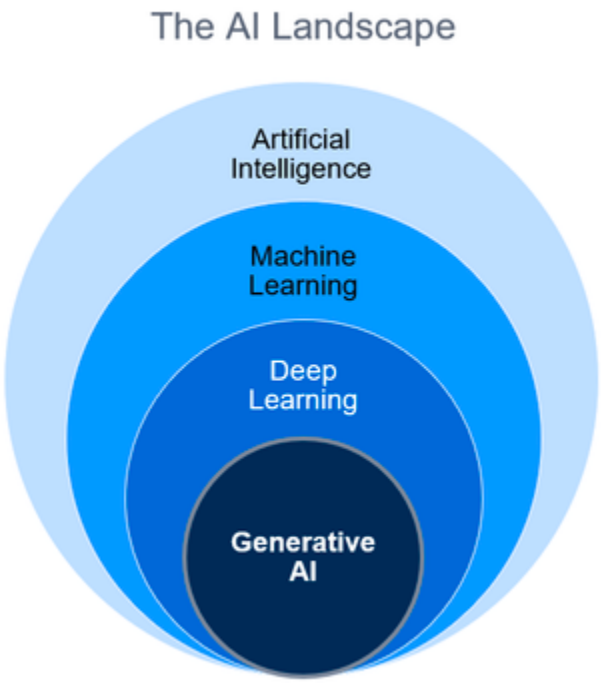

We have already learned about classical ML models, but generative AI is based on another subset of Machine learning which is Deep learning and neural networks.

We don't need to explore the depths of these topics to be able to work with Gen AI, but it is generally good to have a basic understanding of what things are and close the knowledge gap between the two domains. There are many ways to learn a model from data. The Neural Network is one such way.

Traditional machine learning models, such as linear regression and decision trees, **rely on structured data and are designed to perform specific tasks by identifying patterns and relationships within the data.**

Deep learning is a specialized **subset of machine learning** that uses neural networks with many layers (hence "deep") to model complex patterns in large amounts of data. while traditional Machine Learning algorithms have a rather simple structure, such as linear regression or a decision tree. Deep Learning algorithms can be regarded both as a sophisticated and mathematically complex evolution of machine learning algorithms.

> **Neural networks are the foundation of deep learning. They are inspired by the structure and function of the human brain, consisting of interconnected**

nodes (neurons) organized in layers. Neural networks can be simple, with just a few layers (shallow neural networks), or very complex, with many layers (deep neural networks).

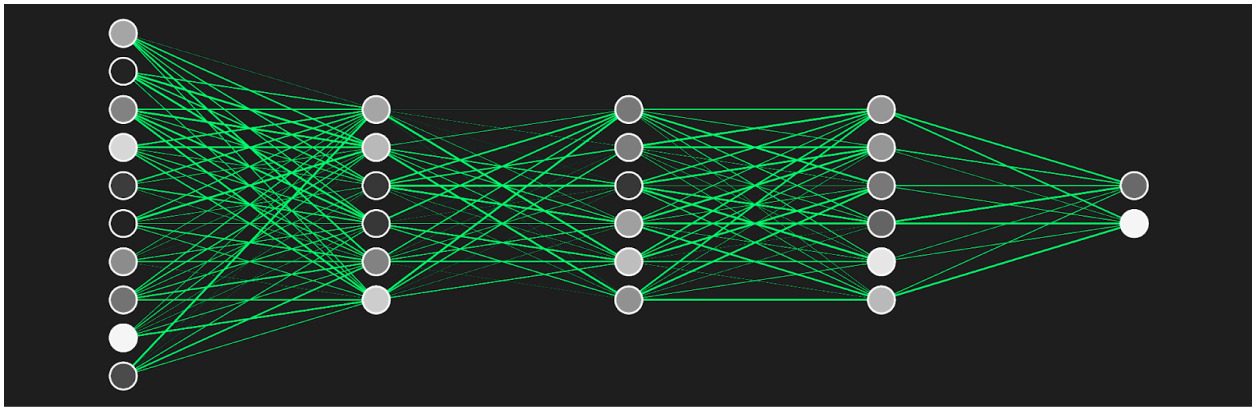

Neural networks are employed across various domains for:

- Identifying objects, faces, and understanding spoken language in applications like self-driving cars and voice assistants.
- Analyzing and understanding human language, enabling sentiment analysis, chatbots, language translation, and text generation.
- Diagnosing diseases from medical images, predicting patient outcomes, and drug discovery.
- Predicting stock prices, credit risk assessment, fraud detection, and algorithmic trading.
- Personalizing content and recommendations in e-commerce, streaming platforms, and social media.
- Powering robotics and autonomous vehicles by processing sensor data and making real-time decisions.
- Monitoring and optimizing manufacturing processes, predictive maintenance, and quality control.
- Analyzing complex datasets, simulating scientific phenomena, and aiding in research across disciplines.

Computer Vision and Natural Language Processing as Subsets of Deep Learning

Within the domain of deep learning, there are specific tasks tailored to particular types of data. Two prominent subsets are Computer Vision (CV) and Natural Language Processing (NLP).

Computer Vision (CV)

Computer Vision is a field of deep learning focused on enabling machines to interpret and make decisions based on visual data. This includes tasks such as:

- **Image Classification**: Categorizing images into predefined classes.
- **Object Detection**: Identifying and locating objects within an image.
- **Image Segmentation**: Dividing an image into segments to identify different objects or regions.
- **Image Generation**: Creating new images from scratch, often using Generative Adversarial Networks (GANs).

Deep learning models, particularly Convolutional Neural Networks (CNNs), are highly effective for CV tasks due to their ability to capture spatial hierarchies in images. CNNs use layers with convolutional filters that automatically learn to detect edges, textures, shapes, and more complex structures as the layers deepen.

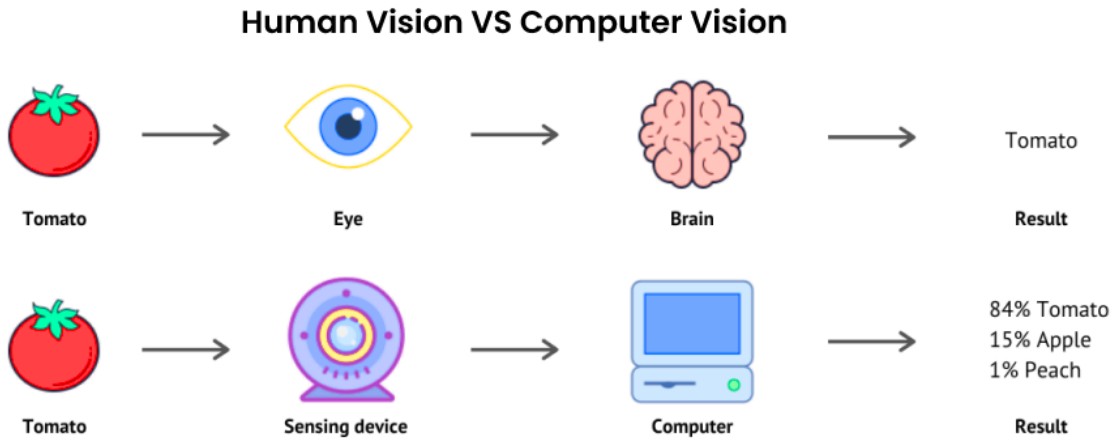

Natural Language Processing (NLP)

Natural Language Processing is a field of deep learning that focuses on the interaction between computers and human language.Natural language processing employs artificial intelligence to accept real-world data, interpret it, and make sense of it in a way that a computer can understand. Key NLP tasks include:

- **Text Classification**: Categorizing text into predefined categories, such as spam detection.
- **Sentiment Analysis**: Determining the sentiment expressed in text, whether positive, negative, or neutral.
- **Language Translation**: Translating text from one language to another.
- **Text Generation**: Creating new, coherent text based on given prompts.
- **Speech Recognition**: Converting spoken language into text.

Deep learning models, especially those based on the Transformer architecture, have revolutionized NLP. Models like GPT (Generative Pre-trained Transformer) and BERT (Bidirectional Encoder Representations from Transformers) use self-attention mechanisms to understand the context and relationships within text, enabling more accurate and sophisticated language understanding and generation.

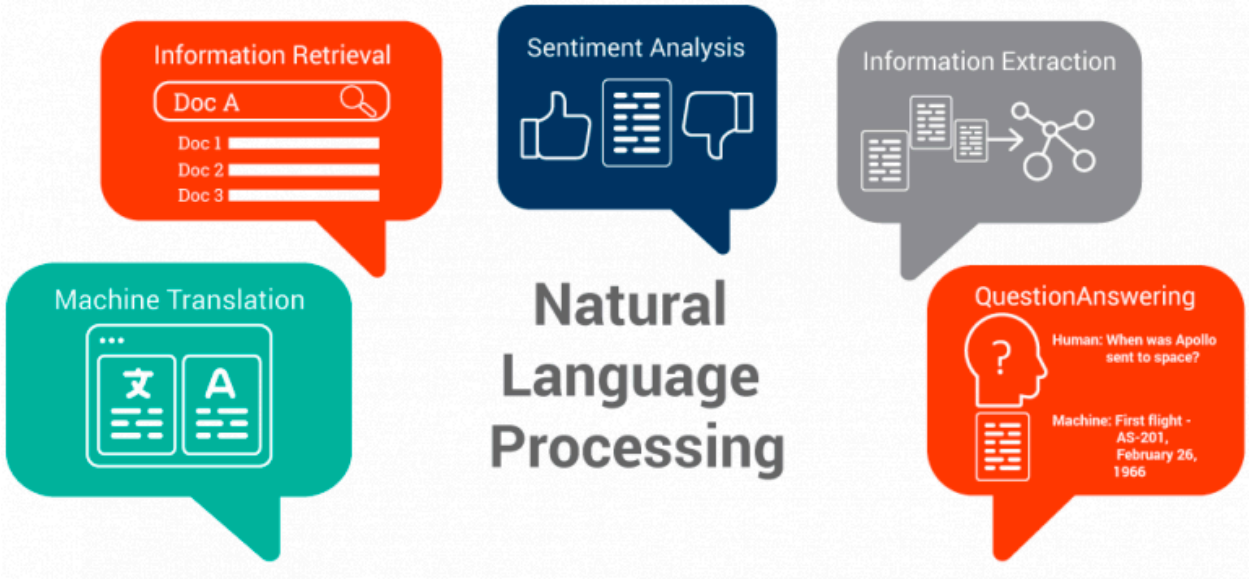

Key Concepts from Deep Learning to Generative AI

At the core, both traditional deep learning models are based on the fundamental principle of learning patterns from data. These models analyze vast amounts of data to discern intricate details, relationships, and patterns that might not be immediately apparent to humans. Here's how they do it:

- **Training Data**: The process starts with training data, a large set of examples relevant to the task at hand. For example, if we're building a model to recognize images of cats, we would provide it with thousands or millions of cat images.
- **Pattern Recognition**: During training, the model identifies various features and patterns in the data. In the case of images, this might involve recognizing edges, textures, shapes, and colors. For text, this might include understanding syntax, semantics, and context.

the model's goal is to accurately predict or classify data based on the patterns it has learned during training. However, the output is limited to these predictions or classifications.

Generative AI: Reproducing Patterns

Generative AI builds upon the foundations of deep learning. While deep learning models typically focus on recognizing patterns and making predictions, Generative AI, on the other hand, takes this a step further. Instead of merely identifying patterns, generative models aim to reproduce or generate new data that reflects the same patterns and details found in the training data. This means creating new instances that are similar to the examples the model was trained on. Without diving into the technical specifics of their architecture, it's useful to understand the general approach generative models take:

1. **Learning from Data**: Just like traditional models, generative models start by learning from a large dataset, capturing the underlying patterns and structures.
2. **Generating New Data**: Using the learned distribution, these models can then generate new data points. For instance, in text generation, the model can predict the next word or sentence based on the preceding text.
3. **Evaluation and Improvement**: Generative models often involve an iterative process where they continuously improve their ability to generate realistic data. Techniques like Generative Adversarial Networks (GANs) use a competitive approach where one part of the model generates data, and another part evaluates it, refining the generation process over time.
- **Creativity and Innovation**: Generative models can create novel content that wasn't explicitly present in the training data but follows the same rules and patterns.

They learn the underlying distribution of a dataset and can generate data points that appear to come from the same distribution. This makes them useful for tasks where creativity or content generation is required.

So the models are still learning the patterns and details of the training data, the difference is that we are now not just identifying them we are also trying to reproduce them.

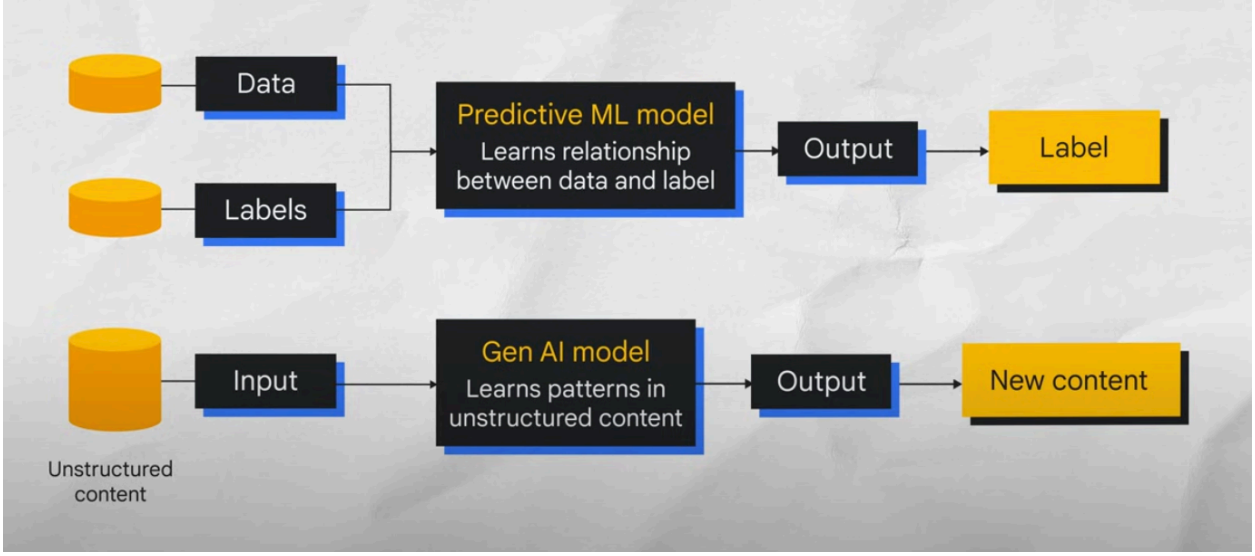

Practical Examples

To illustrate the difference more concretely, let's look at some practical examples:

1. **Image Generation**
    - **Traditional Deep Learning**: A model might be trained to recognize and classify different objects in images (e.g., identifying cats, dogs, cars).
    - **Generative AI**: A model might be trained to generate entirely new images of cats, creating realistic pictures of cats that do not exist in real life.
    - https://stablediffusion.fr/dalle
2. **Text Generation**
    - **Traditional Deep Learning**: A model might be trained to classify the sentiment of a piece of text (e.g., positive, negative, neutral).
    - **Generative AI**: A model might be trained to generate new text, such as writing an essay or creating dialogue for a chatbot, based on the style and content of the training text.
3. **Music Composition**
    - **Traditional Deep Learning**: A model might be trained to classify music genres based on audio features.
    - **Generative AI**: A model might be trained to compose new pieces of music that emulate the style of classical compositions, generating new melodies and harmonies. Eg; https://huggingface.co/spaces/facebook/MusicGen (this is really fun thing to try recording a melody/whistling and adding music to it. The output although is not the best)

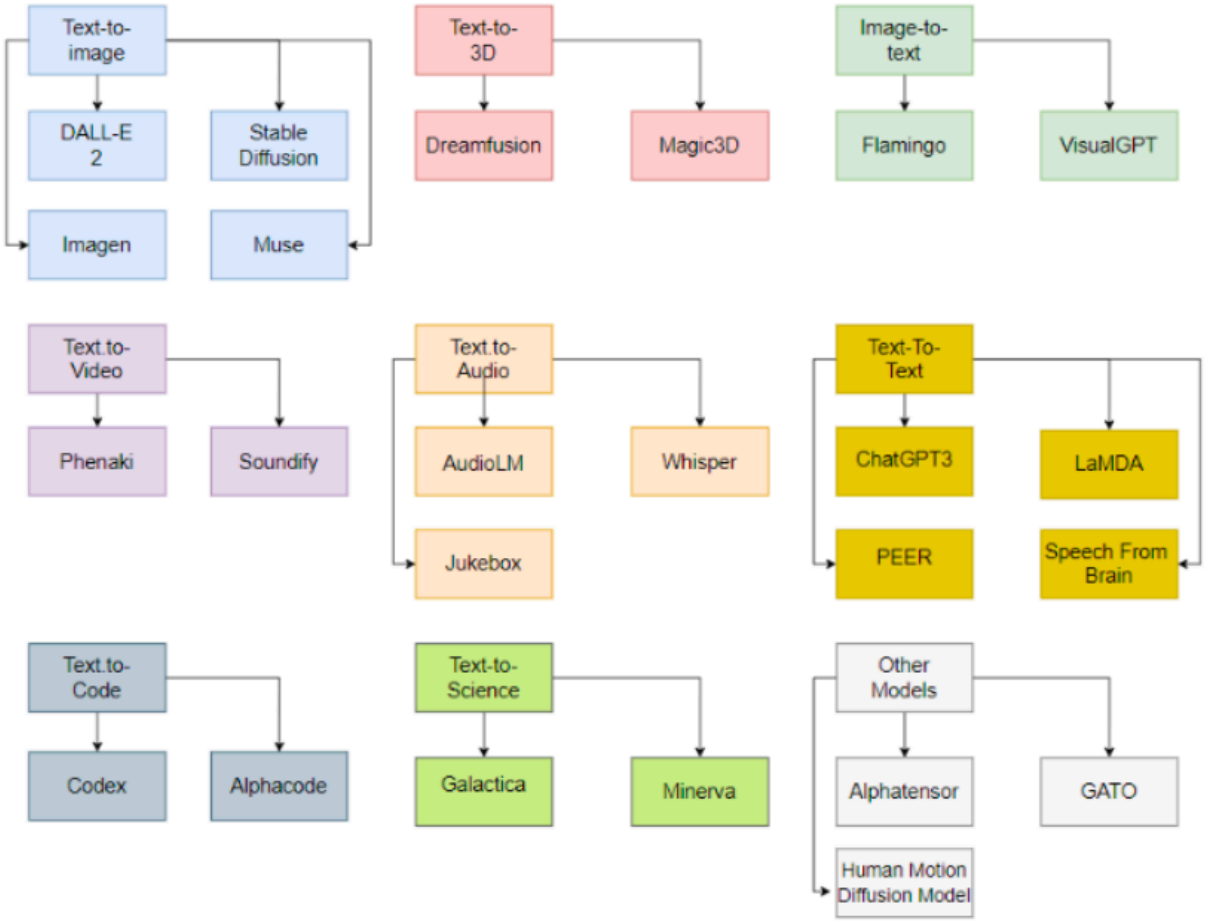

As the field of generative AI continues to evolve, professionals who develop expertise in this area can expect significant career advancement opportunities:

1. **High Demand:** The rapid growth of generative AI is creating a talent shortage, leading to competitive salaries and benefits for skilled professionals.
2. **Cross-industry Relevance:** Expertise in generative AI can be applied across various sectors, providing diverse career paths and opportunities for specialization.
3. **Innovation Potential:** Those working in generative AI have the chance to contribute to groundbreaking technologies that could reshape industries and society.
4. **Leadership Roles:** As companies increasingly adopt generative AI, there will be a need for leaders who understand both the technical and strategic aspects of these technologies.

5. **Entrepreneurship:** The generative AI space is ripe for innovation, offering opportunities for entrepreneurs to create new products, services, and startups.

—-------

Intel