Line of Code

Steps before model training

```
from sklearn.preprocessing import MinMaxScaler

scaler = MinMaxScaler()

df = pd.DataFrame(scaler.fit_transform(df), columns=df.columns)
```

All Numerical Columns.

① cat → Num
② Feature Scaling

Data Notation

samples - m

dimension / # features - d

Predictors - (m, d) matrix X

ith sample - $X^{(i)}$ — one row

jth feature - X_j - one column

True output / target - y^i

Predicted output / target $\hat{y}^i$

Goal of generalization in ML

m samples - $\{x^i, y^i\}_m$

Historical Data → Train an ML Model

$$x^i \xrightarrow[\text{Model}]{\text{ML}} \hat{y}^i \leftarrow \text{Predicted Value}$$

Ideally, $y^i \approx \hat{y}^i$ → Predicted Output

↳ True output

⇒ For all m samples $y^i \approx \hat{y}^i$ → GOOD TRAINING

DO WE WANT TO PERFORM PREDICTIONS ON TRAINGN

DATA??

For training data y is already present?

Why to even do prediction.

Goal of ML: GENERALISATION.

Your model should perform well both on training and NEW data.

True output

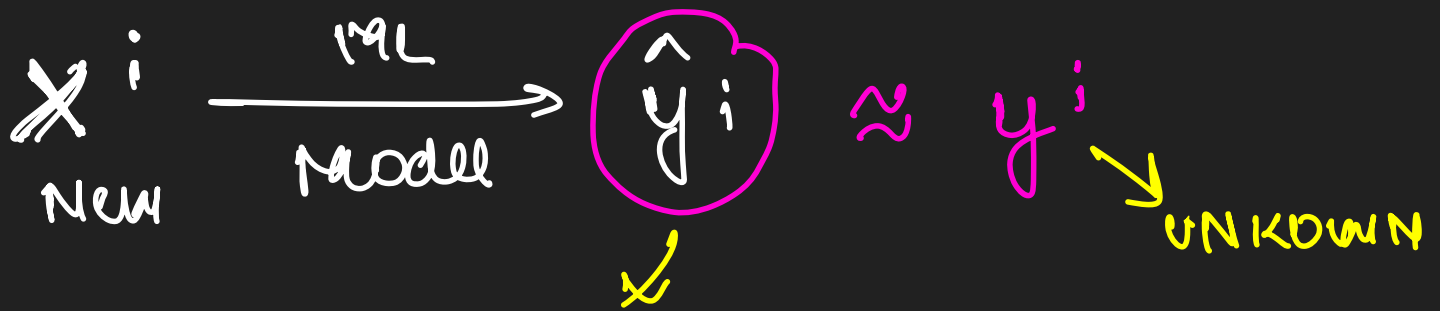

Check if model is GENERALISING??

Problem: No labels for new data.

Solution:

Because we don't know the ground truth (y) for New Sample

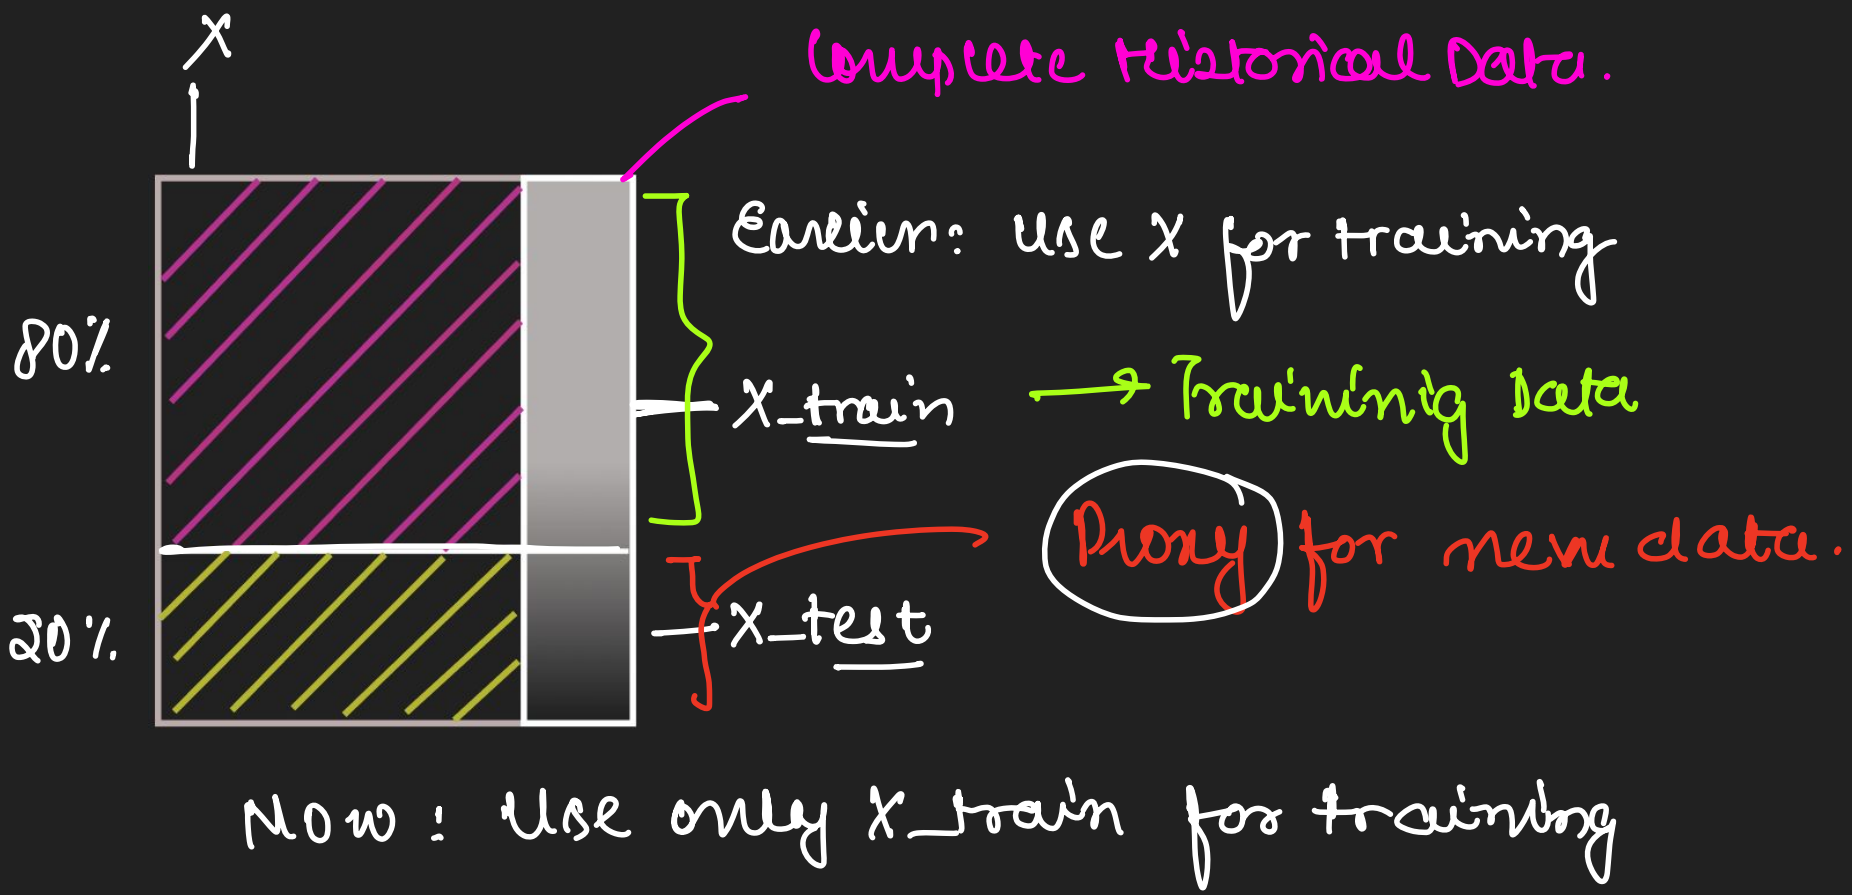

Phases for Model Development

① Training - use training split for model training

$y_{train} \approx \hat{y}_{train}$ - Model is learning

② Testing / Evaluation - test split

$y_{test} \approx \hat{y}_{test}$ - Model is generalizing

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    test_size=0.3,
                                                    random_state=100)
```

0.2

Random split

80% 20%

seed.

200